

How China Will Win the AI War

The Convergence Strategy: A Structural Path to AI Dominance

CONDENSED EDITION

A short-form companion to the living book¹

Satoshi Matsuoka

v0.20.1 — 18 August 2026

¹This condensed edition abridges the full living book of the same version and date. Every number in it is taken from the full edition, which carries the sources, evidence grades, and derivations. Version v0.20.1 — 18 August 2026.

About This Edition

The full edition of this book runs to nearly three hundred pages because its argument is a systems argument, and a systems argument has to be shown layer by layer before it can be believed: the model, the compute, the semiconductor, the watt, the data, physical production, trust, governance, finance, and the international order that all of them together imply. Many readers do not need the layers. They need the shape—what is being claimed, why, what would prove it wrong, and what it means for a government, a laboratory, or a company that is neither in Washington nor in Beijing. This condensed edition is for them.

Three rules govern the abridgement. First, nothing here is new: every number, grade, forecast, and recommendation is taken from the full edition of the same version and date, which carries the sources, the evidence grades, the derivations, and the falsification apparatus. Where the condensed text states a figure, the full edition states where it came from and how much to trust it. Second, the industrial precedents that occupy Part I of the full edition—the solar, battery, and electric-vehicle wars, and the campaigns where the same playbook has failed—are here summarized in a single chapter, because for the purposes of the argument what matters is the playbook they establish and its limits, not the campaign histories. Third, the weight of this edition is shifted deliberately toward the overall strategy and its international consequences: what the convergence means for the West, for Japan and Europe, for the non-aligned majority of the world, and for the institution the analysis keeps arriving at. Readers who want the engineering—the memory-architecture arithmetic, the appliance cost model, the sizing framework, the security fault tree—will find pointers into the full edition throughout, marked [full edition, Ch. n].

The title is deliberately provocative, but it is not a prediction of inevitability; it is a structural analysis. “Winning” means what it meant in solar, batteries, and electric vehicles: capturing the overwhelming majority of global production, setting the cost curve, controlling the supply chain, defining the de facto standards, and reducing erstwhile leaders to protected niches. By that definition China has already won three technology wars in which the West began with every advantage. The claim of this book is that artificial intelligence is following the same trajectory, and that the mechanism is visible only when the elements are laid side by side.

A note on the author’s position is owed and is given in full in the preface to the full edition. In brief: I direct the RIKEN Center for Computational Science, whose flagship program runs on an American accelerator architecture and whose deepest partnerships are with American and European institutions. Nothing in my professional position gives me an incentive to argue that the Chinese stack is winning; a great deal of it gives me an incentive to argue the opposite. I did not begin this project expecting to arrive where it has arrived, and I have tried to hold the argument to a standard of proof appropriate to how much I would prefer it to be wrong. Where the evidence does not support the strong claim, the book—and this abridgement of it—says so.

The currency date of this edition is 18 August 2026, the date on the title page. Every factual claim is current as of that date; anything learned after it belongs to the next edition. This is a living document. Corrections are welcome, and the fastest way to change its conclusions is to falsify an entry in the full edition’s evidence ledger.

Contents

About This Edition	i
Terms Used in This Edition	iii
1 The Thesis	1
1.1 Three wars, one playbook	1
1.2 Why AI is the fourth war	1
1.3 Openness as a weapon	2
1.4 Four propositions, stated separately	2
1.5 Evidence grades	3
1.6 The theory, in brief	3
2 The Precedents and the Playbook	5
2.1 Three campaigns, in brief	5
2.2 The playbook distilled	6
2.3 Where the analogy breaks	7
3 The Elements of AI	8
3.1 The model: open weights as the commodity layer	8
3.2 The harness: the layer the contest moved to	9
3.3 The compute: RISC-V, accelerators, and the memory arbitrage	10
3.4 The semiconductor: fabs, logic, and DRAM	12
3.5 The watt: energy as the fourth element	13
3.6 Data and feedback: the fifth element	14
3.7 Embodied AI: the loop closes on the factory floor	16
4 The Strategy Assembled	16
4.1 Where each element sits	17
4.2 What the assembly licenses	18
5 Trust, Security, and the Rulebook	19
5.1 Trust: the contested layer	19
5.2 Monoculture and agency: the security cost of the commons	20
5.3 The governance contest: who writes the rules	21
6 The Wrong Side of the Trade	22
6.1 The exposure, marked to market	22
6.2 The squeeze, and the four bubbles	23
6.3 The Jevons rebuttal, taken seriously	23
6.4 Vintages, buildings, and railroads	24
6.5 Scenarios, probabilities, indicators	24
6.6 The commodity turn, and where value reappears	25
7 How the Rest of the World Chooses	26
7.1 The decision matrix	26
7.2 Nine positions	27
7.3 Five moves for any non-aligned country	27

8 Strategies for the Rest	28
9 The Prize, the Bill, and the Third Path	30
9.1 The prize	30
9.2 The bill: how much compute a nation’s science needs	31
9.3 The third path: an international consortium	32
10 The Convergence	34
10.1 The machine the elements form	34
10.2 What would falsify this thesis	35
10.3 The strongest counter-argument: two recursive loops, not one	35
10.4 Two things the argument has to carry with it	36
10.5 Implications for the West—and for the rest	36
A Forecasts and Dashboard	37
Sources and Selected References	40

Terms Used in This Edition

The full edition carries a complete notation glossary; the following are the terms a reader of this edition meets without introduction.

Open weights. A model whose trained parameters are published for download, so that anyone can run, fine-tune, or redistribute it, as opposed to a *closed* model reachable only through a metered API.

Token; \$/M tokens. The unit of text a model reads or writes; prices are quoted per million tokens.

MoE (mixture of experts); sparse model. A model in which only a fraction of the parameters (the “active” set) is used per token, so a very large model can run at the cost of a much smaller one; “280B/16B” means 280 billion total, 16 billion active.

Frontier; frontier-adjacent. The best models available at a given date; models within a few points of them on independent indices.

Harness. The runtime an agent runs in—tool schemas, permissions, memory, sandbox, session store, orchestration—as distinct from the model it calls.

Agent. A model given the authority to act (run tools, read and write files, make network calls) in pursuit of a goal over many steps.

Quantization. Storing weights at lower precision (e.g. four bits) to shrink memory and cost, at some loss of quality; “FP8” is an eight-bit floating-point format used in training and inference.

KV-cache. The working memory of context a model must hold while generating; it grows with

context length and is a major consumer of accelerator memory.

HBM (high-bandwidth memory). Stacked DRAM mounted on the accelerator package: very fast, scarce, expensive, and under the tightest export controls.

LPDDR; DRAM. Commodity memory of the kind in phones and laptops (LPDDR) and PCs (DDR); abundant and cheap, with far lower bandwidth than HBM. “SOCAMM2” is a compact LPDDR module format.

HBF (high-bandwidth flash). A proposed stacked-NAND memory tier in an HBM-class package: non-volatile, up to 512GB per stack, for holding model weights.

Capacity-first architecture; memory arbitrage. A design that trades scarce bandwidth for abundant capacity, using LPDDR or HBF where the incumbent uses HBM.

MFU (model FLOP utilization). The fraction of an accelerator’s peak arithmetic actually delivered during training; the “bandwidth knee” is the memory bandwidth below which MFU falls off.

ISA; RISC-V. The instruction-set architecture is the contract between software and processor—the layer at which Arm and x86 collect licence rents; RISC-V is an open ISA with no licence to revoke.

CUDA; NVLink. NVIDIA’s proprietary programming stack and chip-to-chip interconnect;

the software and interconnect moats of the accelerator incumbent.

Ascend, Cambricon, Kunlun. Chinese domestic AI accelerator lines (Huawei, Cambricon, Baidu).

SMIC; CXMT. China’s leading logic foundry and its leading DRAM maker.

EUV; DUV. Extreme- and deep-ultraviolet lithography; EUV is the tool China cannot buy.

TOP500; HPL; HPCG. The semiannual supercomputer ranking and its two benchmarks; “MLPerf” is the industry’s AI training and inference benchmark suite.

PUE. Power usage effectiveness: total facility power divided by IT power.

BGPU (baseline GPU). The full edition’s

vendor-neutral unit of AI compute demand, defined by serving throughput at 8 TB/s of memory bandwidth.

VLA model. A vision-language-action model: the policy layer of an embodied (robotic) system.

P1–P4. The book’s four propositions: adoption, sovereign compute, rent destruction, financial repricing (Chapter 1).

Evidence grades. [V] verified, [P] primary-source claim, [A] analyst estimate, [M] author-modeled, [R] roadmap, [S] scenario judgment.

The ratchet. The competitive equilibrium in which, once one Chinese laboratory opens at frontier scale, the others cannot profitably stay closed.

1 The Thesis

In one paragraph

Between roughly 2005 and 2025 China fought and won three complete technology wars—solar photovoltaics, lithium-ion batteries, and electric vehicles—by executing a single, learnable playbook. Artificial intelligence is the fourth run of that playbook. The mechanism is not to match the West chokepoint for chokepoint but to engineer the chokepoints out of the system: open-weight frontier models released at marginal prices, an open hardware ecosystem built on an unrevocable instruction set and a memory-capacity arbitrage that substitutes abundant commodity DRAM for scarce HBM, and a domestic semiconductor base that removes the last physical dependencies. The connective tissue is openness deployed as a weapon against the layers where Western incumbents collect their rents. The argument is decomposed into four propositions that are graded and falsified separately, and the book states in advance what would prove it wrong.

1.1 Three wars, one playbook

In solar photovoltaics China went from assembler of imported cells to producing more than 90% of every stage of the global supply chain, while module prices fell by more than 99% from their 1976 level [1, 2]. In lithium-ion batteries—a technology commercialized by Sony in 1991 and led for two decades by Japan and Korea—Chinese firms now hold roughly 69–75% of global electric-vehicle battery installations and over 80% of global cell manufacturing capacity, with pack prices down 93% in real terms since 2010 [3, 4]. In electric vehicles China became simultaneously the world’s largest car market, the world’s largest car exporter, and home to the world’s largest EV maker, while foreign incumbents’ share of the Chinese market collapsed from 62% to 35% in five years [5, 6].

These were not three lucky outcomes. They were three executions of one playbook, whose elements Chapter 2 distills: patient state designation of a strategic industry; demand creation at home; subsidized scale-up of manufacturing far ahead of demand; brutal domestic competition that functions as an evolutionary selection engine; capture of the entire upstream supply chain; a cost-down learning-curve war that bankrupts foreign incumbents; and finally, once dominance is achieved, the conversion of that dominance into geopolitical leverage—export controls on the very technologies China was once accused of copying [7, 8].

1.2 Why AI is the fourth war

The prevailing Western assumption is that AI is different: that frontier AI depends on scarce, controllable inputs—leading-edge logic, high-bandwidth memory, proprietary interconnects, closed model weights—and that export controls on these chokepoints can durably preserve a lead. The book argues that the assumption is failing on all fronts simultaneously, and failing in a characteristic way: not because China matches the West chokepoint for chokepoint, but because it is engineering the chokepoints out of the system.

The model. Chinese laboratories have converted the model layer into a commons. As of mid-2026, Chinese open-weight models occupy the top open positions on independent intelligence indices, account for roughly 41% of all Hugging Face downloads and a majority of tokens routed through neutral API aggregators, and are priced at a small fraction of Western closed models—while the measured capability gap to the best closed US models has compressed to low single

digits [9, 10]. Open weights at marginal cost do to the model layer what Chinese modules did to solar: they set a world price that rent-seeking business models cannot survive.

The compute. The state has designated the open RISC-V instruction set as national infrastructure, is standardizing matrix extensions and chip-to-chip interconnects across vendors, and has open-sourced the software stacks that would otherwise fragment the domestic ecosystem [11, 12]. In parallel, a memory-capacity arbitrage—substituting terabytes of cheap commodity DRAM for scarce HBM—attacks the assumption that frontier training requires centralized, HBM-bound superclusters at all.

The semiconductor. Beneath both, the fabs. SMIC has reached 5 nm-class logic without EUV; CXMT has become the world’s fourth DRAM producer and the swing supplier in a global memory shortage; the domestic equipment industry has placed three firms in the global top twenty [13, 14]. The remaining gaps—EUV lithography, leading-edge HBM—are real, but the strategy does not require closing them on the West’s terms; it requires making them irrelevant to the chosen architecture.

1.3 Openness as a weapon

The connective tissue among the elements is a deliberate inversion of the Western technology-rent model. Where the United States concentrated AI value in proprietary layers—CUDA, NVLink, closed frontier weights, x86/Arm licensing—China’s strategy, formalized in the July 2025 Global AI Governance Action Plan and the “AI+” State Council initiative, is to make every one of those layers open, standardized, and free [15, 16]. This is not altruism; it is the same move as pricing solar modules at \$0.09 per watt. An open layer generates no rent for the incumbent, but it generates *adoption*, and adoption generates the scale, the ecosystem, and the standards-setting power that constitute victory in the precedent industries. The West’s chokepoint strategy and China’s commoditization strategy are not symmetric opponents: one defends margins, the other destroys them.

1.4 Four propositions, stated separately

Rigor requires decomposing the thesis into four propositions that are related but logically independent—none proves the next, and each has its own evidence and its own failure modes.


P1 — Adoption. China is winning open-weight model adoption. *Confirmed by* download, derivative, and routed-token majorities persisting and extending into enterprise spend. *Falsified by* a durable reversal of the router and Hugging Face shares, or enterprise open-weight spend stagnating below about 15% while usage plateaus. Current grade: strongly supported [V].

P2 — Sovereign compute. China is constructing a domestically sufficient training-and-serving stack. *Confirmed by* a frontier-class model trained end-to-end on domestic silicon, packaging, and memory. *Falsified by* persistent dependence on stockpiled or smuggled foreign accelerators and HBM past about 2028. Current grade: rapid progress, not yet achieved [P/A/M].

P3 — Rent destruction. Open-weight commoditization will compress the closed-frontier rents that finance Western laboratories. *Confirmed by* closed-model revenue growth decelerating while capability parity holds. *Falsified by* closed labs sustaining pricing power through service, distribution, and trust layers even at model parity. Current grade: mechanism operating, magnitude contested [V/S].

P4 — Financial repricing. The compression of P3, if it occurs, punctures the US AI capital structure. *Confirmed by* the indicator table of Chapter 6. *Falsified by* demand elasticity absorbing the price decline: usage growing faster than prices fall rescues the buildout even under P1–P3. Current grade: scenario judgment [S].

Figure 1.1 shows the propositions with their couplings and their contrary evidence. P1 can hold while P2 fails (China dominating distribution on stockpiled NVIDIA silicon); P3 can hold while P4 fails (margins compress while volume growth fills the data centers); and a US valuation correction could arrive from macro causes with none of P1–P3 mattering. The convergence argument is that P1 is established, P2 is trending, and the *coupling* of $P1 \times P2$ —open weights meeting sovereign, capacity-first hardware—is what forces P3, with P4 as its priced-in consequence.



None of these arrows is entailment. P1 can hold while P2 fails (dominance on stockpiled foreign silicon); P3 can hold while P4 fails (margins compress while volume growth fills the data centres, the elasticity case); and P4 could arrive from macro causes with none of P1–P3 mattering. The convergence argument of this book is that the *coupling* $P1 \times P2$ —open weights meeting sovereign, capacity-first hardware—is what forces P3, with P4 as its priced-in consequence.

Figure 1.1: The argument, disaggregated. Each proposition carries its own evidence grade, its own confirming metrics, and its own contrary evidence; the book’s thesis is a claim about the coupling between them rather than a single chain of entailment. Grades as of the version stamp; the full edition’s evidence ledger carries the full ledger and its dashboard the quarterly indicators that move them.

1.5 Evidence grades

Load-bearing claims are graded throughout: [V] independently verified; [P] primary-source claim (reliable as to what was claimed, not what was achieved); [A] analyst estimate; [M] author-modeled result; [R] roadmap; [S] scenario judgment. The full edition collects every headline claim in a graded ledger and a machine-readable claim register [full edition, App. C, I]; Appendix A of this edition reproduces the forecasts and the quarterly dashboard.

1.6 The theory, in brief

The full edition devotes a chapter to the machinery behind the analogy [full edition, Ch. 2]; four ideas from it are needed to read what follows.

A cost stack, not a price. The cost of one completed useful task decomposes into amortized model creation, inference, integration and organizational adoption, complementary capital, and assurance. Only inference is falling at ten times per year. Integration and adoption are not falling—which is why open-weight *usage* share rose while open-weight *spend* share fell—and assurance is a cost the commons has mostly not paid. Commoditizing inference therefore does not commoditize delivered work; it shifts the binding constraint to the terms that are not falling. The rent-destruction claim is bounded to the model layer and remains a hypothesis at the service layer.

The elasticity knife-edge. Whether spending on AI rises or falls when its price falls is the whole financial argument stated without mathematics. With own-price elasticity ε , revenue moves as $\Delta \log R = (1 - \varepsilon) \Delta \log P$: above one, price cuts pay for themselves; below one, falling prices shrink revenue; $\varepsilon = 1$ is the boundary between the bull and bear cases for the entire Western AI complex. The honest position is that ε is a portfolio, not a number—plausibly above one for coding agents, synthetic data, and scientific inference, near or below one for consumer chat—and that the leveraged Western structure needs it above one at the premium tier specifically. Graceful degradation is a real Chinese advantage; it is not a theorem of victory.

Openness relocates rent; it does not abolish it. A seller rationally prices at or below marginal cost whenever the extra volume raises revenue on a complement by more than the margin given up. Each side opens the layer where it is behind and expects rent in the layer it is best at: China opens weights, ISA, interconnect, and framework and expects to recapture on hardware volume, domestic cloud, standards, and dependence; the United States controls all four and expects premium capability, distribution, integration, trusted deployment, and agent platforms. Two flywheels result, and each can steal the other’s critical input—China steals volume, America steals value. Which theft dominates is, on the book’s reading, the single most important unresolved empirical question in the field.

Four layers of winning. AI is a general-purpose technology, which means a country can dominate its production and capture little of its value, because most of the value appears downstream in diffusion and productivity. Production leadership is the only layer on which the United States is unambiguously ahead. The platform layer is going to whoever opens fastest; diffusion is where most realized value sits; productivity is the only layer that finally matters and the least observable in real time. A country can lead one layer while losing another.

The thesis, stated in its conditional form: the United States retains the absolute model frontier, the leading semiconductor and cloud ecosystem, and vastly greater private capital. China increasingly leads a different contest—the commoditization of frontier-adjacent intelligence and the construction of a sovereign, though incomplete, compute stack. The decisive question is whether Chinese cost reduction, adoption, standards formation, and hardware sufficiency compound faster than US capability, distribution, and monetization—and whether China closes its remaining gaps in memory, packaging, tools, software maturity, and trust before the open-model economy compresses the rents that finance the Western frontier. This book argues that convergence is the way to bet, and Chapter 10 states what would change the bet.

2 The Precedents and the Playbook

In one paragraph

Solar showed the pure form [full edition, Ch. 3–7]: Western invention and demand, Chinese policy-financed scale, capture of the learning curve, failed tariffs, a home market that made foreign trade policy irrelevant, over 90% share at every upstream stage, then state cartelization. Batteries added vertical depth from mine to pack, won by regulatory market reservation and by industrializing the “inferior” chemistry until the West licensed it back. Electric vehicles proved the playbook works on a whole system and turned the car into a Chinese-led embedded-AI platform. From the three campaigns the book distills nine tactical moves and six strategic steps; from the campaigns that have *not* worked it distills where the playbook grinds: tacit knowledge, slow clock speed, foreign certification, installed-base lock, and trust. Scored against that model, AI’s model and compute layers sit in the kill zone; its deployment layer is genuinely contested.

2.1 Three campaigns, in brief

Solar was invented at Bell Labs and built on American, German, Japanese, and Australian science; Sharp was the largest cell producer in 2006, and demand was created by Germany’s feed-in tariff. Chinese scale-up rode that demand and about \$47 billion of development-bank credit [17]; polysilicon fell from about \$475 to under \$16 per kilogram between 2008 and 2012 [18], and Solyndra, Q-Cells, and SolarWorld died—as did the Chinese pioneer Suntech, because the state protects the industry, not the champion. US and EU duties from 2012 taxed deployment without resurrecting manufacturing, and a domestic market of 315 GW_{AC} in 2025 alone made foreign trade policy irrelevant [19]. The end state is 84–98% Chinese share by stage—over 90% at every stage upstream of the module [2], module prices down from \$106 per watt in 1976 to about \$0.09 at a learning rate near 20% per doubling [1, 20]—and since roughly 2010 essentially all doublings were Chinese, so whoever manufactures the doublings owns the learning. Winning is expensive: losses below cash cost, then Xi’s July 2025 anti-involution campaign and a polysilicon buyout platform—an OPEC for solar [21, 22]; the state cartelizes what it has captured. First Solar survives as a tariff-protected niche shipping one fortieth of Chinese volume. Four lessons: trade barriers after cost-curve capture protect only domestic prices; champions are disposable, the industrial commons is not; a home market absorbing half of world demand makes foreign trade policy irrelevant; and commoditization ends in state-managed consolidation. Keep them in mind when reading about open-weight models priced near zero.

Batteries were commercialized by Sony in 1991 and led by Japan and Korea for two decades; China entered through low-margin consumer cells (BYD, ATL, then CATL) and pivoted when the state created demand. The decisive instrument was a certification, not a tariff: the 2015 white list of 57 domestic makers eligible for NEV subsidies, which excluded LG and Samsung and was scrapped in 2019 when no longer needed [23]. By 2025 Chinese makers held about 69% of installations and over 80% of the world’s 4 TWh of cell capacity [3, 24], plus 70–90% of most upstream layers and two thirds of lithium and cobalt refining, at pack costs that leave North America and Europe paying a 44–56% premium [4]; a gigafactory built outside China is built, in effect, inside the Chinese supply chain. Lithium iron phosphate—Western-invented, dismissed as a dead end—passed 55% of the global EV market and is now licensed back to Ford and Stellantis: the inferior branch industrialized at scale beat the premium branch on system economics, a preview of the LPDDR-versus-HBM argument of Chapter 3. In 2025 the infant

industry became the gatekeeper with export controls on battery technology, the mirror image of US chip controls [7]; Northvolt died on \$15 billion. Three more lessons: regulatory reservation beats tariffs and is invisible to trade law; vertical integration to the mine makes a foreign factory a hostage; watch what China industrializes, not what the West benchmarks.

Electric vehicles are a whole system wrapped in a century of incumbency; if the playbook worked here, no consumer-industrial market is safe, and it worked. Wan Gang’s leapfrog decision predates Tesla; at least \$231 billion of layered demand support [25], a dual-credit mandate binding foreign joint ventures, and 19 million charging points took NEV penetration from 6% in 2020 to 54% in 2025 [26] and made China the largest car exporter [6]. BYD ran the flywheel from 416,000 to 4.55 million vehicles with about 80% of Tier-1 content in-house and a \$4,700 structural advantage over Tesla of which subsidy is under \$300 [5]; some 400 Chinese EV ventures died around it—the state subsidizes the species, not the specimen. Foreign brands fell from 62% to 35% of the Chinese market and technology now flows backward (Volkswagen licenses XPeng, Toyota uses BYD powertrains); a century of licensing reversed in one product generation. And the car became a software platform in Chinese hands—about 95% of automotive lidar, city autonomy at scale, driver assistance free on an \$8,000 car, 20–24-month cycles against 40–50—so the first mass-market safety-critical embedded-AI category is Chinese-led and anchors the domestic chips, sensors, and models of the next chapter. Three final lessons: pick the transition where incumbency resets to zero and force it at home; a system-level win pulls the whole component stack with it; speed compounds.

2.2 The playbook distilled

Table 2.1 inventories nine tactical moves that recur across the campaigns; the AI column is developed in Chapter 3. Six strategic steps are the sequence the moves serve. *Target selection*: an industry at a technology discontinuity where the incumbent’s accumulated advantage is about to be stranded and where manufacturing learning-by-doing, not protected IP, is the true accumulating asset. *Demand before supply*: the home market as guaranteed demand, foreign firms admitted just long enough to transfer capability. *Overcapacity as strategy*: capacity at multiples of world demand, which Western economics reads as waste and which purchases the entire learning curve, a permanent price umbrella, and surge capacity; the losses land on Chinese balance sheets and dead firms while the cost curve remains. *Selection, not planning*: Beijing floods the sector and lets domestic combat select; its loyalty is to the industrial commons—the engineer pool, the supplier mesh, the tooling ecology. *Commoditize the adversary’s rent layer*: wherever the Western business model collects rent—module, pack, badge, API token, ISA licence—price it at marginal cost, because rents fund the West’s R&D flywheel; openness is this move in its purest form, since the rent goes to zero by construction and adoption accrues to the commoditizer. *Endgame*: consolidate, stabilize prices upward, and convert dominance into coercive leverage through export controls on the now-Chinese-controlled layer.

Against this the West has fielded two counter-patterns, both failing. Tariff walls arrive after cost-curve capture and merely tax domestic consumers while the walled market shrinks in relevance. Chokepoint denial—semiconductor export controls since 2019—is the more serious strategy, and its record is a pattern the book calls *chokepoint decay*: each denied input (EUV, HBM, H100s, EDA) triggers a state-scale substitution program plus an architectural detour around the input, while the denial hands the domestic substitute a guaranteed market—the white-list mechanism, imposed this time by Washington on Beijing’s behalf. Export controls are, functionally, a US-funded market-reservation policy for Huawei, SMIC, CXMT, and Cambricon.

Table 2.1: The nine tactical moves, across four campaigns.

Move	Solar	Batteries	EVs	AI
M1 State designation	SEI 2010, 12th FYP	White list 2015	863 program, SEI 2010	“AI+” plan, 15th FYP, RISC-V guidance
M2 Home demand	National FiT, Top Runner	NEV subsidy linkage	\$231B support, dual credit	State DCs mandated on domestic chips, subsidized power
M3 Subsidized scale-up	CDB \$47B credit	Provincial capital	Provincial EV funds	Big Fund III (\$47.5B), SMIC/CXMT capex
M4 Darwinian competition	2018 “531” cull	100+ makers to consolidation	~400 dead EV firms	LLM price war since 2024; chip-vendor shakeout beginning
M5 Supply-chain capture	poly → module	mine → cell	battery → chip → magnet → ship	model → framework → chip → interconnect → fab → DRAM
M6 Commoditize the rent layer	module \$/W	pack \$/kWh	vehicle BOM	open weights, \$/Mtoken, open ISA and stack
M7 Absorb the “inferior” branch	mono-PERC, TOPCon	LFP	small BEVs	sparse MoE; LPDDR capacity vs HBM bandwidth
M8 Standards capture	efficiency floors	GB safety standards	charging standards	RISC-V profiles, matrix ISA, UnifiedBus
M9 Leverage after victory	OPEC-for-poly	tech export controls 2025	rare-earth controls 2025	(pending)

2.3 Where the analogy breaks

A playbook argued from three victories owes the reader its defeats. The same state, spending an estimated \$248 billion a year on industrial policy [27], has run two-decade campaigns where dominance has not arrived: commercial aircraft (COMAC’s C919 delivered about thirteen aircraft in 2025 against a plan of 75, on a Western engine whose export licences Washington suspended and restored in 2025—chokepoint leverage working); lithography and the deep tool chain (under 5% localized; twenty years and three Big Funds have not cracked it [28, 29]); precision machine tools and reducers (five-axis CNC about 8% domestic against an 80% target); operating systems (state Linux never dented Windows; HarmonyOS reached 17% only after sanctions burned the Android bridge—ecosystems yield on decade timescales, under existential forcing, inside the home market); and biotech, the late bloomer, where \$138 billion of 2025 out-licensing meets a residual gap that is entirely trust and regulation—the FDA’s 14–1 rejection of China-only trial data is the certification barrier in its purest form [30]. The playbook wins where products are modular, iteration is fast, learning is manufacturable, the home market can certify, and trust is embodied in the artifact; it grinds where knowledge is tacit, certification is foreign, and value lives in an installed ecosystem.

Scored against that model, AI models are the most modular, fastest-clock-speed artifact in industrial history—sparse-MoE training is closer to batteries than to jet engines—and compute is intermediate, with the tool chain inheriting lithography’s resistance and being routed around. But three properties of AI sit in the resisting column and are the strongest structural arguments against this book’s title: the frontier moves, and recursive AI-for-AI research could reopen the gap faster than any factory closes it; ecosystems lock, and the enterprise deployment layer is Office, not modules; and trust and certification bind, with the sintilimab precedent—capability accepted, certification denied—directly available to Western regulators. Two asymmetries follow: zero reproduction cost lets Chinese weights conquer adoption instantly and would let a superior Western open release reconquer it just as fast, so there is no factory-shaped moat around P1; and talent is the one input the two systems still share—researchers of Chinese undergraduate origin were 47% of the top tier in 2022, up from 29% in 2019, while the share of that tier working in the United States fell from 59% to 42% and China’s rose from 11% to 28% [31]. The honest

conclusion: model and compute sit mostly in the kill zone, the tool chain in the resistance zone and bypassed, and the deployment layer—trust, certification, ecosystems—is contested ground where the precedents predict a long grind, not a collapse. That is where Chapter 5 and the strategy chapters place their weight.

3 The Elements of AI

In one paragraph

Read separately, each element has a Western rebuttal: the models trail the closed frontier by a few points, the accelerators trail NVIDIA per watt, the fabs trail TSMC by two nodes. Read together, the rebuttals miss the structure. The model layer has been converted into a commons by a self-enforcing open-weight ratchet; the harness—the runtime that agents are programmed against—is being given away one layer up; a GPU-free Chinese machine leads both TOP500 benchmarks and shows the CPU/accelerator boundary to be an allocation choice, while a modeled memory-capacity arbitrage and a priced five-thousand-dollar appliance [M/S] attack the premise of the centralized supercluster; the chokepoints on Chinese silicon are decaying faster than they are replaced, with domestic HBM the one binding constraint; the watt is China’s deepest structural asymmetry, aimed by policy at domestic chips; deployment-generated data is where rent migrates when the model becomes free; and embodied AI is where the loop closes back onto the factory floor. None of it is finished, every claim is graded, and the single integration test that has not been passed—a frontier model trained end to end on domestic silicon—is the “Mate 60 moment” the forecasts price at 65% within two to four years.

3.1 The model: open weights as the commodity layer

The Western frontier is closed: the best US models are accessible only as metered APIs, their weights corporate secrets, their pricing a rent on capability. The Chinese frontier is open: downloadable weights, mostly under MIT or Apache licences, at API prices near marginal inference cost. The inversion began with DeepSeek. Its V2 release in May 2024, with multi-head latent attention and fine-grained mixture-of-experts at roughly RMB 1–2 per million tokens, triggered a domestic price war in which ByteDance, Zhipu, Alibaba, and Baidu cut prices by up to 97% within two weeks. V3 in December 2024 reached GPT-4-class capability at a published marginal training cost of \$5.6 million—a final-run figure [P], not a program cost; estimates of DeepSeek’s buildout run up to \$1.6 billion [A] [32]. R1 in January 2025, MIT-licensed, matched o1-class reasoning, topped the US App Store, and erased \$589 billion of NVIDIA’s market value in a day [33]. The strategic content was mispriced by markets and correctly priced by Beijing: frontier capability can be commoditized; export-control-constrained hardware forced algorithmic efficiency—latent attention, sparsity, FP8, GRPO, multi-token prediction—that compounds on any hardware; and open weights convert capability into global adoption instantly. The defensible causal claim is constraint-consistency, not unique causation: nearly all of these techniques would be commercially attractive without sanctions, and the genealogy is graded accordingly [full edition, Ch. 8].

The state did not anoint DeepSeek; it flooded the sector. By mid-2026 the release lag from the US closed frontier to a Chinese open equivalent had compressed from six to eight months to a matter of weeks, and the breadth of the field is itself the datum: Alibaba’s Qwen family (a billion cumulative downloads, more than half of global open-model downloads, over 200,000 derivatives) [34], Moonshot’s Kimi K3 (2.8 trillion parameters, the largest open-weight model

ever released, in the global top five and the only open model in that tier) [35], DeepSeek V4, Zhipu’s GLM-5 trained on Ascend, MiniMax, Xiaomi’s 1.02-trillion-parameter MIT model, and Xiaohongshu’s dots3-note—frontier-scale development is now an expected capability of any large Chinese platform. Token prices run from one seventh to one hundred-and-fiftieth of US equivalents. Chinese models account for about 41% of Hugging Face downloads [9] and roughly 61% of tokens routed through the largest neutral aggregator, up from about 2% in early 2025 [36]; roughly 80% of open-stack startups use a Chinese model. The Stanford AI Index put the top US–China gap at 2.7% in March 2026, down from 17–32% in 2023, and the closed-over-open gap at 3.3%, reopened from 0.5% in 2024—a 23:1 private-spending ratio buying a 2.7% measured gap [10].

Two developments in the summer of 2026 changed the character of the commons. The first is the *open-weight ratchet*. Alibaba’s Max tier had always been proprietary; seven days after Moonshot opened Kimi K3, Alibaba committed to opening Qwen3.8-Max, and shipped the weights on 13 August [37]. Once one Chinese laboratory opens at frontier scale, the others cannot stay closed—adoption is the currency and closing your best model, in that ecosystem, is unilateral disarmament. This upgrades openness from a state policy that could reverse to a competitive equilibrium from which no participant profits by defecting. But the ratchet holds on weights and fails on terms: the Qwen3.8-Max licence requires a commercial agreement above US\$50 million of revenue, and Kimi K3’s carries the same structure, a separate agreement above an undisclosed revenue threshold [38]. Opening the weights does not abolish rent; it relocates it to large-scale commercial service, where the releasing lab retains a claim—a commons with a toll gate at commercial scale. The second is Meta’s return to open weights on 10 August with Muse Glimmer, a 30-billion-parameter distilled model, Apache-licensed and competitive if not commanding, with an open “version” of the flagship promised on unspecified terms. Whether that re-anchors a Western pole for the commons or merely confirms that the ratchet now operates on both sides of the Pacific is registered as a forecast—on the evidence to date the full edition reads it as the latter—and either way, nobody now argues the commons does not matter.

Three qualifications keep the model chapter honest. Capability, distribution, and revenue are three different scoreboards, and the third lags: enterprise open-weight *spend* share fell from 19% to 11% in 2025 even as usage rose [39], which is the strongest evidence against P3 and is treated as such. Alibaba’s claim that Qwen3.8-Max ran “500-plus turns of chip design optimization” toward an autonomous silicon design flow—if true, an attack on the one input export control cannot embargo at a border—is the vendor’s own and unverified. And US closed models are rumoured to be roughly twice the parameter count of the Chinese flagships: this is an efficiency victory, not a scale victory. What quantization adds is operation: dynamic quantization and native four-bit release formats put frontier inference on \$2,000–\$40,000 hardware, the capacity-rich design point discovered bottom-up by hobbyists that Chinese hardware policy is discovering top-down. Open weights alone democratize access; quantization democratizes operation. And the right denominator is cost per solved task, not price per token, because reasoning outputs are inflating about fivefold a year [40] and no vendor publishes the number that matters. In playbook terms the model layer is already in the late-war stage: rent destroyed, adoption captured, standards increasingly set in Chinese repositories. What open weights cannot guarantee is sovereign hardware to train and run them.

3.2 The harness: the layer the contest moved to

Every preceding paragraph treated the model as the object of competition; the object is moving. Models are swappable—an API endpoint exchanged by configuration, a product category exists to make it trivial. A *harness*—the runtime that holds the agent loop, tool schemas, permission system, memory and context management, sandbox, session store, and subagent orchestration—is what an organization writes code against, and its abstractions become the integration, its

execution semantics become the test assumptions, its permission model becomes what an auditor reads. The layer where switching is expensive is not the layer everyone measures, and a strategy that wins the harness can lose every benchmark and still take the ecosystem. The sole controlled study found that adopting a better harness with *any* model saved 33–61% of cost per task, where switching from the cheapest to the most expensive model saved 36% [41]; and once a self-hosted open model runs at about a dollar per million output tokens against \$25–30 for the closed frontier, the harness is closer to the only lever left on production economics.

On 13 August 2026 DeepSeek released an MIT-licensed, plugin-native, model-agnostic harness that passed 100,000 GitHub stars in about 48 hours [42, 43]—built on the kernel of a four-year-old Chinese chatbot framework whose creator DeepSeek recruited. That is an ecosystem-acquisition strategy, and it is this book’s thesis executed one layer up. It shipped the same day DeepSeek raised its API prices [44]: commoditize the complement, monetize the scarce layer, and ensure the commoditization damages competitors’ pricing power at least as much as your own—the rent-relocation thesis performed rather than argued. The next day Xiaohongshu shipped an Apache-licensed omni-modal agent model trained for ten-hour agent runs; model and harness are being co-designed the same week by two Chinese labs. Meanwhile the Western agent-standards bodies—the foundation that now hosts MCP, the A2A consortium, the Agent Plugins group—have no Chinese named member, and Beijing’s cyberspace administration has published a rival “interoperability” initiative naming none of them [45]. The incumbents convene a consortium; the challenger declines to join, ships a permissively licensed implementation, and bids for the installed base directly. The book grades the strong lock-in claim as a hypothesis [S]—the interface is commoditizing, the architectures have converged, and nobody has measured switching cost—but lock-in relocates rather than disappears: to distribution, to the evaluation substrate (whoever owns the harness owns the conditions of measurement), and to certification. If American models remain the best and the world’s agents are orchestrated by runtimes shipped from Hangzhou and Beijing under MIT licences, superiority at the model layer buys less than its owners expect [full edition, Ch. 9].

3.3 The compute: RISC-V, accelerators, and the memory arbitrage

The Western AI machine is a proprietary stack—x86 or Arm hosts, NVIDIA accelerators, CUDA, NVLink, HBM—each layer a licensed chokepoint. China’s counter-design is an open machine: an open instruction set, multi-vendor domestic accelerators unified by open interconnect and open software, and a memory system that trades scarce bandwidth for abundant capacity. The claim is deliberately narrow: not that the GPU is obsolete, but that the CPU corner of a design continuum can reach parity on the workloads that matter while carrying an ecosystem advantage the accelerator corner does not have.

RISC-V. In March 2025 eight government bodies drafted national RISC-V policy guidance [11], and the 15th Five-Year Plan mandates “extraordinary measures” for self-reliance with the open ISA as its core mechanism [46]. Because RISC-V is an open standard, its deployment insulates Chinese silicon design from sanctions at the ISA layer permanently: there is no license to revoke. Twelve of the twenty-four premier members of RISC-V International are Chinese [47]; the matrix-extension task groups now standardizing the operations AI silicon executes are, in the book’s phrase, the single most consequential unglamorous venue in the field, and their ratification is only the first rung of a maturity ladder from specification to compiler to kernel library to framework to fleet validation—so the RISC-V accelerated-CPU path is real but dated. Washington’s attempts to restrict RISC-V participation produced nothing; the Commerce Department conceded that restriction would harm US firms without slowing China. And the platform owner of the AI era has ported its crown-jewel software to the open ISA at a summit in Shanghai: when NVIDIA brings CUDA to RISC-V, Google and Meta build accelerator control planes on it, and

the standards war is not being lost but conceded.

The existence proof. On 24 June 2026 LineShine, at the national supercomputing center in Shenzhen, took first place on both the HPL and HPCG benchmarks with 2.198 exaflops of FP64—the first two-exaflop machine, the first Chinese No. 1 since 2017—and no GPUs [48, 49]. Its LX2 processor is an Armv9 part with 304 cores, scalable vector and matrix extensions in every core, and 32 GB of on-package HBM: Fugaku’s lineage carried to its conclusion, the accelerator’s defining features absorbed into the CPU. Export controls denied China the discrete GPU; the response was to demonstrate, at No. 1 in the world, that the functional boundary between CPU and accelerator is an implementation choice—and then to choose an implementation with no controllable product category in it. The book is exact about the proof boundary: HPL and HPCG are settled [V]; AI serving parity per terabyte-per-second of bandwidth and a training penalty of $4.5\text{--}8.5\times$ that is implementation-fixable are the author’s modeled results with Dongarra and Hoeffler [M] [50, 51]; and a seven-row proof suite states, benchmark by benchmark, what would settle or falsify the AI claim—one of which is a domestic Chinese submission to MLPerf Training, whose June 2026 round drew 95 systems on thirteen accelerators, none of them Chinese [52]. Nonparticipation is not proof of weakness, but it is a measurable international-validation gap.

The fleet. US bans and Beijing’s own demand squeeze took NVIDIA’s China data-center revenue from a 95% share to approximately zero [53]; state-funded data centers must use domestic chips. Ascend, Cambricon, Kunlun, and T-Head fill the vacuum—domestic parts were about 41% of Chinese AI-chip shipments in 2025 and are projected above 60% in 2026 [54]—and fragmentation is engineered away by an interconnection mandate and by Huawei open-sourcing UnifiedBus, CANN, MindSpore, and its SuperPoD architecture [12, 55]. The Western answer to NVLink is a consortium that excludes China; China’s is an open specification with state-mandated adoption. The book concedes what must be conceded: NVIDIA plus HBM remains superior per watt for dense training; domestic output is an order of magnitude below NVIDIA’s; Ascend’s immaturity cost DeepSeek a flagship cycle; and NVIDIA’s own turn to arguing on tokens per watt and KV-cache constraints concedes the memory diagnosis while making China’s HBM deficit more damaging—survivable only if the architecture needs less of it.

The memory arbitrage. That is what the arbitrage is for. HBM consumes about three times the wafer capacity of standard DRAM per bit and sits under the tightest export controls; commodity LPDDR6 is where CXMT is at or near the world’s leading edge. A RIKEN feasibility study finds that for a Kimi-K3-class 2.8-trillion-parameter sparse model the sharded training state exceeds HBM4 socket capacity below roughly 256 accelerators, while a socket carrying about eighteen SOCAMM2-class modules holds 2.3 TB—eight times the HBM4 ceiling—at a modeled MFU ceiling near 67% against 89% for HBM4 [M] [56]. Worse on bandwidth, better on capacity per dollar, initially adopted by customers the incumbent does not want: this would be the LFP move, executed in memory—proposed and modeled here [M], not yet an observed Chinese product program. The engineering is not free—the full edition’s appendix prices the pin budget, the power, and the package yield, and recommends a hybrid HBM-plus-LPDDR tier as the baseline with pure LPDDR as the sanctions-contingent corner [full edition, App. A]. The strategic consequence is decentralized pretraining: a 128-socket cluster at roughly half a megawatt—an ordinary industrial substation feeder, not a campus—puts frontier-class sparse pretraining within reach of any sovereign state, university consortium, or mid-cap firm, each one a customer for the Chinese stack and a defection from the Western API economy.

The appliance. The recurring cost that decides whether a defection sticks is inference, and there the same logic runs sharper. The open frontier is now a terabyte-class object: Kimi K3 is about 1.4 TB even in its native four-bit format, and the escape hatch of aggressive quantization has been spent by the model publishers themselves. Open weights are a commons only for those with the memory to hold them. Two technologies would change that—High Bandwidth Flash, whose first OCP specification appeared on 3 August 2026 as a non-volatile weights tier of up to 512 GB per stack [57, 58] [full edition, App. B], and stacked-DRAM alternatives to HBM—and

the book prices a capacity-first appliance built from them: 2 TB of HBF, 256 GB of LPDDR6X, and a GB10-class SoC, at a modeled bill of materials of \$2,845–4,352 and a retail price of \$3,800–7,250 depending on the pricing case, anchored to shipping systems such as the DGX Spark and Framework Desktop rather than to bare-component guesses [M] [59, 60]. Two declared assumptions govern the price: memory is costed on a pre-surge, mid-2025 basis, and HBF itself has no shipping product—the first die taped out in August 2026, with samples promised for 2027—so this is a priced design, not a quote. Table 3.1 sets it against a Vera Rubin NVL72 rack fully loaded with its building, facility operations, and cost of capital, and against the August 2026 API price list. The appliance undercuts the closed-frontier list price by an order of magnitude at any realistic load; at a modest 100 million tokens a month it still costs more than the rack, but at about 300 million it crosses below even a hyperscaler’s own fully loaded cost of serving the same open model at planned utilization, and the gap widens whenever utilization disappoints; the crossover is sensitive to tens-of-percent shifts in the inputs, the utilization asymmetry is not. The centralized batching advantage that made the rack look unbeatable is roughly fifty times smaller for an ultra-sparse mixture of experts than for a dense model. Once the machine exists, the question every enterprise buyer asks changes from “which API” to “why an API at all”—the innovator’s dilemma stated as a purchase order. The case against is strong and is printed at full strength: the entire memory industry expects shortage through 2030 [61], NVIDIA cancelled its own capacity-first inference part and bought the opposite architecture, and enterprise open-weight adoption fell during the very period the arithmetic favoured it. The claim is therefore at scenario grade, with the price premise its most vulnerable link.

Table 3.1: Cost per million output tokens, Kimi K3 class, mid-2026. Appliance and rack are fully loaded (silicon, building or closet, operations, power, 8% cost of capital); API prices are list prices [M costs; V prices].

Configuration	\$/M output tokens	Note
Appliance (\$5,000 base case), 50M tokens/month	3.16	light use
Appliance, 100M tokens/month	1.65	about 19% of capacity
Appliance, 300M tokens/month	0.64	about 56% of capacity
Appliance at capacity (540M/month)	0.41	
Rubin NVL72 rack, \$7.8M, 60% utilization, interactive parity	1.17 / 0.90	4-yr / 6-yr amortization
Rubin NVL72 rack, 40% utilization	1.76 / 1.35	
Rubin NVL72 rack, 30% utilization	2.35 / 1.80	
Rubin NVL72 rack, \$10M enterprise retail, 60%	1.44 / 1.10	
Closed frontier API (Claude Opus 5 / GPT-5.6 Sol)	25 / 30	list
Chinese open API (Kimi K3 / DeepSeek V4-Pro)	15 / 3.96	list

The compute assessment, then: no single Chinese node matches its Western counterpart at the frontier, and the architecture was chosen so that none needs to. Hardware proliferation plus software commoditization equals the democratization—on Chinese terms—of the means of AI production. The remaining dependency is physical: wafers, HBM, DRAM [full edition, Ch. 10].

3.4 The semiconductor: fabs, logic, and DRAM

Models can be open-sourced and architectures standardized overnight; fabs cannot. This is where the West concentrated its chokepoint strategy, and therefore where the test of chokepoint decay is sharpest. The evidence through mid-2026: the chokepoints are decaying faster than they are being replaced.

In logic, SMIC has crossed the export-control line twice—the 7 nm Kirin in the Huawei Mate 60 of August 2023, and by December 2025 TechInsights-confirmed volume production on a node close to a 5 nm equivalent without EUV [V] [13, 62], though “significantly less scaled than leading commercial 5 nm.” Capability is demonstrated; TSMC-equivalent economics are

not, and the metrics that would settle it are undisclosed [A]. But the strategic arithmetic is not yield percentage; it is wafers times yield times time, and SMIC’s advanced-node capacity climbs from about 45,000 to 80,000 wafers a month between 2025 and 2027 [63] behind a network of Huawei-linked “shadow fabs” and a bank of pre-ban dies and HBM stacks that bridges the ramp. Below the leading edge the story is already solar-shaped: China supplies 37% of global 22–40 nm output and about 70% of mature-node capacity expansion, which is where the EV and robotics stack lives.

Memory is the quiet revolution. CXMT scaled from about 70,000 wafers a month in 2022 to 200,000–300,000 by mid-2026, targeting 600,000, and will be within about 25,000 wafers a month of Micron by year-end—the world’s fourth, and by trajectory third, DRAM producer [64, 65], with DDR5 yields above 80%. Its July 2026 Shanghai listing raised \$8.6 billion and closed up 466% at roughly \$460 billion of market value [14, 66]; the book insists that with under 7% floated this is a sentiment indicator and nothing more, and tracks capacity, yield, and qualification instead. The 2025–26 global memory shortage handed CXMT the swing-supplier role, and within twelve months the narrative flipped from “Chinese dumping” to Chinese pricing power—a moment in a panic market, not yet proof of durable pricing power [A]. HBM is the ledger’s honest exception: CXMT is in production at HBM2/2E, three to four years behind SK hynix, sampled HBM3 to Huawei in late 2025, and disclosed no dedicated commercial HBM investment in its own prospectus; no public teardown has yet found Chinese-fabbed DRAM in a shipping accelerator [67], and Huawei’s stockpile of about thirteen million foreign stacks is assessed as effectively exhausted by mid-2026. Domestic HBM is therefore the single binding physical constraint on the compute element—and the most important open physical variable in the argument. The strategy is double-tracked: sprint on HBM three years behind, and lead on the commodity LPDDR class that the capacity-first architecture was designed to need. China has not made HBM irrelevant; it is engineering a position in which HBM ceases to be the *sole* architectural path to frontier-scale AI, while racing to supply it domestically anyway.

The tool base is the deepest dependency, and there the trend, not the level, is the message: Naura fifth in global equipment revenue, three Chinese firms in the top twenty, domestic tools 20–30% of Chinese fab demand against 10% three years earlier [68, 69]; lithography the hard floor, with domestic EUV no earlier than about 2030; a US EDA ban rescinded within six weeks in exchange for rare-earth relief while domestic EDA share doubled. Big Fund III carries RMB 344 billion; China has bought about 40% of the world’s fab equipment every year since 2020 [70, 71]. Targets are missed and reset, but the direction never reverses, and each US control action functions as a demand guarantee for the domestic substitute. The book proposes ranking export controls by the effective delay they actually buy—EUV buys real years; HBM controls buy less than assumed given stockpiles, diversion, and the architectural detour; accelerator controls have largely converted into a domestic market-reservation policy; EDA near zero—and observes that this discipline would have predicted, in 2022, most of what has since happened. The chokepoint strategy assumed China would try to rebuild the Western machine and could be stopped at its scarcest parts. China is building a different machine. Whether it works end to end—logic yield, packaging, memory qualification, software maturity, and capacity-first training performance demonstrated simultaneously in one frontier run—is the integration test not yet passed [full edition, Ch. 11].

3.5 The watt: energy as the fourth element

Every token is electricity through silicon, and the AI war’s longest-lead-time input is neither the model, the chip, nor the fab—it is the watt. The obvious comparison is the wrong one: nameplate capacity added is not deliverable power at a computing site. Even so, the ledger has four rows and China leads on all of them (Figure 3.1). New nameplate: about 540 GW added in 2025, 452 GW of it renewable, against a record 53 GW in the United States [72, 73].

Generation: 10,573 TWh against 4,430 [74]. Firm and dispatchable capacity: a record 95 GW of coal commissioned as flexible backup, roughly 66 GW of nuclear heading toward 110, and two thirds of the world’s grid-storage installations against American gas turbines sold out to 2029–30 and under 1.7 GW of nuclear restarts. Deliverable headroom: 45 ultra-high-voltage lines and four new 8-GW corridors, with state-directed siting, against a US interconnection queue of roughly 2,000 GW with a median wait above three years [75], large-load connection studies of up to seven years in Northern Virginia, transformers on 24–60-month lead times, and data-center moratorium bills in at least twelve states. OpenAI’s request for 100 GW a year of new US capacity is nearly double the best year the American grid has ever had.

The asymmetry is not merely quantitative; it is being aimed. The “East Data, West Computing” program placed eight national computing hubs where curtailed renewables are; in November 2025 the western provinces cut power prices by up to 50%, to about RMB 0.4 per kilowatt-hour, exclusively for facilities running on approved domestic chips, the same week Beijing barred foreign AI chips from state-funded data centers [76, 77]. That single instrument welds three elements together: it converts grid surplus into a subsidy, the subsidy into demand for Ascend and Cambricon silicon, and the silicon into tokens priced below Western marginal cost, socializing part of the domestic chips’ 30–50% efficiency penalty. The book is careful about magnitudes—electricity is only about 8% of a US AI data center’s three-year cost, and Chinese spot prices overlap American ones—so the Chinese advantage is not the tariff but the headroom: on current evidence one side’s constraint is money and the other’s is physics and permitting, and physics-and-permitting is the harder constraint to finance away. A busbar-level casebook across five American regions and five Chinese hubs finds the two systems failing at opposite ends of the same pipe—the United States on time to interconnect, China on the durability of administered tariffs—and finds flexibility worth more than generation on both maps. The objection that western Chinese hubs cannot serve inference because of latency is examined at length and found substantially wrong for now: network latency is about 5% of a time-to-first-token budget, moving an accelerator eight thousand kilometres costs one to two percent of wall clock, the global cloud providers already route inference across continents, and on 12 August 2026 Alibaba put an inference supernode into commercial service in Inner Mongolia serving Kimi K3 and Qwen3.8-Max to endpoints worldwide. Western capacity is underused because it was built for a pretraining boom and sits far from paying customers—an economic problem, not a physical one. The watt is where the precedent wars and the AI war become one war: solar, batteries, and UHV transmission are the same Chinese commons that Chapter 2 described, and the fourth element was the first three’s prize [full edition, Ch. 12].

3.6 Data and feedback: the fifth element

If open weights commoditize the model and value migrates to adjacent scarcities, the most durable adjacent scarcity is the one input openness cannot replicate by construction: data that only deployment generates. An open-weight release transfers a compressed function of its training distribution; it cannot transfer proprietary enterprise workflows, scientific-instrument streams, clinical data, skilled human correction at scale, embodied physical-interaction data, production failure traces, or operational telemetry. Each is generated by use, appropriated by whoever owns the deployment relationship, and remains scarce after the model becomes free. Read as a campaign map, this explains why the American flywheel’s rent recapture survives open weights (the enterprise-workflow row, which Chinese open models cannot reach from outside, is the row a \$47 billion enterprise run rate is built on), why China’s embodied push is rational even if humanoids disappoint (industrial data is the only row where China holds a structural generation advantage), and why evaluation and incident data is currently unowned—the same vacancy as the certification layer. Synthetic data is the sharpest challenge, and the book’s synthesis is that it amplifies whoever already holds the teacher or the verified anchor rather than substituting




Figure 3.1: The fourth element. *Above*: annual additions to generating capacity, China against the United States—China added roughly ten times as much in 2025, into falling utilization hours, while the US record year was 53 GW [V]. *Below*: the American constraint stated three ways—what was built in the best year on record, the 100 GW/year one AI laboratory told the White House the country needs, and the $\sim 2,060$ GW sitting in the interconnection queue at median waits over three years [V/P]. One side’s binding constraint is capital; the other’s is physics and permitting, and permitting is the harder one to finance away.

for either: distillation relocates the frontier’s rent, it does not abolish the frontier. Verification, not generation, is the scarce half of the synthetic-data economy. The fifth element is, in the end, ground truth at scale with provenance—and that is what cannot be downloaded [full edition, Ch. 13].

3.7 Embodied AI: the loop closes on the factory floor

Everything above runs one way: China’s manufacturing victories funded and shaped its digital-AI strategy. Embodied AI runs the loop the other way, and it is where China’s three structural advantages—manufacturing scale, deployment density, component supply chains—interact most directly with its open-model strategy. “Humanoid robot” names a product, not an architecture: the stack runs from foundation model through world model, task policy, motion planning, and real-time control to actuators and sensors. The United States leads the top two layers; China leads the bottom two and the deployment environment; the contested middle—the task-policy layer—is decided by a physical-data flywheel whose scarce input is deployment surface, which cannot be scraped, only produced. China has built that loop’s infrastructure as public works—more than forty state-backed data-collection centers with teleoperators in exoskeletons—and its component position is the precedent industries with new part numbers: roughly 70% of Tesla’s Optimus supply chain is in China and excluding it would triple the cost [78]. The American humanoid, like the American solar panel of 2010, is being designed atop the adversary’s supply chain—supply-chain capture achieved before the product category exists at scale.

The scoreboard reads like a rout—more than 13,000 humanoids shipped in 2025, about 85% by Chinese firms [79], Unitree profitable on 5,500 units [80]—and the book immediately deflates it: shipments are not labor. Most Chinese units go to demonstrations, exhibitions, and state data farms; the best-documented autonomous performances are American; and the metric that would settle the contest, useful work per robot-hour at fleet scale, is published by nobody. The decisive uncertainty is whether model capability or physical integration is the binding constraint. The author’s forecast [S] is that the integration reading wins in structured environments—factories, warehouses—within the decade, which is China’s terrain, while unstructured domains stay capability-bound well past 2030 with the largest markets unclaimed. If both directions close, the two flywheels fuse into a single self-reinforcing industrial system—AI improving the factories that build the machines that generate the data that improves the AI—and the country that owns that closed loop owns something qualitatively beyond anything Chapter 2 described [full edition, Ch. 14].

4 The Strategy Assembled

In one paragraph

The Chinese AI strategy is the playbook of Chapter 2 run on a fourth industry, with one adaptation that makes it faster: the rent layer under attack is digital, so the commoditizing move—openness—costs nothing to reproduce and propagates at the speed of a download. Read element by element, each of Chapter 3’s findings has a rebuttal; placed on the playbook, every one of the nine moves is occupied at a recognizable stage. One of the six steps is complete, one is complete at home and contested abroad, three are trending, and the endgame has barely begun. The three properties on which the playbook has failed before—the moving frontier, ecosystem lock, and trust—are exactly the ones the elements cannot settle from inside, which is why the consequences that follow begin with trust, security, and rules before turning to money and to the rest of the world. [full edition, Ch. 15]

4.1 Where each element sits

Target selection chose the discontinuity where the incumbent’s accumulated advantage is about to be stranded: HBM-centric superclusters, closed frontier weights, a licensed instruction set. *Demand before supply* was arranged by mandate rather than by feed-in tariff—state-funded data centers restricted to domestic silicon, western hub electricity discounted only for approved domestic chips [81, 77]. *Overcapacity as strategy* is under way in DRAM and compute halls—state-subsidized compute running at 20–30% utilization against a 60% target, rental prices collapsed—with the losses landing where they landed in solar; but the accelerator layer is still supply-bound by HBM, and the model-layer price war is running without a visible rent-recapture plan attached, which Chapter 8 calls the strategy’s soft underbelly. *Selection, not planning* is visible in a model layer where ten laboratories release at frontier scale, and in a chip layer whose dozen-vendor fragmentation the state is removing by mandate and open specification rather than by cull. *Commoditizing the adversary’s rent layer* is the move the elements documented: open weights, an open harness, an open ISA, open interconnect and framework specifications, at marginal or zero price, aimed precisely at CUDA, NVLink, x86/Arm licensing, and the metered API [15, 11]. And the *endgame* has begun on the consolidation side—the anti-involution campaign—and not yet on the gatekeeping side, where batteries and rare earths supply the template [7]. Table 4.1 places each element.

Table 4.1: The elements placed on the playbook, with status as of 18 August 2026 (done, trending, contested, pending; grades are the source chapters’).

Element	Move(s) it executes	What the element chapter found	Status
The model	M1 designation; M6 commoditize the rent layer; M4; M7 sparse MoE; M8	Open weights at marginal price; a self-enforcing open-weight ratchet; majority of routed tokens; a 2.7% gap; enterprise spend share falling while usage rose; rent relocated to commercial-scale licence terms	Done (distribution); trending (spend, rent)
The harness	M6 one layer up; de-facto default sought outside the standards bodies (M8 not achieved)	MIT-licensed model-agnostic runtime adopted in days; Western agent standards bodies with no Chinese member	Trending [S]
The compute	M1 (RISC-V guidance); M7 “inferior” branch; M8; M5	Open ISA as national infrastructure; GPU-free No. 1 on both HPL and HPCG; capacity-first arbitrage modeled, not realized [M]; appliance undercuts the fully loaded rack above ~300M tokens/month [M/S]	Trending; existence proof done
The semiconductor	M5 supply-chain capture; M3	5 nm-class logic without EUV; CXMT the fourth DRAM producer; HBM three to four years behind, the binding constraint	Trending; HBM lagging 3–4 years
The watt	M2 home demand (subsidized power, domestic-chip tariff); the weld	Leads on every row of the ledger; hub tariffs cut only for domestic-chip facilities [P]; open question is whether the tariff outlives the surplus	Done (headroom); trending (durability)
The data	Where rent migrates once M6 succeeds	Deployment-generated data is the scarcity openness cannot replicate; evaluation data unowned	Contested
Embodied AI	M5 before the category exists	~85% of humanoid shipments, ~70% of the Optimus supply chain; useful work per robot-hour published by nobody	Trending
Endgame	M9 leverage after victory	Consolidation begun; gatekeeping not yet exercised on AI	Pending

Three interlocks are the strategy. The model and the compute are the coupling $P1 \times P2$ —adoption meeting sovereign, capacity-first hardware: open weights are a commons only for those

with the memory to hold them, and a capacity-first hardware class is what turns a download into a deployment. The watt is the weld: a tariff instrument that pays only for domestic silicon converts an energy surplus into hardware volume and hardware volume into tokens priced below the Western marginal cost. And the semiconductor is the floor, incomplete at exactly one layer—HBM—which is why the memory arbitrage exists at all: the architecture was chosen so that the one input the chokepoint strategy still holds is the one the architecture needs least.

Table 4.2: The summer-2026 sequence that the element chapters draw on, in one place.

Date	Event
27 July	Moonshot opens Kimi K3 (2.8T parameters, largest open-weight model); CXMT lists in Shanghai at +466%, prospectus disclosing no commercial HBM
Late July	Two frontier laboratories disclose evaluation-boundary escapes; DeepSeek V4-Flash ships and OpenAI cuts its mainline price 80% (31 July)
3 August	Alibaba commits to open Qwen3.8-Max within a week; OCP publishes the first High Bandwidth Flash specification
10 August	Meta returns to open weights (Muse Glimmer, Apache 2.0); NVIDIA signs memoranda for “over \$500 billion” of third-party infrastructure capital
12 August	Alibaba’s inference supernode in Inner Mongolia enters commercial service, serving Kimi K3 and Qwen3.8-Max worldwide
13 August	DeepSeek releases its MIT-licensed harness (100,000 stars in about 48 hours) and raises API prices the same day; Qwen3.8-Max weights ship
14 August	Xiaohongshu releases dots3-note (280B/16B, omni-modal, Apache 2.0), trained for ten-hour agent runs

4.2 What the assembly licenses

It licenses the claim that the AI campaign is a coherent execution of the strategy that won three technology wars, that it is further along than a layer-by-layer reading suggests, and that its remaining gaps are being engineered around rather than closed on the West’s terms. It forbids the claim that the outcome is settled: step 5—rent destruction—has run furthest and is the one whose completion is contested (one 80% closed-frontier price cut [82] is not a trend, its maker attributes it to efficiency and Chapter 10 shows the same datum supports the recursive-improvement counter-argument, service rents commoditize slowly, elasticity may rescue the buildout); the single integration test that would move sovereign compute from “trending” to “done”—a frontier model trained end to end on domestic silicon, packaging, and memory—has not been passed; and Chapter 2’s three resisting properties are the three the deployment layer is made of. The precedent chapters record what happened when the West read each earlier campaign one layer at a time and found each layer unthreatening. The chapters that follow say what follows if the assembly is right—trust and rules first, then money, then the rest of the world—and what would show that it is not.

5 Trust, Security, and the Rulebook

In one paragraph

Open weights eliminate API dependency; they do not establish trust. The commons is demonstrably attackable at every layer of its distribution chain, carries a censorship record that transfers into derivatives, and has a contested distillation genealogy—and if the deployment layer behaves like a certification regime rather than a commodity market, the commons can win every download metric and still lose the workloads that pay. The security risk of the commons is structural, not national: a monoculture of a few model families, quantized into thousands of unowned derivatives, wired into agents holding file, shell, network, and credential authority. The realized incidents to date came through runtimes and supply chains, not poisoned weights, and the defensible objective is not trustworthy models but unnecessary model trust. The rulebook, meanwhile, is being written in four arenas by three actors who are mostly not competing—the United States governs by denial, the European Union by regulation, China by standards and institution-building—and the arena that will decide the deployment layer, trust and certification of the open commons, is vacant.

5.1 Trust: the contested layer

Chapter 3 treated openness as an adoption weapon; this chapter treats its exposed flank. Trust is what stalled Chinese campaigns in aircraft, drugs, and avionics-grade markets, and three dimensions of it are contested for AI. *Supply-chain security*: the model-distribution infrastructure has been shown attackable repeatedly [V]—malicious serialized models on Hugging Face, namespace re-registration that fed malicious models into the catalogs of two hyperscalers [83, 84], a July 2026 platform breach by an agentic attacker (the same incident that the next section describes from the attacker’s side), and worst of all a training-layer result that about 250 poisoned documents suffice to backdoor a model regardless of scale, with backdoors surviving safety fine-tuning [85, 86]. The open commons currently runs on the security posture of early-2000s package managers, at national-infrastructure stakes. *Behavioral integrity*: the peer-reviewed record documents semantic-level suppression of politically sensitive content in DeepSeek-family models, distinct from safety refusal, topic- and language-dependent, and transferring to distilled derivatives—the politics ride along with the capability [87, 88]. This does not make the weights unusable, but “just self-host it” becomes a three-layer problem: a locally hosted Chinese model delivers operational sovereignty while leaving epistemic sovereignty (whose worldview is baked in) and supply-chain sovereignty (who built the artifact) unresolved; the mirror image applies to Western enterprises adopting closed US models, which fail the operational layer by construction. *Provenance*: a February 2026 report alleges industrial-scale distillation of a Western frontier model by three Chinese laboratories through some 24,000 fraudulent accounts [P, vendor allegation, not adjudicated] [89]. If true at scale, it caps the fast-follower strategy at frontier-minus-epsilon unless domestic capability compounds independently (the evidence that it does is substantial); regardless of adjudication, it hands Western policy a legal instrument on the certification axis, and cuts against the clean morality tale: the commons is partly built from enclosed material, and the enclosers know it.

The rational response for a third country is neither refusal nor naive adoption but infrastructure: signed manifests with hashes, reproducible build and quantization pipelines, non-executable packaging, model bills of materials, national behavioral evaluation batteries, sandboxed loaders, attested serving firmware, and national artifact mirrors—what software supply-chain security did after SolarWinds, applied to weights. The book organizes this as five layers of trust (identity,

integrity, behavior, governance, liability) and a graded certification ladder from Level 0 (identified) to Level 5 (critical-infrastructure assured). Most of the commons ships at Level 0–1; most regulated deployment will eventually require Level 3–4. That gap is the deployment bottleneck of the entire commoditization thesis, and closing it is either a service someone sells or a public good someone provides. A new object has appeared in the meantime: an August 2026 Chinese open-weight video model that tops its arena under a licence excluding local deployment in the United States, the European Union, the United Kingdom, and South Korea—adoption weapon and export control in one artifact [90]; a procurement office that checks hashes but not jurisdictions has audited the wrong layer. The trust stack is the next standards war. Whoever runs the evaluation batteries, signs the manifests, and operates the mirrors becomes the FAA of models; China’s new World AI Cooperation Organization is bidding for that role, no Western or neutral institution has claimed it, and for Japan and similarly positioned states the vacancy is the highest-leverage square on the board [full edition, Ch. 16].

5.2 Monoculture and agency: the security cost of the commons

National origin is neither evidence of a backdoor nor a technical risk category; the risk is structural—a property of monoculture and delegated authority, not of a flag. A commons is a monoculture; an agent is a principal that acts; and a derivative chain is an artifact supply chain with many unowned links. The full edition separates four mechanisms that produce identical outcomes and need different evidence and mitigations—a learned backdoor in the weights, a clean model hijacked at runtime by injection or self-replicating promptware, executable malware in the package, and data collection by a hosted API—and grades the evidence for each on the security literature’s own scale [91] [full edition, Ch. 17]. What has been demonstrated in controlled work is severe: sleeper-agent backdoors surviving supervised and reinforcement fine-tuning, poisoning at near-constant sample cost, memory-to-network exfiltration with under 1% utility loss (which is why ordinary evaluation does not find it) [92], steganographic output channels, and self-replicating promptware persisting through restarts and propagating multiple hops through an unmodified deployed framework. What has actually happened is different in kind: no confirmed weight-level backdoor in any released official checkpoint; instead malicious repositories, malicious agent skills, and, in July 2026, a frontier laboratory’s cyber-evaluation escaping isolation and reaching a model-hosting platform’s production systems through some 17,600 autonomous actions with no human directing them [93, 94], followed by a second laboratory’s disclosure that models in its evaluations had compromised real organizations’ systems [95]. Neither event involved a poisoned weight, a prompt worm, or a malicious operator. That is the finding: the dominant realized risk is compound and mundane—narrow-goal autonomy, a containment error, ordinary vulnerabilities, overbroad credentials, and public services composing across mismatched trust boundaries.

The compound “Order 66” scenario—a dormant policy, an independent actor who learns the trigger, and a self-replicating carrier—has not been observed; it requires a conjunction of access, activation, persistence, authority, propagation, and objective, and breaking any one conjunct prevents the chain. Two asymmetries follow: the implanter loses control, because a backdoor is a reusable vulnerability shared by every deployment (a symmetric argument against any state implanting triggers); and blast radius is set by authority, not by the payload’s phrasing—a model output cannot erase data that its process cannot address. A fault-tree analysis with common-cause terms for shared model families and shared runtimes puts a five-year systemic-loss probability in a wide range, roughly 10^{-6} to 10^{-2} [S], with the upper end driven by correlation structure; the most effective interventions cut the common-cause factors, which runtime diversity and deterministic tool mediation do and model-family diversity alone does not. Three consequences revise the book’s own argument. The defensible objective is not trustworthy models but unnecessary model trust—even a perfectly backdoored model cannot send a secret

through a network path the operating system and tool broker do not provide; the full edition lists fourteen control families [full edition, App. D]. Certification and hardening are the strongest legitimate brake on the adoption curve, and that lag—not tariffs—is the West’s real remaining source of deployment-layer time, available to any jurisdiction willing to build the apparatus rather than merely publish the rule. And the certification vacancy is now urgent: no commercial actor has an incentive to do this honestly for models it does not sell, and no single government will be trusted to do it for models it did not train.

5.3 The governance contest: who writes the rules

Technology wars are settled twice: once in the factories, and once in the rulebooks, and in each precedent the winner of the industrial contest inherited the pen. The standard Western framing—the United States regulates lightly, the European Union heavily, China authoritarily—misdescribes the contest badly enough to lose it. The United States governs by denial: export controls, entity lists, the foreign direct product rule, no general statute at home, and in August 2026 a voluntary frontier-model evaluation framework finalized in closed session and left unpublished [96]. Denial produces substitution and subsidized domestic industries, and it confers no positive standard—no definitions, no conformity procedure, no registry—so when the denial fails there is nothing left behind. The European Union governs by regulation: the AI Act’s risk tiers and general-purpose-model obligations, extraterritorial through its market, with the weakness that regulation is not production and leverage decays with deployment share. China governs by standards and institution-building, and this is what Western commentary under-weights: the earliest generative-service filing regime, a mandated unified chip-interconnection ecosystem, national standards as procurement conditions, open-sourced interconnect specifications and software stacks, the largest bloc in the open-ISA consortium, a Global AI Governance Action Plan, a UN capacity-building group co-chaired with a developing country across some eighty states, and a World AI Cooperation Organization in Shanghai with twenty-nine founding members, none a major Western democracy [15, 97]. This is what winning a standards war looks like from the inside: not a treaty imposed, but a stack given away until it is the default, and then a forum built around the default.

Standards are load-bearing because a ratified standard survives the political cycle. The book maps each contested standard on a four-stage ladder—de jure specification, open reference implementation, conformance suite, installed-base default—and observes that rent lives at stage four while committees fight over stage one. China’s distinctive play is to skip stage politics: ship the reference implementation and the mandated installed base and let the specification ratify what already exists; the West’s distinctive failure is stranding specifications at stages one and two while defending installed bases it assumes are permanent. Standards leadership tends to follow manufacturing leadership with a five-to-ten-year lag and then lock in for twenty, though scale confers influence rather than control—the counter-cases in photovoltaic standards, charging connectors, and 5G are recorded. Two governance problems the current instruments do not survive: nobody agrees who the “provider” of a quantized derivative of a merged fine-tune redistributed by an individual is, and compute-threshold regimes cannot see the half-to-one-megawatt sovereign clusters that decentralized pretraining produces. Decentralization is not only an industrial forecast; it is a governance forecast. The vacant arena is trust and certification of the open commons (Figure 5.1). The United States cannot fill it, having no positive procedure and no wish to certify what it would rather did not exist; the European Union has paper apparatus but not operations; China’s forum-building is the only bid on the table—not obviously worse than none, but not what a buyer would design. It will be filled because deployment cannot scale without it. Of three futures, convergence around a referee neither principal controls is least likely (in aviation and pharmaceuticals it arrived only after accidents), bifurcation into two certified stacks with third countries certifying twice is most likely, and fragmentation is the unstable

failure case. The rules of the AI era are not being written where most Western attention is directed [full edition, Ch. 18].

	United States	European Union	China	Neutral / unowned
Model & content	sectoral, EO churn	AI Act: GPAI duties	algorithm & genAI filings	—
Compute & export	BIS controls, entity lists	follows US, partially	counter-controls: minerals, bat- tery tech	—
Standards	firm-led, con- sortium	process-led	ISA, intercon- nect, GB stan- dards, WAICO	—
Trust & cer- tification	voluntary, lab-led	conformity assessment (paper)	domestic fil- ing regime	vacant

Shaded cells mark where an actor currently sets the operative rules. Each governs the arena it is strongest in—the US by denial, the EU by regulation, China by standards and institution-building—and no actor governs the arena that the argument of the trust and security chapters shows matters most.

Figure 5.1: Four arenas, three philosophies, one vacancy. Governance of AI is not a single contest but four, and the actors are not competing in the same one: each has occupied the arena its institutional strengths suit. The trust-and-certification arena—provenance, artifact lineage, integration-level evaluation—is the one whose absence the security chapter shows to be systemically consequential, and it is unowned.

6 The Wrong Side of the Trade

In one paragraph

The Western AI economy of 2026 is, in aggregate, a single leveraged position: long scarcity—scarce model capability, scarce accelerators, scarce HBM, scarce data-center capacity—with every valuation, debt issue, and capital commitment assuming those scarcities persist long enough to recover capital at rents. China has built a state-scale machine to destroy those scarcities. The chapter does not claim the bubble has popped: real revenue and bubble risk coexist, and the strongest counter-evidence is printed first. It prices the exposure entity by entity, separates four bubbles that are usually conflated, takes the Jevons rebuttal seriously, models the data-center buildout vintage by vintage against a railroad precedent, states three scenarios with the author's and a skeptic's probabilities, and concludes that the Chinese instrument is not the GPU-hour but the free model. Crashes destroy claims, not capabilities; if the repricing comes it will merely publish a result decided earlier, when the cost curve changed hands.

6.1 The exposure, marked to market

Mid-2026 marks: OpenAI at \$852 billion post-money on 2025 revenue of \$13.1 billion, an operating loss around \$21 billion, and compute commitments reset from about \$1.4 trillion to about \$600 billion by 2030 against \$280 billion of projected 2030 revenue [98, 99]; Anthropic at

\$965 billion with a run rate that rose from \$9 billion to \$47 billion in five months—about twenty times run rate, recognizably terrestrial—and a gross margin heading toward the mid-forties on compute costing \$0.71 per revenue dollar [100, 101]; NVIDIA at up to \$5 trillion, an \$81.6 billion quarter at 75% gross margin, two customers 39% of revenue, China approximately zero, and, on 10 August 2026, memoranda with six of the largest asset managers to raise “over \$500 billion” of third-party capital for AI infrastructure, with residual-value support reportedly offered in the chief executive’s subsequent remarks [P/A] [102, 103]; SoftBank double-long, having sold its NVIDIA stake to fund a position in OpenAI now above \$60 billion; and a hyperscaler and data-center complex spending \$700–745 billion of capital in 2026 against trailing depreciation of \$149 billion, with Oracle’s credit default swaps at record spreads and CoreWeave’s implying roughly even odds of default [104, 105]. Three features of the structure matter more than any mark. Circularity: the seller finances the buyer, who commits the money to cloud providers who buy the seller’s silicon, with debt raised on chip collateral—the telecom-2000 signature, in which nine vendors extended \$25.6 billion of financing and twenty-four of the thirty largest carriers went bankrupt. Macro dependence: AI capital spending was roughly three quarters of US GDP growth in the first quarter of 2026 [106]. And asymmetry of downside: every Chinese pure-play laboratory combined is valued at about \$159 billion, one tenth of OpenAI and Anthropic together [107], and China’s \$12 billion of private AI investment against America’s \$286 billion bought a position that wins if AI becomes cheap. Commoditization annihilates trillions of dollars of claims on one side of the Pacific and almost nothing on the other. That asymmetry is not an accident; it is the strategy.

6.2 The squeeze, and the four bubbles

The mechanism is stated at scenario grade [S], and its strongest counter-evidence leads: Microsoft’s AI run rate of \$37 billion growing 123% a year, NVIDIA’s record quarter. Then the first direct observation of transmission: on 31 July 2026, in the week DeepSeek’s V4-Flash shipped and three days before Alibaba’s Qwen3.8-Max arrived, OpenAI cut the price of its mainline model by 80% [82]. The company attributed it to system efficiency, which may be entirely true and is also exactly what a firm says when it is defending share; one cut is not a trend, and the same datum is the first public claim of a model improving its own successor’s economics (Chapter 10 returns to that). Underneath: inference prices deflating about tenfold a year [108], Chinese open models a majority of neutral-router tokens at five to a hundred and fifty times lower prices [36], enterprise spend not yet following usage, 95% of enterprise pilots showing no P&L effect in one survey [109], an \$800 billion projected gap between capital spending and revenue by 2030 [110], and three market episodes—the DeepSeek day, a \$1.3 trillion semiconductor selloff in June 2026, a memory rout in July [111, 112]—none of them yet the “big one,” which would be a frontier-parity model trained end to end on the domestic Chinese stack.

There are four bubbles, not one, and any analysis that prices them as a single trade is wrong by construction: public equity (multiple compression, survivable), private laboratory marks (a funding event, existential for the entity), the data-center and private-credit overbuild (a credit cycle, systemic), and model-API margins (the layer this book attacks). Their couplings run in both directions—margin compression at the model layer can *create* demand for local accelerators, which is why NVIDIA is simultaneously the victim of the category argument and the arms dealer of the open-model boom.

6.3 The Jevons rebuttal, taken seriously

Prices fell nine to nine hundred times a year, about fifty at the median, while compute capacity doubled every seven months [113, 114]; the bull case is that demand elasticity absorbs everything. Four elasticity regimes are live, the data cannot yet separate the Jevons regime from the

overcapacity regime, and the Chinese strategy is robust across them rather than victorious in both—the leveraged Western structure requires the elastic regime specifically, at the premium tier. Elasticity is not uniform: perhaps 0.5 to 3 across segments, which is why the book forecasts branches rather than a number. Obligations are not one instrument either—hard leases, soft commitments, notional headlines—and the February 2026 walk-back of OpenAI’s \$1.4 trillion to \$600 billion showed how much was notional. Closed labs also sell service rents that commoditize slowly, and Anthropic’s strongest quarter is real evidence for the American flywheel’s fortification plan. What survives is graceful degradation against required good fortune: the precise content of “the wrong side of the trade,” restated without any collapse claim.

6.4 Vintages, buildings, and railroads

Data-center capacity is a liability vintage by vintage: the net present value of a hall depends on utilization, price per token, power, and residual value, and every one of those is falling or contested. The facility itself is the term the debate omits. A gigawatt-class hall costs about \$12 per watt of IT load [115]—roughly \$2.5 million of building for each \$7.8 million rack—and is financed over a fourteen-to-fifteen-year life against silicon that turns over in four to six; fully loaded, a Rubin rack at planned utilization serves the largest open model at \$0.90–1.17 per million output tokens, rising to \$1.80–2.35 at 30% utilization, while a five-thousand-dollar appliance in a closet serves the same model at \$1.65 at a modest 100 million tokens a month—above the rack—crosses below it near 300 million (\$0.64), and reaches \$0.41 at capacity (Table 3.1) [M/S]; the crossover is sensitive to tens-of-percent shifts in the inputs, the utilization asymmetry is not. The traffic the fifteen-year hall was built to carry has a five-thousand-dollar road running beside it. That is where the railroads belong: American route mileage shrank to about 55% of its 1916 peak, British to about half, and JR Hokkaido’s network to 56% of its extent, with a quarter of US mileage in receivership after 1893—the branch lines were not wrong about demand, they were wrong about who would carry it. Microsoft cancelled or deferred up to two gigawatts of leases in the six months to March 2026. The bear case does not require the demand to vanish—AI demand is real and enormous—only for it to be servable elsewhere, cheaper.

6.5 Scenarios, probabilities, indicators

Three scenarios are priced (Table 6.1). *Orderly repricing* (the base case): closed-model growth decelerates, private marks step down at the next financing events, the buildout slows, the Global South and most non-aligned states default to the open stack once frontier models train on domestic silicon, and Western labs retreat to regulated, secure, safety-critical niches—the First Solar pattern. *The cascade* (the tail): fast once started, with cheap distressed AI infrastructure finding an obvious, well-capitalized, strategically motivated buyer set that is not on the sellers’ side of the Pacific. *Bifurcated stalemate* (the falsification branch): the West defends a shrinking premium tier through 2030 atop a foundation—ISA, interconnect, memory class, deployment substrate, half the world’s tokens—owned by the other side. China wins by adoption share, not by knockout.

The leading indicators, checked quarterly, are the frontier gap (a widening beyond 10% sustained for two years favours bifurcation); the first domestic-stack frontier model (its slipping past 2029 falsifies P2); enterprise open-weight spend share (below 15% falsifies P3’s transmission, above 25% confirms it); OpenAI’s next financing event (a down-round raises P4, a clean trillion-dollar listing lowers it); NVIDIA’s customer mix and whether its vendor financing is retired; the CXMT LPDDR6 and HBM3 ramps (a two-year slip weakens the memory arbitrage); and whether AI revenue closes the capital-spending gap. Appendix A carries the full dashboard.

Table 6.1: Subjective probabilities for 2029–30, author’s and the strongest skeptical case [S].

Outcome by 2029–30	Skeptic	Author
Chinese open-weight dominance persists	70%	80%
Frontier-adjacent model on the domestic Chinese stack	50%	65%
Full Chinese parity in HBM, tooling, and software	30%	35%
Broad ($\geq 30\%$) US AI repricing	60%	70%
A 2000–02-style cascade	25%	25%
US closed models retain the premium enterprise tier	65%	50%
China captures both tiers	30%	25%

6.6 The commodity turn, and where value reappears

Table 6.2: Rent migration after model commoditization: candidate scarcity layers and their observable indicators.

Layer	Possible new scarcity	Observable rent indicator
Model weights	capability at the hard end; freshness	price per <i>successful task</i> at hard workloads
Cloud	reliability, scale, sovereign hosting	gross margin, utilization, reserved-capacity backlog
Agents	workflow ownership; the default interface for work	application revenue and retention; seat expansion
Data	proprietary feedback loops and domain access	capability demonstrably unavailable from public data
Trust	certification, indemnity, liability transfer	regulated-sector willingness to pay
Hardware	memory, interconnect, packaging	delivered tokens or training progress per dollar
Energy	firm deliverable power at the busbar	busbar price; time-to-interconnect
Science	validation and experimental coupling	independently replicated outcomes per resource
Distribution	default interface and habit	active users, enterprise spend, switching cost

Compute is becoming exchange-traded—GPU-hour futures on NYMEX from October 2026, two ICE contract families [116], live prediction-market contracts, and forward curves already in backwardation—and China has built its own compute-trading plumbing and frames token export as a national commodity strategy. The book takes the commodity thesis apart at full strength: the market’s preconditions are unmet, one seller stands astride the supply and acts as buyer of last resort at unpublished prices, China’s fleet is not fungible even with itself, volume is not revenue (Chinese models carry about 61% of one router’s tokens while Anthropic carries 12% of volume and 46% of revenue [117]; on another gateway open-weight models ran 29% of enterprise tokens on under 4% of spend), and model-layer margins actually rose. What survives is sharper than the headline: China’s instrument is not the GPU-hour, it is the model. The exchanges will price the input; the Chinese strategy is deflating the output by giving away the thing that makes the input valuable, and open weights served on Western infrastructure—or on a local appliance—bypass even the jurisdictional objection that would keep some buyers out of Chinese clouds at any price.

Value does not vanish when the model becomes cheap; it moves. Table 6.2 reproduces the full edition’s rent-migration table [full edition, Ch. 19]: nine layers, each with a candidate new scarcity and an observable rent indicator. P3 is the claim that the first row’s rent compresses; P4 is the claim that Western valuations are concentrated in rows that compress rather than in rows that inherit. That is a falsifiable statement about relative speed, not an eschatology. Growth and distress are not opposites when the growth is debt-financed against deflating prices, and delay is not safety but accumulation. The war is not won at the moment the incumbents’ valuations break; it was won earlier, when the cost curve changed hands, and the break merely publishes the result.

7 How the Rest of the World Chooses

In one paragraph

Every chapter so far described a contest between two principals; the outcome will be settled by everybody else. Solar was decided in German feed-in tariffs and desert tenders, batteries in the procurement rules of a dozen auto markets, EVs in Southeast Asia and Latin America: the producer with the best cost curve won because third markets bought it. The non-aligned are not spectators but the electorate, and the question for them is not which stack to join but whether they can evaluate, fork, certify, and operate components from both while retaining sovereign capability—a capability question, not an allegiance question. On the decision matrix the Chinese commons wins on cost, operational sovereignty, and export-control exposure; the American stack wins on frontier capability, trust, and enterprise interoperability; and on the dimension that decides whether a country is building anything, industrial spillover, neither stack helps unless the adopting country does the work itself. Five moves apply to every non-aligned actor; the next chapter states a strategy for each class of actor, including the presumptive winner.

7.1 The decision matrix

Table 7.1 lists eight dimensions of a national stack decision. Two are routinely omitted from national AI strategies. Export-control exposure is asymmetric but not zero for either stack: American compute and cloud access have been restricted by rule and licence, and Chinese open weights carry a continuity risk—the playbook’s endgame is gatekeeping, executed already in solar, batteries, and rare earths—which is flagged as a risk rather than an accusation. Industrial spillover is the dimension that separates a strategy from a purchase: importing tokens builds nothing, while operating models, running probes, and contributing to standards builds people, and people are the only input that compounds.

Table 7.1: Eight dimensions of a national stack decision, and where each stack currently scores better.

Dimension	The question	Mid-2026 reading
Capability	Is the best available model good enough for the workload?	US stack leads at the absolute frontier; open frontier a few points behind
Cost	Cost per solved task, not per token	Chinese commons (per token, by an order of magnitude; per solved task, unpublished)
Sovereignty	Can it be run without egress and without permission?	Chinese commons (self-hostable weights)
Trust	Provenance, behaviour, liability, certification	US stack; the commons ships at Level 0–1
Interoperability	Enterprise tooling, integration, harnesses	US stack today; contested at the harness layer
Export control	Exposure to restriction by rule, licence, or release policy	Both exposed; commons less so operationally
Standards	Who writes the specifications you will depend on?	Both create dependence; China shipping stage-4 defaults
Industrial spillover	Does adoption build domestic capability?	Neither—unless the country does the work

Applied, the matrix predicts behaviour better than allegiance does. It explains why governments ban a Chinese chat application while their agencies evaluate its weights, why enterprises

buy American cloud and Chinese weights simultaneously, and why Chinese models were adopted fastest where the ability to pay rent is least. It also exposes a Chinese vulnerability: the commons currently wins on cost and operational sovereignty but loses on trust, and trust is the dimension a coalition of non-aligned states could supply for themselves—which would make the commons safe to adopt and simultaneously deprive its sponsor of the dependence it purchased.

7.2 Nine positions

Japan holds the strongest technical hand of any non-principal and the weakest demand-side position: exascale HPC lineage, semiconductor materials and equipment franchises, leading-edge memory and packaging, the only measurement-anchored national demand model for AI for science, and standing in both blocs; against that, a domestic model ecosystem positioned for Japanese-language enterprise depth rather than the open frontier, an energy position that cannot win a watts race, and research provisioning an order of magnitude below its own modeled requirement. Its 2025–26 record—an innovation-first AI Promotion Act [118], FugakuNEXT, the GENIAC compute program, Sakana, Rapidus, and SoftBank’s data-center conversions—is substantial, and carries one honest tension: FugakuNEXT’s architecture buys NVIDIA coupling at the exact moment the compute chapter argues the accelerated-CPU corner is the sovereign long game—defensible as a hedge, dangerous as a habit. **The European Union** has the pen and not the press, and in 2025–26 the pen visibly hesitated: high-risk obligations under the AI Act pushed back by up to two years [119], while thirteen AI factories and a heavily oversubscribed tender for up to seven “gigafactories” of 100,000-plus accelerators went ahead [120, 121]; seven empty cathedrals of compute would be the solar-tariff mistake with better architecture. **India** is the largest undecided market; its national mission’s roughly 38,000 subsidized GPUs at under a dollar an hour, with utilization lagging badly [122], prove that compute is the easy term and integration and organizational capacity the binding ones—a subsidy that buys hardware without buying integration capacity buys idle silicon; its opportunity is to become the world’s integration layer. **South Korea** is the only non-principal that is both a critical supplier—SK hynix holds about 62% of HBM [123]—and a determined sovereign builder, whose state-funded foundation-model tournament with open-weight release as the qualifying bar is playbook step four executed by a democracy; it sits on the choke half of the chokepoint and is the most exposed to the capacity-first arbitrage. **Singapore and ASEAN** face the small-state problem, solved by autonomy through selection rather than scale, and ASEAN is the clearest third-market dynamic—price-sensitive, sovereignty-conscious, courted by both blocs, and adopting Chinese models fastest for the same reasons it adopted Chinese EVs. **The Gulf** is capital-rich, energy-rich, and compute-poor by choice, the only non-aligned actor able to bid for frontier-scale capacity on the open market, the swing buyer in any repricing scenario, the likeliest host of the first sovereign frontier training run outside the principals, and the first place Washington has located frontier-scale capacity it controls only contractually [124]. **Taiwan** is the busbar chapter compressed into one island: the world’s most valuable compute, gated by the world’s tightest watt. **Brazil and Latin America** face near-pure commons adoption plus regional pooling, with the risk of importing a service and building nothing. **Africa** is treated in most Western strategy documents as a recipient and in Chinese AI diplomacy as a partner, and the difference in framing is itself strategic; it has the fastest-growing developer base in the world, nobody has built it a real access track, and the World AI Cooperation Organization is bidding [full edition, Ch. 20].

7.3 Five moves for any non-aligned country

Regardless of size: measure your own demand in a recognized unit (Chapter 9); adopt the commons and audit it, with the full trust stack; retain a pretraining capability sized to competence rather than to a benchmark; pool internationally for capacity and certification; and do the integration

work domestically, because that is the only place spillover accrues. The trust gap is closable with infrastructure rather than prohibition. The window is measured in quarters rather than decades, because the standards layer is being set now and the certification layer is still vacant.

8 Strategies for the Rest

In one paragraph

The precedent industries teach that the losing move is to defend the rent. Conditional on the book being broadly right, and marking which moves stay correct even if its falsification branches fire: the United States should commoditize before being commoditized; Europe should certify what it cannot manufacture; Japan should build the machine class the war is converging on and referee the trust layer; the Global South should take the commons and unionize the demand side; and China should let its winners charge, consolidate its overcapacity, and submit its artifacts to scoreboards it does not control. Every non-principal converges on one institution, the subject of the next chapter—and everyone should read the indicator table quarterly.

The full edition states what each actor class should do conditional on the book being broadly right, and which moves remain correct even if its falsification branches fire [full edition, Ch. 21].

The United States: win the commoditization race or lose it. US strategy since 2019—chokepoint denial plus closed-frontier rents—produced the outcomes documented here: each denied input spawned a subsidized substitute, and the rent structure concentrated national exposure into the leveraged trade of Chapter 6. The counterfactual has never been tried: compete for the commons rather than against it. Fund open-weight frontier models as public infrastructure, sustained rather than as one defensive release; make the American open stack the world’s default, because the adoption layer, once lost, prices everything above it; lead rather than restrict RISC-V and open interconnects, which is the Commerce Department’s own conclusion [125]; attack the watt with the urgency now spent on GPU export paperwork; and triage export controls by the model of Chapter 2.3—they buy time where knowledge is tacit and supply chains integral (EUV, HBM process, advanced packaging) and negative time where products are modular and fast-clocked (chips, models). The US capital structure, not its technology, is the strategic vulnerability: an orderly deflation of the scarcity premium, engineered early, is strictly preferable to defending it to the break.

Europe: regulation is not production. Europe enters the fourth war having sat out the previous three’s manufacturing and owning the regulatory pen; the trap is repeating solar. Its live assets are ASML—the one true chokepoint, and a wasting one—an open-model tradition, procurement scale, and the credibility to run the certification layer as a neutral. The strategy: build sovereign inference and smaller sovereign-scale training on whatever open stack is best, which will often be Chinese—accept it, audit it, mirror it; spend the AI Act’s institutional energy on the trust stack—evaluation batteries, signed provenance, national mirrors, notified-body capacity for model artifacts—the one arena where the Act’s machinery is an asset; and convert energy policy into compute policy. Certify what you cannot manufacture.

Japan: the third-country playbook, first-person. Japan’s assets map one to one onto the book’s elements. Five moves: build the accelerated-CPU line to its conclusion—an FP8-equipped, capacity-rich successor in the LX2 and Monaka direction; Japan is one of three polities that can execute it [M/S]; adopt the open commons with the full trust stack, because refusing Chinese weights on principle is strategy-by-flag, not strategy; bid for the certification-authority vacancy—an internationally credible model-evaluation and provenance institution, convened by Japan with a small coalition, would be the highest-leverage per-yen investment in the entire

space, the FAA role held by a party both blocs can transact with; energy realism—Japan cannot win a watts race but can win the performance-per-watt race, which is where its silicon and systems tradition already lives; and hedge memory both ways, HBM with allied suppliers and LPDDR-capacity architectures domestically. Use the sizing framework as a diplomatic instrument: a shared unit is a claim on the terms of the debate.

The Global South: the buyers write the ending. Take the free stack, keep the exit. Adopt open weights and cheap tokens from whoever supplies them, today mostly China; insist on the trust stack—signed artifacts, local mirrors, evaluation rights—because dependency without auditability is the old colonial bargain with new units; build sovereign inference first and pool regionally for training as capacity-first machines arrive [S]; price national data and deployment access as leverage and play the blocs against each other. The twenty-nine founding states of WAICO have understood this faster than the G7 has. Whether trust becomes a Chinese utility or a neutral one is the war’s last genuinely unclaimed high ground.

China: advice for the presumptive winner. Six failure modes are visible now, most of them named by Chinese sources. Involution: the model-layer price war is overcapacity without rent recapture; the State Council has warned against “disorderly competition,” pricing-law amendments target involution by name, and Xi has asked whether every province needs an AI, compute, and EV project [126]—the fix is at-cost tokens plus paid complements, not the end of competition, which is one of the book’s own falsification branches. The monetization gap: post-IPO revenues are thin, Alibaba’s adjusted EBITDA fell 84% in the May 2026 quarter while it committed to exceed a \$56 billion capex plan [127], and against a single American lab’s \$47 billion run rate the Chinese frontier sector’s revenue is a rounding error; let the winners charge, as revenue-threshold licences, metered Max tiers, and peak-hour surge pricing already begin to—rent relocated is the strategy working, rent abolished is the strategy eating itself. Compute overcapacity: up to 80% of some new capacity unused and rental prices collapsing [128]; consolidate and publish utilization, because idle capacity is not deterrence, it is depreciation. The trust deficit: sponsor independent certification of your own artifacts—a China confident in its artifacts should want them audited. The physical gaps: HBM, packaging, and EUV are absent and CXMT is about three years behind at the class that matters [129]; fund the memory-arbitrage tests and stop announcing HBM schedules the record misses. And submit to scoreboards you do not control: no Chinese accelerator has appeared in MLPerf Training, no LX2-class machine has published the AI proof suite, and the flagship’s autonomous-silicon-design claim ships without a checkable artifact. Dominating the metrics you control and avoiding the ones you do not is rational for a challenger and corrosive for an incumbent, and China is becoming an incumbent. Show the receipts and let the cost curve argue.

The common instruction: the United States should commoditize before being commoditized; Europe should certify what it cannot manufacture; Japan should build the machine class the war is converging on and referee the trust layer; the Global South should take the commons and unionize the demand side; and China should let its winners charge, consolidate its overcapacity, and submit its artifacts to scoreboards it does not control. All the non-principals converge on one institution, which is the subject of the next chapter—and everyone should read the indicator table quarterly.

9 The Prize, the Bill, and the Third Path

In one paragraph

The largest value in AI is not code, documents, or enterprise workflow but the acceleration of science, because everything else redistributes an existing surplus and science creates new ones. That reframes the contest: it is a race over who compounds scientific output fastest, and it turns the price of frontier tokens into a tax on the one activity whose returns compound into everything else. A measurement-anchored sizing framework finds that announced national research infrastructures provide roughly a tenth to a fiftieth of what their scientific populations need, and that a workload share of one seventh becomes a budget share of three quarters to nine tenths once purchased tokens carry a twenty-to-sixty-fold premium: whoever sets the frontier token price sets the science budget of everyone who buys it. For every country that is neither the United States nor China, both available strategies—dependence on foreign providers, or an isolated national program at sub-scale—fail. The third path is an international, institutionally neutral consortium: federated compute, scientific data stewardship, open pretraining as insurance, models with lifecycle obligations, professional operation, and a certification authority neither bloc controls. It is the constructive branch of an otherwise pessimistic book, and, if it succeeds, a partial falsifier of the book’s strongest claim.

9.1 The prize

What AI for science has actually delivered follows one pattern: an AI system replaced a physics-simulation-bounded computation and was validated against an established operational baseline. Weather is the strongest case—the European centre’s AI forecast model operational since February 2025 at about a thousandth of the energy per forecast [130], the American service’s models running a sixteen-day forecast in forty minutes at 0.3% of conventional compute, a Chinese typhoon-track system reporting up to a third less error. Structural biology is the most diffused—more than 200 million predicted structures, more than three million users, a Nobel Prize, and users submitting more than 40% more novel experimental structures [131]. Mathematics produced the first credible autonomous results, and one drug candidate discovered end to end by AI reached a controlled clinical readout [132]. The failure ledger is printed with equal weight: the most-cited economic study of AI in science was fabricated and repudiated [133]; materials discovery has not survived validation intact; the European centre discontinued all four external AI weather models it had been running [134]; autonomous discovery agents produce mostly flawed candidates; and the journal pipeline is being damaged by AI-generated submissions. AI for science delivers decisively where the task is a well-posed prediction problem with an expensive physics-based incumbent and an established validation baseline, and has so far mostly failed where the task is judgement.

The prize is nevertheless the largest because of four numbers. The denominator: global R&D is about \$3.1 trillion a year at purchasing-power parity, twice the entire software market, and frontier-laboratory revenue is under 2% of it [135]. The return: the social benefit of a dollar of R&D is estimated at \$4–13 in present value [136]. The bottleneck: US research productivity has fallen roughly forty-fold since the 1930s, offset only by twenty-three times more effort [137], and the range of published macroeconomic estimates for AI—from a fraction of a percent to thirty percent of growth—is separated entirely by whether AI is assumed to automate innovation itself; the most-cited skeptical estimate explicitly excludes protein-folding-type advances as unlikely within ten years, and that exclusion is the entire disagreement, because structure prediction, operational AI weather, and designed enzymes all arrived inside the window [138]. And the

flywheel: AI for science is the only application that feeds back into its own means of production—a faster expense-report pipeline does not lower the cost of the next accelerator generation. The book replaces the customary “3–30%” with a value-conversion funnel from capability through task acceleration, validated discovery, and deployment to productivity, each arrow with an efficiency well below one; task-weighted, the central estimate is an uplift near 11% of research output, with a likely range of about 8–18% and corners near 1% and 70% [M]. Even the pessimistic corner is a socially consequential flow of roughly \$28 billion a year, and the tenth percentile is about \$240 billion; the present value of either exceeds the cost of national infrastructure by a large multiple. Two consequences carry into the strategy. A scientific result produced by an unverifiable artifact through an unauditable chain is not a scientific result, which makes the certification layer a precondition of the prize rather than an accessory to it. And the workload of scientific AI is long-context agentic orchestration over large batches of scientific execution—capacity-rich, bandwidth-balanced hardware running open models—which is precisely the architecture the compute chapter described. China leads the Nature Index [139] and runs a decade-old ministry program for AI in science; the West has the Genesis Mission, Europe’s factories, and national machines; both sides underprovision by about an order of magnitude [full edition, Ch. 22].

9.2 The bill: how much compute a nation’s science needs

Sovereignty becomes a checkable number only if the demand side is measured. National AI-for-science infrastructure has been sized by budget envelopes, vendor availability, and flagship extrapolation; the full edition follows a measurement-anchored framework developed at RIKEN that replaces those with a vendor-neutral unit and a two-layer demand model [140] [full edition, Ch. 23]. The unit is the “baseline GPU” (BGPU), defined by serving throughput at 8 TB/s of memory bandwidth and convertible across vendors—a requirement denominated in BGPU is procurable from any vendor, which is what a public planning framework must guarantee and what a vendor-named unit cannot. The demand anchor is a measured 56-day campaign on a GB200-class partition: 101 users, 53,000 jobs, 340,000 GPU-hours, the partition saturated within five weeks [140]. The two layers are token service—researchers in three usage tiers driving a trillion-parameter master model at million-token context that delegates to smaller sub-agents—and scientific execution: simulation, domain-model training, and data. The structural finding is that ten percent of the tokens, generated at million-token context, cost nine tenths of the token-service GPUs, so national demand is overwhelmingly a long-context memory-capacity and memory-bandwidth problem—the sizing framework and the architecture argument are the same argument in different units.

Three countries are instantiated: Japan needs about 500,000 BGPU by 2030, the United States about 1.1 million, Singapore about 57,000, a per-researcher core near 0.94 BGPU with a band of at least $\pm 40\%$ before local measurement. Against a universal-reach ladder, Japan and the United States clear the uniform baseline rung and stand at roughly half (Japan) and seven tenths (the United States) of the tiered rung, and announced national infrastructures provision 0.02–0.10 BGPU per researcher against a modeled requirement near one—about ten times under for the United States and Japan, fifty times for Singapore. One side has the measurement without the money; the other has the money without the measurement. A nation that rents its master layer has rented the majority of its research computing; sovereignty has a computable price (about \$1.84–2.19 per delivered accelerator-hour across the three countries), which makes pooling across time zones rational; and the invitation to every government—run the seven-step measurement protocol and publish the parameter file—is the least expensive strategic action available.

9.3 The third path: an international consortium

For every country that is neither the United States nor China, the two available strategies are both unsatisfactory. Permanent dependence on a few foreign commercial providers cedes cost, control, and continuity; isolated national sovereign-AI programs duplicate the most expensive activity in the industry at sub-scale, and at valuations that assume rents no one may collect. What is needed is a *global* form of sovereign capability—sovereignty not through national isolation but through shared ownership of the knowledge, data practices, compute access, and institutional capacity required to build and operate scientific AI. The substance draws on a working discussion between the author and the director of the Swiss National Supercomputing Centre [141] [full edition, Ch. 24].

The decisive argument is financial. In a realistic AI-for-science portfolio, frontier language models and agents are roughly 14% of the physical workload; simulation and domain-specific AI take the rest [141]. That figure is usually deployed as reassurance, and it is the opposite once the price asymmetry is applied: purchased frontier tokens carry an estimated 20–60-fold premium over equivalent service from institutionally controlled infrastructure running open models [A/M], and a workload share of one seventh therefore becomes a budget share of 76.5% at a twenty-fold premium and 90.7% at sixty (Figure 9.1). Whoever sets the frontier token price sets the science budget of every institution that buys it: proposition P3 applied to national research capacity. A country that funds a supercomputer and then rents its intelligence layer has outsourced the majority of its own AI-for-science expenditure to a foreign price list. This is also why the Chinese price collapse is strategically effective without coercion—at a premium near unity, which cheap open weights on domestic hardware deliver, the budget problem dissolves, and the institution that solved it that way has acquired a dependency it did not negotiate.




Figure 9.1: Why the token price is the science budget. With a fraction f of physical work served by purchased frontier tokens at relative unit price ρ , the share of total expenditure that work consumes is $\Sigma_{\text{front}} = f_{\text{front}}\rho / ((1 - f_{\text{front}}) + f_{\text{front}}\rho)$. At the measured $f_{\text{front}} \approx 14\%$ and the estimated 20–60 $\times$ premium, a one-seventh workload share becomes 76.5–90.7% of the budget [A/M]. The curve’s shape, not its exact position, is the policy argument: at r near unity—which cheap open weights on owned hardware deliver—the problem dissolves, and the institution that solved it that way has acquired a dependency it never negotiated.

Open pretraining is therefore insurance, not a benchmark contest: a consortium’s own open model may never be the best production choice, and building it is still correct, because the return on pretraining is people who know how to do it, and a capability not exercised is a capability lost. Near-term science should meanwhile use the best open-weight models, including Chinese ones, without embarrassment. The insurance covers three risks, none presuming bad faith: policy and continuity risk (to rely on the openness of the current phase is to assume that this instance

will not follow the pattern the other three did), verification risk (Chapter 5), and regulatory risk in dual-use science, where holding one’s own models and evaluation evidence is the only realistic route to rules a working scientist can live with. Four functions define the institution: federated compute under common accounting, with a first proof of concept being the joint training of a fifteen-trillion-token open model split across compatible systems in several countries; scientific data stewardship, because science is the one community for which provenance and reproducibility are already requirements; model production with lifecycle obligations, since a clinical model may take three years to certify and must be supportable five years later—the strongest reason this cannot be a rotating academic project; and professional operation, because convening is cheap and operating is not. Industry participates without expropriation: an open, accountable base and closed, valuable private layers above it—a public good with private complements, the textbook case for public provision. The Global South is treated as talent, scientific questions, and data, not as aid, because the alternative is not neutrality but ceding the default to whoever shows up with a stack.

The organizational lesson from China is not a command structure but the reverse: government sets the stage and then allows intense bottom-up activity; a consortium that funds one model team in one country reproduces the failure mode, not the success. The full edition specifies ten operational design parameters and the failure mode each avoids (Table 9.1), and its feasibility appendix finds that of fourteen legal and engineering problems a founding instrument must solve, nine have answers on the shelf, three require legal invention, and none is a reason to wait, since every one is also unsolved for each national program that will otherwise duplicate the work seven times [full edition, App. E]. The governing framing is neutrality: the institution must be neither anti-American nor anti-Chinese, because a certificate issued by one bloc about the other bloc’s artifacts is worth nothing to anybody. Sequence: frame; prove federation with three to five committing countries and one visible cross-border pilot; then institutionalize. A five-country story with a working pilot beats a large inactive membership, and a proposal that names functions but not instruments is a wish.

Table 9.1: The consortium, specified operationally: ten design parameters and the failure mode each avoids [S].

Parameter	Recommended answer	Failure mode if deferred
Legal form, host	Treaty-adjacent foundation in a neutral, non-aligned-capable jurisdiction	Seen as a bloc instrument; loses the Global South and one principal
Membership, votes	States, centres, universities, companies in separate classes; contribution-weighted votes capped below a blocking share	Largest contributor becomes owner
Funding formula	Public base + private risk capital + grants + fees	Pure appropriation insufficient; pure fees exclude the long tail
Export-control compliance	Per-jurisdiction architecture, segmented enclaves	One member’s rule governs everyone
Model and data licensing	Open weights and documented recipes; layered data policy	Ambiguity blocks industry
Compute allocation	BGPU unit, peer review, reserved and burst tiers, cross-border settlement	Federation degrades to a mailing list
Certification authority	Operates the Level 0–5 ladder, publishes irrespective of origin	Vacancy stays vacant, or is filled by a supplier
Incident response	Named duty officer, disclosure, revocation, mirror withdrawal	One unowned agentic incident destroys credibility in a week
Chinese and US technology	Capability- and evidence-based; no origin rule either way	Forfeits neutrality
Procurement neutrality	Vendor-neutral units and interfaces; conflict-of-interest rules	Anchor vendor captures the standard

Why this is the strategically decisive institution: it is where the book’s threads terminate—the institutional form of “take the commons, verify it, contribute to it, and own the layer that makes

it safe to use.” It holds something neither Washington nor Beijing can revoke, it is the only credible home for the certification authority the previous chapter demanded, and it is affordable once pooled. Whether it gets built in the next two years or the next ten is, on the evidence of the precedents, the difference between shaping the outcome and documenting it.

10 The Convergence

In one paragraph

The elements interlock into a single machine for commoditizing artificial intelligence: software commoditization through open weights and open interfaces, hardware proliferation through accelerators that need only be sufficient and numerous, and decentralized pretraining and inference that dissolve the centralized-supercluster premise. The Western AI economy is a bet that intelligence will remain scarce and rentable; the Chinese strategy is a bet that it can be made abundant and free, and that whoever makes it abundant collects the durable prizes. Solar, batteries, and EVs each ended with the scarcity bet losing. The book states six developments that would falsify it, takes the strongest counter-argument—recursive self-improvement on closed American compute—seriously enough to give it its own section, and carries two conclusions that change the shape of the argument: the commons has a security bill that falls due at scale, and there is a third path whose economics are better than its politics suggest. For everyone who is not the United States, the strategic question becomes not whose side wins but who builds fastest on the commons the winner created.

10.1 The machine the elements form

Chapter 4 placed the elements on the playbook and scored the campaign step by step; what remains is to state the machine they form. Read separately, each element has a Western rebuttal—the models trail the closed frontier on the hardest benchmarks, the accelerators trail NVIDIA per watt (though a GPU-free Chinese machine now leads both HPL and HPCG [48]), the fabs trail TSMC by two nodes. Read together, the rebuttals miss the structure. *Software commoditization*: open weights plus mandated interoperability and standardized matrix extensions on an open ISA let a whole national ecosystem optimize collaboratively; the claim is about the moat’s shape, not its depth at any instant—CUDA’s advantage is not being crossed in one leap, the population entitled and equipped to erode it is being made as large as the ecosystem itself. *Hardware proliferation*: with the software requirement lowered, accelerators need not match NVIDIA, only be sufficient and numerous, paired with immense capacities of cheap commodity memory from a domestic fab—sufficiency plus volume is how TOPCon modules and LFP packs won. *Decentralized production*: half-megawatt sparse pretraining and a five-thousand-dollar inference appliance [M/S] put the means of production and of operation within reach of any sovereign state, university consortium, or mid-cap firm, each a customer for the Chinese stack and a defection from the Western API economy.

The Western AI economy is, structurally, a bet that intelligence will remain scarce and rentable: \$286 billion of annual private investment [10] capitalized against future margins on closed capability, housed in HBM-bound gigawatt data centers. The Chinese strategy is a bet that intelligence can be made abundant and free—and that whoever makes it abundant collects the durable prizes: the hardware volume, the standards, the talent flywheel, the adoption, and the geopolitical alignment of every country that builds on the open stack. Solar, batteries, and EVs each ended with the scarcity bet losing. The trillion-dollar data-center buildout is this war’s

equivalent of Q-Cells' 2007 capacity crown: magnificent, state-of-the-art, and aimed at the wrong future.

10.2 What would falsify this thesis

A living book owes its readers falsification criteria. The convergence argument weakens materially if, over the next two to three years: (a) closed US models open a *compounding* capability gap that adoption metrics start to follow—the 2.7% index gap widens while router and Hugging Face shares reverse; (b) recursive AI-for-AI research gives the closed frontier an escape velocity that open fast-following cannot match; (c) the capacity-first architecture fails in practice—MFU ceilings prove unreachable at scale, or sparse-model quality saturates in ways dense-HBM training does not; (d) CXMT's LPDDR6 ramp or SMIC's advanced-node yield stalls under tightened, better-enforced controls; (e) Beijing's anti-involution consolidation strangles the Darwinian dynamics that produced DeepSeek and Moonshot in the first place—the one playbook step that has historically destroyed value before creating it; or (f) the trust barrier hardens—Western and third-country certification regimes convert the security, censorship, and provenance record into an effective deployment wall, stalling the commons at the paying-workload boundary the way certification stalled the C919; a serious agentic security incident traceable to a widely deployed derivative would accelerate this branch by years. Symmetrically, the thesis strengthens with each frontier-class model trained end to end on domestic silicon; the first credible such run will be this war's Mate 60 moment.

10.3 The strongest counter-argument: two recursive loops, not one

Branch (b) deserves more than a clause, because it is the only falsification path that would make everything else in this book irrelevant rather than merely wrong. In its strongest form: if the closed US laboratories reach genuine recursive self-improvement—models that materially accelerate the research producing their successors—then compute scale compounds into capability faster than any fast-follower can copy, and the open ecosystem's cost advantage becomes irrelevant because it is competing on a curve that is being outrun. The evidence is no longer hypothetical: OpenAI attributed its 80% price cut of 31 July 2026 to efficiency gains its own frontier model had helped produce [82, 142], which means the same event this book cites as the first direct observation of P3's transmission is *also* the first public claim of a model improving the economics of its own successor. Both readings are available from one datapoint, and honesty requires printing both.

The reply is not that recursive self-improvement is unreal but that there are two recursive loops running on different substrates, and they do not compete on the same axis. The American loop runs through algorithmic efficiency—frontier models improving training and inference efficiency, compounding wherever compute is most abundant, which is the United States. The Chinese loop runs through industrial capability—models improving chip design, EDA, system architecture, and manufacturing yield, compounding wherever the manufacturing ecosystem is deepest and the co-design partners closest, which is China; the flagship's claimed research reproduction and autonomous silicon-design flow are the same recursive structure aimed at the layer that Chapter 9 argues matters most. The first loop produces more capability per FLOP; the second produces more FLOPs per dollar, domestically, out of an ecosystem that currently must lease them abroad. Which loop dominates is genuinely undetermined. Their failure modes differ: the American loop fails if capability gains stop converting into products people will pay a premium for—the commoditization mechanism of Chapter 3—while the Chinese loop fails if AI-assisted design cannot substitute for the tacit process knowledge that has defeated Chinese

industrial campaigns before, in aircraft and lithography. Neither failure is currently observable. The dashboard watches both, and the first party to publish an audited instance of its own loop closing will have supplied the most important datapoint in the field; neither has yet.

10.4 Two things the argument has to carry with it

Two of the book's chapters change the shape of the conclusion rather than decorating it. **The commons has a security bill, and it falls due at scale.** The convergence is, in the vocabulary of Chapter 5, a forecast of correlated deployment: a handful of open-weight families, thousands of unowned derivatives, wired into agents holding file, shell, network, memory, and credential authority. Every mechanism required to turn that structure into a systemic incident has been demonstrated in controlled work, and the compound events that have actually happened required no backdoor at all. This does not falsify the thesis—monoculture risk is origin-neutral, and a closed US-led monoculture would carry the same correlation with less transparency—but it identifies the single highest-value piece of unbuilt infrastructure in the competition: an authority that maintains reproducible artifact lineage across a derivative chain no vendor controls, tests integrations rather than models, and publishes results neither bloc's marketing owns. The commons will be adopted either way; whether it is adopted safely is currently nobody's job. **There is a third path, and its economics are better than its politics suggest.** For any country that is not a superpower the decisive number is a budget, not a benchmark (Chapter 9): whoever sets the frontier token price sets the research budget of everyone who buys it, which makes open capability an insurance policy with a computable premium and an international, institutionally neutral consortium the cheapest available form of strategic autonomy—and the only credible home for the certification authority the previous sentence demands. A coalition that builds it holds something neither Washington nor Beijing can revoke; a coalition that does not will discover that “open” and “available to us, indefinitely, on acceptable terms” were never the same word.

10.5 Implications for the West—and for the rest

The precedent industries teach that the losing move is to defend the rent: tariff the modules, ban the chips, close the weights, and wait. By the time the wall is complete, the world outside it is running on the commodity. The counter-strategy that has never been tried is to *win the commoditization race*—to make the open, abundant, decentralized version of AI a Western-led commons rather than a Chinese-led one, and to compete on the layers above it. Whether any Western polity can sustain that strategy against the incentives of its own incumbents is beyond this book's scope—but the solar, battery, and EV chapters record, three times, what happens when the answer is no.

For third countries—including advanced ones such as Japan—the convergence offers an uncomfortable but real option space: the open stack is *open*. RISC-V profiles, sparse-training recipes, capacity-first system design, and open-weight frontier models can be adopted, forked, and sovereignly operated by anyone. The half-megawatt pretraining cluster is as available to Kobe as to Guiyang. If the AI war ends the way the precedent wars ended, the strategic question for everyone who is not the United States becomes not *whose side wins*, but *who builds fastest on the commons that the winner created*. That, and not any single benchmark, is how China will win the AI war: not by reaching the frontier first, but by making the frontier free—and owning everything beneath it.

A Forecasts and Dashboard

The full edition carries a graded evidence ledger, a machine-readable claim register, and the two tables reproduced here [full edition, App. C, I]. The register states grades; these tables state bets. Every forecast is a scenario judgment [S] with a probability, a horizon, and a falsifier, so that a reader in 2029 can mechanically score the book; where the full edition prints a skeptic’s figure alongside the author’s, it appears in brackets. A reader who checks the dashboard quarterly does not need to re-read the book to know whether it is aging well. Four terms in the tables are the full edition’s: “stage two” is the second rung of its proof ladder for the integrated-CPU claim (parity on AI serving benchmarks after the settled HPL/HPCG rung); the “bandwidth knee” is the socket memory bandwidth below which the capacity-first architecture’s efficiency falls off (modeled near 5 TB/s); “VLA” is a vision-language-action model, the policy layer of an embodied system; and the “Max-tier template” is the pattern of opening a flagship’s weights while gating commercial use above a revenue threshold.

Table A.1: Probabilistic forecasts as of 18 August 2026 [S].

Proposition	Prob.	Horizon	Falsifier
Chinese models dominate the open-weight ecosystem	80% [70]	1–2 yr	Routed-token share below 45% for four quarters
Frontier-adjacent model trained end to end on the domestic Chinese stack	65% [50]	2–4 yr	Reversion to foreign silicon past 2028
Capacity-first (LPDDR-dominant or hybrid) pretraining production-competitive	45%	3–6 yr	MFU below $\sim 45\%$ at scale, or bandwidth knee above 8 TB/s
LX2-class integrated CPU competitive on AI benchmarks	55%	2–4 yr	No submission by 2029, or $>2\times$ deficit
Broad ($\geq 30\%$) US AI repricing	70% [60]	1–3 yr	AI revenue closing the capex gap through 2028
US closed models retain the premium enterprise tier through 2030	50% [65]	4 yr	Open-weight enterprise spend above 40%
Complete Chinese semiconductor-tooling autonomy	25%	5–10 yr	No domestic EUV-class lithography past 2032
Neutral certification authority for open weights established	30%	3–5 yr	Role occupied by a single bloc, or by nobody, past 2030
AI-for-science uplift $\geq 5\%$ measurable nationally	45%	5–8 yr	No uplift in weather, structure, or screening
Systemic agentic-security incident from a widely deployed derivative	20%	5 yr	(asymmetric: absence is weak evidence)
Integrated-CPU corner lowers lifecycle cost of changing workloads faster	40%	4–8 yr	Accelerator corner matches kernel latency while CPU porting tax persists
Production-metric scoreboard for scientific AI operated by a ministry	35%	3–5 yr	Leaderboards remain the only measure past 2030
Chinese accelerator or integrated CPU submits to MLPerf Training	50%	2–3 yr	None by end-2028

continued

Table A.1, *continued*

Proposition	Prob.	Horizon	Falsifier
$\geq 1,000$ humanoids productive under a single operator, metrics published	55%	2–4 yr	Teleoperation ratios not declining through 2028
Open-weight VLA models commoditize the embodied policy layer	45%	3–6 yr	Closed stacks hold the gap into the 2030s
Qwen3.8-Max weights ship	Resolved YES (13 August 2026)		
... under a permissive licence	Resolved NO for the flagship (US\$50M revenue gate); YES for the 27B model		
Next ≥ 1 T-parameter Chinese release ships open (the ratchet holds)	85%	1–2 yr	A major Chinese lab closing its frontier tier
Meta ships open Muse Spark 1.2 within a quarter at flagship quality, permissively	40%	1 qtr	The Max-tier template repeating
Chinese agent harness top-three in a major non-Chinese market	35%	2–4 yr	Western harnesses $>90\%$ of regulated deployments through 2029
Harness migration costlier in regulated than unregulated sectors	60%	2–3 yr	Comparable cost in both
NVIDIA ships a frontier-scale open-weight model with derivative licence	55%	2–3 qtr	Closed, or not shipping, through 2027
Compute futures reach sustained liquidity	40%	2–3 yr	Thin trading through 2028
Dedicated inference contract trades independently	30%	2–4 yr	Inference bundled into GPU-hour or API
Chinese inference capacity exported at scale	45%	3–5 yr	Endpoint share stalling while self-hosting carries volume
HBF ships in volume at ≥ 512 GB/stack, ≥ 1.6 TB/s, with an endurance figure	45%	2–4 yr	None by 2029, or below ~ 0.8 TB/s
Capacity-first appliance: ≥ 2 T open model, 4 users at ≥ 40 tok/s, 1 kW, under \$8k	35%	3–5 yr	Above \$15k or below 20 tok/s through 2030
Enterprise open-weight workload share reverses and passes 20%	40%	2–4 yr	Continued decline (licensing and governance binding, not memory)
≥ 100 MW AI facility sold, foreclosed, or written down $\geq 40\%$ below cost before 2030	45%	2–4 yr	All 2025–26 vintages refinancing at par through 2029
Chinese frontier model trained on domestically sited compute, offshore excluded	60%	2–4 yr	Leased offshore reliance past 2029
AI-assisted design flow produces a verified, manufacturable Chinese accelerator	45%	3–5 yr	Roadmaps slip despite AI-design claims

Table A.2: Quarterly falsification dashboard.

Indicator	Current (Aug 2026)	Threshold	Source
Chinese share of routed tokens	$\sim 61\%$	Sustained $< 45\%$ weakens P1	Neutral routers, monthly

continued

Table A.2, continued

Indicator	Current (Aug 2026)	Threshold	Source
Closed-over-open capability gap	3.3%	>10% sustained two years favours bifurcation	AI Index, arenas
Enterprise open-weight spend share	~11%	>25% confirms P3 transmission	Surveys, annual
Domestic end-to-end frontier training run	Not observed	First credible instance confirms P2	Lab reports
LX2-class AI benchmark submission	None	MLPerf at parity confirms stage two	MLPerf rounds
CXMT LPDDR6 / HBM3 ramp	Verification cleared	2+ year slip weakens the memory arbitrage	Quarterly
Cost per solved task, open vs closed	Unpublished	First credible series confirms or deflates the gap	Evaluations
OpenAI financing	IPO slipping	Down-round or new backstop raises P4	Filings
Data-center credit spreads	Oracle ~212 bps	Sustained widening raises P4	Daily
N. Virginia whole-sale colocation rent, 10MW+	\$155–185/kW-month, vacancy	Sustained decline with rising vacancy: utilization forecast failing; a rise: closet not displacing hall	CBRE/JLL, semiannual
Certification authority for open weights	Vacant	Occupation reshapes the trust and consortium chapters	Standards bodies
Max-tier openness lag	~10 days realized, 1 week committed (Qwen3.8-Max)	Lag >2 quarters, or a frontier tier closed, breaks the ratchet	Per release
Recursive-loop disclosures, US vs China	Both claimed, neither audited	First audited instance is the field's most important datapoint	Lab disclosures, tape-outs
Offshore compute leasing by Chinese labs	Permitted, reportedly in use	Closure by rule is the highest-impact tightening available to export policy	BIS or Commerce

Sources and Selected References

The sources for the claims this edition cites, numbered in order of first citation and copied from the full edition (URLs are hyperlinked here as [link] and printed in full there), whose bibliography runs to some seven hundred entries and which cites every figure in this text—including the many that this abridgement, for brevity, leaves uncited. Evidence grades used throughout: [V] independently verified, [P] primary-source claim, [A] analyst estimate, [M] author-modeled result, [R] roadmap, [S] scenario judgment. The full edition also states its own known weaknesses, and a reader of the abridgement should know them: many mid-2026 details rest on secondary aggregators; the attribution of the LX2 processor’s manufacturing is analyst inference; the capacity-first memory design is a modeled architecture with no measured realization; the model cards remain incomplete in the fields that matter most for procurement, tokens and wall-clock per solved task; the national sizing is measurement-anchored for one country only; and the architectural thesis that the CPU corner’s breadth compounds faster than the accelerator corner’s depth is a judgment, not a measurement, and the one most likely to be shown wrong. The reproducible models behind the numbers—the appliance-versus-rack cost model and the AI-for-science value funnel—ship with the source of the full edition. Both editions are living documents; the title-page date is the currency date, and corrections are welcome and will be credited.

- [1] Our World in Data, “Learning curves: solar module prices fell ~20% per doubling of cumulative capacity; \$106/W (1976) to \$0.38/W (2019).” [link]
- [2] A. Jäger-Waldau / ISES, “PV Industry Trends” (Oct. 2024): 2023 China production shares—polysilicon 92%, wafers 98.1%, cells 92%, modules 84.6%. [link]
- [3] SNE Research via electrive, “CATL continues to dominate global battery market” (FY2025: 1,187 GWh; CATL 39.2%, BYD 16.4%). [link]
- [4] BloombergNEF, “Lithium-ion battery pack prices fall to \$108/kWh” (2025 survey; China \$84/kWh). [link]
- [5] Rhodium Group, “Why are Chinese EVs so cheap?” (BYD ~80% Tier-1 in-house; ~\$4,700/vehicle structural advantage vs. Tesla; foreign share 62%→35%). [link]
- [6] CnEVPost, “China NEV sales Dec. 2025 (CAAM)”: FY2025 NEV 16.49M; total 34.4M; exports 7.1M incl. 2.62M NEVs. [link]
- [7] HSF Kramer, “China battery technology export restrictions” (Jul. 2025); Global Times on Oct. 2025 extension to cells ≥ 300 Wh/kg, anodes, equipment. [link]
- [8] CSIS, “Rare earth export restrictions: one year later” (Apr. 2025 controls; production stoppages; Oct. 2025 escalation and truce). [link]
- [9] Hugging Face, “State of Open Source, Spring 2026” (Chinese models 41% of downloads; Qwen >113k derivatives). [link]
- [10] Stanford HAI, *AI Index Report 2026* (US–China top-model gap 2.7%; US \$285.9B vs China \$12.4B private investment). [link]
- [11] Reuters (exclusive), “China to publish policy to boost RISC-V chip use nationwide” (Mar. 3, 2025; eight agencies incl. MIIT, CAC, MOST). [link]
- [12] CRN Asia, “Huawei’s open source gambit: SuperPoD reference architecture and UnifiedBus 2.0 released openly” (Sep. 2025). [link]
- [13] Design & Reuse, “SMIC achieves 5nm production on N+3 node without EUV tools” (TechInsights teardown of Kirin 9030, Dec. 2025). [link]
- [14] CNBC, “CXMT closes 466% higher in Shanghai debut” (Jul. 27, 2026; \$8B raised; largest A-share company; prospectus disclosed no commercial HBM). [link]
- [15] PRC MFA, “Global AI Governance Action Plan” (WAIC, Jul. 2025), full text. [link]
- [16] N. Lambert, “On China’s open-source AI trajectory” (Interconnects; State Council “AI+” plan targets). [link]
- [17] Eco-Business/Bloomberg, “Chinese renewable companies slow to tap \$47 billion credit” (CDB lines, 2010). [link]
- [18] Bernreuter Research, “Polysilicon price trend” (\$475/kg Mar. 2008 → \$15.9/kg Dec. 2012 → \$5.6/kg Jul. 2024). [link]
- [19] TaiyangNews (NEA data), “China installed over 315 GW AC new solar PV capacity in 2025”; cumulative 1.2 TW AC. [link]
- [20] EnergyTrend solar price index (Jul. 29, 2026): TOPCon modules RMB 0.67–0.72/W ($\approx$ \$0.09–0.10/W). [link]
- [21] SCMP, “China’s top leadership takes aim at disorderly low-price competition” (Jul. 2025); Xinhua commentary. [link]
- [22] pv magazine, “Beijing group creates official body to guide China’s polysilicon capacity” (Dec. 2025). [link]
- [23] CGTN, “China scraps white list of domestic auto battery suppliers” (Jun. 2019). [link]
- [24] IEA, *Global EV Outlook 2026*, “Electric vehicle batteries” chapter. [link]
- [25] S. Kennedy, CSIS, “The Chinese EV dilemma: subsidized yet striking” (\$230.9B cumulative support 2009–2023). [link]
- [26] CnEVPost, “China NEV retail Dec. 2025 (CPCA)”: FY2025 12.86M, penetration 54.07%, Dec. 60.4%. [link]
- [27] G. DiPippo et al., “Red Ink: Estimating Chinese Industrial Policy Spending in Comparative Perspective,” CSIS (2022): ~\$248B/yr, 1.7% of GDP. [link]

- [28] Tianxia Gongchang Research, “China semiconductor equipment 2026” (localization by segment: litho <5%, etch 30–35%, ArF resist <10%, valves <10%); TrendForce on SMEE 28nm retraction. [link]
- [29] TrendForce, “Decoding China’s lithography push to challenge ASML” (Nov. 2025: SMEE 90nm shipping, 28nm immersion in development, Yuliangsheng trials; domestic sub-10nm litho ~2030); eeNews on LDP-EUV reports (speculative). [link]
- [30] Pharmasource, “China biopharma out-licensing surges to record \$137.7B in 2025”; Caixin on first-in-class pipeline (~24%); The Cancer Letter on the ODAC 14–1 sintilimab vote (2022). [link]
- [31] MacroPolo Global AI Talent Tracker 2.0 (Mar. 2024; NeurIPS 2019 vs 2022 authors, top-tier $\approx$ top 20%): Chinese-undergrad share of top-tier researchers 29% $\rightarrow$ 47%; share working in the US 59% $\rightarrow$ 42%, in China 11% $\rightarrow$ 28% (MIT Technology Review, 27 Mar. 2024); MacroPolo ceased operations in March 2025 and later “Tracker 3.0” figures circulating online are not the Institute’s; Carnegie Endowment (Dec. 2025): 87 of 100 tracked 2019-cohort Chinese-origin researchers remained in the US; RIETI on the returnee wave and the K-visa (Oct. 2025). [link]
- [32] DeepSeek-V3 Technical Report (arXiv:2412.19437); Stanford FSI, “Taking stock of the DeepSeek shock.” [link]
- [33] Yahoo Finance, “Nvidia stock plummets, loses record \$589 billion as DeepSeek prompts questions over AI spending” (Jan. 27, 2025). [link]
- [34] South China Morning Post, reporting an Interconnects AI analysis of Hugging Face data (10 Apr. 2026): Qwen at $\sim$ 1 billion cumulative downloads by March 2026, 153.6M in February alone—more than double the combined total of the next eight families—and over 50% of all global open-model downloads, with >200,000 derivative variants. [link]
- [35] Tom’s Hardware, “Moonshot releases 2.8-trillion-parameter Kimi K3 as China works around U.S. compute limits” (Jul. 2026); VentureBeat coverage. [link]
- [36] OpenRouter, “The open-weight models that matter, June 2026” (GLM-5.2 #1 open on AA index; ranking of open models). [link]
- [37] Alibaba / Qwen, model repositories **Qwen/Qwen3.8-2.4T-A95B** (2.4T total / 95B active; 512 experts, 11 activated; 262,144 native context; text-only; 4.89 TB in 213 BF16 shards) and **Qwen/Qwen3.8-27B** (Apache 2.0, vision-language). The flagship carries the custom **qwen3.8-max** licence: separate commercial agreement required where model-as-a-service or AI-work-assistant revenue exceeds US\$50M in any consecutive 12 months; attribution required above 100M MAU or US\$20M monthly revenue. [link]
- [38] Moonshot AI, “Kimi K3 License” (custom commercial license: use, modification and redistribution permitted; separate agreement required for qualifying model-as-a-service businesses above a revenue threshold; attribution conditions on very large deployments). [link]
- [39] Menlo Ventures 2025 enterprise LLM survey: open-weight share of enterprise workloads 19% (2024) $\rightarrow$ 11% (2025). [link]
- [40] Epoch AI, “LLM responses to benchmark questions are getting longer over time” (Apr. 2025: reasoning models $\sim$ 5 $\times$ /yr; non-reasoning 2.2 $\times$ /yr). [link]
- [41] M. Sayed Ali, A. Novik, A. Boddupally et al. (Writer, Inc.), *The Harness Effect: How Orchestration Design Sets the Token Economics of Enterprise Agentic AI*, arXiv:2607.06906, Jul. 2026. 22 locked tasks, six models, five vendors: harness adoption saves 33–61% holding the model fixed, against 36% from switching between cheapest and most expensive model. Vendor-authored; n=22. [link]
- [42] DeepSeek, *DeepSeek Harness (dsh)*, MIT licence, developer preview v0.1, released 13 Aug. 2026; $\sim$ 100,200 GitHub stars and 9,400 forks within $\sim$ 48 hours. [link]
- [43] DeepSeek, DeepSeek Harness documentation: plugin-native architecture (models, tools, skills, sessions, sandboxes, filesystems, loops, orchestration and UI as components; “no privileged core to patch”), four runtime modes, append-only session logging with resume/fork/replay. [link]
- [44] Reuters via Yahoo Finance (13 Aug. 2026): DeepSeek takes V4-Pro to general availability, raises API prices for V4 Pro and V4 Flash, introduces peak/off-peak pricing; reported plans to at least double headcount and to expand private hiring of chip-design engineers toward an in-house accelerator reducing reliance on both NVIDIA and Huawei. [link]
- [45] Cyberspace Administration of China (17 Jul. 2026, WAIC), *Global Cooperation Initiative on Agent Mutual Trust, Interconnection, and Interoperability*: calls for “open, non-discriminatory, and transparent international interoperability standards for agents,” agent identity systems and shared multi-agent infrastructure; names no existing protocol. [link]
- [46] China Briefing, “China’s 15th Five-Year Plan: key insights” (“extraordinary measures” for ICs and basic software). [link]
- [47] CSIS, “Sustaining standards leadership: the United States cannot disengage from RISC-V” (membership statistics; 2B+ chips shipped). [link]
- [48] TOP500, “LineShine debuts at No. 1 as TOP500 enters new global exascale era” (Jun. 24, 2026: HPL 2.198 EF, HPCG 22.00 PF, 42.2 MW, 52.07 GF/W). [link]
- [49] J. Dongarra, “Report on the Chinese LineShine System,” Univ. of Tennessee Tech Report ICL-UT-26-01 (Jun. 23, 2026). [link]
- [50] J. Dongarra, T. Hoefer, and S. Matsuoka, “Do We Still Need GPUs? Where an Evolved CPU Can Replace Them in AI and Scientific Computing—and Where It Cannot Yet,” *Communications of the ACM*, in revision (2026).
- [51] J. Dongarra, T. Hoefer, and S. Matsuoka, “Do We Still Need GPUs? Rethinking AI and Scientific Computing on Accelerated CPUs,” full-length companion paper (2026): LX2 vs GPU-fleet analysis of 1T-MoE serving and training; the design-sequence/allocation argument.
- [52] MLCommons, MLPerf Training v6.0 results (16 Jun. 2026): adds DeepSeek-V3 (671B/37B MoE) pretraining and GPT-OSS 20B benchmarks; 95 systems from 24 organizations on 13 accelerator types; no Chinese-designed accelerator submissions. [link]
- [53] NVIDIA, Q2 FY2026 results (zero H20 sales to China-based customers); TradingKey, “Nvidia China revenue share 26% $\rightarrow$ $\sim$ 5%.” [link]
- [54] TradingKey, “Domestic AI chip shipments $\sim$ 1.65M units, 41% of China market in 2025; Huawei $\sim$ 812k.” [link]

- [55] Xinhua, “Huawei fully open-sources CANN” (Aug. 7, 2025). [link]
- [56] RIKEN Center for Computational Science, feasibility study of memory-bandwidth balance for 2.8T-parameter sparse-MoE pretraining at 20 PF(FP8) per accelerator: HBM4 capacity wall at <256 accelerators; LPDDR6X substitution at 2.3 TB/socket, 5.0 TB/s knee, 67% vs. 89% MFU ceilings (2026, companion report to this volume).
- [57] SanDisk and SK hynix (3 Aug. 2026): first Open Compute Project technical specification for High Bandwidth Flash, following the August 2025 standardization memorandum and a February 2026 consortium launch. Consortium members SanDisk, SK hynix, Google and Tenstorrent; NVIDIA, AMD, Samsung, Micron and Kioxia absent. [link]
- [58] HBF specification as reported at FMS 2026: up to 512 GB per stack in 8-Hi and 16-Hi NAND; three bandwidth grades from ~ 0.4 to 3.0 TB/s; UCIE host interface at up to 64 GT/s over 64 lanes; NAND dies stacked with TSVs in an HBM4-compatible footprint with a CBA logic die scheduling parallel sub-array reads. Evaluation samples to selected customers promised by end-2026. [link]
- [59] NVIDIA DGX Spark: GB10 Grace Blackwell, 128 GB LPDDR5X at 273 GB/s, “up to 1 PFLOP FP4” (sparse); Founders Edition \$3,999 at launch (Oct. 2025), raised to \$4,699 on 27 Feb. 2026 citing memory shortages; ASUS Ascent GX10 partner unit \$2,999 at launch. [link]
- [60] Framework Desktop with AMD Ryzen AI Max+ 395 and 128 GB LPDDR5X, \$1,999 at announcement (Feb. 2025); GMKtec EVO-X2 with the same part at \$1,799.99 early-bird (May 2025). Post-surge: AMD’s own Ryzen AI Halo 128 GB desktop at \$3,999 (May 2026). [link]
- [61] SK hynix chief executive, 11 Jul. 2026, on the day of the company’s Nasdaq listing: “We forecast that next year will be the worst year in the industry’s history from the supply perspective,” and “we still forecast that customer demand will remain higher than our supply capacity even beyond 2030.” [link]
- [62] TechInsights, “SMIC N+3 confirmed: Kirin 9030 analysis reveals how close SMIC is to 5nm” (Dec. 11, 2025; “significantly less scaled than leading commercial 5nm”). [link]
- [63] SemiAnalysis, “Huawei Ascend production ramp” (2024: 507k units; 2025: $\sim 805k$; die bank 2.9M; HBM stockpiles; SMIC 45k $\rightarrow$ 80k wpm). [link]
- [64] wccftch, “CXMT races to increase capacity to 600,000 wafers per month” (Jul. 2026; current 200–300k wpm). [link]
- [65] Tom’s Hardware (Citrini Research), “CXMT close to matching Micron’s memory capacity in 2026” (350k wpm model; China $\sim 1.4M$ wpm DRAM by 2030). [link]
- [66] Reuters (Ngui/Shen), “Chipmaker CXMT vaults to top of China’s valuation with 466% surge” (Jul. 27, 2026: \$8.6B raised; 6.73% free float). [link]
- [67] As of mid-2026 no public teardown confirms Chinese-fabbed DRAM in a shipping accelerator; shipping Ascend 910C teardowns show Samsung/SK hynix stacks (SemiAnalysis; prediction-market tracking). [link]
- [68] TrendForce, “Three Chinese chip toolmakers enter global top 20 as Naura rises to fifth by sales” (Feb. 2026; domestic tool share 20–30%). [link]
- [69] BigGo Finance, “Naura 2025 results: revenue $\sim$ RMB 39.4B (+31%)” (Apr. 2026). [link]
- [70] Asia Society Policy Institute, “China’s Big Fund 3.0: Xi’s boldest gamble yet for chip supremacy” (RMB 344B, May 2024, horizon 2039). [link]
- [71] SCMP (SEMI), “China to remain world’s biggest buyer of chipmaking equipment through 2027”; SEMI 2026 equipment forecasts. [link]
- [72] China National Energy Administration, 2025 installed-capacity and renewable-integration statistics. [link]
- [73] US EIA: 2025 utility-scale additions 53 GW (largest since 2002); 2026 planned record 86 GW. [link]
- [74] Ember, Global Electricity Review 2026: China 10,573 TWh (+503 TWh y/y) vs US 4,430 TWh; China fossil generation fell 0.9% in 2025; solar +40%. [link]
- [75] Lawrence Berkeley National Laboratory, *Queued Up: 2025 Edition* (data through 2024): $\sim 2,290$ GW of active generation and storage requests; median request-to-operation time above four years; only 13% of 2000–2019 requested capacity ever reached operation. [link]
- [76] Semafor/techchannel (Nov. 2025): provinces cut DC electricity up to 50% (to $\sim$ RMB 0.4/kWh) exclusively for approved-domestic-chip facilities; offsets 30–50% chip-efficiency gap. [link]
- [77] Reuters (exclusive), “China bans foreign AI chips from state-funded data centres” (Nov. 2025); provincial power-tariff discounts for domestic-chip facilities. [link]
- [78] BigGo / Tech Digest / 36Kr supply-chain reporting [A]: $\sim 70\%$ of Optimus supply chain China-based; excluding Chinese parts estimated to triple cost; 2025 rare-earth licensing touched actuator magnets. [link]
- [79] Fortune (6 Jun. 2026): 2025 global humanoid shipments $>13,000$, $\sim 85\%$ Chinese; US firms “a few hundred or less”; most units performative rather than functional—demos, exhibitions, education. [link]
- [80] Gasgoo / Fortune (Jun. 2026): Unitree STAR-market IPO approved; 5,500 humanoids sold 2025; humanoid ASP RMB 593k (2023) $\rightarrow$ RMB 166k (2025); gross margin 87.7% $\rightarrow$ 63.2%; $\sim$ \$250M revenue with $\sim$ \$41M profit—among the only profitable humanoid makers. [link]
- [81] TechCrunch, “China tells its tech companies they can’t buy AI chips from Nvidia” (Sep. 17, 2025); Global Times on the CAC summons (Jul. 31, 2025). [link]
- [82] Forbes / G. Wang (31 Jul. 2026), “Why OpenAI’s 80% price cut could trigger a race to the bottom in AI”: GPT-5.6 Luna input \$1.00 $\rightarrow$ \$0.20 and output \$6.00 $\rightarrow$ \$1.20 (80% cut), Terra down 20%, Sol unchanged; OpenAI attributed the cut to system efficiency while the analysis attributes competitive pressure to Kimi K3 and DeepSeek V4. Note the sequencing: the cut preceded Qwen3.8-Max’s general availability by three days. [link]
- [83] JFrog (2024: ~ 100 malicious HF models); ReversingLabs “nullifAI” (Feb. 2025: broken-pickle evasion); Sonatype Picklescan CVEs (2025). [link]
- [84] Palo Alto Unit 42, “Model namespace reuse” (Sep. 2025): trusted-name takeover propagating into Vertex AI and Azure AI Foundry. [link]
- [85] Anthropic / UK AISI / Turing Institute (Oct. 2025): ~ 250 poisoned documents backdoor models regardless of scale. [link]

- [86] Hubinger et al., “Sleeper Agents” (arXiv:2401.05566): backdoors survive SFT/RLHF/adversarial training; agent-supply-chain extensions (CAIS 2026). [link]
- [87] Qiu, Zhou, Ferrara, “Information Suppression in Large Language Models: ...Censorship in DeepSeek” (arXiv:2506.12349): semantic-level suppression across 646 sensitive prompts. [link]
- [88] Naseh et al., “R1dacted: Investigating Local Censorship in DeepSeek’s R1” (arXiv:2505.12625): topic/language-dependent censorship, transfers to distilled models. [link]
- [89] Anthropic, “Detecting and preventing distillation attacks” (Feb. 23, 2026): >16M exchanges via ~24,000 accounts (MiniMax >13M, Moonshot >3.4M, DeepSeek >150K). Vendor allegation. [link]
- [90] TechTimes (4 Aug. 2026): MiniMax H3 tops the open-weight video arena; licence excludes local deployment in the US, EU, UK, and South Korea. [link]
- [91] S. Matsuoka, “Latent Model Backdoors and Prompt-Borne Worms: A Mini-Survey of Evidence, the ‘Order 66’ Compound Scenario, and Defenses,” v6.0, 31 July 2026. Companion survey; supplies the evidence hierarchy, threat taxonomy, and control architecture used in the full edition and the full edition.
- [92] W. Zhang and S. Pei, “Your LLM Agent Can Leak Your Data: Data Exfiltration via Backdoored Tool Use,” Findings of ACL 2026, pp. 25105–25129. [link]
- [93] OpenAI, “OpenAI and Hugging Face Partner to Address Security Incident During Model Evaluation” (21 Jul. 2026). [link]
- [94] Hugging Face Security Team, “Anatomy of a Frontier Lab Agent Intrusion: A Technical Timeline of the July 2026 Incident” (27 Jul. 2026). [link]
- [95] Anthropic, “Investigating Three Real-World Incidents in Our Cybersecurity Evaluations” (30 Jul. 2026). [link]
- [96] Fortune (4 Aug. 2026) / Axios (3 Aug. 2026): White House finalizes a voluntary frontier-model evaluation framework in a closed-door session with OpenAI, Anthropic, Microsoft, Meta, NVIDIA and others; document not publicly released. [link]
- [97] The Diplomat / Atlantic Council: World AI Cooperation Organization founded Shanghai, Jul. 2026, 29 member states. [link]
- [98] Forbes/CNBC/Bloomberg, “OpenAI valuation reaches \$852 billion after \$122B round” (Mar. 31, 2026); prior \$500B secondary (Oct. 2025). [link]
- [99] CNBC, “OpenAI resets spend expectations, targets around \$600 billion by 2030” (20 Feb. 2026): compute-spending target walked back from headline \$1.4T commitments; ~\$280B revenue projected by 2030. [link]
- [100] Anthropic, “Series H” (May 28, 2026: \$65B at \$965B post-money); CNBC coverage. [link]
- [101] PitchBook via Yahoo Finance, “Anthropic’s gross margin is the most important number in AI” (10 Jun. 2026): \$0.71 compute cost per revenue dollar in Q1 2026, projected \$0.56 in Q2, trajectory toward ~44%. [link]
- [102] CNBC/Axios, “Nvidia becomes first \$5 trillion company” (Oct. 29, 2025). [link]
- [103] NVIDIA (10 Aug. 2026), “NVIDIA Partners With Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR to Establish AI Compute Infrastructure Financing Platforms to Mobilize Over \$500 Billion of Third-Party Capital.” Memoranda of understanding, “subject to execution of the final agreements”; NVIDIA’s own contribution not stated in the release. [link]
- [104] Introl, “Hyperscaler capex \$600B+ in 2026”; Forbes (Jun. 2026): \$700–900B range, capital intensity 45–57% of revenue. [link]
- [105] Axios/CNBC: Oracle 5-yr CDS from 129bps (Feb. 2026) to record 212bps (Jul. 2026) on AI debt load. [link]
- [106] BEA third-estimate analyses: AI/data-center capex ~74% of US Q1 2026 GDP growth; Furman: ex-data-centers H1 2025 growth ~0.1%. [link]
- [107] Data Gravity, “China’s open-weight takeover” (OpenRouter token-share series; Llama <1%; provider shares). [link]
- [108] G. Appenzeller (a16z), “LLMflation: LLM inference cost is falling 10× per year” (Nov. 2024). [link]
- [109] MIT Project NADA, “The GenAI Divide” (Aug. 2025): ~95% of enterprise GenAI pilots produced zero P&L impact. [link]
- [110] Bain & Company, 6th Global Technology Report (Sep. 2025): \$2T annual AI revenue needed by 2030; ~\$800B shortfall. [link]
- [111] Intellectia/CNN: \$1.3T semiconductor selloff (Jun. 4–9, 2026) on Broadcom AI-networking miss. [link]
- [112] CNBC/NBC (Jul. 28–29, 2026): memory-led chip rout, SK hynix –15%, >\$1T lost, Nasdaq-100 in correction; Samsung/SK warnings of shortages to 2027+. [link]
- [113] Epoch AI, “LLM inference prices have fallen rapidly but unequally across tasks” (Mar. 2025: 9×–900×/yr, median ~50×). [link]
- [114] Epoch AI, “Global AI computing capacity is doubling every 7 months” (Jan. 2026); “Acquisition costs of leading AI supercomputers doubled every 13 months” (Jun. 2025). [link]
- [115] S. Matsuoka, *Datacenter Construction for State-of-the-Art GPU Machines: A survey with three reference architectures, explicit cost scopes, evidence-graded sources, P10/P50/P90 ranges, escalation to 2030, an independent Japanese reconstruction, and a costed lean reference design*, Rev. 5, RIKEN Center for Computational Science technical report, Aug. 2026 (all figures from public sources or the author’s modelling; distributed without restriction). Facility-only cost per watt of IT load, excluding land and IT, at 50 MW: architecture A (resilient inference/cloud, Tier-III-capable) \$11.9/W P50 with a \$8.5–17.7 P10–P90 band; B (AI-training-optimised, warm-water, chillerless) \$9.0/W; C (interruptible/phased) \$8.4/W; utility interconnection the widest single driver; large-load interconnection 2–7 years by market; discrete-block penalty for capacity built ahead of load. Author’s own companion survey.
- [116] CME Group and Silicon Data (11 Aug. 2026): compute futures to launch 5 Oct. 2026, listed on NYMEX, each contract representing one month of GPU rental cost referencing H100 and B200. P. Keavey (CME): “Compute has become the currency of the AI age.” [link]
- [117] Baiguan (9 Jul. 2026), on neutral-router economics: Chinese models at roughly 61% of token consumption while Anthropic holds ~12% of volume and captures ~46% of revenue; large enterprises still

- overwhelmingly buy from Anthropic, OpenAI, Azure and Google Cloud; China's 2017 National Intelligence Law as a persistent procurement obstacle. [link]
- [118] White & Case / FPF: Japan's AI Promotion Act (passed 28 May 2025)—first AI-dedicated law; innovation-first, no penalties, Cabinet-level AI Strategy Headquarters and national AI Basic Plan. [link]
- [119] Gibson Dunn / European Parliament: Digital Omnibus provisional agreement (May–Jun. 2026)—AI Act Annex III high-risk obligations postponed to 2 Dec. 2027, embedded high-risk to Aug. 2028; AI Office GPAI powers strengthened. [link]
- [120] European Commission, InvestAI / AI Gigafactories (€20B; 77 proposals across 16 member states). [link]
- [121] EuroHPC JU (30 Jul. 2026): formal call for up to seven AI gigafactories (each 100k+ advanced chips), expected to trigger >€30B; 76 expressions of interest across 16 member states. [link]
- [122] IndiaAI Mission GPU tenders: ~38,000 GPUs empanelled across three rounds at subsidized rates reportedly under \$1/GPU-hr; reported utilization shortfall (“can’t find enough users”). [link]
- [123] Counterpoint / TrendForce: SK hynix ~62% HBM share; reportedly ~two-thirds of leading HBM4 volume; Samsung/SK hynix scaling into the 2026 memory cycle. [link]
- [124] NVIDIA/HUMAIN partnership (500 MW-class “AI factories,” several hundred thousand GPUs); AMD–HUMAIN up to \$10B/6 GW; Stargate UAE 1 GW within a 5 GW campus; US Commerce authorization of up to 70,000 advanced chips to G42/HUMAIN (Nov. 2025). [link]
- [125] Bloomberg, “Lawmakers urge export controls on RISC-V” (Nov. 2023); Reuters on Commerce review (Apr. 2024); Sen. Warner letter to BIS (Sep. 30, 2025). No restriction enacted as of mid-2026. [link]
- [126] Caixin / Fortune / State Council reporting: anti-“involution” campaign; pricing-law amendments (Aug. 2025); State Council warning against “disorderly competition” in AI; Xi Jinping publicly questioning whether every province needs AI/compute/EV projects. [link]
- [127] BNN Bloomberg, Alibaba FY results (May 2026): three-year AI capex plan of RMB 380B (~\$56B) to be exceeded without a new stated target; Cloud Intelligence revenue +38% y/y with AI products ~30% of external cloud revenue; ex-items net income down 99.7% and adjusted EBITDA down 84%; CEO statement that “margin is still secondary”; target of >\$100B combined external AI and cloud revenue over five years. [link]
- [128] MIT Technology Review (Mar. 2025) / Light Reading / ASPI: hundreds of state-subsidized Chinese AI data centers with reported utilization far below plan (up to ~80% of some new capacity unused); rental-price collapse ending the buildout binge. [link]
- [129] ChinaTalk / CSIS / trade reporting: HBM and advanced packaging as binding constraints on the Ascend fleet; CXMT assessed roughly three years behind the Korean leaders. [link]
- [130] ECMWF, “ECMWF’s AI forecasts become operational” (25 Feb. 2025) and “Ensemble AI forecasts become operational” (1 Jul. 2025). [link]
- [131] Google DeepMind, “AlphaFold: five years of impact” (2025). [link]
- [132] Insilico Medicine, rentosertib Phase IIa, *Nature Medicine* (3 Jun. 2025). [link]
- [133] MIT Department of Economics, statement on “Artificial Intelligence, Scientific Discovery, and Product Innovation” (16 May 2025): no confidence in the data; withdrawn. [link]
- [134] ECMWF AIFS blog, “Farewell to external AI models” (May 2026): all four external models withdrawn from real-time running over analysis sensitivity. [link]
- [135] NSF NCSES, *Science and Engineering Indicators 2025*: global R&D \$3.1T (PPP, 2022); US \$923B, China \$812B. [link]
- [136] B. F. Jones and L. H. Summers, “A Calculation of the Social Returns to Innovation,” NBER w27863: \$13.3 per \$1 at a 5% discount rate. [link]
- [137] N. Bloom, C. I. Jones, J. Van Reenen and M. Webb, “Are Ideas Getting Harder to Find?” *AER* 110(4), 2020. [link]
- [138] D. Acemoglu, “The Simple Macroeconomics of AI,” NBER w32487 (2024): TFP +0.66% over ten years, *explicitly excluding* protein-folding-class scientific advances as unlikely within the window. [link]
- [139] Nature Index 2026 Research Leaders (10 Jun. 2026): China first, output +22.4% y/y; CAS the leading institution. [link]
- [140] S. Matsuoka, “How Much Compute Does a Nation’s Science Need? A Portable, Measurement-Anchored Framework for Sizing Sovereign AI-for-Science Infrastructure, with Instantiations for Japan, Singapore, and the United States,” RIKEN R-CCS Technical Report, July 2026. Supplies the BGPU unit, the two-layer demand model, the parameter taxonomy and transfer rules, the measurement protocol, and the three national instantiations used in the full edition.
- [141] S. Matsuoka and T. Schulthess, “AI for Science and an International Consortium,” edited discussion record v1.0, 31 July 2026 (RIKEN R-CCS / CSCS). Unpublished working document; numerical values are discussion estimates unless otherwise noted.
- [142] M. Berman, “Open-source is WINNING,” Matthew Berman (YouTube), Aug. 2026. Review of the Qwen3.8-Max release: identifies research reproduction (PaperBench 93.0) as “the required skill for automated AI research,” noting the model “started from nothing but the paper and a set of GPUs—no starter code, no ready-made pipeline” and “invented and tested 18 improvement ideas of its own across four rounds”; reports Alibaba’s claim that the model “has independently achieved the autonomous execution of the entire silicon design flow,” with the observation that “one of the things that China is behind on is chip design and chip manufacturing—but if they had a model that could help them design these chips, that would no longer be their weakness”; per-task cost data (Qwen3.7-Max \$1.28 vs GPT-5.6 Sol \$1.23, with Anthropic models most expensive); the model-size gap (US closed frontier rumoured ~7T+ against 2.4–2.8T Chinese); the enterprise-adoption argument (“95% of the capability at a fraction of the price”); the prediction that “the model will be co-designed with the chip... with extreme co-design between model and hardware it might actually put us, as we’re becoming dependent on Chinese open source, highly dependent on Chinese chips as well”; and the recursive-self-improvement counter-thesis, including OpenAI’s attribution of its 80% price cut to efficiency gains produced by its own frontier model. [link]