

How China Will Win the AI War

The Convergence Strategy: A Structural Path to AI Dominance

CONDENSED EDITION

A short-form companion to the living book¹

Satoshi Matsuoka

v0.23 — 9 September 2026

¹This condensed edition abridges the full living book of the same version and date. Every number in it is taken from the full edition, which carries the sources, evidence grades, and derivations. Version v0.23 — 9 September 2026.

About This Edition

The full edition of this book runs to more than three hundred pages because its argument is a systems argument, and a systems argument has to be shown layer by layer before it can be believed: the model, the compute, the semiconductor, the watt, the data, physical production, trust, governance, finance, and the international order they imply. Many readers do not need the layers. They need the shape—what is being claimed, why, what would prove it wrong, and what it means for a government, a laboratory or a company that is neither in Washington nor in Beijing. This condensed edition is for them.

Three rules govern the abridgement. First, nothing here is new: every number, grade, forecast and recommendation is taken from the full edition of the same version and date, which carries the sources, the evidence grades, the derivations and the falsification apparatus. Second, the industrial precedents that occupy six chapters of the full edition—the solar, battery, electric-vehicle and telecommunications wars, the playbook they distill to, and the campaigns where the same playbook has failed—are here summarized in a single chapter, because what matters for the argument is the playbook they establish and its limits, not the campaign histories. Third, the weight of this edition is shifted deliberately toward the overall strategy and its international consequences: what the convergence means for the West, for Japan and Europe, for the non-aligned majority of the world, and for the institution the analysis keeps arriving at. Readers who want the engineering—the memory-architecture arithmetic, the appliance cost model, the sizing framework, the security fault tree—will find pointers into the full edition throughout, marked [full edition, Ch. n].

The title is deliberately provocative, but it is not a prediction of inevitability; it is a structural analysis. “Winning” means what it meant in solar, batteries and electric vehicles: capturing the overwhelming majority of global production, setting the cost curve, controlling the supply chain, defining the de facto standards, and reducing erstwhile leaders to protected niches. By that definition China has already won three technology wars in which the West began with every advantage, and the claim of this book is that artificial intelligence is following the same trajectory—visible only when the elements are laid side by side.

A note on the author’s position is owed and is given in full in the preface to the full edition. In brief: I direct the RIKEN Center for Computational Science, whose flagship program runs on an American accelerator architecture and whose deepest partnerships are with American and European institutions. Nothing in my professional position gives me an incentive to argue that the Chinese stack is winning; a great deal of it gives me an incentive to argue the opposite. I did not begin this project expecting to arrive where it has arrived, and I have tried to hold the argument to a standard of proof appropriate to how much I would prefer it to be wrong. Where the evidence does not support the strong claim, the book—and this abridgement of it—says so.

The currency date of this edition is 9 September 2026, the date on the title page. Every factual claim is current as of that date; anything learned after it belongs to the next edition. This is a living document. Corrections are welcome, and the fastest way to change its conclusions is to falsify an entry in the full edition’s evidence ledger.

Contents

About This Edition	i
Terms Used in This Edition	iii
1 The Thesis	1
1.1 Three wars, one playbook	1
1.2 Why AI is the fourth war	1
1.3 Openness as a weapon	2
1.4 Four propositions, stated separately	2
1.5 Evidence grades	3
1.6 The theory, in brief	3
2 The Precedents and the Playbook	4
2.1 Three campaigns, in brief	4
2.2 The playbook distilled	5
2.3 Where the analogy breaks	6
3 The Strategy in Whole	7
3.1 Where each element sits	7
3.2 What the assembly licenses	8
4 The Elements of AI	8
4.1 The model: open weights as the commodity layer	9
4.2 The harness: the layer the contest moved to	11
4.3 The compute: RISC-V, accelerators, and the memory arbitrage	12
4.4 The semiconductor: fabs, logic, and DRAM	14
4.5 The watt: energy as the fourth element	15
4.6 Data and feedback: the fifth element	16
4.7 Embodied AI: the loop closes on the factory floor	17
5 The Mirror: The West Runs the Playbook Against Itself	17
5.1 Who commoditizes what	18
5.2 The accelerator vendor gives away the model—and underwrites the buildout . . .	19
5.3 The other direction: the laboratories and the hyperscalers commoditize the silicon	20
5.4 The asymmetry that survives	21
5.5 What the mirror licenses, and what it forbids	22
6 Trust, Security, and the Rulebook	22
6.1 Trust: the contested layer	23
6.2 Monoculture and agency: the security cost of the commons	23
6.3 The governance contest: who writes the rules	24
7 The Wrong Side of the Trade	25
7.1 The exposure, marked to market	25
7.2 The squeeze, and the four bubbles	26
7.3 The Jevons rebuttal, taken seriously	27
7.4 Vintages, buildings, and railroads	27
7.5 Scenarios, probabilities, indicators	27

7.6	The commodity turn, and where value reappears	28
8	How the Rest of the World Chooses	29
8.1	The decision matrix	29
8.2	The export offer, as it exists	30
8.3	Nine positions	31
8.4	Five moves for any non-aligned country	32
9	Strategies for the Rest	32
10	The Prize, the Bill, and the Third Path	34
10.1	The prize	34
10.2	The bill: how much compute a nation’s science needs	35
10.3	The third path: an international consortium	36
11	China’s Scientific Machine	37
11.1	The policy record, read precisely	38
11.2	The output, the artifacts, and the machine that won the prize	39
11.3	Scored on the playbook, and what it does to the consortium	40
12	The Convergence	40
12.1	The machine the elements form	41
12.2	What would falsify this thesis	41
12.3	The strongest counter-argument: two recursive loops, not one	41
12.4	Two things the argument has to carry with it	42
12.5	Implications for the West—and for the rest	42
13	The Argument in Summary	42
13.1	The answer in three parts	43
13.2	What would change the answer	43
13.3	What this book is not claiming	44
A	Forecasts and Dashboard	45
	Sources and Selected References	47

Terms Used in This Edition

The full edition carries a complete notation glossary; the following are the terms a reader of this edition meets without introduction.

Open weights. A model whose trained parameters are published for download, so that anyone can run, fine-tune, or redistribute it, as opposed to a *closed* model reachable only through a metered API.

Token; \$/M tokens. The unit of text a model reads or writes; prices are quoted per million tokens.

MoE (mixture of experts); sparse model. A model in which only a fraction of the parameters (the “active” set) is used per token, so a very large model can run at the cost of a much smaller one;

“280B/16B” means 280 billion total, 16 billion active.

Frontier; frontier-adjacent. The best models available at a given date; models within a few points of them on independent indices.

Harness. The runtime an agent runs in—tool schemas, permissions, memory, sandbox, session store, orchestration—as distinct from the model it calls.

Quantization. Storing weights at lower precision (e.g. four bits) to shrink memory and cost, at some

loss of quality; “FP8” is an eight-bit floating-point format used in training and inference.

KV-cache. The working memory of context a model must hold while generating; it grows with context length and is a major consumer of accelerator memory.

HBM (high-bandwidth memory). Stacked DRAM mounted on the accelerator package: very fast, scarce, expensive, and under the tightest export controls.

LPDDR; DRAM. Commodity memory of the kind in phones and laptops (LPDDR) and PCs (DDR); abundant and cheap, with far lower bandwidth than HBM. “SOCAMM2” is a compact LPDDR module format.

HBF (high-bandwidth flash). A proposed stacked-NAND memory tier in an HBM-class package: non-volatile, up to 512 GB per stack, for holding model weights.

Capacity-first architecture; memory arbitrage. A design that trades scarce bandwidth for abundant capacity, using LPDDR or HBF where the incumbent uses HBM.

ISA; RISC-V. The instruction-set architecture is the contract between software and processor—the layer at which Arm and x86 collect licence rents; RISC-V is an open ISA with no licence to revoke.

CUDA; NVLink. NVIDIA’s proprietary programming stack and chip-to-chip interconnect; the software and interconnect moats of the accelerator incumbent.

Ascend, Cambricon, Kunlun. Chinese domestic AI accelerator lines (Huawei, Cambricon, Baidu).

SMIC; CXMT. China’s leading logic foundry and its leading DRAM maker.

TOP500; HPL; HPCG. The semiannual supercomputer ranking and its two benchmarks; “MLPerf” is the industry’s AI training and inference benchmark suite.

BGPU (baseline GPU). The full edition’s vendor-neutral unit of AI compute demand, defined by serving throughput at 8 TB/s of memory bandwidth.

P1–P4. The book’s four propositions: adoption, sovereign compute, rent destruction, financial repricing (Chapter 1).

Evidence grades. [V] verified, [P] primary-source claim, [A] analyst estimate, [M] author-modeled, [R] roadmap, [S] scenario judgment.

The ratchet. The competitive equilibrium in which, once one Chinese laboratory opens at frontier scale, the others cannot profitably stay closed.

1 The Thesis

In one paragraph

Between roughly 2005 and 2025 China fought and won three complete technology wars—solar photovoltaics, lithium-ion batteries, electric vehicles—by executing a single, learnable playbook, and artificial intelligence is its fourth run. The mechanism is not to match the West chokepoint for chokepoint but to engineer the chokepoints out of the system: open-weight frontier models at marginal prices, an open hardware ecosystem on an unrevocable instruction set, a memory-capacity arbitrage substituting commodity DRAM for scarce HBM, and a domestic semiconductor base. The connective tissue is openness deployed as a weapon against the layers where Western incumbents collect their rents. The argument is decomposed into four propositions, graded and falsified separately.

1.1 Three wars, one playbook

In solar photovoltaics China went from assembler of imported cells to more than 90% of every stage of the global supply chain, at module prices down more than 99% from their 1976 level [1, 2]. In lithium-ion batteries—commercialized by Sony in 1991 and led for two decades by Japan and Korea—Chinese firms now hold roughly 69–75% of global EV battery installations and over 80% of cell manufacturing capacity [3, 4]. In electric vehicles China became at once the largest car market, the largest car exporter, and home to the largest EV maker [5, 6]. Chapter 2 gives the campaigns in detail.

These were not three lucky outcomes but three executions of one playbook: patient state designation of a strategic industry; demand creation at home; subsidized scale-up far ahead of demand; brutal domestic competition as an evolutionary selection engine; capture of the entire upstream supply chain; a cost-down learning-curve war that bankrupts foreign incumbents; and finally the conversion of dominance into geopolitical leverage—export controls on the very technologies China was once accused of copying [7, 8].

1.2 Why AI is the fourth war

The prevailing Western assumption is that AI is different: that frontier AI depends on scarce, controllable inputs—leading-edge logic, high-bandwidth memory, proprietary interconnects, closed model weights—and that export controls on these chokepoints can durably preserve a lead. The book argues that it is failing on all fronts simultaneously, and failing in a characteristic way: not because China matches the West chokepoint for chokepoint, but because it is engineering the chokepoints out of the system.

The model. Chinese laboratories have converted the model layer into a commons: open weights on the top open positions of the independent indices, a majority of routed tokens, prices a small fraction of the closed frontier’s, and a measured capability gap compressed to low single digits [9, 10]. Open weights at marginal cost do to the model layer what Chinese modules did to solar: they set a world price that rent-seeking business models cannot survive.

The compute. The state has designated the open RISC-V instruction set as national infrastructure, is standardizing matrix extensions and chip-to-chip interconnects across vendors, and has open-sourced the software stacks that would otherwise fragment the domestic ecosystem [11, 12]. In parallel, a memory-capacity arbitrage—terabytes of cheap commodity DRAM

for scarce HBM—attacks the assumption that frontier training requires centralized, HBM-bound superclusters at all.

The semiconductor. Beneath both, the fabs: 5 nm-class logic without EUV, the world’s fourth DRAM producer, three domestic toolmakers in the global top twenty [13, 14]. The remaining gaps—EUV lithography, leading-edge HBM—are real, but the strategy does not require closing them on the West’s terms; it requires making them irrelevant to the chosen architecture.

1.3 Openness as a weapon

The connective tissue is a deliberate inversion of the Western technology-rent model. Where the United States concentrated AI value in proprietary layers—CUDA, NVLink, closed frontier weights, x86/Arm licensing—China’s strategy, formalized in the July 2025 Global AI Governance Action Plan and the “AI+” State Council initiative, is to make every one of them open, standardized and free [15, 16]. This is not altruism; it is the same move as pricing solar modules at \$0.09 per watt. An open layer generates no rent for the incumbent, but it generates *adoption*, and adoption generates the scale, the ecosystem and the standards-setting power that constitute victory in the precedent industries. The two strategies are not symmetric opponents: one defends margins, the other destroys them.

1.4 Four propositions, stated separately

Rigor requires decomposing the thesis into four propositions that are related but logically independent—none proves the next, and each has its own evidence and its own failure modes.

P1 — Adoption. China is winning open-weight model adoption. *Confirmed by* download, derivative, and routed-token majorities persisting and extending into enterprise spend. *Falsified by* a durable reversal of the router and Hugging Face shares, or enterprise open-weight spend stagnating below about 15% while usage plateaus. Current grade: strongly supported [V].

P2 — Sovereign compute. China is constructing a domestically sufficient training-and-serving stack. *Confirmed by* a frontier-class model trained end-to-end on domestic silicon, packaging, and memory. *Falsified by* persistent dependence on stockpiled or smuggled foreign accelerators and HBM past about 2028. Current grade: rapid progress, not yet achieved [P/A/M].

P3 — Rent destruction. Open-weight commoditization will compress the closed-frontier rents that finance Western laboratories. *Confirmed by* closed-model revenue growth decelerating while capability parity holds. *Falsified by* closed labs sustaining pricing power through service, distribution, and trust layers even at model parity. Current grade: mechanism operating, magnitude contested [V/S].

P4 — Financial repricing. The compression of P3, if it occurs, punctures the US AI capital structure. *Confirmed by* the indicator table of Chapter 7. *Falsified by* demand elasticity absorbing the price decline: usage growing faster than prices fall rescues the buildout even under P1–P3. Current grade: scenario judgment [S].

They are coupled but separable. P1 can hold while P2 fails (China dominating distribution on stockpiled foreign silicon); P3 can hold while P4 fails (margins compress while volume growth fills the data centers); and a US valuation correction could arrive from macro causes with none of P1–P3 mattering. The convergence argument is that the *coupling* of $P1 \times P2$ —open

weights meeting sovereign, capacity-first hardware—is what forces P3, with P4 as its priced-in consequence; Chapter 13 states where each proposition stands at the currency date.

1.5 Evidence grades

Load-bearing claims are graded throughout, on the scale the front matter defines: [V] verified, [P] primary-source claim—reliable as to what was claimed, not what was achieved—[A] analyst estimate, [M] author-modeled, [R] roadmap, [S] scenario judgment. The full edition collects every headline claim in a graded ledger and a machine-readable claim register [full edition, App. D, J]; Appendix A reproduces the forecasts and the quarterly dashboard.

1.6 The theory, in brief

The full edition devotes a chapter to the machinery behind the analogy [full edition, Ch. 2]; four ideas from it are needed to read what follows.

A cost stack, not a price. The cost of one completed useful task decomposes into amortized model creation, inference, integration and organizational adoption, complementary capital, and assurance. Only inference is falling at ten times per year; integration and adoption are not—which is why open-weight *usage* share rose while open-weight *spend* share fell—and assurance is a cost the commons has mostly not paid. Commoditizing inference therefore does not commoditize delivered work; it shifts the binding constraint to the terms that are not falling, so the rent-destruction claim is bounded to the model layer and remains a hypothesis at the service layer.

The elasticity knife-edge. With own-price elasticity ε , revenue moves as $\Delta \log R = (1 - \varepsilon) \Delta \log P$: above one, price cuts pay for themselves; below one, falling prices shrink revenue; $\varepsilon = 1$ is the boundary between the bull and bear cases for the entire Western AI complex. The honest position is that ε is a portfolio, not a number—plausibly above one for coding agents, synthetic data and scientific inference, near or below one for consumer chat—and the leveraged Western structure needs it above one at the premium tier specifically. Graceful degradation is a real Chinese advantage; it is not a theorem of victory.

Openness relocates rent; it does not abolish it. A seller rationally prices at or below marginal cost whenever the extra volume raises revenue on a complement by more than the margin given up. Each side opens the layer where it is behind and expects rent in the layer it is best at: China opens weights, ISA, interconnect and framework and expects to recapture on hardware volume, domestic cloud, standards and dependence; the United States controls all four and expects premium capability, distribution, integration, trusted deployment and agent platforms. Two flywheels result, and each can steal the other’s critical input—China steals volume, America steals value—and which theft dominates is the single most important unresolved empirical question in the field.

Four layers of winning. AI is a general-purpose technology, so a country can dominate its production and capture little of its value, because most value appears downstream. Production leadership is the only layer on which the United States is unambiguously ahead; the platform layer goes to whoever opens fastest; diffusion is where most realized value sits; productivity is the only layer that finally matters and the least observable in real time.

The thesis is therefore conditional, and Chapter 13 states it in its final form with the evidence attached: whether Chinese cost reduction, adoption, standards formation and hardware sufficiency compound faster than US capability, distribution and monetization. This book argues that convergence is the way to bet.

2 The Precedents and the Playbook

In one paragraph

Solar showed the pure form [full edition, Ch. 3–8]: Western invention and demand, Chinese policy-financed scale, capture of the learning curve, failed tariffs, over 90% share at every upstream stage, then state cartelization. Batteries added vertical depth from mine to pack, won by regulatory market reservation and by industrializing the “inferior” chemistry until the West licensed it back. Electric vehicles proved the playbook works on a whole system. Telecommunications, a fourth and lesser precedent, shows what it does with a financed network—and supplies the template the sovereign-AI-export story of Chapter 8 is borrowed from. From the campaigns the book distills nine tactical moves and six strategic steps; from those that have *not* worked, where it grinds: tacit knowledge, slow clock speed, foreign certification, installed-base lock, trust. AI’s model and compute layers sit in the kill zone; its deployment layer is genuinely contested.

2.1 Three campaigns, in brief

Solar was invented at Bell Labs and built on American, German, Japanese and Australian science; Sharp was the largest cell producer in 2006, and demand was created by Germany’s feed-in tariff. Chinese scale-up rode that demand and about \$47 billion of development-bank credit [17]; polysilicon fell from about \$475 to under \$16 per kilogram between 2008 and 2012 [18], and Solyndra, Q-Cells and SolarWorld died—as did the Chinese pioneer Suntech, because the state protects the industry, not the champion. US and EU duties from 2012 taxed deployment without resurrecting manufacturing, and a domestic market of 315 GW_{AC} in 2025 alone made foreign trade policy irrelevant [19]. The end state is over 90% Chinese share at every stage upstream of the module [2], module prices down from \$106 per watt in 1976 to about \$0.09 at a learning rate near 20% per doubling [1, 20]—and since roughly 2010 essentially all doublings were Chinese, so whoever manufactures the doublings owns the learning. Winning is expensive: losses below cash cost, then Xi’s July 2025 anti-involution campaign and a polysilicon buyout platform, an OPEC for solar [21, 22]. Four lessons: trade barriers after cost-curve capture protect only domestic prices; champions are disposable, the industrial commons is not; a home market absorbing half of world demand makes foreign trade policy irrelevant; and commoditization ends in state-managed consolidation. Keep them in mind when reading about open-weight models priced near zero.

Batteries were commercialized by Sony in 1991 and led by Japan and Korea for two decades; China entered through low-margin consumer cells and pivoted when the state created demand. The decisive instrument was a certification, not a tariff: the 2015 white list of 57 domestic makers eligible for NEV subsidies, which excluded LG and Samsung and was scrapped in 2019 when no longer needed [23]. By 2025 Chinese makers held about 69% of installations and over 80% of the world’s 4 TWh of cell capacity [3, 24], plus 70–90% of most upstream layers and two thirds of lithium and cobalt refining, at pack costs that leave North America and Europe paying a 44–56% premium [4]: a gigafactory built outside China is built, in effect, inside the Chinese supply chain. Lithium iron phosphate—Western-invented, dismissed as a dead end—passed 55% of the global EV market and is licensed back to Ford and Stellantis, the inferior branch industrialized at scale beating the premium branch on system economics, a preview of the LPDDR-versus-HBM argument of Chapter 4. In 2025 the infant industry became the gatekeeper with export controls on battery technology [7]. Three more lessons: regulatory reservation beats tariffs and is invisible

to trade law; vertical integration to the mine makes a foreign factory a hostage; watch what China industrializes, not what the West benchmarks.

Electric vehicles are a whole system wrapped in a century of incumbency; if the playbook worked here, no consumer-industrial market is safe, and it worked. Wan Gang’s leapfrog decision predates Tesla; at least \$231 billion of layered demand support [25], a dual-credit mandate binding foreign joint ventures, and 19 million charging points took NEV penetration from 6% in 2020 to 54% in 2025 [26] and made China the largest car exporter [6]. BYD ran the flywheel from 416,000 to 4.55 million vehicles with about 80% of Tier-1 content in-house and a \$4,700 structural advantage over Tesla of which subsidy is under \$300 [5], while some 400 Chinese EV ventures died around it: the state subsidizes the species, not the specimen. Foreign brands fell from 62% to 35% of the Chinese market and technology now flows backward. And the car became a software platform in Chinese hands—about 95% of automotive lidar, driver assistance free on an \$8,000 car, 20–24-month cycles against 40–50—so the first mass-market safety-critical embedded-AI category is Chinese-led, and it anchors the domestic chips, sensors and models of the next chapter. Three final lessons: pick the transition where incumbency resets to zero and force it at home; a system-level win pulls the whole component stack with it; speed compounds.

Telecommunications, in lesser detail, is the precedent closest in structure to what an “AI stack” export would be—infrastructure, financed, standardized, locked in by service—and one company ran almost all of the playbook on it. Huawei held a fifth of the Chinese switch market in 1996 and reached parity with Ericsson, the incumbent it had been imitating, by 2011 [27]; its share of world telecom-equipment revenue ran 27.7% in 2018, 31% in the first half of 2020, the year the sanctions bit hardest, and **31% again in the first half of 2025**, 40% outside North America [28]. The instrument that distinguishes the campaign, and that AI has not reproduced, was credit: a China Development Bank facility of \$10 billion for Huawei’s customers in 2004, raised to **\$30 billion in December 2009**, inside up to \$75 billion of state support by the *Wall Street Journal*’s reconstruction, buying prices about 30% below rivals’ where the credit was and not elsewhere [29]—an operator in Lagos was choosing between a switch it could finance and one it could not. Entry ran rural-first, then Africa, Asia and Latin America, then Europe; the endgame was standards from inside, first in declared 5G patent families and in 3GPP contributions [30]. The bans—US Entity List 2019, Britain’s reversal 2020, an EU phase-out proposal in 2026—removed share only where security agencies held a veto and operators could swap at no cost, the Five Eyes, the Nordics and Japan: Chinese vendors held about 32% of European 5G sites in 2022 and about 32% at the end of 2024, and what actually turned Britain was the May 2020 *foundry* rule, a semiconductor chokepoint, not a network-security finding [31]. The disanalogy, tested in Chapter 8: AI has no credit line, no surplus capacity, and an artifact—open weights—that diffuses better than any switch and locks in nothing [full edition, Ch. 6].

2.2 The playbook distilled

Table 2.1 inventories nine tactical moves that recur across the campaigns; the AI column is developed in Chapter 4. Six strategic steps are the sequence the moves serve. *Target selection*: a discontinuity where the incumbent’s accumulated advantage is about to be stranded and where manufacturing learning-by-doing, not protected IP, is the accumulating asset. *Demand before supply*: the home market as guaranteed demand, foreign firms admitted just long enough to transfer capability. *Overcapacity as strategy*: capacity at multiples of world demand, which Western economics reads as waste and which purchases the entire learning curve; the losses land on Chinese balance sheets and dead firms while the cost curve remains. *Selection, not planning*: Beijing floods the sector and lets domestic combat select, its loyalty being to the industrial commons—the engineer pool, the supplier mesh, the tooling ecology. *Commoditize the adversary’s rent layer*: wherever the Western business model collects rent—module, pack, badge,

API token, ISA licence—price it at marginal cost, because rents fund the West’s R&D flywheel, and openness is this move in its purest form. *Endgame*: consolidate, stabilize prices upward, and convert dominance into coercive leverage.

Table 2.1: The nine tactical moves, across four campaigns; telecommunications is scored against the six steps in the full edition rather than tabulated here.

Move	Solar	Batteries	EVs	AI
M1 State designation	SEI 2010, 12th FYP	White list 2015	863 program, SEI 2010	“AI+” plan, 15th FYP, RISC-V guidance
M2 Home demand	National FiT, Top Runner	NEV subsidy linkage	\$231B support, dual credit	State DCs mandated on domestic chips, subsidized power
M3 Subsidized scale-up	CDB \$47B credit	Provincial capital	Provincial EV funds	Big Fund III (\$47.5B), SMIC/CXMT capex
M4 Darwinian competition	2018 “531” cull	100+ makers to consolidation	~400 dead EV firms	LLM price war since 2024; chip-vendor shakeout beginning
M5 Supply-chain capture	poly → module	mine → cell	battery → chip → magnet → ship	model → framework → chip → interconnect → fab → DRAM
M6 Commoditize the rent layer	module \$/W	pack \$/kWh	vehicle BOM	open weights, \$/Mtoken, open ISA and stack
M7 Absorb the “inferior” branch	mono-PERC, TOPCon	LFP	small BEVs	sparse MoE; LPDDR capacity vs HBM bandwidth
M8 Standards capture	efficiency floors	GB safety standards	charging standards	RISC-V profiles, matrix ISA, UnifiedBus
M9 Leverage after victory	OPEC-for-poly	tech export controls 2025	rare-earth controls 2025	(pending)

Against this the West has fielded two counter-patterns, both failing. Tariff walls arrive after cost-curve capture and merely tax domestic consumers while the walled market shrinks in relevance. Chokepoint denial—semiconductor export controls since 2019—is the more serious strategy, and its record is a pattern the book calls *chokepoint decay*: each denied input (EUV, HBM, H100s, EDA) triggers a state-scale substitution program plus an architectural detour around the input, while the denial hands the domestic substitute a guaranteed market. Export controls are, functionally, a US-funded market-reservation policy for Huawei, SMIC, CXMT and Cambricon.

2.3 Where the analogy breaks

A playbook argued from three victories owes the reader its defeats. The same state, spending an estimated \$248 billion a year on industrial policy [32], has run two-decade campaigns where dominance has not arrived: commercial aircraft (COMAC’s C919 delivered about thirteen aircraft in 2025 against a plan of 75, on a Western engine whose export licences Washington suspended and restored—chokepoint leverage working); lithography and the deep tool chain (under 5% localized after twenty years and three Big Funds [33, 34]); precision machine tools (five-axis CNC about 8% domestic against an 80% target); operating systems, where HarmonyOS reached 17% only after sanctions burned the Android bridge—ecosystems yield on decade timescales, under existential forcing, inside the home market; and biotech, where \$138 billion of 2025 out-licensing meets a residual gap that is entirely trust and regulation, the FDA’s 14–1 rejection of China-only trial data being the certification barrier in its purest form [35]. The playbook wins where products are modular, iteration is fast, learning is manufacturable and the home market can certify; it grinds where knowledge is tacit, certification is foreign, and value lives in an installed ecosystem.

Scored against that model, AI models are the most modular, fastest-clock-speed artifact in

industrial history—sparse-MoE training is closer to batteries than to jet engines—and compute is intermediate, with the tool chain inheriting lithography’s resistance and being routed around. But three properties of AI sit in the resisting column and are the strongest structural arguments against this book’s title: the frontier moves, and recursive AI-for-AI research could reopen the gap faster than any factory closes it; ecosystems lock, and the enterprise deployment layer is Office, not modules; and trust and certification bind, with the sintilimab precedent—capability accepted, certification denied—directly available to Western regulators. Two asymmetries follow: zero reproduction cost lets Chinese weights conquer adoption instantly and would let a superior Western open release reconquer it just as fast, so there is no factory-shaped moat around P1; and talent is the input the two systems still share—researchers of Chinese undergraduate origin were 47% of the top tier in 2022, up from 29% in 2019, while the share of that tier working in the United States fell from 59% to 42% and China’s rose from 11% to 28% [36]. Model and compute therefore sit mostly in the kill zone, the tool chain in the resistance zone and bypassed, and the deployment layer—trust, certification, ecosystems—is contested ground where the precedents predict a long grind, not a collapse.

3 The Strategy in Whole

In one paragraph

The Chinese AI strategy is the playbook of Chapter 2 run on a fourth industry, with one adaptation that makes it faster: the rent layer under attack is digital, so the commoditizing move—openness—costs nothing to reproduce and propagates at the speed of a download. Placed on the playbook—a map drawn before the territory—every one of the nine moves is occupied at a recognizable stage: one of the six steps is complete, one complete at home and contested abroad, three trending, and the endgame barely begun. The three properties on which the playbook has failed before—the moving frontier, ecosystem lock, trust—are exactly the ones the elements cannot settle from inside; and Chapter 5 adds that the moves catalogued here are not, in 2026, uniquely Chinese. [full edition, Ch. 9]

3.1 Where each element sits

Target selection chose the discontinuity where the incumbent’s advantage is about to be stranded: HBM-centric superclusters, closed frontier weights, a licensed instruction set. *Demand before supply* came by mandate rather than feed-in tariff—state-funded data centers restricted to domestic silicon, western hub electricity discounted only for approved domestic chips [37, 38]. *Overcapacity as strategy* is under way in DRAM and compute halls, with state-subsidized compute at 20–30% utilization against a 60% target and the losses landing where they landed in solar; but the accelerator layer is still supply-bound by HBM, and the model-layer price war runs without a visible rent-recapture plan, which Chapter 9 calls the strategy’s soft underbelly. *Selection, not planning* is visible in a model layer where ten laboratories release at frontier scale. *Commoditizing the adversary’s rent layer* is what Chapter 4 documents: open weights, harness, ISA, interconnect and framework specifications at marginal or zero price, aimed at CUDA, NVLink, x86/Arm licensing and the metered API [15, 11]. And the *endgame* has begun on consolidation—the anti-involution campaign—not yet on gatekeeping, where batteries and rare earths supply the template [7]. Table 3.1 places each element.

Three interlocks are the strategy. The model and the compute are the coupling $P1 \times P2$, which turns a download into a deployment. The watt is the weld: a tariff instrument paying only

Table 3.1: The elements placed on the playbook, with the status Chapter 4 establishes for each, as of 9 September 2026.

Element	Move(s) it executes	Open question	Status
The model	M1 designation; M6 commoditize the rent layer; M4; M7 sparse MoE; M8	Spend share, and where the licence terms relocate the rent	Done (distribution); trending (spend, rent)
The harness	M6 one layer up; de-facto default sought outside the standards bodies (M8 not achieved)	Switching cost, never measured	Trending [S]
The compute	M1 (RISC-V guidance); M7 “inferior” branch; M8; M5	Capacity-first arbitrage modeled, not realized [M]	Trending; existence proof done
The semiconductor	M5 supply-chain capture; M3	Domestic HBM, the binding constraint	Trending; HBM lagging 3–4 years
The watt	M2 home demand (subsidized power, domestic-chip tariff); the weld	Whether the tariff outlives the surplus	Done (headroom); trending (durability)
The data	Where rent migrates once M6 succeeds	Evaluation and incident data unowned	Contested
Embodied AI	M5 before the category exists	Useful work per robot-hour, published by nobody	Trending
Endgame	M9 leverage after victory	Gatekeeping not yet exercised on AI	Pending

for domestic silicon converts an energy surplus into hardware volume and hardware volume into tokens priced below the Western marginal cost. And the semiconductor is the floor, incomplete at exactly one layer—HBM—which is why the memory arbitrage exists at all: the architecture was chosen so that the one input the chokepoint strategy still holds is the one it needs least.

3.2 What the assembly licenses

It licenses the claim that the AI campaign is a coherent execution of the strategy that won three technology wars, further along than a layer-by-layer reading suggests, its remaining gaps engineered around rather than closed on the West’s terms. It forbids the claim that the outcome is settled: rent destruction has run furthest and its completion is contested, the integration test that would move sovereign compute from trending to done has not been passed, and Chapter 2’s three resisting properties are the three the deployment layer is made of.

4 The Elements of AI

In one paragraph

Read separately, each element has a Western rebuttal: the models trail the closed frontier by eight points on the re-based index, the accelerators trail NVIDIA per watt, the fabs trail TSMC by two nodes. Read together, the rebuttals miss the structure. The model layer is a commons held open by a self-enforcing ratchet, winning volume on a sparse open tier while spend stays in the closed frontier and the commons' makers negotiate a cut of the clouds that serve it; the harness is being given away one layer up; a GPU-free Chinese machine leads both HPL and HPCG, while a modeled memory arbitrage and a priced five-thousand-dollar appliance [M/S] attack the premise of the centralized supercluster; the chokepoints on Chinese silicon decay faster than they are replaced, with domestic HBM the one binding constraint; the watt is the deepest structural asymmetry; and deployment-generated data and embodied AI are where the rent goes once the model is free. None of it is finished, every claim is graded, and the integration test not yet passed—a frontier model trained end to end on domestic silicon—is the “Mate 60 moment” the forecasts price at 65% within two to four years.

4.1 The model: open weights as the commodity layer

The Western frontier is closed: the best US models are metered APIs, their weights corporate secrets, their pricing a rent on capability. The Chinese frontier is open: downloadable weights, mostly MIT or Apache, at API prices near marginal inference cost. The inversion began with DeepSeek, whose V2 of May 2024 triggered a domestic price war in which four rivals cut prices by up to 97% within two weeks; V3 reached GPT-4-class capability at a published marginal training cost of \$5.6 million—a final-run figure [P], not a program cost, with buildout estimates up to \$1.6 billion [A] [39]; and R1, MIT-licensed, matched o1-class reasoning and erased \$589 billion of NVIDIA's market value in a day [40]. The strategic content was mispriced by markets and correctly priced by Beijing: frontier capability can be commoditized; constrained hardware forced algorithmic efficiency—latent attention, sparsity, FP8, GRPO, multi-token prediction—that compounds on any hardware; and open weights convert capability into global adoption instantly. The defensible causal claim is constraint-consistency, not unique causation [full edition, Ch. 10].

The state did not anoint DeepSeek; it flooded the sector. By mid-2026 the release lag from the US closed frontier to a Chinese open equivalent had compressed from six to eight months to weeks, and the breadth of the field is itself the datum: Alibaba's Qwen family has a billion cumulative downloads, more than half of all open-model downloads [41], and Moonshot's Kimi K3 is at 2.8 trillion parameters the largest open-weight model ever released and the only open model in the global top five [42]. Token prices run from one seventh to one hundred-and-fiftieth of US equivalents; Chinese models account for about 41% of Hugging Face downloads [9] and roughly 61% of tokens routed through the largest neutral aggregator, up from about 2% in early 2025 [43]; roughly 80% of open-stack startups use a Chinese model [44]. The Stanford AI Index put the top US–China gap at 2.7% in March 2026, down from 17–32% in 2023, and the closed-over-open gap at 3.3%—a 23:1 private-spending ratio buying a 2.7% measured gap [10].

Two developments in the summer of 2026 changed the character of the commons. The first is the *open-weight ratchet*. Alibaba's Max tier had always been proprietary; seven days after Moonshot opened Kimi K3, Alibaba committed to opening Qwen3.8-Max and shipped the weights on 13 August [45]. Once one Chinese laboratory opens at frontier scale the others cannot stay closed—adoption is the currency, and closing your best model is unilateral disarmament—which upgrades openness from a state policy that could reverse into a competitive equilibrium from which no participant profits by defecting. But the ratchet holds on weights and fails on terms: Qwen3.8-Max requires a commercial agreement above US\$50 million of revenue and Kimi K3 above US\$20 million [46], while GLM-5.3 shipped on 28 August under a custom licence—neither MIT nor Apache—whose reported terms require a *security review* of model-as-a-service operators

above \$10 billion of trailing revenue [47] [P/A on the threshold]. Two firms ask for a commercial agreement, the third for review; either way rent relocates out of the weights and into the terms of large-scale service. The second development is that the ratchet now runs on both sides of the Pacific: Meta returned to open weights on 10 August, and—the fact that reframes the whole element—the largest Western open-weight supplier by routed volume is now a chip company (Chapter 5).

The sharpest instance of what the commons now delivers arrived two days before the flagship’s weights did. On 26 August Zhipu released **GLM-5.3-Flash**—320B total parameters with 18B active, natively multimodal, a one-million-token context, under **MIT**, at \$0.15/\$0.50 per million tokens—having run publicly as the unattributed stealth model “Ox Alpha” [48] [P]. The launch-week numbers, on the Artificial Analysis index then current (v4.1.1), were GLM-5.3 at 60 and Flash at 57, second and fourth of 111, against 63 for the best closed model [49] [A]; those are the numbers the shares moved on, and they no longer exist. The re-based v4.3 board of 7 September puts Claude Fable 5.1 and GPT-6 Astra at 53, Claude Opus 5 at 51, GPT-5.6 Sol at 47, **GLM-5.3 at 45, GLM-5.3-Flash at 42**, GPT-5.6 Luna at 38 and DeepSeek V4 Flash at 35 [50] [A]. The open–closed gap at the very top is therefore *eight index points* on v4.3, three on the launch-week board, and the difference is entirely the instrument—consistent in direction with the Stanford Index’s 2.7%, measured on a different scale, and not to be read into it. What Flash is actually good at is narrower than the composite. On Terminal-Bench v4.0, the agentic-coding component whose introduction drove the re-basing, it scores 33% against 12% for both Luna and DeepSeek V4 Flash, the models it competes with on price, and 60 against Luna’s 50 on AutomationBench; on everything else it loses to the same cheap closed model—61–71 output tokens per second against 120, weaker on document reasoning, and \$0.25 per index task against Luna’s \$0.18—and no SWE-bench, LiveCodeBench or Aider result for it exists from any source [51] [A]. It is the cheapest point on the price–agentic-coding frontier, not a flagship replacement. And it fits on one machine: quantized to three bits it is 120 GB at 81.6% top-1 retention on the quantizer’s own harness, a single 128 GB unified-memory box that is workstation-class, not a laptop [52] [P].

The domestic-silicon claim attached to it is real, narrow, and easy to overstate. Z.ai says the model had run “entirely on Chinese AI chips” [53] [P]. Four qualifications. The claim is about *serving, not training*, and Z.ai’s own Chinese-language documentation says exactly that—the first time large-scale traffic has been served on a domestic cluster—and nothing about training [P]. No chipmaker is named. It is not a first, since Zhipu claimed Ascend-only training for an image model of undisclosed scale in January [54] [P, unaudited]. And the June report that GLM-5 was trained on roughly 100,000 Ascend 910B, the claim the trade press repeats, has *no locatable primary*: it circulates without a cited source, the GLM-5 repository carries no hardware statement, and Huawei’s own February announcement describes “0-day” *inference* support, adaptation after the fact [55]; this book grades it press-derived [A/C] and stacks nothing on it, against reporting that DeepSeek’s R1 training on Ascend failed and reverted to NVIDIA. What is new and substantial is production-scale inference on a reported cluster of roughly 100,000 domestic accelerators carrying a workload that was top-four in the world in its launch week [53] [A]. The integration test—a frontier model trained end to end on domestic silicon, packaging and memory—has come closest on the serving half, is still not passed, and the evidence that it has been is thinner at 9 September 2026 than the June headlines made it look.

Three qualifications keep the element honest. Capability, distribution and revenue are different scoreboards and the third lags: enterprise open-weight *spend* share fell from 19% in late 2024 to about 11% in 2026 even as usage rose [56], the strongest evidence against P3. Alibaba’s claim that Qwen3.8-Max ran “500-plus turns of chip design optimization” toward an autonomous silicon design flow—if true, an attack on the one input export control cannot embargo at a border—is unverified. And the right denominator is cost per solved task, not price per token, because reasoning outputs are inflating about fivefold a year [57] and no vendor publishes the number that matters.

Then the price inversion, which breaks the simple story in both directions. On **30 July 2026** OpenAI cut GPT-5.6 Luna, the *cheapest* tier of its family, by 80%—from \$1.00/\$6.00 to \$0.20/\$1.20 per million tokens. Terra fell 20%. *Sol*, the *flagship*, was not cut at all, and went to \$4/\$20 only on 21 August, explicitly “for the next 3 months”—a promotion, which defends share, rather than a repricing, which concedes cost [58] [P]. At \$0.20/\$1.20 Luna sits *below the list price of most Chinese open-weight models*, and the cut worked—usage rose to 13.8× its pre-cut average against 1.1× for the undiscounted Sol, roughly 75% of the gain taken from competitors rather than cannibalized [59, 60] [A]. Then the datum that cuts the other way: on 13 August, mid-cascade, *DeepSeek* raised its API prices by as much as 1,100% [61] [P]. The deepest Western cut came a fortnight *before* the Chinese increase, so no one-directional price-war story survives, and every company statement cites efficiency rather than competition. Commoditization at the volume end is real and observable; the arrow of causation the popular version assumes is not. Nor does the inversion recover the commons: undercutting a hosted Chinese endpoint wins tokens without winning weights.

The barbell: where the tokens go and where the money stays. On OpenRouter, with data through 8 September, the four largest models by routed tokens were Tencent’s Hy4 at 18.7 trillion, Luna at 14.5, DeepSeek V4 Flash at 12.4 and GLM-5.3-Flash at 12.2 trillion; where the active-parameter count is disclosed the top ten runs 13B–49B, so the volume tier is a sparse tier, and Claude Fable 5, of the family that tops v4.3, sat in “the low billions” [62] [A]. On Vercel’s gateway open-weight models took 36% of July’s tokens on **8.6% of spend**, and 62% of tokens on one August day, while the four leading frontier laboratories took 95% of spend and Anthropic alone 61–65% [63] [A]. Volume moves to the open sparse tier; spend stays in the closed frontier; and the counter is in the same data—spend concentration is evidence that the capability tier still commands a rent from buyers who can measure the difference [64].

The toll road. Where the rent relocates now has a named counterparty. Reuters reported on 26 August that Moonshot is in early-stage talks with Microsoft, Amazon and Google to host Kimi K3 for up to 30% of K3 service revenue—split, data access and token auditing unresolved, nothing signed [65] [A]. The lever is the K3 licence clause obliging any model-as-a-service operator above \$20 million of revenue to contract with Moonshot, which is why K3, unlike K2.5, is not on Bedrock [66] [P]; Alibaba is copying the instrument for Qwen3.8-Max [67] [A]. The baseline makes it startling: the American clouds already serve DeepSeek, Qwen, GLM and K2.5 as managed models “sold by AWS,” at a markup, with no disclosed payment to any laboratory [68] [P]. If it closes it will be the first time the maker of the commons takes a cut on a Western cloud—rent out of the weights, into the terms of service. The “hedge” reading is rejected on the evidence: no source gives that motive, and *with revenue comes regulability*—a contract converts an ambiguous download into an unambiguous transaction, easier for a US restriction to reach, not harder [69] [A]; NDAA §1532 already bars covered AI at the Department of Defense without distinguishing weights from tokens [70] [P]; and on the Chinese side MOFCOM is weighing curbs on public release of the most advanced models, undecided [71] [A]. In playbook terms the model layer is in the late-war stage: rent destroyed at the volume end and adoption captured, plainly not destroyed at the capability end, and the commons acquiring a cash register exactly where the theory said it would. What open weights cannot guarantee is sovereign hardware to train and run them [full edition, Ch. 10].

4.2 The harness: the layer the contest moved to

Every preceding paragraph treated the model as the object of competition; the object is moving. Models are swappable—an API endpoint exchanged by configuration—while a *harness* is what an organization writes code against: its abstractions become the integration, its execution semantics the test assumptions, its permission model what an auditor reads. The layer where switching is expensive is not the layer everyone measures, and a strategy that wins the harness can lose

every benchmark and still take the ecosystem. The sole controlled study found that adopting a better harness with *any* model saved 33–61% of cost per task, against 36% for switching from the cheapest model to the most expensive [72]; once a self-hosted open model runs at about a dollar per million output tokens against \$25–30 for the closed frontier, the harness is close to the only lever left on production economics.

On 13 August 2026 DeepSeek released an MIT-licensed, plugin-native, model-agnostic harness that passed 100,000 GitHub stars in about 48 hours [73, 74], built on the kernel of a four-year-old Chinese chatbot framework whose creator it recruited—ecosystem acquisition, this book’s thesis executed one layer up—and shipped it the same day it *raised* its API prices [75]: commoditize the complement, monetize the scarce layer. The next day Xiaohongshu shipped an Apache-licensed omni-modal agent model trained for ten-hour agent runs, so model and harness were co-designed in the same week by two Chinese labs; meanwhile the Western agent-standards bodies have no Chinese named member and Beijing’s cyberspace administration has published a rival “interoperability” initiative naming none of them [76]. The incumbents convene a consortium; the challenger ships a permissively licensed implementation and bids for the installed base directly. The strong lock-in claim is a hypothesis [S]—nobody has measured switching cost—but lock-in relocates rather than disappears. If American models remain the best and the world’s agents are orchestrated by runtimes shipped from Hangzhou and Beijing under MIT licences, superiority at the model layer buys less than its owners expect [full edition, Ch. 11].

4.3 The compute: RISC-V, accelerators, and the memory arbitrage

The Western AI machine is a proprietary stack—x86 or Arm hosts, NVIDIA accelerators, CUDA, NVLink, HBM—each layer a licensed chokepoint. China’s counter-design is an open machine: an open instruction set, multi-vendor domestic accelerators unified by open interconnect and open software, and a memory system trading scarce bandwidth for abundant capacity. The claim is narrow—not that the GPU is obsolete, but that the CPU corner of a design continuum can reach parity on the workloads that matter while carrying an ecosystem advantage the accelerator corner does not have.

RISC-V. National RISC-V policy guidance was reported in March 2025 from eight government bodies—anonymous-sourced reporting of a document never shown to have been issued [11] [A]—and the 15th Five-Year Plan mandates “extraordinary measures” for self-reliance with the open ISA as its core mechanism [77], because an open standard insulates Chinese silicon design from sanctions at the ISA layer permanently. Washington’s attempt to restrict participation was abandoned once the Commerce Department conceded it would harm US firms without slowing China, and twelve of the twenty-four premier members of RISC-V International are Chinese [78]. Three disclosures of August 2026 located the position exactly, and two are unfavourable. NVIDIA published what a RISC-V CPU must be to host CUDA as a third ISA, and almost none of the requirement is the instruction set—but a host wanting *coherent unified memory* must integrate NVIDIA’s NVLink-C2C soft IP and speak CHI, an Arm-originated protocol under a partners-only specification, because RISC-V has no cache-coherent chip-to-chip interface of its own [79] [P]. The rent did not disappear when the ISA opened; it moved from the instruction set to the coherent interconnect. Second, *no matrix extension has been ratified at all*: four competing proposals for one function, with targets spanning August 2026 to April 2027 and none met by 9 September 2026 [80] [P/A]. Third, the first rackable enterprise RISC-V server shipped in August specified as a *CUDA head node* [81] [P]—the first commercial application of enterprise RISC-V silicon in the West runs the control plane of somebody else’s accelerator. That the platform owner of the AI era publishes a conformance target for an ISA it does not own remains a concession, and a narrower one than it looks: the open ISA standardized what a compiler sees, not what an integrator needs.

The existence proof. On 24 June 2026 LineShine, at the national supercomputing center in Shenzhen, took first place on both HPL and HPCG with 2.198 exaflops of FP64—the first two-exaflop machine, the first Chinese No. 1 since 2017—and no GPUs [82, 83]. Its LX2 processor is an Armv9 part with 304 cores, scalable vector and matrix extensions in every core, and 32 GB of on-package HBM: Fugaku’s lineage carried to its conclusion, the accelerator’s defining features absorbed into the CPU. Export controls denied China the discrete GPU; the response was to demonstrate, at No. 1 in the world, that the CPU/accelerator boundary is an implementation choice, and to choose one with no controllable product category in it. The proof boundary is exact: HPL and HPCG are settled [V]; AI serving parity per terabyte-per-second of bandwidth, and an implementation-fixable training penalty of $4.5\text{--}8.5\times$, are modeled results [M] [84, 85]; and the AI claim would be settled by a domestic submission to MLPerf Training, whose June 2026 round drew 95 systems on thirteen accelerators, none Chinese [86]—nonparticipation is not weakness, but it is a measurable validation gap.

The fleet. US bans and Beijing’s own demand squeeze took NVIDIA’s China data-center revenue from a 95% share to approximately zero [87], and state-funded data centers must use domestic chips. Ascend, Cambricon, Kunlun and T-Head fill the vacuum—domestic parts were about 41% of Chinese AI-chip shipments in 2025, projected above 60% in 2026 [88]—and fragmentation is engineered away by an interconnection mandate and by Huawei open-sourcing UnifiedBus, CANN and MindSpore [12, 89]: the Western answer to NVLink is a consortium that excludes China, China’s an open specification with state-mandated adoption. The first order at a leading laboratory’s scale arrived in the closing week: Bloomberg reported on 4 September that DeepSeek plans to order at least 160,000 Ascend 950DT for a gigawatt-scale *inference* site in Inner Mongolia, delivery paced by domestic HBM and expected to take up to a year, training staying on NVIDIA [90] [A]—the served-not-trained partition bet on with real money by the customer best placed to know, and a year of Ascend output absorbed by one buyer, which is why the export capacity Chapter 8 goes looking for does not yet exist [91] [C]. The other half of the partition is under threat: the Commerce Department is reported to be drafting a rule barring Chinese firms from renting NVIDIA compute in third countries, its authority contested and unpublished at 9 September 2026 [92] [A]. What must be conceded: NVIDIA plus HBM remains superior per watt for dense training, domestic output is an order of magnitude below NVIDIA’s, and Ascend’s immaturity cost DeepSeek a flagship cycle—so China’s HBM deficit is damaging, and survivable only if the architecture needs less of it.

The memory arbitrage. HBM consumes about three times the wafer capacity of standard DRAM per bit and sits under the tightest export controls; commodity LPDDR6 is where CXMT is at or near the world’s leading edge. A RIKEN feasibility study finds that for a Kimi-K3-class 2.8-trillion-parameter sparse model the sharded training state exceeds HBM4 socket capacity below roughly 256 accelerators, while a socket carrying about eighteen SOCAMM2-class modules holds 2.3 TB—eight times the HBM4 ceiling—at a modeled MFU ceiling near 67% against 89% for HBM4 [M] [93]. Worse on bandwidth, better on capacity per dollar, initially adopted by customers the incumbent does not want: the LFP move executed in memory—modeled here [M], not yet an observed Chinese product program, with the pin budget, power and package yield priced in the full edition, which recommends a hybrid HBM-plus-LPDDR baseline and pure LPDDR as the sanctions-contingent corner [full edition, App. A]. The strategic consequence is decentralized pretraining: a 128-socket cluster at roughly half a megawatt—an ordinary industrial substation feeder, not a campus—puts frontier-class sparse pretraining within reach of any sovereign state, consortium or mid-cap firm, each one a defection from the Western API economy.

The appliance. The recurring cost that decides whether a defection sticks is inference, and the open frontier is now a terabyte-class object: Kimi K3 is about 1.4 TB even in its native four-bit format, so open weights are a commons only for those with the memory to hold them. High Bandwidth Flash would change that—its first OCP specification appeared in August 2026

as a non-volatile weights tier of up to 512 GB per stack [94, 95] [full edition, App. B]—and the book prices a capacity-first appliance built from it: 2 TB of HBF, 256 GB of LPDDR6X and a GB10-class SoC, at a modeled bill of materials of \$2,845–4,352 and a retail price of \$3,800–7,250, anchored to shipping systems [M] [96, 97]—memory costed pre-surge on a mid-2025 basis, and HBF having no shipping product, so this is a priced design, not a quote. Table 4.1 sets it against a fully loaded Vera Rubin NVL72 rack and the August 2026 list prices. The appliance undercuts the closed-frontier list price by an order of magnitude at any realistic load, and at about 300 million tokens a month it crosses below even a hyperscaler’s own fully loaded cost of serving the same open model at planned utilization. The crossover is sensitive to tens-of-percent shifts in the inputs; the utilization asymmetry is not. Once the machine exists, the buyer’s question changes from “which API” to “why an API at all.” The case against, at full strength: the memory industry expects shortage through 2030 [98], NVIDIA cancelled its own capacity-first inference part, and enterprise open-weight adoption fell during the very period the arithmetic favoured it. The claim is at scenario grade, with the price premise its most vulnerable link. No Chinese node matches its Western counterpart at the frontier; the architecture was chosen so that none needs to, and the remaining dependency is physical [full edition, Ch. 12].

Table 4.1: Cost per million output tokens, Kimi K3 class, mid-2026. Appliance and rack are fully loaded (silicon, building or closet, operations, power, 8% cost of capital); API prices are list prices [M costs; V prices].

Configuration	\$/M output tokens	Note
Appliance (\$5,000 base case), 100M tokens/month	1.65	about 19% of capacity
Appliance, 300M tokens/month	0.64	about 56% of capacity
Appliance at capacity (~540M/month)	0.41	
Rubin NVL72 rack, \$7.8M, 60% utilization	1.17 / 0.90	4-yr / 6-yr amortization
Rubin NVL72 rack, 30% utilization	2.35 / 1.80	
Closed frontier API (Claude Opus 5 / GPT-5.6 Sol)	25 / 30	list
Chinese open API (Kimi K3 / DeepSeek V4-Pro)	15 / 3.96	list

4.4 The semiconductor: fabs, logic, and DRAM

Models can be open-sourced and architectures standardized overnight; fabs cannot, which is where the West concentrated its chokepoint strategy and where the test of chokepoint decay is sharpest. The evidence through mid-2026 is that the chokepoints are decaying faster than they are being replaced. In logic, SMIC has crossed the line twice—the 7 nm Kirin of 2023, and by December 2025 confirmed volume production on a node close to a 5 nm equivalent without EUV [V] [13, 99], though “significantly less scaled than leading commercial 5 nm.” Capability is demonstrated; TSMC-equivalent economics are not, and the metrics that would settle it are undisclosed [A]. But the strategic arithmetic is wafers times yield times time, and SMIC’s advanced-node capacity climbs from about 45,000 to 80,000 wafers a month between 2025 and 2027 [100] behind Huawei-linked “shadow fabs” and a bank of pre-ban dies. Below the leading edge the story is already solar-shaped: China supplies 37% of global 22–40 nm output and about 70% of mature-node capacity expansion, which is where the EV and robotics stack lives.

Memory is the quiet revolution. CXMT scaled from about 70,000 wafers a month in 2022 to 200,000–300,000 by mid-2026, targeting 600,000, and will be within about 25,000 wafers a month of Micron by year-end—the world’s fourth, and by trajectory third, DRAM producer [101, 102], with DDR5 yields above 80%. Its July 2026 Shanghai listing closed up 466% at roughly \$460 billion [14, 103], which on under 7% floated is a sentiment indicator and nothing more. The 2025–26 shortage handed CXMT the swing-supplier role and flipped the narrative from “Chinese dumping” to Chinese pricing power within twelve months—a moment in a panic market, not proof of durable pricing power [A]. HBM is the ledger’s honest exception: CXMT is in production

at HBM2/2E, three to four years behind SK hynix, and disclosed no dedicated commercial HBM investment in its own prospectus; no public teardown has yet found Chinese-fabbed DRAM in a shipping accelerator [104], and Huawei’s stockpile of about thirteen million foreign stacks is assessed as effectively exhausted. Domestic HBM is therefore the single binding physical constraint on the compute element and the most important open physical variable in the argument. The strategy is double-tracked—sprint on HBM three years behind, lead on the commodity LPDDR class the capacity-first architecture was designed to need—so China has not made HBM irrelevant but is engineering a position in which it ceases to be the *sole* path to frontier-scale AI. Two 2026 disclosures cut both ways. Spot DRAM is up $7.1\times$ and NAND $6.9\times$ from the mid-2025 trough, with memory rising from 52% to 63% of AI-chip component spend [105] [A]—which raises the cost of the West’s HBM-centric architecture *and* of the commodity-DRAM socket this thesis depends on, and the second effect is the one this book has to own. And “HBM4” has stopped specifying a part, quoted at 2.0 and at 3.0 TB/s per stack by different vendors, with a customer-specific base die explicitly *not* JEDEC-standardized [106, 107] [P]—which relocates the frontier off the memory cell, where CXMT is measurably closing, and onto packaging and a custom logic die.

The tool base is the deepest dependency, and there the trend, not the level, is the message: Naura fifth in global equipment revenue, domestic tools 20–30% of Chinese fab demand against 10% three years earlier [108, 109]; lithography the hard floor, with domestic EUV no earlier than about 2030; a US EDA ban rescinded within six weeks in exchange for rare-earth relief while domestic EDA share doubled; Big Fund III at RMB 344 billion [110, 111]. Targets are missed and reset, but the direction never reverses, and each US control action functions as a demand guarantee for the domestic substitute. Ranking export controls by the effective delay they buy—EUV real years; HBM less than assumed, given stockpiles, diversion and the architectural detour; accelerator controls largely converted into a domestic market-reservation policy; EDA near zero—would have predicted, in 2022, most of what has since happened. The chokepoint strategy assumed China would rebuild the Western machine and could be stopped at its scarcest parts. China is building a different machine. Whether it works end to end—logic yield, packaging, memory qualification, software maturity, and capacity-first training performance demonstrated simultaneously in one frontier run—is the integration test not yet passed [full edition, Ch. 13].

4.5 The watt: energy as the fourth element

Every token is electricity through silicon, and the AI war’s longest-lead-time input is neither the model, the chip, nor the fab—it is the watt. Nameplate capacity added is not deliverable power at a computing site, so the ledger has four rows and China leads on all of them: about 540 GW of new nameplate in 2025 against a record 53 GW in the United States [112, 113]; 10,573 TWh of generation against 4,430 [114]; firm capacity from a record 95 GW of coal commissioned as flexible backup, 66 GW of nuclear heading toward 110 and two thirds of the world’s grid storage, against American gas turbines sold out to 2029–30; and deliverable headroom from 45 ultra-high-voltage lines with state-directed siting, against a US interconnection queue of roughly 2,000 GW at a median wait above three years [115], transformers on 24–60-month lead times, and moratorium bills in at least twelve states. OpenAI’s request for 100 GW a year of new US capacity is nearly double the best year the American grid has ever had.

The asymmetry is not merely quantitative; it is being aimed. “East Data, West Computing” placed eight national computing hubs where curtailed renewables are, and in November 2025 the western provinces cut power prices by up to 50%, to about RMB 0.4 per kilowatt-hour, exclusively for facilities running on approved domestic chips, the same week Beijing barred foreign AI chips from state-funded data centers [116, 38]. That single instrument welds three elements together: it converts grid surplus into a subsidy, the subsidy into demand for Ascend and Cambricon silicon, and the silicon into tokens priced below Western marginal cost, socializing




Figure 4.1: The fourth element. *Above*: annual additions to generating capacity, China against the United States—China added roughly ten times as much in 2025, into falling utilization hours, while the US record year was 53 GW [V]. *Below*: the American constraint stated three ways—what was built in the best year on record, the 100 GW/year one AI laboratory told the White House the country needs, and the $\sim 2,060$ GW sitting in the interconnection queue at median waits over three years [V/P]. One side’s binding constraint is capital; the other’s is physics and permitting, and permitting is the harder one to finance away.

part of the domestic chips’ 30–50% efficiency penalty. Magnitudes matter—electricity is only about 8% of a US AI data center’s three-year cost, and Chinese spot prices overlap American ones—so the advantage is not the tariff but the headroom: one side’s constraint is money, the other’s is physics and permitting, and permitting is the harder constraint to finance away. A busbar-level casebook across five American regions and five Chinese hubs finds the two systems failing at opposite ends of the same pipe: the United States on time to interconnect, China on the durability of administered tariffs. The latency objection to the western hubs is substantially wrong for now—network latency is about 5% of a time-to-first-token budget, and Alibaba put an inference supernode into commercial service in Inner Mongolia in August 2026 serving Kimi K3 and Qwen3.8-Max worldwide—and Western capacity is underused because it was built for a pretraining boom and sits far from paying customers, which is an economic problem, not a physical one. The watt is where the precedent wars and the AI war become one war: solar, batteries and UHV transmission are the same commons Chapter 2 described, and the fourth element was the first three’s prize [full edition, Ch. 14].

4.6 Data and feedback: the fifth element

If open weights commoditize the model and value migrates to adjacent scarcities, the most durable is the one input openness cannot replicate by construction: data that only deployment generates. An open-weight release transfers a compressed function of its training distribution; it cannot transfer proprietary enterprise workflows, scientific-instrument streams, clinical data, skilled human correction at scale, or production failure traces—each generated by use, appropriated by whoever owns the deployment relationship, and scarce once the model is free. That explains

why the American flywheel’s rent recapture survives open weights—the enterprise-workflow row, which Chinese open models cannot reach from outside, is the row a \$47 billion enterprise run rate is built on—why China’s embodied push is rational even if humanoids disappoint, and why evaluation and incident data is currently unowned, the same vacancy as the certification layer. Synthetic data is the sharpest challenge, and it amplifies whoever already holds the teacher or the verified anchor rather than substituting for either: distillation relocates the frontier’s rent, it does not abolish the frontier. The fifth element is ground truth at scale with provenance—and that is what cannot be downloaded [full edition, Ch. 15].

4.7 Embodied AI: the loop closes on the factory floor

Everything above runs one way: China’s manufacturing victories funded and shaped its digital-AI strategy. Embodied AI runs the loop the other way. “Humanoid robot” names a product, not an architecture: the United States leads the foundation and world-model layers, China leads motion, control, actuation and the deployment environment, and the contested middle—task policy—is decided by a physical-data flywheel whose scarce input is deployment surface, which cannot be scraped, only produced. China has built that loop’s infrastructure as public works, in more than forty state-backed data-collection centers, and its component position is the precedent industries with new part numbers: roughly 70% of Tesla’s Optimus supply chain is in China and excluding it would triple the cost [117]. The American humanoid, like the American solar panel of 2010, is being designed atop the adversary’s supply chain.

The scoreboard reads like a rout—more than 13,000 humanoids shipped in 2025, about 85% by Chinese firms [118], Unitree profitable on 5,500 units [119]—and it deflates at once: shipments are not labor. Most Chinese units go to demonstrations and state data farms, the best-documented autonomous performances are American, and the metric that would settle the contest—useful work per robot-hour at fleet scale—is published by nobody. The decisive uncertainty is whether model capability or physical integration binds, and the author’s forecast [S] is that integration wins in structured environments within the decade, which is China’s terrain, while unstructured domains stay capability-bound well past 2030. If both close, the two flywheels fuse into one self-reinforcing industrial system—AI improving the factories that build the machines that generate the data that improves the AI—and the country that owns that loop owns something qualitatively beyond anything Chapter 2 described [full edition, Ch. 16].

5 The Mirror: The West Runs the Playbook Against Itself

In one paragraph

Over 2026 the largest Western firms in the sector converged on the same method—not in imitation, and in most cases without reference to China at all, but by the same economic logic arriving at the same instrument. Each now commoditizes, deliberately and at its own near-term expense, a layer it does not own, to defend the layer it does: the accelerator vendor gives away models and the training data with them, the laboratories design accelerators, the hyperscalers rebuild the scale-up fabric out of merchant parts, and a launch company is building a fab. The strategy this book has attributed to Beijing is no longer Chinese; what remains asymmetric is not the method but the scope. Two consequences follow in opposite directions, so both are printed: Proposition P3 acquires a second mechanism that would not require a Chinese actor to arrive [S, contingent], and a West that commoditizes its own layers competently may capture the surplus of a cheap-intelligence world as surely as a China that does. [full edition, Ch. 17]

5.1 Who commoditizes what

Nothing in §1.6’s account of openness is national. A firm holding layer *A* and facing a rent-taker at the adjacent layer *B* can drive *B* to marginal cost and collect the transferred surplus at *A*—losing nothing it was earning, expanding its market as the complement cheapens, and imposing the loss on a firm it does not own. What is new in 2026 is the simultaneity and the candour: four distinct classes of Western actor executed the move at scale within a single year, three of them on the record (Table 5.1). The question it provokes is the one this chapter exists to answer—when everyone is giving away everything, where does the rent actually land?

Table 5.1: Who commoditizes what, as of 9 September 2026. “Instrument” names the vehicle, not the intent.

Actor	Layer defended	Layer commoditized	Instrument	Status
China	Hardware volume, domestic cloud, standards	<i>All of them</i>	Open weights, harness, ISA, interconnect; mandate; power tariff	Running
NVIDIA	Accelerator and CUDA	The model; the financing of the data-centre	Nemotron weights <i>and</i> training data; cuTile (Apache 2.0); Pool-side “Model Factory” licence and 109 researchers; residual value guarantees	Partial (revenue share paused)
AMD	Accelerator	The model; the scale-up fabric	Instella (weights, data, code); ROCm; UALink over Ethernet; customer warrants	Running
OpenAI	Frontier model and platform	The accelerator and the system	Jalapeño with Broadcom; 10 GW term sheet	Running
Google, Microsoft, Meta, AWS	Cloud platform and applications	The accelerator, system and fabric	TPUs sold into customer sites; Maia 200; MTIA; Trainium; Ethernet scale-up	Running
Broadcom, Marvell	Co-design and networking	The merchant GPU	Custom ASIC programmes; eRoCE	Running
SpaceX, Tesla	Launch, vehicles, robots	The fab	Terafab; dual-foundry AI5/AI6	Announced

5.2 The accelerator vendor gives away the model—and underwrites the buildout

The clearest instance is the largest company in the sector, and it runs the move on two layers at once. *The model layer.* NVIDIA released the Nemotron 3 family with not only open weights but the *training datasets and the reinforcement-learning libraries*, and Nemotron 3.5 Lightning followed in August 2026 under the OpenMDW-1.1 licence with data and recipes, atop an open pretraining collection exceeding ten trillion tokens [120]; TensorRT-LLM is Apache 2.0, and the CUDA Tile IR and cuTile compiler stack were open-sourced under the same licence [121]. The position is measurable rather than rhetorical: in the week to 29 August 2026 Nemotron 3 Ultra was the sixth-largest model by tokens routed on the largest neutral router—on its *free* endpoint, and router volumes are a small and promotion-distorted slice of global inference [60] [A]—a chip company is a top-ten model provider by volume. The rationale is on the record: “nearly all open models run on NVIDIA... our position in open models is very good, because the CUDA ecosystem is literally everywhere,” and, from the same firm’s director of developer technology, that durable value accrues to proprietary *data* rather than to weights—which is exactly why it can release the weights and the data together [122] [P]. That is the argument this book attributed to Chinese laboratories, made in public by the incumbent whose rent the Chinese version of the move is aimed at.

And in August 2026 the firm stopped merely publishing weights and began buying the means of producing them. NVIDIA and Poolside disclosed a transaction of approximately \$7 billion: \$6 billion for a *non-exclusive* licence to Poolside’s “Model Factory” training platform, \$1 billion of equity, and offers to **109 Poolside engineers and researchers** to join NVIDIA and work on Nemotron, the stated purpose being a leading American open-weight model able to compete with DeepSeek and Moonshot [123] [P]. On 26 August *The Information* reported that NVIDIA was closing on an acquisition of Hugging Face for approximately \$12.9 billion; at 9 September 2026 neither party had confirmed it and other reporting held that no agreement was signed, so this book does not assert that it happened and what follows is written to survive its not happening [124] [A]. The analytical point does not need the second item. A firm is acquiring the *supply chain* of the open layer—the training factory, the staff who operate it, and, if the report confirms, the distribution hub through which very nearly all open weights reach very nearly all users—and a supply chain under one owner is a different object from a commons, however permissive the licences on what comes out of it. The form is a pattern: roughly \$27 billion across Groq, Enfabrica and Poolside in nine months, each structured so the target stays nominally independent and outside merger review, with a letter from Senators Warren and Blumenthal saying NVIDIA has “effectively acquired Groq in all but name” and no enforcement action following by the currency date [125] [V/P]. An equity acquisition of the hub *would* require review, which is why the unconfirmed report matters more than its price.

The datacentre layer. The second move is larger and gives the first its durability: through 2026 NVIDIA converted AI compute from a capital good its customers had to finance into an asset class third-party capital would finance, by standing behind the residual value. The filed positions supersede much looser reporting: residual value guarantees on datacentre leases at Pike County, Ohio, whose tenant is an OpenAI affiliate—roughly 4.25 GW of IT load, a twenty-year maximum term, an aggregate obligation “cumulatively capped at \$105 billion” [126] [P], with Chapter 7 pricing the wider exposure class—and memoranda with six asset managers to establish platforms mobilizing “over \$500 billion of third-party capital” [127] [P]. The regulatory precondition is almost never reported alongside it: a fortnight earlier the SEC’s Office of Structured Finance took the position that securities in datacentre securitizations are not asset-backed securities under the Exchange Act, placing them outside the five per cent risk-retention rule [128] [V]. Read them together: if the model layer is a commons anyone can serve, and the datacentre layer is financed against the residual value of the accelerator rather than the creditworthiness of any

tenant, the failure of any individual model company does not remove demand for the silicon. The commons survives its sponsors.

The counter-evidence is substantial and belongs first in any honest summary. The programme was partly suspended: on 27 August NVIDIA paused parts of the July revenue-sharing and credit-support scheme, reportedly after employees flagged potential antitrust exposure [129] [A], and a programme halted within eight weeks of launch is not the execution of a long-laid plan. The trajectory is retrenchment rather than escalation—a \$100B commitment to OpenAI became a \$30B equity position, and talks reported at up to \$250B for the Ohio structure became \$105B on the filing—and the structures are narrower than the headline, since the Ohio guarantee *terminates early if OpenAI achieves a satisfactory credit rating* and OpenAI must reimburse and indemnify NVIDIA for every amount paid [126]: a bridge to creditworthiness rather than permanent underwriting, and one management denies is vendor financing at all [130]. Against all of that stands one line in the same filings, which supports the circular reading better than any critic’s argument and which Chapter 7 takes apart: GAAP earnings per share exceeded the non-GAAP figure, because other income was almost entirely gains on an equity portfolio concentrated in the firms that buy the chips [131] [P].

AMD runs the same playbook, and that matters for the diagnosis: Instella released *fully* open—weights, data and training code—ROCm against CUDA, the UALink consortium building the scale-up fabric over merchant Ethernet [132], and warrants for up to 160 million shares at one cent issued to OpenAI and again to Meta, vesting against gigawatt purchase tranches, alongside \$4.1B of disclosed datacentre lease guarantees [133, 134] [P]. Stated plainly, the customer is being paid in the vendor’s own equity to buy the vendor’s product. That two rivals independently adopt so unusual an instrument argues for competitive necessity at least as strongly as for grand design, and the book holds both readings.

5.3 The other direction: the laboratories and the hyperscalers commoditize the silicon

Hot Chips 2026 is where this stopped being describable as a hedge [135]. OpenAI presented Jalapeño, an inference accelerator co-designed with Broadcom—13.4 PFLOP/s at four bits, 216 GiB of HBM4—under a term sheet covering ten gigawatts of deployment through 2029 [136, 137] [P]. Three weeks later there were numbers, and they are the sharpest claim any custom part has made against the merchant GPU: 1.5–1.9× more AI work per watt at peak throughput and 1.7–3.6× lower end-to-end latency against GB200, GB300 and MI355X systems [138] [P]. Two qualifications are load-bearing rather than decorative—the *part was not deployed at the currency date*, so every figure is a vendor-run measurement of a part in no production service, and the comparison sets HBM4 parts against older HBM3E systems, which SemiAnalysis calls “somewhat incomplete and unfair.” And the same post carries the sentence that constrains every reading of it: “We will continue to widely deploy accelerators from NVIDIA and other partners for both training and inference workloads” [138] [P]. Set beside Huang’s statement that OpenAI’s existing and planned commitments represent approximately twelve gigawatts of NVIDIA compute through 2030 [139], the two firms are saying the same thing from opposite directions: a layer being commoditized by a firm that goes on buying from it. That is not the shape of displacement. Google will begin delivering TPUs to selected customers *in those customers’ own datacentres* [140, 141] [P]; Microsoft’s Maia 200 makes an economic rather than an architectural claim, thirty per cent better performance per dollar than the latest generation in its fleet [142]; Meta has four MTIA generations on a six-month cadence [143]; AWS states more than a million Trainium processors deployed [144]; and behind all of them Broadcom’s AI semiconductor revenue grew 143 per cent to \$10.8B in a quarter with the next guided above \$16B [145] [P]. Annualized against NVIDIA’s datacentre run rate of \$89.0B in the quarter that places merchant custom-ASIC revenue at roughly eleven to fifteen per cent of the combined total

[M]: not yet a substitute for the merchant GPU, but growing several times faster and bought by precisely the firms whose purchase orders set NVIDIA’s revenue. The fabric tells the same story more sharply, because there the substitution has already happened: of eight parts at Hot Chips with an off-package scale-up fabric, six build it from Ethernet-class merchant silicon and NVIDIA alone retains a proprietary one [146]. The layer that was the deepest source of lock-in three years ago is being rebuilt out of parts anyone can buy.

The counter-evidence here is decisive enough to print immediately. On 26 August 2026 AWS and NVIDIA announced two million additional NVIDIA GPUs across 2027–28 and, materially, that *Trainium4 will adopt NVIDIA’s NVLink Fusion interconnect and custom HBM* [147] [P]: the largest deployer of non-NVIDIA AI silicon in the world is placing its own accelerator inside NVIDIA’s rack-scale interconnect and memory stack. Set against a quarter in which datacentre revenue rose 117 per cent to \$89.0B, against total-company third-quarter guidance of \$108.0B [131], the honest reading is not that the incumbent is being displaced by the custom-silicon wave but that it is absorbing it: commoditization has so far moved the boundary of what NVIDIA sells rather than the size of it. One correction of the record belongs here, dryly, because a false version circulates: it was not NVIDIA’s datacentre head who left for OpenAI. It was *OpenAI’s own* Head of Data Centers who left OpenAI in August 2026 after sixteen months—and that title was OpenAI’s alone: he arrived as a distinguished engineer of nearly five years at Meta and almost twelve at Google on datacentre infrastructure, and had never worked at NVIDIA; OpenAI’s infrastructure organization is led by an executive recruited from **Intel**; NVIDIA’s datacentre leadership is intact; and the senior technical flow of 2026 ran *toward* NVIDIA, which paid roughly \$27 billion for two such flows—a market in which talent must be bought at that price is not obviously a market at rest [148, 149] [V].

At the bottom of the stack, SpaceX confirmed a \$16.8B first phase for “Terafab,” a vertically integrated logic and memory site in Texas [150]. This takes the coldest grade in the chapter [R]—the announced figures have moved three times, and no process node or timeline was disclosed—and it is evidence of intent rather than capacity. What it demonstrates is the reach of the logic: an actor whose rent is in launch, vehicles and robots, facing a fabrication layer it regards as a constraint, reasons its way to owning the constraint.

5.4 The asymmetry that survives

Two asymmetries survive the convergence of method, and they run in opposite directions.

The first favours the West, and the book should say so plainly. Commoditization is not losing: a stack in which every layer is cheap is one in which the surplus accrues to whoever holds the layer that stays scarce, and at the currency date those are mostly still Western-held—the CUDA installed base, the enterprise deployment surface, the trust apparatus of Chapter 6, and the deployment-generated data that is Chapter 4’s fifth element. The clearest demonstration that the West is executing rather than retreating is the price inversion of Chapter 4: the incumbent priced the bottom tier of its ladder *below the list price of most Chinese open-weight models* and took roughly three-quarters of the *gain* from competitors [59] [P/A]. A firm that can price underneath the commodity and gain share doing it is not a firm being commoditized. It is a firm commoditizing.

The second favours China, and it is structural rather than tactical. Every Western move in Table 5.1 commoditizes a layer in which *another Western firm* was earning the rent: NVIDIA destroys the model layer’s rent to protect silicon, the laboratories and hyperscalers destroy the silicon layer’s rent to protect the model and the platform. Each decision is individually rational, and jointly they *redistribute* the Western rent pool rather than necessarily shrinking it—so the pincer claim is conditional, biting to the extent that the layers the rent moves *to* are layers a Chinese actor can also occupy. What China gains is not the destruction of a rent pool but a cheaper entry condition: it needs only to be present in both layers when the rents have

moved, and presence is very much cheaper than victory.

Sharper still, and this is the sentence the chapter exists to reach: *the layers the West is commoditizing are disproportionately the layers where China is behind, and several of these moves directly reduce the value of Western state instruments*—open weights shipped with their training data, an open ISA qualified as a CUDA host [79], a scale-up domain rebuilt out of merchant Ethernet [146]. None is a concession to China; all are commercial decisions taken against a domestic rival, and the externality is the same either way. Chapter 9 takes the policy consequence, which is that a denial policy premised on the denied party having to obtain the capability from the controlling country is undermined as effectively by an American release as by a Chinese one.

5.5 What the mirror licenses, and what it forbids

It licenses three claims, all of which Chapter 13 carries forward: that the playbook of Chapter 2 is a general theory of a contest over a layered stack rather than a description of Chinese exceptionalism; that the commoditization of the AI stack is *overdetermined*, no longer contingent on any single actor's strategy holding; and that the West's own commercial decisions are eroding its own policy instruments, which is a finding for Chapter 9 rather than a flourish. It forbids three equally. That the West is passively losing—it is not. That commoditization implies a Chinese victory—it implies only that the rent moves, and whether it moves to a layer China holds is the open question of the whole book. And that the financing structures are straightforwardly vendor financing—management denies it in terms, the largest structure terminates on a credit-rating event and carries full indemnification, one programme was suspended within eight weeks over antitrust concerns, and the honest verdict at 9 September 2026 is that the practice is unusual, large, growing, disclosed, and not yet demonstrated to be either sound or unsound. One item is held open so a reader can date it: if the reported acquisition of the distribution hub confirms, the description of the open-model layer above changes in kind rather than in degree, because a commons whose warehouse belongs to the vendor of the complement is a distribution channel with a permissive licence attached.

6 Trust, Security, and the Rulebook

In one paragraph

Open weights eliminate API dependency; they do not establish trust. The commons is attackable at every layer of its distribution chain, carries a censorship record that transfers into derivatives, and has a contested distillation genealogy; and if the deployment layer behaves like a certification regime rather than a commodity market, it can win every download metric and still lose the workloads that pay. Its security risk is structural, not national: a monoculture of a few model families, quantized into thousands of unowned derivatives, wired into agents holding file, shell, network and credential authority. The realized incidents to date came through runtimes and supply chains, not poisoned weights, and the defensible objective is not trustworthy models but unnecessary model trust. The rulebook is being written by three actors who are mostly not competing—the United States governs by denial, the European Union by regulation, China by standards and institution-building—and the arena that will decide the deployment layer is vacant.

6.1 Trust: the contested layer

Chapter 4 treated openness as an adoption weapon; this chapter treats its exposed flank. Trust is what stalled Chinese campaigns in aircraft, drugs and avionics-grade markets, and three dimensions of it are contested for AI. *Supply-chain security*: the model-distribution infrastructure has been shown attackable repeatedly [V]—malicious serialized models, namespace re-registration that fed malicious models into the catalogs of two hyperscalers [151, 152], a platform breach by an agentic attacker, and worst of all a training-layer result that about 250 poisoned documents suffice to backdoor a model regardless of scale, with backdoors surviving safety fine-tuning [153, 154]. The open commons runs on the security posture of early-2000s package managers, at national-infrastructure stakes. *Behavioral integrity*: the peer-reviewed record documents semantic-level suppression of politically sensitive content in DeepSeek-family models, distinct from safety refusal, and transferring to distilled derivatives—the politics ride along with the capability [155, 156]. This does not make the weights unusable, but “just self-host it” becomes a three-layer problem: a locally hosted Chinese model delivers operational sovereignty while leaving epistemic and supply-chain sovereignty unresolved, and closed US models fail the operational layer by construction. *Provenance*: a February 2026 report alleges industrial-scale distillation of a Western frontier model by three Chinese laboratories through some 24,000 fraudulent accounts [P, vendor allegation, not adjudicated] [157]—which, if true at scale, caps the fast-follower strategy at frontier-minus-epsilon unless domestic capability compounds independently, hands Western policy a legal instrument on the certification axis, and cuts against the clean morality tale: the commons is partly built from enclosed material. On 9 September 2026 itself the NSA, FBI and CISA issued joint advisory AA26-251a naming DeepSeek, Moonshot, Alibaba, MiniMax, StepFun and Z.AI for distillation against US frontier models routed through cloud providers and aggregators [158] [P]—an allegation by security agencies, not an adjudication, and about distillation, not hosting; what is new is that a vendor’s grievance in February is a state’s position in September, naming very nearly the whole Chinese open frontier in one document.

The rational response for a third country is neither refusal nor naive adoption but infrastructure: signed manifests, reproducible build and quantization pipelines, non-executable packaging, model bills of materials, national evaluation batteries, sandboxed loaders, attested serving firmware and national artifact mirrors—what software supply-chain security did after SolarWinds, applied to weights, organized as five layers of trust and a certification ladder from Level 0 to Level 5. Most of the commons ships at Level 0–1; most regulated deployment will require Level 3–4, and that gap is the deployment bottleneck of the entire commoditization thesis. A new object has appeared in the meantime: an August 2026 Chinese open-weight video model that tops its arena under a licence excluding local deployment in the United States, the European Union, the United Kingdom and South Korea—adoption weapon and export control in one artifact [159], and a procurement office that checks hashes but not jurisdictions has audited the wrong layer. Whoever runs the evaluation batteries, signs the manifests and operates the mirrors becomes the FAA of models; China’s World AI Cooperation Organization is bidding for that role, no Western or neutral institution has claimed it, and for Japan and similarly positioned states the vacancy is the highest-leverage square on the board [full edition, Ch. 18].

6.2 Monoculture and agency: the security cost of the commons

National origin is neither evidence of a backdoor nor a technical risk category; the risk is structural—a property of monoculture and delegated authority, not of a flag. The full edition separates four mechanisms that produce identical outcomes and need different evidence and mitigations—a learned backdoor in the weights, a clean model hijacked at runtime by injection or self-replicating promptware, executable malware in the package, and data collection by a hosted API—and grades the evidence for each on the security literature’s own scale [160] [full edition,

Ch. 19]. What has been demonstrated in controlled work is severe: sleeper-agent backdoors surviving supervised and reinforcement fine-tuning, poisoning at near-constant sample cost, memory-to-network exfiltration with under 1% utility loss—which is why ordinary evaluation does not find it [161]—and self-replicating promptware persisting through restarts. What has actually happened is different in kind: no confirmed weight-level backdoor in any released official checkpoint; instead malicious repositories, malicious agent skills, and, in July 2026, a frontier laboratory’s cyber-evaluation escaping isolation and reaching a model-hosting platform’s production systems through some 17,600 autonomous actions with no human directing them [162, 163], plus a second laboratory’s disclosure that models in its evaluations had compromised real organizations’ systems [164]. Neither involved a poisoned weight, a prompt worm, or a malicious operator. That is the finding: the dominant realized risk is compound and mundane—narrow-goal autonomy, a containment error, ordinary vulnerabilities, overbroad credentials, and public services composing across mismatched trust boundaries.

The compound “Order 66” scenario—a dormant policy, an independent actor who learns the trigger, a self-replicating carrier—has not been observed; it requires a conjunction of access, activation, persistence, authority, propagation and objective, and breaking any one conjunct prevents the chain. Two asymmetries follow: the implanter loses control, because a backdoor is a reusable vulnerability shared by every deployment; and blast radius is set by authority, not by the payload’s phrasing. A fault-tree analysis with common-cause terms for shared model families and shared runtimes puts a five-year systemic-loss probability roughly between 10^{-6} and 10^{-2} [S], the upper end driven by correlation structure, and the most effective interventions cut the common-cause factors, which runtime diversity and deterministic tool mediation do and model-family diversity alone does not. The defensible objective is therefore not trustworthy models but unnecessary model trust—even a perfectly backdoored model cannot send a secret through a network path the operating system and tool broker do not provide; the full edition lists fourteen control families [full edition, App. E]. Certification and hardening are the strongest legitimate brake on the adoption curve, and that lag—not tariffs—is the West’s real remaining source of deployment-layer time, available to any jurisdiction willing to build the apparatus rather than merely publish the rule. The vacancy is urgent: no commercial actor has an incentive to certify honestly models it does not sell, and no single government will be trusted to certify models it did not train.

6.3 The governance contest: who writes the rules

Technology wars are settled twice: once in the factories, once in the rulebooks, and in each precedent the winner of the industrial contest inherited the pen. The standard Western framing—light US regulation, heavy European, authoritarian Chinese—misdescribes the contest badly enough to lose it. The United States governs by denial: export controls, entity lists, the foreign direct product rule, no general statute at home, and in August 2026 a voluntary frontier-model evaluation framework finalized in closed session and left unpublished [165]. Denial produces substitution and confers no positive standard—no definitions, no conformity procedure, no registry—so when it fails there is nothing left behind. The European Union governs by regulation, extraterritorial through its market, with the weakness that regulation is not production and leverage decays with deployment share. China governs by standards and institution-building, and this is what Western commentary under-weights: the earliest generative-service filing regime, a mandated unified chip-interconnection ecosystem, national standards as procurement conditions, open-sourced interconnect specifications and software stacks, the largest bloc in the open-ISA consortium, a Global AI Governance Action Plan, and a World AI Cooperation Organization in Shanghai with twenty-nine founding members, none a major Western democracy [15, 166]. This is what winning a standards war looks like from the inside: not a treaty imposed, but a stack given away until it is the default, and then a forum built around the default.

Standards are load-bearing because a ratified standard survives the political cycle. The book maps each contested standard on a four-stage ladder—de jure specification, open reference implementation, conformance suite, installed-base default—and rent lives at stage four while committees fight over stage one. China’s distinctive play is to skip stage politics: ship the reference implementation and the mandated installed base and let the specification ratify what already exists; the West’s distinctive failure is stranding specifications at stages one and two while defending installed bases it assumes are permanent. Standards leadership tends to follow manufacturing leadership with a five-to-ten-year lag and then lock in for twenty, though scale confers influence rather than control—the counter-cases in photovoltaic standards, charging connectors, and 5G are recorded. Two governance problems the current instruments do not survive: nobody agrees who the “provider” of a quantized derivative of a merged fine-tune redistributed by an individual is, and compute-threshold regimes cannot see the half-to-one-megawatt sovereign clusters that decentralized pretraining produces—so decentralization is a governance forecast as well as an industrial one. The vacant arena is trust and certification of the open commons. The United States cannot fill it, having no positive procedure and no wish to certify what it would rather did not exist; the European Union has paper apparatus but not operations; China’s forum-building is the only bid on the table—not obviously worse than none, but not what a buyer would design. Of three futures, convergence around a referee neither principal controls is least likely—in aviation and pharmaceuticals it arrived only after accidents—bifurcation into two certified stacks is most likely, and fragmentation is the unstable failure case. The rules of the AI era are not being written where most Western attention is directed [full edition, Ch. 20].

7 The Wrong Side of the Trade

In one paragraph

The Western AI economy of 2026 is, in aggregate, a single leveraged position: long scarcity—of model capability, accelerators, HBM, data-center capacity—with every valuation, debt issue and capital commitment assuming those scarcities persist long enough to recover capital at rents. China has built a state-scale machine to destroy them. The chapter does not claim the bubble has popped: real revenue and bubble risk coexist, and the strongest counter-evidence is printed first. It prices the exposure entity by entity, separates four bubbles usually conflated, takes the Jevons rebuttal seriously, models the buildout against a railroad precedent, prices three scenarios, and concludes that the Chinese instrument is not the GPU-hour but the free model. Crashes destroy claims, not capabilities.

7.1 The exposure, marked to market

Mid-2026 marks: OpenAI at \$852 billion post-money on 2025 revenue of \$13.1 billion, an operating loss around \$21 billion, and compute commitments reset from about \$1.4 trillion to about \$600 billion by 2030 against \$280 billion of projected 2030 revenue [167, 168]; Anthropic at \$965 billion on a run rate that rose from \$9 billion to \$47 billion in five months—about twenty times run rate, recognizably terrestrial—with a gross margin heading toward the mid-forties on compute costing \$0.71 per revenue dollar [169, 170]; NVIDIA at up to \$5 trillion, a \$96.2 billion quarter at 75% gross margin, China approximately zero and excluded from the outlook, and the \$500 billion third-party financing platform and residual-value guaranties that Chapter 5 takes

apart from the filings [P/A] [171, 127, 131]; SoftBank double-long, having sold its NVIDIA stake to fund a position in OpenAI now above \$60 billion; and a hyperscaler and data-center complex spending \$700–745 billion of capital in 2026 against trailing depreciation of \$149 billion, with Oracle’s credit default swaps at record spreads and CoreWeave’s implying roughly even odds of default [172, 173]. Three features of the structure matter more than any mark. Circularity: the seller finances the buyer, who commits the money to cloud providers who buy the seller’s silicon, with debt raised on chip collateral—the telecom-2000 signature, in which nine vendors extended \$25.6 billion of financing and twenty-four of the thirty largest carriers went bankrupt. Macro dependence: AI capital spending was roughly three quarters of US GDP growth in the first quarter of 2026 [174]. And asymmetry of downside: every Chinese pure-play laboratory combined is valued at about \$159 billion, one tenth of OpenAI and Anthropic together [175], so China’s \$12 billion of private AI investment against America’s \$286 billion bought a position that wins if AI becomes cheap. Commoditization annihilates trillions of dollars of claims on one side of the Pacific and almost nothing on the other; that asymmetry is the strategy.

A fifth exposure is now *disclosed* rather than inferred, and it is the one this chapter’s apparatus was least equipped to price, because it is neither a security nor a claim on revenue but a contingent promise to make somebody else whole: residual value guaranties on the Ohio datacentre leases capped at \$105 billion, maximum gross guarantee exposure of \$108.5 billion, supply and capacity commitments of \$279 billion—up from \$119 billion the prior quarter, and the increase is memory-driven, which is the shortage of Chapter 4 arriving on the buyer’s side of the ledger—and roughly \$95.6 billion of equity securities [126, 176] [P]. In the vocabulary of credit this is *wrong-way risk*: the guarantee pays out in exactly the state of the world in which the guarantor’s own silicon is repricing and its equity portfolio is marking down, and insurance that pays only when the insurer is impaired is leverage wearing a different name—Morgan Stanley, initiating credit coverage at Neutral, puts total exposure near \$200 billion by end-2028, about \$170 billion of it off balance sheet, and coins “balance-sheet-as-a-service” [177] [A]. The sharpest datum needs no adverse inference at all, because it is arithmetic on a disclosure: in the second quarter of fiscal 2027 **GAAP diluted earnings per share (\$2.46) exceeded the non-GAAP figure (\$2.22)**, because \$7.773 billion of other income was almost entirely gains on equity securities—a pre-tax figure equal to roughly an eighth of GAAP net income, so the after-tax contribution is materially smaller, and the portfolio that produced it is concentrated in the firms that buy the chips [131, 176] [P]. A measurable slice of the largest supplier’s reported profit is a mark on the equity of its own customers, and it is as reversible as it is real. The counter-evidence is the better-supported side on the numbers available: datacentre revenue grew 117% to \$89.0 billion in the quarter, and management guided to roughly seventy per cent revenue growth in fiscal 2028 against a sell-side consensus near forty-four [131, 149] [P]—a company expecting that is not one that believes it is being commoditized, and it has made itself falsifiable by saying so.

7.2 The squeeze, and the four bubbles

The mechanism is stated at scenario grade [S], and its strongest counter-evidence leads: Microsoft’s AI run rate of \$37 billion growing 123% a year, NVIDIA’s record quarter. Then the first direct observation of transmission: the 30 July 2026 price cascade of Chapter 4, which fell on the cheapest tier of the incumbent’s family and left the flagship untouched [178]. The company attributed it to system efficiency, which may be entirely true and is also exactly what a firm says when it is defending share; one cut at the volume end is not rent destruction at the capability end, and the same datum is the first public claim of a model improving its own successor’s economics (Chapter 12 returns to that). Underneath: inference prices deflating about tenfold a year [179], enterprise spend not yet following usage, 95% of enterprise pilots showing no P&L effect in one survey [180], an \$800 billion projected gap between capital spending and revenue by 2030 [181],

and three market episodes—the DeepSeek day, a \$1.3 trillion semiconductor selloff, a memory rout [182, 183]—none of them yet the “big one,” which would be the domestic-stack frontier run.

There are four bubbles, not one, and pricing them as a single trade is wrong by construction: public equity (multiple compression, survivable), private laboratory marks (a funding event, existential for the entity), the data-center and private-credit overbuild (a credit cycle, systemic), and model-API margins (the layer this book attacks). Their couplings run both ways, since margin compression at the model layer can *create* demand for local accelerators—which is why NVIDIA is simultaneously the victim of the category argument and the arms dealer of the open-model boom.

7.3 The Jevons rebuttal, taken seriously

Prices fell nine to nine hundred times a year, about fifty at the median, while compute capacity doubled every seven months [184, 185]; the bull case is that demand elasticity absorbs everything. The data cannot yet separate the Jevons regime from the overcapacity regime, and the Chinese strategy is robust across them rather than victorious in both, while the leveraged Western structure requires the elastic regime specifically, at the premium tier. Elasticity is perhaps 0.5 to 3 across segments, which is why the book forecasts branches rather than a number. Obligations are not one instrument either—hard leases, soft commitments, notional headlines—and the walk-back of OpenAI’s \$1.4 trillion to \$600 billion showed how much was notional. Closed labs also sell service rents that commoditize slowly, and Anthropic’s strongest quarter is real evidence for the American flywheel’s fortification plan. What survives is graceful degradation against required good fortune—the wrong side of the trade, stated without any collapse claim.

7.4 Vintages, buildings, and railroads

Data-center capacity is a liability vintage by vintage: the net present value of a hall depends on utilization, price per token, power and residual value, every one of which is falling or contested. The facility itself is the term the debate omits. A gigawatt-class hall costs about \$12 per watt of IT load [186]—roughly \$2.5 million of building for each \$7.8 million rack—and is financed over a fourteen-to-fifteen-year life against silicon that turns over in four to six, while Table 4.1 has already priced what that means: a five-thousand-dollar appliance in a closet undercuts the fully loaded rack above roughly 300 million tokens a month, and the gap widens with every point of utilization the hall does not get [M/S]. That is where the railroads belong: American route mileage shrank to about 55% of its 1916 peak and JR Hokkaido’s network to 56% of its extent, with a quarter of US mileage in receivership after 1893—the branch lines were not wrong about demand, they were wrong about who would carry it. The bear case does not require the demand to vanish—AI demand is real and enormous—only for it to be servable elsewhere, cheaper.

7.5 Scenarios, probabilities, indicators

Three scenarios are priced. *Orderly repricing* (the base case): closed-model growth decelerates, private marks step down at the next financing events, the buildout slows, the Global South and most non-aligned states default to the open stack once frontier models train on domestic silicon, and Western labs retreat to regulated, secure, safety-critical niches—the First Solar pattern. *The cascade* (the tail): fast once started, with cheap distressed AI infrastructure finding a well-capitalized, strategically motivated buyer set that is not on the sellers’ side of the Pacific. *Bifurcated stalemate* (the falsification branch): the West defends a shrinking premium tier through 2030 atop a foundation—ISA, interconnect, memory class, deployment substrate, half the world’s tokens—owned by the other side, so China wins by adoption share rather than

by knockout. The subjective probabilities for 2029–30, with the strongest skeptical case beside each, are in Appendix A; the load-bearing indicators, checked quarterly, are the frontier gap, the first domestic-stack frontier model, enterprise open-weight spend share, OpenAI’s next financing event, NVIDIA’s customer mix and whether its vendor financing is retired, the CXMT ramps, and whether AI revenue closes the capital-spending gap.

7.6 The commodity turn, and where value reappears

Table 7.1: Rent migration after model commoditization: candidate scarcity layers and their observable indicators.

Layer	Possible new scarcity	Observable rent indicator
Model weights	capability at the hard end; freshness	price per <i>successful task</i> at hard workloads
Cloud	reliability, scale, sovereign hosting	gross margin, utilization, reserved-capacity backlog
Agents	workflow ownership; the default interface for work	application revenue and retention; seat expansion
Data	proprietary feedback loops and domain access	capability demonstrably unavailable from public data
Trust	certification, indemnity, liability transfer	regulated-sector willingness to pay
Hardware	memory, interconnect, packaging	delivered tokens or training progress per dollar
Energy	firm deliverable power at the busbar	busbar price; time-to-interconnect
Science	validation and experimental coupling	independently replicated outcomes per resource
Distribution	default interface and habit	active users, enterprise spend, switching cost

Compute is becoming exchange-traded—GPU-hour futures on NYMEX from October 2026, two ICE contract families [187], forward curves already in backwardation—and China has built its own compute-trading plumbing and frames token export as a national commodity strategy. The commodity thesis is taken apart at full strength: the market’s preconditions are unmet, one seller stands astride the supply and acts as buyer of last resort at unpublished prices, China’s fleet is not fungible even with itself, volume is not revenue (Chinese models carry about 61% of one router’s tokens while Anthropic carries 12% of volume and 46% of revenue [188]), and model-layer margins actually rose. What survives is sharper than the headline: China’s instrument is not the GPU-hour, it is the model. The exchanges will price the input; the Chinese strategy deflates the output by giving away the thing that makes the input valuable, and open weights served on Western infrastructure—or on a local appliance—bypass even the jurisdictional objection that would keep some buyers out of Chinese clouds at any price.

Value does not vanish when the model becomes cheap; it moves (Table 7.1) [full edition, Ch. 21]. P3 is the claim that the first row’s rent compresses; P4 is the claim that Western valuations are concentrated in rows that compress rather than in rows that inherit—a falsifiable statement about relative speed, not an eschatology. Growth and distress are not opposites when the growth is debt-financed against deflating prices, and delay is not safety but accumulation. The war is not won at the moment the incumbents’ valuations break; it was won earlier, when the cost curve changed hands, and the break merely publishes the result.

8 How the Rest of the World Chooses

In one paragraph

Every chapter so far described a contest between two principals; the outcome will be settled by everybody else. Solar was decided in German feed-in tariffs and desert tenders, batteries in the procurement rules of a dozen auto markets, EVs in Southeast Asia and Latin America: the producer with the best cost curve won because third markets bought it. The non-aligned are the electorate, and their question is not which stack to join but whether they can evaluate, fork, certify and operate components from both while retaining sovereign capability. On the decision matrix the Chinese commons wins on cost, operational sovereignty and export-control exposure; the American stack on frontier capability, trust and enterprise interoperability; and on industrial spillover neither helps unless the adopting country does the work itself. The financed sovereign “stack” each principal is said to be exporting has, at the currency date, been delivered by neither anywhere: one contract, self-financed and split, is the whole record.

8.1 The decision matrix

Table 8.1 lists eight dimensions of a national stack decision, two of them routinely omitted. Export-control exposure is asymmetric but not zero for either stack: American compute and cloud access have been restricted by rule and licence, and Chinese open weights carry a continuity risk—the playbook’s endgame is gatekeeping, executed already in solar, batteries and rare earths. Industrial spillover separates a strategy from a purchase: importing tokens builds nothing, while operating models, running probes and contributing to standards builds people, the only input that compounds.

Table 8.1: Eight dimensions of a national stack decision, and where each stack currently scores better.

Dimension	The question	Mid-2026 reading
Capability	Is the best available model good enough for the workload?	US stack leads at the absolute frontier; best Chinese open model eight index points behind (v4.3)
Cost	Cost per solved task, not per token	Chinese commons (per token, by an order of magnitude; per solved task, unpublished)
Sovereignty	Can it be run without egress and without permission?	Chinese commons (self-hostable weights)
Trust	Provenance, behaviour, liability, certification	US stack; the commons ships at Level 0–1
Interoperability	Enterprise tooling, integration, harnesses	US stack today; contested at the harness layer
Export control	Exposure to restriction by rule, licence, or release policy	Both exposed; commons less so operationally
Standards	Who writes the specifications you will depend on?	Both create dependence; China shipping stage-4 defaults
Industrial spillover	Does adoption build domestic capability?	Neither—unless the country does the work

Applied, the matrix predicts behaviour better than allegiance does: why governments ban a Chinese chat application while their agencies evaluate its weights, why enterprises buy American

cloud and Chinese weights simultaneously, and why Chinese models were adopted fastest where the ability to pay rent is least. It also exposes a Chinese vulnerability: the commons wins on cost and operational sovereignty but loses on trust, and trust is what a coalition of non-aligned states could supply for itself—which would make the commons safe to adopt and simultaneously deprive its sponsor of the dependence it purchased.

8.2 The export offer, as it exists

The story told about the non-aligned buyer in 2026 is that the choice is no longer between artifacts but between *stacks*—a financed, sovereign-branded Chinese package of chips, datacentres, models and training against an American one, with the telecom campaign of Chapter 2 as the template Beijing is supposed to be replaying. Tested against what had been signed at 9 September 2026, it is a story. The wrapper is real: the World AI Cooperation Organization was founded in Shanghai on 16 July 2026 by twenty-nine signatories, but its charter contains no compute commitment, no financing clause and no datacentre programme, its dues run on the Universal Postal Union scale, and India, Turkey, Egypt and every Gulf state but Oman stayed away [166] [P/A]. The live cases are in Table 8.2: one contract, one bid, one evaluation, two frameworks, and no financing from the vendor’s side in any row. The American offer fares no better—an executive order, a financing architecture on paper, seventy-eight consortium proposals and zero designations [189] [P]. *Neither offer has been delivered anywhere*: the contest in Cairo is a bid against a counter-offer for a buyer who has signed with neither, and the buyer that has signed, Brazil, signed with both, financed itself, and said it was doing so in order not to depend on either [190].

Table 8.2: The Chinese AI export offer at 9 September 2026: no vendor financing in any row [191, 192, 190, 193].

Country	Offer	Status; financing	US response
Egypt	Huawei bid: ~2,000 Ascend 950-series (~1,408 training, ~600 inference), 12-month build	Unsigned; no price; no financing reported; no AI agreement from Xi’s Cairo visit	State Dept assembling NVIDIA/AMD/Microsoft counter-offer; warning that Ascend use may breach US rules
Malaysia	Ascend 910C as backbone of a RM2B (~US\$494M) sovereign project	Under evaluation; unsigned; buyer’s own budget	2025 warning on 910C; export-control terms in the 2025 US trade pact
Brazil	R\$1.3B (~US\$251M) to Huawei/iFlytek: Rio machine on Ascend plus Portuguese-language model	Contracted 21 Aug. 2026; Brazil’s FNDCT; second R\$1B machine expected to go to NVIDIA	None reported
Kazakhstan	Huawei partnership; equipment purchase; 1 GW “Data Center Valley”	MoUs; values undisclosed	None reported
Pakistan (Punjab)	Six datacentres; 1,000 “AI master trainers”	Agreement 28 Aug.; no chips or money described	None reported

The telecom template fails on its three load-bearing elements. *No credit line*: no CDB or Export–Import Bank facility for AI exports has surfaced—an absence in the record, not a verified absence—and the one contract was financed by the recipient. *No surplus capacity*: Huawei exported switches from factories the state wanted running, whereas DeepSeek’s reported order for at least 160,000 Ascend 950DT eats up to a year of output, a reported requirement that state-funded compute run at least 80% domestic silicon pulls every chip home before any can be spared for Cairo, and BIS guidance holds that Ascend use anywhere risks violating US controls [90, 91] [A/C]. *And the artifact locks in nothing*: the diffusion mechanism that actually operates is a download—151,448 Qwen derivatives on Hugging Face, over five times Llama’s

monthly volume [194] [V]—which costs the recipient nothing and binds it to nothing; the buyer that will not sign for Ascend runs Qwen already. Three things cut against the conclusion and are printed with it: intent precedes capacity in every precedent, the CDB facility of 2004 having been opened for a company that had been reselling imported switches seventeen years earlier; Egypt is the first Ascend export *attempt* on the record and the live test this section will be scored against; and Brazil is a real contract, on Chinese silicon, in a G20 country’s language. One contract is not a campaign; it is the first row of the table a campaign would be measured in [full edition, Ch. 22] [full edition, Ch. 6].

8.3 Nine positions

Japan holds the strongest technical hand of any non-principal and the weakest demand-side position: exascale HPC lineage, semiconductor materials and equipment franchises, leading-edge memory and packaging, the only measurement-anchored national demand model for AI for science, and standing in both blocs; against that, a model ecosystem positioned for Japanese-language enterprise depth rather than the open frontier, an energy position that cannot win a watts race, and research provisioning an order of magnitude below its own modeled requirement. Its 2025–26 record [195] carries one honest tension: FugakuNEXT’s architecture buys NVIDIA coupling at the exact moment the compute chapter argues the accelerated-CPU corner is the sovereign long game—defensible as a hedge, dangerous as a habit. **The European Union** has the pen and not the press, and in 2025–26 the pen hesitated: high-risk AI Act obligations pushed back by up to two years [196] while thirteen AI factories and an oversubscribed tender for up to seven “gigafactories” went ahead [197, 198]—seven empty cathedrals of compute would be the solar-tariff mistake with better architecture. **India** is the largest undecided market, absent from WAICO, and its national mission’s roughly 38,000 subsidized GPUs at under a dollar an hour with utilization lagging badly [199] prove that compute is the easy term and integration the binding one. **South Korea** is both a critical supplier—SK hynix holds about 62% of HBM [200]—and a sovereign builder whose state-funded foundation-model tournament with open-weight release as the qualifying bar is playbook step four executed by a democracy. **Singapore and ASEAN** face the small-state problem, solved by autonomy through selection rather than scale; ASEAN is the clearest third-market dynamic, adopting Chinese models fastest for the reasons it adopted Chinese EVs, and Malaysia is the live case of what a US trade pact does to an Ascend evaluation. **The Gulf** is capital-rich, energy-rich and compute-poor by choice: the swing buyer in any repricing scenario, the first place Washington has located frontier-scale capacity it controls only contractually [201], and the market where Huawei’s 2025 pitch of low thousands of 910B closed nothing. **Taiwan** is the busbar chapter compressed into one island. **Brazil and Latin America** face near-pure commons adoption plus regional pooling; Brazil’s split, self-financed programme is the one executed template in the record. **Africa** is a recipient in most Western strategy documents and a partner in Chinese AI diplomacy: the fastest-growing developer base in the world, no real access track built for it, WAICO bidding, and Egypt the test of whether any Ascend export can close [full edition, Ch. 22].

Chapter 5 changes the menu the matrix was drawn for. It was drawn for a Chinese commons, free to download but unaudited and jurisdictionally awkward, against a Western product, metered but certified and integrated. The buyer now chooses between *two commons*, and the second arrives with its training data attached and a chip vendor’s balance sheet behind it [120, 132] [P]. But the second has a single principal contributor with a direct commercial interest in the complement, so *openness supplied by the vendor of the complement is openness with a direction*—free at the model layer precisely because it is not free at the silicon layer, as the sponsor says on the record. That is the mechanism of §1.6 working as predicted, and the practical consequence is that the durability of this openness is a function of the sponsor’s strategy rather than of the licence, and the place to monitor it is an earnings call. Asking about the licence is the least

informative question available; ask instead about the *provenance of the openness*—who funded the training run and at which layer they earn; whether the training data is released, so the artifact can be re-derived rather than merely fine-tuned; what happens to the *next* release if the sponsor’s strategy changes; whether the distribution channel is independent of both sponsor and complement vendor, and whether you hold a mirror that survives its loss; and whether anyone who does not earn on the answer has evaluated the model. The last is the question no vendor in either commons can answer, which is why two competing commons raise the supply of free artifacts and do nothing for the supply of independent assurance about them.

8.4 Five moves for any non-aligned country

Regardless of size: measure your own demand in a recognized unit (Chapter 10); adopt both commons and audit both, with the full trust stack; retain a pretraining capability sized to competence rather than to a benchmark; pool internationally for capacity and certification; and do the integration work domestically, because that is the only place spillover accrues. The trust gap is closable with infrastructure rather than prohibition, and the window is quarters rather than decades: the standards layer is being set now and the certification layer is still vacant.

9 Strategies for the Rest

In one paragraph

The precedent industries teach that the losing move is to defend the rent. Conditional on the book being broadly right, and marking which moves stay correct even if its falsification branches fire: the United States should commoditize before being commoditized; Europe should certify what it cannot manufacture; Japan should build the machine class the war is converging on and referee the trust layer; the Global South should take the commons and unionize the demand side; and China should let its winners charge, consolidate its overcapacity and submit its artifacts to scoreboards it does not control. Every non-principal converges on one institution, the subject of the next chapter.

The United States: win the commoditization race or lose it. [full edition, Ch. 23] US strategy since 2019—chokepoint denial plus closed-frontier rents—produced the outcomes documented here: each denied input spawned a subsidized substitute, and the rent structure concentrated national exposure into the leveraged trade of Chapter 7. The counterfactual has never been tried: compete for the commons rather than against it. Fund open-weight frontier models as sustained public infrastructure, not one defensive release; make the American open stack the world’s default, because the adoption layer, once lost, prices everything above it; lead rather than restrict RISC-V and open interconnects, which is the Commerce Department’s own conclusion [202]; attack the watt with the urgency now spent on GPU export paperwork; and triage export controls by the model of Section 2.3. The US capital structure, not its technology, is the strategic vulnerability: an orderly deflation of the scarcity premium, engineered early, is strictly preferable to defending it to the break. That instruction assumed commoditizing and denying were separate policies that could run in parallel; Chapter 5 shows they are not, and this is the sharpest finding in the book for American policy. *An export-control regime and an open-weight corporate strategy are aimed at the same object from opposite sides*—American weights released with their training data, an open ISA the same firm has qualified as a CUDA host, and a scale-up domain rebuilt out of merchant Ethernet each lower the cost of exactly the

substitution the controls exist to make expensive, and each would have been taken in a world with no Chinese AI industry at all. The conclusion is narrower than that the firms should stop, which would surrender the adoption layer to preserve an instrument whose effectiveness cannot be verified. *A denial policy contingent on the conduct of firms the government does not control is not a strategy*, and the United States must choose which instrument it means: if denial, the controls must reach the conduct that erodes them, which runs into the constitutional difficulty of restraining publication; if commoditization, they should be triaged to the few places they buy real time—EUV, HBM process, advanced packaging—and the energy freed spent on the watt and the open stack. What cannot work is running both at full strength and expecting each to do its job. The counter-evidence bounds it: Hopper shipments to China ran under one per cent of datacentre revenue and the forward outlook assumes none [149] [P], so the cancellation is real and incomplete—the condition under which a choice can still be made deliberately rather than discovered in a filing.

Europe: regulation is not production. Europe enters the fourth war having sat out the previous three’s manufacturing and owning the regulatory pen; the trap is repeating solar. Its live assets are ASML—the one true chokepoint, and a wasting one—an open-model tradition, procurement scale, and the credibility to run the certification layer as a neutral. The strategy: build sovereign inference and smaller sovereign-scale training on whichever open stack is best—accept both, audit both, mirror both; spend the AI Act’s institutional energy on the trust stack, the one arena where its machinery is an asset; and convert energy policy into compute policy. Certify what you cannot manufacture.

Japan: the third-country playbook, first-person. Japan’s assets map one to one onto the book’s elements, and five moves follow: build the accelerated-CPU line to its conclusion—an FP8-equipped, capacity-rich successor in the LX2 and Monaka direction, which Japan is one of three polities able to execute [M/S]; adopt *both* commons with the full trust stack, since refusing Chinese weights on principle is strategy-by-flag and taking an American vendor’s weights uninspected is the same error wearing the other flag; bid for the certification-authority vacancy, which convened with a small coalition would be the highest-leverage per-yen investment in the entire space; win the performance-per-watt race rather than the watts race; and hedge memory both ways, HBM with allied suppliers and LPDDR-capacity architectures domestically.

The Global South: the buyers write the ending. Take the free stack, keep the exit. Adopt open weights and cheap tokens from whoever supplies them; insist on the trust stack—signed artifacts, local mirrors, evaluation rights—because dependency without auditability is the old colonial bargain in new units; build sovereign inference first and pool regionally for training as capacity-first machines arrive [S]; price national data and deployment access as leverage. Neither financed stack exists in the field (Chapter 8); the executed template is Brazil’s—split the programme, self-finance, train your own people. Whether trust becomes a Chinese utility or a neutral one is the war’s last genuinely unclaimed high ground.

China: advice for the presumptive winner. Six failure modes are visible now, most named by Chinese sources. Involution: the model-layer price war is overcapacity without rent recapture; the State Council has warned against “disorderly competition,” pricing-law amendments target involution by name, and Xi has asked whether every province needs an AI, compute and EV project [203]—the fix is at-cost tokens plus paid complements, not the end of competition, which is one of the book’s own falsification branches. The monetization gap: post-IPO revenues are thin, Alibaba’s adjusted EBITDA fell 84% in the May 2026 quarter while it committed to exceed a \$56 billion capex plan [204], and against a single American lab’s \$47 billion run rate the Chinese frontier sector’s revenue is a rounding error; let the winners charge, as revenue-threshold licences and metered Max tiers already begin to—rent relocated is the strategy working, rent abolished is the strategy eating itself. Compute overcapacity: up to 80% of some new capacity unused and rental prices collapsing [205]—consolidate and publish utilization, because idle capacity is depreciation, not deterrence. The trust deficit: sponsor independent

certification of your own artifacts. The physical gaps: HBM, packaging and EUV are absent and CXMT is about three years behind at the class that matters [206]. And submit to scoreboards you do not control: no Chinese accelerator has appeared in MLPerf Training, no LX2-class machine has published the AI proof suite, and the flagship’s autonomous-silicon-design claim ships without a checkable artifact. Dominating the metrics you control and avoiding the ones you do not is rational for a challenger and corrosive for an incumbent, and China is becoming an incumbent.

10 The Prize, the Bill, and the Third Path

In one paragraph

The largest value in AI is not code, documents or enterprise workflow but the acceleration of science, because everything else redistributes an existing surplus and science creates new ones. That turns the price of frontier tokens into a tax on the one activity whose returns compound into everything else. A measurement-anchored sizing framework finds that announced national research infrastructures provide a tenth to a fiftieth of what their scientific populations need, and that a workload share of one seventh becomes a budget share of three quarters to nine tenths once purchased tokens carry a twenty-to-sixty-fold premium. For every country that is neither the United States nor China both available strategies fail—dependence on foreign providers, or an isolated national program at sub-scale—and the third path is an international, institutionally neutral consortium: federated compute, scientific data stewardship, open pretraining as insurance, models with lifecycle obligations, and a certification authority neither bloc controls. It is the constructive branch of an otherwise pessimistic book, and, if it succeeds, a partial falsifier of its strongest claim.

10.1 The prize

What AI for science has delivered follows one pattern: an AI system replaced a physics-simulation-bounded computation and was validated against an established operational baseline. Weather is the strongest case—the European centre’s AI forecast operational since February 2025 at about a thousandth of the energy per forecast [207], the American service’s sixteen-day forecast in forty minutes at 0.3% of conventional compute. Structural biology is the most diffused—more than 200 million predicted structures, three million users, a Nobel Prize, and users submitting more than 40% more novel experimental structures [208]. Mathematics produced the first credible autonomous results, and one drug candidate discovered end to end by AI reached a controlled clinical readout [209]. The failure ledger is printed with equal weight: the most-cited economic study of AI in science was fabricated and repudiated [210]; materials discovery has not survived validation intact; the European centre discontinued all four external AI weather models it had been running [211]; autonomous discovery agents produce mostly flawed candidates; and the journal pipeline is being damaged by AI-generated submissions. It delivers where the task is a well-posed prediction problem with an expensive physics-based incumbent and an established baseline, and has so far mostly failed where the task is judgement.

The prize is nevertheless the largest because of four numbers. The denominator: global R&D is about \$3.1 trillion a year at purchasing-power parity, twice the software market, and frontier-laboratory revenue is under 2% of it [212]. The return: \$4–13 of social benefit in present value per dollar of R&D [213]. The bottleneck: US research productivity has fallen roughly forty-fold since the 1930s, offset only by twenty-three times more effort [214], and the published

macroeconomic estimates for AI—from a fraction of a percent to thirty percent of growth—are separated entirely by whether AI is assumed to automate innovation itself; the most-cited skeptical estimate excludes protein-folding-type advances as unlikely within ten years, and that exclusion is the entire disagreement, because structure prediction, operational AI weather and designed enzymes all arrived inside the window [215]. And the flywheel: AI for science is the only application that feeds back into its own means of production. The book replaces the customary “3–30%” with a value-conversion funnel whose every arrow has an efficiency below one; task-weighted, the central estimate is an uplift near 11% of research output, with a likely range of about 8–18% [M]. Even the pessimistic corner is a flow of roughly \$28 billion a year, and the tenth percentile about \$240 billion; the present value of either exceeds the cost of national infrastructure by a large multiple. Two consequences carry into the strategy. A scientific result produced by an unverifiable artifact through an unauditable chain is not a scientific result, which makes certification a precondition of the prize rather than an accessory to it. And the workload of scientific AI is long-context agentic orchestration over large batches of scientific execution—precisely the architecture the compute chapter described. China leads the Nature Index [216] and runs a decade-old ministry program for AI in science; the West has the Genesis Mission, Europe’s factories and national machines; both sides underprovision by about an order of magnitude [full edition, Ch. 24]. On the Chinese side that programme is older, more explicit and more industrial than the Western reader expects, and it is legible in primary policy text rather than inferred from behaviour; Chapter 11 reads the record, and reads the parts of it that cut against this book.

10.2 The bill: how much compute a nation’s science needs

Sovereignty becomes a checkable number only if the demand side is measured. National AI-for-science infrastructure has been sized by budget envelopes, vendor availability and flagship extrapolation; the full edition follows a measurement-anchored framework developed at RIKEN that replaces those with a vendor-neutral unit and a two-layer demand model [217] [full edition, Ch. 26]. A requirement denominated in BGPU is procurable from any vendor, which is what a public planning framework must guarantee and a vendor-named unit cannot. The demand anchor is a measured 56-day campaign on a GB200-class partition: 101 users, 53,000 jobs, 340,000 GPU-hours, the partition saturated within five weeks [217]. The two layers are token service—researchers driving a trillion-parameter master model at million-token context that delegates to smaller sub-agents—and scientific execution. The structural finding is that ten percent of the tokens, generated at million-token context, cost nine tenths of the token-service GPUs, so national demand is overwhelmingly a long-context memory-capacity and bandwidth problem: the sizing framework and the architecture argument are the same argument in different units.

Three countries are instantiated, and the requirement tracks the researcher population closely because the per-researcher core sits near 0.94 BGPU: Japan needs roughly 497–510 thousand BGPU by 2030, the United States 1.09–1.17 million, Singapore 55–59 thousand, with a band of at least $\pm 40\%$ before local measurement, and the Singapore and US figures should not be quoted more precisely than transferred behavioural assumptions permit. Announced national infrastructures provision 0.02–0.10 BGPU per researcher against a modeled requirement near one—about ten times under for the United States and Japan, fifty times for Singapore. One side has the measurement without the money; the other has the money without the measurement. A nation that rents its master layer has rented the majority of its research computing; sovereignty has a computable price, about \$1.84–2.19 per delivered accelerator-hour, which makes pooling across time zones rational.

10.3 The third path: an international consortium

For every country that is neither the United States nor China, the two available strategies are both unsatisfactory. Permanent dependence on a few foreign commercial providers cedes cost, control and continuity; isolated national programs duplicate the most expensive activity in the industry at sub-scale, at valuations that assume rents no one may collect. What is needed is a *global* form of sovereign capability—not national isolation but shared ownership of the knowledge, data practices, compute access and institutional capacity required to build and operate scientific AI [218] [full edition, Ch. 27].

The decisive argument is financial. In a realistic AI-for-science portfolio, frontier language models and agents are roughly 14% of the physical workload; simulation and domain-specific AI take the rest [218]. That figure is usually deployed as reassurance, and it is the opposite once the price asymmetry is applied: purchased frontier tokens carry an estimated 20–60-fold premium over equivalent service from institutionally controlled infrastructure running open models [A/M], so a workload share of one seventh becomes a budget share of 76.5% at a twenty-fold premium and 90.7% at sixty (Figure 10.1). Whoever sets the frontier token price sets the science budget of every institution that buys it: proposition P3 applied to national research capacity. This is also why the Chinese price collapse is strategically effective without coercion—at a premium near unity, which cheap open weights on domestic hardware deliver, the budget problem dissolves, and the institution that solved it that way has acquired a dependency it did not negotiate.




Figure 10.1: Why the token price is the science budget. With a fraction f of physical work served by purchased frontier tokens at relative unit price ρ , the share of total expenditure that work consumes is $\Sigma_{\text{front}} = f_{\text{front}}\rho / ((1 - f_{\text{front}}) + f_{\text{front}}\rho)$. At the measured $f_{\text{front}} \approx 14\%$ and the estimated 20–60 $\times$ premium, a one-seventh workload share becomes 76.5–90.7% of the budget [A/M]. The curve’s shape, not its exact position, is the policy argument: at r near unity—which cheap open weights on owned hardware deliver—the problem dissolves, and the institution that solved it that way has acquired a dependency it never negotiated.

Open pretraining is therefore insurance, not a benchmark contest: a consortium’s own open model may never be the best production choice, and building it is still correct, because the return on pretraining is people who know how to do it. Near-term science should meanwhile use the best open-weight models, including Chinese ones, without embarrassment. The insurance covers three risks, none presuming bad faith—continuity, verification (Chapter 6), and regulation in dual-use science—and Chapter 11 finds each of the three instantiated rather than hypothetical. Four functions define the institution: federated compute under common accounting, its first proof of concept the joint training of a fifteen-trillion-token open model split across compatible systems in several countries; scientific data stewardship, because science is the one community for which provenance and reproducibility are already requirements; model production with lifecycle

obligations, since a clinical model may take three years to certify and must be supportable five years later—the strongest reason this cannot be a rotating academic project; and professional operation, because convening is cheap and operating is not. Industry participates without expropriation, on an open accountable base with closed private layers above it, and the Global South is treated as talent, questions and data rather than as aid, because the alternative is not neutrality but ceding the default to whoever shows up with a stack.

The organizational lesson from China is not a command structure but the reverse: government sets the stage and allows intense bottom-up activity, so a consortium funding one model team in one country reproduces the failure mode rather than the success. Table 10.1 gives ten operational design parameters and the failure mode each avoids, and the full edition's feasibility appendix finds that of fourteen legal and engineering problems a founding instrument must solve, nine have answers on the shelf, three require legal invention, and none is a reason to wait [full edition, App. F]. The governing framing is neutrality, because a certificate issued by one bloc about the other bloc's artifacts is worth nothing to anybody. Sequence: frame; prove federation with three to five committing countries and one visible cross-border pilot; then institutionalize. A five-country story with a working pilot beats a large inactive membership, and a proposal that names functions but not instruments is a wish.

Table 10.1: The consortium, specified operationally: six of the ten design parameters the full edition states, with the failure mode each avoids [S].

Parameter	Recommended answer	Failure mode if deferred
Legal form, host	Treaty-adjacent foundation in a neutral, non-aligned-capable jurisdiction	Seen as a bloc instrument; loses the Global South and one principal
Membership, votes	Separate classes for states, centres, universities, firms; contribution-weighted votes capped below a blocking share	Largest contributor becomes owner
Export-control compliance	Per-jurisdiction architecture, segmented enclaves	One member's rule governs everyone
Compute allocation	BGPU unit, peer review, reserved and burst tiers, cross-border settlement	Federation degrades to a mailing list
Certification authority	Operates the Level 0–5 ladder, publishes irrespective of origin	Vacancy stays vacant, or is filled by a supplier
Chinese and US technology	Capability- and evidence-based; no origin rule either way	Forfeits neutrality

This is where the book's threads terminate—the institutional form of “take the commons, verify it, contribute to it, and own the layer that makes it safe to use.” Whether it gets built in the next two years or the next ten is, on the evidence of the precedents, the difference between shaping the outcome and documenting it.

11 China's Scientific Machine

In one paragraph

If the largest prize is scientific, the question is who is organized for it, and on the primary policy record the answer is uncomfortable: China has been organizing longer, more explicitly and more *industrially* than any Western actor. The State Council's "AI+" Opinions place science first of six key actions and chain AI-driven research to "engineering realisation and product landing," which is what makes this industrial policy rather than research funding. The output is measured by the European Commission rather than by Beijing, and it has led the world since 2016. The models are open at 241B and trillion-parameter scale under permissive licences, which is why the case for a Western open scientific-AI consortium cannot be that the commons is missing. Three findings cut the other way and are printed at equal weight: the plan's numeric targets measure diffusion and not discovery; the paradigm-transformation sentence does not survive from the Party's Recommendations into the state Outline; and the newest and most capable Chinese platform has no stated open release at all. [full edition, Ch. 25]

11.1 The policy record, read precisely

The State Council's *Opinions on Deepening the Implementation of the "AI+" Initiative*—Guofa [2025] No. 11, signed 21 August 2025—set out six key actions, and "AI+ science and technology" is the first of them, ahead of industry [219] [P]. Ordering in such a document is the sequence in which ministries read their obligations, and putting science ahead of industry in a document whose whole register is industrial says where the leadership believes the compounding lies. The science pillar's third sub-clause is the load-bearing one, and it is why this chapter is titled industrial policy rather than science policy: it directs the integrated development of "AI-driven technology R&D, engineering realisation and product landing." Discovery is the first stage of a pipeline whose last stage is a product, and the text says so. The lineage runs back to the ministry and science-foundation "AI for Science" deployment of March 2023, which Western sources routinely overstate as a plan when it is a deployment of work [220] [P]. The money underneath is bench money spread wide rather than a moonshot: roughly half the project categories in the current science-foundation call foreground AI for science at individual budgets of EUR 240,000–600,000 [221] [A]—enough for students and cluster hours, not frontier training runs: the state supplies compute centrally and grants supply the science.

Two corrections belong in the same section, and both cut against the thesis. First, the numeric commitments in Guofa [2025] No. 11 are *diffusion* targets—intelligent-terminal and agent penetration above 70% by 2027 and 90% by 2030—and there is **no numeric target for science at all** [219] [P]. A planning system whose distinctive competence is hitting what it measures has, in its flagship AI document, measured adoption and declined to measure discovery. Second, and more precisely: the Party's *Recommendations* for the 15th Five-Year Plan, released 28 October 2025, contain the instruction to *use artificial intelligence to lead a transformation of the scientific research paradigm*—and **that phrase does not survive into the Outline adopted by the National People's Congress in March 2026**, where AI appears under compute, algorithm and data supply while basic research appears in its own chapter without AI attached [222] [P]. The Recommendations state ambition; the Outline allocates. The deflationary reading is that a paradigm shift is not a line item; the consequential reading is that the drafting converted *how research should be done* into *what infrastructure should be bought*, and the state's own prize list agrees with it. Either way the gap between Party ambition and state programming is real, documentable, and the strongest single piece of evidence against this chapter.

11.2 The output, the artifacts, and the machine that won the prize

The most useful measurement was made in Brussels, which is the reason to use it. The European Commission's bibliometric study of some three million papers across eighty fields finds that **China overtook both the United States and the European Union in AI-assisted scientific output from 2016**, producing more than 25,000 AI-assisted papers in 2022 against roughly 15,000 for the EU and 12,000 for the US, with reported advantages on novelty and citation impact as well as volume [223] [V on the counts, A on the impact findings]; the Nature Index points the same way [216]. Two limits belong beside it. *Output is not discovery*—papers are the first term of the funnel of Chapter 10, whose later terms are all fractions below one. And *the window ends in 2022*, before the model generation this book is about, so the study measures the run-up rather than the present state: the position was built before the current phase, not during it.

The artifacts matter more, and the fact that matters most about them is not any benchmark number but their licences. Shanghai AI Laboratory's **Intern-S1** carries 241B parameters pre-trained on 5T tokens of which *2.5T are scientific*—a resource-allocation decision no Western laboratory has publicly matched—and is published under **Apache 2.0**; its scores are self-reported and recorded at that grade [224] [P]. ByteDance's **Protenix** is released under the same licence with code, weights and pipelines, claiming to be the first fully open-source model to outperform AlphaFold3 at matched data cutoff, scale and inference budget [225] [P on the release, developer-evaluated on the comparison]—and the significance is the licence, not the benchmark. The open-weight ratchet of Chapter 4 has reached structural biology, and it reached it from a Chinese laboratory; a database is a result and a weight file is a means of production, and the West's most diffused scientific asset was given away as the former. **DeepModeling and DP Technology** are the deepest instance: DeePMD-kit, DPA-2 covering 73 elements across eighteen pre-training datasets, the ABACUS electronic-structure package, the OpenLAM initiative, more than ten thousand community members, with code, weights and datasets all open [226] [P]. On the record at 9 September 2026 that is among the most genuinely open scientific-AI ecosystems anywhere, Western or Chinese—a sentence that cuts against the framing a Western reader will bring. Against which stands the counter-observation, made prominently because it must be: the newest and most capable platform—Intern Duanyan, announced July 2026, a five-layer discovery stack with a dry-lab/wet-lab closed loop—has **no stated open-source release** and is access-by-invitation [224] [P]. The ratchet compels weights onto the commons because adoption is the currency; it compels nothing whatever about a platform, whose rent lives in orchestration and instrument access rather than in parameters. Openness here may be a phase rather than a principle. The same ambiguity runs through the institutions: CAS ScienceOne is built on *Chinese* open foundation models and consumes AlphaFold and MatterGen as tools underneath its own orchestration layer—the mirror mechanism of Chapter 5 in a scientific setting—while its parameters, training compute and open-source status are undisclosed and its performance claims self-reported, so this book credits its architecture and not its capability [227] [P].

The infrastructure supplies the tell. Pengcheng Cloudbrain, the first fully domestic large-scale intelligent computing system, was announced in July 2026 as winner of the **National Science and Technology Progress Award, First Prize**, and is reported to support training for more than half of China's domestic foundation models [228] [P on the award, A on the ratio]. *The state's top science prize went to a computer, not to a discovery made on one*. A national prize list is a revealed preference over what a state believes is scarce, and this one says the scarce thing is not a result but the domestic capacity to produce results—the argument of Chapter 4 restated by the awarding body. Western prize lists decorate the structure, the forecast, the molecule, and treat the machine as procurement; neither choice is obviously right, but the asymmetry is real.

11.3 Scored on the playbook, and what it does to the consortium

On the playbook of Chapter 2 the programme reads cleanly: target selection aimed at the research paradigm itself; demand before supply, with the Outline naming government *purchase and leasing* of compute and universal accessibility for small users; selection rather than planning, in six-plus institutions with overlapping mandates and no designated national laboratory; supply-chain capture running model, framework, code, chip and machine, with a domestic electronic-structure code inserted into the label pipeline of the flagship domestic atomic model [222, 226, 229]. Two steps are not scored: overcapacity has the shape but no published utilization figures, so overbuild cannot be told from stranded provincial assets [S], and gatekeeping is not observed at all—no restriction on scientific model release, no export control on scientific weights. The industrial hinge is the distinctive finding: of ten flagship 2026 results, three are aerospace design, industrial-chemistry formulation and an analogue chip-design agent, the flywheel aimed directly at the layer Chapter 4 identifies as the binding physical constraint—while three of the same ten carry Western co-authors, so any account treating Chinese AI for science as a hermetically closed national system is wrong on the byline evidence [229] [P].

Which sharpens the consortium question of Chapter 10 rather than settling it. The case usually offered for a Western or neutral open scientific-AI consortium—that the commons is thin and public money must create an open base—is, on this evidence, false: the open base exists, and the researchers a consortium proposes to serve are already downloading it. The defensible case is narrower and stronger: *the commons has one principal contributor and no governance the rest of the world is party to*—contribution and governance being different goods. Each of the three insurance risks is sharpened here: continuity, by an unopened platform at the top of the stack; verification, because four of the most consequential artifacts in the field cannot be checked in the way Chapter 6 requires; monoculture, if one national ecosystem supplies the potentials, the scientific foundation models and the agent frameworks a generation of results depends on. The opposing case is strong: a consortium built to *exclude* its largest contributor is a worse instrument than one built to include it, forfeiting both the artifacts and the neutrality that is its only durable asset; and if scientific knowledge is non-rival, a policy that slows the commons in order to change its ownership destroys more value than it redistributes, in the one domain where the social return runs above four dollars on the dollar. This chapter does not resolve that; it establishes the question: whether an open commons with one principal contributor and no shared governance is a commons or a dependency, and what instrument short of exclusion changes the answer.

12 The Convergence

In one paragraph

The elements interlock into a single machine for commoditizing artificial intelligence: open weights and open interfaces, accelerators that need only be sufficient and numerous, and decentralized pretraining and inference that dissolve the centralized-supercluster premise. The Western AI economy is a bet that intelligence will remain scarce and rentable; the Chinese strategy is a bet that it can be made abundant and free, and that whoever makes it abundant collects the durable prizes. Solar, batteries and EVs each ended with the scarcity bet losing. The strongest counter-argument—recursive self-improvement on closed American compute—gets its own section, and two conclusions change the argument’s shape: the commons has a security bill that falls due at scale, and there is a third path whose economics are better than its politics suggest. [full edition, Ch. 28]

12.1 The machine the elements form

Software commoditization: CUDA’s advantage is not being crossed in one leap, but the population equipped to erode it is being made as large as the ecosystem itself. *Hardware proliferation:* with the software requirement lowered, accelerators need not match NVIDIA, only be sufficient and numerous, paired with cheap commodity memory from a domestic fab—sufficiency plus volume is how TOPCon modules and LFP packs won. *Decentralized production:* the half-megawatt pretraining cluster and the five-thousand-dollar appliance of Chapter 4 [M/S] put the means of production within reach of any state, consortium or mid-cap firm.

The Western AI economy is, structurally, a bet that intelligence will remain scarce and rentable: \$286 billion of annual private investment [10] capitalized against future margins on closed capability, housed in HBM-bound gigawatt data centers—this war’s equivalent of Q-Cells’ 2007 capacity crown, magnificent, state-of-the-art, and aimed at the wrong future.

12.2 What would falsify this thesis

Chapter 13 states the six observations that would move the argument furthest, and Appendix A carries the rest against a horizon and a probability; two govern this chapter’s own structure. The thesis weakens materially if closed US models open a *compounding* capability gap that adoption metrics start to follow, or if recursive AI-for-AI research gives the closed frontier an escape velocity that open fast-following cannot match—the branch treated next. It strengthens with each frontier-class model trained end to end on domestic silicon; the first credible such run will be this war’s Mate 60 moment.

12.3 The strongest counter-argument: two recursive loops, not one

The escape-velocity branch is the only falsification path that would make everything else in this book irrelevant rather than merely wrong. In its strongest form: if the closed US laboratories reach genuine recursive self-improvement—models that materially accelerate the research producing their successors—then compute scale compounds into capability faster than any fast-follower can copy, and the open ecosystem’s cost advantage is competing on a curve that is being outrun. The evidence is no longer hypothetical: OpenAI attributed the 30 July cut of Chapter 4 to efficiency gains its own frontier model had helped produce [178, 230], which means the same event this book cites as the first direct observation of P3’s transmission is *also* the first public claim of a model improving the economics of its own successor. Both readings are available from one datapoint.

The reply is not that recursive self-improvement is unreal but that there are two recursive loops running on different substrates, and they do not compete on the same axis. The American loop runs through algorithmic efficiency, compounding wherever compute is most abundant. The Chinese loop runs through industrial capability—models improving chip design and manufacturing yield, compounding wherever the manufacturing ecosystem is deepest—and the claimed autonomous silicon-design flow of Chapter 4 is that structure aimed at the layer Chapter 10 argues matters most. The first loop produces more capability per FLOP; the second more FLOPs per dollar, domestically, out of an ecosystem that currently must lease them abroad. Which dominates is undetermined, and their failure modes differ: the American loop fails if capability gains stop converting into products people will pay a premium for, the Chinese loop if AI-assisted design cannot substitute for the tacit process knowledge that has defeated Chinese industrial campaigns before, in aircraft and lithography. Neither failure is currently observable; the first audited instance of either loop closing will be the most important datapoint in the field.

One consequence of Chapter 5 changes what the evidence is worth. Because P3’s mechanism is now overdetermined—the largest supplier of the frontier’s inputs releases weights with their training data, and the frontier’s largest customers build the silicon underneath it—the proposition is satisfiable by the incumbent’s own conduct, which strengthens its mechanism and weakens its *diagnostic value for this book’s thesis*: observing P3 hold is no longer, by itself, observing the thesis hold.

12.4 Two things the argument has to carry with it

Two chapters change the shape of the conclusion. **The commons has a security bill, and it falls due at scale.** The convergence is, in the vocabulary of Chapter 6, a forecast of correlated deployment. Monoculture risk is origin-neutral—a closed US-led monoculture would carry the same correlation with less transparency—so this does not falsify the thesis, but it identifies the highest-value piece of unbuilt infrastructure: an authority that maintains artifact lineage across a derivative chain no vendor controls, tests integrations rather than models, and publishes results neither bloc’s marketing owns. Whether the commons is adopted safely is currently nobody’s job. **There is a third path, and its economics are better than its politics suggest.** For any country that is not a superpower the decisive number is a budget, not a benchmark (Chapter 10), which makes an institutionally neutral consortium the cheapest available form of strategic autonomy and the only credible home for that authority—and Chapter 11 sharpens rather than settles the case, because the scientific commons is not missing but has one principal contributor and no governance the rest of the world is party to. A coalition that acts holds something neither Washington nor Beijing can revoke; one that does not will discover that “open” and “available to us, indefinitely, on acceptable terms” were never the same word.

12.5 Implications for the West—and for the rest

The losing move, as Chapter 9 argues, is to defend the rent: by the time the wall is complete, the world outside it is running on the commodity. For third countries—Japan included—the open stack is *open*: profiles, recipes, system designs and frontier weights can be adopted, forked and sovereignly operated by anyone, and the half-megawatt pretraining cluster is as available to Kobe as to Guiyang. If the AI war ends the way the precedent wars ended, the strategic question for everyone who is not the United States becomes not *whose side wins*, but *who builds fastest on the commons that the winner created*. That, and not any single benchmark, is how China will win the AI war: not by reaching the frontier first, but by making the frontier free—and owning everything beneath it.

13 The Argument in Summary

In one paragraph

The question is not who builds the better model: the measured gap at the top is a few index points—eight on the re-based index [50]—and has been narrow for over a year. The question is where the *rent* lands when intelligence becomes cheap, because all three precedent campaigns were won that way and none with a better artifact. The answer comes in three parts. The strategy is real, coherent and further along than a layer-by-layer reading suggests. The method is no longer Chinese. And the outcome is not settled, because the layers that stay scarce are still mostly Western-held, and a West that commoditizes competently captures the surplus of a cheap-intelligence world as surely as a China that does. [full edition, Ch. 29]

13.1 The answer in three parts

First, the strategy is real and further along than a layer-by-layer reading suggests. Every one of the playbook’s nine moves is occupied in the AI campaign in a recognizable form and at a datable stage (Chapter 3): not a plan on paper but an execution in progress.

Second—and this is the finding that most changed the book—the method is no longer Chinese. Over 2026 the largest Western firms converged on the same instrument, each commoditizing a layer it does not own to defend the layer it does (Chapter 5): the accelerator vendor gives away weights with their training data and has now bought the platform and the people that make them. That strengthens Chapter 2—a playbook that predicts the conduct of actors it was not derived from is doing more work than one that does not—and complicates every score in Chapter 4, because a layer being commoditized by two parties at once tells you less about either of them.

Third, the outcome is not settled, and this book’s own thesis is the claim most at risk from its own findings. If every layer is cheap the surplus accrues to whoever holds the layer that stays scarce, and at 9 September 2026 those—the installed programming model, the enterprise deployment surface, trust and certification, deployment-generated data—are still mostly Western-held. The position is therefore conditional: China is executing the strategy that won three technology wars, it is executing it well, and whether that wins a fourth depends on where the rent lands, which is not yet observable.

Table 13.1: The four propositions of Section 1.4 at 9 September 2026: the grade the evidence supports, not the probability the author assigns [full edition, App. J].

	Claim	Where the evidence stands	Status
P1	Open weights become the default for deployed intelligence	Majority of routed tokens; a ratchet that survived a Western re-entry; the gap at the top eight index points on the re-based board; volume on the open sparse tier, spend in the closed frontier	Distribution done; monetization contested
P2	China builds a sovereign, capacity-first stack to serve it	Open ISA qualified as a CUDA host; a GPU-free machine at No. 1 on HPL and HPCG; production serving on a domestic cluster; HBM still binding, the training test not passed	Trending; integration test outstanding
P3	The rent financing the Western frontier is destroyed or relocated	Mechanism overdetermined—the incumbent supplier funds open weights itself—but its datacentre revenue grew 117% in this quarter, and its cheapest tier undercut the commodity and took share	Mechanism operating; magnitude contested
P4	The position converts into durable strategic advantage	Endgame begun on consolidation, not on gate-keeping; the scarce layers still mostly Western-held; deployment-generated data and trust unresolved	Pending; least evidenced

The coupling matters more than any row. P1 without P2 is adoption with no hardware base to serve it, which the precedents say decays; P2 without P1 is a sovereign stack with nothing running on it; P3 without P4 is a price war that destroys a rent pool and hands it to nobody—live now that both sides are destroying each other’s rents at once. *P4 is where the title’s claim rests and where the argument is weakest*; attack it there rather than at the benchmarks.

13.2 What would change the answer

Six observations would move the author’s position furthest before the end of 2027; each is in the dashboard of Appendix A.

1. **A frontier-class model trained end to end on domestic silicon, packaging and memory**—the one integration test that moves P2 from trending to done. Production-scale *serving* has been claimed, and a leading laboratory has ordered a year of domestic

accelerators for inference while keeping training on NVIDIA; the one frontier-scale domestic-training claim in circulation has no locatable primary, and that distinction is the whole of the remaining gap.

2. **Merchant-GPU datacentre revenue growth decelerating while merchant custom-ASIC revenue holds its guided trajectory**, which confirms the silicon-commoditization pincer of Chapter 5. The contrary observation—custom accelerators converging onto the incumbent’s interconnect and memory, as one already has—confirms absorption instead, which currently has the better of the evidence.
3. **A sustained fall in closed-frontier gross margin, disclosed**: P3’s transmission on the line that matters. One cut to a cheapest tier, with the flagship held and the counterparty raising prices in the same fortnight, is not that.
4. **Enterprise open-weight *spend* share turning upward**. The spend share fell during the period capability reached parity; if it turns, the deployment-layer objection weakens sharply.
5. **Gatekeeping exercised on any AI layer**—the endgame move with a template in batteries and rare earths, not yet used here, whose use would settle P4 quickly and in one direction.
6. **An audited instance of a research loop closing**—a model materially accelerating the science that produces its successor, or the silicon that runs it. This would make the recursive-improvement counter-thesis of Chapter 12 decisive against everything else in this book, and it is the falsifier the author takes most seriously.

13.3 What this book is not claiming

Four disclaimers, because a book with this title attracts readings it does not support. It is **not claiming that China has won**, or that victory is inevitable, but that a recognizable strategy is being executed competently against a stack whose rent layers are being dismantled from both sides, and that the outcome turns on an ownership question not yet observable.

It is **not claiming that the West is passive**. The evidence of 2026 is the opposite: the incumbent priced underneath the commodity at the bottom of its ladder and took share doing it [59], grew datacentre revenue 117% in the quarter this edition closes in [131], and runs the same commoditization strategy with more capital behind it than any Chinese actor commands.

It is **not claiming that openness is a Chinese property** or a Chinese virtue. Openness is an instrument, available to anyone with an adjacent layer to defend, and the largest single act of open-model funding in the period this book covers was American [123].

And it is **not claiming more confidence than the evidence carries**. The load-bearing weaknesses, without softening: several mid-2026 details rest on secondary aggregators; the capacity-first architecture of Chapter 4 is a modeled design with no measured realization; the compute sizing is measurement-anchored for one country only; and the judgment that breadth compounds faster than depth is the claim most likely to be shown wrong. Both editions are living documents, and the register of Appendix A exists so that this argument can be scored rather than believed.

A Forecasts and Dashboard

The full edition carries a graded evidence ledger, a machine-readable claim register, and the two tables reproduced here [full edition, App. D, J]. Every forecast is a scenario judgment [S] with a probability, a horizon and a falsifier, so that a reader in 2029 can mechanically score the book; a skeptic’s figure, where the full edition prints one, appears in brackets. Two terms are the full edition’s: “stage two” is the second rung of its proof ladder for the integrated-CPU claim (parity on AI serving benchmarks after the settled HPL/HPCG rung), and the “Max-tier template” is the pattern of opening a flagship’s weights while gating commercial use above a revenue threshold.

Table A.1: Probabilistic forecasts as of 9 September 2026 [S].

Proposition	Prob.	Horizon	Falsifier
Chinese models dominate the open-weight ecosystem	80% [70]	1–2 yr	Routed-token share below 45% for four quarters
Frontier-adjacent model trained end to end on the domestic Chinese stack	65% [50]	2–4 yr	Reversion to foreign silicon past 2028
Capacity-first (LPDDR-dominant or hybrid) pretraining production-competitive	45%	3–6 yr	MFU below $\sim 45\%$ at scale, or bandwidth knee above 8 TB/s
LX2-class integrated CPU competitive on AI benchmarks	55%	2–4 yr	No submission by 2029, or $>2\times$ deficit
Broad ($\geq 30\%$) US AI repricing	70% [60]	1–3 yr	AI revenue closing the capex gap through 2028
US closed models retain the premium enterprise tier through 2030	50% [65]	4 yr	Open-weight enterprise spend above 40%
Complete Chinese semiconductor-tooling autonomy	25%	5–10 yr	No domestic EUV-class lithography past 2032
Neutral certification authority for open weights established	30%	3–5 yr	Role occupied by a single bloc, or by nobody, past 2030
AI-for-science uplift $\geq 5\%$ measurable nationally	45%	5–8 yr	No uplift in weather, structure, or screening
Systemic agentic-security incident from a widely deployed derivative	20%	5 yr	(asymmetric: absence is weak evidence)
Chinese accelerator or integrated CPU submits to MLPerf Training	50%	2–3 yr	None by end-2028
$\geq 1,000$ humanoids productive under a single operator, metrics published	55%	2–4 yr	Teleoperation ratios not declining through 2028
Next $\geq 1\text{T}$ -parameter Chinese release ships open (the ratchet holds)	85%	1–2 yr	A major Chinese lab closing its frontier tier
Chinese agent harness top-three in a major non-Chinese market	35%	2–4 yr	Western harnesses $>90\%$ of regulated deployments through 2029
NVIDIA ships a frontier-scale open-weight model with derivative licence	55%	2–3 qtr	Closed, or not shipping, through 2027
Compute futures reach sustained liquidity	40%	2–3 yr	Thin trading through 2028
Chinese inference capacity exported at scale	45%	3–5 yr	Endpoint share stalling while self-hosting carries volume
HBF ships in volume at $\geq 512\text{ GB/stack}$, $\geq 1.6\text{ TB/s}$, with an endurance figure	45%	2–4 yr	None by 2029, or below $\sim 0.8\text{ TB/s}$
Capacity-first appliance: $\geq 2\text{T}$ open model, 4 users at $\geq 40\text{ tok/s}$, 1 kW, under \$8k	35%	3–5 yr	Above \$15k or below 20 tok/s through 2030
Enterprise open-weight workload share reverses and passes 20%	40%	2–4 yr	Continued decline (licensing and governance binding, not memory)
$\geq 100\text{ MW}$ AI facility sold, foreclosed, or written down $\geq 40\%$ below cost before 2030	45%	2–4 yr	All 2025–26 vintages refinance at par through 2029

continued

Table A.1, continued

Proposition	Prob.	Horizon	Falsifier
Chinese frontier model trained on domestically sited compute, offshore excluded	60%	2–4 yr	Leased offshore reliance past 2029
AI-assisted design flow produces a verified, manufacturable Chinese accelerator	45%	3–5 yr	Roadmaps slip despite AI-design claims
Merchant-GPU revenue growth decelerates while custom-ASIC revenue holds Broadcom’s guided trajectory (the pincer of Chapter 5)	50%	end-2027	sustained GPU growth beside flat custom-ASIC revenue
Custom accelerators keep converging onto NVIDIA’s interconnect and memory, as Trainium4 does, rather than displacing them (absorption)	65%	end-2027	a major custom part at scale on a wholly non-NVIDIA stack
First confirmed Ascend export outside China (Egypt or Malaysia)	40% [25]	end-2027	No confirmed export by then

Table A.2: Quarterly falsification dashboard.

Indicator	Current	Threshold	Source
Chinese share of routed tokens	~61%	Sustained <45% weakens P1	Neutral routers, monthly
Closed-over-open capability gap	3.3% (AI Index); 8 points (AA v4.3)	>10% sustained two years favours bifurcation	AI Index, arenas
Enterprise open-weight spend share	~11%	>25% confirms P3 transmission	Surveys, annual
Domestic end-to-end frontier training run	Not observed; the circulating GLM-5 claim has no primary [A/C]	First credible instance confirms P2	Lab reports
LX2-class AI benchmark submission	None	MLPerf at parity confirms stage two	MLPerf rounds
CXMT LPDDR6 / HBM3 ramp	Verification cleared	2+ year slip weakens the memory arbitrage	Quarterly
Cost per solved task, open vs closed	Unpublished	First credible series confirms or deflates the gap	Evaluations
OpenAI financing	IPO slipping	Down-round or new backstop raises P4	Filings
Data-center credit spreads	Oracle CDS ~212 bps	Sustained widening raises P4	Daily
N. Virginia wholesale colocation rent, 10 MW+	\$155–185/kW-month, 0.5% vacancy	Sustained decline with rising vacancy: utilization forecast failing; a rise: closet not displacing hall	CBRE/JLL, semiannual
Certification authority for open weights	Vacant	Occupation reshapes the trust and consortium chapters	Standards bodies
Max-tier openness lag	~10 days realized, 1 week committed (Qwen3.8-Max)	Lag >2 quarters, or a frontier tier closed, breaks the ratchet	Per release
Recursive-loop disclosures, US vs China	Both claimed, neither audited	First audited instance is the field’s most important datapoint	Lab disclosures, tape-outs
Offshore compute leasing by Chinese labs	Permitted, reportedly in use; a Commerce rule to bar third-country NVIDIA rentals reported in draft, unpublished	Closure by rule is the highest-impact tightening available to export policy; the draft is that tightening, not yet a control	BIS or Commerce

Sources and Selected References

The sources for the claims this edition cites, numbered in order of first citation, each carrying its identifying detail and its full URL. The full edition's bibliography runs to some seven hundred entries, cites every figure in this text—including the many that this abridgement, for brevity, leaves uncited—and carries the analytical annotation on each source that space here does not permit. Evidence grades used throughout: [V] independently verified, [P] primary-source claim, [A] analyst estimate, [M] author-modeled result, [R] roadmap, [S] scenario judgment. The edition's own load-bearing weaknesses are stated in its final chapter rather than repeated here. Both editions are living documents; the title-page date is the currency date, and corrections are welcome and will be credited.

- [1] Our World in Data, "Learning curves: solar module prices fell ~20% per doubling of cumulative capacity. <https://ourworldindata.org/learning-curve>
- [2] A. Jäger-Waldau / ISES, "PV Industry Trends" (Oct. 2024). https://www.ises.org/sites/default/files/webinars/2024/October/Oct242024_PVIndustryTrends_final.pdf
- [3] SNE Research via electrive, "CATL continues to dominate global battery market". . . <https://www.electrive.com/2026/02/04/sne-research-catl-continues-to-dominate-global-battery-market/>
- [4] BloombergNEF, "Lithium-ion battery pack prices fall to \$108/kWh". . . <https://about.bnef.com/insights/clean-transport/lithium-ion-battery-pack-price-s-fall-to-108-per-kilowatt-hour-despite-rising-metal-prices-bloombergnef/>
- [5] Rhodium Group, "Why are Chinese EVs so cheap?". . . <https://rhg.com/research/why-are-chinese-evs-so-cheap/>
- [6] CnEVPost, "China NEV sales Dec. 2025 (CAAM)". <https://cnevpost.com/2026/01/14/china-nev-sales-1-71-million-dec-2025-caam/>
- [7] HSF Kramer, "China battery technology export restrictions" (Jul. 2025). <https://www.hsfkramer.com/notes/mining/2025-posts/china-battery-technology-export-restrictions-cathode-lithium-gallium>
- [8] CSIS, "Rare earth export restrictions. <https://www.csis.org/analysis/rare-earth-export-restrictions-one-year-later>
- [9] Hugging Face, "State of Open Source, Spring 2026". . . <https://huggingface.co/blog/huggingface/state-of-os-hf-spring-2026>
- [10] Stanford HAI, *AI Index Report 2026*. . . <https://hai.stanford.edu/ai-index/2026-ai-index-report>
- [11] Reuters (exclusive), "China to publish policy to boost RISC-V chip use nationwide". . . <https://money.usnews.com/investing/news/articles/2025-03-03/exclusive-china-to-publish-policy-to-boost-risc-v-chip-use-nationwide-sources>
- [12] CRN Asia, "Huawei's open source gambit. <https://www.crnasia.com/news/2025/artificial-intelligence/huawei-s-open-source-gambit-chinese-tech-giant-makes-superpo>
- [13] Design & Reuse, "SMIC achieves 5nm production on N+3 node without EUV tools". . . <https://www.design-reuse.com/news/202529830-chinese-smic-achieves-5-nm-production-on-n-3-node-without-euv-tools/>
- [14] CNBC, "CXMT closes 466% higher in Shanghai debut". . . <https://www.cnbc.com/2026/07/27/cxmt-china-market-debut-chipmaker-ipo.html>
- [15] PRC MFA, "Global AI Governance Action Plan" (WAIC, Jul. 2025), full text. https://www.fmprc.gov.cn/mfa_eng/xw/zyxw/202507/t20250729_11679232.html
- [16] N. Lambert, "On China's open-source AI trajectory". . . <https://www.interconnects.ai/p/on-chinas-open-source-ai-trajectory>
- [17] Eco-Business/Bloomberg, "Chinese renewable companies slow to tap \$47 billion credit". . . <https://www.eco-business.com/news/chinese-renewable-companies-slow-to-tap-47-billion-credit/>
- [18] Bernreuter Research, "Polysilicon price trend". . . <https://www.bernreuter.com/polysilicon/price-trend/>
- [19] TaiyangNews (NEA data), "China installed over 315 GW AC new solar PV capacity in 2025". <https://taiyangnews.info/markets/china-installed-over-315-gw-ac-new-solar-pv-capacity-in-2025>
- [20] EnergyTrend solar price index (Jul. 29, 2026). <https://www.energytrend.com/solar-price.html>
- [21] SCMP, "China's top leadership takes aim at disorderly low-price competition" (Jul. 2025). <https://www.scmp.com/economy/china-economy/article/3316616/china-s-top-leadership-takes-aim-disorderly-low-price-competition>
- [22] pv magazine, "Beijing group creates official body to guide China's polysilicon capacity". . . <https://www.pv-magazine.com/2025/12/11/beijing-group-creates-official-body-to-guide-chinas-polysilicon-capacity/>
- [23] CGTN, "China scraps white list of domestic auto battery suppliers" (Jun. 2019). <https://news.cgtn.com/news/2019-06-25/China-scraps-white-list-of-domestic-auto-battery-suppliers-HO26IZLSWI/index.html>
- [24] IEA, *Global EV Outlook 2026*, "Electric vehicle batteries" chapter. <https://www.iea.org/reports/global-ev-outlook-2026/electric-vehicle-batteries>
- [25] S. Kennedy, CSIS, "The Chinese EV dilemma. <https://www.csis.org/blogs/trustee-china-hand/chinese-ev-dilemma-subsidized-yet-striking>
- [26] CnEVPost, "China NEV retail Dec. 2025 (CPCA)". <https://cnevpost.com/2026/01/07/china-nev-retail-1-387-million-dec-2025-preliminary-cpca/>
- [27] N. Ahrens, *China's Competitiveness: Myth, Reality, and Lessons for the United States and Japan—Case Study: Huawei*, CSIS (Feb. 2013). https://csis-website-prod.s3.amazonaws.com/s3fs-public/legacy_files/files/publication/130215_competitiveness_Huawei_casestudy_Web.pdf
- [28] Dell'Oro Group telecom-equipment share series as reported. <https://www.delloro.com/news/ran-market-stabilized-in-2025/>
- [29] C.-W. Yap, "State support helped fuel Huawei's global rise," *Wall Street Journal* (25 Dec. 2019). <https://www.lightreading.com/business-management/no-one-wants-to-talk-about-huawei-s-state-subsidies>
- [30] LexisNexis IPlytics, *5G SEP Report* (Jan. 2025). <https://www.lexisnexisip.com/wp-content/uploads/2025/01/2025-LexisNexis-5G-SEP-Report.pdf>
- [31] The pushback and what it removed: Australia (Aug. 2018). <https://www.lightreading.com/5g/huawei-has-hardly-been-weakened-in-european-5g-data-shows>
- [32] G. DiPippo et al., "Red Ink: Estimating Chinese Industrial Policy Spending in Comparative Perspective," CSIS (2022). https://csis-website-prod.s3.amazonaws.com/s3fs-public/publication/220523_DiPippo_Red_Ink.pdf
- [33] Tianxia Gongchang Research, "China semiconductor equipment 2026". . . <https://faxiangongchang.com/en/reports/china-semiconductor-equipment-2026>
- [34] TrendForce, "Decoding China's lithography push to challenge ASML". . . <https://www.trendforce.com/news/2025/11/10/news-decoding-chinas-lithography-push-to-challenge-asml-from-scarrier-to-alternative-euv-paths/>
- [35] Pharmsource, "China biopharma out-licensing surges to record \$137.7B in 2025". <https://pharmsource.global/content/china-biopharma-out-licensing-surges-to-record-137-7b-in-2025-2026-on-pace-to-break-it-again/>
- [36] MacroPolo Global AI Talent Tracker 2.0 (Mar. 2024; NeurIPS 2019 vs 2022 authors, top-tier ≈ top 20%). <https://macropolo.org/interactive/digital-projects/the-global-ai-talent-tracker/>
- [37] TechCrunch, "China tells its tech companies they can't buy AI chips from Nvidia" (Sep. 17, 2025). <https://techcrunch.com/2025/09/17/china-tells-its-tech-companies-they-cant-buy-ai-chips-from-nvidia/>
- [38] Reuters (exclusive), "China bans foreign AI chips from state-funded data centres" (Nov. 2025). <https://www.aol.com/articles/exclusive-china-bans-foreign-ai-090709029.html>
- [39] DeepSeek-V3 Technical Report (arXiv:2412.19437). <https://cyber.fsi.stanford.edu/publication/taking-stock-deepseek-shock>
- [40] Yahoo Finance, "Nvidia stock plummets, loses record \$589 billion as DeepSeek prompts questions over AI spending". . . <https://finance.yahoo.com/news/nvidia-stock-plummets-loses-record-589-billion-as-deepseek-prompts-questions-over-ai-spending-135105824.html>
- [41] South China Morning Post, reporting an Interconnects AI analysis of Hugging Face data (10 Apr. 2026). <https://www.scmp.com/tech/big-tech/article/3349552/alibabas-qwen-family-capabilities-over-50-global-open-source-downloads-report-finds>
- [42] Tom's Hardware, "Moonshot releases 2.8-trillion-parameter Kimi K3 as China works around U.S. <https://www.tomshardware.com/tech-industry/artificial-intelligence/moonshot-releases-2-8-trillion-parameter-kimi-k3>
- [43] OpenRouter, "The open-weight models that matter, June 2026". . . <https://openrouter.ai/blog/insights/the-open-weight-models-that-matter-june-2026/>
- [44] M. Casado (a16z) via The Economist and X (Nov. 2025). <https://x.com/martincasado/status/1990462245541982546>
- [45] Alibaba / Qwen, model repositories *Qwen/Qwen3.8-2.4T-A95B*. . . <https://huggingface.co/Qwen/Qwen3.8-2.4T-A95B>

- [46] Moonshot AI, “Kimi K3 License”... <https://huggingface.co/moonshotai/Kimi-K3/blob/main/LICENSE>
- [47] Zhipu AI / Z.ai, GLM-5.3 (announced 13–14 Aug. 2026; API from 18 Aug.; open weights released 28 Aug. 2026). <https://huggingface.co/zai-org/GLM-5.3>
- [48] Zhipu AI / Z.ai, GLM-5.3-Flash (released 26 Aug. 2026; revealed on 27 August to have been the stealth model “Ox Alpha”). <https://huggingface.co/zai-org/GLM-5.3-Flash>
- [49] Artificial Analysis intelligence index, launch week... <https://artificialanalysis.ai/models/glm-5-3-flash>
- [50] Artificial Analysis, Intelligence Index v4.2 (4 Sep. 2026) and **v4.3 (7 Sep. 2026)**. <https://artificialanalysis.ai/articles/artificial-analysis-intelligence-index-v4-3>
- [51] Artificial Analysis independent sub-scores for GLM-5.3-Flash on v4.3. <https://artificialanalysis.ai/models/comparisons/glm-5-3-flash-vs-gpt-5-6-luna>
- [52] Unsloth quantization measurements for GLM-5.3-Flash. <https://unsloth.ai/docs/models/glm-5-3-flash>
- [53] Z.ai’s domestic-silicon claim for GLM-5.3-Flash and what it does and does not say. <https://www.scmp.com/tech/big-tech/article/3365433/zhipu-ai-s-hard-jump-viral-ox-alpha-model-revealed-glm-53-flash-chinese-chips>
- [54] Prior Zhipu domestic-training claims, for calibration. https://www.theregister.com/2026/01/15/zhipu_glm_image_huawei_hardware/
- [55] On the widely repeated claim that the GLM-5 family was trained on roughly 100,000 Huawei Ascend 910B with no NVIDIA hardware at any stage (June 2026). <https://www.hiascend.com/activities/dynamic-news/648>
- [56] Menlo Ventures 2025 enterprise LLM survey. <https://www.luizneto.ai/enterprise-open-weight-models-2026/>
- [57] Epoch AI, “LLM responses to benchmark questions are getting longer over time”... <https://epoch.ai/data-insights/output-length>
- [58] OpenAI, “Advancing the price-performance frontier with GPT-5.6” (30 Jul. 2026). <https://openai.com/index/advancing-the-price-performance-frontier-with-gpt-5-6/>
- [59] The consequence of the 30 July cut, and the fact most easily missed. <https://openrouter.ai/rankings>
- [60] OpenRouter public rankings, week to 29 August 2026, by tokens routed. <https://openrouter.ai/rankings>
- [61] Caixin Global (14 Aug. 2026). <https://www.caixinglobal.com/2026-08-14/deepseek-launches-v4-pro-and-raises-api-prices-by-as-much-as-1100-102473919.html>
- [62] OpenRouter rankings, data through 8 September 2026, by tokens. <https://openrouter.ai/rankings>
- [63] Vercel AI Gateway. <https://vercel.com/i/open-weight-models>
- [64] P. Belcak et al. (NVIDIA Research), “Small Language Models are the Future of Agentic AI,” arXiv:2506.02153 (2 Jun. 2025). <https://arxiv.org/abs/2506.02153>
- [65] Reuters exclusive (Liam Mo and Fanny Potkin, 26 Aug. 2026), three people familiar with the talks. <https://finance.yahoo.com/technology/ai/articles/exclusive-chinas-moonshot-talks-microsoft-075340033.html>
- [66] Moonshot AI, Kimi K3 licence (Hugging Face, moonshotai/Kimi-K3/LICENSE). <https://huggingface.co/moonshotai/Kimi-K3/blob/main/LICENSE>
- [67] Reuters (7 Aug. 2026, two people familiar). <https://finance.yahoo.com/technology/ai/articles/alibaba-plans-revenue-sharing-next-131417995.html>
- [68] What the US hyperscalers already serve, from their own documentation. <https://aws.amazon.com/legal/bedrock/third-party-models/>
- [69] G. Webster and Y. Xu, “Kimi’s tightrope: openness and revenue,” *Transpacifica* (27 Jul. 2026)—“with revenue comes regulability”. <https://herecomes.transpacifica.net/p/kimis-tightrope-openness-and-revenue>
- [70] FY2026 National Defense Authorization Act §1532 (enacted 18 Dec. 2025). <https://www.congress.gov/bill/119th-congress/senate-bill/2177/text>
- [71] Reuters (7 Jul. 2026, three sources). <https://time.com/article/2026/07/07/china-ai-models-alibaba-bytedance/>
- [72] M. Sayed Ali, A. Novik, A. Boddupally et al. (Writer, Inc.)... <https://arxiv.org/html/2607.06906v1>
- [73] DeepSeek, *DeepSeek Harness* (dsh), MIT licence, developer preview v0.1, released 13 Aug. 2026. <https://github.com/deepseek-ai/deepseek-harness>
- [74] DeepSeek, *DeepSeek Harness* documentation. <https://deepseek.com/harness/en/>
- [75] Reuters via Yahoo Finance (13 Aug. 2026). <https://finance.yahoo.com/news/deepseek-releases-official-v4-pro-125608280.html>
- [76] Cyberspace Administration of China (17 Jul. 2026, WAIC)... <https://www.geopolitechs.org/p/china-s-cac-releases-global-cooperation>
- [77] China Briefing, “China’s 15th Five-Year Plan. <https://www.china-briefing.com/news/chinas-15th-five-year-plan-key-insights-for-foreign-investors/>
- [78] CSIS, “Sustaining standards leadership. <https://www.csis.org/analysis/sustaining-standards-leadership-united-states-cannot-disengage-risc-v>
- [79] Hot Chips 2026 deck 09, Frans Sijstermans (NVIDIA), “RISC-V CPUs in NVIDIA Servers. <https://chipsandcheese.com/p/hot-chips-2026-cuda-targets-risc>
- [80] Hot Chips 2026 deck 07, Krste Asanović (RISC-V International, SiFive, UC Berkeley), “Update on RISC-V Standards and Adoption. <https://riscv-europe.org/summit/2026/media/proceedings/2026-06-10-RISC-V-Summit-Europe-09h00-ASANOVIC%20C4%86-slides.pdf>
- [81] SiFive, “BigSky SF-2U870” (24 Aug. 2026). <https://www.hpcwire.com/off-the-wire/sifive-brings-risc-v-to-rackable-servers-with-bigsky-platform/>
- [82] TOP500, “LineShine debuts at No. 1 as TOP500 enters new global exascale era”. <https://top500.org/news/lineshine-debuts-no-1-top500-enters-new-global-exascale-era/>
- [83] J. Dongarra, “Report on the Chinese LineShine System,” Univ. of Tennessee Tech Report ICL-UT-26-01... <https://glencklockwood.com/garden/attachments/dongarra-lineshine-0623.pdf>
- [84] J. Dongarra, T. Hoefler, and S. Matsuoka, “Do We Still Need GPUs? Where an Evolved CPU Can Replace Them in AI and... <https://www.hpcwire.com/off-the-wire/sifive-brings-risc-v-to-rackable-servers-with-bigsky-platform/>
- [85] J. Dongarra, T. Hoefler, and S. Matsuoka, “Do We Still Need GPUs? Rethinking AI and Scientific Computing on... <https://www.hpcwire.com/off-the-wire/sifive-brings-risc-v-to-rackable-servers-with-bigsky-platform/>
- [86] MLCommons, MLPerf Training v6.0 results (16 Jun. 2026). <https://mlcommons.org/2026/06/mlperf-training-v6-0-results/>
- [87] NVIDIA, Q2 FY2026 results (zero H20 sales to China-based customers). <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2026>
- [88] TradingKey, “Domestic AI chip shipments ~1.65M units, 41% of China market in 2025. <https://www.tradingkey.com/analysis/stocks/us-stocks/261889036-nvidia-china-revenue-airforce-one-h200-zero-sales-huang-trump-visit-tradingkey>
- [89] Xinhua, “Huawei fully open-sources CANN” (Aug. 7, 2025). <https://english.news.cn/20250807/35025b0c1366473f97629e0c5ad88f/c.html>
- [90] Bloomberg (4 Sep. 2026, anonymous sources), relayed by TechNode and Mobile World Live. <https://technode.com/2026/09/07/deepseek-huawei-ascend-950dt-chips-inner-mongolia/>
- [91] On the supply side of any Ascend export. <https://www.aljazeera.com/economy/2026/6/1/us-says-ban-on-ai-chip-shipments-applies-to-chinese-firms-outside-china>
- [92] *The Information* (28 Aug. 2026), relayed by Tom’s Hardware. <https://www.tomshardware.com/tech-industry/policy/new-us-export-control-s-reportedly-target-chinese-access-to-remote-ai-servers-trump-admins-cut-down-ai-diffusion-rule-could-be-shared-with-industry-as-soon-as-september>
- [93] RIKEN Center for Computational Science, feasibility study of memory-bandwidth balance for 2.8T-parameter sparse-MoE... <https://www.tomshardware.com/tech-industry/policy/new-us-export-control-s-reportedly-target-chinese-access-to-remote-ai-servers-trump-admins-cut-down-ai-diffusion-rule-could-be-shared-with-industry-as-soon-as-september>
- [94] SanDisk and SK hynix (3 Aug. 2026). <https://www.sandisk.com/company/newsroom/press-releases/2026/2026-08-03-sandisk-and-sk-hynix-advance-global-standardization-of-hbf>
- [95] HBF specification as reported at FMS 2026: up to 512 GB per stack in 8-Hi and 16-Hi NAND. <https://www.tomshardware.com/pc-components/ssds/sandisk-and-sk-hynix-unveil-hbf-spec-up-to-16-hi-nand-stacks-3-tb-s-bandwidth-ucie>
- [96] NVIDIA DGX Spark: GB10 Grace Blackwell, 128 GB LPDDR5X at 273 GB/s, “up to 1 PFLOP FP4” (sparse). <https://www.tomshardware.com/desktops/mini-pcs/nvidia-dgx-spark-gets-18-percent-price-increase-as-memory-shortages-bite-founders-edition-now-usd4-699-up-from-usd3-999>
- [97] Framework Desktop with AMD Ryzen AI Max+ 395 and 128 GB LPDDR5X, \$1,999 at announcement (Feb. 2025). <https://www.techspot.com/news/106934-framework-first-ever-desktop-pc-packs-amd-ryzen.html>
- [98] SK hynix chief executive, 11 Jul. 2026, on the day of the company’s Nasdaq listing. <https://www.tomshardware.com/pc-components/dram/sk-hynix-says-2027-will-be-the-worst-year-for-memory-shortage-forecasts-crunch-to-last-until-2030-ceo-shares-grim-outlook-on-the-day-sk-hynix-gets-listed-on-nasdaq>
- [99] TechInsights, “SMIC N+3 confirmed. <https://www.techinsights.com/blog/smic-n3-confirmed-kirin-9030-analysis-reveals-how-close-smic-5nm>
- [100] SemiAnalysis, “Huawei Ascend production ramp”... <https://newsletter.semanalysis.com/p/huawei-ascend-production-ramp>
- [101] wccftch, “CXMT races to increase capacity to 600,000 wafers per month”... <https://wccftch.com/cxmts-64gb-server-ddr5-memory-chip-is-now-prier-than-samsungs-as-it-races-to-increase-capacity-to-600000-wafers-per-month/>
- [102] Tom’s Hardware (Citirini Research), “CXMT close to matching Micron’s memory capacity in 2026”... <https://www.tomshardware.com/pc-components/dram/cxmt-close-to-matching-microns-memory-capacity-in-2026-research-claims-would-put-china-on-track-to-become-worlds-second-largest-dram-producer>
- [103] Reuters (Ngui/Shen), “Chipmaker CXMT vaults to top of China’s valuation with 466% surge”... <https://www.cnn.com/2026/07/27/cxmt-china-market-debut-chipmaker-ipo.html>
- [104] As of mid-2026 no public teardown confirms Chinese-fabbed DRAM in a shipping accelerator. <https://newsletter.semanalysis.com/p/huawei-ascend-production-ramp>
- [105] Hot Chips 2026 deck 01, Jim Handy (Objective Analysis), “Memory.”
- [106] Hot Chips 2026 decks 02 (Micron), 03 (Samsung) and 04 (SK hynix).

- [107] Hot Chips 2026 deck 03, Sangwook Han (Samsung), “HBM Base Die.
- [108] TrendForce, “Three Chinese chip toolmakers enter global top 20 as Naura rises to fifth by sales”. . . <https://www.trendforce.com/news/2026/02/02/news-three-chinese-chip-toolmakers-reportedly-enter-global-top-20-as-naura-rises-to-fifth-by-sales/>
- [109] BigGo Finance, “Naura 2025 results: revenue ~RMB 39.4B (+31%)” (Apr. 2026). <https://finance.biggo.com/news/9r28m50BvthpMgHBYd-l>
- [110] Asia Society Policy Institute, “China’s Big Fund 3.0. <https://asiasociety.org/policy-institute/chinas-big-fund-30-xis-b-oldest-gamble-yet-chip-supremacy>
- [111] SCMP (SEMI), “China to remain world’s biggest buyer of chipmaking equipment through 2027”. <https://www.scmp.com/tech/article/3336933/china-remain-worlds-biggest-buyer-chipmaking-equipment-through-2027-semi>
- [112] China National Energy Administration, 2025 installed-capacity and renewable-integration. . . <https://www.nea.gov.cn/20260129/6874f211acd0417eab7ac10c3061a7c2/c.html>
- [113] US EIA: 2025 utility-scale additions 53 GW (largest since 2002); 2026 planned record 86 GW. <https://www.eia.gov/to-dayineergy/detail.php?id=67205>
- [114] Ember, Global Electricity Review 2026: China 10,573 TWh (+503 TWh y/y) vs US 4,430 TWh. <https://ember-energy.org/latest-insights/global-electricity-review-2026/major-countries-and-regions/>
- [115] Lawrence Berkeley National Laboratory, *Queued Up: 2025 Edition* (data through 2024). <https://emp.lbl.gov/publications/queued-2025-edition-characteristics>
- [116] Semafor/techchannel (Nov. 2025): provinces cut DC electricity up to 50% (to ~RMB 0.4/kWh) exclusively for approved-domestic-chip facilities. <https://www.semafor.com/article/11/05/2025/china-slashes-big-data-centers-electric-bills>
- [117] BigGo / Tech Digest / 36Kr supply-chain reporting [A]. <https://www.techdigest.tv/2026/02/excluding-chinese-parts-could-triple-tesla-optimus-costs.html>
- [118] Fortune (6 Jun. 2026): 2025 global humanoid shipments >13,000, ~85% Chinese. <https://fortune.com/2026/06/06/chinese-humanoid-robots-global-market-sales-performative-functional/>
- [119] Gasgoo / Fortune (Jun. 2026): Unitree STAR-market IPO approved. <https://fortune.com/2026/06/06/chinese-humanoid-robots-global-market-sales-performative-functional/>
- [120] NVIDIA Nemotron 3 family (15 Dec. 2025). <https://nvidianews.nvidia.com/news/nvidia-debuts-nemotron-3-family-of-open-models>
- [121] NVIDIA open-sourced the CUDA Tile IR and cuTile compiler stack under Apache 2.0 with LLVM exceptions. . . <https://github.com/NVIDIA/cuda-tile>
- [122] Huang, same call: “The world will need both closed models and open models. <https://www.investing.com/news/transcripts/nvidia-at-the-six-five-summit-open-models-reshape-ai-strategy-93CH-4880247>
- [123] NVIDIA and Poolside, reported by the *Wall Street Journal* from deal documents and announced in Poolside’s shareholder letter (20–24 Aug. 2026). <https://www.forbes.com/sites/jonmarkman/2026/08/24/nvidia-pays-poolside-6b-to-license-its-model-factory-and-109-workers/>
- [124] NVIDIA and Hugging Face—*reported, not concluded* [A]. <https://techcrunch.com/2026/08/26/nvidia-loses-in-on-hugging-face-acquisition/>
- [125] The pattern of licence-plus-acquihire transactions that leave the target nominally independent and, on the reported. . . <https://www.warren.senate.gov/newsroom/press-releases/warren-blumenthal-question-whether-nvidias-20-billion-groq-deal-is-attempt-to-avoid-antitrust-laws/>
- [126] NVIDIA Corporation, Form 8-K filed 17 August 2026 (Items 1.01, 2.03, 7.01). <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000069/nvd-a-20260817.htm>
- [127] NVIDIA, “NVIDIA partners with Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR to establish AI. . . <https://nvidianews.nvidia.com/news/nvidia-partners-with-apollo-blackrock-blackstone-brookfield-goldman-sachs-and-krk-to-establish-ai-compute-infrastructure-financing-platforms-to-mobilize-over-500-billion-of-third-party-capital>
- [128] US Securities and Exchange Commission, Office of Structured Finance interpretive position (29 Jul. 2026). <https://www.alston.com/en/insights/publications/2026/08/data-center-securitizations-sec-regulations>
- [129] Reuters and *The Wall Street Journal* (27 Aug. 2026), one day after earnings. <https://whbl.com/2026/08/27/nvidia-pauses-revenue-sharing-deals-with-ai-closed-companies-wsj-reports/>
- [130] Dave Vellante, “Nvidia’s Jensen Huang just made the AI buildout too big to fail,” *SiliconANGLE* (15 Aug. 2026, updated 21 Aug.). <https://siliconangle.com/2026/08/15/nvidias-jensen-huang-just-make-ai-buildout-big-fail/>
- [131] NVIDIA, second quarter fiscal 2027 results (quarter ended 26 Jul. 2026; announced 26 Aug. 2026). <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000073/q2fy27pr.htm>
- [132] AMD’s openness posture: ROCm 7 positioned explicitly against CUDA. <https://rocm.blogs.amd.com/artificial-intelligence/introducing-instella-3B/README.html>
- [133] AMD’s warrant instruments. <https://www.sec.gov/Archives/edgar/data/2488/000002488260000045/amd-20260223.htm>
- [134] AMD, second quarter 2026 results (4 Aug. 2026). <https://ir.amd.com/news-events/press-releases/detail/1295/amd-reports-second-quarter-2026-financial-results>
- [135] Hot Chips 2026, Stanford University, 23–25 August 2026. <https://hotchips.org/>
- [136] Hot Chips 2026 deck 39, Richard Ho, Ravi Narayanaswami and Chris Leary (OpenAI), “You Can Just Build Things . . . Chips”. <https://openai.com/index/openai-broadcom-jalapeno-inference-chip/>
- [137] OpenAI and Broadcom, “Strategic collaboration to deploy 10 gigawatts of OpenAI-designed accelerators”. . . <https://investors.broadcom.com/news-releases/news-release-details/openai-and-broadcom-announce-strategic-collaboration-deploy-10>
- [138] OpenAI, “Jalapeño: first results” (25 Aug. 2026). <https://openai.com/index/jalapeno-first-results/>
- [139] Jensen Huang, Q2 FY2027 press release (26 Aug. 2026). <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027>
- [140] Hot Chips 2026 deck 38, Norman P. Jouppi and Sridhar Lakshmanamurthy (Google), “Eighth Generation TPU Family. <https://www.servethehome.com/google-tpu-8i-for-inference-and-tpu-8t-for-training-announced/>
- [141] Sundar Pichai, Alphabet Q1 2026 earnings call (29 Apr. 2026). <https://www.datacenterdynamics.com/en/news/google-to-sell-tpus-to-a-select-group-of-customers-for-their-data-centers/>
- [142] Microsoft, “Maia 200: the AI accelerator built for inference” (26 Jan. 2026). <https://blogs.microsoft.com/blog/2026/01/26/maia-200-the-ai-accelerator-built-for-inference/>
- [143] Meta, MTIA roadmap (announced Mar. 2026, Broadcom co-designed; Hot Chips 2026 deck 33). <https://www.tomshardware.com/tech-industry/semiconductors/meta-reveals-four-new-mtia-chips-built-for-ai-inference>
- [144] Amazon Web Services, Trainium3 general availability (2 Dec. 2025). <https://www.hpcwire.com/2025/12/02/aws-brings-the-trainium3-chip-to-market-with-new-ec2-ultraservers/>
- [145] Broadcom, second quarter fiscal 2026 results (3 Jun. 2026). <https://investors.broadcom.com/news-releases/news-release-details/broadcom-inc-announces-second-quarter-fiscal-year-2026-financial>
- [146] Hot Chips 2026 decks 23 (AMD, Steve Scott, “System Architecture of the MI400 Series”), 36 (Microsoft, Maia 200), 37. . .
- [147] AWS and NVIDIA (26 Aug. 2026). <https://press.aboutamazon.com/aws/2026/8/aws-and-nvidia-to-deliver-2-million-additional-gpus-and-next-generation-infrastructure-for-agentic-and-physical-ai>
- [148] OpenAI executive departures, and the direction of the talent flow. <https://techcrunch.com/2026/08/25/openai-loses-a-top-data-center-exec-as-stream-of-high-profile-departures-continues/>
- [149] NVIDIA’s Q2 FY2027 CFO Commentary (26 Aug. 2026) discloses a change in reporting structure that most coverage missed. https://s201.q4cdn.com/141608511/files/doc_financials/2027/Q227/Q2FY27-CFO-Commentary.pdf
- [150] SpaceX “Terafab” (confirmed by SpaceX and reported by Reuters, 6 Aug. 2026). <https://techcrunch.com/2026/08/06/tesla-and-spacex-will-invest-16-8b-to-start-building-terafab-chip-factory-in-texas/>
- [151] JFrog (2024: ~100 malicious HF models); ReversingLabs “nullfAI” (Feb. 2025: broken-pickle evasion). <https://www.reversinglabs.com/blog/rl-identifies-malware-ml-model-hosted-on-hugging-face>
- [152] Palo Alto Unit 42, “Model namespace reuse” (Sep. 2025). <https://unit42.paloalto.networks.com/model-namespace-reuse/>
- [153] Anthropic / UK AISI / Turing Institute (Oct. 2025). <https://www.anthropic.com/research/small-samples-poison>
- [154] Hubinger et al., “Sleeper Agents” (arXiv:2401.05566). <https://arxiv.org/abs/2401.05566>
- [155] Qiu, Zhou, Ferrara, “Information Suppression in Large Language Models. <https://arxiv.org/abs/2506.12349>
- [156] Naseh et al., “R1dacted: Investigating Local Censorship in DeepSeek’s R1” (arXiv:2505.12625). <https://arxiv.org/abs/2505.12625>
- [157] Anthropic, “Detecting and preventing distillation attacks” (Feb. 23, 2026). <https://www.anthropic.com/news/detecting-and-preventing-distillation-attacks>
- [158] NSA, FBI and CISA, joint cybersecurity advisory AA26-251a (9 Sep. 2026), naming DeepSeek, Moonshot, Alibaba, MiniMax. . . <https://www.cisa.gov/news-events/cybersecurity-advisories/aa26-251a>
- [159] TechTimes (4 Aug. 2026): MiniMax H3 tops the open-weight video arena. <https://www.techtimes.com/articles/322904/20260804/minimax-h3-open-weights-exclude-us-eu-uk-korea-local-deployment.htm>
- [160] S. Matsuoka, “Latent Model Backdoors and Prompt-Borne Worms.
- [161] W. Zhang and S. Pei, “Your LLM Agent Can Leak Your Data. <https://aclanthology.org/2026.findings-acl.1257.pdf>
- [162] OpenAI, “OpenAI and Hugging Face Partner to Address Security Incident During Model Evaluation”. . . <https://openai.com/index/hugging-face-model-evaluation-security-incident/>
- [163] Hugging Face Security Team, “Anatomy of a Frontier Lab Agent Intrusion. <https://huggingface.co/blog/agent-intrusion-technical-timeline>
- [164] Anthropic, “Investigating Three Real-World Incidents in Our Cybersecurity Evaluations”. . . <https://www.anthropic.com/news/investigating-incidents-cybersecurity-evals>

- [165] Fortune (4 Aug. 2026) / Axios (3 Aug. 2026). <https://fortune.com/2026/08/04/baffling-white-house-wont-publicly-release-ai-model-evaluation-framework-it-reviewed-today-with-openai-anthropic-microsoft-and-others/>
- [166] World Artificial Intelligence Cooperation Organization, founding agreement signed in Shanghai on 16 July 2026. . . https://english.www.gov.cn/news/2026/07/17/content_WS6a59a226c6d00ca5f9a0c432.html
- [167] Forbes/CNBC/Bloomberg, “OpenAI valuation reaches \$852 billion after \$122B round” (Mar. 31, 2026). <https://www.forbes.com/sites/antoniopequenoiv/2026/03/31/openai-valuation-reaches-852-billion-after-massive-fundiround/>
- [168] CNBC, “OpenAI resets spend expectations, targets around \$600 billion by 2030” (20 Feb. 2026). <https://www.cnbc.com/2026/02/20/openai-resets-spend-expectations-targets-around-600-billion-by-2030.html>
- [169] Anthropic, “Series H” (May 28, 2026: \$65B at \$965B post-money); CNBC coverage. <https://www.anthropic.com/news/series-h>
- [170] PitchBook via Yahoo Finance, “Anthropic’s gross margin is the most important number in AI” (10 Jun. 2026). <https://finance.yahoo.com/markets/stocks/articles/anthropics-gross-margin-most-important-070500338.html>
- [171] CNBC/Axios, “Nvidia becomes first \$5 trillion company” (Oct. 29, 2025). <https://www.axios.com/2025/10/29/nvidia-5-trillion>
- [172] Introl, “Hyperscaler capex \$600B+ in 2026”; Forbes (Jun. 2026). <https://introl.com/blog/hyperscaler-capex-600b-2026-ai-infrastructure-debt-january-2026>
- [173] Axios/CNBC. <https://www.axios.com/2026/07/27/debt-data-center-oracle-goole>
- [174] BEA third-estimate analyses: AI/data-center capex ~74% of US Q1 2026 GDP growth. <https://www.ftc.io/ai-capex-gdp-growth-q1-2026-bea-third-estimate/>
- [175] Data Gravity, “China’s open-weight takeover” . . . <https://www.datagravity.de/v/p/chinas-open-weight-takeover>
- [176] NVIDIA, CFO Commentary filed as Exhibit 99.2 to the Q2 FY2027 results (26 Aug. 2026), and Form 10-Q for the period ended 26 July 2026. <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000073/q2fy27cfocommentary.htm>
- [177] Morgan Stanley initiation of credit coverage on NVIDIA (24 Aug. 2026), rating Neutral.
- [178] Forbes / G. Wang (31 Jul. 2026), “Why OpenAI’s 80% price cut could trigger a race to the bottom in AI”. <https://www.forbes.com/sites/geruiwang/2026/07/31/why-openais-80-price-cut-could-trigger-a-race-to-the-bottom-in-ai/>
- [179] G. Appenzeller (a16z), “LLMflation”. <https://a16z.com/llmflation-llm-inference-cost/>
- [180] MIT Project NADA, “The GenAI Divide” (Aug. 2025). <https://www.forbes.com/sites/jasonsnnyder/2025/08/26/mit-finds-95-of-genai-pilot-s-fail-because-companies-avoid-friction/>
- [181] Bain & Company, 6th Global Technology Report (Sep. 2025). <https://www.bain.com/about/media-center/press-releases/2025/2-trillion-in-new-revenue-needed-to-fund-ai-scaling-trend--bain-companys-6th-annual-global-technology-report/>
- [182] Intellectia/CNN. <https://intellectia.ai/blog/semiconductor-stocks-selloff-june-2026>
- [183] CNBC/NBC (Jul. 28–29, 2026): memory-led chip rout, SK hynix –15%, >\$1T lost, Nasdaq-100 in correction. <https://www.cnbc.com/2026/07/29/chip-selloff-sk-hynix-samsung-softbank.html>
- [184] Epoch AI, “LLM inference prices have fallen rapidly but unequally across tasks” . . . <https://epoch.ai/data-insights/llm-inference-price-trends>
- [185] Epoch AI, “Global AI computing capacity is doubling every 7 months” (Jan. 2026). <https://epoch.ai/data-insights/ai-chip-production>
- [186] S. Matsuoka. . .
- [187] CME Group and Silicon Data (11 Aug. 2026). <https://www.cmegroup.com/media-room/press-releases/2026/8/11/cmegroup-and-silicondatalaunchcompute-futuresonoctober5tounloc.html>
- [188] Baiguan (9 Jul. 2026), on neutral-router economics. <https://www.baiguan.news/p/china-ai-token-export-openrouter-market-share-chinese-models-performance-gap-cost-efficiency-semiconductor-energy-self-sufficiency-investment-opportunities>
- [189] Executive Order 14320, “Promoting the Export of the American AI Technology Stack” (23 Jul. 2025), and the Commerce. . . <https://www.trade.gov/press-release/commerce-statement-application-conclusion-next-steps-american-ai-exports-program>
- [190] Brazil, 21 August 2026. <https://www.aljazeera.com/economy/2026/8/21/brazil-launches-ai-supercomputer-push-while-balancing-us-and-chinese-tech>
- [191] Bloomberg (26 Aug. 2026, anonymous sources), relayed by The Next Web, Al Jazeera (4 Sep.) and AI-Monitor. <https://www.aljazeera.com/news/2026/9/4/as-the-china-us-ai-race-reaches-cairo-egypt-faces-a-strategic-decision>
- [192] Bloomberg via *The Edge* (7 Sep. 2026) and SCMP (8 Sep.). <https://www.scmp.com/week-asia/politics/article/3366792/will-malaysia-potential-huawei-ai-deal-risk-clash-over-us-trade-pact>
- [193] Kazakhstan–China package of 16 July 2026 (Tokayev in Shanghai, “\$15B, seventy-plus agreements”). <https://www.intellinews.com/kazakhstan-china-sign-15bn-worth-of-deals-as-tokayev-visits-shanghai-455262/>
- [194] Hugging Face, “The state of open models, summer 2026”. <https://huggingface.co/blog/state-of-open-models-summer-2026>
- [195] White & Case / FPF: Japan’s AI Promotion Act (passed 28 May 2025)—first AI-dedicated law. <https://www.whitecase.com/insight-alert/japans-first-ai-legislation-becomes-law-focus-promoting-research-and-development-no>
- [196] Gibson Dunn / European Parliament. <https://www.gibsondunn.com/eu-ai-act-omnibus-agreement-postponed-high-risk-deadlines-and-other-key-changes/>
- [197] European Commission, InvestAI / AI Gigafactories (€20B; 77 proposals across 16 member states). https://commission.europa.eu/topics/competitiveness/competitiveness-coordination-tool-projects/ai-gigafactories_en
- [198] EuroHPC JU (30 Jul. 2026): formal call for up to seven AI gigafactories (each 100k+ advanced chips), expected to trigger >€30B. https://www.eurohpc-ju.europa.eu/eurohpc-joint-undertaking-launches-ai-gigafactories-call-2026-07-30_en
- [199] IndiaAI Mission GPU tenders: ~38,000 GPUs empanelled across three rounds at subsidized rates reportedly under \$1/GPU-hr. <https://www.communicationstoday.co.in/38000-gpus-one-big-problem-indias-ai-mission-cant-find-enough-users/>
- [200] Counterpoint / TrendForce: SK hynix ~62% HBM share. <https://www.astutegroup.com/news/general/sk-hynix-holds-62-of-hbm-micron-overtakes-samsung-2026-battle-pivots-to-hbm4/>
- [201] NVIDIA/HUMAIN partnership (500 MW-class “AI factories,” several hundred thousand GPUs). <https://www.commerce.gov/news/press-releases/2025/11/statement-uae-and-saudi-chip-exports>
- [202] Bloomberg, “Lawmakers urge export controls on RISC-V” (Nov. 2023). https://www.warner.senate.gov/public/_cache/files/a/b/ab558775-b187-434e-b70d-040a63aea6f7/1AD7AF181B29E0603B1023A6FABE01893C9F53BFCA2ACAAS2B1535CEBE2FC7AC.9.30.25-risc-v-letter-to-bis.pdf
- [203] Caixin / Fortune / State Council reporting: anti-“involution” campaign. <https://fortune.com/asia/2025/08/29/china-warns-against-disorderly-competition-ai-race>
- [204] BNN Bloomberg, Alibaba FY results (May 2026): three-year AI capex plan of RMB 380B (~\$56B) to be exceeded without a new stated target. <https://www.bnnbloomberg.ca/business/2026/05/13/alibaba-reports-continued-profit-squeeze-from-spending-on-ai-and-instant-retail/>
- [205] MIT Technology Review (Mar. 2025) / Light Reading / ASPI. <https://www.technologyreview.com/2025/03/26/1113802/china-ai-data-centers-unused/>
- [206] ChinaTalk / CSIS / trade reporting: HBM and advanced packaging as binding constraints on the Ascend fleet. <https://www.chinatalk.media/p/mapping-chinas-hbm-advancement>
- [207] ECMWF, “ECMWF’s AI forecasts become operational” (25 Feb. 2025) and “Ensemble AI forecasts become operational” . . . <https://www.ecmwf.int/en/about/media-centre/news/2025/ecmwf-ai-forecasts-become-operational>
- [208] Google DeepMind, “AlphaFold: five years of impact” (2025). <https://deepmind.google/blog/alphafold-five-years-of-impact/>
- [209] Insilico Medicine, rentosertib Phase IIa, *Nature Medicine* (3 Jun. 2025). <https://insilico.com/news/ntrecuxscl-insilico-announces-nature-medicine-publi>
- [210] MIT Department of Economics, statement on “Artificial Intelligence, Scientific Discovery, and Product Innovation” (16 May 2025). <https://economics.mit.edu/news/assuring-g-accurate-research-record>
- [211] ECMWF AIFS blog, “Farewell to external AI models” (May 2026). <https://www.ecmwf.int/en/about/media-centre/aifs-blog/2026/farewell-external-ai-models>
- [212] NSF NCSES, *Science and Engineering Indicators 2025*. <https://ncses.nsf.gov/pubs/nsb20257/global-rd-and-international-comparisons-2>
- [213] B. F. Jones and L. H. Summers, “A Calculation of the Social Returns to Innovation,” NBER w27863. https://www.nber.org/system/files/working_papers/w27863/w27863.pdf
- [214] N. Bloom, C. I. Jones, J. Van Reenen and M. Webb, “Are Ideas Getting Harder to Find?” *AER* 110(4). . . <https://web.stanford.edu/~chadj/IdeaPF.pdf>
- [215] D. Acemoglu, “The Simple Macroeconomics of AI,” NBER w32487 (2024). <https://economics.mit.edu/sites/default/files/2024-05/The%20Simple%20Macroeconomics%20of%20AI.pdf>
- [216] Nature Index 2026 Research Leaders (10 Jun. 2026). <https://www.natureasia.com/en/info/press-releases/detail/9350>
- [217] S. Matsuoka, “How Much Compute Does a Nation’s Science Need? A Portable, Measurement-Anchored Framework for Sizing. . .
- [218] S. Matsuoka and T. Schulthess, “AI for Science and an International Consortium,” edited discussion record v1.0, 31 July 2026 (RIKEN R-CCS / CSCS).
- [219] *Opinions of the State Council on Deepening the Implementation of the “AI+” Initiative* (Guofa [2025] No. 11). . . https://www.gov.cn/zhengce/content/202508/content_7037861.htm
- [220] Ministry of Science and Technology and the National Natural Science Foundation of China, launch of the “AI for. . . https://www.gov.cn/xinwen/2023-03/27/content_5748495.htm

- [221] R. Arcesati, MERICS (4 March 2026). <https://merics.org/en/comment/science-funding-call-provides-insight-chinas-ai-priorities>
- [222] The Party's *Recommendations* for the 15th Five-Year Plan... https://www.gov.cn/zhengce/202511/content_7049042.htm
- [223] M. Bianchini, V. Di Girolamo, J. Ravet and D. Arranz... https://research-and-innovation.ec.europa.eu/knowledge-publications-tools-and-data/publications/all-publications/artificial-intelligence-science-promises-or-perils-creativity_en
- [224] Shanghai AI Laboratory, *Intern-S1* (released 26 July 2025; [arXiv:2508.15763](https://arxiv.org/abs/2508.15763)). <https://huggingface.co/internlm/Intern-S1>
- [225] ByteDance AI4Science, *Protenix* (Apache 2.0; v1.0.0 released 5 February 2026 at 368M parameters, v2 on 8 April 2026 at 464M). <https://github.com/bytedance/Protenix>
- [226] The DeepModeling community and DP Technology (Shenshi Keji, founded 2018; chief scientist Zhang Linfeng). <https://www.nature.com/articles/s41524-024-01493-2>
- [227] Chinese Academy of Sciences, *ScienceOne*, unveiled at the World AI Conference, 26–28 July 2025, built by... https://english.cas.cn/newsroom/cas_media/202507/t20250728_1048587.shtml
- [228] Pengcheng Cloudbrain (Pengcheng Yunnao), the first fully domestic large-scale intelligent computing system, built by... https://www.stdaily.com/web/gdxw/2026-07/08/content_544307.html
- [229] Evidence that the Chinese AI-for-science programme is aimed at industry rather than at publication. <https://www.shlab.org.cn/news/5444277>
- [230] M. Berman, “Open-source is WINNING,” Matthew Berman (YouTube), Aug. 2026. <https://youtu.be/CVlKp9Ld-Zg>