

How China Will Win the AI War

The Convergence Strategy: A Structural Path to AI Dominance

A Living Book¹

Satoshi Matsuoka

v0.23 — 9 September 2026

¹This is a living document, revised as the underlying industrial and technical facts evolve. Version v0.23 — 9 September 2026. Corrections and additions are welcome.

Preface

The title of this book is deliberately provocative, but it is not a prediction of inevitability—it is a structural analysis. “Winning” here does not mean producing the single best frontier model in any given quarter. It means what winning meant in solar photovoltaics, in lithium-ion batteries, and in electric vehicles: capturing the overwhelming majority of global production, setting the cost curve, controlling the supply chain, defining the de facto standards, and reducing erstwhile leaders to protected niches sustained by tariffs and subsidies. By that definition China has already won three consecutive technology wars in which the West, at the outset, held every advantage—the science, the patents, the firms, and the markets.

The central claim of this book is that artificial intelligence is now following the same trajectory, and that the mechanism is visible in three separable elements: *the model* (open-weight frontier models released at marginal prices), *the compute* (an open, standardized hardware ecosystem built on RISC-V, domestic accelerators, and a memory-capacity arbitrage that substitutes abundant commodity DRAM for scarce HBM), and *the semiconductor* (the fabs, logic, and DRAM base that removes the last physical chokepoints). Each element is driven by the same national strategy of *openness as a weapon*: open weights, open instruction sets, open interconnect specifications, and open software stacks, deployed to commoditize precisely the layers where Western incumbents collect their rents.

Chapter 2 supplies the theoretical machinery—the production economics of intelligence, AI as a general-purpose technology, and the industrial-organization account of what openness does to rents—for readers who want the argument’s engine before its evidence. Part I then surveys the three precedent industries—how the war was actually fought and won in each—adds, in lesser detail, the telecommunications campaign as the precedent closest in structure to an exported AI stack, and then, in a chapter of equal weight, the campaigns where the same playbook has *not* worked. Part II opens by stating the strategy in whole—each element placed on the playbook of Part I and the campaign scored against it—and then analyses the elements of AI separately: the model and the harness that has become its contested layer, the compute, the semiconductor, the watt, the data, and—closing the loop back into Part I’s terrain—embodied AI and physical production; it closes with the finding that the largest Western firms now run the same playbook against one another. Part III turns to consequences—trust and agentic security, governance, the financial exposure of the incumbents, how the non-aligned majority of the world chooses, and the strategies available to every actor class. Part IV takes up the largest prize on the table and the one with the weakest constituency: why AI for science dominates the value case, how China has organized for it, how much compute a nation’s science actually needs, and the international institution that the analysis keeps arriving at. Part V assembles the convergence argument and then states the whole book in short form, with its odds and its falsifiers attached. Throughout, I have tried to keep the analysis anchored in numbers with sources, and to flag clearly where evidence is contested or preliminary. As a living book, chapters will be revised and extended as events develop; readers should treat quantitative claims as dated to the version stamp above.

The book carries an apparatus that a reader should be able to hold it to. Selected headline claims are tagged inline with an evidence class (Section 1.5); the complete graded ledger—which is the authoritative record, and the fastest route to this book’s weakest link—appears as Appendix D, with a machine-readable claim register and a quarterly falsification dashboard in Appendix J.

Author-modeled results are disclosed as such wherever they are load-bearing, and the four core propositions are stated, coupled, and falsified separately in Section 1.4. The book also carries an explicit *currency date*, and that date is defined to be the date of the edition itself—9 September 2026, the date on the title page. Every factual claim in the stable chapters is current as of that date; anything the author learned after it belongs to the next edition, not to this one. Defining the cutoff as the edition date is the only convention that cannot contradict itself: a book whose declared freeze precedes its own completion will sooner or later discuss an event it has ruled out of scope. Appendix H carries the most recent development cycle in detail, because fast-moving material benefits from being quarantined in one place where it can be replaced wholesale; but it is inside the currency date, not outside it.

A condensed edition of this book—about fifty pages, with Part I summarized in a single chapter and the weight shifted toward the overall strategy and its international consequences—is published alongside it, from the same source tree, under the same version stamp. Nothing in it is new: every number, grade, and forecast is drawn from this full edition, which remains the authoritative record and the place where the sources are.

A note on the author’s position, and on motive

Scholarly convention asks an author writing on a contested subject to state where he stands in relation to it. Here the convention matters more than usual, because the position is unusually entangled with the subject—and because, on the face of it, the position argues against the conclusion.

I direct the RIKEN Center for Computational Science, the institution that built and operates Fugaku, the machine that swept all four major supercomputing benchmark lists in 2020 and 2021; and I direct the FugakuNEXT program that will succeed it toward the end of this decade. I hold a professorship at the Institute of Science Tokyo. For some three decades my professional life has been spent inside high-performance computing—designing machines and the processors in them, running procurements, arguing about architecture in public and in committee, and building the software and scientific applications that have to justify the expenditure afterwards.

That career has been conducted, almost entirely, in partnership with Western technology and Western institutions. FugakuNEXT is being built by Fujitsu around the MONAKA-X processor tightly coupled to NVIDIA accelerators over NVLink Fusion—a design decision that places Japan’s flagship national program on the American accelerator platform for the coming decade, and one I defended and continue to defend on its engineering merits [356]. My center’s relationships with NVIDIA, AMD, Intel, Arm, and the American and European system vendors are longstanding, active, and cordial; a substantial part of my working life consists of technical argument with their architects, which is among the more enjoyable things about the job. In January 2026 RIKEN, Fujitsu, and NVIDIA entered a joint partnership with Argonne National Laboratory on AI for science and next-generation computing [392]. In June 2026 Japan became the *first international partner* in the United States’ Genesis Mission—a one-billion-dollar, five-year commitment, half from each nation, binding twelve DOE national laboratories to twelve Japanese research institutions across quantum information, fusion, biotechnology, advanced materials, particle physics, and autonomous laboratories, with RIKEN among the central Japanese participants [390]. Within Japan I have been one of the principal advocates for AI for Science as a national priority, through RIKEN’s TRIP and TRIP-AGIS platforms and through the compute-sizing framework that Chapter 26 of this book is built on.

The reader should weigh all of that as a declaration of interest, and should notice which way it points. My institution buys Western technology and will go on buying it. My flagship program runs on an American accelerator architecture. My laboratory’s deepest international partnerships are with American and European institutions, and I value them. Several of this book’s load-bearing technical arguments rest on work I coauthored or wrote—the accelerated-CPU

analysis with Jack Dongarra and Torsten Hoefler [184, 185], the national AI-for-science sizing framework [305], the agentic-backdoor survey behind Chapter 19 [291], and the international-consortium discussion record developed with Thomas Schulthess [306]—and the economic-impact study discussed in Chapter 24 was commissioned by my own institution, a fact disclosed again at the point of use. *If these interests bias this book, they bias it against its thesis.* Nothing in my professional position gives me an incentive to argue that the Chinese stack is winning. A great deal of it gives me an incentive to argue the opposite, and it would have been considerably more comfortable to do so.

Why write it, then. Not for lack of commentary: there is a great deal of analysis of Chinese AI progress and of the competition between China and the United States and its allies. The problem is that almost all of it is *fragmentary*. The model community sees benchmarks and licences. The semiconductor community sees export controls and fab capacity. The financial community sees valuations and capital expenditure. The security community sees supply-chain risk. The policy community sees five-year plans and ministerial documents. Each of these communities sees its own layer with real expertise and the adjacent layers dimly, if at all—and the strategy this book describes is *invisible at any single layer*. It becomes visible only when the model, the compute, the semiconductor, the watt, the data, and physical production are laid side by side, because it is a systems strategy, and a systems strategy is precisely the thing that fragmentary analysis cannot resolve. An open instruction set is unremarkable on its own; so is a memory-capacity trade, so is an open-weight release, so is a provincial electricity tariff. Assembled, they are a machine for changing where cost and rent sit in an entire industry—which is exactly what the solar, battery, and electric-vehicle campaigns turned out to be, and exactly what was missed at the time, layer by layer, by people who each understood their own layer perfectly well.

I should also report an experience rather than only an argument. I did not begin this project expecting to arrive where it has arrived. Each time I have gone looking for the layer at which the Chinese position falls apart—the software stack, the memory wall, the packaging, the energy, the training economics, the scientific applications—I have more often found additional evidence of coherence than the decisive gap I expected, and the gaps I did find have tended to be narrower and to be closing faster than the public discussion assumes. That has been an unwelcome finding, professionally and personally, and I have tried throughout to hold it to a standard of proof appropriate to how much I would prefer it to be wrong. Where the evidence does not support the strong claim—and there are several such places, marked as such—the book says so.

Which produces the obligation. The vantage point here is not better than anyone else's; it is merely unusual. Very few people are positioned to look at processor architecture, semiconductor supply chains, national computing programs, model economics, and international scientific collaboration at the same time, with direct professional access to most of them, and with an obligation to no vendor and no government's messaging. Having had that view, I do not think it is defensible to leave what it shows in conference corridors, closed briefings, and program-committee conversations. Documenting it is the ordinary duty of a researcher who has done the analysis, and so is committing to maintain it: this is a living document, with graded evidence, stated probabilities, explicit falsification criteria, and a standing willingness to correct the manuscript against its own thesis wherever the record demands it. I intend to keep updating it as an academic and not as an advocate, which means reporting the evidence as it arrives regardless of which side it favours, and being publicly wrong when the record says so.

A note on perspective: this book is written not to celebrate the outcome it describes, nor to counsel despair, but because strategy made without a correct model of the adversary's strategy fails. The West's repeated pattern across solar, batteries, and EVs was not a failure of invention—it invented nearly everything—but a failure to take Chinese industrial strategy seriously until the cost curve had already been captured. The purpose of writing down the mechanism is to make that mistake harder to repeat.

Contents

Preface	i
A note on the author’s position, and on motive	ii
Notation, Abbreviations, and Glossary	xiii
1 The Thesis	1
1.1 Three wars, one playbook	1
1.2 Why AI is the fourth war	1
1.3 Openness as a weapon	2
1.4 Four propositions, stated separately	2
1.5 Evidence grades	3
2 A Theory of the Contest	5
2.1 The production economics of intelligence	5
2.1.1 The cost stack	5
2.1.2 Four learning curves, not one	6
2.1.3 Efficiency price cuts versus strategic price cuts	6
2.1.4 The demand side, and why it decides everything	6
2.1.5 From own-price elasticity to differentiated-system profit	7
2.2 AI as a general-purpose technology	8
2.2.1 Four layers of winning	9
2.2.2 Why this strengthens the AI-for-Science case	10
2.3 Industrial organization: openness relocates rent, it does not abolish it	10
2.3.1 The rent-relocation table	10
2.3.2 The concepts the book uses, named	11
2.3.3 The dual flywheel	11
2.4 Forecasting discipline	13
I The Precedents: Three Wars Already Won	14
3 Solar Photovoltaics: The Template	15
3.1 The West invents, the West leads (1954–2007)	15
3.2 The first scale-up and the first bankruptcies (2005–2013)	15
3.3 Trade defenses fail (2012–2018)	16
3.4 Domestic demand as strategic ballast (2011–2026)	16
3.5 The end state: total supply-chain capture	16
3.6 Winning is expensive: the involution phase	17
3.7 The Western rebuild that wasn’t	18
3.8 Lessons carried forward	18

4	Lithium-Ion Batteries: Vertical Depth	19
4.1	From Japan's invention to China's entry (1991–2011)	19
4.2	The white list: protection by regulation, not tariff (2015–2019)	19
4.3	The end state: shares and costs	19
4.4	The LFP counterattack: innovation where the West wasn't looking	21
4.5	From protection to gatekeeping	21
4.6	The Western casualty list	21
4.7	Lessons carried forward	21
5	Electric Vehicles: The System-Level Win	22
5.1	The leapfrog decision (2001–2012)	22
5.2	The demand machine (2013–2022)	22
5.3	BYD: the vertically integrated apex predator	23
5.4	The incumbents' collapse and the reverse technology flow	23
5.5	The EV as AI system: the bridge to Part II	24
5.6	Lessons carried forward	24
6	Telecommunications: The Network War, in Lesser Detail	25
6.1	The rise, measured	25
6.2	The entry pattern: encircle the cities from the countryside	25
6.3	The weapon: credit	26
6.4	The endgame: standards from the inside	26
6.5	The pushback, and what it removed	26
6.6	What transfers to AI, and what does not	27
7	The Playbook Distilled	28
7.1	The mechanism in general form	29
7.2	Two Western counter-patterns, both failing	29
8	Where the Analogy Breaks	30
8.1	The unfinished campaigns	30
8.1.1	Commercial aircraft and jet engines	30
8.1.2	Lithography, and the deep tool chain	30
8.1.3	Machine tools, reducers, and the precision economy	30
8.1.4	Operating systems and software ecosystems	31
8.1.5	Biotech: the surprise late bloomer	31
8.2	The success–failure model	31
8.3	Scoring AI against the model	32
II	The Elements of AI	33
9	The Strategy in Whole: The Elements on the Playbook	34
9.1	The strategy in one paragraph	34
9.2	Where each element sits	35
9.3	Scoring the campaign against the six steps	35
9.4	Where the analogy tells us to look	37
9.5	What the assembly does and does not license	38

10 The Model: Open Weights as the Commodity Layer	39
10.1 The DeepSeek shock and the price war (2024–2025)	39
10.2 An innovation genealogy	40
10.3 The ecosystem survey: eight frontiers, one commons	41
10.3.1 The open-weight ratchet: why the commons is now an equilibrium	42
10.3.2 The ratchet crosses the Pacific: Meta’s return to open weights	45
10.3.3 GLM-5.3, “Ox Alpha,” and the serving claim	46
10.3.4 The barbell: where the tokens go and where the money stays	49
10.3.5 AI-designed silicon: why the compute dependence is transitional	50
10.4 The toll road: the licence becomes a claim on the clouds	52
10.5 Three scoreboards, not one	53
10.6 Quantization: frontier models on commodity hardware	58
10.6.1 The dynamic-quantization breakthrough	58
10.6.2 The commodity-hardware frontier	59
10.6.3 The metric that matters: cost per solved task	60
10.6.4 The economics: LLMflation	60
10.7 The price inversion	60
10.8 Openness as declared national strategy	63
10.9 Compute constraints as forcing function	63
11 The Harness: The Layer the Contest Moved To	65
11.1 The layer acquires a name	65
11.2 The harness is a larger economic lever than the model	66
11.3 What China shipped in August	66
11.4 The standards layer, and who is not in the room	68
11.5 The strongest case against this chapter	69
11.6 The thesis, narrowed to what survives	69
12 The Compute: RISC-V, Accelerators, and the Memory Arbitrage	71
12.1 State-mandated architectural independence: RISC-V	71
12.1.1 The unresolved centre: four proposals for one function	72
12.1.2 The readiness number	73
12.2 LineShine and the LX2: the GPU-free machine takes No. 1	74
12.2.1 The proof boundary: what this machine establishes, and what it does not	75
12.2.2 The AI proof suite: what would settle it	77
12.3 The template: from LX2 to a RISC-V accelerated CPU, in two variants	78
12.4 The accelerator fleet and the unified ecosystem	79
12.4.1 Expelling NVIDIA, adopting the fleet	79
12.4.2 The unified interconnection ecosystem	80
12.4.3 Ethernet takes the scale-up domain, and the standards body loses the wire	81
12.4.4 The incumbent roadmap, and what it does and does not settle	82
12.5 What the argument is not: the GPU is not obsolete	83
12.5.1 Breadth against depth	84
12.6 Architecture theory: what to report, and why	86
12.7 The software and operations moat, made measurable	87
12.8 Democratizing AI through the LPDDR6 hardware arbitrage	88
12.8.1 The HBM scaling wall	88
12.8.2 CXMT and the LPDDR6 disruption	89
12.8.3 The trade-off: capacity over peak bandwidth	89
12.8.4 Strategic consequence: decentralized pretraining	90
12.9 The memory unlock: what happens if capacity becomes cheap	91
12.9.1 The constraint, priced	92

12.9.2	Two technologies, precisely described	92
12.9.3	The machine, and whether it works	93
12.9.4	Head to head: an appliance against a Rubin rack, and both against the API price list	94
12.9.5	Why this would matter more than it appears	100
12.9.6	The case against, which is strong	100
12.9.7	Where China would gain, stated narrowly	101
12.9.8	The dilemma this creates	101
12.10	The West defects to the open ISA: NVIDIA, Google, Meta	102
12.11	Assessment	104
13	The Semiconductor: Fabs, Logic, and DRAM	105
13.1	Logic: 5 nm without EUV	105
13.2	DRAM: CXMT and the swing-supplier position	106
13.2.1	The HBM ledger	106
13.2.2	The 2026 memory wall is priced, not clocked	107
13.3	HBM4 stopped being a specification	108
13.3.1	A 1.5× spread inside one nominal generation	108
13.3.2	The base die became a logic die	109
13.4	Packaging: the chokepoint hiding inside the memory ledger	110
13.5	A sanctions-elasticity model	110
13.6	Equipment, EDA, materials: the long ladder	111
13.7	The national-open logic at the silicon layer	112
14	The Watt: Energy as the Fourth Element	114
14.1	The supply asymmetry	114
14.2	Energy as industrial policy: the token factory's power bill	117
14.2.1	The value of flexible computing	117
14.3	The busbar casebook: ten regions, one table	118
14.3.1	Busbar-level deliverability: the unit the argument must be priced in . . .	118
14.4	Does the data centre have to be near you? The locality question, examined . . .	121
14.4.1	The latency budget, decomposed	121
14.4.2	The revealed preference of the industry	121
14.4.3	What the Chinese policy document actually says	122
14.4.4	The decisive counter-example, three days old	122
14.4.5	Where locality genuinely binds—stated at full strength	123
14.4.6	The claim that survives, and the one that should replace it	123
14.5	Where the watt meets the architecture	124
15	Data and Feedback: The Fifth Element	125
15.1	What open weights do not commoditize	125
15.2	The synthetic-data question	126
15.3	Seven questions that decide the element	126
16	Embodied AI: The Loop Closes on the Factory Floor	128
16.1	The embodied stack, separated	128
16.2	The physical-data flywheel, and who is building it	129
16.3	The industrial bill of materials	129
16.4	The deployment scoreboard, and what it does not show	131
16.5	Metrics that matter: the embodied proof suite	132
16.6	Comparative national positions	132
16.7	Falsifiers	133

17 The Mirror: The West Runs the Playbook Against Itself	134
17.1 The move, stated generally	134
17.2 The accelerator vendor commoditizes the model, and underwrites the buildout	135
17.2.1 The model layer	135
17.2.2 The datacentre layer	137
17.2.3 The counter-evidence, which is substantial	138
17.2.4 The quarter, as the filings state it rather than as it was reported	139
17.2.5 AMD runs the same playbook, which matters for the diagnosis	139
17.3 The laboratories and the hyperscalers commoditize the silicon	140
17.3.1 Jalapeño as a direct challenge, and the sentence that constrains how it can be read	141
17.3.2 And the counter-evidence here is decisive enough to be printed first in any honest summary	142
17.4 And one actor commoditizes the fab	142
17.5 The same method, and the asymmetry that survives it	143
17.5.1 What this does to Proposition P3	144
17.6 What the mirror licenses, and what it forbids	145
 III The Consequences	 146
18 Trust: The Contested Layer of the Open Commons	147
18.1 The supply chain is an attack surface	147
18.2 Behavioral integrity: the censorship record	148
18.3 Provenance and the distillation war	148
18.4 The sovereign trust stack	149
18.4.1 Five layers of trust	149
18.4.2 A graded certification ladder	149
 19 Monoculture and Agency: The Security Cost of the Commons	 151
19.1 Evidence grades for security claims	151
19.2 What has been demonstrated (Level A)	152
19.3 What has actually happened (Level C)	153
19.4 The compound scenario, and why it is conditional	154
19.5 Why this book's mechanism raises the stakes	155
19.6 From possibility to probability	156
19.6.1 Model monoculture is not the only monoculture	158
19.7 Consequences for the strategy	158
 20 The Governance Contest: Rules, Standards, and Who Writes Them	 160
20.1 Three governing philosophies	161
20.2 Standards are the load-bearing arena	162
20.3 The open-weight governance problem	163
20.4 Dual-use science: the governance failure mode that would hurt most	163
20.5 The vacant arena, and why it will be filled	164
20.6 Three governance futures	164
 21 The Wrong Side of the Trade	 166
21.1 The exposure map	166
21.2 The squeeze mechanism	168
21.2.1 The transmission argument no longer requires China	169
21.3 Four bubbles, not one	170

21.4	The Jevons rebuttal, taken seriously	171
21.4.1	Not all obligations are the same instrument	171
21.4.2	The elasticity that decides everything	172
21.5	Failure modes, entity by entity	172
21.5.1	OpenAI: the leveraged pure-play	172
21.5.2	Anthropic: better economics, same trade—and the strongest quarter anyone has printed	174
21.5.3	NVIDIA: three exposures compounding, one quarter at a time	175
21.5.4	The fifth exposure: guarantees, residuals, and a portfolio that marks with its customers	176
21.5.5	Google: the hedged incumbent	178
21.5.6	SoftBank and Arm: the double long, fully assembled	179
21.5.7	The data-center complex: capacity as liability, vintage by vintage	179
21.6	Scenarios and timelines	182
21.6.1	Scenario I: Orderly repricing (base case)	182
21.6.2	Scenario II: The cascade	182
21.6.3	Scenario III: Bifurcated stalemate	182
21.6.4	The author's probabilities and the skeptic's	183
21.6.5	Indicators, 2026–2030	183
21.7	The commodity turn: what happens when inference is traded	183
21.7.1	The instruments are real	183
21.7.2	Why the incumbent stands in the way	184
21.7.3	If it succeeds: the solar mechanism, applied to tokens	185
21.7.4	The case against, at full strength	186
21.7.5	What survives, and it is not what the headline suggested	187
21.8	Where value reappears after the model becomes cheap	188
21.9	What survives	189
22	The Non-Aligned: How the Rest of the World Chooses	190
22.1	The decision matrix	190
22.2	Two commons, not one	190
22.3	The export offer, as it actually exists	192
22.3.1	The wrapper	193
22.3.2	The three live cases, graded	193
22.3.3	The two offers, side by side	194
22.3.4	Where the telecom template transfers, and where it does not	196
22.4	Nine positions	197
22.5	What the matrix predicts	199
22.6	The non-aligned strategy, stated once	200
23	Strategies for the Rest	201
23.1	The United States: win the commoditization race or lose it	201
23.2	Europe: regulation is not production	202
23.3	Japan: the third-country playbook, first-person	202
23.4	The middle powers: what two commons are worth	203
23.5	The Global South: the buyers write the ending	203
23.6	China: advice for the presumptive winner	204
23.7	The common instruction	206

IV AI for Science: The Largest Prize	208
24 The Largest Prize: Why AI for Science Dominates the Value Case	209
24.1 What has actually been demonstrated	209
24.2 The failure ledger	210
24.3 Why the prize is nevertheless the largest one	211
24.3.1 The denominator	211
24.3.2 The return	211
24.3.3 The bottleneck	212
24.3.4 The flywheel	213
24.4 The value-conversion funnel	213
24.4.1 The arithmetic, done properly	214
24.4.2 The measured analogue: what national-scale HPC actually returned . . .	215
24.5 Measuring the scientific frontier: a production metric, not another benchmark . .	217
24.6 The architecture of the prize	218
24.7 Who is positioned	218
24.8 The strategic restatement	218
25 China's Scientific Machine: AI for Science as Industrial Policy	220
25.1 The policy record, read precisely	220
25.2 The output, measured by someone other than China	222
25.3 The institutional layer	223
25.4 The models, and the fact that matters most about them	224
25.5 The infrastructure, and the prize the state actually gave	226
25.6 The programme, scored on the playbook	227
25.7 What this does to the consortium proposal, and what it does not	229
26 How Much Compute Does a Nation's Science Need?	231
26.1 The measurement gap	231
26.2 A vendor-neutral unit, and a conversion rule	232
26.2.1 From scalar to workload-weighted vector	233
26.3 The two-layer demand model	233
26.4 Typed parameters and transfer rules	235
26.5 What three countries find	237
26.6 The universal-reach ladder, and the order-of-magnitude gap	238
26.7 What this means for the argument of this book	238
27 The Third Path: An International Consortium for Open Scientific AI	240
27.1 The economics: why the token price is the science budget	240
27.2 Open pretraining as insurance, not as a benchmark contest	242
27.3 Design: a shared platform with friendly competition	242
27.4 What already exists, and what is missing	243
27.5 Industry without expropriation, and the layered data policy	244
27.6 The long tail and the Global South as strategy	244
27.7 Sequence, and the test of seriousness	245
27.8 The institution, specified operationally	245
27.9 Why this is the strategically decisive institution	245

V	Synthesis and Summary	248
28	The Convergence	249
28.1	Assembling the elements	249
28.2	What would falsify this thesis	250
28.2.1	The strongest counter-argument: two recursive loops, not one	250
28.3	Two things the argument has to carry with it	251
28.4	Implications for the West—and for the rest	252
29	The Argument in Summary	253
29.1	The question, and the answer in one page	253
29.2	The four propositions, and where they now stand	254
29.3	What would change the answer	255
29.4	What to do about it	255
29.5	What this book is not claiming	256
A	The LPDDR6X Socket: Engineering Reality Check	257
A.1	The pin budget	257
A.2	What the trade buys, restated precisely	257
A.3	The reconciled design point	258
A.4	The hybrid tier, which is probably the better design	259
A.5	Tensor placement and migration policy for the hybrid	259
A.6	Package yield: where capability and competitiveness diverge	261
A.7	The package-complete checklist	261
A.8	The falsification program	261
B	High Bandwidth Flash: A Technical Dossier	262
B.1	What it is, and why this book cares	262
B.2	Lineage: three earlier attempts to make flash behave like memory	263
B.3	The announcement record, February 2025 to August 2026	263
B.4	The specification, as published, and what it omits	265
B.5	Power, energy, and the physics of reading flash fast	265
B.6	Endurance: the number nobody has printed	267
B.7	Cost structure: cheap bits, expensive stack	267
B.8	The perceived advantages, and the case against each	268
B.9	The competitive landscape	269
B.10	The research literature	270
B.11	What the main text assumes, and what would falsify it	270
C	The Hot Chips 2026 Ledger	273
C.1	The methodological warning: disclosure has collapsed unevenly	273
C.2	The data-centre CPUs	274
C.3	The accelerators	275
C.4	Four schools of thought for beating the GPU	275
C.5	FP64 bifurcated rather than disappeared	277
C.6	What the conference says about openness	278
C.7	What to watch	278
D	Evidence Ledger	280
E	Deployment Control Architecture	283
F	Consortium Feasibility: The Legal and Engineering Homework	285

G Living Data Appendix: Flagship Model Cards	287
H Current Cycle: From Models to Harnesses	289
H.1 The chronology	289
H.2 Qwen3.8-Max: what is claimed, what is verified	289
H.3 The second week: 4–11 August	290
H.4 The third week: 11–15 August	291
H.5 The financing turn: the same week, in capital markets	293
H.6 The memory question, entered into the record	293
H.7 The Pareto observation	293
H.8 What this changes, and what it does not	294
I Chronology	296
J Claim Register and Falsification Dashboard	298
J.1 Register schema	298
J.2 Probabilistic forecasts	299
J.3 Quarterly falsification dashboard	303
Methods, Sources, and Versioning	305

Notation, Abbreviations, and Glossary

This book draws on economics, computer architecture, semiconductor manufacturing, electrical-grid engineering, security research, and science policy, each of which brings its own vocabulary and its own habits of abbreviation. Several of them use the same letters for different things. What follows is the reference apparatus: the notation conventions, then the labelling schemes, then the mathematical symbols by chapter, then the acronyms, then the terms of art. Nothing here is required to follow the argument—every quantity is also defined where it is first used—but a reader who arrives at Chapter 26 wondering what a BGPU is should not have to hunt for it.

Notation conventions

GPT. Economists have used **GPT** for *general-purpose technology* since the early 1990s. Throughout this book that sense is always written out in full; the bare acronym GPT always denotes the OpenAI model family.

The suffix discipline. Six disciplines cannot share twenty-six letters without collision, and relying on context to separate them is not good enough for a document meant to be checked rather than read once. The rule adopted here is therefore mechanical: *no bare single letter carries two different meanings across the book*. Where two fields would have claimed the same letter, one of three moves is made, and the reader can rely on all three.

First, a roman suffix names the referent. A subscript set in upright roman type is a label, not an index: f_{heavy} is the heavy-user fraction, f_{front} the frontier-workload share, e_{INT8} the INT8 efficiency, d_{duty} the activity-duty product, r_{disc} the discount rate, T_{exit} the exit horizon. Italic subscripts, by contrast, are running indices (i, j, k, g, t, v) and mean what indices normally mean.

Second, a superscript names the physical dimension where two quantities share both letter and subscripts. Electrical power is written $P_{\bullet}^{\text{el}}$ throughout Chapter 12 and Appendix A— $P_{\text{facility}}^{\text{el}}$, $P_{\text{memory}}^{\text{el}}$ —so that bare P can remain the quality-adjusted price of Chapter 2 and Chapter 21 without ambiguity. Gross funnel output is U^{gross} , distinguishing it from the net uplift U that the same chapter reports.

Third, a distinct letter or an alternative alphabet is chosen outright where suffixing would be clumsy. Trust in the differentiated-competition model is τ , not T , leaving T for time and token streams. Fixed development cost is K_i , not F_i , leaving F for arithmetic rates and facility equivalents. Decode throughput is Θ_m and Θ_s , not R_m and R_s , leaving R for revenue and rent. Fleet availability is α , not a , leaving a for addressability in the AI4S funnel. Service-tier population share is φ_i , not s_i , leaving s for speed-up. The frontier-expenditure share is Σ_{front} , leaving the science-production metric $\mathcal{S}$ as the only S -family object. In the security fault tree the common-cause *events* are set in calligraphic type— $\mathcal{L}_{\text{sys}}$, $\mathcal{M}$, $\mathcal{B}$, with beta factors $\beta_{\mathcal{M}}$ and $\beta_{\mathcal{B}}$ —so that an event is never mistaken for a scalar.

Two conventions are deliberately *not* disambiguated because they are unambiguous by type. Probability is always written with the operator $\text{Pr}(\cdot)$ rather than a letter. And the letter C is

cost everywhere in the mathematics; the claim identifiers CR-01...CR-12 of Appendix J are an alphanumeric label, never a variable.

Labelling schemes

The book runs several labelled systems. They are independent, and this is the index:

Evidence grades [V] [P] [A] [M] [R] [S] — how well a claim is supported: **V** verified against primary or benchmark evidence; **P** primary claim by a party with direct knowledge (vendor, filing, government); **A** analyst estimate; **M** author-modeled; **R** roadmap or schedule; **S** scenario judgment. Defined in Section 1.5; the complete ledger is Appendix D.

Propositions P1–P4 — the four load-bearing claims, stated and falsified separately in Section 1.4: adoption (P1), sovereign compute (P2), rent destruction (P3), financial repricing (P4).

Security evidence Levels A–D — Chapter 19: **A** controlled technical demonstration; **B** observation on a real system or released artifact; **C** confirmed in-the-wild incident; **D** hypothesis or policy inference. Distinct from the evidence grades above.

Certification Levels 0–5 — Chapter 18: the graded assurance ladder from hash-identified artifact (0) to critical-infrastructure assured (5).

Parameter classes [U] [M] [T] [P] and [F] — Chapter 26 only, and *not* the evidence grades: **U** universal, **M** must-measure, **T** transferred, **P** policy choice, **F** future-scenario modifier. These classify model *inputs* by portability across borders.

Standards-maturity stages 1–4 — Chapter 20: de jure specification, open reference implementation, conformance suite, installed-base default.

Playbook moves M1–M9 and steps 1–6 — Chapter 7: M1–M9 are the tactical inventory of Table 7.1; steps 1–6 are the strategic sequence they serve.

Claim IDs CR-01–CR-12 and falsification branches (a)–(f) — Appendix J and Section 28.2 respectively.

Mathematical symbols

Symbol	Meaning (and where introduced)
Economics of intelligence (Chapter 2, Chapter 21)	
C_{task}	delivered cost of one <i>completed task</i> , decomposed into seven terms below
$c_{\text{pre}}, c_{\text{post}}$	amortized pretraining and post-training cost
c_{inf}	inference cost—the only term falling at $\sim 10\times/\text{year}$
$c_{\text{int}}, c_{\text{org}}$	integration and organizational-adoption cost
c_{comp}	complementary capital (pipelines, evaluation, process redesign, staff)
c_{trust}	assurance and certification cost
P	quality-adjusted price: price of a fixed amount of <i>useful work</i> , not of a token
Q	quantity of delivered work per period
R	revenue, $R = PQ$
ϵ	own-price elasticity of demand, $\epsilon = -\partial \log Q / \partial \log P > 0$; $\epsilon > 1$ is elastic

Symbols, continued

Symbol	Meaning
$\Delta \log X$	proportional (percentage) change in X
i, j	indices for competing differentiated systems (own and rival)
q, τ, I	capability, trust/assurance, and integration/distribution attributes of a system; τ (not T) is trust, by the suffix discipline above
π_i	profit of system i
MC_i, K_i	marginal cost and fixed (development) cost of system i ; K rather than F to keep F for arithmetic rates
$R_i^{\text{complement}}$	rent earned in adjacent layers because system i is widely used
NPV_v	net present value of data-centre vintage v ; u_t utilization in year t , q_t delivered throughput, P_t realized price, r_{disc} discount rate, T_{exit} exit horizon

Architecture and serving (Chapter 12, Appendix A)

T_{workflow}	end-to-end wall-clock time = $T_{\text{host}} + T_{\text{accelerator}} + T_{\text{movement}} + T_{\text{synchronization}} + T_{\text{orchestration}}$
B_{token}	bytes of memory traffic per generated token (six terms; Eq. 12.2)
$N_{\text{effective batch}}$	concurrent sequences over which one weight read is amortized
$F_{\text{achieved}}, F_{\text{token}}$	sustained arithmetic rate, and arithmetic work per token
$\text{BW}_{\text{achieved}}$	delivered memory bandwidth under the real access pattern
η	delivered fraction of peak memory bandwidth; swept 0.53–0.70
e_{INT8}	sustained INT8 matrix-engine efficiency (threshold $\gtrsim 0.38$)
$C_{\text{lifecycle}}$	nine-term total cost of owning a platform (Eq. 12.4)
$P_{\text{facility}}^{\text{el}}$	total facility <i>electrical</i> power; IT load plus cooling and overhead. The ^{el} superscript marks every power term, reserving bare P for price
Y_{package}	advanced-packaging yield, a product over all dies and interfaces
$P_{\text{memory}}^{\text{el}}$	full memory-subsystem power (DRAM, PHY, controller, PMIC, refresh, board, idle); likewise $P_{\text{socket}}^{\text{el}}, P_{\text{network}}^{\text{el}}, P_{\text{storage}}^{\text{el}}, P_{\text{cooling}}^{\text{el}}, P_{\text{overhead}}^{\text{el}}$
Δt_{eff}	effective delay in years purchased by an export control

Security (Chapter 19)

$\mathcal{L}_{\text{sys}}$	the systemic-loss <i>event</i> ; a fault tree over compromise, deployment, trigger, propagation, containment. Its probability is written $\Pr(\mathcal{L}_{\text{sys}})$
p_i	per-gate probabilities of the fault tree
$\mathcal{M}, \mathcal{B}$	common-cause events: monoculture depth ($\mathcal{M}$), and broad agent permissions ($\mathcal{B}$). Calligraphic, because they are events rather than scalars
$\beta_{\mathcal{M}}, \beta_{\mathcal{B}}$	beta factors—the share of a gate’s failure attributable to each common cause

AI for science (Chapter 24)

a	addressability: share of research effort reachable by current AI
s	speed-up on the addressable share; $(s-1)$ is the gain over baseline
v	validation survival: fraction of accelerated output that withstands scrutiny
d	deployment conversion: fraction of validated results that reach use
ℓ_j	lag-and-friction discount (calendar delay, organizational adjustment, regulation)

Symbols, continued

Symbol	Meaning
w_j	share of global research effort in task class j
U, U_{AI4S}	incremental uplift to realized research output; task-weighted form in Eq. 24.2
U^{gross}	gross converted output, $U^{\text{gross}} = asvd$ —contrast the net uplift $U = a(s-1)vd$. Not to be confused with the capacity requirements $G_\bullet$ of Chapter 26
$\mathcal{S}$	scientific production metric: a <i>vector</i> of validated outcomes over scalar cost
National sizing (Chapter 26, Chapter 27)	
BGPU	Baseline GPU: the AI-serving throughput of an 8 TB/s-HBM reference accelerator. A planning unit, not a product
BGPU _{BW}	separate bandwidth currency for classical simulation; <i>never</i> summed with serving-BGPU
γ_g, δ_g	conversion factor for accelerator g , and its benchmark-adjustment band
H_{yr}	effective delivered hours per installed BGPU-year ($= \alpha \times 8,760 = 7,884$)
α	fleet availability: fraction of wall-clock time an installed device is usable in production; 0.90 at the reference fleet. Written α , not a , to keep a for addressability
N	open-tier researcher population
φ_i, u_i, r_i	service-tier population share (φ , not s , which is speed-up), duty factor, and active output-token rate in tier i
T_{nat}	national output-token stream
k	BGPU cost per token per second of the canonical serving configuration
Θ_m, Θ_s	master and sub-agent decode throughput in tokens per second (Θ , not R , which is revenue); BW delivered bandwidth, L_m master resident context length, b_{KV} KV-cache bytes per token of context
$G_{\text{tok}}, G_{\text{sci}}$	token-service and scientific-execution capacity requirements
W_{ag}	annual agentic-execution workload in delivered BGPU-hours
$f_{\text{heavy}}, d_{\text{duty}}, w_{\text{heavy}}$	heavy-user fraction, activity-duty product, and per-heavy-user daily driver
χ_{sim}	simulation share: fraction of agentic execution that is simulation rather than training
$n_{\text{prog}}, w_{\text{prog}}$	count and annual cost of persistent-surrogate programs
$F_{\text{eq}}, g_{\text{fac}}$	large-instrument facility equivalents, and installed BGPU each requires
$\mathbf{B}_g, \omega_k$	device capability vector for accelerator g (seven components, each normalized to the reference device) and the declared national workload mix over those components
$\Sigma_{\text{front}}(f_{\text{front}}, \rho)$	frontier share of AI4S expenditure at frontier-workload share f_{front} and frontier price premium ρ

Abbreviations

Abbreviation	Expansion
Models and machine learning	
LLM	large language model
MoE	mixture of experts—a model where each token activates only a small subset of parameters
MLA	multi-head latent attention; GQA grouped-query attention; SWA sliding-window attention
KV cache	key–value cache: the stored attention state of the context so far
DSA	DeepSeek sparse attention
GRPO	group relative policy optimization (critic-free reinforcement learning)
MTP	multi-token prediction
QAT	quantization-aware training
FP8, FP4, BF16, INT8, INT4, MXFP4, UE8M0	numeric formats: floating-point and integer, at the stated bit widths; MXFP4 is a block-scaled 4-bit float; UE8M0 an 8-bit scale format targeted at domestic Chinese silicon
VLA	vision–language–action model (robotics)
API	application programming interface
RLHF	reinforcement learning from human feedback
HLE, GPQA, AIME, SWE-bench, MMLU	capability benchmarks (Humanity’s Last Exam; graduate-level Q&A; a mathematics olympiad set; software-engineering task suites)
TTC	test-time compute (parallel or extended reasoning at inference)
Elo	a relative rating derived from pairwise preference comparisons
Hardware, architecture, and manufacturing	
ISA	instruction set architecture—the contract between software and processor
ABI	application binary interface
RISC-V	an open, royalty-free ISA; RVV its vector extension, RVA23 a profile
SVE, SME	Arm’s scalable vector and scalable matrix extensions
IME, AME, VME	RISC-V integrated, attached, and vector–matrix extensions
HBM	high-bandwidth memory—stacked DRAM on the processor package
LPDDR	low-power double-data-rate memory—commodity mobile-class DRAM
DRAM	dynamic random-access memory
SOCAMM	a JEDEC module standard putting LPDDR on a server module
PHY, PMIC	physical-layer interface circuitry; power-management integrated circuit
TSV	through-silicon via—a vertical interconnect through a die
EUV, DUV	extreme- and deep-ultraviolet lithography
EDA	electronic design automation—chip-design software
CUDA	NVIDIA’s proprietary parallel-computing platform; CANN Huawei’s analogue
NVLink, UnifiedBus (UB), UALink	accelerator interconnect specifications
NPU, TPU	neural / tensor processing unit
MFU	model FLOP utilization—achieved fraction of peak arithmetic throughput

Abbreviations, continued

Abbreviation	Expansion
HPL, HPCG	the compute-bound and bandwidth-bound supercomputing benchmarks; TOP500 and Green500 the resulting lists
FLOP/s	floating-point operations per second; EF, PF, TF exa-, peta-, tera-
TOPS	integer operations per second (trillions)
PUE	power usage effectiveness—facility power divided by IT power
TCO	total cost of ownership
Energy and grid	
UHV, UHVDC	ultra-high-voltage (direct-current) transmission
BESS	battery energy storage system
RTO / ISO	regional transmission organization; PJM, ERCOT, MISO, SPP, CAISO are the US regional grid operators
CHP	combined heat and power
Institutions, policy, and finance	
AI4S	AI for Science
MIIT, CAC, MOST, NDRC	Chinese ministries: Industry and Information Technology; Cyberspace Administration; Science and Technology; National Development and Reform Commission
FYP	Five-Year Plan; SEI Strategic Emerging Industries; GB <i>Guobiao</i> national standards
NEV	new-energy vehicle; FiT feed-in tariff; LFP lithium iron phosphate
BIS, EAR	US Bureau of Industry and Security; Export Administration Regulations
WAICO	World AI Cooperation Organization
DOE, ANL, ORNL, BNL	US Department of Energy and its Argonne, Oak Ridge, and Brookhaven laboratories
ECMWF	European Centre for Medium-Range Weather Forecasts; AIFS its AI forecasting system
ARR	annual recurring revenue (an extrapolated run rate, not trailing revenue)
CDS	credit default swap; bps basis points (hundredths of a percent)
RPO	remaining performance obligations—contracted but unrecognized revenue
SPV	special-purpose vehicle (off-balance-sheet financing)
capex, NPV, FCF	capital expenditure; net present value; free cash flow
PPP	purchasing power parity

Terms of art

Prefill and decode The two phases of transformer inference. *Prefill* processes the whole input prompt at once and is compute-bound; *decode* generates output one token at a time, re-reading model weights and attention state for each, and is memory-bandwidth-bound. Almost every architectural argument in this book turns on the difference.

Roofline; operational (arithmetic) intensity; ridge point A performance model in which achievable speed is bounded by $\min(\text{peak compute}, \text{intensity} \times \text{bandwidth})$. *Operational intensity* is arithmetic performed per byte of memory traffic; the *ridge point* is the intensity

at which the binding constraint switches from bandwidth to compute. Workloads below the ridge are memory-bound.

The knee The bandwidth below which a design’s achievable utilization falls sharply—in this book, the ~ 5 TB/s requirement of the capacity-first memory architecture.

All-reduce, all-to-all Collective communication patterns. *All-reduce* combines values across all devices and returns the result to each; *all-to-all* exchanges distinct data between every pair, and is what mixture-of-experts routing requires.

Expert parallelism, expert imbalance Distributing a mixture-of-experts model’s experts across devices. *Imbalance* is uneven token routing, which idles devices and is the sparse workload’s most under-reported failure mode.

Wright’s law / learning curve The empirical regularity that unit cost falls by a constant percentage for each doubling of cumulative production. The engine of every Part I precedent.

Jevons paradox The observation, from Jevons on coal, that efficiency gains can *raise* total consumption of a resource when demand is sufficiently elastic. In this book’s terms, the case $\epsilon > 1$.

Involution (*neijuan*) The Chinese term for self-defeating internal competition: firms in a saturated sector expending ever more effort for shrinking margins. The “anti-involution” campaign is the state’s attempt to consolidate such sectors.

Tape-out; known-good die; chipkill Releasing a completed chip design to manufacturing; a die tested as functional before packaging; a memory-protection scheme surviving the failure of an entire device.

Pareto frontier The set of options none of which can be improved on one dimension without worsening another—here, the cost-versus-capability frontier of available models.

Busbar; firm power; curtailment; interconnection queue A *busbar* is the physical junction where generation meets the transmission network, and the point at which delivered price and available capacity are actually determined. *Firm* power is capacity guaranteed available when called; *curtailment* is generation deliberately not taken, usually because transmission or demand is unavailable; the *interconnection queue* is the backlog of projects awaiting permission to connect. A *capacity auction* procures firm capacity for a future delivery year.

Vintage A cohort of data-centre capacity built in the same period, sharing hardware generation, contract terms, and financing cost—the unit at which returns should be assessed.

Neocloud A specialist provider renting accelerator capacity, typically financed by debt secured on the hardware itself.

Abliteration A community technique that identifies and removes a model’s refusal or suppression behaviour by editing activations, without full retraining.

Sim-to-real The transfer of robot policies trained in simulation to physical hardware, and the gap that transfer exposes.

Estimand The quantity a study actually estimates, as distinct from the quantity of interest—the distinction that makes national compute estimates comparable or not.

Notified body Under EU law, an organization designated to assess conformity before a product may be placed on the market. The institutional form this book argues is missing for model artifacts.

Chokepoint decay This book’s term for the erosion of a control’s effectiveness as stockpiles, diversion, substitution, and architectural redesign accumulate.

Mate 60 moment Shorthand for a public demonstration that a control was ineffective, after the 2023 Huawei handset built on domestic 7 nm-class silicon. Used here for the first frontier model trained end-to-end on a domestic stack.

Chapter 1

The Thesis

1.1 Three wars, one playbook

Between roughly 2005 and 2025, China fought and won three complete technology wars. In solar photovoltaics, it went from assembler of imported cells to producing over 90% of every stage of the global supply chain, while module prices fell by more than 99% from their 1976 level [1, 2]. In lithium-ion batteries, a technology commercialized by Sony in 1991 and led for two decades by Japan and Korea, Chinese firms now hold roughly 69–75% of global EV battery installations and over 80% of global cell manufacturing capacity, with pack prices down 93% in real terms since 2010 [39, 40, 44]. In electric vehicles, China became simultaneously the world’s largest car market, the world’s largest car exporter, and home to the world’s largest EV maker, while foreign incumbents’ share of the Chinese market collapsed from 62% to 35% in five years [64, 61].

A fourth campaign, in telecommunications equipment, ran the same playbook against a network rather than a commodity and won everywhere that trust was not a procurement criterion; Chapter 6 scores it in lesser detail, because it is the precedent the AI export offer is most often compared to and the one whose central instrument—sovereign vendor credit—AI has not yet reproduced. These were not lucky outcomes. They were executions of a single, learnable playbook, whose elements Chapter 7 distills: patient state designation of a strategic industry; demand creation at home; subsidized scale-up of manufacturing far ahead of demand; brutal domestic competition that functions as an evolutionary selection engine; capture of the entire upstream supply chain; a cost-down learning-curve war that bankrupts foreign incumbents; and finally, once dominance is achieved, the conversion of that dominance into geopolitical leverage—export controls on the very technologies China was once accused of copying [49, 78].

1.2 Why AI is the fourth war

The prevailing Western assumption is that AI is different: that frontier AI depends on scarce, controllable inputs—leading-edge logic, high-bandwidth memory (HBM: stacked DRAM mounted on the accelerator package), proprietary interconnects, closed model weights—and that export controls on these chokepoints can durably preserve a lead. This book argues the assumption is failing on all three fronts simultaneously, and failing in a characteristic way: not because China matches the West chokepoint-for-chokepoint, but because it is engineering the chokepoints out of the system.

The model (Chapter 10). Chinese laboratories have converted the model layer into a commons. As of mid-2026, Chinese open-weight models occupy the top open positions on independent intelligence indices, account for roughly 41% of all Hugging Face downloads and a majority of tokens routed through neutral API aggregators (application programming interfaces—metered, hosted access to a model), and are priced at a small fraction of

Western closed models—while the measured capability gap to the best closed US models has compressed to low single digits [85, 84, 94]. Open weights at marginal cost do to the model layer what Chinese modules did to solar: they set a world price that rent-seeking business models cannot survive.

The compute (Chapter 12). The state has designated the open RISC-V instruction set architecture (ISA—the contract between software and processor, and the layer at which Arm and x86 collect licence rents) as national infrastructure, is standardizing matrix/vector extensions and chip-to-chip interconnects across vendors, and has open-sourced the software stacks (CANN, MindSpore, UnifiedBus specifications) that would otherwise fragment the domestic ecosystem [106, 132, 131]. In parallel, a memory-capacity arbitrage—substituting terabytes of cheap LPDDR (low-power DDR, the commodity mobile-class DRAM of phones and laptops) for scarce HBM—attacks the assumption that frontier training requires centralized, HBM-bound superclusters at all (Section 12.8).

The semiconductor (Chapter 13). Beneath both layers, the fabs. SMIC has reached 5 nm-class logic without EUV; CXMT has become the world’s fourth DRAM producer and the swing supplier in a global memory shortage; the domestic equipment industry has placed three firms in the global top twenty [139, 148, 154]. The remaining gaps—EUV lithography, leading-edge HBM—are real, but the strategy does not require closing them on the West’s terms; it requires making them irrelevant to the chosen architecture.

1.3 Openness as a weapon

The connective tissue among the three elements is a deliberate inversion of the Western technology-rent model. Where the United States concentrated AI value in proprietary layers—CUDA, NVLink, closed frontier weights, x86/Arm licensing—China’s strategy, formalized in the July 2025 Global AI Governance Action Plan and the “AI+” State Council initiative, is to make every one of those layers open, standardized, and free [99, 101]. This is not altruism; it is the same move as pricing solar modules at \$0.09 per watt. An open layer generates no rent for the incumbent, but it generates *adoption*, and adoption generates the scale, the ecosystem, and the standards-setting power that constitute victory in the precedent industries. The West’s chokepoint strategy and China’s commoditization strategy are not symmetric opponents: one defends margins, the other destroys them.


1.4 Four propositions, stated separately

Rigor requires decomposing the thesis into four propositions that are related but logically independent—none proves the next, and each has its own evidence and its own failure modes:

- P1 — Adoption.** China is winning open-weight model adoption. *Confirmed by:* download, derivative, and routed-token majorities persisting and extending into enterprise spend. *Falsified by:* a durable reversal of the OpenRouter/Hugging Face shares, or enterprise open-weight spend stagnating below ~15% while usage metrics plateau. Current grade: strongly supported [V].
- P2 — Sovereign compute.** China is constructing a domestically sufficient training-and-serving stack. *Confirmed by:* a frontier-class model trained end-to-end on domestic silicon, packaging, and memory. *Falsified by:* persistent dependence on stockpiled or smuggled foreign accelerators and HBM past ~2028. Current grade: rapid progress, not yet achieved [P/A/M].

- P3 — Rent destruction.** Open-weight commoditization will compress the closed-frontier rents that finance Western laboratories. *Confirmed by:* closed-model gross revenue growth decelerating while capability parity holds; margin compression propagating up the stack. *Falsified by:* closed labs sustaining pricing power through service, distribution, and trust layers even at model parity. Current grade: mechanism operating, magnitude contested [V/S].
- P4 — Financial repricing.** The compression of P3, if it occurs, punctures the US AI capital structure. *Confirmed by:* the indicator table of Chapter 21. *Falsified by:* demand elasticity (the Jevons case, §21.4) absorbing the price decline—AI getting cheap faster than usage grows is the bear case; usage growing faster than prices fall is the bull case, and it rescues the buildout even under P1–P3. Current grade: scenario judgment [S].

Figure 1.1 shows the propositions with their couplings and their contrary evidence. P1 can hold while P2 fails (China dominating distribution on stockpiled NVIDIA silicon); P3 can hold while P4 fails (margins compress while volume growth fills the data centers); and a US valuation correction could arrive from macro causes with none of P1–P3 mattering. The convergence argument of this book is that P1 is established, P2 is trending, and the *coupling* of $P1 \times P2$ —open weights meeting sovereign, capacity-first hardware—is what forces P3, with P4 as its priced-in consequence. Chapters 10–21 take the propositions in order.



None of these arrows is entailment. P1 can hold while P2 fails (dominance on stockpiled foreign silicon); P3 can hold while P4 fails (margins compress while volume growth fills the data centres, the elasticity case of §21.4); and P4 could arrive from macro causes with none of P1–P3 mattering. The convergence argument of this book is that the *coupling* $P1 \times P2$ —open weights meeting sovereign, capacity-first hardware—is what forces P3, with P4 as its priced-in consequence.

Figure 1.1: The argument, disaggregated. Each proposition carries its own evidence grade, its own confirming metrics, and its own contrary evidence; the book’s thesis is a claim about the coupling between them rather than a single chain of entailment. Grades as of the version stamp; Appendix D carries the full ledger and Appendix J the quarterly indicators that move them.

1.5 Evidence grades

A claim’s evidentiary weight is never self-evident from its phrasing, and so load-bearing claims are graded throughout: [V] independently verified (teardown, benchmark submission, audited filing, multiple independent confirmations); [P] primary-source claim (a company or government statement, reliable as to what was claimed, not what was achieved); [A] analyst estimate; [M] author-modeled result (including companion feasibility studies); [R] roadmap; [S] scenario

judgment. Appendix D collects the headline claims in one graded ledger. Where a section rests on the author's own companion studies—notably the memory-architecture analyses of Chapter 12—the text says so explicitly: those are modeled architectures, not measured production systems.

The thesis, stated in its conditional form: the United States retains the absolute model frontier, the leading semiconductor and cloud ecosystem, and vastly greater private capital. China increasingly leads a different contest—the commoditization of frontier-adjacent intelligence and the construction of a sovereign, though incomplete, compute stack. The decisive question is whether Chinese cost reduction, adoption, standards formation, and hardware sufficiency compound faster than US capability, distribution, and monetization—and whether China closes its remaining gaps in memory, packaging, tools, software maturity, and trust before the open-model economy compresses the rents that finance the Western frontier. This book argues that convergence is the way to bet; the chapters that follow document why, and Chapter 28 states what would change the bet.

The reader who finishes Part I will recognize, in Part II, every move already made once, twice, or three times before—and, in Chapter 8, the places where the precedents genuinely do not transfer.

Chapter 2

A Theory of the Contest

The preceding chapter stated a thesis and four propositions. This one supplies the theory they rest on, because a book that argues from three industrial precedents to a fourth needs to say *why* the analogy should transfer, and an argument built on the phrase “openness as a weapon” needs to explain what openness does to rents rather than merely asserting that it destroys them. Four bodies of theory are load-bearing: the production economics of intelligence, the literature on general-purpose technologies, industrial-organization and ecosystem theory, and the discipline of probabilistic forecasting. Readers who want the narrative can skip to Part I and return here when the argument makes a claim they doubt; readers who suspect the book of motivated reasoning should start here, because this is where its machinery is exposed.

2.1 The production economics of intelligence

Treat delivered intelligence as an output with a production function and a demand curve, and most of the confusion in the public debate resolves.

2.1.1 The cost stack

The cost of delivering one unit of *useful* AI work—not one token, but one completed task—decomposes into seven layers, each with its own learning curve and its own competitive dynamics:

$$C_{\text{task}} = \underbrace{c_{\text{pre}} + c_{\text{post}}}_{\text{amortized model creation}} + \underbrace{c_{\text{inf}}}_{\text{inference}} + \underbrace{c_{\text{int}} + c_{\text{org}}}_{\text{integration and adoption}} + \underbrace{c_{\text{comp}}}_{\text{complementary capital}} + \underbrace{c_{\text{trust}}}_{\text{assurance}} .$$

Frontier pretraining (c_{pre}) is a fixed cost amortized over deployment volume—which is precisely why an actor that gives models away and captures volume elsewhere can rationally spend on it, and why an actor that must recover it from metered access cannot cut price below the amortization line. Post-training and synthetic-data generation (c_{post}) has become a large and rising share, and it is where reinforcement learning converts compute into capability without a new pretraining run. Inference (c_{inf}) is the term everyone quotes and the only one falling at 10× per year. Integration and organizational adoption ($c_{\text{int}}, c_{\text{org}}$) are the terms enterprise buyers actually experience, and they are *not* falling—which is the single best explanation for the awkward fact of Chapter 10 that open-weight usage share rose while open-weight *spend* share fell. Complementary capital (c_{comp})—data pipelines, evaluation harnesses, process redesign, skilled staff—is the term the productivity literature identifies as the reason general-purpose technologies show up in the statistics years after they show up in the laboratories. And assurance (c_{trust}), the subject of Chapters 18 and 19, is a cost the commons has so far mostly *not* paid, which is why its apparent price advantage overstates its delivered advantage in regulated settings.

Two consequences follow immediately. First, *commoditization of c_{inf} does not commoditize C_{task}* ; it merely shifts the binding constraint to the terms that are not falling. This is why this

book’s claim about rent destruction is bounded to the model layer and stated as a hypothesis at the service layer. Second, the terms differ in who bears them: a national research system pays c_{inf} and c_{org} but can socialize c_{trust} through a shared certification body, which is the economic argument for Chapter 27 stated in one line.

2.1.2 Four learning curves, not one

Part I’s precedent industries had a single dominant learning curve—manufacturing—and Wright’s law described it. AI has at least four, and they move at different speeds and reward different actors: *hardware* learning (fab yield, packaging, system integration: the classic curve, and the one China’s industrial machine is built for); *algorithmic* learning (architecture, sparsity, precision, attention—roughly a threefold annual efficiency contribution on published estimates, and the curve on which Chinese laboratories have contributed most per dollar); *software and kernel* learning (compilers, serving stacks, collective libraries—the curve where the incumbent’s seven-year CUDA lead lives, and where the LineShine class is furthest behind); and *operational* learning (fleet reliability, utilization, scheduling, failure recovery—the curve nobody publishes and every hyperscaler competes on). A country can lead on two and lose on two. Chapters 10 and 12 argue China leads the algorithmic curve and is buying the hardware curve, while the software and operational curves remain the West’s strongest defensive positions—which is a more precise statement of this book’s thesis than any claim about model rankings.

2.1.3 Efficiency price cuts versus strategic price cuts

A discipline the book owes its reader: a falling price is not evidence of a falling cost. Price declines from genuine efficiency (c_{inf} falling) are durable and reprice the industry permanently. Price declines from strategic subsidy—state-discounted power, below-cost token pricing to buy adoption, capacity subsidized by a provincial balance sheet—are reversible, and their strategic function is to acquire the volume that generates the learning that eventually makes them durable. Part I shows this sequence completing three times: solar, batteries, and EVs all passed through a below-cost phase whose losses landed on Chinese balance sheets and whose learning stayed. The correct reading of Chinese token pricing is therefore neither “dumping, therefore fake” nor “cheap, therefore superior technology,” but *a subsidy purchasing a position on a learning curve*—with the empirical question being whether the curve delivers before the subsidy exhausts patience.

2.1.4 The demand side, and why it decides everything

The whole financial argument of this book turns on one question that can be stated without any mathematics at all: *when the price of AI falls, does the money spent on AI go up or down?* If buyers respond to cheaper tokens by buying so many more of them that total spending rises, the incumbents’ capital expenditure is vindicated. If they buy more but not enough more, total spending falls, and every valuation premised on the old price collapses with it. Everything below is that sentence made precise.

Three quantities are needed. Let P be the *quality-adjusted price* of delivered AI—the price of a fixed amount of useful work, not of a token, so that a model which halves its price while doubling the tokens it burns per task has not moved P at all. Let Q be the *quantity* of that work purchased per period. Let $R = P \times Q$ be the resulting *revenue* of the sellers, which is what valuations discount.

The responsiveness of quantity to price is the *own-price elasticity of demand*, written ϵ and defined so that it is a positive number:

$$\epsilon = - \frac{\partial \log Q}{\partial \log P} > 0.$$

In words: ϵ is the percentage rise in quantity produced by a one-percent fall in price. (The logarithms simply convert “percentage change” into arithmetic that adds up; the minus sign makes ϵ positive, because quantity moves opposite to price. An ϵ of 2 means a 10% price cut brings a 20% volume increase.) The convention is stated explicitly because the sign is inconsistently defined in the literature, and a sign error inverts the conclusion.

Since $R = PQ$, taking logarithms gives $\log R = \log P + \log Q$, and differencing gives the relation this book uses throughout, where $\Delta \log X$ denotes a proportional (percentage) change in X :

$$\Delta \log R = \Delta \log P + \Delta \log Q = (1 - \epsilon) \Delta \log P.$$

Read the result rather than the derivation. If $\epsilon > 1$, the bracket $(1 - \epsilon)$ is negative, so a *falling* price ($\Delta \log P < 0$) produces *rising* revenue: demand is *elastic*, and price cuts pay for themselves. If $\epsilon < 1$, demand is *inelastic*, the bracket is positive, and falling prices shrink the industry’s revenue even as its output grows. The knife-edge at $\epsilon = 1$ is exactly the boundary between the bull and bear cases for the entire Western AI complex, which is why Chapter 21 returns to it.

The refinement that single-number arguments on both sides miss is that ϵ is not a number but a *portfolio*. It is plausibly well above one for coding agents, synthetic-data generation, video, and the scientific inference of Chapter 24, where cheaper tokens unlock work that was previously uneconomic and demand is bounded only by budget. It is plausibly near or below one for consumer chat and routine summarization, where consumption is bounded by human attention rather than by price—nobody reads twice as many chatbot replies because they became free. Aggregate behaviour is a spending-weighted mixture of segment elasticities, which is why neither camp’s single figure predicts it, and why this book treats the aggregate as unresolved rather than assumed.

2.1.5 From own-price elasticity to differentiated-system profit

The own-price relation just derived answers “what happens when *the* price falls,” and that is the wrong question for this book’s third proposition. There is no single AI product with a single price. There is a Chinese open-weight model that can be downloaded, an American closed model sold through a metered interface, and a dozen intermediate offerings, and the strategic question is not how buyers respond to a price cut but *which system they choose when the alternatives differ in price, capability, trustworthiness, and ease of integration all at once*. That is a question about substitution between differentiated products, and it needs a different instrument. The distinction is not a refinement; a book that argues P3 from own-price elasticity alone is answering a question nobody in this market actually faces.

Index the competing systems by i , and let j stand for its rivals—so P_i is the price of system i and P_j the price of the competing system, and likewise for every other quantity. Four attributes decide a purchase:

P — **price** what the buyer pays per unit of delivered work, as in §2.1.4.

q — **capability** how good the system is at the buyer’s actual task: the quality dimension that lets a more expensive system win anyway.

τ — **trust and assurance** whether the system’s provenance, behaviour, and liability can be established to the standard the buyer’s sector requires—the subject of Chapter 18, and the reason a regulated buyer may refuse a free model outright.

I — **integration and distribution** how much work it takes to connect the system to the buyer’s existing data, tools, and processes, and whether it is already in front of the user—the c_{int} and c_{org} terms of the cost stack (§2.1).

Demand for system i then depends on all eight quantities, its own and its rival’s:

$$Q_i = Q_i(P_i, P_j, q_i, q_j, \tau_i, \tau_j, I_i, I_j).$$

The content of that expression is entirely in what it refuses to omit. A price-only model deletes q , τ , and I , and therefore predicts that a free model with adequate capability sweeps the market. The real world contains buyers who will not deploy an unauditable artifact at any price, and buyers for whom switching cost exceeds several years of licence fees, which is why open-weight *usage* share and open-weight *spending* share have moved in opposite directions (Chapter 10).

Now the seller’s side. Write π_i for the profit of system i , MC_i for its *marginal cost* (the cost of serving one more unit of work—for AI, essentially inference), and K_i for its *fixed cost* (model development: the pretraining run, the research programme, the salaries, incurred before the first unit is sold). Then

$$\pi_i = \underbrace{(P_i - MC_i) Q_i}_{\text{gross margin on the model itself}} - \underbrace{K_i}_{\text{development}} + \underbrace{R_i^{\text{complement}}}_{\text{rent from everything adjacent}},$$

where $R_i^{\text{complement}}$ collects the profits the seller earns *elsewhere* because this system exists and is widely used: cloud services it pulls through, hardware it sells, advertising it enables, proprietary data it accumulates, applications and support it attaches, integration it bills for.

That third term is the book’s central insight rendered as an accounting identity. Openness relocates rent rather than necessarily destroying it, and here is the mechanism: a seller can rationally drive P_i down to MC_i , or below it, *whenever the extra volume Q_i that results raises $R_i^{\text{complement}}$ by more than the margin given up*. Giving the model away is not charity and not dumping; it is a bet that the complement is where the scarcity will sit. Both flywheels of §2.3.3 are strategies for maximizing $R^{\text{complement}}$, differing only in which complement each side believes will be scarce in 2030—and Table 21.4 in Chapter 21 is simply this term, itemized.

The same formulation disciplines this book’s own thesis, and it corrects a tempting overreach: it is *not* true that “the Chinese position wins under either elasticity regime.” A low-cost position fails strategically if pricing stays below full cost while K_i cannot be recovered; if trust and integration terms (τ , I) dominate the purchase decision so that price undercuts do not move Q ; if thin margins starve the operational-learning loop that was the point of buying volume; or—the American flywheel’s theft, now expressible in the notation—if US application providers build atop Chinese weights, capturing the customer relationship and its $R^{\text{complement}}$ while the weights earn nothing, leaving the open stack with tonnage and no margin. What the elasticity regimes do determine is asymmetric *exposure*: the leveraged Western capital structure requires $\epsilon > 1$ at the premium tier specifically, while the Chinese position degrades more gracefully under the unfavourable regime because its costs are lower and its rent expectations sit in hardware, standards, and infrastructure rather than in model-layer margin. Graceful degradation is a real advantage; it is not a theorem of victory, and Chapter 21’s indicators are what decide the empirical question.

2.2 AI as a general-purpose technology

A terminological warning. Economists abbreviate *general-purpose technology* as **GPT**, a usage that predates the model family sharing those initials by roughly three decades. The collision is unfortunate and unavoidable, since both terms are standard in their fields and this book needs both. *Throughout this book, “GPT” in the economic sense—and only in that sense—is written out in full as “general-purpose technology.” The bare acronym GPT always refers to the OpenAI model family* (GPT-4, GPT-5.6, and so on). Where the phrase “GPT literature” appears, it means the economics literature cited in this section.

The solar–battery–EV analogy is productive and incomplete, and Chapter 8 already states its limits. The deeper correction is categorical: those are manufactured product systems, whereas AI is also a *general-purpose technology*—a technology that is pervasively applicable across sectors rather than confined to one, that continues improving for decades after introduction, and that

provokes complementary innovation in the industries adopting it, so that its economic effect arrives mostly through what *other* sectors build with it [358, 359]. Electricity, the semiconductor, the internal-combustion engine, and the Internet are the canonical instances. That places AI closer to those than to the photovoltaic module, and it changes what “winning” can mean: a country can dominate the production of a general-purpose technology and still capture little of its value, because most of the value appears downstream in the hands of whoever deploys it best.

2.2.1 Four layers of winning




Figure 2.1: Winning has four meanings, and they are won separately. The press scores layer 1; the general-purpose-technology literature locates most realized value at layers 3 and 4, arriving late and only after complementary investment [360]. Positions shown are this book’s assessment as of the version stamp [S]; layer 4 is genuinely unresolved for both parties, which is the honest state of the question.

A general-purpose-technology contest is decided at four distinct layers, and a country can lead one while losing another (Figure 2.1):

Production leadership — who makes the best models and the most hardware. This is the layer the press scores, and the only one on which the United States is unambiguously ahead.

Platform leadership — who sets the standards, the interfaces, the reference implementations, and the default stack. Chapter 20 argues this layer is going to whoever opens fastest, and that the incumbent is conceding it a layer at a time.

Diffusion leadership — who actually deploys the technology across firms, laboratories, hospitals, and government. This is where the general-purpose-technology literature locates most of the realized value, and where the productivity J-curve says the returns arrive late and only after complementary investment [360].

Productivity leadership — who converts the first three into measured economic and scientific output. This is the only layer that ultimately matters, and it is the least observable in real time.

Two asymmetric failure modes follow, and both are historically attested. A country can *lead the core technology and underperform* if its institutions fail to diffuse it—the standard reading of the productivity paradox, and a live risk for an American ecosystem whose AI value

Table 2.1: The four layers of general-purpose-technology winning as a matrix: definition, positions, and what to watch.

Layer	Definition	US position	China position	Observable metric	Grade
Production	best models, most hardware	leads absolute frontier	leads open frontier; 41% domestic chip share	AI Index gap; shipment shares	V
Platform	standards, interfaces, default stack	CUDA moat; conceding ISA layer	RISC-V, UB, CANN opened; state-mandated interop	standards adoption; derivative-model share	V/A
Diffusion	actual deployment across the economy	enterprise pilots failing at high rates	46% global model share; Global South default	routed tokens; production deployments	V/A
Productivity	measured economic and scientific output	unmeasured; curve lag	J-unmeasured; survey signals only	SDG-sector productivity series; Eq. 24.3	M/S

is concentrated in a handful of firms and whose enterprise pilots reportedly fail at high rates. And a country can *fail to dominate the core technology and still capture enormous gains* through complementary organizational and industrial change—which is, precisely, the strategy this book attributes to China, and the reason the “but the best model is American” rebuttal misses the argument.

2.2.2 Why this strengthens the AI-for-Science case

The general-purpose-technology framing does something else important: it explains why AI for Science (AI4S) is categorically different from other application markets rather than merely larger. In general-purpose-technology terms, science is where the *innovation complementarity* operates at maximum strength—AI improving the production of new materials, drugs, energy systems, algorithms, and semiconductor technology, each of which feeds back into the technology’s own production function. Chapter 24 makes that argument empirically; here it is placed in its theoretical home. A general-purpose technology that accelerates the production of general-purpose technologies is the mechanism behind every explosive-growth model in the literature, and also the reason the sober estimates that exclude it arrive at small numbers.

2.3 Industrial organization: openness relocates rent, it does not abolish it

This book’s central slogan—openness as a weapon—requires the discipline that industrial-organization theory imposes. Open layers do not destroy profit; they *commoditize complements* so that rent reappears in an adjacent, controlled layer. The correct strategic question is therefore not whether a layer is open but where the opener expects the rent to return.

2.3.1 The rent-relocation table

Read the last row as the actual strategic bet each side is making. China opens the layers where it is behind or where openness buys adoption, and expects to collect in volume, standards, and infrastructure—the layers its industrial system is best at. The United States keeps the layers where it is ahead and expects to collect in capability, distribution, and integration—the layers its corporate system is best at. *Neither is being altruistic and neither is being naive*; they are making opposite bets about which layer will be scarce in 2030. This book’s thesis is a bet on the Chinese reading, and Chapter 28 states what would show it wrong.

Table 2.2: Where each side opens, and where it expects rent to reappear. Opening a layer is a claim about the adjacent layer.

Layer	China: open or controlled?	United States: open or controlled?
Model weights	Open (MIT/Apache flagships)	Controlled (metered closed frontier)
Instruction set	Open (RISC-V, national policy)	Controlled via Arm/x86 licensing, but conceding—CUDA now hosts on RISC-V
Interconnect / cluster spec	Open (UnifiedBus published; state-mandated interop)	Controlled (NVLink; consortium alternatives exclude China)
Framework / kernel stack	Open (CANN, MindSpore open-sourced)	Controlled (CUDA; the deepest moat)
<i>Expected rent recap- ture</i>	hardware volume, domestic cloud, national standards, data-centre and energy construction, embedded AI products, dependence of the adopting stack	premium model capability, distribution and habit, enterprise integration, trusted deployment, proprietary data, agent platforms, cloud orchestration


2.3.2 The concepts the book uses, named

Several ideas recur throughout and are worth naming in their standard forms, so that readers can connect the argument to the literature it draws on: the *industrial commons*—the shared pool of engineers, suppliers, and tooling whose loss makes re-entry infeasible, which is what Part I claims the West surrendered three times; *ecosystem competition*, where the unit of rivalry is a coalition rather than a firm [366]; *modularity and vertical integration*, which determine where value can be captured and which Baldwin and Clark’s design-rules framework formalizes [365]; *platform envelopment*, the move by which an incumbent in one layer absorbs an adjacent one—exactly what CUDA-on-RISC-V is, read charitably; *complements and bottlenecks*, the reason commoditizing your complement raises your own margin; *standards capture* and *architectural control points*, the subject of Chapter 20; *dynamic capabilities*, the firm-level ability to reconfigure that explains why some incumbents survive transitions [364]; and *disruptive versus sustaining trajectories*, which is the correct frame for the LPDDR-versus-HBM argument—a technology that is worse on the incumbent’s primary metric (bandwidth), better on a neglected one (capacity per dollar), and initially adopted by customers the incumbent does not want [363]. Schumpeter supplies the background claim that this is how growth normally happens [361], and Porter the reminder that structure, not virtue, determines who profits [362].

2.3.3 The dual flywheel

The contest can now be drawn (Figure 2.2). The American flywheel runs *capability* → *premium pricing* → *capital formation* → *larger clusters* → *stronger frontier models* → *distribution and application revenue*. The Chinese flywheel runs *open deployment* → *volume* → *learning and optimization* → *lower cost* → *broader adoption* → *standards and hardware demand* → *further deployment*. Each is internally coherent, each has historical precedent, and—the point that matters—they *can coexist*. The contest is not necessarily winner-take-all.

What makes it a contest rather than two parallel businesses is that each flywheel can appropriate the other’s critical input. The Chinese wheel steals the American wheel’s *volume*: application demand routed to open weights is demand that does not fund the next frontier run, and the capital-formation step is the American wheel’s most fragile link (Chapter 21). The American wheel steals the Chinese wheel’s *value*: closed agents capturing the highest-value workflows leave the open stack serving high-volume, low-margin traffic, which is the classic fate



Each wheel is internally coherent and both can turn at once; the contest is which one appropriates the other's critical input first.

Figure 2.2: Two coherent strategies, two thefts. The American wheel converts capability into price into capital into more capability; the Chinese wheel converts openness into volume into learning into lower cost into more adoption. Neither is irrational and neither requires the other to fail. The strategic contest is the pair of dashed arrows: whether open weights absorb enough application volume to starve the capital-formation step, or whether closed agents capture enough of the high-value workflows to leave the open stack with tonnage and no margin. Both are observable in the indicators of Chapter 21.

of a commoditized layer, and which would leave China with the tonnage and America with the profit. Both thefts are observable in the indicators of Chapter 21, and which one dominates is, on this book's reading, the single most important unresolved empirical question in the field.

2.4 Forecasting discipline

The final piece of theory is methodological. A book making contingent claims about a fast-moving system should state them as forecasts with horizons, probabilities, and falsifiers—not because the numbers are precise, but because the alternative permits confidence in one part of the argument to leak into another. Appendix J therefore carries every load-bearing proposition with a probability, a horizon, an evidence grade, a next verification event, and an explicit falsifier; Chapter 21 prints two probability sets side by side—the author's and the strongest skeptical case against it—rather than one; and Chapter 28 states the six developments that would falsify the thesis outright. The purpose is not false precision. It is that a reader in 2028 should be able to establish, mechanically, which parts of this book were right.

Part I

The Precedents: Three Wars Already Won

Chapter 3

Solar Photovoltaics: The Template

Solar PV is the template war: the first, the most complete, and the one whose mechanics were most visibly repeated in batteries and EVs. Every structural element of the playbook appears here in its purest form, and the final state—above 90% Chinese share at every stage of the supply chain—defines what “winning” means in this book.

3.1 The West invents, the West leads (1954–2007)

The silicon solar cell was invented at Bell Labs in 1954; the modern PV industry was built on American, German, Japanese, and Australian science. As late as 2006, Sharp of Japan was the world’s largest cell producer; in 2007 Germany’s Q-Cells overtook it at 370 MW of annual production [3]. Demand was created almost entirely by Western policy: Germany’s 2004 EEG amendment set rooftop feed-in tariffs at 57.4 euro cents per kWh, guaranteed for twenty years, igniting the first global solar boom [4]. The scientific pipeline that would matter most, however, ran through Australia: Shi Zhengrong, a PhD student of Martin Green at UNSW, returned to China to found Suntech Power in Wuxi in 2001—with municipal-government seed capital.

3.2 The first scale-up and the first bankruptcies (2005–2013)

Suntech listed on the NYSE in December 2005—the first private Chinese company to do so—and Shi briefly became the richest man in China; by end-2011 Suntech’s capacity had reached 2 GW [5]. Chinese producers, feeding on German and later Italian and Spanish FiT demand, scaled cell and module manufacturing at a pace no Western firm attempted. The financing was policy: in 2010 alone, the China Development Bank extended approximately \$47 billion in credit lines to renewable-energy manufacturers, including \$7.8 billion to Suntech, \$5.6 billion to Yingli, and \$4.7 billion each to Trina and JA Solar [6]. In October 2010 the State Council designated new energy one of seven Strategic Emerging Industries, and the 12th Five-Year Plan for Solar PV set explicit industrial targets [7].

The polysilicon price collapse did the rest. Spot polysilicon peaked near \$475/kg in March 2008; by December 2012 it was below \$16/kg [8]. Western and Japanese incumbents, structured for high-price niche markets, were annihilated by the combination of Chinese capacity and collapsing input costs: Solyndra failed in 2011 despite a \$535 million US loan guarantee; Q-Cells—world number one four years earlier—entered insolvency in April 2012 and was sold to Korea’s Hanwha; SolarWorld, the petitioner of the US trade cases, followed in 2017 [9, 10, 11].

Notably, the shakeout consumed the Chinese pioneers too. In March 2013 Suntech defaulted on \$541 million of convertible bonds—the first mainland Chinese default on US bonds—and its Wuxi subsidiary entered bankruptcy [5]. This is a recurring and under-appreciated feature of the playbook: China does not protect its champions individually; it protects the *industry*, and lets

domestic competition kill firms freely. The capacity, the workforce, and the supplier network survive the corpses.

3.3 Trade defenses fail (2012–2018)

The Western response was tariffs. The US Commerce Department’s October 2012 determinations set anti-dumping margins of 18–32% on major Chinese producers (250% country-wide) plus countervailing duties of 15% [12]. The EU imposed provisional duties in June 2013, scheduled to average 47.6%, but—under Chinese pressure on other trade fronts—settled within weeks for a minimum-import-price undertaking [13]. The 2018 US Section 201 safeguard added a 30% tariff [14].

None of it worked. Chinese producers moved cell and module assembly to Southeast Asia, moved up the value chain, and—decisively—kept the learning curve. Tariffs taxed Western deployment without resurrecting Western manufacturing; by the time the tariff walls were fully built, there was no meaningful Western industry left inside them to protect.

3.4 Domestic demand as strategic ballast (2011–2026)

From 2011 China built a domestic FiT-driven market that eventually dwarfed the export markets, insulating its manufacturers from Western trade policy entirely. The “531” shock of May 2018—an overnight cut of FiTs and subsidized quotas—was itself a deliberate stress test that culled weaker producers [15]. The Top Runner auctions (2015–2017) tied market access to efficiency floors, forcing PERC and mono-crystalline technology from niche to standard [16]. Domestic installations reached 277 GW_{AC} in 2024 and 315 GW_{AC} in 2025 (roughly 415 GW_{DC}, about 60% of world additions), bringing cumulative Chinese capacity to 1.2 TW—by itself more than the entire world had installed as of 2022 [17, 18].

3.5 The end state: total supply-chain capture

Table 3.1 shows the 2023 production shares by stage. The IEA’s assessment is that shares of polysilicon, ingot, and wafer capacity are converging on 95%; nine of the top ten polysilicon producers are Chinese (Germany’s Wacker, the survivor, fell to eighth); and Chinese manufacturing costs run 10% below India, 20% below the US, and 35% below Europe [19, 20]. Global module capacity reached ~ 1.8 TW/yr in 2025 against ~ 0.6 – 0.7 TW of demand [21]; China exported 268 GW of modules in 2025, and cell exports grew 73% as tariff-protected markets imported the layer below the tariff line instead [22, 23].

Table 3.1: Chinese share of global PV production by stage, 2023, and price collapse.

Stage	China share 2023	Note
Polysilicon	92%	93.5% in 2024; 9 of top 10 firms
Wafers	98.1%	near-monopoly
Cells	92%	cell exports +73% in 2025
Modules	84.6%	~ 1.8 TW/yr global capacity
Module price	\$106/W (1976) $\rightarrow$ \$0.09–0.10/W (2026)	
Learning rate	$\sim 20\%$ price decline per doubling of cumulative capacity	

The cost curve is the weapon. Across five decades the module price fell from \$106 per watt to roughly nine cents—a 99.9% decline—following a Wright’s-law learning rate of about 20% per doubling of cumulative production [1, 24]. Whoever manufactures the doublings owns the

learning; from roughly 2010 onward, essentially all doublings were Chinese. Figure 3.1 shows both halves of the result: the curve, and the stage-by-stage capture at its end.




Figure 3.1: The template war. *Above*: module price against cumulative installed capacity, both log scales—a Wright’s-law decline of roughly 20% per doubling, from \$106/W in 1976 to \$0.09–0.10/W in 2026 [1, 24]. Whoever manufactures the doublings owns the learning; from about 2010 essentially all doublings were Chinese. *Below*: 2023 production shares by stage [2]. Values [V]; the price series mixes contemporaneous reporting with index reconstruction and is directional at any single point.

3.6 Winning is expensive: the involution phase

Victory did not end the war economics. With capacity at triple demand, 2024–2025 prices fell below cash cost across the chain: polysilicon dipped under \$4.50/kg, and five Chinese majors forecast combined 2025 net losses of RMB 29–33 billion (\$4.1–4.7 billion) [25]. Beijing’s response is instructive for the AI chapters: in July 2025 Xi Jinping personally launched the “anti-involution” campaign against “disorderly low-price competition”; the six largest polysilicon producers organized a ~RMB 50 billion fund to buy out and retire a third of national capacity; and in December 2025 an official “storage and integration platform” was chartered to manage polysilicon supply—an OPEC-for-solar [26, 27, 28]. Having captured the industry, the state now moves to cartelize it and harvest stable rents. Meanwhile technology leadership has also passed to China: LONGi holds the world perovskite-tandem efficiency record (35.5%, July 2026), and Trina leads global perovskite patent filings [29, 30].

3.7 The Western rebuild that wasn't

The US Inflation Reduction Act briefly built module-assembly capacity from 8 GW to 56.5 GW (mid-2025), but cell capacity barely materialized, and the July 2025 OBBBA phase-out of deployment credits cut ten-year US installation forecasts by 17–41% [31, 32]. First Solar survives—profitably—as a CdTe niche protected by tariffs and its non-silicon technology; it produced 16.1 GW in 2025, roughly *one fortieth* of Chinese module shipments [33]. That ratio, four decades after the US invented the technology, is the terminal scoreboard of the first war.

3.8 Lessons carried forward

Four lessons from solar recur throughout this book. *First*, trade barriers imposed after cost-curve capture protect nothing but domestic prices. *Second*, the champions are disposable; the industrial commons is not. *Third*, a subsidized home market large enough to absorb half of world demand makes foreign trade policy irrelevant. *Fourth*, the endgame of commoditization is state-managed consolidation—dominance first, margins later. Keep all four in mind when reading about open-weight models priced near zero.

Chapter 4

Lithium-Ion Batteries: Vertical Depth

If solar demonstrated the cost-curve war, lithium-ion batteries demonstrated something further: the capture of an entire vertical—from mine to refinery to cathode to cell to pack—such that even a competitor who builds a gigafactory outside China still builds it, in effect, *inside* the Chinese supply chain.

4.1 From Japan’s invention to China’s entry (1991–2011)

Sony commercialized the lithium-ion battery in 1991, building on Asahi Kasei’s and John Goodenough’s chemistry; Japan then Korea (LG Chem, Samsung SDI) led for two decades [34]. China’s entry came through consumer electronics: BYD, founded in Shenzhen in 1995 with RMB 2.5 million and twenty employees, was by 2002 the world’s largest nickel-cadmium battery maker with 65% of global production, supplying Motorola and Nokia [35]. CATL’s lineage is similar: Robin Zeng’s ATL (1999) built lithium-polymer cells for phones (and was sold to Japan’s TDK); in 2011 Zeng spun the EV battery business out as CATL in Ningde [36]. The pattern—enter at the low-margin consumer end, accumulate manufacturing capability, then pivot into the strategic segment when the state creates demand—recurs in Part II.

4.2 The white list: protection by regulation, not tariff (2015–2019)

The decisive policy instrument was not a tariff but a certification. From 2015, NEV purchase subsidies applied only to vehicles using batteries from an MIIT-approved “white list” of 57 manufacturers—all domestic; LG and Samsung, which had just built Chinese plants, were excluded [37]. For four years, the world’s largest and fastest-growing EV market was reserved, in effect, for CATL, BYD, and their domestic rivals. The list was scrapped in June 2019—precisely when the champions no longer needed it. Subsidy design was also used to steer chemistry: density-linked subsidies (2017–2019) pushed nickel-based NCM when that was deemed strategic, then cost-linked criteria (2020) resurrected LFP [38]. Regulation as industrial policy, adjusted dynamically, with foreign firms as designated losers.

4.3 The end state: shares and costs

The 2025 scoreboard (Table 4.1): a 1.19 TWh market, with CATL and BYD alone above 55% and Chinese makers collectively at 69% of installations and ~75% of deployment; China holds over 80% of the world’s 4 TWh of cell manufacturing capacity, against 6–7% each for the EU

Table 4.1: Global EV battery installations, full-year 2025 (SNE Research) [39].

Maker	GWh	Share (%)
CATL (CN)	464.7	39.2
BYD (CN)	194.8	16.4
LG Energy Solution (KR)	108.8	9.2
SK On (KR)	—	3.7
Panasonic (JP)	—	3.7
Samsung SDI (KR)	—	2.4
World total	1,187.0	100.0
All Chinese makers	~820	~69

and US [39, 40]. Upstream the shares are deeper still: roughly 85% of anode capacity, 82% of electrolytes, 74% of separators, 70% of cathodes; over 90% of graphite processing; two-thirds of lithium and cobalt refining; and, through Tsingshan and partners, control of three-quarters of Indonesia’s nickel refining [41, 42, 43].




Figure 4.1: The second cost-curve war. *Above*: BloombergNEF’s volume-weighted average pack price, real terms, 2010–2025 [44]; the 2022 uptick is the lithium spike. *Below*: the 2025 regional gap—North American and European packs run 44% and 56% above Chinese [44]. All values [V].

The cost curve replicated solar’s. BNEF’s pack price series fell from \$1,474/kWh (2010) to \$108/kWh (2025), a 93% real decline; the China-specific price is \$84/kWh against \$121 in North America and \$131 in Europe—a 44–56% Western cost penalty at the pack level [44].

4.4 The LFP counterattack: innovation where the West wasn't looking

The most strategically instructive episode is lithium iron phosphate. LFP was invented in the West (Goodenough's patent, 1997) and dismissed by Western industry as an energy-density dead end; forecasts gave it 15–40% of the 2030 market. Chinese engineering—BYD's cell-to-pack Blade architecture, CATL's equivalent—recovered the density penalty at pack level, and LFP's cost, safety, and cycle-life advantages did the rest: LFP passed 50% of the *global* EV market in 2024 and 55% in 2025 [38, 40]. The West now licenses the technology it invented back from China: Ford's Michigan plant operates under CATL license, and a CATL–Stellantis JV is building a €4.1 billion LFP plant in Spain [45]. The lesson for AI: the “inferior” technology branch, industrialized ruthlessly at scale, beats the premium branch on system economics—a preview of the LPDDR-versus-HBM argument of Chapter 12.

China is now ahead on the next branches too: CATL's Naxtra sodium-ion line entered mass production in December 2025 (175 Wh/kg, 10,000 cycles, -40°C operation); second-generation Shenxing LFP charges 520 km in five minutes; BYD's platform charges at one megawatt [46, 47, 48]. Solid state remains genuinely open, but the IEA's sober assessment is prototypes industry-wide, with BYD targeting 2027 [40].

4.5 From protection to gatekeeping

The arc completed in 2025: the infant industry became the gatekeeper. In July 2025 MOFCOM placed LFP/LMFP cathode process technology and lithium-extraction technology under export license; in October the controls extended to high-energy cells ($\geq 300\text{ Wh/kg}$), anode materials, and the manufacturing equipment itself [49]. China now restricts the export of battery *technology* exactly as the US restricts semiconductors—the mirror image the West did not anticipate when it assumed technology flowed only one way.

4.6 The Western casualty list

The counterfactual competitors did not fail for lack of capital. Northvolt—Europe's champion, \$15 billion raised from VW, BMW, Goldman—collapsed into the largest industrial bankruptcy in modern Swedish history (November 2024 Chapter 11; Swedish bankruptcy March 2025) having produced 79.8 MWh in three quarters of 2023, roughly *one five-thousandth* of CATL's annual output [50]. A123 had failed the same way a decade earlier, its assets bought by Wanxiang [51]. By mid-2025, 46 GWh of announced US cell capacity had been cancelled; LGES posted its first quarterly loss in three years and cut capex 30%; Panasonic cut its North American output projections [52, 53]. Grid storage, the next battlefield, is already decided: China deployed $\sim 65\%$ of global BESS GWh in 2025, and storage packs from Chinese LFP oversupply price at \$70/kWh [54, 44].

4.7 Lessons carried forward

Batteries add three lessons to solar's four. *Fifth*: regulatory market reservation (the white list) is more effective than tariffs and invisible to trade law. *Sixth*: vertical integration to the mine means a foreign competitor's factory is a hostage, not a rival—its cathodes, graphite, and equipment are Chinese regardless of its flag. *Seventh*: the dismissed-technology branch (LFP) is where cost-curve wars are won; watch what China industrializes, not what the West benchmarks.

Chapter 5

Electric Vehicles: The System-Level Win

EVs are the capstone precedent because the automobile is not a component but a *system*—the most complex mass-manufactured product in existence, wrapped in brands, dealer networks, safety regulation, and a century of incumbent advantage. If the playbook worked here, no consumer-industrial market is safe from it. It worked here.

5.1 The leapfrog decision (2001–2012)

The strategic insight predates Tesla. Wan Gang—an Audi engineer in Germany through the 1990s—persuaded Beijing around 2000 that China could never catch Toyota and VW in internal combustion, but could *leapfrog* them if the industry changed energy carrier; EV powertrains entered the state 863 R&D program in 2001, and Wan became Minister of Science and Technology in 2007 [55]. The early demand instruments were clumsy and mostly failed—the 2009 “Ten Cities, Thousand Vehicles” pilots missed most targets; fewer than 500 EVs were sold nationally in 2009 [56]. The state persisted anyway: NEVs were designated a Strategic Emerging Industry in 2010, and purchase subsidies, plate privileges, and procurement began compounding.

5.2 The demand machine (2013–2022)

Cumulative support to the sector, 2009–2023, is estimated by CSIS at \$230.9 billion—purchase subsidies, tax exemptions, infrastructure, and R&D, a conservative figure excluding cheap credit and land [57]. The instruments were layered: purchase subsidies (wound down by end-2022, precisely as the market became self-sustaining); purchase-tax exemption extended through 2027; free “green plates” in cities where an ICE plate costs tens of thousands of dollars at auction; and the 2017 dual-credit mandate that forced every manufacturer, foreign JVs included, to produce NEVs or buy credits from those (BYD, Tesla) who did [58]. Beneath it all, infrastructure at a scale the West still does not comprehend: 19.3 million charging points by late 2025, of which 4.6 million public—roughly eighteen times the US public network [59].

The result is the fastest substitution event in automotive history (Figure 5.1): NEV retail penetration went from ~6% (2020) to 27.6% (2022) to 47.9% (2024) to 54.1% full-year 2025, crossing 60% in December 2025 [60]. China has been the world’s largest car market for seventeen consecutive years; in 2023 it passed Japan as the world’s largest auto exporter, and exported 7.1 million vehicles in 2025 (2.6 million NEVs, doubling year-on-year); in June 2026 monthly exports passed one million for the first time, more than half NEVs [61, 62].




Figure 5.1: Two curves, one mechanism. New-energy-vehicle retail penetration in China against the share of the Chinese market held by foreign brands, 2020–2025 [V] [60, 64]. December 2025 monthly NEV penetration passed 60%. The technology transition that reset incumbency and the collapse of incumbency are the same event, viewed from two sides.

5.3 BYD: the vertically integrated apex predator

BYD is the pattern’s exemplar. Battery maker first (Chapter 4), it acquired a moribund automaker in 2003 and spent fifteen years unremarkably; then the flywheel engaged: 416 thousand vehicles in 2020, 4.27 million in 2024, 4.55 million in 2025—taking the global BEV crown from Tesla by over 600,000 units [63]. The weapon is vertical integration: BYD makes approximately 80% of its Tier-1 components in-house—cells, motors, power electronics including IGBT and SiC dies from BYD Semiconductor—avoiding an estimated \$2,369 per vehicle in supplier margins; Rhodium’s teardown accounting puts BYD’s structural cost advantage over Tesla at ~\$4,700 per vehicle, of which direct subsidy is under \$300 [64]. The Seagull retails under \$8,000 domestically. BYD even integrated logistics: an eight-ship ro-ro fleet, flagship capacity 9,200 vehicles, able to export a million cars a year [65]; and factories in Thailand, Brazil, Hungary, and Turkey place final assembly inside tariff walls.

Around the apex predator, an ecosystem of siege: Geely (3.0 million, owner of Volvo, Zeekr, and Lotus), Chery (1.34 million exported in 2025—one vehicle every 23 seconds), SAIC/MG (first Chinese brand past one million cumulative European sales), Xiaomi (289,000 orders for the YU7 in its first hour; 600,000 cumulative deliveries within 22 months of entering the car industry), XPeng, NIO, Li Auto, and Huawei’s HIMA alliance [66, 67, 68]. Behind them, a graveyard: roughly 400 Chinese EV ventures died between 2018 and 2025 [69]. The state subsidizes the species, not the specimen.

5.4 The incumbents’ collapse and the reverse technology flow

Foreign brands’ share of the Chinese market fell from 62% (2020) to 35% (2025). GM, at 4.04 million Chinese sales in 2017, sold 1.8 million in 2024 and took over \$5 billion in China write-downs; VW’s China deliveries fell from 4.2 million (2019) to 2.7 million (2025) [70, 71]. The humiliation is qualitative, not just quantitative—the technology now flows backward: VW paid \$700 million for 5% of XPeng and will license XPeng’s autonomous-driving stack for its China EVs from 2026; Audi builds on a SAIC platform; Stellantis paid €1.5 billion for 20% of Leapmotor and distributes its cars; Toyota’s China BEVs carry BYD powertrains [72]. A century of licensing ran the other way; it took one product generation to reverse it.

Western trade response repeated the solar script, with the same result so far: US 100% tariffs (May 2024) seal a market Chinese makers had barely entered; EU countervailing duties of 17–35%

(October 2024) are already being negotiated down toward a minimum-price regime while BYD’s Hungarian, Turkish, and Thai plants tunnel under them [73, 74].

5.5 The EV as AI system: the bridge to Part II

The EV war matters to this book’s argument for one further reason: the automobile became a software and AI platform in Chinese hands. Chinese suppliers hold $\sim 95\%$ of global ADAS LiDAR shipments (Hesai, RoboSense); city-level navigate-on-autopilot deployed at scale in China from 2024; BYD fitted its “God’s Eye” ADAS across its entire line down to the \$8,000 Seagull at zero additional cost [75, 76]. McKinsey’s 2026 assessment: Chinese vehicles are “frequently more advanced than their Western counterparts” in software and automated driving, developed on 20–24-month cycles against the West’s 40–50, with $\sim 30\%$ lower BOM and capex [77]. In other words: the first mass-market, safety-critical, embedded-AI product category is already Chinese-led—and it is the demand anchor for the domestic AI chips, LiDAR, and models of Part II. Rare-earth motor magnets, where China’s April 2025 export controls idled Ford and Suzuki lines within weeks, previewed the supply-chain leverage now being assembled in AI inputs [78].

5.6 Lessons carried forward

The EV war contributes the final lessons. *Eighth*: pick the technology transition where incumbency resets to zero, and force the transition by regulation at home. *Ninth*: a system-level product won end-to-end pulls its entire component stack to dominance with it (batteries, LiDAR, power semiconductors, magnets)—and the converse: dominance in components makes the next system-level product cheaper to win. *Tenth*: speed is a weapon; a $2\times$ development-cycle advantage compounds every generation. All ten lessons are now being applied to AI.

Chapter 6

Telecommunications: The Network War, in Lesser Detail

The three preceding chapters each took an industry from Western invention to Chinese supply-chain capture and priced every step. This chapter is deliberately shorter, and it is different in kind. Solar modules, battery cells and electric vehicles are *commodities*: the campaign in each case commoditized the artifact and then owned what the artifact had to be made from. Telecommunications equipment is a *network*—sold to a few hundred operators rather than millions of buyers, installed for decades, serviced continuously, and bound to standards that are written in committee rather than set by price. The playbook nevertheless ran here, one company executed almost all of it, and the record is complete enough to score. It belongs in Part I for two reasons. It is the precedent closest in structure to what an “AI stack” export would be—infrastructure, financed, standardized, locked in by service—and it is the case most often invoked, in 2026, to describe what China is now supposed to be doing with sovereign AI. Part III tests that invocation against what has actually been signed. This chapter supplies the template it is being tested against, and states at the end where the template does not transfer.

6.1 The rise, measured

Huawei was founded in 1987 as a reseller of imported switches and had about a fifth of the Chinese switch market by 1996. By 2006 revenue was US\$8.5B, by 2009 US\$21.8B, and by 2011 the company was at parity with Ericsson, the incumbent it had been imitating fifteen years earlier [774] [A]. The global share series is the one that matters, and it is worth printing whole because its shape is the argument: 27.7% of world telecom-equipment revenue in 2018, 28.1% in the first half of 2019, 29% in the second quarter of 2019 against Nokia’s 15.7% and Ericsson’s 13.1%, and 31% in the first half of 2020—the year the sanctions bit hardest [773] [A]. Revenue peaked at RMB 891B in 2020, fell to RMB 636B in 2021 after the Honor divestiture and the loss of the handset business’s chip supply, and recovered to RMB 880.9B (about US\$126B) in 2025 with net profit of RMB 68B and R&D at 21.8% of revenue [772] [P]. The company operates in more than 170 countries and regions by its own count [772].

6.2 The entry pattern: encircle the cities from the countryside

The sequence is the playbook’s target-selection and demand-before-supply moves executed against a market whose incumbents had chosen, rationally, not to serve its margins. Huawei went first where Ericsson and Alcatel would not go: the Chinese countryside, with more than two hundred staff across the counties of a single northeastern province against three or four for the incumbent; then Russia, Thailand, Brazil and South Africa; Ethiopia in 2003; a US\$200M CDMA network

in Nigeria in 2005; the Netherlands and Germany by the middle of the decade; and only then the Tier-1 European operators, with British approval in 2005 as the pivot into the markets that set reference prices for everyone else [774, 778] [A]. The Nortel episode belongs here as a caution rather than a proof: a 2004 intrusion traced to Shanghai removed some 1,488 documents including optical and next-generation wireless research, and Huawei underbid Nortel by about 40% on a 2008 Canadian contract; but a 2014 University of Ottawa study found no evidence that the intrusion caused Nortel’s 2009 bankruptcy, and this book does not assert what the record does not [778] [A].

6.3 The weapon: credit

The instrument that distinguishes this campaign from solar and batteries, and that AI has not yet reproduced, is vendor finance at sovereign scale. China Development Bank opened a US\$10B credit facility for Huawei’s customers in 2004 and raised it to **US\$30B in December 2009**, with a further US\$600M from the Export–Import Bank; roughly US\$10B had been drawn by 2013, and a 2011 European Union inquiry put the facility at about ten times Huawei’s operating cash flow [774, 775] [A]. The *Wall Street Journal*’s reconstruction in 2019 put total state support at up to US\$75B—some US\$46B in loans, credit lines and other financing from state lenders, about US\$25B in tax savings over 2008–18, plus grants and discounted land [775] [A]. Huawei disputed the characterization, and the counter-data must be printed with it: the same reporting has Huawei’s prices about 30% below rivals’ in developing markets, but a 2017 European study found them only slightly lower in Europe and a 2019 study found them about 25% *higher* at home [775]. The discount was where the credit was. An operator in Lagos or Jakarta was not choosing between two switches; it was choosing between a switch it could finance and one it could not, and the financing came from the vendor’s state.

This is the demand-before-supply move of Chapter 7 run with a bank rather than a feed-in tariff, and it is the most exportable instrument in the playbook because it requires nothing of the recipient except a signature. It is also the instrument whose absence, in the AI campaign as of 9 September 2026, Chapter 22 records.

6.4 The endgame: standards from the inside

The consolidation-and-gatekeeping step took a form the earlier campaigns did not need. Rather than build a parallel standards body, Huawei captured position inside the existing one: first in declared 5G patent families and first in 3GPP technical contributions across 2015–24, with Chinese-headquartered firms above 40% of declared 5G families [776] [A]. A vendor with the largest installed base *and* the largest share of the specification is not a supplier the operator can leave at the next generation. That is move M8 of the playbook executed in the only arena where it counts for a network, and it is the move the AI campaign has attempted and not yet achieved—Chapter 11 records Western agent-standards bodies with no Chinese member, and Chapter 12 an open instruction set with no ratified matrix extension.

6.5 The pushback, and what it removed

The bans came in sequence: Australia in August 2018, the US Entity List on 16 May 2019, the EU’s 5G toolbox in January 2020, the United Kingdom’s reversal on 14 July 2020, Sweden that October and, in January 2026, a Commission proposal for mandatory phase-out of high-risk suppliers [777] [P/A]. The cost was real and mostly Western: the British removal was priced at up to £2B and up to three years’ delay; the American “rip and replace” programme had about US\$5B appropriated and 53 of 126 projects complete by June 2026 against a deadline it had already missed [777] [A].

What the bans removed is the finding. Huawei’s global share was 31% in the first half of 2020 and **31% again in the first half of 2025**—21% outside China and 40% outside North America [773] [A]. Chinese vendors’ share of European 5G sites was about 32% in mid-2022 and about 32% at the end of 2024, with Germany at 59%; of twenty-three European states that legislated, seven imposed restrictions [777] [A]. Omdia counts Huawei the leading vendor in three of five world regions, third in Europe, and absent from North America [773]. The campaign failed where the security agencies had a veto and the operators could swap vendors without cost—the Five Eyes, the Nordics, Japan—and nowhere else. And the instrument that actually turned Britain was not a network-security finding but the May 2020 foundry rule: a semiconductor chokepoint, which is to say the Western counter-pattern that Chapter 7 calls chokepoint denial, applied at the one layer where it still held [777].

6.6 What transfers to AI, and what does not

Scored against the six steps, telecommunications is a complete campaign: target selection (the unserved periphery), demand before supply (sovereign credit), overcapacity as strategy (the 30% discount), selection (Huawei over ZTE, over a dozen domestic rivals), commoditization of the incumbents’ rent (equipment margins compressed to the point that Nortel died and Alcatel and Lucent merged to survive), and an endgame of standards position and installed base that the bans could dent but not reverse. On the success–failure model of Chapter 8 it sits in the kill zone on modularity and codification and *outside* it on trust and certification, which is exactly the partition the outcome shows: won everywhere trust was not a procurement criterion, lost everywhere it was.

Three things transfer directly to the AI campaign, and Parts II and III document each. The rural-first entry pattern is the non-aligned adoption channel of Chapter 22: usage arrives before procurement, in markets the incumbents price for someone else. The standards-from-inside endgame is the move the harness and compute chapters record as attempted. And the demonstration that a chokepoint at the silicon layer can reverse a network position built over two decades is the strongest single argument, in this book, for taking the semiconductor chapter’s constraint seriously.

Three things do not transfer, and Chapter 22 tests each against the record at 9 September 2026. *The credit is absent.* Chapter 22 finds no sovereign credit facility for AI infrastructure exports in the record, and the one contracted deployment on Chinese silicon abroad financed by the recipient. *The capacity is absent.* Huawei exported switches because it had surplus factories and a state that wanted them running; in AI the same company is supply-constrained at the accelerator, a single domestic customer’s order absorbs a year of output, and Beijing’s own mandates pull the chips home. *And the artifact does not lock in.* A telecom network is hardware plus service plus debt, and it binds the operator for a generation. Open weights are the opposite object: they cost the recipient nothing and bind the recipient to nothing, and they are the instrument through which Chinese AI actually diffuses—one hundred and fifty thousand derivatives of a single model family, five times the download volume of its nearest Western rival [771] [V]. That is a more effective diffusion mechanism than Huawei ever had and a far weaker rent-capture mechanism, and the difference between the two is the reason this chapter is shorter than the three before it. The telecom war shows what the playbook does with a network. The AI war, so far, is being fought with a commons.

Chapter 7

The Playbook Distilled

Before turning to AI, it is worth compressing the three surveys into the invariant structure. Table 7.1 aligns the moves across the three won wars and their emerging AI equivalents; the telecommunications campaign of Chapter 6 fits every row with one instrument the table does not show, sovereign vendor credit, and is omitted from it only because the table is about commodities; the rest of this short chapter states the mechanism in general form.

Table 7.1: The playbook across four wars. AI-column entries are developed in Part II.

Move (M1–M9)	Solar	Batteries	EVs	AI
M1 State designation	SEI 2010; 12th FYP	white list 2015	863 program; SEI 2010	“AI+” plan; 15th FYP; RISC-V guidance
M2 Home demand creation	national FiT; Top Runner	NEV subsidy linkage	\$231B support; dual credit	state data centers mandated domestic chips; subsidized power
M3 Subsidized scale-up	CDB \$47B credit lines	Big-Fund-style provincial capital	provincial EV funds	Big Fund III (\$47.5B); SMIC/CXMT capex
M4 Darwinian domestic competition	2018 “531” cull	100+ cell makers → consolidation	~400 dead EV firms	large-language-model (LLM) price war since 2024; chip-vendor shakeout beginning
M5 Supply-chain capture	poly → wafer → cell → module	mine → refine → cathode → cell	battery → chip → magnet → ship	model → framework → chip → interconnect → fab → DRAM
M6 Commoditize the rent layer	module \$/W	pack \$/kWh	vehicle BOM	open weights; \$/Mtoken; open ISA/stack
M7 Absorb the “inferior” branch	mono-PERC, then TOPCon	LFP	BEV small cars	sparse MoE; LPDDR capacity vs. HBM bandwidth
M8 Standards capture	efficiency floors	GB safety standards	charging/swap standards	RISC-V profiles; matrix ISA; UnifiedBus; SuperPoD reference
M9 Leverage after victory	OPEC-for-poly	tech export controls 2025	rare-earth controls 2025	(pending)

7.1 The mechanism in general form

Table 7.1 inventories nine *tactical moves* (M1–M9), which recur in different orders across campaigns. The six *steps* below are the strategic sequence those moves serve. References elsewhere in this book name the move or step explicitly to keep the two numberings apart.

Step 1. Target selection. Choose an industry at a technology discontinuity, where the incumbent’s accumulated advantage (ICE drivetrains, PERC lines, HBM-centric clusters) is about to be stranded, and where manufacturing learning-by-doing—not protected IP—is the true accumulating asset.

Step 2. Demand before supply. Use the home market—the largest single market on Earth for nearly everything—as guaranteed demand: subsidies, mandates, procurement reservation. Foreign firms are admitted precisely long enough to transfer capability (JV requirements, white-list exclusions), then outcompeted.

Step 3. Overcapacity as strategy. Fund capacity at multiples of world demand. Western economics reads this as waste; strategically it is the purchase of (a) the entire learning curve, (b) a permanent price umbrella too low for any entrant to survive under, and (c) wartime surge capacity. The losses land on Chinese balance sheets, provincial governments, and dead firms; the cost curve remains.

Step 4. Selection, not planning. Beijing does not pick the winning firm; it floods the sector and lets domestic combat select. Suntech died, LONGi rose; hundreds of EV makers died, BYD rose. The state’s loyalty is to the industrial commons—the engineer pool, the supplier mesh, the tooling ecology.

Step 5. Commoditize the adversary’s rent layer. Wherever the Western business model collects rent—the module, the pack, the badge, the API token, the ISA license—price it at marginal cost. Rents fund the West’s R&D flywheel; destroying the rent destroys the flywheel. Openness (open weights, open ISAs, open specs) is this move in its purest form: the rent goes to zero *by construction*, and adoption—the true currency—accrues to the commoditizer.

Step 6. Endgame: cartelize and gatekeep. After victory, consolidate capacity (anti-involution), stabilize prices upward, and convert dominance into coercive leverage via export controls on the now-Chinese-controlled layer.

7.2 Two Western counter-patterns, both failing

Against this, the West has fielded two responses. *Tariff walls* (solar 2012, EVs 2024) arrive after cost-curve capture and merely tax domestic consumers while the walled market shrinks in relevance. *Chokepoint denial* (semiconductor export controls since 2019) is the more serious strategy and is the direct subject of Part II; its record so far is a pattern this book calls *chokepoint decay*: each denied input (EUV, HBM, H100s, EDA) triggers a state-scale substitution program plus an architectural detour around the input, while the denial itself hands the domestic substitute a guaranteed market—the white-list mechanism, imposed this time by Washington on Beijing’s behalf. Export controls are, functionally, a US-funded market-reservation policy for Huawei, SMIC, CXMT, and Cambricon. Whether the decay outruns the West’s lead is the empirical question of the next four chapters.

Chapter 8

Where the Analogy Breaks

A playbook argued from three victories owes the reader its defeats. The same state, spending comparable money under the same five-year-plan machinery—roughly \$248B per year of industrial policy, 1.7% of GDP, more than any comparison country in absolute or relative terms [A] [240]—has run campaigns for two decades in sectors where dominance has *not* arrived. This chapter surveys them, extracts the variables that separate success from failure, and then applies the resulting model to AI honestly: some of AI sits squarely in the playbook’s kill zone, and some of it carries exactly the properties that have defeated the playbook elsewhere.

8.1 The unfinished campaigns

8.1.1 Commercial aircraft and jet engines

COMAC’s C919 is the anti-BYD. Nineteen years after program launch, deliveries ran ~ 13 aircraft in 2025 against an original plan of 75—roughly two per month—while over 40% of core content, including the LEAP-1C engine ($\sim 30\%$ of aircraft cost), remains Western [V] [241]. The 2025 episode in which Washington suspended and then restored LEAP export licenses demonstrated live chokepoint leverage of exactly the kind that failed in solar [241]. The domestic CJ-1000A engine, in flight test since 2025 with certification targets that independent analysts call years optimistic [P/A], and EASA’s stated three-to-six-year horizon for European certification, bound the escape [P] [242, 243]. Domestic market share: $\sim 7\%$ against a Made-in-China-2025 target of 10% [162].

8.1.2 Lithography, and the deep tool chain

Chapter 13 treated lithography as the hard floor; here it serves as the control case. Against etch at 30–35% localization and cleaning above 50%, lithography stands below 5%; SMEE ships at 90 nm; the one domestic immersion-DUV tool in fab testing resembles ASML’s 2008-generation machine; and a claimed 28 nm breakthrough was publicly retracted by SMEE’s own shareholder [A/P] [244, 157]. Beneath the scanners, the pattern repeats fractally: ArF photoresist below 10% domestic, precision vacuum valves below 10% against VAT’s 60–70% global share, electron microscopes over 90% imported [244, 245]. Twenty years and three Big Funds have not cracked it.

8.1.3 Machine tools, reducers, and the precision economy

High-end five-axis CNC machine tools: $\sim 8\%$ domestic against an 80% target, with Japan and Germany holding the balance [V] [246]. Robot harmonic reducers: Harmonic Drive Systems held $\sim 85\%$ of the world market as late as 2023, Nabtesco $\sim 60\%$ of RV reducers—though here the siege is visibly progressing (Leaderdrive at 15% global; Chinese reducers now inside Tesla’s

and Unitree’s humanoids), and domestic industrial-robot share crossed 50% in 2024 from 18% in 2015 [247, 246]. Precision scientific instruments: high-end domestic share barely 17%, with ~\$17B of annual imports [245].

8.1.4 Operating systems and software ecosystems

The purest network-effects sector shows the playbook’s slowest grind—and its most interesting partial win. Two decades of state Linux (Kylin and descendants) never dented Windows. Yet HarmonyOS, forced into existence by US sanctions exactly as this book’s mechanism predicts, reached 17% of mainland smartphone OS share in 2026—ahead of iOS for six consecutive quarters—with the Android-compatibility bridge burned (HarmonyOS NEXT) and eight million registered developers [V/P] [248]. WPS reached ~600M monthly active users [248]. The lesson cuts both ways: ecosystems yield eventually, but on decade timescales, only under existential forcing, and so far only inside the protected home market.

8.1.5 Biotech: the surprise late bloomer

Pharmaceuticals show what the lagging sectors look like the year they stop lagging. Chinese biopharma out-licensing to Western majors hit \$137.7B across 186 deals in 2025—ten times 2021—with Chinese assets roughly a third of global in-licensing spend; first-in-class pipeline share reached ~24%, second only to the US [V/A] [249]. The residual gap is trust and regulation, not invention: the FDA’s 14–1 rejection of China-only trial data for sintilimab marks the certification barrier in its purest form [249]. Biotech in 2025 looks like batteries in 2015—which is precisely why the trust chapter (Chapter 18) matters for AI.

8.2 The success–failure model

Table 8.1: What separates playbook victories from twenty-year grinds.

Variable	Favors the playbook	Defeats or delays it
Product architecture	Modular, decomposable (cells, modules, packs)	Integral, tightly coupled (engines, scanners)
Clock speed	Fast iteration; learning curve lives on the factory floor	Slow cycles; the frontier barely moves (litho optics) or moves via decade-long certification
Knowledge type	Codified process engineering, scale learning	Tacit craft accumulated in one supply chain (Zeiss, hot-section metallurgy, five-axis)
Regulatory gate	Home market can self-certify (EVs, panels)	Foreign safety/efficacy certification controls access (EASA/FAA, FDA)
Demand lever	State can create/reserve demand at scale	Install base and network effects lock incumbents (OS, Office, engines’ fleets)
Trust requirement	Commodity performance is verifiable at purchase	Buyer must trust decades of provenance (avionics, drugs, <i>models?</i>)

The synthesis across the scholarly assessments—the USCC’s Made-in-China-2025 scorecard, CSIS’s EV-versus-aviation comparison, Chan’s clock-speed argument [A] [162, 250, 251]—is compact: *the playbook wins where products are modular, iteration is fast, learning is manufacturable, the home market can certify, and trust is embodied in the artifact; it grinds where knowledge is tacit, certification is foreign, and value lives in an installed ecosystem.*

8.3 Scoring AI against the model

Now apply Table 8.1 to the three elements, without thumb on scale.

Favoring the playbook: AI models are the most modular, fastest-clock-speed artifact in industrial history—weights are digital goods with zero reproduction cost, the frontier turns over in months, capability is substantially verifiable at the benchmark (with caveats), and the learning curve (data, recipes, kernels) is largely codified and, in China’s case, deliberately published. Sparse-MoE training is closer to batteries than to jet engines. The compute layer is intermediate: accelerators and CPUs are modular chiplet products on fast cycles—LX2 demonstrates exactly that—while the tool chain beneath them (Chapter 13) inherits lithography’s tacit-knowledge resistance, which is why this book’s architecture argument routes around the tool chain rather than through it.

Resisting the playbook: three properties of AI sit in the right-hand column, and they are the strongest structural arguments against this book’s title. First, *the moving frontier*: unlike solar’s fixed physics, the capability target relocates several times a year; a fast follower is never finished, and a discontinuity (recursive AI-for-AI R&D) could reopen the gap faster than any factory closes it. Second, *ecosystem lock*: the enterprise deployment layer—distribution, workflow integration, habits, procurement—is Office, not modules; Chapter 21 takes this seriously as the closed labs’ real moat. Third, *trust and certification*: an AI model deployed in government, finance, or defense is closer to an avionics component or a drug than to a battery—provenance, auditability, and alignment with the buyer’s values are part of the product. The sintilimab precedent—capability accepted, certification denied—is directly available to Western regulators as an AI playbook, and censorship and security findings (Chapter 18) give it factual ammunition.

Two further asymmetries deserve the record. Digital goods cut both ways: zero reproduction cost is what lets Chinese weights conquer adoption instantly—and what would let a single superior Western open release reconquer it just as fast; there is no factory-shaped moat around P1. And the talent flow, the deepest input, currently splits the difference: researchers of Chinese origin constitute roughly half of top-tier AI authorship (47% by undergraduate origin in MacroPolo’s 2022 cohort, from 29% in 2019), the US still hosts the majority of the established cohort (87 of 100 tracked 2019-cohort researchers remained in 2025) though its share of where top-tier researchers *work* fell from 59% to 42% between the two cohorts while China’s rose from 11% to 28%, and the inflow of new Chinese talent to the US is visibly thinning under visa friction while Beijing’s K-visa recruits in the opposite direction [V/A] [252]. Talent is the one input where the two systems still share a supply chain.

The honest conclusion: the model and compute elements sit mostly in the playbook’s kill zone; the tool chain sits in its resistance zone (and is being bypassed rather than assaulted); and the *deployment* layer—trust, certification, ecosystems—is genuinely contested ground where the precedents predict a long grind, not a collapse. That is exactly where the next three chapters, and the strategy chapter, will place their weight.

Part II

The Elements of AI

Chapter 9

The Strategy in Whole: The Elements on the Playbook

Part I distilled a playbook from three campaigns and then, in a chapter of equal weight, marked the terrain on which the playbook has failed. Part II takes the elements of AI one at a time—the model and the harness that has become its contested layer, the compute, the semiconductor, the watt, the data, and physical production—and each chapter closes with its own assessment, before a final chapter (Chapter 17) reports the development that changes what all of them mean: over 2026 the largest Western firms converged on the same method, each commoditizing a layer it does not own in order to defend the layer it does. Before it does, this chapter states the strategy in whole, because the elements are easier to read once the reader knows where each sits: it places every element on the six steps and nine moves of Chapter 7, marks which steps are complete, which are trending, and which have not begun, and states why the consequences of Part III are the ones they are. It is a map drawn before the territory. Nothing in it is new evidence; every claim below is argued and graded in the chapter it points to, and where a figure is drawn from Part III it is cited. A reader who prefers evidence before conclusion may go directly to Chapter 10 and return here afterwards; the map reads the same in either order. It is placed first because the precedent chapters record what happened when each of the earlier campaigns was read one layer at a time and each layer found unthreatening, and this book does not want to invite that reading.

9.1 The strategy in one paragraph

The Chinese AI strategy is the playbook of Chapter 7 run on a fourth industry, with one adaptation that the earlier campaigns did not need and that makes this one faster: the rent layer being attacked is digital, so the commoditizing move—openness—costs nothing to reproduce and propagates at the speed of a download. Target selection chose the discontinuity where the incumbent’s accumulated advantage is about to be stranded (HBM-centric superclusters, closed frontier weights, a licensed instruction set). Demand before supply was arranged by mandate rather than by feed-in tariff (state-funded data centres restricted to domestic silicon; western hub electricity discounted only for approved domestic chips) [117, 118]. Overcapacity as strategy is being funded across accelerators, DRAM, and compute halls at multiples of current demand, with the losses landing where they landed in solar. Selection, not planning, is visibly at work in a model layer where ten laboratories release at frontier scale, and in a chip layer whose dozen-vendor fragmentation the state is now removing by mandate and open specification rather than by cull (Chapter 12). Commoditizing the adversary’s rent layer is the move the whole of Part II documents: open weights, an open harness, an open ISA, open interconnect and framework specifications, at marginal or zero price, aimed precisely at CUDA, NVLink, x86/Arm licensing, and the metered API [99, 101, 106, 131]. And the endgame—consolidation

and gatekeeping—has begun on the consolidation side (the anti-involution campaign [26]) and not yet on the gatekeeping side, where the earlier campaigns supply the template (battery-technology and rare-earth export controls of 2025 [49, 78]). The West’s two counter-patterns, tariff walls and chokepoint denial, are behaving as Chapter 7 said they would: the first is irrelevant to a digital good, and the second is decaying by substitution and by architectural detour.

9.2 Where each element sits

Table 9.1 maps the elements onto the moves, with the status that the chapter on each element establishes. The reading that matters is not any single row but the fact that every row is occupied—that the same nine moves the earlier campaigns used can each be found, in a recognizable form and at a recognizable stage, in the AI campaign.

Three of the rows deserve a sentence each on how they interlock, because the interlock is the strategy. The model and the compute are the coupling that Chapter 1 called $P1 \times P2$ —adoption meeting sovereign, capacity-first hardware: open weights are a commons only for those with the memory to hold them, and a capacity-first hardware class—the commodity-DRAM socket of Section 12.8, the appliance of Section 12.9.4—is what turns a download into a deployment. The watt is the weld: a tariff instrument that pays only for domestic silicon converts an energy surplus into hardware volume and hardware volume into tokens priced below the Western marginal cost, which is the demand-before-supply step of the playbook executed with electricity rather than a feed-in tariff. And the semiconductor is the floor beneath both, incomplete at exactly one layer—HBM—which is why the memory arbitrage exists at all: the architecture was chosen so that the one input the chokepoint strategy still holds is the one the architecture needs least.

9.3 Scoring the campaign against the six steps

Step 1, target selection: complete. The discontinuity was correctly identified and is being exploited on three fronts at once. The book’s own theory chapter (Section 2.3) says why the choice was correct: rent is being relocated, not abolished. China opens the layers where it is behind or where openness buys adoption—weights, ISA, interconnect, framework—and expects to recapture on hardware volume, domestic cloud, standards, and dependence; the United States keeps the layers where it is ahead.

Step 2, demand before supply: complete at home, contested abroad. At home the state has arranged demand by mandate and by tariff. Abroad the demand is real but not yet paying: Chinese models carry roughly 61% of one neutral router’s tokens and a fraction of its revenue [84, 538], enterprise open-weight spend share fell during the year usage rose, and the Global South adoption channel exists as usage before it exists as procurement. This is the one step whose foreign half is where the trust and certification barriers of Chapter 8 bite, and Part III’s opening chapters are about that half.

Step 3, overcapacity as strategy: under way in DRAM and compute halls, and already producing the step’s characteristic pathology—with two differences from the precedents. State-subsidized compute is running at 20–30% utilization against a 60% target, by some reports up to 80% of new capacity idle, and rental prices have collapsed [537, 422]. That is what the step looked like from inside during solar’s 2011–2013 and EVs’ 2018–2025, and the losses land where they landed then. But two things differ. In the accelerator layer the constraint is still supply, not demand—Ascend output is bounded by HBM (Chapter 13)—so overcapacity there has not yet been possible; and Chapter 23 records that the model-layer price war is step 3 running without a visible step-5 rent-recapture plan attached, which is the strategy’s soft underbelly rather than its purchase of a learning curve.

Step 4, selection: working in the model layer, beginning in silicon. The model layer is the purest instance of Darwinian domestic competition the playbook has ever produced—the

Table 9.1: The elements of Part II placed on the playbook of Part I, with the status the element chapters establish. Status as of 9 September 2026: **done** (the move has produced its intended effect), **trending** (visibly under way, effect not yet secured), **contested** (the evidence divides), **pending** (not yet begun or not yet observable). Grades are those of the source chapters.

Element	Playbook move(s) it executes	What the element chapter finds	Status
The model (Ch. 10)	M1 designation (“AI+”, 15th FYP); M6 commoditize the rent layer; M4 Darwinian competition; M7 sparse MoE; M8 standards	Open weights at marginal price; a self-enforcing open-weight ratchet; majority of routed tokens, ~41% of downloads, a 2.7% measured gap; enterprise spend share falling while usage rose; rent relocated to commercial-scale licence terms rather than abolished	Done (distribution); trending (production spend, rent)
The harness (Ch. 11)	M6 one layer up; a de-facto default sought outside the standards bodies (M8 not achieved)	MIT-licensed model-agnostic runtime shipped and adopted in days; Western agent-standards bodies with no Chinese member; lock-in relocating to distribution, evaluation, certification	Trending [S]
The compute (Ch. 12)	M1 (RISC-V guidance); M7 absorb the “inferior” branch; M8 standards; M5 supply chain	Open ISA as national infrastructure; a GPU-free machine at No. 1 on both HPL and HPCG; capacity-first memory arbitrage modeled, not yet realized [M]; a priced appliance that undercuts the fully loaded rack above about 300M tokens a month [M/S]	Trending; existence proof done
The semiconductor (Ch. 13)	M5 supply-chain capture; M3 sub-sidized scale-up	5 nm-class logic without EUV; CXMT the fourth DRAM producer and swing supplier; domestic tools 20–30% of fab demand; HBM three to four years behind the merchant part <i>and</i> outside the co-design relationships that now define the frontier part, and still the single binding physical constraint	Trending; HBM lagging three to four years
The watt (Ch. 14)	M2 home demand (subsidized power, domestic-chip tariff); the weld between elements	Leads on nameplate, generation, firm capacity, deliverable headroom; hub tariffs cut only for domestic-chip facilities [P], converting grid surplus into demand for domestic silicon and cheap tokens; the binding question is whether an administered tariff outlives the surplus	Done (headroom, aiming); trending (durability)
The data (Ch. 15)	Where rent migrates once M6 succeeds	Deployment-generated data is the scarcity openness cannot replicate; China holds the industrial row, the West the enterprise-workflow row; evaluation data unowned	Contested
Embodied AI (Ch. 16)	M5 supply-chain capture before the category exists; the loop back to Part I	~85% of humanoid shipments, ~70% of the Optimus supply chain; deployment surface as public works; useful work per robot-hour published by nobody	Trending
Endgame (Ch. 7 step 6)	M9 leverage after victory	Consolidation begun (anti-involution); gatekeeping not yet exercised on AI; template exists in batteries and rare earths	Pending

state did not anoint DeepSeek; it flooded the sector, and a dozen laboratories now release at frontier scale on a cadence of weeks. The chip-vendor shakeout is starting. The one policy that could damage this step is the state’s own anti-involution consolidation, and it is registered as a falsification branch in Chapter 28.

Step 5, commoditize the rent layer: the step this campaign has run furthest, and the one whose completion is contested. Every proprietary Western layer has an open Chinese counterpart shipped at marginal price. What has not yet happened is the second half of the step—the destruction of the rent that finances the Western frontier. The first direct observation of transmission arrived on 30 July 2026 with an 80% cut to the *cheapest tier* of the closed frontier’s ladder—the flagship held unchanged that day, and reduced only three weeks later, by a fifth, and then only for three months (Section 10.7)—in the week Kimi K3’s weights and DeepSeek V4-Flash shipped [403]; the company attributes the cut to efficiency, and Section 28.2.1 shows that the same datum supports the recursive-improvement counter-argument, so both readings must be printed. One cut is not a trend, closed labs also sell service rents that commoditize slowly, and elasticity may rescue the buildout even under falling prices. Proposition P3 is graded “mechanism operating, magnitude contested,” and Chapter 21 prices what follows if it completes.

Step 6, endgame: begun on one side. Consolidation has started in the model-price war and in compute; gatekeeping has not been exercised on any AI layer. Chapter 22 treats the continuity risk this implies for every country adopting the commons as a risk rather than an accusation, and Chapter 27 prices the insurance.

9.4 Where the analogy tells us to look

Chapter 8 scored AI against the variables that separate the playbook’s victories from its defeats and concluded that the model and compute elements sit mostly in the kill zone—modular, fast-clocked, codified, self-certifiable, benchmark-verifiable—while three properties resist: the moving frontier, ecosystem lock, and trust and certification. The chapters that follow confirm the placement. The model layer behaves exactly as a modular, fast-clocked artifact behaves under the playbook: release lag compressed from months to weeks, adoption followed price, standards migrated to Chinese repositories. The compute layer behaves as an intermediate case: chiplet products and open specifications iterate quickly, while the tool chain beneath them behaves like lithography always has and is routed around rather than conquered. And the three resisting properties are precisely the ones the elements cannot settle from inside—the top US–China gap is 2.7% while the closed-over-open gap has reopened to 3.3%, and either could compound; the enterprise deployment layer is still Office, not modules; and the commons ships at trust Level 0–1 into a world whose paying workloads will require Level 3–4.

That is why, once the elements have been examined, Part III opens where it opens. The deployment layer is where the precedents predict a long grind rather than a collapse, and the deployment layer is made of trust, security, and rules. Chapter 18 treats the exposed flank of openness—supply-chain integrity, behavioural integrity, provenance—and the certification ladder that would close it. Chapter 19 treats the security bill the commons has not paid, which is structural rather than national and falls due at exactly the correlated scale that success implies. Chapter 20 treats the rulebook, which the earlier campaigns’ winners inherited and which is being written now in arenas most Western attention is not directed at. Only after the deployment layer is understood does the book turn to the rent side of the ledger—Chapter 21, which prices what happens to a capital structure long scarcity when step 5 completes—and then to the electorate that will decide the outcome: the non-aligned majority of the world (Chapter 22), the strategies available to every actor class including the presumptive winner (Chapter 23), the scientific prize that dominates the value case (Chapter 24), the compute a nation’s science actually needs (Chapter 26), and the institution that all of those threads terminate in (Chapter 27).

9.5 What the assembly does and does not license

Assembled, the elements license one claim and forbid another. They license the claim that the AI campaign is a coherent execution of the same strategy that won three technology wars, that it is further along than a layer-by-layer reading suggests, and that its remaining gaps—HBM, packaging, tools, software maturity, trust—are being engineered around rather than closed on the West’s terms. They forbid the claim that the outcome is settled. On the evidence the following chapters lay out, one step of six is complete, one is complete at home and contested abroad, three are trending, and the endgame has barely begun; the single integration test that would move P2 from “trending” to “done”—a frontier-class model trained end to end on domestic silicon, packaging, and memory—has not been passed; and the three properties on which the playbook has failed before are the three on which Part III has to do its work. The chapters of Part II now supply the evidence, element by element. One caution belongs here rather than at the end, because it governs how the whole of Part II should be read: the moves catalogued above are not, in 2026, uniquely Chinese. Chapter 17 finds the same instrument in Western hands, aimed by each actor at whichever adjacent layer is taking the rent it wants. The playbook of Part I turns out to be a general theory of a contest over a layered stack rather than a description of Chinese exceptionalism, which strengthens Part I and complicates every score in Table 9.1. Part III then says what follows if the assembly is right, and what would show that it is not.

Chapter 10

The Model: Open Weights as the Commodity Layer

The first element of the AI war is the model itself. The Western frontier is closed: the best US models are accessible only as metered APIs, their weights corporate secrets, their pricing a rent on capability. The Chinese frontier is open: weights downloadable, licenses mostly permissive (MIT, Apache 2.0)—though, as §10.3 details, the newest flagships have begun attaching custom commercial terms that are themselves evidence for this book’s rent-relocation thesis—and API prices that until recently sat at or near marginal inference cost. That last proposition is the one this chapter has to qualify hardest, because the price sheets of July and August 2026 stopped obeying it: the American incumbent’s cheapest tier now undercuts most of the Chinese frontier on list price, while the Chinese price leader has *raised* prices by as much as elevenfold. Section 10.7 works through what that inversion does and does not mean. This chapter surveys the whole picture—the labs, the releases, the algorithmic innovations, the quantization ecosystem that is collapsing the last barrier between frontier capability and commodity hardware, the price line at which the two ecosystems now meet, and—new at this currency date—the licence clause by which the makers of the commons have begun to negotiate a share of the American clouds’ revenue from serving it (Section 10.4).

10.1 The DeepSeek shock and the price war (2024–2025)

The opening move replicated the solar price collapse, compressed into eighteen months. DeepSeek—a 2023 spin-out of the High-Flyer quantitative fund—released V2 in May 2024 with Multi-head Latent Attention and a fine-grained MoE at roughly RMB 1–2 per million tokens; within two weeks ByteDance, Zhipu, Alibaba, and Baidu had cut prices by up to 97%, Baidu making its mainline models free within four hours of Alibaba’s cut [79]. V3 (December 2024)—671B parameters, 37B active, and the first open large-scale model natively trained in FP8—reached GPT-4-class capability with a published marginal training cost of \$5.6 million (2.79 million H800-hours), a figure legitimately contested as excluding R&D and infrastructure, but directionally transformative either way [80]. R1 (January 2025) matched OpenAI’s o1-class reasoning, was released under MIT license, topped the US iOS App Store, and erased \$589 billion of NVIDIA’s market capitalization in a single trading day—the largest one-day single-company loss in US market history [81]. The peer-reviewed Nature version disclosed that R1’s reinforcement-learning stage cost about \$294,000 [82].

The strategic content of the shock was mispriced by markets and correctly priced by Beijing: it demonstrated that frontier capability could be commoditized, that export-control-constrained hardware forced *algorithmic* efficiency (MLA, extreme sparsity, FP8 training, GRPO, multi-token prediction) which then compounds on any hardware, and that open weights convert capability into global adoption instantly.

A cost-accounting discipline applies to every training-cost figure in this book. A “\$5.6M training run” is final-run marginal compute [P]; the full program cost adds failed and ablation runs, post-training and RL, synthetic-data inference, data acquisition and cleaning, personnel, hardware depreciation, networking, power, and the prior research that produced the reusable architecture. DeepSeek’s own disclosed hardware position ($\sim 50,000$ H800s; buildout estimates to \$1.6B [A]) bounds the true denominator [80]. The strategic point survives the correction—marginal-cost transparency itself resets buyer expectations—but the book does not treat run cost as program cost.

10.2 An innovation genealogy

A release chronology shows cadence but not authorship, and the strategically important question is which architectural innovations are Chinese-origin, which were independently reproduced, and—the question this book cares about most—which *reduce a hardware requirement*. Table 10.1 attempts the genealogy.

Table 10.1: Innovation genealogy of the open frontier. “Hardware effect” names the resource the technique economizes—the column that connects this chapter to Chapter 12.

Technique	First at scale	Origin / diffusion	Hardware effect
Multi-head latent attention (MLA)	DeepSeek V2, 2024-05	Chinese-origin; widely adopted	KV cache to $\sim 2\%$ of dense: <i>memory capacity and bandwidth</i>
Fine-grained + shared-expert MoE	DeepSeek V2/V3	Chinese-origin refinement of a Western idea	activated fraction to 1.8–3%: <i>FLOPs and bandwidth</i>
Native FP8 pretraining	DeepSeek V3, 2024-12	first open large-scale demonstration	<i>memory and compute</i> ; sets the domestic-chip numeric format
GRPO and critic-free RL	DeepSeek R1, 2025-01	Chinese-origin; broadly reproduced	removes critic model: <i>post-training compute</i>
Multi-token prediction	DeepSeek V3	Chinese-origin at scale	<i>tokens per step</i> ; speculative decode
Sparse attention (DSA)	DeepSeek V3.2, 2025-09	Chinese-origin	$O(L^2) \rightarrow O(Lk)$: <i>long-context compute</i> —the decisive lever for CPU-class prefill
INT4 / MXFP4 QAT release formats	Kimi K2 Thinking 2025-11; K3 2026-07	Chinese-origin at frontier scale; gpt-oss parallel	<i>serving memory and bandwidth</i> ; 4-bit as release format
Hybrid linear attention at frontier scale	Kimi K3, 2026-07; GLM-5.3-Flash, 2026-08	Chinese-origin	<i>long-context memory traffic</i> ; Zhipu claims $\sim 4.4\times$ KV-cache reduction against its own flagship [P, vendor]
Expert-parallel dispatch/communication overlap	DeepSeek V3 infra	Chinese-origin engineering	<i>interconnect</i> : tolerates weaker fabric

Read down the right-hand column, but read it carefully, because it is easy to overstate what this table shows, in both its arithmetic and its causality. The arithmetic first: it is *not* true that almost none of these techniques economizes dense FLOPs—the table itself refutes it. Fine-grained MoE reduces active FLOPs directly; FP8 reduces compute cost as well as memory; critic-free RL removes an entire model’s worth of post-training compute; multi-token prediction and speculative decoding change compute utilization per emitted token. The accurate reading is one of *emphasis*, not exclusion: the portfolio is concentrated unusually strongly on memory capacity, memory bandwidth, interconnect efficiency, and low-precision execution—the resources that became disproportionately scarce under export control—while the compute savings it also delivers are the kind any laboratory, constrained or not, would pursue.

The causality second. Efficiency is valuable to everyone; MoE, low precision, sparse attention, and communication overlap were all being advanced by unconstrained Western laboratories in the same period, and most would likely have progressed without any export controls at all. The

defensible claim is therefore *constraint-consistency*, not unique causation: export controls are a plausible accelerant and a selection pressure that shaped which efficiencies were prioritized and how aggressively, but not the sole origin of the techniques. Table 10.2 states the questions that separate strong from weak co-design evidence, and applies them.

Table 10.2: A causality matrix for the genealogy: what would make “co-design under sanction” strong or weak for a given technique, with the two clearest verdicts from Table 10.1.

Diagnostic question	Interpretive value	Verdict on examples
Did the technique predate the relevant restriction?	Weakens a sanctions-origin claim	MoE, low precision: yes—weak
Was it independently pursued by unconstrained labs?	Shows general value, not uniquely Chinese adaptation	MoE, FP8, sparse attention: yes—weak
Does it specifically match a domestic hardware limitation?	Strengthens co-design attribution	MLA, dispatch/comm overlap: yes—strong
Is there an explicit domestic numeric-format, compiler, or processor target?	Direct evidence of hardware–algorithm co-design	UE8M0 FP8 “for next-generation domestic chips”: yes—strongest
Would it remain commercially attractive without sanctions?	Separates national adaptation from ordinary cost reduction	nearly all: yes—caps every claim above

On this matrix, UE8M0 FP8 with a declared domestic-silicon target is strong co-design evidence; generic MoE adoption is weak; MLA and the interconnect-tolerance engineering sit in between, matching domestic limitations specifically enough to carry weight. The genealogy’s strategic conclusion, stated at the strength the matrix actually supports: the *direction* of Chinese architectural emphasis is exactly the direction that makes the hardware of Chapter 12 viable, and export controls demonstrably sharpened that emphasis even where they did not originate it.

10.3 The ecosystem survey: eight frontiers, one commons

Consistent with playbook step 4 (selection, not planning; move M4), the state did not anoint DeepSeek; it flooded the sector. Table 10.3 compresses the release cadence of the major laboratories; the notes that follow give each lab’s trajectory. Two structural observations frame the table. First, the cadence itself is the weapon: the interval between a US closed-frontier release and its Chinese open-weight counterpart compressed from six–eight months in 2024 to “a matter of weeks” by mid-2026 [95]. Second—stated in its properly qualified form—open-weight release has become the dominant international distribution strategy of the leading Chinese laboratories, *although license terms range from permissive Apache/MIT grants to custom commercial licenses that preserve control over high-revenue deployment*. The table therefore carries a license column, because the license is strategic data, not legal boilerplate: Kimi K3, the largest open-weight model ever released, ships under a custom license that permits use, modification, redistribution, and derivatives but requires a separate commercial agreement from qualifying model-as-a-service businesses above a revenue threshold, with attribution conditions on very large deployments [463]. That is this book’s own industrial-organization thesis enacted in a license file: opening the weights does not abolish rent—it relocates it to large-scale commercial service, where the releasing lab retains a claim. By 9 September 2026 the pattern has three frontier-scale instances with three different thresholds and two different instruments, the newest of which substitutes a security review for a fee (§10.3.3).

Two terminological disciplines apply throughout. These are *open-weight* releases, not open-source: with few exceptions the training code, data, and full recipes are not published, which is why R1 spawned an entire replication literature [253]. And the correct narrow claim is that leading Chinese laboratories have made open-weight release the dominant *international distribution strategy* for their frontier-adjacent families—not that every Chinese frontier system

is open (several are not), nor that the West releases nothing (gpt-oss, Gemma, OLMo, and NVIDIA’s Nemotron line are real, if outclassed at the top end).

Two disciplines govern this table. *First, a currency date:* every fact in the stable chapters is current as of 9 September 2026—the date of this edition, printed on its title page—and nothing learned after that date is retrofitted into the analysis. Freezing the data at a date *earlier* than the edition invites exactly the contradiction a reader would expect—a declared July cutoff sitting a few pages from an August release—so the rule adopted here is the only one that cannot contradict itself: *the cutoff is the edition date*. Correspondingly, Appendix H carries the most recent cycle in detail so that the fastest-moving material sits in one replaceable place rather than being threaded through the argument. Applying the rule, the table above runs through the release of GLM-5.3’s weights on 28 August and the reveal, the day before, that the stealth model “Ox Alpha” was GLM-5.3-Flash [711, 712]; Qwen3.8-Max sits at its 3 August announcement [393, 396], and DeepSeek’s V4-Flash where its actual 31 July release date puts it [401]. No analytical claim in this book depends on which model holds the crown in any given week; the three-artifact structure—stable monograph, machine-readable claim register, living appendix—exists precisely so that the argument does not chase the cadence it describes. *Second, exact naming:* the open Qwen3.6 releases are the 27B dense and 35B-A3B variants under Apache 2.0 [464]—“Qwen3.6” is not one frontier-scale flagship, and the table says so.

Second, the breadth of the table is itself the datum. The rows are not the work of one national champion and a few followers. DeepSeek, Alibaba, Moonshot, Zhipu, MiniMax, Baidu, Tencent and Huawei are expected entrants; the last two rows added in 2026 are not. Xiaomi—a handset and appliance maker—open-sourced a 310B/15B omni-modal model and a 1.02-trillion-parameter text model under MIT in April [667] [V]. Xiaohongshu—a lifestyle and commerce platform of more than 400 million monthly users, preparing a Hong Kong listing at a reported valuation above \$70B—released *dots3-note Preview* on 14 August under Apache 2.0: 280B total and 16B active parameters, text, image, video and audio input, a 512K context, and a reinforcement-learning method (TEMPO, Chapter 11) built for agent runs of ten hours or more; the same laboratory’s dots-note-3.0 had scored a perfect 42/42 at the 2026 International Mathematical Olympiad three weeks earlier, and Huawei’s Ascend stack carried the open model on day zero [662, 664, 665, 666] [V/A]. What that pair of rows records is that frontier-scale model development in China has become a capability a large technology platform is *expected* to have, rather than the specialty of a handful of laboratories—which is the industrial-base observation of Part I arriving at the model layer, and a stronger fact about the ecosystem than any single release in the table.

10.3.1 The open-weight ratchet: why the commons is now an equilibrium

The Qwen3.8-Max release deserves more than a chronology entry, because it converts this book’s central claim from a description of state strategy into something considerably stronger: a description of a *competitive equilibrium that no participant can profitably defect from*.

Start with the fact pattern. Across the Qwen family’s history the numbered open releases (7B, 72B, the Qwen3 family, the 3.6 dense and A3B variants) shipped under Apache 2.0 and drove the download statistics that make Qwen the most-adopted open family on earth—roughly a billion cumulative downloads by March 2026, 153.6 million in February alone, more than half of all global open-model downloads, and over 200,000 derivative variants [406]. The *Max* tier, meanwhile—Qwen2.5-Max, Qwen3.5-Omni, Qwen3.6-Plus, Qwen3.7-Max—had been consistently proprietary: API and chatbot only. That was the settled pattern, and a reasonable observer in June 2026 would have predicted it continuing.

It did not continue. On 27 July, Moonshot open-sourced Kimi K3’s weights at 2.8 trillion parameters. Seven days later Alibaba announced Qwen3.8-Max and committed, publicly and unambiguously, on its own official account: “*Next week, the open weights of Qwen3.8-Max will be released, and Qwen3.8-27B is also going open-weights*” [393]. The first Max-tier model in the family’s history to carry an open-weight commitment carried it within a week of a rival opening

Table 10.3: Chinese open-weight frontier releases, 2024–August 2026 (selected; parameters as total/active where MoE; benchmark figures are vendor-reported [P] unless noted). License classes: Apache = Apache 2.0; MIT; custom = custom commercial license (weights downloadable, high-revenue service deployment separately controlled). “AA index” is the Artificial Analysis Intelligence Index, always with its version: v4.1.1 is the launch-week board of August 2026, v4.3 the re-based board of 7 September [758]; scores are not comparable across versions. Current as of 9 September 2026 (the edition date); the most recent cycle is analyzed in Appendix H.

Date	Lab	Model	License	Significance
2024-05	DeepSeek	V2 (236B/21B)	MIT-class	MLA + fine-grained MoE; triggers national LLM price war
2024-12	DeepSeek	V3 (671B/37B)	MIT-class	native FP8 pretraining; \$5.6M marginal cost claim
2025-01	DeepSeek	R1 (671B/37B)	MIT	o1-class reasoning; the NVIDIA crash
2025-04	Alibaba	Qwen3 (0.6B–235B/22B)	Apache	family of eight, 36T tokens, hybrid reasoning
2025-06	Baidu	ERNIE 4.5 (to 424B)	Apache	wholesale open-sourcing of a former closed flagship
2025-07	Moonshot	Kimi K2 (1T/32B)	modified MIT	largest open model to date; agentic focus
2025-08	DeepSeek	V3.1	MIT	hybrid think/non-think; UE8M0 FP8 “for next-generation domestic chips”
2025-09	DeepSeek	V3.2 (DSA)	MIT	sparse attention: $O(L \cdot k)$; API cut >50% again
2025-11	Moonshot	K2 Thinking (1T/32B)	modified MIT	native INT4 QAT; beats GPT-5 on HLE (44.9 vs 41.7)
2026-02	Zhipu	GLM-5 (744B/40B)	MIT	SWE-bench Verified 77.8%; Ascend-only training <i>reported</i> in the press (Jun. 2026), no primary located [A/C]
2026-02	MiniMax	M2.5 (230B/10B)	custom	#1 Multi-SWE-Bench; \$0.15/\$1.20 per Mtok
2026-04	DeepSeek	V4 (1.6T/49B)	MIT	1M context; $\sim 150\times$ cheaper than GPT-5.5 output
2026-04	Alibaba	Qwen3.6 (27B dense; 35B-A3B)	Apache	Qwen crosses 1B cumulative downloads
2026-04	Xiaomi	MiMo-V2.5 (310B/15B); V2.5-Pro (1.02T/42B)	MIT	omni-modal input, 1M context; a handset maker ships a trillion-parameter open model [667]
2026-06	MiniMax	M3 (~ 428 B/23B)	custom	multimodal; 1M context; AA index 44 at release (June board, pre-re-basing)
2026-06	Zhipu	GLM-5.2 (~ 753 B/40B)	MIT	#1 open model on the June Artificial Analysis board
2026-07	Moonshot	Kimi (2.8T/ ~ 104 B)	K3 custom	largest open model ever; MXFP4; global top-5 overall
2026-07	DeepSeek	V4-Flash (~ 280 B/13B)	MIT	weights shipped; 1M context; \$0.14/\$0.28 per Mtok; gains from post-training alone
2026-08	Alibaba	Qwen3.8-Max (2.4T/ ~ 95 B)	custom	weights shipped ; \$50M revenue-threshold licence; text-only; 262K native context (§10.3.1)
2026-08	Alibaba	Qwen3.8-27B	Apache	vision-language companion; the permissive half of the release
2026-08	Zhipu	GLM-5.3 (~ 744 B/40B)	custom	weights released 28 Aug.; glm-5.3 licence, reported security review above \$10B MaaS revenue; AA index 45 on v4.3 (60, #2 of 111, at launch on v4.1.1) (§10.3.3)
2026-08	Zhipu	GLM-5.3-Flash (320B/18B)	MIT	the stealth model “Ox Alpha”; natively multimodal, 1M context; AA index 42 on v4.3 at \$0.25 per task (57, #4, at launch on v4.1.1); <i>served</i> on undisclosed domestic silicon (§10.3.3)
2026-08	Meta	Muse Glimmer (30B dense)	Apache	Meta’s return to open weights after 16 months; open Spark 1.2 promised (§10.3.2)
2026-08	Xiaohongshu	dots3-note (280B/16B)	Preview Apache	text/image/video/audio input, 512K context; TEMPO long-horizon RL for agents; a social platform ships an agent model [662]

at comparable scale—alongside a published blog with the full benchmark table, and a complete project trace on GitHub for the flagship autonomy demonstration [394].

The causal reading is hard to avoid, and it is the most important structural observation in this chapter. *Once one Chinese laboratory opens at frontier scale, the others cannot stay closed.* The mechanism is not ideology and not state compulsion; it is the arithmetic of the strategy every Chinese lab is running. Their competitive position is built on adoption—downloads, derivatives, developer mindshare, standards gravity, and the deployment volume that Chapter 2’s flywheel converts into learning and cost advantage. Against a rival whose frontier-scale weights are downloadable, a metered-only flagship forfeits precisely the currency the strategy is denominated in. Whatever premium the closed tier collects is smaller than the adoption it surrenders, because the adoption is what the whole apparatus—hardware demand, standards position, national ecosystem—was built to capture. Closing your best model, in that ecosystem, is unilateral disarmament.

This is a *ratchet*, and its properties matter. Each open release at a new capability level raises the floor every competitor must match; the floor never falls, because no lab gains by being the one that retreats behind an API while its rivals ship weights. Defection is individually irrational even when it would be collectively profitable—the exact inverse of the Western equilibrium, where each frontier laboratory’s rivalry is over premium capability sold at metered access, and closing is individually rational for each. Two ecosystems, two stable equilibria, each internally coherent, each self-enforcing without a planner. The Chinese one produces a commons; the American one produces a rent.

Three consequences follow. *First*, the thesis is more robust than a policy-grounded account of openness would be. Grounding Chinese openness in the Five-Year Plan, MIIT guidance, and WAICO—the natural reading, and the common one—invites the obvious rebuttal that policy can reverse. The ratchet does not depend on Beijing’s continued preference. Even if the state’s enthusiasm cooled tomorrow, a Chinese lab that closed its frontier tier would hand the developer base, the derivative ecosystem, and the standards position to whichever rival stayed open. *Second*, it predicts the pattern forward: the next frontier-scale Chinese release will be open, and the one after that, not because anyone decreed it but because the first defector loses—a prediction the next release tested within the month, when Zhipu published GLM-5.3’s weights on 28 August under a custom licence and its distilled companion under MIT (§10.3.3). *Third*, it refines rather than contradicts the rent-relocation argument of Section 2.3: rent still relocates—to licence terms (Kimi K3’s revenue threshold), to hosted service, to hardware, to the complements of Table 2.2—but it relocates *out of the weights*, and the ratchet is why it cannot relocate back.

The register carries the near-term test, and it has now resolved—in two directions at once, which is why it was registered as two forecasts rather than one. *The weights shipped.* Within roughly ten days of the announcement, Alibaba published the 2.4-trillion-parameter model to Hugging Face—as Qwen3.8-2.4T-A95B, not under the product name “Max”—at 4.89 TB across 213 BF16 shards, with an FP8 variant, followed on 15 August by Qwen3.8-27B [484, 485]. The 90% forecast resolves **YES**. *The licence is not permissive.* The flagship carries a custom qwen3.8-max licence requiring a separately negotiated commercial agreement from any licensee whose model-as-a-service or AI-work-assistant business exceeds **US\$50 million of revenue in any consecutive twelve months**, with an additional attribution obligation above 100 million monthly active users or US\$20 million of monthly revenue; internal use whose outputs are not shared with third parties is unrestricted [484]. Only the 27B companion is Apache 2.0. The 55%→30% permissive-licence forecast therefore resolves **NO** on the model that matters, and the correction is worth printing loudly because much of the trade press reported the release as uniformly Apache-licensed and was wrong [485].

Two further precisions the record requires. The Max-tier model is *text-only*, not multimodal as the pre-release press consensus had it; the 27B is the vision-language variant. And the licence contains *no geographic restriction*—the US, EU, UK and South Korea exclusion that circulated in

early August belongs to MiniMax’s H3 video model, a different company and a different artifact, and conflating them would misstate both.

So the ratchet holds on weights and fails on terms, which is precisely the discrimination the two forecasts were built to make. The mechanism compels *weights onto the commons*, because adoption is the currency; it does not compel Apache terms, because rent relocation operates on exactly that margin. What Alibaba has built is a commons with a toll gate at commercial scale: free to researchers, to enterprises deploying internally, and to every derivative-builder below fifty million dollars—which is to say free to essentially the entire population that generates ecosystem gravity—and chargeable precisely where a rival could resell it as a service. Kimi K3 set the pattern with a twenty-million-dollar threshold; Alibaba raised the number and kept the structure. This is not openness retreating. It is openness that has learned where its own rent lives, and Section 2.3’s rent-relocation table predicted the destination; Section 10.4 follows the clause to the counterparty it has now found.

10.3.2 The ratchet crosses the Pacific: Meta’s return to open weights

On 10 August 2026 the ratchet acquired its first American datapoint. Meta released the weights of *Muse Glimmer*—a 30B-parameter dense multimodal model with a $\sim 128\text{K}$ context, distilled from its closed flagship and engineered to run on a single consumer GPU—under Apache 2.0, the most permissive licence Meta has ever attached to a frontier-family artifact, and a deliberate break from the custom community licences of the Llama era [473, 475]. Zuckerberg announced the release personally on X and, in a 6,500-word letter the same day, recommitted Meta to “American open source”—with a companion promise that an open-weight *version* of the flagship Muse Spark 1.2 would follow “in the coming weeks” [474]. The sequence behind the reversal is worth restating precisely, because it is the argument of this chapter run in mirror image: Llama collapsed below 1% of routed open-model volume; Meta shelved its Behemoth flagship, pivoted to closed development through late 2025, and shipped Muse Spark closed in April 2026; and then, sixteen months after its last open release—with Chinese families at roughly 61% of routed tokens and the K3/V4-Flash/Qwen3.8 barrage dominating the commons—it returned to open weights [473, 84].

Does this accelerate the Chinese position or begin to reverse it? Both readings deserve their strongest form, and the evidence to date assigns them different time horizons.

The re-anchoring case. An Apache-licensed American model at genuine quality, backed by Meta’s distribution and compute, gives Western enterprises and governments what the trust chapter (Chapter 18) argues they lack: a way to consume open weights without consuming Chinese provenance. If the promised open Spark 1.2 ships at competitive frontier quality, the commons stops being singularly Chinese—and with it the monoculture argument of Chapter 19 weakens, the certification problem eases for Western deployers, and the derivative ecosystem acquires a second pole. Zuckerberg’s own framing is explicitly this: American open source must lead, and US policy should reduce the friction preventing it [474]. For the non-aligned world of Chapter 22, a two-pole commons is strictly better than a one-pole commons, and this book has argued throughout that a contested commons is the healthiest available outcome.

The acceleration case. Three facts cut the other way, and at present they cut deeper. First, *what is actually open*: the small distilled model shipped; the flagship is a promise, of an unspecified “version,” under an unspecified licence—precisely the pattern this chapter documented for Alibaba’s Max tier, but starting from sixteen months further behind in derivative ecosystems, fine-tune tooling, and deployed mindshare. Second, *measured quality*: on the independent Artificial Analysis index as it stood in August (v4.1.1, before the September re-basings) Glimmer scored 35—above Gemma 4 31B at 30, below Qwen3.6 27B Reasoning at 38 and Kimi K2.5 at 36—with an 82% measured hallucination rate against 49% for the comparable Qwen [475]. A rule of method applies to that number and to every index number that follows: Artificial Analysis re-based the index twice in the week before 9 September 2026, on 4 September (v4.2)

and 7 September (v4.3), and the two re-basings moved every model’s score by ten to fifteen points and re-ordered the top, so this book treats *rank within a version* as informative and treats any comparison of scores across versions as meaningless, and prints the version beside every score [758]. The entry is competitive, not commanding; it raises the floor of the commons without capturing it. Third, and structurally most important, *the ratchet interpretation*: Meta’s return is what the mechanism of Section 10.3.1 predicts any lab whose strategy depends on adoption will eventually do. It is therefore *evidence for the ratchet as a mechanism*—competitive compulsion, not policy preference, now drives openness on both sides of the Pacific—while leaving entirely open who owns the commons the ratchet enlarges. A commons contested by an American entrant validates the terrain Chinese strategy chose; it vindicates the fast-follower bet that value would migrate out of the weights; and every argument of Chapter 2 about rent relocation now applies to Meta too, which must monetize adjacent to weights it gives away, against incumbents seven releases into that same discipline.

The falsifiable statement, registered in Appendix J: the re-anchoring case requires the open Spark 1.2 to ship within the committed window, at flagship-representative quality, under terms permissive enough to seed derivatives—three conditions, each observable within a quarter. If they hold, the openness scoreboard genuinely re-opens, and later editions of this book will say so. If the pattern instead repeats the Max-tier template—small model open, flagship deferred, terms restrictive—then August 2026 will read as the week the American ecosystem conceded that the Chinese strategy was correct and began running it from behind. Either way, one proposition was strengthened in August: *nobody now argues the commons does not matter*. The contest has moved from whether to open to who opens more, faster, at higher quality—which is the contest the convergence strategy was built to win, on terrain where the incumbent advantage of distribution matters less than the compounding advantages of cadence and cost.

10.3.3 GLM-5.3, “Ox Alpha,” and the serving claim

The last fortnight before the currency date supplied the sharpest available test of three propositions this chapter has been building—the ratchet, the deployable frontier, and the sovereignty of the stack underneath—and it resolved them in three different directions, which is why it is worth taking slowly.

The ratchet held on weights and moved again on terms. Zhipu announced GLM-5.3 on 13–14 August, opened the API on 18 August, and published the weights on 28 August: ~744B total parameters with ~40B active, 78 layers, 256 routed experts with 8 active per token, a one-million-token context [711] [P]. The weights shipped, as §10.3.1’s mechanism predicts they must. The licence, however, is neither MIT nor Apache 2.0 but a custom *glm-5.3* instrument whose reported terms require a security review of model-as-a-service operators above \$10 billion of trailing revenue [711] [P/A on the reported threshold]. That is the third frontier-scale instance of the same structure and the most revealing of the three, because it changes the instrument as well as the number: Moonshot set a \$20 million revenue threshold and Alibaba a \$50 million one, both requiring a negotiated *commercial* agreement, while Zhipu sets a threshold two orders of magnitude higher and asks not for money but for *review*. Rent relocated to licence terms in the first two cases; in the third, what is retained at commercial scale is a governance hook. Either way the direction is the one Section 2.3 predicted—out of the weights, into the terms of large-scale service—and the population below the threshold, which is to say every researcher, every enterprise deploying internally, and every derivative-builder on Earth, gets the artifact free.

The scoreboard moved, and the cheap model moved it further than the flagship. On 26 August Zhipu released GLM-5.3-Flash, and on 27 August confirmed it had been running publicly as the unattributed stealth model “Ox Alpha”: 320B total parameters with 18B active, hybrid sparse and linear attention, the first natively multimodal member of the GLM-5 series, trained on a 30-trillion-token multimodal corpus, a one-million-token context—and released under **MIT**, at \$0.15/\$0.50 per million tokens, roughly 9× below the flagship [712] [P]. In its launch week, on

the Artificial Analysis index then current (v4.1.1), GLM-5.3 scored 60, second of 111 evaluated models, and GLM-5.3-Flash scored 57, fourth, at \$0.09 of cost per task against the flagship's \$1.40, with the best closed model, Claude Opus 5 at maximum reasoning effort, at 63 [713] [A]. Those are the numbers the trade press carried and the numbers the shares moved on, and they are also numbers that no longer exist: the re-based v4.3 board of 7 September puts Claude Fable 5.1 and GPT-6 Astra (both at maximum effort) at 53, Claude Opus 5 at 51, Muse Spark 1.3 at 48, GPT-5.6 Sol at 47, **GLM-5.3 at 45, GLM-5.3-Flash at 42**, Qwen3.8-2.4T-A95B at 40, GPT-5.6 Luna at 38, DeepSeek V4 Pro at 36 and DeepSeek V4 Flash at 35 [758] [A]. Under the rule of method stated in §10.3.2, the launch-week rank is informative and the launch-week score is not, and the gaps this book prints are the v4.3 ones. Two readings follow, and the second is the one that matters. The open-closed gap at the very top is *eight index points* on v4.3—three on the launch-week board, which is the figure that circulated, and the difference is entirely the instrument: the re-basing replaced Terminal-Bench 2.1 with v4.0, swapped in AutomationBench-AA, dropped GPQA Diamond and raised the private-set weight, and the top of the board re-ordered accordingly. Eight is consistent in direction with the Stanford Index's 2.7% but is measured on a different instrument with a different scale, so the two numbers corroborate a shape rather than confirming each other, and neither should be read into the other. And the distilled tier is still the interesting one, though the claim has to be made more carefully than the launch-week board invited: in its launch week the 320B/18B model tied GPT-5.6 Sol at high reasoning effort and beat Gemini 3.7 Flash; on v4.3 it sits eleven points below the closed leaders and five below Sol, but above Qwen3.8's 2.4-trillion-parameter flagship, above GPT-5.6 Luna and above both DeepSeek V4 tiers—an artifact less than half the size of its own flagship, three points behind it, at \$0.25 per index task against the flagship's \$2.01 [758, 759] [A]. One caveat belongs in print: the launch-week Flash score circulated in Z.ai's own materials before it appeared on the public board, which is a provenance irregularity rather than a refutation, and this book prints the number with that history attached [713]. Adoption followed at the speed the ecosystem now runs at—346,516 Hugging Face downloads in the first month, 62 quantized derivatives, a listing on the Dell Enterprise Hub—and Zhipu's Hong Kong-listed shares rose sharply on the reveal [712, 715] [V]. What that 320B/18B model does to the hardware requirement is the subject of §10.6, and it is the single most consequential measurement in this chapter.

What the small model is actually good at, stated exactly. The composite index hides the shape of the result, and the independent v4.3 sub-scores are more useful than the headline. On Terminal-Bench v4.0—the agentic-coding component whose introduction drove the re-basing—GLM-5.3-Flash scores **33%**, against 12% for GPT-5.6 Luna and 12% for DeepSeek V4 Flash, the two models it competes with on price, and 42% for its own flagship; on AutomationBench-AA it scores 60% against Luna's 50%. On everything that is not agentic coding it loses to the same cheap closed model: 15% against 24% on GDP.pdf, 15% against 21% on CritPt, 61–71 output tokens per second against Luna's 120, 34.5 seconds to the first answer token, and a cost per index task of \$0.25 against Luna's \$0.18 [759] [A]. The precise finding is therefore not that a distilled Chinese model has caught the closed volume tier; it is that GLM-5.3-Flash is *the cheapest point on the price-agentic-coding frontier*, and not a general replacement for the flagship tier—an 18B-active model that does one economically important thing at a third of the cost of anything that does it better, and that is slower, worse-calibrated and less knowledgeable than the closed model priced beside it at everything else. The vendor's own card claims considerably more, 84.3 on Terminal-Bench 2.1 and 63.4 on DeepSWE v1.1, and neither figure has been replicated by anyone; no SWE-bench Verified, SWE-bench Pro, LiveCodeBench or Aider result for the model exists from any source, vendor or independent, so the coding claim rests on two independent agentic sub-scores and nothing else, and this book says so rather than filling the gap with the model card [759] [A/C on the vendor figures]. The hardware statement that belongs beside the capability statement is equally narrow: the model does not run on a consumer GPU or on a 32–64 GB machine; its floor is roughly 100–128 GB of unified memory at one to three bits,

retaining 71–82% of its unquantized top-1 accuracy on the quantizer’s own harness (§10.6) [714] [P]—a single workstation-class machine, not a laptop.

The domestic-silicon claim is real, narrow, and easy to overstate—so state it narrowly. Z.ai’s own announcement says the model was “previously previewed as Ox Alpha, *running* entirely on Chinese AI chips”; founder Jie Tang adds, “Powered by pure Chinese chips” [715] [P]. Four qualifications are load-bearing and none is a quibble. *First, the claim is about serving, not training:* Z.ai has made no training-hardware statement for GLM-5.3 or for Flash, and implicator.ai’s formulation is exact—served is not trained. *Second, no chipmaker is named:* Counterpoint’s Ivan Lam speculates Ascend plus other domestic vendors, which is analyst inference and not disclosure, and CNBC reported that it could not verify the claim independently [A]. *Third, it is not a first:* Zhipu stated that GLM-Image was trained end to end on Huawei Atlas hardware in January 2026, with no server count or timing disclosed, and the GLM-5 family was reported in June 2026 to have been trained on roughly 100,000 Ascend 910B with MindSpore and no NVIDIA hardware at any stage [716]. *Fourth—and this is the qualification that changes the grade of the third—the June claim has no locatable primary.* It circulates through TechTimes, Tom’s Hardware and Decrypt without a cited source; the GLM-5 repository carries no hardware statement; and Huawei’s own Ascend announcement of February 2026 describes “0-day” *inference* support and post-release training guides for GLM-5, which is adaptation after the fact, not a training claim [763]. This book therefore grades the GLM-5 training claim as press-derived [A/C], not as an unaudited company statement, and does not stack anything on it; the only Zhipu domestic-*training* claim with a primary behind it is an image model of undisclosed scale [P, unaudited]. Z.ai’s own Chinese-language documentation for Flash is consistent with that reading, because it claims exactly what the English announcement claims and no more: the phrase is *shouci zai daquimo liuliang zhong ... tigong fuwu*—the first time large-scale traffic has been *served* on a domestic cluster—and it says nothing about training [P]. Against all of it stands the reporting that DeepSeek’s R1 training did not complete reliably on Ascend and reverted to NVIDIA [104]—which is precisely why the training/serving distinction carries weight: the one domestic training attempt with independent reporting attached to it is the one that failed. What *is* new, and it is substantial, is production-scale inference: a reported cluster of roughly 100,000 domestic accelerators serving 62 trillion tokens over the preview week while carrying a workload that ranked fourth in the world on the launch-week (v4.1.1) board [715] [A]. Circulating reports of “100 trillion tokens per day” are irreconcilable with router data and are not used here.

Score it against the test this book actually set. Proposition P2 is confirmed by a frontier-class model *trained* end to end on domestic silicon, packaging and memory—claim CR-09, the “Mate 60 moment” of AI. August 2026 supplies the serving half of that stack at production scale, on a workload that was top-four in the world in its launch week, under a company statement that is unusually specific about what it does *not* say. That is the closest the integration test has come to being passed, and it is not a pass—and the fourth qualification above makes the failure cleaner, not closer: no training claim has been made for these models; the one training claim for a predecessor that has a primary is an image model with no scale disclosed; the frontier-scale training claim that the trade press repeats has no primary at all; no vendor is named; and the memory and packaging inside that cluster are undisclosed. A frontier model trained end to end on domestic silicon, packaging and memory has not been passed as a test, and the evidence that it has is thinner at 9 September 2026 than the June headlines made it look. CR-09 therefore remains “not yet observed” and this book still does not assert it. The revision the evidence does support is narrower and still consequential: the *inference* leg of sovereign compute has moved from roadmap to production, which is the leg that carries the deployed token volume, sets the standards, and pays the bills—and it has done so on the workload the whole world is currently ranking.

10.3.4 The barbell: where the tokens go and where the money stays

The router and gateway data of the fortnight before 9 September 2026 make an argument that the index tables cannot, and it is the spine of this book’s reading of the small model. Take the volume side first. On OpenRouter, with data through 8 September, the four largest models by routed tokens were Tencent’s Hy4 preview at 18.7 trillion, GPT-5.6 Luna at 14.5 trillion, DeepSeek V4 Flash at 12.4 trillion and **GLM-5.3-Flash at 12.2 trillion**; the rest of the top ten were DeepSeek V4 Flash 0423, MiniMax M3, Hy3, NVIDIA’s Nemotron 3 Ultra, Xiaomi’s MiMo-V2.5 and GLM-5.3 itself, at 3.4–5.0 trillion each. Where the active-parameter count is disclosed at all, the top ten runs 49B, 13B, 18B and below: the volume tier is a sparse-mixture tier, and the frontier flagships are not in it. An OpenRouter engineer, quoted in the same week, put Claude Fable 5—of the closed family that tops the v4.3 board—in “the low billions” of monthly output tokens on the platform against roughly 12 trillion each for Luna and GLM-5.3-Flash, on a platform whose total volume has risen more than twenty-five-fold in a year and doubled in the last month [760] [A]. Router volumes are a small, promotion-distorted slice of global inference and are read here as a direction and not a census; but three orders of magnitude is a direction.

Now the spend side, from a gateway that sees the invoices as well as the tokens. On Vercel’s AI Gateway, open-weight models took 36% of July’s tokens on **8.6% of spend**, and 62% of tokens on a single day, 22 August; 81% of tokens ran on models that did not exist six months earlier; and the gateway puts the self-hosting break-even “past 50 million tokens per day.” In the same data the four leading frontier laboratories still take 95% of spend, and Anthropic alone 61–65% of it, at roughly 4.4 times the average token price [761] [A]. That is a barbell: the sparse open tier is winning *volume* and losing *spend*, and the closed frontier is doing the reverse, and the two halves of the market are separating rather than converging. It is also, read against Section 2.3, exactly the rent-relocation pattern this book predicted, arriving in the accounts: the model layer at the volume end has been commoditized to a price at which a gateway’s customers move to whichever sparse model shipped last month, and the rent has moved up—into the capability tier, into the harness (Chapter 11), into the hardware, and, as §10.4 shows, into the terms on which the commons is served at scale. Morgan Stanley’s formulation of the consequence is the correct one and the correctly limited one: open weights “may really kill... the idea that the foundation model layer can sustain high rents”—the rent, not the compute demand, which the same barbell shows growing on both arms [762] [A].

The small-model argument has a literature, and the literature cuts both ways in the way the data do. Belcak and colleagues at NVIDIA Research argued in 2025 that most agentic sub-tasks are repetitive, narrow and non-conversational, and are therefore better served by small models with a large model held in reserve—a thesis the router data now look like a demonstration of, and one whose author’s own laboratory then shipped a 120B agentic flagship, which is not small [762] [A]. Against it stands the “Hierarchy of Agentic Capabilities” study of January 2026: on 150 realistic workplace tasks, “weaker models struggle with fundamental tool use and planning,” and even the best models fail about 40% of them [762] [A]. The Flash sub-scores of §10.3.3 are the two findings in one model: a third of Terminal-Bench v4.0 is the best result at its price and a long way from a solved workload, and the composite index still puts it eleven points below the closed leaders. The counter the book has to print is therefore the strongest one available, and it is in the same data as the thesis: spend concentration in the frontier is evidence that the capability tier still commands a rent, that buyers who can measure the difference pay 4.4 times the average price for it, and that the 40% of workplace tasks nobody’s model completes is where that rent is earned. What the barbell does not show is which arm grows faster, because token volume and spend are being measured on different customers—the volume arm on a router whose users self-select for price, the spend arm on a gateway whose largest customers self-select for capability—and no source at 9 September 2026 reports both on one population. The book’s position is correspondingly narrow: the volume tier is Chinese, sparse and open, and no longer commands a rent; the capability tier is American, metered and closed, and still does; and the

strategic question of Chapter 21 is not whether the first is true but how long the second lasts.

10.3.5 AI-designed silicon: why the compute dependence is transitional

The Qwen3.8-Max release supplies what may be the single most consequential datum in this chapter, and it bears directly on the one element of the Chinese stack that Part II has consistently graded weakest.

Alibaba’s own announcement lists, among Qwen3.8-Max’s headline capabilities, “system-level autonomous planning with closed-loop adaptive learning, driving *500+ turns of chip design optimization*” [393], and its accompanying material claims the model has “independently achieved the autonomous execution of the entire silicon design flow” [395]. The observation that makes this strategic rather than merely impressive was put most directly by an independent Western commentator reviewing the release: *one of the things China is behind on is chip design and chip manufacturing—but a model that can help design those chips means that is no longer the weakness* [395].

Set that against the architecture of Part II and the loop closes with unusual precision. China holds an open instruction set with no revocable licence (Section 12.1); domestic logic at 5 nm-class without EUV; a memory strategy that routes around HBM (Section 12.8); an open interconnect and software stack; a demonstrated exascale integrated-CPU machine (Section 12.2); and a national EDA push against the one remaining Western duopoly. The binding constraints on that stack have never been ambition or capital—they have been *design throughput and tooling*: the accumulated human expertise and proprietary software that turn an architecture into a manufacturable die. That is exactly the input an autonomous silicon-design flow attacks, and it is the input least protected by export control, because a design capability that a Chinese lab trains cannot be embargoed at a border.

This is the flywheel of Section 24.3.4 operating on the element the book had marked as slowest. AI improving the design of the chips that run the AI is the recursive loop in its most literal form, and it arrives—if the claims survive scrutiny—aimed squarely at the gap between the open-frontier position China already holds and the sovereign-compute position it does not yet. The corollary for this book’s evidence base is that the current arrangement, in which Chinese labs reportedly train on NVIDIA capacity leased offshore in Singapore and Malaysia while serving domestically [404], should be read as *transitional rather than structural*. It is what a fast-moving ecosystem does while its own silicon matures; it is not a stable dependence, and the co-design mechanism that ends it is now visible in a shipped product rather than a roadmap. Claim CR-09—a frontier-class model trained end-to-end on a domestic stack—remains “not yet observed,” and this book still does not assert it; §10.3.3 shows how far the *serving* half of that stack has come and why the training half is still not in evidence. But the pathway to it has shortened, and the forecast register is revised accordingly.

The most striking corroboration comes from a commentator who reaches the book’s conclusion while explicitly not wanting to. Reviewing the release for a Western audience, Berman traces the adoption logic—US enterprises “can get 95% of the capability and pay a fraction of the price,” and for the vast majority of use cases “the absolute frontier just isn’t necessary and in fact might be overkill,” so they will choose the open models, which are Chinese—and then follows it one step further to the conclusion Part II of this book is built on: *“eventually what will likely happen is the model will be co-designed with the chip to be the most efficient and the most optimized for that specific chip... with extreme co-design between model and hardware it might actually put us, as we’re becoming dependent on Chinese open source, highly dependent on Chinese chips as well”* [395]. That is the convergence thesis of this book, stated independently, by an observer who calls the resulting geopolitical exposure “not my opinion—it is just a fact,” and who is manifestly unhappy about it. When the mechanism this book has spent Part II describing is arrived at independently by someone who would prefer it were false, the argument has stopped being idiosyncratic.

Two disciplines keep this section honest. *The claims are the vendor’s*: 500 turns of chip-design optimization and an autonomous silicon flow are Alibaba’s descriptions of Alibaba’s model, and the appropriate verification event—an independently reproduced tape-out-grade design, or third-party evaluation of the flow—has not occurred. The GitHub project trace published alongside the autonomous-coding demonstration is a genuine step toward auditability and more than most Western frontier claims carry; the chip-design claim has no equivalent artifact yet. *And the counter-argument is real*: the same review notes that the leading US closed models are rumoured at roughly seven trillion parameters, against 2.4–2.8T for the Chinese frontier—roughly twice the scale, which requires the best accelerators in the world and a great many of them [395]. The Chinese achievement is therefore properly read as *capability parity at a third of the parameters and a fraction of the price*, which is an efficiency victory rather than a scale victory—and efficiency victories are exactly what Part I’s precedents are made of, but they are not the same thing as removing the compute constraint. Section 28.2.1 takes the strongest form of that counter-argument seriously.

Kimi K3 deserves one more paragraph, because it is the cleanest single demonstration of the five-frontier framework of §10.5. Its own technical report confirms 2.8T total parameters, 104B activated, a one-million-token context, native vision—and explicitly acknowledges that the model still trails the strongest proprietary systems overall [462]. Here is the open frontier and the absolute frontier in a single primary source: the largest, most capable open-weight artifact ever released, self-described as behind the closed leaders. Both of this book’s core observations—that China leads the open frontier decisively, and that the absolute frontier has not dissolved—are stated by the releasing laboratory itself.

Alibaba (Qwen) is the adoption champion: the most-downloaded open family on Earth, crossing one billion cumulative Hugging Face downloads in July 2026 with over 113,000 derivative models—more than Google and Meta combined, and roughly 40% of all new LLM derivatives on the Hub [85, 166]. **Moonshot (Kimi)** is the scale champion: K3’s 2.8 trillion parameters (896 experts, ~1.8% activation, hybrid linear attention, native MXFP4, 1M context) make it the largest open-weight model ever released, ranking in the global top five overall—the only open model in that tier [88]. **DeepSeek** is the efficiency champion, each generation shipping an attention or sparsity innovation the rest of the field adopts within months. **Zhipu (GLM)**—which listed in Hong Kong in January 2026 as “the world’s first publicly traded foundation-model company”—holds the open-weight quality crown on independent indices, at 45 on the re-based v4.3 Artificial Analysis index against 53 for the best closed systems (60 against 63 on the launch-week v4.1.1 board), and is reported in the trade press—with no primary that this book could locate—to have trained GLM-5 on Ascend, a claim graded [A/C] and separated carefully in §10.3.3 from the serving claim Z.ai actually makes for GLM-5.3-Flash [90, 713, 758, 763, 92]. Around them: **MiniMax** (M-series, agentic coding), **StepFun** (throughput-optimized), **Xiaomi** (MiMo, which took the largest single-provider share on OpenRouter within months of entry), **ByteDance** (Doubao at 345M MAU and >120 trillion tokens/day; China’s top-rated Arena model), and **Tencent Hunyuan** [83, 96]. Zhipu and MiniMax IPO’d in Hong Kong in January 2026; Moonshot raised at a \$20B valuation [92], and by September had filed confidentially for a Hong Kong listing of its own at a reported pre-money valuation near \$50B [757] [A]. DeepSeek, which had taken no outside capital at all, took its first in June 2026: roughly \$7.4B at a valuation above \$50B, with the state’s National AI Industry Investment Fund—the “Big Fund”—as the sole holder of direct equity and voting rights and Tencent and CATL in a non-voting partnership behind it [733] [A]; a later report of a \$74B valuation conflicts with that account and is not used here. The LLM price war is the “531” cull in progress: token prices at 1/7th to 1/150th of US equivalents, margins near zero, survivors selected by capability-per-yuan [93].

10.4 The toll road: the licence becomes a claim on the clouds

Section 10.3.1 ended with a licence clause and a prediction: that rent, driven out of the weights by the ratchet, would relocate to the terms of large-scale service, and that the releasing laboratory would retain a claim there. In the last fortnight of August the claim acquired a counterparty, and the counterparty is the American hyperscaler.

The fact pattern, at the strength the sourcing supports. On 26 August Reuters reported, on three anonymous sources, that Moonshot AI is in *early-stage* talks with Microsoft, Amazon and Google over agreements under which the three clouds would host Kimi K3 and Moonshot would take a share of the revenue—the laboratory seeking **up to 30% of K3-related service revenue**, on terms said to match those offered to its large direct customers. The revenue split, data access and the auditing of token counts were all unresolved; the three clouds declined comment and Moonshot did not respond; nothing was signed, and Reuters’ own formulation was that there is “no certainty” agreements will result. Smaller platforms had already signed: Chinasoft International disclosed a “token revenue-sharing and joint innovation agreement” with Moonshot on 20 July, percentages undisclosed [747, 748] [A]. Every account says *revenue*, not profit, and none attributes to the talks any motive beyond monetization; this book does not add one.

What forces the negotiation is the licence, not diplomacy. The reason a hyperscaler must talk to Moonshot at all is a clause that did not exist for K2. The K3 licence provides that a licensee operating a model-as-a-service business whose aggregate revenue exceeds \$20 million over any consecutive twelve months “must enter into a separate agreement with Moonshot AI before using the Software or its derivative works for any commercial purpose,” with a display requirement above 100 million monthly users or \$20 million of monthly revenue [748] [P]. That is consistent with K3, unlike K2.5, not being on Amazon Bedrock, and why it is reachable in Microsoft Foundry only through Fireworks AI as an intermediary partner model, at \$3.30/\$16.50 per million tokens [750, 748] [P]. Alibaba is copying the instrument: Reuters reported on 7 August that large commercial users of Qwen3.8-Max will be required to share revenue, the rate under negotiation and the design “modelled on Moonshot’s,” with the \$50 million threshold of the **qwen3.8-max** licence as the lever [749] [A]. The clause is the mechanism; the hyperscaler is the counterparty; and the rent Section 2.3 said would leave the weights has arrived, in draft, in a contract.

Against what already exists. The talks are startling only against the baseline, so the baseline has to be stated from the clouds’ own documentation. Amazon Bedrock has served DeepSeek-R1 since January 2025 and, since February 2026, DeepSeek V3.2, MiniMax M2.1, GLM 4.7 and 4.7 Flash, Kimi K2.5 and Qwen3 Coder Next as fully managed models; Google Vertex serves Kimi K2 Thinking, MiniMax M2, DeepSeek-V3.2, GLM 4.7 and GLM 5; Microsoft Foundry lists DeepSeek in its “sold by Azure” category. AWS’s third-party model terms state that DeepSeek and Qwen serverless models “are sold by AWS” and that customer data is not shared with the model providers. No cloud discloses any payment to any Chinese laboratory for serving permissively licensed weights, and Bedrock lists DeepSeek V3.2 at roughly two to five times the laboratory’s own price [750] [P]. The American hyperscalers, in other words, are already the largest Western resellers of the Chinese commons, at a markup, for nothing. What the Moonshot talks would be—if they close—is the first time the maker of the commons takes a cut of the commons’ service revenue on a Western cloud: a toll booth built on somebody else’s road, and the sharpest instance yet of the result Section 2.3 predicted, because the weights remain free and the rent is collected one layer up, from the party with the customers.

The “hedge” reading, stated and rejected on the evidence. It is tempting to read the talks as insurance—a Chinese laboratory buying protection against a Western shutdown by making American clouds its partners, or the clouds buying protection against a Chinese one by contracting for supply. No source attributes either motive, and the direction of the legal exposure runs the other way. Webster and Xu’s analysis of the K3 licence is the only serious treatment located,

and its title is its argument: *with revenue comes regulability*. Downloading open weights left it genuinely ambiguous whether any transaction with a Chinese laboratory had occurred; “if a US company needs a contract with Moonshot to provide the inference tokens,” the picture changes, under a supply-chain rule “it’s harder to argue there’s no transacting going on,” and if Moonshot were sanctioned “the contract to provide inference is clearly an instance of doing business” [751] [A]. The instruments that would give that exposure teeth are partly enacted and partly pending. Section 1532 of the FY2026 National Defense Authorization Act, in force since December 2025, requires the Department of Defense to exclude “covered AI”—defined by reference to DeepSeek and High-Flyer by name, to the Consolidated Screening and 1260H lists, and to covered-nation domicile—and does not distinguish self-hosted weights from cloud-served tokens; the No Adversarial AI Act remained in committee at 9 September 2026 [752] [P]. Treasury Secretary Bessent was reported in July to have threatened Entity-List designation of Chinese laboratories [A]; Axios reported the same month a draft executive order that would make US hosts of Chinese models guarantee their security and accept liability, not issued at the currency date [751] [A]. And on 9 September 2026 itself the NSA, FBI and CISA issued a joint advisory naming DeepSeek, Moonshot, Alibaba, MiniMax, StepFun and Z.ai as running distillation campaigns against US frontier models, with the traffic routed through cloud providers and third-party aggregators—an advisory about distillation rather than hosting, but one that names the prospective toll-road partner and the clouds in the same paragraph [754] [P]. A revenue contract with a US hyperscaler is therefore not a shield against any of that; it is the document that would convert an ambiguous download into an unambiguous transaction. On the Chinese side the picture is the mirror image and equally unresolved: the Ministry of Commerce spent June and July consulting Alibaba, ByteDance and Z.ai on options that include barring public release of the most advanced models, restricting overseas access and criminalizing model leakage, with “nothing decided” and no timeline, and Bloomberg reported in May that passport-surrender rules first applied at DeepSeek had been extended to Moonshot, StepFun and ByteDance [753] [A]. If there is a shutdown the toll road hedges against, then, it is the Chinese one—a contracted revenue stream from Redmond, Seattle and Mountain View being a harder thing for Beijing to close than a download link. That is this book’s inference and not anyone’s reporting, and it is graded accordingly [S].

What it does to the thesis. The commons is acquiring a cash register, and the two facts that matter are that it remains a commons and that the register sits exactly where the theory said it would. Below the threshold nothing changes: every researcher, every enterprise deploying internally, every derivative-builder and every service business under \$20 million of revenue uses K3 on the terms it was released under, and the download counts of Table 10.4 continue to accrue. Above it, the same artifact is a toll road, and the toll is levied not on the user but on the reseller with the largest customer base on earth. That is “openness relocates rent” in its most literal form yet, and it is the result of Section 2.3 and §10.3.1 with a named counterparty and a percentage attached. The counter is the size of the register. Moonshot’s reported annualized revenue was \$300 million in June, with no geographic split disclosed, against a frontier laboratory’s run rate two orders of magnitude larger [757] [A]; 30% of a hyperscaler’s K3 line is a claim on a market that NVIDIA’s Poolside purchase and its Nemotron line are built to commoditize from the other side (Chapter 17); and a negotiation that Reuters describes as early-stage, unsigned and silent on its three hardest terms is evidence of intent, not of income. The toll road is real as a mechanism and unproven as a business, and the ratio between those two statements is the ratio this chapter has found at every layer of the model element.

10.5 Three scoreboards, not one

Before the numbers, a distinction that resolves most of the argument about whether China has “caught up.” There are *five* different frontiers, and conflating them produces both triumphalism

and complacency:

1. **The absolute frontier** — the highest demonstrated capability under controlled evaluation, irrespective of access or price. The United States leads this, and the lead is real.
2. **The open frontier** — the highest capability available as downloadable weights under usable licence terms. China leads this, and not narrowly.
3. **The deployable frontier** — the highest capability an institution can *operate reliably at acceptable cost*, after quantization, serving constraints, security review, and licence compatibility. China increasingly leads this.
4. **The economic frontier** — the system with the lowest cost per *successful task* for a given workload (§10.6.3). This is frequently neither principal’s flagship: for high-volume routine work it is often an older, smaller model that nobody writes about.
5. **The scientific frontier** — the system that produces the greatest validated scientific value per unit of national resource. This is the frontier Chapter 24 argues matters most, and the one nobody currently measures.

	leader today	decides what
1. Absolute best under controlled evaluation	United States	prestige, hardest tasks
2. Open best downloadable under usable licence	China	developer default
3. Deployable best operable reliably at acceptable cost	China (rising)	industrial outcomes
4. Economic lowest cost per <i>successful task</i>	often neither: an older, smaller model	the majority of tokens
5. Scientific validated scientific value per unit of national resource	unmeasured	long-run national power

Apparent contradictions dissolve here. China can trail frontier 1 while leading 2 and 3. A US closed system can dominate frontier 5 for a given workload despite far higher token prices, because scientific value per resource depends on capability at the hard end rather than price at the volume end. And frontier 4 is frequently defined by a model nobody writes about.

Figure 10.1: Five frontiers, won separately. Only the first two are systematically measured; the third and fourth are measurable but rarely reported (§10.6.3); the fifth—the one Chapter 24 argues matters most for national outcomes—has no accepted metric at all, which is both a gap in the literature and an invitation to whoever builds the first credible one. Positions are this book’s assessment as of the version stamp [S].

Many apparent contradictions dissolve under this framework (Figure 10.1). China may trail the absolute frontier while leading the open and deployable ones. A US closed system may dominate the *scientific* frontier for a particular workload despite far higher token prices, because scientific value per resource depends on capability at the hard end rather than on price at the volume end. And an eighteen-month-old open model may define the economic frontier for the majority of tokens actually generated in the world. Frontier 3 is where the two halves of this chapter meet: a model that ranked fourth in the world in its launch week (v4.1.1) and sits eleven

index points below the closed leaders on the re-based v4.3 board, and that fits, quantized, in the memory of one workstation (§10.3.3, §10.6) is a capability an institution can operate without a vendor relationship, a network egress, or a procurement fight, which is a different kind of lead from the one the index tables measure. That China may lead frontiers 2, 3, and 4 without leading 1 *strengthens* this book’s thesis rather than weakening it, because industrial victory in every precedent of Part I depended on what a buyer could actually install, at what price, under what terms—not on the laboratory maximum. The rest of this section reports frontiers 1 and 2, which are measured; Chapters 18 and 26 supply what is known about 3 and 4; frontier 5 is, at present, an open measurement problem and an invitation.

Three independent measurement families bear on the adoption proposition (P1), and they must be kept apart, because they measure different things and currently disagree in one important place. *Family one—capability*: index and arena measurements of model quality. *Family two—distribution*: downloads, derivatives, and routed tokens. *Family three—production and revenue*: enterprise deployments and paid spend. Families one and two tell one story (Table 10.5); family three lags it, and honesty requires displaying the lag.

On capability: the Stanford AI Index puts the top US–China gap at 2.7% as of March 2026, down from double digits in 2023—while also finding that the closed-over-open gap *reopened* to 3.3% (from 0.5% in August 2024) and that six of the top ten arena models are closed [94]. Both facts are true at once: China reached the top competitive tier; the absolute closed frontier did not dissolve. A second instrument agrees on the shape and, at 9 September 2026, disagrees with itself on the size: on the Artificial Analysis index the best open-weight Chinese model, GLM-5.3, scores 45 on the re-based v4.3 board against 53 for the best closed models, Claude Fable 5.1 and GPT-6 Astra—*eight* index points—where the launch-week v4.1.1 board of a fortnight earlier had it at 60 against 63, three points, and the 320B/18B GLM-5.3-Flash at 42 on v4.3 sits eleven behind the leaders, having been fourth in the world and level with GPT-5.6 Sol on the board it launched on [713, 758] [A]. The Stanford figure and the Artificial Analysis figure are different instruments on different scales, and a 2.7% gap on one is not three points or eight points on the other; two instruments agreeing that the gap is small and closed-favouring is corroboration of a direction, not confirmation of a magnitude, and neither instrument measures the hard agentic tail where the closed lead is widest—which is exactly the tail the September re-basing weighted more heavily, and the re-ordering it produced is itself evidence of where that lead lives. Benchmark figures additionally carry protocol variance that vendor tables hide—Kimi’s own documentation shows 65.8% on SWE-bench Verified single-attempt becoming 71.6% with parallel test-time compute [87], and reasoning-model response lengths are inflating $\sim 5\times$ per year [V] [254], so a model ten times cheaper per token that spends ten times the tokens per solved task has gained nothing per task. On distribution and production, Table 10.4 states what each metric can and cannot support.

One adoption channel none of those six metrics registers is the silicon industry’s own. At Hot Chips 2026, nine separate decks reported inference economics on Chinese open-weight models: d-Matrix on GLM 5.2 and Kimi K3 serving rates; OXMIQ building its cost model on Kimi-K2; NVIDIA measuring agentic throughput on DeepSeek-V4-Pro in the Rubin deck and using DeepSeek V3 training-step time as the headline multi-tenancy result in the Spectrum-X deck; Intel plotting every published frontier model from Llama 2 70B to Kimi K2 to establish its capacity-versus-bandwidth thesis; SambaNova quoting MiniMax M2.7; OpenAI reporting DeepSeek R1 on GB300 [723] [V]. No Chinese chip was named anywhere in the conference and no Chinese vendor presented; what the Chinese laboratories supplied was the *workload*. American silicon vendors benchmark and publish against Chinese weights because those are the frontier-scale weights they can obtain, run, and quote in public—which makes Chinese model artifacts a de facto reference load in the design and marketing of Western accelerators, a form of standard-setting that no download counter captures and that runs in the opposite direction to the procurement bans of Chapter 18.

Table 10.4: The adoption-metric matrix: what each measure proves.

Metric	Measures	Principal limitation
HF downloads (41% Chinese)	Distribution, developer interest	Automation/CI inflation; no proof of production use [85]
Derivatives (>113k Qwen)	Ecosystem reuse	Many experimental or dormant
OpenRouter tokens (~61% mid-2026)	Usage on one neutral router	Platform-specific mix; not a census [84]
API revenue	Monetized hosted demand	Misses self-hosting entirely
Enterprise deployment (DeepSeek >26k orgs [P])	Operational adoption	Vendor-survey selection [167]
Enterprise <i>spend</i> share	Follows the money	Fell 19%→11% in 2025—the countertrend [168]

Table 10.5: The model element, mid-2026 snapshot.

Capability	Adoption
US–China top-model gap: 2.7% (Stanford AI Index 2026), from 17.5–31.6% in 2023 [94]	Chinese models: 41% of all Hugging Face downloads (trailing year); China passed the US in cumulative downloads [85]
Release lag behind US frontier: from 6–8 months to “a matter of weeks” [95]	OpenRouter routed tokens: Chinese models ~2% (early 2025) → ~61% (mid-2026); Meta/Llama below 1% [84, 96]
Top open-weight models: GLM-5.3 and GLM-5.3-Flash, Kimi K3, DeepSeek V4, MiniMax M3—all Chinese; the best US open entry is still an NVIDIA model, sixth on the router at 5.62T weekly tokens (Chapter 17) [84, 720]	a16z: of startups pitching with open-source stacks, ~80% use a Chinese model; DeepSeek counts >26,000 enterprise deployments [97, 167]




Figure 10.2: The model element in two pictures. *Above*: share of tokens routed through a neutral aggregator by author origin; Chinese-origin models passed US-origin models in June 2026, having been at $\sim 2\%$ eighteen months earlier [V, platform-specific] [84, 96]. *Below*: the two gaps the Stanford AI Index measures. The national gap closed to 2.7%; the closed-over-open gap *re-opened* to 3.3% [V] [94]. Both facts are true, and the book’s argument requires holding them together: China reached the top competitive tier, and the absolute closed frontier did not dissolve. Intermediate points are interpolated between reported measurements.

Two further facts frame the table. On hard, contamination-resistant agentic benchmarks the closed-open gap runs larger than headline indices suggest [89]; and the investment asymmetry is staggering—US private AI investment ran \$286B in 2025 against China’s \$12B (a figure that understates Chinese state capital, per the Index’s own caveat) [94]: a 23:1 private-spending ratio buying a 2.7% measured gap. Recall the solar lesson: Sharp and Q-Cells also outspent Suntech per watt of R&D. The question is not who spends most for the peak; it is who owns the volume, the derivatives, the developers, and the price floor. Meta’s Llama—the West’s open standard-bearer—collapsed below 1% of routed volume and retreated toward closed development; OpenAI’s answer, gpt-oss (August 2025), was outclassed within weeks [98]. The family-three lag—spend share falling even as usage share soared—is the strongest current evidence *against* P3, and Chapter 21 treats it as such; the next section explains the mechanism by which the curves are nonetheless converging.

10.6 Quantization: frontier models on commodity hardware

The most under-analyzed force multiplier of the open-weight strategy is quantization. Open weights alone democratize *access*; quantization democratizes *operation*. The claim must be stated with more care than it usually is. For selected models, quantizers, and benchmark suites, community tooling has reduced storage by 60–90% while preserving much of benchmark performance; degradation is highly task-, layer-, context-, and calibration-dependent, and becomes material at extreme 1–2-bit precision. What is robust is narrower and still consequential: 4-bit has become a *release* format rather than a lossy afterthought, and 2-bit-class dynamic quants retain enough capability to be operationally useful on hardware an order of magnitude cheaper. Any serious evaluation reports a *quality vector*—mathematics, coding, multilingual reasoning, long-context retrieval, tool use, calibration, refusal behaviour, safety—alongside memory savings, throughput, and energy, rather than a single retention number.

10.6.1 The dynamic-quantization breakthrough

The emblematic actor is Unsloth, an eight-person company whose “dynamic” GGUF quantizations allocate precision non-uniformly—MoE expert layers driven to 1–2 bits, attention and dense layers held at 4–6 bits, calibrated on curated token sets against KL divergence. Days after R1’s release, Unsloth’s 1.58-bit dynamic quant cut the 671B model from 720 GB to 131 GB (an 80% reduction) while a naive 1.58-bit quant produced gibberish; the dynamic version retained ~70% of a demanding code-generation benchmark and the 2.2-bit version 92% [169]. A precision discipline applies to all such figures: retention is task-dependent, and at 1–2 bits it degrades unevenly across mathematics and code, factual recall, long-context retrieval, agentic tool reliability, multilinguality, and safety behavior—the numbers quoted here are specific quants on specific benchmarks [P], not a general law that “90% compression costs single digits.” The robust claims are narrower and sufficient: 4-bit is now effectively a release format (below), and 2-bit-class dynamic quants preserve enough capability to be operationally useful on hardware an order of magnitude cheaper. The method scaled with the frontier: Kimi K2 (1T) to 245 GB at 1.8 bits; K2 Thinking to 245 GB at 1 bit with “~85% accuracy retention”; GLM-5 (744B) to 176–241 GB; Kimi K3 (2.8T, 1.56 TB native) to 594 GB at ~79% retention and 861 GB at ~90% [170]. Unsloth alone counts ~10 million monthly model downloads [171]; the broader quantization stack—llama.cpp/GGUF, AWQ (MLSys 2024 best paper), GPTQ, EXL3, bitsandbytes/QLoRA, vLLM’s quantized serving matrix—is now standard infrastructure [172, 173].

Crucially, the labs closed the loop by shipping *natively* quantized frontier weights: DeepSeek in FP8 since V3, Kimi K2 Thinking in INT4 via quantization-aware training (2× speedup at near-FP16 quality), K3 in MXFP4, OpenAI’s gpt-oss in MXFP4 [87, 88, 98]. Four-bit inference is no longer a lossy afterthought; it is the release format, co-designed with training. This matters

doubly for Chapter 12: native low-precision weights quarter the memory-capacity and bandwidth demand of serving—which is exactly the resource the LPDDR-based architecture supplies in abundance.

10.6.2 The commodity-hardware frontier

The consequence is that frontier-class inference now runs on hardware measured in single-digit thousands of dollars. The documented data points span three classes. *CPU servers*: a \$2,000 used-EPYC build (512 GB DDR4) runs the full 671B R1 at 3.5–4 tok/s; KTransformers (Tsinghua, SOSP 2025) reached 286 tok/s prefill and 12+ tok/s decode for R1/V3 on one RTX 4090 plus a dual-Xeon board with 1 TB of DRAM—a few-thousand-dollar machine serving the model that crashed NVIDIA’s market cap [174, 175]. *Unified-memory desktops*: a single Mac Studio M3 Ultra (512 GB unified at 819 GB/s, ~\$10k) runs R1 4-bit at 17–18 tok/s under 200 W; two of them ran Kimi K2.5 (1T, native INT4) at 22 tok/s; four in an RDMA-over-Thunderbolt cluster (~\$40k, 1.5 TB unified) served K2 Thinking at 25 tok/s under 450 W [176]. *Dedicated local-AI boxes*: NVIDIA’s DGX Spark (128 GB unified, \$4,700) and AMD’s Strix Halo (128 GB, \$2,300) both run 120B-class models at 34–39 tok/s, and 30+ laptop models ship with local-AI accelerators in fall 2026 [177]. llama.cpp’s MoE CPU-offload flag (August 2025) made the pattern general: one consumer GPU for attention, cheap system DRAM for the experts [178].

Note what this hardware list is, architecturally: *capacity-rich, bandwidth-modest memory attached to modest compute*—the LPDDR6X design point of Section 12.8, discovered independently and bottom-up by the enthusiast market. The sparse-MoE model family and the capacity-first hardware family are co-evolving at every scale from the \$2,000 desktop to the national cluster.

The sharpest instance yet is GLM-5.3-Flash (§10.3.3), and it is worth quoting at measurement precision because the temptation to overstate it is considerable. Unsloth’s published quantizations of the 320B/18B model run from UD-IQ1_S at 93.1 GB, retaining 70.9% top-1 accuracy against the unquantized reference, through UD-Q2_K_XL at 108.7 GB and 78.3%, to **UD-IQ3_XXS at 120.4 GB and 81.6%** and UD-Q4_K_XL at 199.7 GB and 92.2%; the BF16 weights are 641.6 GB. Minimum combined RAM plus VRAM is 100 GB at one bit, 128–150 GB at three bits and 162–210 GB at four; Unsloth states that the three-bit quantization “works on 128GB devices like a Mac or NVIDIA DGX Spark,” and reports measured decode on four DGX Spark units at native FP8 of 28–34 tok/s on short prompts and ~43 tok/s at 50K-token prompts [714] [P]. The comparison that gives the numbers their force is the flagship: full GLM-5.3 at FP8 requires eight H200- or H20-class accelerators. Three disciplines keep the claim honest. This is *workstation and small-server* territory—a single 128 GB unified-memory box of the class priced above at \$2,300–\$4,700 for the three-bit quantization, and the four-Spark configuration that produced the measured decode rates is a ~\$19,000 machine—and emphatically not a laptop; the accuracy figures are Unsloth’s own top-1 retention measurements on their own harness, not an independent evaluation across the quality vector this section insists on; and there is no verified throughput measurement on a Mac Studio for this model, so the memory footprint is the claim and the four-Spark decode rate is the only throughput number printed here. With those caveats stated, the strategic content survives all of them: a model that ranked fourth in the world in its launch week (v4.1.1) and stands eleven index points from the closed leaders on the re-based v4.3 board, above the closed volume tier it is priced against, now fits in the memory of a machine a single engineering team can requisition without a procurement cycle—which is the *deployable* frontier of §10.5 arriving, in a specific configuration, at a specific measured size.

10.6.3 The metric that matters: cost per solved task

Price per million tokens is the wrong denominator, and this book uses it only where nothing better exists. The economically meaningful quantity is the cost of a *completed* unit of work:

$$C_{\text{solved}} = \frac{C_{\text{input}} + C_{\text{output}} + C_{\text{tools}} + C_{\text{retries}}}{p_{\text{success}}},$$

where parallel sampling counts every sampled token, and where wall-clock latency and maximum concurrency are reported alongside. The correction is not pedantic: reasoning-model output lengths have been inflating roughly fivefold per year [254], so a model advertised at a tenth the price per output token that spends ten times the reasoning tokens per solved task has delivered no saving at all, and a headline “150 times cheaper” can survive translation into delivered task economics or vanish entirely. Wherever this book quotes a price ratio, read it as an upper bound on the true economic advantage until C_{solved} is published—which, as Appendix G records, essentially no vendor yet does.

10.6.4 The economics: LLMflation

Layered on falling hardware requirements is the steepest price deflation of any industrial input in history: inference cost at constant capability falls roughly $10\times$ per year—GPT-3-level capability fell from \$60 to \$0.06 per million tokens in three years [179]. Self-hosting break-evens are now measured in months for any serious deployment (roughly one month of payback at 50B tokens/month) [180]. The strategic geometry is the solar module’s: a Western closed-model business must recover \$100B-scale capex from a price that Chinese open weights plus community quantization reset toward marginal cost every quarter. What that geometry does not predict—and what the next section documents—is which party arrives at the low price first. Chapter 21 prices the consequences.

10.7 The price inversion

The obvious next step in that argument is the one the evidence does not take. If open weights reset price toward marginal cost, the expected consequence is a Western frontier laboratory forced downward—and in July and August 2026 a Western frontier laboratory did move downward, sharply and in public. But the details of the move, and the direction the Chinese price leader moved at the same time, break the simple version of the story, and this section prints the break because the book’s discipline is to print the datum that cuts against it.

Start with what was actually cut. On 30 July 2026—the press cycle ran on the 31st, which is why the date is often given a day late—OpenAI reduced the price of **GPT-5.6 Luna**, the cheapest tier of the family and the successor to the “nano” class, by 80%: from \$1.00/\$6.00 to \$0.20/\$1.20 per million input/output tokens. Terra was cut 20%, to \$2/\$12. *Sol, the flagship, was not cut at all*, holding at \$5/\$30 [717] [P]. Three weeks later, on 21 August, Sol went to \$4/\$20—explicitly “for the next 3 months,” which makes it a promotion rather than a repricing, and the distinction matters because a promotion is a defense of share while a repricing is a concession about cost. What the headline “80% price cut” describes, then, is a deep cut at the bottom of the ladder, a shallow one in the middle, and nothing at the top: the shape of a firm defending the volume end of its business while leaving the capability rent intact.

The consequence is the fact this chapter most needed and the one most easily missed. At \$0.20/\$1.20, GPT-5.6 Luna is *cheaper on list price than most Chinese open-weight models*: below GLM-5.3, below Kimi K3, below Qwen3.8-Max, and below DeepSeek V4-Pro on input. On the combined million-in, million-out basis of Table 10.6, Luna costs \$1.40 against GLM-5.3’s \$5.80, Kimi K3’s \$18 and Claude Opus 5’s \$30 [718] [P/A]. The incumbent did not merely match the commodity price; at the bottom of its ladder it went underneath it. Two precisions keep that

Table 10.6: The price inversion at 9 September 2026: list price in US dollars per million tokens, ordered by the combined cost of one million input plus one million output tokens—a crude but consistent basis, and an upper bound on economic advantage rather than a cost per solved task (§10.6.3). “Open” means downloadable weights, for which the listed price is the vendor’s hosted endpoint and the marginal price to a self-hoster is hardware. † DeepSeek’s July list, superseded by the increase of 13 August 2026 [719]. ‡ Sol was \$5/\$30 until 21 August 2026—held unchanged through the 30 July cut—when it was cut to \$4/\$20 for a stated three months [717].

Model	Access	Origin	Input	Output	1M+1M
DeepSeek V4-Flash†	open (MIT)	CN	0.14	0.28	0.42
GLM-5.3-Flash	open (MIT)	CN	0.15	0.50	0.65
MiniMax M2.5	open (custom)	CN	0.15	1.20	1.35
GPT-5.6 Luna	closed API	US	0.20	1.20	1.40
GLM-5.3	open (custom)	CN	1.40	4.40	5.80
Qwen3.8-Max	open (custom)	CN	2.00	6.00	8.00
GPT-5.6 Terra	closed API	US	2.00	12.00	14.00
Kimi K3	open (custom)	CN	3.00	15.00	18.00
GPT-5.6 Sol‡	closed API	US	4.00	20.00	24.00

from becoming a slogan. The inversion holds against the *frontier-scale* open models and not against the distilled tiers—GLM-5.3-Flash at \$0.65 combined and DeepSeek V4-Flash at its July list of \$0.42 both sit below Luna, and both are downloadable, so their marginal price to an organization with the hardware of §10.6 is not a list price at all. And per-token list prices remain an upper bound on economic advantage rather than a measure of delivered work: a cheaper token that buys more tokens per solved task has bought nothing (§10.6.3).

The cut worked, on the only evidence that can show it. Between 27 July and 14 August, Luna’s usage rose to 13.8× its pre-cut average and Terra’s to 5.6×, against 1.1× for the undiscounted Sol, with roughly 75% of the gain taken from competitors rather than cannibalized from the vendor’s own higher tiers [718] [A]. That is behavioural evidence of genuine cross-price elasticity at the volume end of the market—buyers moved, and they moved from somewhere else. The router data corroborate the shape: in the week to 29 August the fastest-growing Western entries on OpenRouter were its two cheapest metered tiers—GPT-5.6 Luna at 6.75 trillion tokens, up 31% after the cut, and Gemini 3.7 Flash at 4.05 trillion, up 156%—on a board still led by a Chinese open model, Ox Alpha at 20.7 trillion and up 216%, with DeepSeek V4 Flash at 12.3T and Xiaomi’s MiMo-V2.5 at 10.0T behind it [720] [A]. Router volumes are a small and promotion-distorted slice of global inference and are read here as a direction, not a census. One entry on that board belongs to neither camp as this chapter has drawn them: the sixth-largest model by routed tokens, at 5.62 trillion and up 27%, is *NVIDIA*’s Nemotron 3 Ultra, released free and open-weight by the chip vendor whose accelerators serve most of the rest of the list. Chapter 17 takes up what it means that the hardware incumbent is now the largest Western open-weight supplier.

Now the datum that cuts against the thesis. On 13 August 2026 DeepSeek—the laboratory that started the price war, the reference point for every “150 times cheaper” comparison in this literature, and the firm whose pricing this book has treated as the commodity floor—*raised* its API prices by as much as 1,100%, launching V4-Pro with peak and off-peak tiering; Caixin called it “a break from the prolonged price war among China’s AI developers” [719] [P]. No rationale was stated. The sequencing has to be read carefully, and it forbids the reading that would be most convenient in both directions: the deepest Western cut came on 30 July, two weeks *before* the Chinese increase, so the cut cannot have been a response to it and the increase cannot have been a response to the cut. Candidate explanations exist—capacity scarcity making peak-hour tokens genuinely more expensive to serve, a deliberate move up-market as the distilled tiers take the volume, or a research house declining to subsidize inference it would rather spend

on training—but each is inference, none is stated, and this book selects among them only to the extent of grading them all [M] and noting that the tiering itself is a capacity signal rather than a margin signal. What does not survive is the simple story in which cheap Chinese models force Western prices down: the price leader reversed direction mid-cascade, while the Western incumbent went below it. Commoditization at the volume end is real and observable in Table 10.6; the arrow of causation implied by the popular version of the story is not.

What the inversion does *not* do is recover the commons, and the reason is the structural point of §10.3.1 restated in price terms. A metered tier competes with a downloadable one only for the buyers who intend to rent tokens; for everyone who intends to run the model—the self-hosters of §10.6, the fine-tuners, the sovereign deployers of Chapter 22, the enterprises whose objection to a Chinese endpoint is jurisdictional rather than financial—the relevant price of an open model is not its vendor’s list price at all, and no cut to a closed API changes it. Undercutting GLM-5.3’s hosted endpoint therefore wins tokens without winning weights: it defends the volume tier of the rental business while leaving untouched the derivative ecosystems, the fine-tuning tooling, the quantization conventions and the standards gravity that Table 10.4 measures and that this chapter has argued are the currency the contest is actually denominated in. The two readings of the cut are correspondingly uncomfortable in opposite directions, and this book takes neither on faith: if \$0.20/\$1.20 is comfortably above cost, then the rent at the volume end was thin long before any Chinese release and the deflation argument of Chapter 21 must rest on the capability tier instead; if it is not, the incumbent is buying share at precisely the tier a commons with a floor at zero is best equipped to take. The evidence available at 9 September 2026 does not discriminate between them, which is why the dashboard watches the sequence of cuts rather than any instance.

The attribution discipline this book applies to Chinese vendor claims has to be applied symmetrically here, and it is uncomfortable. *No OpenAI or Anthropic executive has attributed a 2026 price cut to Chinese competition on the record* [V]. Every company statement cites efficiency—including the striking and entirely unaudited claim, in OpenAI’s own announcement of the 30 July cut, that GPT-5.6 Sol autonomously optimized its production kernels, “reducing serving costs by 20% and improving token-generation efficiency by 15%” [717] [P, vendor]. The China attribution that appears throughout the trade press, and that this chapter’s own argument would be pleased to accept, is press inference. The strongest adjacent primary statement is Altman’s, and its provenance has to be stated, because it is not the pricing announcement: both remarks come from a podcast appearance on *Invest Like the Best* on 29 July, the day before the cut. He said “We better be the greatest and the cheapest,” and, of Kimi K3 specifically, “You get a better deal today, at least at a particular like latency, using OpenAI’s models than Kimi”—and the middle clause is load-bearing, because it restricts the comparison to a latency band rather than asserting a general price advantage. Read it precisely: that is a public commitment to undercut the Chinese open models, which is strong evidence that they are the competitive reference point—and it is not a concession that they caused the cut. The distinction is the difference between a firm that is watching a rival and a firm that has been forced by one, and the record supports the first without establishing the second.

Nor can this chapter claim the cuts are below cost, because nobody has shown it. The available unit economics are all reported via *The Information* and unconfirmed by the companies: OpenAI at \$5.7B of first-quarter 2026 revenue against \$3.7B of cash burn, with adjusted gross margin around 33%—though the reporting conflicts even on direction, one series having the figure fall from 46% to 33% and another rise from 33% to 39%, which is itself the point; Anthropic at a \$65B run rate at the end of July with *API gross margin reported above 80%* [721] [A]. *No credible published estimate of loss per token exists, and no analyst asserts that these prices are below cost*—“inference loses money” is a claim about model mix, free tiers and training amortization, not a universal one, and the 80% figure sitting next to the 33% figure shows how much of the disagreement is really about what is being counted. The honest statement of what the price line

shows is therefore narrow: at the volume end, metered pricing has converged toward a level a commons with a floor at zero makes indefensible, and the American incumbent reached that level first and voluntarily; at the capability end, the rent is intact, unmatched, and—on the one margin figure anyone has published—very large. That is Proposition P3 confirmed at the bottom of the ladder and unproven at the top, which is exactly where Chapter 21 takes it up.

10.8 Openness as declared national strategy

What began as firm-level catch-up strategy is now explicit state doctrine. The Global AI Governance Action Plan (WAIC, July 2025) commits China to “cross-border open-source communities,” open sharing of “basic resources,” and lowering “the thresholds of technological innovation”—accompanied by the World AI Cooperation Organization, headquartered in Shanghai and founded in July 2026 with 29 member states (none a major Western democracy), a UN capacity-building group co-chaired with Zambia spanning 80 countries, and BRICS AI centers [99, 100, 764]. The State Council’s “AI+” plan targets self-controllable base models and a thousand deployed industry cases; draft policy counts open-source contribution toward academic credit [101]. As Tsinghua’s Liu Zhiyuan put it: “Open source has become politically correct” [102]. The export product is not a model; it is a *stack*—weights, tooling, quantizations, and cheap tokens—aimed at the Global South, where two-thirds of EV sales already run on Chinese LFP and the AI adoption default is being set the same way. Even Western institutions that ban the DeepSeek *app* on data-sovereignty grounds turn to self-hosted Chinese weights: Microsoft evaluated a fine-tuned DeepSeek V4 for a Copilot agent in June 2026 [181].

10.9 Compute constraints as forcing function

Export controls shaped this element rather than preventing it. Denied H100-class bandwidth and capacity, Chinese labs innovated in exactly the directions that reduce dependence on both: extreme MoE sparsity (K3 activates 1.8% of weights per token; V4-Pro 3%), attention compression (MLA, DSA, KV caches at ~2% of dense), low-precision training (FP8, INT4, MXFP4), and multi-token prediction [103, 88]. The 2025 episode in which DeepSeek’s next flagship was delayed by failed training runs on Huawei Ascend—then reverted to NVIDIA hardware while Ascend took inference—shows the transition is real but incomplete [104], and it is the reason this book grades the two halves of that transition separately. The report, five months later, that Zhipu had trained the GLM-5 family entirely on Ascend would show how fast the *training* half is completing if it had a primary behind it; it does not—the claim circulates in the trade press without a source, the repository carries no hardware statement, and Huawei’s own announcement describes post-release inference adaptation—so this book grades it [A/C], and the one Zhipu training claim that does have a primary, GLM-Image, discloses no scale and has not been audited [90, 763, 716]. What is on the record by 9 September 2026 is the other half—a model that was globally top-four in its launch week *served* at production scale on a domestic cluster, with no chipmaker named and no training claim attached (§10.3.3)—and the distinction is where the whole element turns, because inference is where the deployed volume, the standards and the revenue sit, while training is where the export controls actually bite. The design direction is now locked in: Chinese frontier models are co-evolving with the domestic hardware of Chapter 12—sparse, capacity-hungry, bandwidth-frugal, natively low-precision—which is precisely the profile that CPU-class and LPDDR-based systems serve best. DeepSeek’s adoption of the UE8M0 FP8 scale format “for next-generation domestic chips” made the co-design explicit [105].

The model layer, in playbook terms, is already in the late-war stage, though the ledger has to be read line by line rather than declared. Adoption: captured. Standards—weight formats, sparsity recipes, RL pipelines, quantization conventions, and now the reference workloads against which American accelerators are benchmarked in public—increasingly set in Chinese repositories.

Rent: destroyed at the volume end of the market, where Section 10.7 finds list prices converged and the Western incumbent priced underneath the Chinese frontier, but plainly not at the capability end, where the only published margin figure in the record is above 80%. What open weights cannot yet guarantee is the hardware to run and train them sovereignly—and this element closes with the serving half of that guarantee demonstrated at production scale and the training half resting on a press claim without a primary. That is the work of the next two elements.

Chapter 11

The Harness: The Layer the Contest Moved To

Every chapter so far has treated the model as the object of competition. This chapter argues that the object is moving, and that the move is the most strategically consequential development of 2026—more consequential than any individual release, because it changes what winning a release would even mean.

The claim in one sentence: *models are swappable; harnesses are programmed to*. A model sits behind an API call and can be exchanged for another with a configuration change—an entire product category, the LLM gateway, exists to make exactly that substitution trivial. A *harness*—the runtime that holds the agent loop, the tool schemas, the permission system, the memory and context management, the sandbox, the session store, the subagent orchestration—is something an organization writes code *against*. Its abstractions become the shape of the integration. Its execution semantics become the assumptions the tests encode. Its permission model becomes the artifact an auditor reads. If that is right, then the layer where switching is expensive is not the layer everyone is measuring, and a strategy that wins the harness can lose every benchmark and still take the ecosystem.

11.1 The layer acquires a name

A useful sign that an industry has discovered a layer is that its competitors independently start using the same word for it. Through 2026 all three of the leading agentic laboratories published engineering material on *harnesses* as such: OpenAI on harness engineering in an agent-first world, Anthropic on harness design for long-running application development, and—in August—DeepSeek, in a paper that cites both [489]. Three laboratories that agree on almost nothing agree that this is a distinct artifact with its own design discipline [V].

What is inside it is mostly not the model call. The most careful public anatomy of a production harness, an analysis of Claude Code by researchers at MBZUAI and UCL, puts the finding plainly: “the core of the system is a simple while-loop that calls the model, runs tools, and repeats. *Most of the code, however, lives in the systems around this loop*: a permission system with seven modes and an ML-based classifier, a five-layer compaction pipeline for context management, four extensibility mechanisms (MCP, plugins, skills, and hooks), a subagent delegation and orchestration mechanism, and append-oriented session storage” [495]. A widely repeated corollary—that only 1.6% of the codebase is AI decision logic—should be handled with more care than it usually receives: the paper cites it as a *community estimate* derived from line-counting a leak-derived bundle containing generated and minified code, and does not present it as a measurement [A]. The qualitative claim survives the quantitative one, and the qualitative claim is enough.

11.2 The harness is a larger economic lever than the model

The strongest evidence that this layer matters is not rhetorical but experimental, and as of this edition there is exactly one controlled study of it. Researchers at Writer held twenty-two audited evaluation tasks, six foundation models from five vendors across three capability tiers, and identical judges, prices, and prompts constant, and varied only the orchestration layer: a baseline production loop against a purpose-built harness [494]. Blended across models, the harness cut cost per task by 41%, tokens per task by 38%, and median latency by 44%, while raising quality per dollar by 82%, at maintained task completion.

The sentence that matters for this book is the comparison the authors themselves draw: *keeping any model and adopting the harness saved between 33% and 61%, while switching between the cheapest and most expensive model under baseline conditions saved only 36%* [V, single study]. Every model in the panel got substantially cheaper regardless of vendor or tier. If that result generalizes—and twenty-two tasks from the vendor of the winning harness is a caveat that must be printed beside it—then the harness is a *larger* lever on production economics than the choice of model, which inverts the entire public conversation about which layer is competitive.

There is a second, quieter consequence. A harness is also the instrument by which models are measured: the same model produces different results under different harnesses, and system-prompt budgets across harnesses vary by nearly two orders of magnitude [496]. Whoever owns the harness owns the conditions of measurement—which is worth remembering the next time a scoreboard is cited, including in this book.

There is a third consequence, and it compounds rather than competes with the first two. Chapter 12’s priced bill of materials (§12.9.4) shows a capacity-first local appliance running the largest open-weight model in existence at a little over a dollar per million output tokens, against twenty-five to thirty for the closed frontier—and, at realistic utilization, below a hyperscaler’s own marginal cost of serving the identical model. If that holds, model cost is heading toward a rounding error at the margin for an organization willing to run its own hardware, which does not shrink the harness argument, it *sharpens* it: the Writer study’s 33–61% saving from adopting a harness while holding the model constant was already larger than the 36% saving from switching models under a fixed harness; once the model’s own marginal cost approaches zero, the harness is not merely the larger of two levers, it is closer to the *only* lever left on production economics. An organization that owns both the cheap inference tier and the harness that drives it captures the saving twice, which is the strongest economic reason—stronger than the licence, stronger than the star count—to expect DeepSeek Harness and its peers (Table 11.1) to be paired with exactly the kind of local capacity Chapter 12 prices, rather than deployed against a metered API.

11.3 What China shipped in August

On 13 August 2026 DeepSeek released *DeepSeek Harness*, an MIT-licensed agent harness whose organizing principle is that *everything is a plugin*: models, tools, skills, sessions, sandboxes, filesystems, loops, orchestration, and interface are all components, with no privileged core to patch [486, 487]. It is model-agnostic by construction—Anthropic, OpenAI, Bedrock, Azure, Gemini and DeepSeek endpoints are all supported—it speaks MCP through a client plugin (tools only; resources and prompts are not bridged), and it ships subagent bridges that let it drive Claude Code and OpenAI Codex as subordinate components [488, 491]. Within roughly forty-eight hours the repository passed 100,000 GitHub stars [V]. For scale: LangGraph, the most-downloaded open agent framework, has accumulated roughly 34,000 stars over more than a year.

The accompanying paper is stranger and more interesting than a release note. *A Programming Paradigm for Spatiotemporal Composability*, an eighty-eight-page draft from Peking University and DeepSeek, is a programming-languages paper, not a machine-learning one [489]. Its subject is

a meta-framework called Cordis, and its two contributions are *temporal composability*—unloading a component provably returns the system to the state it would have occupied had the component never loaded, via *reversible effects* in which every context mutation carries a runtime-tracked inverse—and *spatial composability*, in which components declare dependencies and are reactively notified as those dependencies change, via *reactive coefficients*. The lineage is Moggi on monads, Plotkin and Power on algebraic effects, Petricek and Orchard on coefficients; the novelty claim is lifting these static type-system ideas into runtime mechanisms. Two honest qualifications belong here, both stated by the authors: the paper is a self-hosted preprint “under active revision” rather than a refereed publication, and its validation case study is observational rather than controlled. It is also worth being precise that the paper describes *Cordis*, the substrate; the harness is a product built on it, and applying the formalism to self-evolving harnesses is named as future work, not as an achieved result.

The provenance is where the strategic content sits. Cordis is not new and did not originate at DeepSeek: it is the kernel of Koishi, a four-year-old Chinese open-source chatbot framework with over four thousand community-contributed plugins, and the paper’s first author is Koishi’s creator [489, 493]. DeepSeek did not build a plugin ecosystem; it recruited the maintainer of one, formalized his runtime with a university partner, and shipped it under MIT. That is an ecosystem-acquisition strategy rather than a model strategy, and it is this book’s thesis executed one layer up.

The complements logic is exact, and it is the mechanism of Chapter 2 rather than a new one. DeepSeek gave the harness away under the most permissive licence available *on the same day* it raised API prices for V4-Pro and V4-Flash and introduced peak/off-peak pricing [492, 509]. Commoditize the complement, monetize the scarce layer, and—because the harness is model-agnostic—ensure that the commoditization damages competitors’ pricing power at least as much as your own. The *South China Morning Post* read the move as “shifting the global competition from model intelligence to one that focuses on how seamlessly an AI agent can connect to real-world software” [490], which is the thesis of this chapter stated by a newspaper.

The next day supplied the model-side counterpart, and it came from an unexpected direction. On 14 August Xiaohongshu’s dots studio released *dots3-note Preview* under Apache 2.0—280B parameters with 16B active, text, image, video and audio input, a 512K context—and the release note’s centre of gravity was not a benchmark but a training method: TEMPO, “test-time-scaled value estimation with macro-step policy optimization,” a reinforcement-learning scheme for agent trajectories running ten hours and more [662] [V for the release; A for the method, which awaits its technical report]. The problem it attacks is credit assignment over long horizons: an agent that works for ten hours and receives one reward at the end cannot learn which of thousands of decisions mattered. TEMPO divides the trajectory into macro-steps and, at each boundary, has the *same* model switch from actor to critic—reasoning, using tools and spending test-time compute to estimate the expected value of the state it has reached—so that dense intermediate signal replaces a single terminal one; the premise, stated in the release material, is that judging progress is easier than making it [663] [A]. Two things follow for this chapter. A model trained to evaluate its own intermediate states is a model trained *for* the harness loop, and the harness that logs, sandboxes and replays those states is where such training data comes from; the layer and the model are being co-designed, in the same week, by two Chinese laboratories. And the economics sharpen the cost argument of Section 12.9.4 rather than softening it: an agent that makes thousands of model calls per task multiplies whatever price gap exists between a closed API and an open model run locally, so the shift from chat to long-running agents makes the open commons’s cost advantage more consequential, not less. Huawei’s Ascend stack published day-zero deployment support for the model [665] [A].

DeepSeek is not alone, and the inventory matters more than any single entry (Table 11.1). What the table shows is a bloc shipping runtimes, sandboxes, command-line harnesses, and agent platforms—infrastructure rather than weights—across at least four laboratories and one state

initiative.

Table 11.1: Chinese agent-infrastructure artifacts as of 9 September 2026. The pattern is runtimes and protocols, not only models.

Artifact	Organization	Type	Licence	Grade
DeepSeek Harness (dsh)	DeepSeek	agent harness, plugin-native	MIT	V
Cordis	DeepSeek / PKU	meta-framework + formal paper	open	V
OpenSandbox	Alibaba → neutral org	sandbox runtime for agents	Apache 2.0	V
Qwen Code, Qoder CLI	Alibaba	CLI coding harnesses	Apache 2.0	A
Qianwen Open Platform	Alibaba	third-party agent platform	proprietary	V
AgentLoop / Teams / Run	Alibaba Cloud	agent operations stack	proprietary	V
Kimi Code CLI	Moonshot	CLI harness, ACP-compatible	MIT	A
dots3-note Preview + TEMPO	Xiaohongshu	agent-trained open model; long-horizon RL method	Apache 2.0	V/A
CodeBuddy Code	Tencent	CLI harness, ACP-compatible	—	A
Dify	LangGenius	agent/RAG platform, 144k stars	open	A
Agent interoperability initiative	CAC (state)	standards-diplomacy instrument	—	V

11.4 The standards layer, and who is not in the room

The governance picture is the sharpest geopolitical fact in this chapter, and it is a fact about membership lists. Anthropic donated the Model Context Protocol to the Linux Foundation’s newly created Agentic AI Foundation in December 2025, co-founded with Block and OpenAI, with Google, Microsoft, AWS, Cloudflare and Bloomberg as supporting members; MCP at donation reported more than ten thousand active public servers and over ninety-seven million monthly SDK downloads [498]. Google’s Agent2Agent protocol, donated to the Linux Foundation, passed 150 organizations in its first year [499]. In August 2026 a further standard, Agent Plugins 1.0, arrived with a technical steering committee drawn from Amazon, Cursor, Microsoft, OpenAI, Vercel and Google [500].

No Chinese company is a named member of any of them [V, by absence]. And on 17 July 2026 the Cyberspace Administration of China released a “Global Cooperation Initiative on Agent Mutual Trust, Interconnection, and Interoperability,” calling for “open, non-discriminatory, and transparent international interoperability standards for agents” and for shared agent identity and interface-compatibility infrastructure—while naming no existing protocol, neither MCP nor A2A [502]. Two weeks later DeepSeek shipped an MIT harness that reached a hundred thousand stars in two days.

Chapter 20 argued that standards are where governance becomes irreversible, and gave China’s standards strategy as the arena the West systematically under-weights. This is that argument arriving at the agent layer, and the structure is the familiar one from Part I: the incumbent bloc convenes a consortium of incumbents; the challenger declines to join, ships a permissively licensed implementation, and bids for the installed base directly. Whether that

succeeds is unknown. That it is the same play, run again, on a layer nobody was watching a year ago, is not.

11.5 The strongest case against this chapter

This book’s discipline requires the counter-argument at full strength, and here it is unusually strong—strong enough that the thesis must be narrowed rather than merely defended.

First, the interface is actively being commoditized, right now, by the incumbents. Agent Plugins 1.0—announced on 6 August 2026—is a vendor-neutral package format for distributing skills and MCP servers across agent clients “without requiring file rearrangement or manifest rewriting,” governed independently and developed in public [500]. Six of the largest firms in the industry are standardizing precisely the interface this chapter calls sticky. A thesis that harness lock-in is structural has to explain why the participants are dismantling it.

Second, the architecture is not a secret. Four independent teams—Claude Code in TypeScript, Codex CLI in Rust, Aider in Python, and OpenClaw—converged on the same structural pattern: an outer loop, tool execution, result capture, staged permissions [495]. Convergent design is evidence that the architecture is obvious, and obvious architectures are not moats.

Third, the migration evidence points to weeks, not quarters. The best-documented public account of leaving a framework—Octomind’s departure from LangChain—reports rigid abstractions, debugging overhead, and a loss of iteration speed, but quantifies no switching cost [497]. More pointedly, OpenAI’s Codex shipped a one-line importer for Claude Code configuration: the cost was real enough to automate, and small enough to be automated in one command [496]. Skills appear to transfer better than the thesis wants: one reported comparison found a skill trained under one harness scoring 80.4 when transferred versus 81.8 trained natively [A, one skill, one benchmark pair].

Fourth, the builder of the most successful harness in the world denies the premise. Boris Cherny, the creator of Claude Code, has described the harness as “the thinnest possible wrapper” and located “all the secret sauce” in the model [496]. Against him, an engineer at Obvix Labs supplies the opposing epigram—“the model is interchangeable, the harness is not.” Both cannot be right, and the person with the most direct evidence is on the other side of this chapter’s argument.

Fifth, and most awkward for a chapter about Chinese lock-in strategy: Alibaba is doing the opposite. OpenSandbox—the agent execution sandbox released in March 2026, not August—was moved out of Alibaba’s GitHub organization into neutral governance, carries an OpenSSF best-practices badge, is listed on the CNCF landscape, and lists Claude Code, Gemini CLI, LangGraph and Google ADK among its supported integrations [503]. That is a firm commoditizing the execution layer beneath every harness, American ones included, and donating it away—the Kubernetes playbook, not a capture play. Any account that treats the Chinese move as uniformly lock-in-seeking is contradicted by its largest participant.

Sixth, and decisively for intellectual honesty: nobody has measured the thing this chapter asserts. There is no published, methodologically sound study comparing the cost of switching harnesses against the cost of switching models. The Writer experiment measures operating economics, not migration cost. Enterprise surveys establish that multi-model deployment is now the norm—over three-quarters of teams run multiple models—but do not measure framework-switching rates at all [504]. This chapter therefore advances a *well-motivated hypothesis with strong circumstantial support*, and it should be read at that grade [S], not as a finding.

11.6 The thesis, narrowed to what survives

What survives the counter-argument is not the strong claim but a sharper one. The technical interface *is* commoditizing, quickly, and MCP plus Agent Plugins will likely make tool definitions

and skill packages portable across clients. Lock-in does not therefore disappear; it *relocates*—which is the rent-relocation mechanism of Section 2.3 applied to infrastructure rather than to weights. It relocates to three places the standards do not touch.

To distribution. Agent Plugins standardizes packaging, not distribution. Each platform still controls its own marketplace, so an author of a valuable skill is not merely building to a standard—they are choosing which gatekeeper reaches enterprise buyers. The switching cost becomes commercial rather than technical, which is harder to automate away than a config importer. As one practitioner put it, if useful behavior ends up inside vendor namespaces, “the standard will look open while lock-in just moves elsewhere” [501].

To the evaluation substrate. If the harness determines measured model quality, then the harness that an ecosystem standardizes on defines the scoreboard that ecosystem believes. DeepSeek used its harness internally to benchmark its own coding models before releasing it [A]. Owning the instrument is a durable position that no packaging standard addresses.

To certification, which is where this argument meets Chapter 18. For unregulated development, the transfer evidence suggests re-validation is cheap. For regulated deployment it is not, and for a reason that has nothing to do with performance regression: the artifact an auditor examines is the permission model, the sandbox boundary, the audit log, and the tool-invocation trace—all of which are harness properties, not model properties. Re-certifying an agent system after a harness migration is paperwork, evidence, and elapsed calendar time at an assessor, and the certification ladder of Section 18.4.2 prices none of it against the model. This is the version of the switching-cost argument that most plausibly holds at scale, and it is testable: if harness migrations in regulated sectors run materially longer than in unregulated ones, the mechanism is real.

Three implications follow for this book’s propositions. P1 (adoption) acquires a second front: adoption of Chinese *infrastructure* is now measurable independently of adoption of Chinese weights, and the star counts say it is moving faster. P3 (rent destruction) is reinforced—a hundred-thousand-star MIT harness released the same day as an API price increase is the rent-relocation thesis performed rather than argued. And the book’s central asymmetry sharpens into its most uncomfortable form: *a bloc can lose the model contest and still win the ecosystem*, because the thing an enterprise cannot cheaply change is not which model answers the call but which runtime made the call, logged it, sandboxed it, and can prove to an auditor what it did. If American models remain the best—and on today’s independent indices they are—and the world’s agents are orchestrated by runtimes shipped from Hangzhou and Beijing under MIT licences, then superiority at the model layer buys less than its owners expect. That is the claim this chapter registers, at scenario grade, with its falsifiers stated in Appendix J.

Chapter 12

The Compute: RISC-V, Accelerators, and the Memory Arbitrage

The second element is the machine the models run on. The Western AI machine is a proprietary stack—x86/Arm hosts, NVIDIA accelerators, CUDA, NVLink, HBM—each layer a licensed chokepoint. China’s counter-design, now visible end to end, is an *open* machine: open instruction sets, multi-vendor domestic accelerators unified by open interconnect and open software, and a memory system that trades scarce bandwidth for abundant capacity. In June 2026 this ceased to be a projection: an all-CPU, GPU-free Chinese machine took the No. 1 position on both of the world’s canonical supercomputing benchmarks. The claim this book draws from that is narrower than the headlines suggest and, properly stated, more interesting: not that the GPU is obsolete—it is not, and Section 12.5 says so plainly—but that the CPU corner of the design continuum can reach performance parity while carrying a software and application-ecosystem advantage the accelerator corner cannot easily replicate. This chapter treats the layers in turn—the instruction set, the machine that proved the architecture, the template it establishes, the accelerator fleet, the memory arbitrage, and the quiet defection of the Western ecosystem itself to the open ISA.

12.1 State-mandated architectural independence: RISC-V

China is actively organizing its domestic industry to circumvent the Western ISA monopolies—Arm and x86—that form the bedrock of Western computing. In March 2025, eight government bodies including MIIT, the CAC, and MOST were reported to have jointly drafted the first national policy guidance for RISC-V adoption nationwide—the only source is an anonymous-sourced Reuters exclusive describing a document in preparation, and no evidence has appeared that it was ever formally issued [106] [A]; the 15th Five-Year Plan (2026–2030) mandates “extraordinary measures” for integrated circuits and basic software self-reliance, with accelerated deployment of the open RISC-V instruction set as a core mechanism [107]. Because RISC-V is an open standard, its deployment insulates Chinese silicon design from sanctions at the ISA layer permanently: there is no license to revoke.

The execution is characteristically whole-of-nation. The Chinese Academy of Sciences fields one of the world’s largest RISC-V development organizations—over 600 hardware and 400 software researchers—whose open-source XiangShan cores now benchmark at Arm Neoverse-class performance (Kunminghu: 6-wide decode, four RVV 1.0 vector units, ~15 SPECint/GHz), with the openRuyi OS as the first RVA23-native software platform [108, 109]. Alibaba’s T-Head ships the XuanTie C930 server-class RISC-V CPU; SpacemiT’s K3 pairs an eight-core RISC-V cluster with a 60-TOPS AI engine [110, 111]. Within RISC-V International (Swiss-incorporated precisely to be sanction-proof), 12 of 24 premier members are Chinese [112]. The RISC-V Summit China 2026 (Shenzhen, October 18–20) is projected to host over 450 exhibitors and 3,600 industry professionals, its call for topics explicitly focused on ISA standardization of matrix and vector

extensions across the full AI computing stack [113]. Those extensions—the IME, VME, and AME task groups—are intended to give the open ISA native matrix operations, and Chinese vendors are the largest bloc shaping them [114]. Their schedules and roles must be stated separately, because they are commonly presented as sharing one near-term ratification horizon and they do not: IME (integrated matrix extension—matrix operations executed within the vector unit) targeted ratification in August 2026; AME (attached matrix extension—a tile engine with its own architectural state, the closest analogue to Arm’s SME and the extension class the accelerated-CPU template of Section 12.3 actually requires) carries a board-ratification target of April 2027; VME (vector–matrix interaction) carries a board-ratification target of November 2026 [468, 469]; Section 12.1.1 records the state of the ratified set at the currency date, which is empty, with the qualification that IME’s target fell thirteen days before it and no announcement had appeared. And ratification is only the first rung of a maturity ladder the strategy must climb end to end:

specification → hardware → compiler lowering → kernel library
→ framework integration → fleet validation,

with the AME plan itself treating simulators, ABI definition, GCC/LLVM support, and basic kernels as staged milestones and optimized software as an explicitly later problem [469]. The RISC-V accelerated-CPU path is therefore real but dated: the ISA foundation completes across 2026–2027, and the software rungs above it are measured in further years—which is exactly the Fugaku-calibrated interval Section 12.2.1 describes, now with its clock start pinned to primary-source schedules. Vendors such as EVAS Intelligence are already shipping full-stack RISC-V AI “supernode” systems—a self-developed cloud AI chip with block-quantized FP8, 3.2 TB/s per-chip interconnect, and claimed scaling to 10^5 -card clusters [115]. Two billion RISC-V chips have shipped; current forecasts project 36 billion units annually by 2031, with AI-accelerator SoCs alone accounting for roughly 70% of RISC-V silicon revenue [198].

The strategic reading: the ISA layer is being moved from the rent column to the commons column, exactly as model weights were in Chapter 10. What the ISA strategy needed, as of 2025, was an existence proof that a domestically fielded, non-x86, non-NVIDIA machine could stand at the absolute top of world computing. It arrived ahead of schedule—on Arm rather than RISC-V, which, as Section 12.3 argues, is precisely the point. Two concessions have to be entered before that machine is described, because together they bound how fast the open ISA can receive it: the arithmetic the ISA has not yet standardized, and the software that does not yet pass its own tests.

12.1.1 The unresolved centre: four proposals for one function

The standards body’s own August 2026 status report is the best available account of where the open ISA is strong and where it is not, and it should be read in that order. Strong: RVA23 was ratified in October 2024 with vector *and* hypervisor support mandatory rather than optional, RVA23.1 is near ratification, and the Server Platform specification—the boot, firmware and discovery layer on which any rackable machine depends—was ratified in May 2026 [674] [P]. That is a fast-converging platform layer, and Section 12.10 shows what it bought. Not strong, and this is the sharpest single weakness in the open-ISA position: *no matrix extension has been ratified at all*. What exists instead is one fast-track dot-product proposal and three competing task groups pursuing three different answers to one question—where the matrix state lives.

Four proposals for one function, with ratification targets spanning August 2026 to April 2027; IME’s own target date of 27 August 2026 fell three days after the conference update this section draws on and a fortnight before this edition’s currency date; no ratification announcement had appeared by 9 September 2026, and this book records the gap rather than closing it [A]. What is certain is that at the update presented on 24 August the ratified set was empty [674]

Table 12.1: RISC-V matrix arithmetic as of 9 September 2026: four proposals, one function, and no ratified matrix extension announced—with the qualification that IME’s 27 August target fell a fortnight before this edition’s currency date [674, 114, 468, 469].

Proposal	Where the matrix state lives	Status
Zvbdota	not a matrix extension: a vector dot product	fast-track proposal
IME	inside the vector registers; no new architected state	task group; 27 Aug. 2026 target, no announcement by 9 September 2026
VME	vector registers plus a new accumulator register file	task group; Nov. 2026 target
AME	separate matrix register state; the SME analogue	task group; April 2027 board target

[P]. The comparison that matters is not with a rival open standard but with the shipping alternatives: NVIDIA offers one tensor-core datapath and Intel one AMX, and a buyer of either knows what instruction its kernels will be written against for the life of the machine. An HPC or AI procurement that buys RISC-V in 2026 is buying vector-only arithmetic, with the matrix path to be settled later by a process whose three candidates are not source-compatible with one another. This is the fragmentation cost of openness, and it is not this book’s characterization: RISC-V International’s own material now carries “optimization guidance” options—a way of encoding *performance expectations* rather than function into a profile—which concedes that a functional-only ISA string does not tell a vendor how to ship a competitive core, nor a compiler which of two conforming machines it is actually targeting [674].

The consequence for this chapter is specific, and is better stated against its own argument than around it. The accelerated-CPU template of Section 12.3 requires an AME-class engine with its own architectural state; that is the proposal with the *latest* target of the three and the largest software surface, because new architected state means new context-switch, ABI, debugger and unwinder work before a single kernel is tuned. The maturity ladder above therefore does not begin in 2026 for the extension the template actually needs—it begins when that extension ratifies, and every rung is measured from there. China’s position inside this process is genuinely privileged, since Chinese vendors are the largest bloc in the rooms where the choice is made; but a bloc large enough to shape three proposals is also a bloc large enough to keep any one of them from winning quickly, and nothing in the open-governance model settles the question as fast as a single vendor simply deciding it.

12.1.2 The readiness number

The second concession is measured, published by a party with no incentive to flatter it, and it is the most useful single number in this chapter about how far the open ISA sits from the machine room. Canonical carried an RVA20 riscv64 baseline in Ubuntu from 20.04 LTS through 25.04, moved to RVA23 at 25.10, and ships 26.04 LTS with RVA23 as the long-term baseline—an unusually aggressive baseline move, and one that does not orphan the installed base, since RVA20 silicon retains a supported path on 24.04 LTS for up to ten years under Ubuntu Pro [675] [P]. The number is this. Canonical’s own figure is that roughly 95% of archive packages *build* on riscv64—and only **55.1%** of tested packages *pass*, against 92.3% on amd64 and 92.5% on arm64, on roughly one seventh the test coverage of those architectures [675] [P]. Canonical attributes the gap to emulation: the tests run on emulated hosts, because server-grade RVA23 hardware to run them on did not exist.

That is what “the software rungs are measured in years” looks like when somebody measures

it. It is also, precisely, the quantity the lifecycle decomposition of Eq. (12.4) calls $C_{\text{regression}}$ —the recurring cost of keeping a working stack working—expressed as a pass rate rather than as an adjective, and it is exactly the kind of series the AI-HPC lifecycle benchmark proposed in Section 12.7 would produce and MLPerf structurally cannot. Two qualifications keep it from being read as a verdict. First, the deficit is attributed to test infrastructure rather than to the architecture, which makes it the kind of gap that closes quickly once the infrastructure changes—and Canonical’s board farm runs on Alibaba and SpacemiT hardware, so Chinese silicon is already inside the Western distribution’s test loop. Second, and cutting the other way, RVA23 is a *richer* baseline than the arm64 target the same distribution ships against, which remains plain ARMv8-A with no mandatory optional extensions: a distribution compiled to RVA23 may assume vector and hypervisor support that its Arm build cannot. That is a real architectural advantage of the younger ISA, and it converts into a delivered one on the day the silicon arrives.

It arrived in the week before this edition’s currency date. SiFive announced the BigSky SF-2U870 on 24 August 2026: 32 P870-D cores at 2.0 GHz, 256 GB of DDR5-5600, 64 lanes of PCIe Gen5, in a 2U chassis, fully RVA23 compliant and shipping with Ubuntu 26.04 LTS and RHEL 10—the first rackable enterprise RISC-V server [676] [P]. What it is *not* deserves equal clarity: available in limited quantities, with no price, no SPEC result and no named customer tape-out. Futurum’s assessment is the one worth printing—BigSky is “primarily a software porting vehicle rather than revenue-generating hardware”—and that is exactly why it belongs here rather than in a product survey [676] [A]. It breaks the circularity the pass-rate number describes: the distributions could not test on hardware because no server-grade hardware was rackable, and the hardware was not orderable because the software did not pass. A porting vehicle a customer can rack today moves the multi-year software phase of every subsequent SiFive-IP-based custom SoC off the critical path—which is this book’s own claim that the software interval is the binding one, made by a vendor with money at stake. Section 12.10 takes up the other half of the same announcement, which is what BigSky was designed to sit next to.

12.2 LineShine and the LX2: the GPU-free machine takes No. 1

On June 24, 2026, at ISC in Hamburg, the 67th TOP500 list opened a new era twice over. LineShine (*Ling Sheng*, “intelligent splendor”), built by the National Supercomputing Center in Shenzhen, debuted at **No. 1 on HPL at 2.198 EFLOP/s**—the first machine ever to sustain two exaflops of FP64, the first Chinese No. 1 since 2017, and, most consequentially for this book, **a machine containing no GPUs whatsoever** [182, 183]. It simultaneously took **No. 1 on HPCG** (22.00 PF, ahead of El Capitan’s 17.41 and Fugaku’s 16.00)—the memory-bandwidth-bound benchmark—while posting 7.92 EF on mixed-precision HPL-MxP, at 42.2 MW and 52.1 GF/W, in the same Green500 band as the US GPU flagships [182, 185].

The processor behind it, the Armv9 **LX2**, is the architectural pivot of this chapter: 304 cores per socket, every core carrying both SVE vector and SME matrix engines, with 32 GB of on-package HBM at 4 TB/s plus a large DDR5 tier—60.3 TF FP64, 240 TF BF16, 960 TOPS INT8 per socket at ~690 W [183, 185]. It is Fugaku’s architectural lineage (A64FX: the first general-purpose CPU with on-package HBM) carried to its logical conclusion: the accelerator’s defining features—wide vectors, matrix engines, HBM, low precision—absorbed *into the CPU instruction set*. Precision about what this is: the LX2 is an *integrated accelerated processor*, GPU-free in product taxonomy rather than free of accelerator-class silicon. The claim it supports is correspondingly precise—not that accelerator silicon is unnecessary, but that the discrete accelerator *product category*, with the separate vendor, separate memory space, separate software monopoly, and separate margin structure attached to it, is an implementation choice rather than an architectural necessity. Its designer is undisclosed; analysts read Huawei’s fingerprints in the software stack, SMIC 7 nm-class as the process candidate “by elimination,” and Chinese

state media claim the machine runs on “China’s first domestically developed HBM”—a claim examined in Chapter 13 [186, 187].

The quantitative meaning of this machine for AI—not just simulation—is worked out in the study by Dongarra, Hoefler, and Matsuoka, “Do We Still Need GPUs?”, in its short form [184] and full-length companion [185], and the book adopts its findings as load-bearing. Four of them matter here. *First*, for the traditional HPL/HPCG workload space the question is settled on the public record [V]: a CPU machine now leads both the compute-bound and the bandwidth-bound FP64 benchmarks at GPU-class energy efficiency—while transformer training and high-concurrency serving remain, as the study itself insists, projections pending direct measurement [M]. *Second*, for frontier AI serving—tested at the hardest available point, a 1T-parameter MoE (Kimi-K2) at 256K context—decode reduces to a bandwidth roofline on which *an LX2 socket serves the same users per TB/s as a GPU*, and prefill reduces to stated, achievable thresholds on INT8 matrix efficiency and interconnect; the all-CPU machine lands at roughly $1.3\text{--}1.4\times$ a B200 fleet’s per-stream energy under central assumptions [185]. *Third*, training on the shipping part costs $4.5\text{--}8.5\times$ a GPU fleet’s energy per token at matched BF16—but every term of that deficit (the missing FP8 path, engine width, fabric ratio) is an implementation choice, not an architectural constant; an FP8-equipped successor cuts the central gap to $\sim 3\times$, and a part with GPU-parity matrix throughput trains at parity to first order. *Fourth*, and most generally: with HBM now common to both device classes, “CPU” versus “GPU” has collapsed into a single design continuum parameterized by die-area allocation between low-precision matrix throughput and high-precision generality—*either corner can be provisioned to do the other’s job* [184, 185].

Jack Dongarra’s own verdict, delivered alongside his technical report on the system, states the strategic conclusion this book would otherwise have to argue: China “can adapt to develop its own version of technology as good as—or maybe even better than—existing technology, despite US export controls” [188, 183]. Export controls denied China the discrete GPU; the response was to demonstrate, at No. 1 in the world, that the functional boundary between CPU and accelerator is an implementation choice—and then to choose an implementation with no controllable product category in it.


12.2.1 The proof boundary: what this machine establishes, and what it does not

That last sentence is the strongest one this book will make about LineShine, and it is worth stating exactly where its support ends—because the argument survives a much narrower reading than its enthusiasts give it, and is more useful stated narrowly.

The disciplined formulation is two-stage. *Stage one, established*: discrete-GPU packaging is not required for world-leading HPL and HPCG performance at competitive energy efficiency [V]. *Stage two, argued but not yet measured*: the same integration extends to AI serving, and—with different matrix and memory provisioning—to training [M]. LineShine does *not* yet validate large-transformer training efficiency, expert-routing throughput and tail latency, distributed optimizer scaling, collective performance under AI communication patterns, kernel coverage and software productivity, serving throughput at production concurrency, or reliability and checkpoint/restart behaviour across multi-week runs. HPL and HPCG are not transformer benchmarks, and no amount of eminence in the rankings makes them so. Figure 12.1 draws the line.

Having conceded that, three observations push back on the deflationary reading, and together they are why this book still treats the machine as a hinge rather than a curiosity.

First, the inference from stage one to the serving half of stage two is shorter than “modeled” suggests. Autoregressive decode over a sparse mixture-of-experts is memory-bound by an order of magnitude—its arithmetic intensity sits far below any current socket’s compute-to-bandwidth ridge—so the number of concurrent streams a socket serves is, to first order, a *ratio* of delivered bandwidth to per-user KV traffic, not a fitted performance model. What is genuinely modeled



Stage-1 claim (established): discrete-GPU packaging is not required for world-leading HPL and HPCG performance at competitive energy. **Stage-2 claim (argued, not yet measured):** the same integration extends to AI serving, and—with different matrix and memory provisioning—to training.

The book’s *strategic* claim requires only stage 1 plus the die-area allocation argument, which is independent of this machine.

Figure 12.1: What LineShine proves, and what it does not. The strongest form of the architectural argument is not that a CPU has been shown to train transformers—it has not—but that the functional boundary between CPU and accelerator has become an allocation choice, of which LineShine is the first at-scale demonstration in the direction the incumbent does not occupy. Stated in the two-stage form this book adopts; measured values from the June 2026 TOP500 submission and the accompanying technical report [182, 183], modelled values from the companion analysis [185].

is a single efficiency coefficient—written η throughout this book and in the claim register, and defined as the *fraction of a socket’s peak memory bandwidth that is actually delivered to the serving workload*, so that achieved bandwidth is η times nominal. The companion analysis sweeps η over 0.53–0.70 rather than asserting a value, and every serving conclusion in this chapter is reported across that band [185]. The prerequisite that a CPU can *drive* on-package HBM rather than merely attach it is demonstrated on deployed hardware in this exact lineage—STREAM parity with the contemporary GPU generation on A64FX in 2018, 450 GB/s measured per memory domain on the LX2, and a world-record HPCG that is precisely a delivered-bandwidth benchmark. Prefill, the compute-bound half, is reduced not to an assertion but to *stated numerical thresholds*—sustained INT8 efficiency above roughly 0.38, INT8-eligible fraction above 0.90, achieved all-reduce above 9–38 GB/s—which is the form a falsifiable claim takes.

Second, the software gap is a measured historical quantity rather than an unknown. The same lineage supplies the calibration: Fugaku led the benchmark lists from 2020, yet a performant AI stack for its architecture arrived years later. The correct prior for an architecture whose hardware prerequisites are demonstrated and whose software ecosystem is young is therefore not “unproven” but “capability present, stack lagging by a measurable interval”—which is exactly how the West should read the LX2, and exactly the interval during which a determined national programme closes it.

Third, and most importantly, the strategic claim does not rest on this machine at all. The load-bearing argument of Section 12.3 is the die-area allocation argument: once HBM is common to both device classes, what distinguishes a future CPU from a future GPU is how compute silicon is divided between low-precision matrix throughput and high-precision generality, with injection bandwidth provisioned in ratio. That argument is architectural, not empirical, and it is corroborated from three independent directions at once. LineShine approaches the continuum from the CPU corner. NVIDIA’s own product line approaches it from the GPU corner—Grace CPUs, LPDDR-based memory subsystems and SOCAMM modules, CUDA ported to RISC-V hosts (Section 12.10). And the hyperscalers’ custom accelerators approach it from a third

direction entirely: systolic and dataflow designs, none of them SIMT, converging on similar allocations without sharing a device category. Three actors with different constraints, different customers, and different starting points are converging on one design space. *A category that three independent parties are dissolving is not a category.*

There is also a fairness point worth naming, because it recurs whenever a challenger architecture appears. The demand that LineShine prove transformer-training competitiveness before the architectural argument counts applies a standard the incumbent never met: GPUs were adopted for deep learning because they were available, programmable, and fast enough, years before anyone demonstrated they were the optimal allocation—and their FP64 capability has since been substantially removed in the tensor-heavy parts, which is itself an admission that the allocation is a choice. The correct question is not whether the CPU corner has proven itself against a standard the GPU corner was never held to; it is which corner the workload of the next decade rewards, and Chapters 24 and 26 argue that long-context agentic scientific orchestration rewards the capacity-rich, generality-preserving corner more than the dense-FLOP one.

12.2.2 The AI proof suite: what would settle it

Because the second stage is contested, this book states what would resolve it. Table 12.2 specifies a measurement programme for LineShine or any LX2-class successor. It is deliberately constructed so that each row can be run by the machine’s operators and published without disclosing anything about the processor’s provenance, and so that a failure on any row is a genuine falsification of the corresponding claim in this book.

Table 12.2: Proposed AI proof suite for an LX2-class integrated accelerated processor. Each row states the claim at risk and the outcome that would falsify it.

Measurement	Claim at risk	Falsified if
MLPerf Inference (closed division), 1T-class MoE	serving parity per TB/s	delivered streams per TB/s fall materially below GPU submissions at matched precision
Long-context decode at 256K, full KV traffic	the bandwidth-roofline argument	achieved fraction of peak collapses under the serving placement (aggregate $\neq$ per-domain)
Prefill at production concurrency	the stated INT8 thresholds	sustained e_{INT8} stays below 0.38 with a mature kernel library
All-to-all expert routing at scale	the fabric-in-ratio discipline	achieved collectives fall below the stated break-evens on the declared layout
MLPerf Training, and a multi-week frontier-class run	the time-indexed training verdict	energy per trained token stays worse than $\sim 4\times$ a contemporary GPU fleet after an FP8-capable generation
Checkpoint/restart and failure-domain behaviour at full state	multi-week operability	mean time between interruptions times restart cost makes long runs uneconomic
Energy per <i>accepted</i> training token, and per solved agentic task	the whole economic case	the delivered-work comparison reverses the roofline comparison

A protocol note binds the whole table: every row should be run on both a dense 70B-class model *and* a 1T-class sparse model, at short and 256K context, across batch and sequence-concurrency sweeps, reporting median and tail latency, achieved bandwidth *by memory domain*

(aggregate figures hide placement failures), energy measured at the system boundary rather than the socket, and the exact parallel layout and software versions—because Eq. (12.2) makes every headline number conditional on precisely those choices.

Until those numbers exist, the honest position is the one printed in Figure 12.1: chokepoint decay is demonstrated in its strongest observed form for the HPL/HPCG workload space—the chokepoint not breached but *defined out of the machine*—and argued, with stated thresholds and a stated experimental programme, for the AI workload space that matters more.

12.3 The template: from LX2 to a RISC-V accelerated CPU, in two variants

This section presents an author-modeled architecture and strategy analysis [M/R/S], not a measured system or an announced product. The LX2 settles feasibility of the class; what generalizes is the *template*. The D–H–M analysis reduces the accelerated CPU to a short design sequence [185]: the memory system is chosen first, and its bandwidth fixes the (moderate) count of strong general-purpose cores needed to saturate it; scalable vector/matrix extensions then set per-core matrix throughput independently of core count—engine width becomes *implementation-selectable within physical and system constraints*, because the ISA no longer fixes it (“selectable,” not “arbitrary”: width remains bounded by register and tile-state bandwidth, issue capacity, die area, power, timing closure, delivered memory bandwidth, context-switch cost, thermals, and interconnect injection—a qualification the term must always carry); interconnect injection is provisioned in ratio to the matrix rate; and one address space unifies the whole. Nothing in that sequence is Arm-specific. SVE and SME are simply the first shipping instances of scalable vector and matrix ISAs; RISC-V’s ratified RVV 1.0 is the second scalable vector instance, and the AME extension now being standardized—with Chinese vendors as the largest bloc in the room, IME targeted for August 2026 and AME for April 2027 (Section 12.1)—is the open ISA’s SME equivalent. XiangShan-class open cores already carry four RVV units per core; an SME-class outer-product engine attached to such a core is an integration exercise of exactly the kind the LX2 has now de-risked, not an invention. For China the substitution is strictly preferable: the LX2 still sits on an Arm architectural license—a Western chokepoint in principle, however grandfathered in practice—whereas its RISC-V successor sits on nothing revocable at all. This is why the LX2’s arrival on Arm is not a detour from the RISC-V strategy but its proof-of-concept: the national ISA plan (Section 12.1) supplies the destination, and LineShine supplies the demonstrated microarchitecture to port to it.

The template’s second degree of freedom is the memory system, and here the two halves of this chapter join. Because the design sequence starts from the memory tier, one tape-out family yields two machines:

The inference-centric variant keeps the LX2’s own memory choice—on-package HBM at GPU-generation bandwidth—and provisions the matrix engines to the serving thresholds the D–H–M study states (sustained INT8 efficiency $\gtrsim 0.38$, achieved all-reduce cleared in ratio) [185]. This is the machine for the token factories of Chapter 10: decode at parity per TB/s with the GPU it replaces, prefill hidden inside the decode tier, FP64 retained for the converged AI-plus-simulation workload. Its binding input is HBM supply—the subject of Chapter 13.

The training-centric variant substitutes the HBM tier with massive LPDDR6X capacity—the arbitrage of Section 12.8—accepting the bandwidth knee in exchange for the 2.3 TB/socket that eliminates activation recomputation, dissolves the sub-256-node capacity wall, and sources its memory from the one tier where a domestic supplier (CXMT) sits at the world’s leading edge rather than four years behind it. Provision the same SME-class engines with

an FP8 path—the single highest-leverage addition the D–H–M training analysis identifies [185]—and the pretraining energy deficit falls from categorical to a low single-digit factor, paid for in memory that costs a fraction of HBM per bit and is not export-controllable.

Same cores, same engines, same fabric discipline, same software stack; only the memory tier and the compute-to-bandwidth ratio differ. The dual-variant template is the hardware mirror of the model strategy: the HBM variant serves what exists, the LPDDR6X variant trains what is next, and both are buildable on open IP, on domestic 7 nm-class logic, with domestic memory. A national program that executes it holds, in one product family, the replacement for the NVIDIA inference fleet *and* the decentralized pretraining capability of Section 12.8—with no Western license anywhere in the bill of materials.

One negative result bounds this template usefully, because it establishes that nobody else is building it. Every datacenter CPU presented in 2026 shipped *uniform* performance cores—Intel’s Diamond Rapids at 256 identical P-cores, with its efficiency-core line moved out to a separate product; Fujitsu’s MONAKA at 144; Arm’s AGI at 136; NVIDIA’s Vera at 88; IBM’s Z at 11—and agentic orchestration was given as the justification in near-identical language across vendors [672] [P]. What those cores do *not* carry is the other half of this section’s template. Vera has no NPU and no matrix engine at all, because its matrix unit sits at the far end of a 1.8 TB/s coherent NVLink-C2C link to a GPU; Arm’s AGI has no SME, no matrix unit and no quoted AI throughput; Fujitsu defers SME2 to MONAKA’s 1.4 nm successor. Only Intel ships a matrix unit inside the socket, AMX with FP8. And no CPU in the 2026 server track carries HBM. The Western server CPU of 2026 is therefore converging on the *opposite* allocation from the LX2: maximum general-purpose throughput for orchestration, with every matrix operation and every high-bandwidth byte deliberately placed on the far side of a link. Two things follow. The LX2 is more singular than the continuum argument by itself suggests—not one instance of an industry trend but the only shipping large-scale processor that took the integrated corner seriously. And the memory arbitrage of Section 12.8 has nowhere to live in a Western socket: if the matrix engines are across a link, the capacity must be provisioned against the accelerator, in the tier and the packaging the accelerator vendor controls. That is a structural reason the capacity-first move is available to an integrator who owns both ends of the link and awkward for one who owns neither.

12.4 The accelerator fleet and the unified ecosystem

12.4.1 Expelling NVIDIA, adopting the fleet

The accelerator story of 2023–2026 is a pincer: US export controls squeezed NVIDIA’s China products from above while Beijing squeezed demand for them from below. The A800/H800 were banned in October 2023; the H20 in April 2025 (a \$4.5B write-off), un-banned in August 2025 in exchange for an unprecedented 15% revenue share to the US Treasury—whereupon the CAC summoned NVIDIA over alleged backdoors, banned major platforms from buying, and regulators ordered that new state-funded data centers use domestic AI chips exclusively, with provinces cutting electricity tariffs up to 50% for domestic-chip facilities [116, 117, 118]. NVIDIA’s China data-center revenue went to approximately zero in its Q2 FY2026; Jensen Huang’s own scoreboard: “from 95% market share to near zero,” followed by his November 2025 statement that “China is going to win the AI race,” walked back within hours [119, 120].

Into the vacuum, the fleet. Huawei’s Ascend line shipped ~507,000 units in 2024 and ~805,000 in 2025 (~650k of them 910C-class), with a disclosed three-year public roadmap—950PR/950DT in 2026 with Huawei *self-developed* HBM (HiBL 1.0, HiZQ 2.0), 960 in 2027, 970 in 2028 [121, 122]. Cambricon’s revenue grew 453% in 2025 to its first annual profit, targeting 500,000 accelerators in 2026; Moore Threads’ Shanghai IPO popped ~400% on debut; Biren listed in Hong Kong; Baidu lit a 30,000-chip Kunlun P800 cluster; Alibaba’s T-Head PPU matched the H20 closely

enough that China Unicom deployed 16,384 of them in a single facility, and a 10,000-chip cluster of Alibaba’s “Zhenwu” followed in April 2026 [123, 124, 125, 126]. In May 2026, for the first time, nine domestic AI chips from seven vendors entered the “secure and reliable” government procurement catalogue [127]. Domestic chips took roughly 41% of Chinese AI-chip shipments in 2025 (~1.65M units) and are projected to pass 60% in 2026 [128]. China’s aggregate intelligent-computing capacity reached 2,185 EFLOPS by mid-2026 at 71% utilization [129].

The fleet’s first order at the scale of a leading laboratory arrived in the first week of September, and it is worth reading for its partition rather than its size. Bloomberg reported on 4 September, from anonymous sources, that DeepSeek plans to order at least 160,000 Ascend 950DT accelerators for a gigawatt-scale *inference* site at Ulanqab in Inner Mongolia, at an implied cost of roughly \$2.6B—a derived figure, not a disclosed one—with delivery paced by Huawei’s and the domestic HBM chain’s capacity and expected to take up to a year; training, on the same reporting, remains on NVIDIA hardware [755] [A]. The reading this chapter draws is precise. The largest domestic laboratory is committing its serving fleet to domestic silicon at gigawatt scale while keeping its training runs on the incumbent’s, which is exactly the served-not-trained partition that the proof boundary of Section 12.2.1 turns on: stage one of the integration argument—domestic parts carry the memory-bound decode workload at production concurrency—is being bet on with real money by the customer best placed to know, and stage two—that the same parts can carry a frontier pretraining run—is being conspicuously not bet on by the same customer, in the same order. It is also a datum on the memory concession of Section 12.4.4: an order whose delivery clock is set by HBM rather than by logic is the Ascend ramp’s binding constraint stated as a lead time. And it has a consequence outside this chapter. One customer absorbs a year of output. The export capacity that Chapter 22 goes looking for—Ascend for Cairo, for Kuala Lumpur, for Rio—does not exist yet, because the home fleet is consuming the production run, and a reported mandate that state-funded compute run at least 80% domestic silicon points the next run the same way [770] [C]. The counter is that DeepSeek’s previous Ascend engagement, in 2025, ended in failed training runs and a reversion to NVIDIA [104] [A]—the software-maturity cost that Section 12.7 prices; a year of delivery is also a year of CANN hardening on the most demanding serving workload in the country, and the order is the strongest available evidence that the laboratory judged that cost survivable the second time.

The other side of the same partition is the compute DeepSeek and its peers still rent abroad, and the closing fortnight also moved that. *The Information* reported on 28 August that the Commerce Department is drafting a rule to block Chinese firms from renting NVIDIA compute in third countries, reportedly triggered by Moonshot’s training of Kimi K3 on servers in Thailand; the draft may reach industry in September, its legal authority is contested, and it had not been published at 9 September 2026 [756] [A]. The dashboard in Appendix J carries offshore leasing as an indicator, with the annotation that closure by rule would be the highest-impact tightening available to export policy; that tightening is now reported to be in draft. Two things belong beside it. The reporting records the authority as disputed, and a rule whose object is a tenant in a third country has to be enforced against the vendor and the host rather than the tenant—which is why this book grades it as a draft and not as a control. And if it issues, it selects for precisely the partition the Ulanqab order already shows: inference at home on domestic silicon because it can be, training abroad on NVIDIA because it must be, and the second half of that sentence is the half a remote-access rule is written against. Section 10.3.5 in Chapter 10 carries the argument that the dependence is transitional; a rule of this kind would compress the transition rather than settle it.

12.4.2 The unified interconnection ecosystem

The fleet’s weakness—fragmentation across a dozen incompatible vendors—is being engineered away by mandate and by open specification. MIIT policy directs the creation of a unified AI-chip computing interconnection ecosystem, so that domestic accelerators from competing vendors

interoperate within shared clusters (the Computing Power Interconnection action plan; the “AI+” implementation documents) [130]. Huawei supplied the concrete substrate by *open-sourcing* the crown jewels: the UnifiedBus 2.0 interconnect specification, released openly so that “10,000+ NPUs work as one machine,” with 300+ Atlas 900 SuperPoDs already deployed on UB 1.0; the full CANN software stack (its CUDA analogue, seven years in development), open-sourced in August 2025; MindSpore; and the complete SuperPoD reference architecture [131, 132]. The announced Atlas 950 SuperPoD couples 8,192 Ascend 950DTs into 8 FP8-EFLOPS with 16 PB/s of interconnect; the 960-generation SuperCluster targets over one million NPUs [122]. One must apply the usual discount to vendor roadmaps—but the *structure* is the point: where the West’s answer to NVLink is a consortium specification (UALink) that excludes China—and which, as Section 12.4.3 shows, arrived in 2026 not as a fabric of its own but as a protocol carried over merchant Ethernet—China’s answer is an open specification with state-mandated adoption. This top-down standardization actively dismantles the need for proprietary software stacks like CUDA, enforcing interoperability across the domestic supply chain—the Top Runner program, replayed at the cluster level. LineShine’s own LingQi fabric—1.6 Tb/s injection per node, ≥ 3.5 Pb/s bisection—demonstrates the discipline at full machine scale [183].

12.4.3 Ethernet takes the scale-up domain, and the standards body loses the wire

The Western half of this contest moved in 2026, and it moved toward the same principle—open, multi-vendor, merchant—by an entirely different route. Scale-up, the load/store domain inside which accelerators address one another’s memory, has been NVLink’s private territory and the single largest source of lock-in in the stack, because it is the layer at which a buyer’s accelerator choice becomes irreversible. The 2026 record shows it being taken by Ethernet.

Table 12.3: Scale-up fabric by part, 2026 [P, vendor-presented]. Among the parts presented with an off-package scale-up fabric, the Ethernet-class designs are the clear majority; NVIDIA’s NVLink is the only proprietary scale-up fabric a third party can buy [681].

Part	Scale-up fabric	Domain
AMD MI455X	UALink over Ethernet on merchant 512-port 200G switch ASICs; no scale-up NIC in the tray	72 GPUs, load/store
Microsoft Maia 200	one Ethernet protocol from chip outward	6,144 accelerators, two-tier Clos
SambaNova SN50	ten integrated 800G Ethernet ports per socket	64 sockets, two merchant switches
Meta MTIA 400	RoCE on dedicated I/O chiplets	rack scale
NVIDIA	NVLink 6, proprietary, TFLOPS in-network compute	130 rack scale

NVIDIA alone retains a proprietary scale-up fabric, and pays for it with in-network compute nobody else has to match [681] [P]. AMD’s claim is the one to test: that NVLink-class scale-up semantics—a coherent load/store domain, not a message-passing network—can be built out of commodity Ethernet switch silicon. If that holds at production scale it removes the largest lock-in in the stack, not by legislating against it but by making the alternative a line item on a merchant price list. Microsoft’s deck makes the reasoning explicit in a way vendor material rarely does: the slide is titled “Why Unified?” and lists a standard switch ecosystem as the reason [681].

The negative result in the same track is more instructive still, and it states in its sharpest available form a pattern this book keeps arriving at from different directions. Every mechanism the Ultra Ethernet Consortium spent three years arguing for shipped in production silicon in 2026—packet spraying, out-of-order placement with in-order delivery, selective retransmission, receiver-driven credit congestion control, packet trimming, programmable transport. Not one of them shipped as UET. They shipped as *RoCEv2 extensions*: Broadcom brands its transport eRoCE, describes it as “MRC++,” and preserves the verbs ecosystem so that NCCL, RCCL and MPI run unmodified—“the ecosystem cost is zero.” Neither Broadcom’s nor NVIDIA’s deck presents a UET transport or a libfabric stack, and NVIDIA’s networking deck does not mention the consortium at all [682] [P]. *The consortium standardized the mechanism set and lost the wire format.*

That is the finding, and it cuts against the reflex that a good open specification defeats a proprietary incumbent: the installed base beats the standards body, and the winning move is to smuggle the new mechanisms in under the old ABI rather than ask the world to relink. Two corollaries follow for the Chinese position, and they run in opposite directions. The favourable one: UnifiedBus’s advantage over UALink was never that it is the better document—it is that it arrives with a mandated installed base attached, which is the only thing that has ever settled a fabric war. The unfavourable one: the same logic is precisely what protects CUDA, and it is why $C_{\text{migration}}$ in Eq. (12.4), rather than any specification, is the quantity a challenger actually has to attack. An open standard that requires the customer to rewrite the layer above it has already lost to a proprietary one that does not.

12.4.4 The incumbent roadmap, and what it does and does not settle

The obvious objection to everything in this chapter is that NVIDIA is not standing still. Its published cadence runs Blackwell to Rubin—in full production in the second half of 2026, at 288 GB of HBM4 and roughly 13 TB/s—to Rubin Ultra, to *Feynman* in 2028 [510]. What is officially known about Feynman is short and worth stating exactly, because supply-chain reporting around it is abundant and mostly unverifiable: at GTC 2026 NVIDIA named the generation, placed it in 2028, and disclosed its companion parts—a new GPU, the Rosa CPU, an LP40 inference processor, BlueField-5, CX10, and Kyber racks in *both* copper and co-packaged-optical scale-up variants, the first time NVIDIA has committed to CPO at rack scale [511]. Everything beyond that—the process node, the memory generation, the interconnect bandwidth, the rack power—circulates as supply-chain rumor, and this book declines to print it as fact [R].

Four observations follow, and they do not point where the objection assumes.

First, the roadmap is slipping on packaging, not on silicon. The Kyber rack for Rubin Ultra reportedly moved from 2027 to 2028 on a signal-integrity problem in a 78-layer orthogonal midplane, with the interim two-rack configuration cancelled after customer resistance; NVIDIA’s public response was that its roadmap is intact [512] [A]. Feynman’s reported differentiators—3D stacking, co-packaged optics, custom memory—are the same class of advanced-packaging risk that just cost the previous generation a year. A roadmap is a schedule of intentions, and this one has already demonstrated its error bars.

Second, and more consequentially, NVIDIA itself has stopped arguing on the axis the objection assumes. The GTC 2026 metrics were tokens per watt, tokens per second per user, and cost per token—not FLOPS. The company introduced attention-FFN disaggregation precisely because, in its own framing, GPU utilization on attention barely improves with batch size when the workload is bounded by loading the KV cache, and introduced new storage tiers that treat the KV cache as a first-order architectural constraint [511, 513]. That is this book’s argument, made by the incumbent: the binding constraint at inference is memory capacity and bandwidth, not arithmetic. A faster arithmetic roadmap does not settle a memory-bound contest, and Section 12.8’s capacity-first thesis is strengthened rather than weakened by NVIDIA conceding the diagnosis.

Third, the incumbent has now shipped the heterogeneous phase split, and its own GPU is not the right chip for both halves of it. Under NVIDIA authorship, the 2026 conference presented the Groq-derived LP30 LPU: no DRAM at all, 0.5 GB of on-die SRAM per chip aggregating to 128 GB across a 256-chip rack at roughly 40 PB/s and 315 PFLOPS of FP8; a cycle-exact VLIW ISA on a single global clock, with no memory barriers and no kernel boundaries; every chip its own router at 350 ns SRAM to SRAM [703] [P]. NVIDIA acquired Groq’s assets and acquired its leadership in a transaction reported near \$20B, completed in late December 2025, and Groq 3 LPX entered full production on 24 August 2026. The architecture is arresting, but the commercially significant content of the deck is the disaggregation set: prefill on Rubin and decode on LP30; attention on Rubin and feed-forward on LP30; Rubin verifying while LP30 drafts—claimed at roughly $3\times$ to $5\times$ the interactivity of Rubin alone [703] [P, vendor-stated, unaudited]. NVIDIA has also open-sourced the CUDA Tile IR and cuTile under Apache 2.0, in December 2025, and stated at the conference that it is extending CUDA to cover the LPU as a supported target; outlets characterize the depth of that integration differently, from an LPU extension of CUDA to a fully supported target, and this book declines to choose between them [702] [A].

The reading for this chapter is not that NVIDIA has conceded anything to China. It is that the strongest confirmation of the die-area-allocation argument of Section 12.3 now comes from the company whose product category that argument dissolves: the correct allocation changes *within a single inference request*, between prefill and decode, and NVIDIA is selling the heterogeneity rather than denying it. Heterogeneous phase-split serving is therefore no longer a research idea a challenger might exploit; it is a shipping product line the incumbent reached first, with the software integration—the hard part, and the part Section 12.7 prices—absorbed into CUDA rather than left to a consortium. And the memory bet inside it is the *opposite* of Section 12.8’s: half a gigabyte of SRAM per chip is a bandwidth-extreme, capacity-nil allocation, and the same company is now running that bet and the LPDDR-based Vera line simultaneously. Section 12.9 treats the revealed preference this implies as the strongest available argument *against* this chapter’s capacity-first thesis, and pointing at the disaggregation slide does not answer it.

Fourth, the honest concession. The same diagnosis makes China’s memory deficit *more* damaging, not less. If capacity and bandwidth decide inference economics, then HBM is the chokepoint that matters, and it is the one where the gap is widest: analyst work puts CXMT at HBM2-class with HBM3 samples near 50% yield, against an incumbent shipping HBM4 in volume, and reports that Huawei’s Ascend ramp is bounded by memory rather than by logic capacity [514, 121] [A]. This book’s position is not that the deficit is illusory. It is that a deficit in one input is survivable when the architecture is chosen to need less of it—which is the entire content of the capacity-first argument—and that the relevant comparison is deployed tokens per watt at the busbar, not peak FLOPS on a slide. On that comparison the American roadmap advantage and the Chinese energy-and-deployment advantage are running in opposite directions, and neither has yet been measured against the other in public. That measurement, not the next architecture name, is what would settle the question.

12.5 What the argument is not: the GPU is not obsolete

Two claims must be kept apart, and the enthusiasm around this architecture routinely blurs them. The first—*the architectural continuum is real*—is strongly supported: by LineShine, by Grace, by TPUs and dataflow accelerators, by integrated HBM processors, and by the progressive redistribution of die area among scalar, vector, matrix, memory, and communication functions across every vendor’s roadmap. The second—that *the CPU endpoint will become economically optimal for frontier AI*—is unproven, and depends on software maturity, sustained matrix efficiency, interconnect, manufacturing economics, deployment reliability, and total cost of ownership. The disciplined statement is this:

LineShine demonstrates that the traditional CPU–GPU category boundary is no longer fundamental. It does not yet demonstrate that any particular point on the resulting design continuum is economically optimal for frontier AI.

Nothing in this chapter should be read as predicting the obsolescence of discrete accelerators. The GPU corner is not a legacy position: it holds the deepest kernel libraries in computing, seven years of compounded framework integration, tuned collectives, the largest installed base of trained engineers, and—not least—an operational track record at fleet scale that no CPU-side machine has yet attempted for a multi-week frontier run. A book that predicted the disappearance of that would be making the same category error it accuses others of making about Chinese industrial capability.

The interesting claim is different, and it is positive rather than eliminative: *the CPU corner can reach performance parity on the workloads that matter, and if it does, it arrives carrying an ecosystem advantage that the accelerator corner does not have.*

12.5.1 Breadth against depth

The two corners have opposite software strengths, and naming them correctly is more useful than ranking them. But first, a framing discipline the comparison itself demands: the practical alternatives are not “a CPU” and “a GPU.” They are an *integrated accelerated-CPU system* and a *heterogeneous CPU–GPU system* with coherent or semi-coherent memory and a decade of orchestration maturity—and the heterogeneous system does not need to acquire the application universe, because it can keep irregular control, I/O, meshing, and legacy code on its host while accelerating selected kernels. The honest comparison is therefore between complete workflows:

$$T_{\text{workflow}} = T_{\text{host}} + T_{\text{accelerator}} + T_{\text{movement}} + T_{\text{synchronization}} + T_{\text{orchestration}}, \quad (12.1)$$

where every term is wall-clock time for one complete unit of work: T_{host} is time spent on the general-purpose processor (I/O, meshing, control flow, sparse assembly, the irregular tail); $T_{\text{accelerator}}$ is time in the accelerated kernels themselves; T_{movement} is data transfer between memory spaces; $T_{\text{synchronization}}$ is waiting at barriers and collectives; and $T_{\text{orchestration}}$ is scheduling, launch, and runtime overhead. The terms are additive only to the extent they are not overlapped—a well-tuned heterogeneous system hides part of T_{movement} behind $T_{\text{accelerator}}$, and the equation should be read as an upper bound that good engineering erodes rather than as an identity. The integrated design wins exactly where the removal of the movement, synchronization, and orchestration terms outweighs the discrete accelerator’s kernel-depth advantage on $T_{\text{accelerator}}$. For kernel-dominated dense training, the last term structure favors the discrete accelerator and will for years. For the interleaved simulation–analysis–inference workflows of agentic science (Chapter 24), the middle terms are a large and growing fraction of wall-clock, and they are the terms the integrated design deletes. Everything below should be read against Eq. (12.1), not against a device-category comparison.

A second discipline: “runs natively” is not one property but four, and the breadth advantage is strongest on the first two:

Binary/ABI portability — existing binaries and build systems execute unchanged. The integrated CPU wins this outright; the heterogeneous system wins it too, on its host side.

Source portability — code recompiles without restructuring. The integrated corner wins broadly; accelerator offload requires directive or kernel-language surgery.

Performance portability — the recompiled code actually exploits the matrix engines. *Not* automatic on either corner: legacy scientific code may compile on an accelerated CPU, while exploiting SME/AME-class engines still requires library substitution, data-layout changes, mixed-precision reformulation, compiler maturity, and sometimes algorithm redesign. This is the rung where the maturity ladder of Section 12.1 bites both sides.

Operational portability — schedulers, monitoring, checkpointing, failure handling, and operator skills transfer. The CPU corner’s genuine decades-deep advantage; conversely, mature framework abstractions preserve much application-level investment on the GPU side even when the vendor retunes internal kernels.

The accelerator corner’s advantage is *depth*: an extraordinarily well-optimized vertical stack for a narrow set of operations—dense and block-sparse matrix multiply, attention, collectives, the reduction and normalization kernels around them—refined over a decade with enormous investment, integrated into every framework, and backed by profiling and debugging tools built for that stack specifically. On the kernels that matter for a transformer, this is a genuine and durable moat, and the LineShine class is years behind it (Section 12.2.1).

The CPU corner’s advantage is *breadth*, and it is under-priced because it is invisible in benchmark tables:

One address space, no host–device split. There is no explicit data movement, no separate allocator, no pinned-memory dance, no kernel-launch boundary, and no class of bug that exists only because two memories must be kept coherent by hand. For workflows that interleave simulation, analysis, and inference—which is what agentic science actually is (Chapter 24)—this removes the dominant source of both engineering effort and lost performance.

The existing software universe runs natively. Fifty years of accumulated scientific, numerical, systems, and enterprise code executes without porting. In a typical scientific application the accelerated kernels are a small fraction of the source; the rest is I/O, meshing, control flow, sparse assembly, string handling, and glue—and on the accelerator corner that remainder either runs on the host anyway or must be rewritten. The CPU corner pays no porting tax on the 90% of code nobody writes papers about.

The irregular tail is native. Sparse operations, indirect addressing, branch-heavy control flow, dynamic data structures, mixed-precision solvers, graph traversal, and irregular communication are what much of computational science actually consists of, and they are exactly where accelerators are weakest. As AI itself turns sparser—mixture-of-experts routing, dynamic depth, retrieval, tool dispatch—this stops being a scientific-computing footnote and becomes an AI property.

Mature operations. Operating systems, schedulers, containers, filesystems, debuggers, profilers, checkpointing, RAS, and fleet management are decades matured on the CPU corner. Multi-week reliability is one of the seven open questions in Table 12.2, and it is the question the CPU corner is best equipped to answer.

No kernel-language lock-in. Code written to a scalable vector and matrix ISA is portable across vendors implementing that ISA—which is the entire point of Section 12.1, and the layer at which the rent of Chapter 2’s Table 2.2 currently sits.

One fleet instead of two. A converged machine serves simulation, AI training, AI serving, and data analytics from one procurement, one operations team, and one power envelope. For a national laboratory or a mid-size country, the elimination of the two-fleet problem is worth more than a percentage of peak.

The strategic bet the CPU corner represents is therefore a bet about *which kind of software capital compounds faster*. Depth is a body of kernels that must be substantially rewritten every architecture generation anyway—and the industry rewrites it, repeatedly, whenever precision formats, attention mechanisms, or sparsity patterns change. Breadth is the accumulated application capital of half a century, and it does not have to be rewritten at all if the matrix engines arrive inside the ISA. If matrix throughput becomes an implementation choice rather than a

device category—which is the argument of Section 12.3—then the corner that already holds the application universe has to acquire a decade of kernel depth, while the corner that holds the kernel depth has to acquire the application universe (or, per Eq. (12.1), keep hosting it and pay the movement and orchestration terms forever). Both are hard. This book’s judgement, stated as a judgement [S], is that the first is the more tractable problem, because kernels are a finite and well-specified target with a large and motivated workforce, whereas the application universe is neither finite nor specified.

The finite-kernel claim itself deserves an immediate qualification: the kernel set is finite at any moment but the target moves—across precision modes, sparsity patterns, expert-routing schemes, dynamic shapes, multimodal operators, graph compilation, distributed layouts, optimizer variants, numerical-reproducibility requirements, and failure recovery. “Finite” therefore means *bounded and specifiable per generation*, not fixed. The empirically decidable form of the bet is: *which side reduces the total lifecycle cost of supporting a changing workload portfolio faster?* That question is now a registered forecast in Appendix J, with observable indicators (time-to-competitive-kernel for each new attention or sparsity mechanism on each corner; fraction of new model architectures served within a fixed window (say six months) of release; operator-hours per delivered petaflop-month), rather than a rhetorical flourish.

That is also why the strategic reading of LineShine survives the deflationary correction. The machine’s significance is not that it retires the GPU. It is that a serious industrial actor has now demonstrated, at the top of the world rankings, that the ecosystem-rich corner of the continuum can be built at exascale—and if that corner also reaches AI parity, the competition between corners will be decided not by peak FLOPs but by which one is easier to program, operate, and converge. Those are the axes on which the CPU corner has never been behind.

12.6 Architecture theory: what to report, and why

Because this chapter argues from architecture rather than from benchmarks, it owes the reader the analytical vocabulary it is using. Every claim here reduces to the roofline framework: performance is bounded by $\min(\text{peak compute}, \text{operational intensity} \times \text{achieved bandwidth})$, and the workload’s operational intensity—FLOPs per byte of memory traffic—determines which bound binds [367]. Autoregressive decode over a sparse MoE has an operational intensity an order of magnitude below any current socket’s ridge point, which is not a modelling choice but an arithmetic property of reading weights once per generated token; that is why decode reduces to bandwidth and why the per-TB/s comparison of Section 12.2.1 is a ratio rather than a fit. Prefill sits above the ridge, which is why it reduces to matrix-engine thresholds instead. The memory-wall literature explains why the ridge point has drifted upward for thirty years and why capacity-first designs are a rational response; the communication-avoiding-algorithms literature explains why the fabric-in-ratio discipline of Section 12.3 is a first-order requirement rather than a detail; and energy-roofline analysis explains why the right unit for the comparisons of Chapter 14 is joules per delivered token rather than watts per socket.

One refinement makes the decode-parity claim of Section 12.2.1 properly conditional. “Streams per TB/s” is comparable across devices only under matched conditions, because the per-token byte traffic is itself a sum of terms that batching, caching, and placement all move:

$$B_{\text{token}} = \frac{B_{\text{active weights}}}{N_{\text{effective batch}}} + B_{\text{KV read}} + B_{\text{KV write}} + B_{\text{activations}} + B_{\text{routing}} + B_{\text{collective}}, \quad (12.2)$$

Each term is bytes moved per generated token: $B_{\text{active weights}}$ is the weight traffic for the experts actually activated, divided by $N_{\text{effective batch}}$ (the number of concurrent sequences over which one weight read is amortized—“effective” because expert routing means different sequences in a batch may need different experts, so the achieved amortization is below the nominal batch size); $B_{\text{KV read}}$ and $B_{\text{KV write}}$ are traffic to and from the key-value cache holding the attention state of

the context so far; $B_{\text{activations}}$ is intermediate tensor traffic; B_{routing} is the metadata moved to dispatch tokens to experts; and $B_{\text{collective}}$ is the per-token share of cross-device communication under the declared parallelism layout. Throughput is then bounded on both sides:

$$\text{tokens/s} \leq \min\left(\frac{F_{\text{achieved}}}{F_{\text{token}}}, \frac{\text{BW}_{\text{achieved}}}{B_{\text{token}}}\right), \quad (12.3)$$

where F_{achieved} is the arithmetic rate the device actually sustains on this workload (FLOP/s, not peak), F_{token} is the arithmetic work required per generated token, and $\text{BW}_{\text{achieved}}$ is delivered memory bandwidth under the workload’s real access pattern (bytes/s, again not peak). The left-hand ratio is the compute bound and the right-hand ratio the bandwidth bound; whichever is smaller is what the machine delivers. At low batch on a sparse MoE the amortized active-weight term dominates B_{token} , pushing the system onto the bandwidth bound; at extreme context, KV traffic dominates instead; and MLA, GQA, prefix sharing, expert caching, and quantization each move the boundary by shrinking one term or another. The per-TB/s parity of Section 12.2.1 is therefore asserted only under matched active bytes per token, effective precision, cache behavior, batching, expert residency, KV representation, context length, routing, and collective pattern—which is exactly what the proof suite’s protocol requires disclosed, and why MLPerf alone (which may not exercise the long-context sparse regime that carries the strategic claim) does not settle it.

Accordingly, for every design point this book proposes or evaluates, the quantities that should be reported—and which Appendix A now carries or names as owed—are: delivered bandwidth under representative access patterns rather than peak; all-to-all and collective injection requirements at the declared parallelism layout; power per *delivered* token; performance under expert imbalance, which is the sparse workload’s most under-reported failure mode; checkpoint and restart overhead at full state; utilization loss from failure and maintenance; and software maturity and operator labour, which are costs even when they are not line items. A design compared on peak FLOP/s and nominal bandwidth alone has not been compared.

12.7 The software and operations moat, made measurable

The breadth-versus-depth bet of Section 12.5.1 was registered as a forecast about lifecycle cost; that cost must therefore stop being a qualitative residual and become an object with a definition. Here it is:

$$\begin{aligned} C_{\text{lifecycle}} = & C_{\text{hardware}} + C_{\text{energy}} + C_{\text{porting}} + C_{\text{optimization}} \\ & + C_{\text{regression}} + C_{\text{operations}} + C_{\text{downtime}} + C_{\text{migration}} + C_{\text{decommission}}, \end{aligned} \quad (12.4)$$

where C_{hardware} and C_{energy} are the acquisition and electricity costs that procurement compares, C_{porting} is the one-time cost of making the workload run at all, $C_{\text{optimization}}$ the cost of making it run well, $C_{\text{regression}}$ the recurring cost of keeping it running well across toolchain releases, $C_{\text{operations}}$ the staffing of the fleet, C_{downtime} the value of lost availability, $C_{\text{migration}}$ the cost of moving the workload off the platform at end of life, and $C_{\text{decommission}}$ the disposal and data-transfer cost of retiring it. The first two are what procurement compares; the remaining seven are where platform advantage actually lives. CUDA’s moat, decomposed, is a claim about the seven: lower porting cost (everything already runs), lower optimization cost (kernels exist), lower regression burden (the vendor absorbs it each release), mature operations tooling, less downtime, and—the term nobody prices until they pay it—migration cost *away*, which is the lock-in rent collected at the end of the relationship rather than during it.

For each candidate architecture, the measurable schedule that would let a buyer compare Eq. (12.4) across platforms: time to first correct execution of a representative workload; time to 50%, 75%, and 90% of attainable performance (the performance-portability curve, which

distinguishes “compiles” from “performs”—the four portability types of Section 12.5.1 made quantitative); source lines modified and architecture-specific code fraction; compiler and library defects encountered per port; person-months of optimization; regression burden across toolchain releases; operator staffing per thousand nodes; unplanned-interruption rate; checkpoint/restart loss; performance portability across hardware generations; and the cost of migrating one production workload off the platform. The comparative cases that would anchor the table span the industry’s whole experience: NVIDIA/CUDA (the incumbent baseline), AMD/ROCm (the best-documented catch-up attempt), Huawei Ascend/CANN/MindSpore (the open-sourced challenger), A64FX/Fugaku (the direct historical calibration for the LX2 class—and the source of this book’s “capability present, stack lagging by a measurable interval” prior), OpenXLA and the compiler-abstraction bet, the RISC-V vector/matrix toolchains as they mature through the Section 12.1 ladder, and an LX2-class machine itself when access permits.

Two institutional observations complete the section. MLPerf provides the time-to-train anchor—and its v6.0 round now includes a DeepSeek-V3 MoE pretraining benchmark, making it directly relevant to the sparse workloads this book cares about [378]—but MLPerf measures the *result* of software maturity, not its cost: a heroic two-hundred-engineer port scores identically to a recompile. The complementary instrument the field lacks, and which this book now formally proposes as a deliverable of the Chapter 27 consortium, is an *AI-HPC lifecycle benchmark*: a standardized workload portfolio whose reported metrics are the schedule above—engineering effort, defect counts, regression burden, operational staffing—audited across platforms on a fixed cadence. It would do for the software moat what TOP500 did for FLOPs: convert a marketing claim into a time series. Publication of such a series is now a named verification event for the breadth-versus-depth forecast in Appendix J; until it exists, every claim about “ecosystem maturity”—including this book’s—is graded [S].

12.8 Democratizing AI through the LPDDR6 hardware arbitrage

The conventional Western paradigm assumes frontier training requires terabytes-per-second of High Bandwidth Memory—a technology where allied firms (SK hynix, Samsung, Micron) hold a commanding lead and supply is structurally tight (HBM consumes roughly $3\times$ the wafer capacity of standard DRAM per bit [133]). US export controls accordingly treat HBM as a chokepoint. China’s most interesting countermove is not to duplicate HBM—it is to sidestep it by exploiting massive capacities of low-cost commodity memory, accepting a bandwidth penalty and engineering around it in software. This is the LFP move (Chapter 4), executed in memory—and it is the training-centric variant of the Section 12.3 template.

12.8.1 The HBM scaling wall

Analysis of large-scale Mixture-of-Experts training exposes the capacity limitation of the HBM-centric architecture. A feasibility study by the RIKEN Center for Computational Science of pretraining a 2.8-trillion-parameter MoE of the Kimi-K3 class finds that, under standard production-grade kernel fusion, sustaining a 20-PF(FP8) accelerator requires approximately 14.5 TB s^{-1} of memory bandwidth per accelerator—beyond even HBM4 provisioning [134]. The binding constraint, however, is capacity, not bandwidth: on a 128-accelerator cluster, the sharded model state alone ($\sim 306\text{ GB}$ per accelerator) exceeds the 288 GB carried by an eight-stack HBM4 package. One correction is owed here, and it sharpens rather than softens the result. 288 GB is a *packaging convention*, not a JEDEC ceiling: AMD’s MI455X ships 432 GB from twelve stacks, and at that capacity the 306 GB per-accelerator state fits (Section 13.3) [679]. The floor therefore moves rather than vanishing—from roughly 136 accelerators at eight stacks to roughly 91 at twelve, for the same aggregate $\sim 39\text{ TB}$ of state. The precise statement, with the discipline the result demands, is this: *under the stated Kimi-K3-class model, optimizer representation,*

activation policy, and parallelization assumptions, an eight-stack HBM4 configuration has a capacity floor near 136 accelerators and a twelve-stack configuration one near 91—beneath it the Western architecture does not merely slow down but ceases to function [134]. What the correction does to the argument is worth stating, because it cuts in the book’s favour rather than against it: the floor is now a function of *which vendor’s HBM4 the buyer bought*, not of the generation, which is precisely the finding of Section 13.3—“HBM4” has stopped specifying a part, and a capacity floor that depends on a procurement choice is a weaker foundation for a national architecture than one that depends on physics. The threshold is configuration-specific, not universal: it moves with optimizer state, precision, sharding strategy, sequence length, expert placement, recomputation policy, and checkpoint design. It remains a strong result because the stated configuration is the *standard* production configuration for the model class.

12.8.2 CXMT and the LPDDR6 disruption

ChangXin Memory Technologies (CXMT) supplies the memory class on which the alternative rests. Reports place its first LPDDR6 part—12,800 Mbit s⁻¹, 16 Gbit per die—at the tail end of R&D verification in mid-2026, with sampling underway and mass production targeted for 2H26, which would place CXMT at or ahead of SK hynix’s LPDDR6 timeline: a Chinese memory firm reaching a leading-edge commodity node *first*. This claim is graded [A]—analyst-reported roadmap evidence—because no official product announcement or customer qualification has been published; the falsification test is simply whether a qualified part ships on the reported schedule, and Appendix J tracks it [135]. The JEDEC SOCAMM2 module standard (LPDDR-on-module for AI servers, already in mass production for NVIDIA’s own Grace/Vera-Rubin platforms; SK hynix began mass production of 192 GB SOCAMM2 modules in April 2026 [136, 470]) provides the packaging: with approximately 18 SOCAMM2 modules, a single socket reaches ~2.3 TB of directly attached capacity—eight times the HBM4 ceiling, at a fraction of the cost per bit [134].

A second discipline separates what is *observed* from what is *proposed*. What the evidence supports: a capacity-first LPDDR architecture is a plausible and rational Chinese countermove, because it maps onto exactly the memory class where domestic capability is strongest and sidesteps the class where it is weakest. What the evidence does not yet support: that this architecture is an *implemented* national countermove, or a product program that CXMT is leading. The design in this section and Appendix A is this book’s own engineering analysis [M] of what the pieces permit—not a report of a Chinese program, and the text is written so the two cannot be confused.

12.8.3 The trade-off: capacity over peak bandwidth

Before the trade is stated, the design hierarchy has to be fixed—and it is this book’s own reconciled topology, rather than any preference for the novel corner, that fixes it. Appendix A’s design point shows the pure-LPDDR socket clearing its ~5 TB/s requirement only under streaming-dominant access; under random expert-fetch behavior it misses by 10–30%. The engineering ranking is therefore: (1) *the hybrid HBM+LPDDR configuration as the baseline*—most of the capacity advantage, a quarter of the pin burden, graceful degradation; (2) *pure LPDDR as a higher-risk, sanctions-contingent research configuration*—the corner a fully export-controlled program is pushed into, and the configuration whose feasibility milestone (the 5 TB/s knee test) should be funded first precisely because it is the one in doubt; (3) *HBM-only as the incumbent performance baseline*. What follows analyzes the pure-LPDDR corner because it is the *strategically novel* one, not because it is the recommended one—the sanctions geography, not the engineering, is what selects it.

The arbitrage rests on a calculated hardware–software trade, and the entire subsection is an author-modeled result [M] pending the experimental program of Appendix A. With 2.3 TB per socket, activation recomputation can be eliminated entirely, cutting both FLOP and memory-

traffic requirements; under aggressive kernel fusion the per-accelerator bandwidth requirement drops to a knee of $\sim 5.0 \text{ TB s}^{-1}$. An LPDDR6X system provisioned at 7.0 TB s^{-1} peak delivers that achieved bandwidth, yielding a Model FLOP Utilization ceiling of about 67%, against 89% for the HBM4 configuration [134]. The engineering price of that 7 TB/s socket is itself nontrivial—on the order of 4,400 active data pins at LPDDR6 rates, several times the disclosed bandwidth of NVIDIA’s own LPDDR-based Vera CPU (1.2 TB/s) [257]; Appendix A does the arithmetic and states the validation milestones. What the MFU sacrifice buys: cost-per-bit an order of magnitude better; capacity-per-socket that *enables* small-cluster training at all; and supply from a domestic fab ramping toward 300,000+ wafers per month (Chapter 13). For the stated model class and parallelism layout, below 256 accelerators HBM4 capacity cannot hold the sharded state at all—there, LPDDR6X-class capacity is not the cheaper option but the enabling one.




Figure 12.2: Why capacity is the arbitrage. *Above*: capacity against bandwidth for shipped parts [V] and for the modeled LPDDR6X socket [M]. The dashed line is the per-accelerator model state for the reference 2.8T-parameter sparse-MoE run at 128 nodes ($\sim 306 \text{ GB}$): everything to its left cannot hold the model at all, which is the sense in which the HBM path has a cluster-size floor. *Below*: the trade the substitution buys—22 points of MFU ceiling surrendered for an eightfold capacity gain and an order-of-magnitude better cost per bit. Curve shape is modeled, not measured; Appendix A states the experiments that would confirm or refute it.

12.8.4 Strategic consequence: decentralized pretraining

The capacity advantage attacks the assumption that frontier pretraining requires centralized, gigawatt-class superclusters. Scale honesty first, because the intuitive figure for a machine of this class errs badly in the conservative direction. A 128-socket machine at 20 PF(FP8) per socket is a 2.56-EF FP8 installation. Its power is *not* tens of megawatts:

$$P_{\text{facility}}^{\text{el}} = \underbrace{N_{\text{sockets}} P_{\text{socket}}^{\text{el}} + P_{\text{network}}^{\text{el}} + P_{\text{storage}}^{\text{el}}}_{\text{IT load}} + \underbrace{P_{\text{cooling}}^{\text{el}} + P_{\text{overhead}}^{\text{el}}}_{\text{facility overhead, i.e. (PUE-1) \times IT load}},$$

where N_{sockets} is the number of accelerated processors, $P_{\text{socket}}^{\text{el}}$ the power of one processor plus its attached memory tier, $P_{\text{network}}^{\text{el}}$ and $P_{\text{storage}}^{\text{el}}$ the interconnect and storage draw, and the last two terms the cooling and electrical losses that power-usage effectiveness (PUE) summarizes as a multiplier on IT load. Table 12.4 groups the terms exactly this way.

Table 12.4: Power accounting for a 128-socket capacity-first pretraining cluster [M]. A “socket” is one accelerated processor plus its attached LPDDR6X memory tier; nodes carry two sockets.

Term	Value	Basis
Compute silicon per socket	2.0–2.6 kW	~2× a B200-class part at 20 PF(FP8)
LPDDR6X tier per socket (2.3 TB @ 7 TB/s)	0.3–0.5 kW	~5–8 pJ/bit at full stream
Socket total	2.3–3.1 kW	
128 sockets (64 dual-socket nodes)	295–397 kW	
Network (multi-rail, 1.6 TB/s per socket)	25–40 kW	optics plus switching
Storage and head nodes	15–30 kW	checkpoint tier at full state
IT load	335–467 kW	
Cooling and facility overhead (PUE 1.15–1.30)	+50–140 kW	liquid-cooled assumption
Facility total	~0.4–0.6 MW	8–12 racks

Half a megawatt is an ordinary industrial substation feeder, not a campus—which strengthens the decentralization argument rather than weakening it, and is the arithmetic a skeptical reader should check first. Even a 256-socket machine lands near 1 MW. What that buys is *smaller sovereign-scale*: a national-laboratory or mid-cap-company object; “democratization” here means the pretraining club expanding from three hyperscalers to every state and serious institution that wants in, and the capacity threshold itself is workload-specific, not a universal wall. Within that scope the claim stands [M]: sovereign states, national labs, and mid-size companies can execute full pretraining of frontier-class sparse models on 128–256-socket (64–128-node) clusters, because LPDDR6 provides the state capacity and Chinese sparse-model recipes (Chapter 10) provide the bandwidth frugality. Note the co-design loop closing at every scale: the same capacity-first, bandwidth-modest architecture that the \$2,000 enthusiast rig discovered bottom-up (Section 10.6) reappears here as national strategy—Chinese models are trained sparse and bandwidth-lean *because* of the hardware environment, and Chinese hardware is provisioned capacity-rich *because* the models reward it. The Western trillion-dollar bet on centralized HBM-bound data centers is thus exposed on two flanks at once: its scarce input (HBM) is bypassed, and its scale premise (only hyperscalers can pretrain) is dissolved. Hardware proliferation plus software commoditization equals the democratization—on Chinese terms—of the means of AI production.

12.9 The memory unlock: what happens if capacity becomes cheap

Every argument in this book about open weights has an unstated premise that this section makes explicit and then interrogates: *that somebody can afford to run them*. Open weights are a commons only for those with the memory to hold them. And in 2026 that memory became, abruptly, the most expensive component in computing.

12.9.1 The constraint, priced

The surge is not a trend but a step change. Conventional DRAM contract prices rose 93–98% quarter-on-quarter in 1Q26, 58–63% in 2Q26, and 13–18% in 3Q26—compounding to roughly $3.6\times$ in three quarters, on top of a $\sim 172\%$ rise across 2025 [563] [A]. Retail followed: a 32 GB DDR5 kit went from roughly \$85 in mid-2025 to \$369–432 in 2026; 64 GB DDR4 registered ECC modules, the staple of a self-built inference server, went from under \$110 to \$500 used and \$835 new [564] [A]. HBM4 is reported near \$2 per gigabit—about \$16 per gigabyte, so on the order of \$1,000 for a 64 GB stack—with forecasts of \$4–5 per gigabit in 2027 [565] [A].

One product decision states the consequence better than any price series. On 6 March 2026 Apple *withdrew* the 512 GB option from the Mac Studio and raised the 96-to-256 GB upgrade from \$1,600 to \$2,000 [566] [V]. The most capable consumer machine for running large open-weight models locally got *worse* in 2026, because of memory. The commons expanded and the ability to hold it contracted, in the same year.

Set that against what the commons now weighs. Kimi K3 ships natively in MXFP4—quantization-aware from the fine-tuning stage, so there is no further compression to recover—and the published download is roughly 1.4 TB against 5.6 TB at BF16 [567] [V]. DeepSeek V4-Pro at 1.6 T parameters is roughly 800 GB at four bits; GLM-5 at 744 B is roughly 372 GB [M, arithmetic]. *The frontier of the open commons is a terabyte-class object*, and the escape hatch of aggressive quantization has already been spent by the model publishers themselves.

The second tier is different in kind. Key-value cache is written every token, and its size is architectural: DeepSeek’s own hardware paper reports 70.3 KB per token for its multi-head latent attention against 516.1 KB for Llama-3.1-405B’s grouped-query attention—a seven-fold difference [568] [P]. At the one-million-token contexts these models now advertise, that is roughly 72 GB for the MLA model and 528 GB for the GQA one. The same paper states the underlying scissors: memory demands growing by more than 1000% per year against high-bandwidth memory capacity growing under 50%.

Two tiers, then, with opposite properties: weights are enormous, static, and read-only; context is smaller, volatile, and written constantly. The industry has been trying to serve both from the same expensive resource.

12.9.2 Two technologies, precisely described

High Bandwidth Flash is the more directly relevant, and it is now a published specification rather than a concept. SanDisk and SK hynix signed a standardization memorandum in August 2025, formed a consortium under the Open Compute Project in February 2026, and released the first OCP technical specification on 3 August 2026: NAND dies stacked with through-silicon vias in a package matching HBM4’s footprint and interposer slot, connected to the host over UCle rather than PCIe—which is the whole idea, since removing the NVMe controller and the PCIe link is what yields roughly a hundredfold bandwidth improvement over conventional flash. The specification describes up to 512 GB per stack across three bandwidth grades from about 0.4 to 3.0 TB/s [569, 570] [V]. SanDisk’s own fact sheet takes the roadmap further: 512 GB per sixteen-die stack at 1.6 TB/s in the first generation, stack capacities “set to reach up to 1 TB” in the second and 1.5 TB in the third at 3.2 TB/s, with per-generation power falling to $0.8\times$ and then $0.64\times$ of the first—and, in the line that matters most for what follows, “*non-volatile, no refresh power needed*” [586] [P]. A terabyte of weights is therefore *one* second-generation stack; two stacks hold the entire open commons at four bits.

Four qualifications must travel with those numbers, and a book that prints the first without the rest is doing marketing. The bandwidth and capacity figures are the two vendors’ development targets, not conformance requirements; the specification carries no version number, no conformance procedure and no certification scheme. *No endurance, latency, power or price figure has been published at all* [571] [A]. Analyst estimates put HBF latency in microseconds

against HBM’s hundred nanoseconds, minimum read granularity roughly 128 times larger, and power at about twice HBM4 per unit—with cost per gigabyte perhaps tenfold lower but cost per *stack* comparable, because advanced packaging rather than the die dominates [572] [A]. And the consortium’s membership is as informative as its specification: SanDisk, SK hynix, Google and Tenstorrent are in; NVIDIA, AMD, Samsung, Micron and Kioxia are not. Nothing is sampling; the first HBF die taped out in August 2026, product samples are now promised for 2027, and volume is generally placed at 2028 or later. Appendix B sets out the full record—lineage, announcements, the specification and its omissions, power, endurance, cost structure, the debate and the literature—so that every assumption this section and Section 12.9.4 make about the technology can be checked against it.

The suitability question resolves cleanly, and in the direction that matters. Flash tolerates weights—deterministic, read-only, sequential-friendly—and cannot tolerate key–value cache, which would write every token into a medium with finite program-erase endurance. *HBF is a weights tier, not a memory replacement*, and its own vendors position it as complementary to HBM rather than as a substitute.

ZAM, the second technology, is a capacity architecture, and the public reporting under-describes it. SAIMEMORY Corporation was incorporated in December 2024 as a SoftBank subsidiary; its development partner is *Intel*, its foundry partner Powerchip, and its Series A investors Fujitsu, the Development Bank of Japan and RIKEN, with a NEDO grant [573, 574] [V]. (This author directs a RIKEN centre; RIKEN’s investment is disclosed here, and what follows draws in part on direct knowledge of the programme.) ZAM—Z-Angle Memory—is stacked DRAM built on Intel’s next-generation DRAM bonding work, itself validated under a US Department of Energy advanced-memory programme. Its architectural moves are two: a “via-in-one” through-silicon via that replaces HBM’s many TSVs and frees die area for cells, and a vertical thermal path that conducts heat upward through the stack rather than trapping it—which is what removes the sixteen-to-twenty-layer ceiling that bounds conventional stacking. The reported headline is power, roughly 40–50% below comparable HBM; the *purpose* is layer count. Once the thermal ceiling is lifted, capacity comes from stacking height rather than from lithographic density, which is why the target is 512 GB per chip—eight times an HBM4 stack—and why the economic logic is that such stacks can be built on mature, cheaper DRAM nodes rather than at the leading edge; Intel’s own statement of the programme names “optimizes memory costs” among its objectives [574] [P/A]. A VLSI Symposium paper in June 2026 reported functionally validated nine-layer silicon at about 0.25 Tb/s/mm² and under 0.7 pJ/bit of data-movement energy [575] [P]. That is a first step, not a ceiling: SoftBank’s timeline is a prototype in the fiscal year to March 2028 and commercialization in FY2029, and the point of the vertical thermal path is precisely that it removes the reason stacks stop at sixteen—the design targets *sixty or more layers*, which is where the 512 GB figure comes from and why the economics turn on mature nodes stacked high rather than leading-edge nodes stacked shallow [P, author’s knowledge; not yet in the public record]. ZAM is therefore the DRAM-side counterpart of HBF’s flash-side move: capacity by height on cheap nodes, rather than bandwidth by density on expensive ones. It is later than HBF and its capacity claims are further from demonstrated silicon, and the book grades it accordingly—but it is a cost-per-bit technology, and reading it as merely a low-power HBM variant would miss the point of the design.

12.9.3 The machine, and whether it works

The architecture the two tiers imply is straightforward: a processor with roughly a terabyte of HBF holding weights, and a few hundred gigabytes of LPDDR6X holding context, key–value cache and activations. That tier is not speculative: JEDEC published LPDDR6 in July 2025 and in April 2026 previewed the extension roadmap explicitly aimed at data centres—512 GB densities and an LPDDR6-based SOCAMM2 module standard in development with a stated upgrade path from today’s LPDDR5X SOCAMM2, the module NVIDIA already ships on Grace

and Vera—so the LPDDR6X-on-SOCAMM2 tier is where the industry’s own standards work is going, and Section 12.8 has already argued its bandwidth case [576, 585, 577] [V]. Silicon in this direction already exists: Qualcomm’s AI200 carries 768 GB of LPDDR per card, and the AI250 adds near-memory compute for 2027.

Does the bandwidth arithmetic close? At four-bit weights, a decode step reads active parameters only—52 GB per token for Kimi K3’s 104 B active, about 24.5 GB for DeepSeek V4-Pro, 20 GB for GLM-5. Two 512 GB stacks at SanDisk’s first-generation 1.6 TB/s each give 3.2 TB/s aggregate, which yields roughly 62 tokens per second on a 2.8-trillion-parameter model for a single user, 131 for DeepSeek V4-Pro, 160 for GLM-5 [M, ideal ceiling with no read amplification]. Sixty tokens per second is faster than a person reads. *The concept is not arithmetically fanciful; it is bandwidth-sufficient with margin.*

Energy deserves precise treatment, because it is easy to get wrong in the direction of false alarm by carrying a datacentre metric over to a device that is not a datacentre. The datacentre literature is discouraging: measured flash reads cost about 102.4 pJ/bit against DRAM’s 4.2, and a peer-reviewed analysis finds that streaming mixture-of-experts weights from commodity SSDs raises per-token energy 3.8 to 12.5 times, concluding that flash needs roughly a tenfold improvement to be energy-competitive [578] [V]. Two things follow, and only one of them is a problem. Commodity NVMe genuinely cannot do this job at speed: at 60 tokens per second on Kimi K3, weight-read power alone would be roughly 2.5 kW—outside any distributed envelope. HBF, by removing the NVMe controller and PCIe serialization in favour of on-package UCIe, is estimated at about twice HBM4’s energy per bit read [572] [A]—which puts the same workload near 250 W against HBM’s 125 W [M].

But the SSD-offload critique is a *datacentre* critique, in which memory is powered continuously and energy is a dominant recurring cost at scale, and it does not transfer to a distributed machine, for two reasons the fact sheet makes plain. First, *HBF is non-volatile*: it draws essentially nothing when not being read, while a terabyte of DRAM must refresh continuously whether or not inference is running. A distributed appliance does not run at full load around the clock; at a duty cycle of one-third, HBF’s twofold active penalty is offset if HBM’s standby draw across the sixteen idle hours equals about half its active read power across the eight busy ones—a condition sixteen refreshing HBM4 stacks plausibly meet [M]. Over a day, the two are near parity. Second, and more simply, *the difference is small in absolute terms*. Even at full duty the HBF penalty is on the order of 100–250 W: roughly 70–180 kWh a month, which is tens of dollars at commercial tariffs [M]. The service such a machine displaces is metered at thousands to tens of thousands of dollars a month for a small or mid-sized enterprise. The energy delta is two to three orders of magnitude below the bill it eliminates, and it is comfortably within a workstation’s thermal budget. *Energy per bit decides whether the architecture is elegant, not whether it is viable*. What decides viability is that HBF ships at its specified capacity and bandwidth, and that is the number this section waits on.

12.9.4 Head to head: an appliance against a Rubin rack, and both against the API price list

Assertion is cheap; the question is what such a machine would cost to build and to run, against what it displaces. This subsection prices it, on stated assumptions, and the model that produces every number below ships with the manuscript source so that any reader can change an input and watch the conclusion move.¹

The appliance. Two terabytes of HBF as four 512 GB first-generation stacks at 1.6 TB/s each—6.4 TB/s aggregate; 256 GB of LPDDR6X on a 512-bit bus at roughly 900 GB/s for context,

¹`models/appliance-vs-rubin.py` in the source archive. All memory prices are taken on a *pre-surge*, mid-2025 basis—the book’s declared assumption for this analysis, since the argument concerns what the architecture does once memory prices normalize.

key-value cache and activations; a GB10-class system-on-chip with UCIE ports for the flash stacks and about 400 TFLOPS of achieved dense FP4 for prefill; a 2.5D package carrying the SoC and the four stacks; and the ordinary remainder of a workstation. The comparison model is Kimi K3—2.8T parameters, 104B active, sixteen of 896 experts routed per token, native MXFP4—because it is the largest and among the strongest open-weight models in existence and the hardest case for a local machine [567, 587].

Decode is memory-bound, so throughput follows from bytes per step. Each routed expert is roughly 1.5 GB at four bits; the attention, shared-expert and embedding parameters read on every step are about 27 GB; a single-user step is therefore about 52 GB and a single user decodes at roughly 120 tokens per second from 6.4 TB/s. Under the mixture-of-experts saturation arithmetic of the next subsection—uniform routing, which is the pessimistic case—four concurrent users each touch a growing union of experts, and the machine delivers about **52 tokens per second to each of four users**, roughly 210 aggregate; eight users see about 30 each [M]. Key-value traffic at four users and 128K context in FP8 is under 400 GB/s, comfortably inside the LPDDR budget, and the K3 architecture’s hybrid linear attention keeps the resident state small—under 16 GB for eight 128K sequences [587]. The weak point is prefill: 128K tokens of fresh prompt takes about a minute at 400 TFLOPS, 32K about seventeen seconds. Agentic loops mostly reuse cached prefix and pay only for the increment, and prefill can be offloaded to a shared service (Section 14.4), but a long cold prompt is where a rack beats a box.

Power at full load is roughly 620 W—about 360 W of that in flash reads at an assumed 7 pJ/bit, twice a plausible HBM4 figure, plus a 200 W SoC—and about 45 W at rest, since the flash draws nothing and only the LPDDR self-refreshes [M]. That is 3 J per generated token: a workstation, not a rack.

The bill of materials. Every component below is anchored to a finished, shipping, pre-surge system rather than to a guess, because a bare component price is easy to inflate quietly and a finished-product price is not. Three pre-surge, mid-2025-basis reference points do the anchoring. *Finished GB10-class systems:* the ASUS Ascent GX10—the same “up to 1 PFLOP FP4” GB10 silicon as DGX Spark, 128 GB of LPDDR5X, no HBF at all—sold at \$2,999 at launch; NVIDIA’s own Founders Edition launched at \$3,999 [592]. *Finished Strix-Halo-class systems:* the Framework Desktop, AMD Ryzen AI Max+ 395 with 128 GB of LPDDR5X, sold at \$1,999 at announcement, and the GMKtec EVO-X2 on the same silicon at \$1,799.99 early-bird [593]. Both are complete, retail-margined, finished computers with no HBF tier, so they bound what the SoC, board, PSU, case and cooling should cost *before* a single flash stack is added. *Finished NVMe retail:* a 1 TB drive ran \$65–85 in Q1 2025 and \$45–90 across late 2025, before the 2026 NAND shortage roughly doubled prices [669]—and a retail SSD carries a controller ASIC, a PCB, a connector, an enclosure and a full retail margin, none of which a bare HBF stack carries, since HBF has no controller and no PCIe: it is read directly over UCIE from inside the package. That is the basis for the stack price itself. SanDisk’s own positioning is that HBF stack cost runs “similar” to HBM [586], while raw TLC NAND at mid-2025 spot was about four cents a gigabyte, putting ~\$21 of die in a 512 GB stack [594]—so the die is a rounding error and TSV stacking, the logic die and test dominate the rest. Against a finished 1 TB NVMe retailing at \$65–100, a bare 512 GB HBF stack costing “a couple hundred dollars at most”—\$150 to \$250 across the three cases below, never approaching HBM parity—is the defensible *ceiling*, not merely a plausible base case: it undercuts a finished consumer SSD despite the added TSV-stacking step, because it skips the controller, the PCB, the enclosure and the retail margin that the SSD carries. LPDDR is priced at \$3.3 to \$4.5 per gigabyte, the pre-surge LPDDR5X range, with a modest premium toward the top of that band for LPDDR6X being the newer part [595]. The SoC and remaining-system lines are benchmarked directly against the same GX10/Framework anchors—a full finished computer at \$2,000 to \$3,000 with no HBF tier at all leaves little room for a bare SoC line priced any higher—with one exception: the 2.5D package stays anchored to

CoWoS-class advanced-packaging estimates for carrying four stacks alongside the SoC, since none of the three reference systems does that packaging step at all and the finished-system comparison therefore has nothing to say about it [596]. Margins run from PC-OEM to Apple-class.

Table 12.5: Capacity-first appliance: bill of materials and price on a pre-surge memory basis, anchored against finished-system pricing (ASUS Ascent GX10 \$2,999; Framework Desktop \$1,999; 1 TB NVMe retail \$65–100). 2 TB HBF (4 stacks), 256 GB LPDDR6X, GB10-class SoC, 2.5D package [M].

Component	Optimistic	Base	Conservative
HBF, 4×512 GB stacks	\$600	\$800	\$1,000
LPDDR6X, 256 GB	\$845	\$973	\$1,152
SoC (GB10-class, UCIe, ~ 1 PF FP4)	\$500	\$650	\$850
2.5D package, SoC + 4 stacks	\$600	\$700	\$800
Board, PSU, chassis, cooling, storage	\$300	\$400	\$550
Bill of materials	\$2,845	\$3,523	\$4,352
Gross margin	25%	30%	40%
Price	\$3,800	\$5,000	\$7,250

For scale: the DGX Spark launched at \$3,999 with 128 GB of LPDDR5X and a GB10 before its memory-driven rise to \$4,699; the ASUS Ascent GX10 partner box on the same silicon at \$2,999; the Framework Desktop at \$1,999 with 128 GB on Strix Halo; and the Mac Studio at \$9,499 with 512 GB of unified LPDDR5X before Apple withdrew the configuration [592, 593, 566]. A \$5,000 base price for sixteen times the Spark’s weight capacity in a non-volatile tier those systems do not have at all, twice its LPDDR, and comparable arithmetic sits close to the GX10 and below the Mac Studio: inside the range those finished products actually sold at pre-surge, not above it.

The rack. The central comparator is a Vera Rubin VR200 NVL72: seventy-two GPUs at 288 GB of HBM4 and about 13 TB/s each, drawing about 210 kW at the rack [591]. What it costs must be stated with care, because a rack has no retail price and the two sides of this comparison must be priced at the same point in the chain. Morgan Stanley’s May 2026 estimate of \$7.8M per rack, with roughly \$2M of it in memory, is labelled a bill of materials but is a costed component list *at the prices a cloud provider pays*—GPUs at about \$55,000 each, which is NVIDIA’s selling price with its roughly 75% gross margin already inside, not its cost of goods—and the report of the note says so explicitly [590] [A]. It is therefore an end-user transaction price for a volume buyer, and it is the right counterpart to the appliance’s retail price, which is likewise a bill of materials plus a 25–40% maker’s margin: neither side is priced at bare component cost, and the asymmetry to avoid—an appliance at retail against a rack at NVIDIA’s cost—is not present. What *is* present is dispersion in what different buyers pay for the same rack. Trade reporting in March 2026 put ODM quotes for the VR200 NVL72 at \$5–7M without warranty and “as much as \$8.8M,” with NVIDIA-supplied trays now about 90% of a system’s cost and the integrator’s margin thin [590] [A]; a hyperscaler buying tens of thousands of racks pays near the bottom of that range, while an enterprise or a neocloud buying a few racks through an OEM, with warranty and support, pays a good deal more—published GB300 NVL72 figures alone span \$3.7–4M from one analyst to \$6–6.5M in trade reporting, and figures near \$10M have been quoted for fully supported systems [590] [A/R]. The model therefore runs the rack twice: at \$7.8M, the hyperscaler’s price, and at \$10M, an enterprise-retail price with support, which is the price the buyer of a \$5,000 appliance would actually face if it bought a rack instead. Its Kimi K3 throughput is extrapolated from published measurements: on B200-class silicon, a one-trillion-parameter Kimi at 32B active decodes at roughly 3,200–3,900 tokens per second per GPU at 30–90 tokens per second per user, and Moonshot’s own vLLM deployment reports “2K+

tokens per GPU-second” for K3 on GB300 NVL72 [589, 588]. Scaling for K3’s threefold-larger active parameter count and Rubin’s 1.6-fold bandwidth advantage gives about 3,000 tokens per second per GPU in throughput mode (roughly 25 tokens per second per user) and about 1,800 at interactive parity with the appliance’s 50 [M].

The building, and the closet. A rack is not a cost; a rack *in a building that can power and cool it* is, and the two sides of this comparison differ most sharply in exactly that term, so it is priced explicitly rather than folded into an overhead line. A rack of this class cannot go into an existing server room: it needs a purpose-built facility with a dedicated substation, liquid-cooling plant, and—for the stringent-SLA inference this comparison is about—concurrently maintainable power. What such a facility costs, per watt of IT load and excluding both land and the IT itself, is documented in a companion survey by this author, whose Tier-III-capable “resilient inference/cloud” reference architecture prices at \$11.9 per watt at the median with a \$8.5–17.7 planning band, dominated by electrical plant and utility interconnection [597] [M/A; the self-citation discipline of the preface applies]; the independent anchors sit on top of it—Epoch AI’s one-gigawatt model puts construction at \$12 of a \$38-per-watt all-in campus, JLL’s 2026 shell-and-core average is \$11.3M per megawatt, and Turner & Townsend’s Silicon Valley index is \$13.3 per watt on a narrower scope [598, 599, 600] [A]. At \$12 per watt a 210 kW rack carries about **\$2.5 million of building**, a third again on top of the \$7.8M of silicon, and the building has a fourteen-to-fifteen-year life against the silicon’s four to six [598, 601]—so it will outlive three generations of the racks it was built for. Operating it costs, before a kilowatt-hour is bought, roughly \$300 per kilowatt of IT load per year in staff, maintenance contracts, security, insurance, property tax and water: Epoch’s \$0.3B per gigawatt-year, KPMG’s 2026 benchmark of \$237–353, an independent US synthesis near \$320 [598, 602] [A]—about \$63k per rack-year—and the equipment inside it erodes on its own schedule, VRLA batteries in four to six years, UPS modules in ten to twelve, chillers and cooling towers in about twenty, and JLL’s own guidance is that two-thirds of facilities need significant renovation within ten years and nearly all within fifteen [603, 599] [A]. And capital is not free: Meta’s Hyperion campus was financed at 6.58%, CoreWeave’s notes at 9 to 9.75% [604, 605] [V]; the model charges both sides of the comparison an 8% cost of capital. A cross-check that the bottom-up number is not invented: facility annuity plus operating cost comes to about \$142 per kilowatt-month, and CBRE’s asking rent for 10 MW-plus wholesale capacity in Northern Virginia in the second half of 2025 was \$155–185 per kilowatt-month, power passed through at cost [606] [V]—the difference is the landlord’s margin, which is what a hyperscaler avoids by owning and a neocloud pays. Two things the model does *not* price, both cutting against the rack: the two-to-seven-year wait for a large-load grid interconnection that the same survey documents as the binding schedule constraint, and the discrete-block penalty of a facility built ahead of the racks that will fill it [597].

The closet is the other side of the ledger, and it is nearly empty. A 620-watt appliance goes into the room the enterprise already has, on the circuit it already has, cooled by the air conditioning it already runs; the model charges it a \$500 UPS, a small-room PUE of 1.3 (worse than a hyperscale hall’s 1.2), the same 8% cost of capital, and no operations labour at all on the stated assumption that the agent stack administers itself—an assumption the harness argument of Chapter 11 makes plausible and this book flags as an assumption [M]. Fully loaded, then: the rack’s operator pays about **\$2.9M per rack-year** at four-year silicon amortization—\$2.35M silicon, \$0.29M building, \$0.06M facility operations, \$0.17M power—and \$2.2M on a six-year silicon life; at 60% utilization that is \$0.70 and \$0.54 per million output tokens in throughput mode and **\$1.17 and \$0.90 at interactive parity** [M]. The honest reading of the decomposition is that at full utilization the building is a minority of the per-token cost—about 12% including its operations, because silicon that depreciates in four years dominates a building that depreciates in fifteen. What the building changes is not the margin at full load but everything else: the absolute capital committed, the fifteen years it is committed for, the years before the first token

can be served, and—decisively—what happens to the per-token number when utilization does not arrive, which is a question the closet never has to answer, because it was already there.

The comparison. Table 12.6 puts the three things side by side: what the appliance costs per million output tokens at realistic monthly loads, what the rack costs its operator fully loaded, at planned and at disappointed utilization, and what the API price list charges.

Table 12.6: Cost per million *output* tokens, Kimi K3 or nearest frontier equivalent, both sides fully loaded at an 8% cost of capital. Appliance at base price, four-year life, 3%/yr maintenance, \$500 closet UPS, small-room PUE 1.3, US commercial tariff. Rack at four- or six-year silicon life, \$12/W facility over fifteen years, \$300/kW-yr facility operations, PUE 1.2, hyperscale tariff. API prices are August 2026 list, output tokens only; input and caching would add to both sides. Independent capability index (Artificial Analysis, **v4.1.1, the August 2026 board**) in brackets where available [505]; the index was re-based on 7 September (v4.3), moving every score by ten to fifteen points, so bracketed scores compare only with each other and not with any v4.3 figure [758] [M for costs; V for prices].

Option	\$/M output	Note
<i>Local appliance, Kimi K3 [60], 4 users at ~52 tok/s</i>		
at 50M tokens/month	\$3.16	\$158/mo; light office use
at 100M tokens/month	\$1.65	\$165/mo; ~19% of capacity
at 300M tokens/month	\$0.64	\$191/mo; agentic, ~56% of capacity
at capacity (~540M/month)	\$0.41	\$224/mo
optimistic / conservative price, 100M	\$1.30 / \$2.26	\$3.8k / \$7.3k machine
<i>Rubin VR200 NVL72, Kimi K3, operator's fully loaded cost, 60% utilization</i>		
throughput mode, 4-yr / 6-yr	\$0.70 / \$0.54	~25 tok/s per user
interactive parity, 4-yr / 6-yr	\$1.17 / \$0.90	~50 tok/s per user; \$2.9M/\$2.2M per yr
of which building + facility ops	\$0.14	12%; \$2.5M capex, \$63k/yr
interactive parity at 40% utilization	\$1.76 / \$1.35	demand disappoints
interactive parity at 30% utilization	\$2.35 / \$1.80	the stranded-vintage case
at \$10M enterprise retail, parity, 60%	\$1.44 / \$1.10	via OEM with support; \$3.5M rack-yr
at \$10M enterprise retail, parity, 30%	\$2.89 / \$2.19	
<i>API list price, output tokens</i>		
Claude Opus 5 [63]	\$25.00	
GPT-5.6 Sol [61]	\$30.00	
Grok 4.6 [61]	\$6.00	
Gemini 3.1 Pro	\$12.00	≤200K context
Kimi K3, Moonshot API [60]	\$15.00	
Qwen3.8-Max	\$6.00	
DeepSeek V4-Pro	\$3.96 / \$1.98	peak / off-peak, from 16 Aug

Three readings follow, in ascending order of importance.

First, against the closed frontier the appliance is not close; it is a different order of magnitude. At 100 million output tokens a month—about forty tokens per second sustained, a modest agentic workload—the machine costs \$165 a month all-in, closet and cost of capital included, and \$1.65 per million tokens, against \$25–30 per million for the two leading closed models. The one-time \$5,000 is recovered against Claude Opus 5 in about two months at that load, and in about three weeks at 300 million. The capability gap between the model it runs and the models it displaces is three index points on the v4.1.1 board the table reports; on the re-based v4.3 board the gap between the best closed models and the best open Chinese model is eight (Section 10.3.3), and the arithmetic of the table does not move with it.

Second, against the rack, the appliance does more than reach parity: at realistic utilization its fully loaded cost falls below the rack's own—and the reason is the batching arithmetic that

follows. A hyperscaler serving the same model on Rubin at the same 50 tokens per second per user (interactive parity), building and capital charge included, carries a cost of \$0.90 to \$1.17 per million tokens at its planned 60% utilization, and even in pure-throughput mode (25 tokens per second per user, more tokens but a worse experience) that floor is \$0.54 to \$0.70. The appliance costs more than the rack at light load—\$3.16 per million at 50 million tokens a month, since it is idle most of the time and utilization is the whole story for a fixed-cost box—but crosses below the rack’s interactive-parity cost by 300 million tokens a month (\$0.64 against \$0.90–1.17), matches the rack’s four-year throughput floor at the same load, and at full capacity (\$0.41 per million) undercuts even the rack’s most favorable six-year-amortization throughput number. And the rack’s numbers are its numbers *if the demand arrives*. Utilization is symmetric in the arithmetic and asymmetric in the world: the appliance’s owner chose to buy it because the demand was already there, whereas the rack’s owner built a fifteen-year facility on a forecast, and at 40% utilization the rack’s parity cost is \$1.35–1.76, at 30%—the stranded-vintage case Chapter 21 prices—\$1.80–2.35, three to four times the appliance’s cost at the same tokens per second per user. And that is the hyperscaler’s rack. The enterprise that is actually choosing between a closet and a rack pays enterprise retail for the rack, near \$10M with support, at which its parity cost is \$1.10–1.44 at planned utilization and \$2.19–2.89 at 30%—the appliance’s crossover comes at a lower load, and the gap at disappointed utilization widens to four to five times. Two things keep this honest. It is a unit-cost crossing, not a capacity claim: one appliance serves four concurrent users at ~52 tokens per second each, a rack serves thousands, and nothing here says a fleet of appliances is a substitute for a rack’s aggregate throughput. And note what has *not* happened on the rack’s side of the ledger: it has not amortized its weight reads across a thousand users the way it would for a dense model, because with sixteen of 896 experts routed the batch reads most of the model anyway. The centralized machine’s traditional cost advantage on this workload has not merely narrowed to a utilization advantage—at the appliance’s own full utilization, it inverts, and it inverts further the moment the datacentre’s utilization forecast is wrong.

Third, and this is the strategic content: the API price list sits above both by a multiple that the appliance makes visible. DeepSeek’s own new peak price for a model with half K3’s active parameters is four dollars; the rack’s fully loaded cost of serving K3, building included, runs roughly fifty cents to a dollar and change; the closed frontier charges twenty-five to thirty. The gap between the second and third numbers is the model-layer margin that Chapter 21 documents rising above 70%. The appliance is a \$5,000 device whose function is to let a small enterprise opt out of that margin entirely—and once the machine exists, the pricing question every enterprise buyer asks changes from “which API” to “why an API at all.” That is the innovator’s dilemma stated as a purchase order.

What the model leaves out, stated so it can be checked. On the appliance side: no software labour and no operations labour beyond what the agent stack does for itself; no redundancy; prefill latency on long cold prompts; a memory basis that is pre-surge by construction—at 2Q26 LPDDR prices near \$19 per gigabyte the LPDDR6X alone would cost within a few hundred dollars of what the whole machine costs here; and HBF endurance, which is benign for a machine running one model but must be budgeted for one that reloads a 1.4 TB checkpoint frequently. On the rack side: the landlord’s margin above the bottom-up facility cost, which a neocloud pays and a hyperscaler avoids; the interconnection wait and the discrete-block penalty, both unpriced and both against the rack; the possibility that Rubin outperforms the extrapolation; and, in the other direction, that a rack idling at 30% draws somewhat less than nameplate power, which the model does not credit it. None of these changes the order of magnitude against the closed price list; most of them move the appliance-versus-rack comparison, in one direction or the other, by tens of percent—and the appliance-versus-rack crossover described above is exactly the kind of result that a tens-of-percent shift could move back the other way at planned utilization, which is why the model ships with the manuscript

rather than only its conclusion. What no tens-of-percent shift moves is the utilization asymmetry, and that is the term Chapter 21 takes up.

12.9.5 Why this would matter more than it appears

If the machine works, what changes is not convenience but the structure of the industry, because centralized serving earns its advantage from three amortizations, and cheap capacity attacks all three.

The first is *the shared model copy*: one expensive resident model serving thousands of users. Cheap capacity does not eliminate the advantage but collapses its magnitude, since the scarce asset stops being scarce. The second is *amortized accelerators and scale-up networks*: if weights live in a cheap capacity tier, the machine’s cost migrates from memory to the processor and the surrounding infrastructure—which is precisely the migration this book has argued favours a manufacturing ecosystem over a design-rent ecosystem. The third is *amortized training cost*, and open weights already dissolved that one; a local runner pays none of it.

The most quantitative version of the argument, and the one this book has not seen made elsewhere, concerns batching. A centralized server’s decisive advantage is that it reads each weight once and amortizes it across a large batch: at batch 512 on a *dense* model, per-token weight traffic falls by a factor of 512. But the frontier open models are not dense—they are ultra-sparse mixtures of experts, and Kimi K3 activates sixteen of 896 experts per token. Under uniform routing, the expected number of *distinct* experts touched at batch B is $E[1 - (1 - k/E)^B]$, so the batch reads nearly the whole model while generating only B tokens. The effective amortization is about $1.1\times$ at batch 8, $2.5\times$ at batch 128, and roughly $9\times$ at batch 512—against $512\times$ for a dense model of the same size [M].

The centralized batching advantage on weight bandwidth is roughly fifty times smaller for an ultra-sparse mixture of experts than for a dense model. Three caveats keep it honest: real routing is skewed and token-correlated, which improves amortization above this floor; at large batch the bottleneck migrates from bandwidth to arithmetic, so the advantage changes form rather than vanishing; and the datacentre still amortizes the *capital* cost of memory across users regardless of bandwidth. But the direction is unambiguous, and the irony is sharp: the architecture Chinese laboratories adopted *because* they were compute-constrained is the architecture that most erodes the centralized serving advantage.

12.9.6 The case against, which is strong

The price premise runs against the entire industry. SK hynix’s chief executive said on the record in July 2026 that 2027 would be “the worst year in the industry’s history from the supply perspective” and that demand would exceed supply “even beyond 2030” [579] [V]. Micron and independent analysts concur. A section predicated on falling memory prices must acknowledge that every supplier expects the opposite.

There is one genuine asymmetry, and it is the strongest price evidence available: *NAND is heading toward surplus while DRAM stays scarce*. Industry forecasts put NAND in a modest deficit for 2026 but turning to surplus in the second half of 2027 [580] [A]. If HBF’s cost base is NAND and its competitor’s is DRAM, HBF’s relative position improves through 2027–28 even if nothing else changes. The argument of this section is therefore properly stated as conditional and long-dated, not as a forecast of imminent relief.

The dominant vendor bet the other way, with money. In 2026 NVIDIA removed Rubin CPX—a deliberately cheap, GDDR7-based, capacity-first inference part—from its roadmap, with supply-chain reporting indicating no memory or substrate orders were placed; and it acquired Groq’s assets and acquired its leadership in a transaction reported near \$20B, completed in late December 2025, taking on the *opposite* architecture: a few hundred megabytes of SRAM at extreme bandwidth [581, 582] [A/V]. By August 2026 that bet was a shipping product line

with a disaggregation roadmap attached (Section 12.4.4), which strengthens the objection rather than dating it. The stated logic was that demand had shifted from long-context capacity toward time-to-first-token. This book cannot wave that away. The available rebuttal is real but partial: time-to-first-token is a *serving-provider* metric, produced by competition for interactive users at scale, and a single-user on-premises machine optimizes a different objective. But the revealed preference of the best-informed buyer in the industry currently runs against capacity-first inference silicon, and that belongs on the record before any rebuttal.

The premise that memory is the binding constraint is contradicted by the adoption data. Enterprise open-weight share of workloads fell from 19% in late 2024 to 13% in mid-2025 to about 11% in 2026—*declining across exactly the period in which open-model capability reached parity* [583] [A]. The constraints enterprises actually report are licensing terms, governance and provenance, support, and total cost of ownership. If capability parity did not move enterprises toward open weights, it is not obvious that a memory-price decline would.

And cheaper memory helps the incumbent too. More capacity per accelerator means larger batches and better amortization; every party promoting HBF is promoting it to hyperscalers, and the consortium’s two silicon partners are datacentre companies. A local box also idles nights and weekends, so its effective cost per token is its nameplate divided by its duty cycle—while a serving provider backfills idle capacity with training, evaluation and batch work.

12.9.7 Where China would gain, stated narrowly

The general claim—that cheap memory favours Chinese integrators—is weak, because an input-price fall is global and lowers costs for every system builder. Two specific mechanisms are stronger.

The first is that *HBF’s input requirements map onto Chinese capabilities better than HBM’s do*. HBF needs high-layer-count 3D NAND at scale and mature wafer-to-wafer hybrid bonding. YMTC has both: it entered the top three NAND makers by shipment share for the first time in Q2 2026 at 14%, displacing Kioxia, and its Xtacking process *is* hybrid bonding, in mass production since 2018 [584] [A]. In HBM, China is years behind; in the specific technology stack HBF requires, it is not. Reporting that YMTC has disclosed HBF development plans exists but is single-sourced and this book grades it accordingly [R].

The second is the batching mechanism above: the ultra-sparse mixture-of-experts architectures Chinese laboratories converged on under compute scarcity are precisely the architectures for which centralized batching amortizes worst. That is not a manufacturing advantage but an architectural one, and it was arrived at for unrelated reasons.

Against both: YMTC’s entity-listing bars qualification into Western enterprise servers, which is where the margin is; and China lags in exactly the advanced packaging that analysts say dominates HBF stack cost, so cheap NAND dies do not automatically yield cheap HBF stacks.

12.9.8 The dilemma this creates

Suppose the conditional holds—HBF or a successor delivers terabyte-class weight capacity at flash economics and DRAM-class energy per bit, on the 2028–30 horizon. The consequence for the frontier laboratories is not competition but an *innovator’s dilemma* in its textbook form. Their business model prices intelligence per token, and its unit economics depend on serving many users from one expensive resident copy. A world in which near-frontier capability runs on a capital-cost appliance in a customer’s own building does not reduce that revenue—it reprices the service against a substitute whose marginal cost is electricity. The incumbent cannot follow without cannibalizing the metered business that funds its training runs, and cannot ignore it without ceding the deployment layer.

That is precisely the structure this book documented in Part I. The Western incumbents in solar, batteries and vehicles did not lose because their products were worse; they lost because

the basis of competition moved from a dimension they led to a dimension a manufacturing ecosystem led, and their business models could not follow the move. The claim here is registered at scenario grade, with three observable conditions: HBF or equivalent capacity silicon shipping at its specified capacity and bandwidth, with the endurance figure that a weights tier requires; a memory-price trajectory that turns, or at least a NAND-DRAM divergence that persists; and enterprise open-weight adoption reversing its decline. *None of the three has yet occurred*, and the falsifiers are in Appendix J.

12.10 The West defects to the open ISA: NVIDIA, Google, Meta

The final structural fact of the compute element is that the Western ecosystem is not defending the proprietary-ISA order—it is quietly joining the open one. The evidence spans every layer of the US stack.

A taxonomy first, because “independence” has five distinct layers and conflating them flatters the thesis: (1) host-ISA independence—what RISC-V directly delivers; (2) accelerator-ISA and microarchitecture independence; (3) framework and compiler portability; (4) scale-up/scale-out fabric independence; (5) memory and packaging independence. RISC-V settles only layer 1; the larger AI rents live in layers 2–5, which is where CANN/UB open-sourcing, the matrix-extension standardization, and the memory strategy of this chapter respectively operate. Unit forecasts deserve the same discipline: most RISC-V units are embedded controllers, and unit share is not server-revenue share—though the SHD revenue decomposition ($\sim 70\%$ of projected 2031 RISC-V revenue from AI-accelerator SoCs) shows the value migrating to exactly the contested layer [198].

NVIDIA is the most striking case. Its GPUs have carried embedded RISC-V management cores since 2016—ten to forty per chip, roughly one billion cores shipped in 2024 alone [189]. In July 2025, at the RISC-V Summit *China* in Shanghai, NVIDIA announced that CUDA itself is being ported to RISC-V host CPUs, making the open ISA the third sanctioned CUDA host after x86 and Arm—explicitly enabling “a RISC-V CPU to be the main application processor” of a CUDA system [190]. Two readings of this move are available and both should be on the record. The moat-extension reading: NVIDIA is porting CUDA *onto* the open host layer precisely to keep the accelerator rent (layers 2–3) intact—the RISC-V CPU runs Linux and drivers while the GPU keeps the compute [191]. The concession reading: the platform owner no longer expects the proprietary host order to hold, and is repositioning its rent one layer up—at a summit in Shanghai, weeks after its China market went to zero. The readings agree on the structural fact this book needs: the foundation layer is going open with both superpowers’ participation, and the contest has moved to the layers above it. In January 2026 SiFive announced NVLink Fusion integration for datacenter-class RISC-V subsystems [192]; RISC-V International’s own five-year CEO, Calista Redmond, departed in December 2024—for NVIDIA [193].

In August 2026 that announcement became a specification, and what the specification says is the most important new fact in this section. NVIDIA stated publicly what a RISC-V CPU must be in order to host CUDA alongside x86 and Arm—and *almost none of the requirements are about the instruction set* [673] [P]. RVA23 compliance is asked for, and then the deck points at the platform: the RISC-V Server SoC and Server Platform specifications, the Boot and Runtime Services ratified by RISC-V International in August 2025, and ACPI, identified as critical at the very start of the port—because ACPI is how an operating system discovers the ISA string, the interrupt controllers, the timers and the PCI topology without being rebuilt for each machine. The CUDA-specific asks reduce to two: coherent PCIe I/O, so that host-to-device copies need no cache flush or invalidate, and full-speed PCIe peer-to-peer with the root complex preserving standard ordering. That is the list. It vindicates the layer taxonomy above—layer 1 turns out to be nearly free, and what a platform owner actually demands is firmware, discovery and I/O coherence—and it is a warning to any programme that treats ISA ratification as the milestone that matters. *The barrier to hosting CUDA on an open ISA was never the ISA.*

Table 12.7: What NVIDIA asks of a RISC-V CPU, by tier [673] [P]. Almost none of it is the instruction set, and the requirement document itself is stated to be available to partners only.

Tier	Requirement	Layer it sits in
ISA	RVA23 compliance	instruction set
Platform	Server SoC and Server Platform specs; Boot and Runtime Services (ratified Aug. 2025); ACPI for discovery	firmware, boot, discovery
CUDA over PCIe	coherent PCIe I/O, no flush or invalidate on copies; full-speed peer-to-peer, standard root-complex ordering	host I/O
NVLink Fusion	NVIDIA’s NVLink-C2C soft IP behind C2C hubs, speaking CHI	coherent chip-to-chip

The sting is in the second tier, and it is the sharpest irony this chapter has to report. Hosting CUDA over PCIe is the entry requirement; a custom CPU that wants *coherent unified memory* with the GPU—the configuration that makes a host-plus-accelerator system behave like the single address space this chapter has argued for throughout—must join NVLink Fusion, and to do that it must integrate NVIDIA’s NVLink-C2C soft IP behind C2C hubs and speak CHI, an Arm-originated coherence protocol [673] [P]. It must speak CHI because RISC-V has no standard cache-coherent chip-to-chip interface of its own to speak instead. And the requirement document specifying all of this is stated to be available to partners only. Assemble the three facts and the shape is unmistakable: the open ISA enters the CUDA ecosystem through a proprietary link, speaking a protocol that originates with the ISA it was meant to displace, under a specification that is not public. SiFive appears on NVIDIA’s NVLink Fusion partner slide alongside Arm, Fujitsu and Intel [673]. None of this refutes the defection thesis—the port is real and the host layer is genuinely opening—but it locates the rent precisely. The rent has moved from the instruction set to the coherent interconnect and the partner list, which is layer 4 of the taxonomy above, and RISC-V has no ratified answer to put at layer 4. Section 12.1.1’s missing matrix extension and this missing coherence interface are the same gap seen twice: the open ISA has standardized what a compiler sees and not what a system integrator needs.

Google co-founded the RISC-V Foundation in 2015 and the RISE software consortium in 2023, ships RISC-V in its Titan security silicon, deploys SiFive X280 vector cores inside its TPU ecosystem as the programmable front-end to its matrix units, and committed Android to the architecture [194]. **Meta**’s MTIA accelerator family—four generations disclosed through March 2026, scaling to 10 PFLOPS parts deployed 72 to a rack—is built on RISC-V processing elements throughout [195]. **Microsoft** is the conspicuous holdout, its Azure silicon remaining Arm-based (Cobalt) with no public RISC-V commitment—a gap worth watching precisely because it is now singular [196]. Attempts in Washington to restrict RISC-V collaboration—congressional letters in 2023, 2024, and 2025 noting that Chinese firms hold nearly half the consortium’s board seats—have produced no enforceable action, for the reason the Commerce Department itself conceded: the standard is open, Swiss-domiciled, and restricting US participation would harm US firms without slowing China at all [197].

The other half of the BigSky announcement of Section 12.1.2 belongs here, because it is the entry tier above made into a product. SiFive’s first rackable RISC-V server is specified to act as a *CUDA head node* for NVIDIA GPU clusters [676] [P]—which is to say that the first commercial application of enterprise RISC-V silicon in the West is running the control plane of somebody else’s accelerator. It is a real market and a rational first product; it is not host-ISA independence delivering accelerator independence, and this section’s own taxonomy is what says so.

The structural observation is harder on the Western story than the defection narrative allows. Three Western vendors set out to build RISC-V server silicon. SiFive shipped a server. Ventana was absorbed into Qualcomm without ever bringing its first product to market—Veyron V1 never reached customers and V2 slipped. Tenstorrent is the subject of reported, unconfirmed talks for a Qualcomm acquisition at up to \$10B [677] [A/R]. Consolidation into a single acquirer—whose other motive for RISC-V is its own licensing dispute with Arm—is arguably the dominant Western RISC-V story of the period, and it is not a story about independent competition producing an open ecosystem; it is a story about an open ISA being used as leverage inside a bilateral licensing fight. Two absences complete the picture from the other direction: no Chinese RISC-V vendor presented at the 2026 conference, and no Chinese chip of any kind is named anywhere in its programme [672]. The Western industry consolidates its RISC-V vendors in public; the Chinese one does not appear in public at all, while Chinese boards sit in the test farm of the Western distribution that ships the ISA (Section 12.1.2). No comfort should be drawn from that silence in either direction: absence from a US conference programme in 2026 is a fact about the conference at least as much as about the industry.

The pattern matters more than any single datum: the ISA layer is tipping the way the model layer tipped. Forecasts now put RISC-V at $\sim 25\%$ of processor units by 2030, with AI as the driver [198, 199]. When the platform owner of the AI era ports its crown-jewel software to the open ISA at a summit in Shanghai, the standards war is not being lost; it is being conceded. What remains proprietary in the Western stack—the GPU microarchitecture, CUDA’s implementation, Arm’s cores, and now the coherent chip-to-chip link and the partner list that gates it—sits on top of a foundation both superpowers now agree will be open. The correction the 2026 evidence forces is that the rent moved rather than evaporated: winning layer 1 is necessary, cheaper than anyone expected, and by itself worth very little. Only one of them has organized its national industry to own that foundation, and only the layers above it will decide what owning it is worth.

12.11 Assessment

Honesty about the current balance, and about what this chapter does *not* claim (Section 12.5): NVIDIA-plus-HBM remains the superior machine per watt for dense frontier training today (CloudMatrix-class systems reach parity by brute force at $\sim 4\times$ the power [137]); Ascend software maturity cost DeepSeek a flagship cycle [104]; domestic accelerator output in 2026 (roughly 1.5–2M die-equivalents) remains an order of magnitude below NVIDIA’s [121, 138]; and the LX2’s AI-serving numbers are, as its own analysts insist, projections from published specifications pending direct measurement [185]. But the structural position has inverted over the past year. A GPU-free Chinese machine leads both HPL and HPCG; the CPU-versus-accelerator distinction has been shown, quantitatively, to be an allocation choice rather than a category; the template is portable to an ISA no one can revoke, in two memory variants matched to the two halves of the AI workload; and the Western ecosystem itself—NVIDIA first among them—is porting its software to that ISA. The remaining dependency of the compute element is physical: wafers, HBM stacks, and DRAM bits. That is the third element.

Chapter 13

The Semiconductor: Fabs, Logic, and DRAM

The third element is the physical base. Models can be open-sourced and architectures standardized overnight; fabs cannot. This is where the West’s chokepoint strategy concentrated—EUV denial since 2019, successive equipment and HBM controls, entity lists—and therefore where the empirical test of *chokepoint decay* (Chapter 7) is sharpest. The evidence through mid-2026: the chokepoints are decaying faster than they are being replaced.

13.1 Logic: 5 nm without EUV

SMIC crossed the line that export controls were designed to hold twice. The Huawei Mate 60’s 7 nm Kirin 9000S (August 2023) was the first shock; by December 2025, TechInsights teardowns confirmed volume production on SMIC’s N+3 node—a scaled evolution of its 7 nm-class technology that TechInsights describes as a key indicator of how close SMIC is to a true 5 nm-equivalent node without EUV, while remaining “significantly less scaled than leading commercial 5 nm” from TSMC or Samsung [V] [139, 258]. Capability is demonstrated; TSMC-equivalent economics are not, and the comparison metrics that would settle it—density, yield by die size, wafer cost, cycle time—remain undisclosed [A]. The costs are real: multi-patterning burns wafer capacity, N+2 yields are estimated at 40–55% for large dies (Cambricon’s biggest reportedly ~20%), and margins show the strain [140, 123]. But the strategic arithmetic is not yield percentage; it is wafers times yield times time. SMIC’s advanced-node capacity is climbing from ~45,000 wafers/month (2025) toward 60,000 (2026) and 80,000 (2027) [121]; SMIC posted record 2025 revenue of \$9.3B at 93.5% utilization [141]. Around it, Huawei’s “shadow fab” network—PXW, SwaySure, Pengjin, SiEn, and others identified by open-source analysis, roughly eleven fabs of which eight are 300 mm—adds capacity outside the visible SMIC envelope [142]. Huawei’s die bank (over 2.9 million Ascend dies procured from TSMC before enforcement closed the loophole, plus ~13 million HBM stacks stockpiled) bridges the gap while domestic capacity ramps [121].

Below the leading edge, the mature-node story is already solar-shaped: China holds 37% of global 22–40 nm output in 2026 (42% projected by 2028) and drives nearly 70% of global 300 mm mature-node capacity expansion, with trailing-edge price pressure forcing Taiwanese incumbents through the familiar sequence of supplier squeezes and order migration [143]. Every power stage, microcontroller, and sensor of the EV and robotics stack rides on these nodes—and recall (Chapter 5) that Chinese automakers are directed toward 100% domestic chips by 2027 [144].

13.2 DRAM: CXMT and the swing-supplier position

Memory is the quiet revolution. CXMT grew from $\sim 70,000$ wafers/month in 2022 to 200,000–300,000 by mid-2026, targeting 600,000 via new Hefei and Shanghai fabs; independent modeling puts it within $\sim 25,000$ wafers/month of Micron by end-2026—which would make China’s national champion the world’s fourth, and trajectory-wise third, DRAM producer [145, 146]. Its DDR5 yields reached 80%+; it publicly launched DDR5-8000 and LPDDR5X-10667 in November 2025; and its July 2026 Shanghai IPO raised \$8.6B and closed +466% at a $\sim \$460$ B market capitalization—instantly the largest A-share company [147, 148, 259]. Three disciplines keep that number honest. First, only 6.73% of shares floated, so the print measures state-channelled investor appetite in a thin market rather than industrial value—first-day market capitalization is a *sentiment* indicator and nothing more [V] [259]. The industrial indicators are the ones to track instead: IPO proceeds and their allocation, wafer capacity by node, revenue, gross margin, yield, customer qualification, and the capex plan. Second, the prospectus disclosed *no dedicated commercial HBM investment*, which sits awkwardly beside the narrative—this book’s included—that CXMT HBM3 is an imminent strategic bridge. The contradiction should be stated rather than smoothed: either HBM development is funded outside the disclosed lines, buried inside general R&D, or less mature than the reporting implies. All three readings are live, and the resolution is a disclosure this book does not have [A]. Third, the swing-supplier evidence (CXMT server modules pricing above Samsung’s) is a moment in a panic market, not proof of durable pricing power: qualification status, channel inventory, specification, shortage timing, and regional procurement all move that number.

That last fact is the essential one, and it deserves the full accounting, because domestic HBM is now the single binding physical constraint on the compute element—and the most important open variable in this book.

13.2.1 The HBM ledger

The state of Chinese HBM as of mid-2026, stated carefully by evidence grade. *In production*: CXMT HBM2 since 2H24, 8-high HBM2E since 1H25—roughly three to four years behind SK hynix, on a ~ 16 nm-class cell with a $>30\%$ cost-per-bit disadvantage, and today under 2% of CXMT’s $\sim 265,000$ monthly wafer starts [149, 148]. *Sampled*: CXMT delivered HBM3 samples to Huawei in October 2025, with large-scale HBM3 manufacturing slated for 2026 and HBM3E for 2027; Korean-sourced reports have CXMT dedicating $\sim 20\%$ of capacity to an HBM3 line [200]. *Announced*: Huawei’s Ascend 950PR ships in 2026 with self-developed “HiBL 1.0” HBM (128 GB, 1.6 TB/s) and the 950DT with “HiZQ 2.0” (144 GB, 4 TB/s)—whose DRAM dies are fabbed by an undisclosed partner, with CXMT and the entity-listed Wuhan Xinxin (XMC, $\sim 3,000$ HBM wafers/month with SJSemi packaging) the documented candidates [122, 201]. *Claimed*: Chinese state media assert LineShine runs on “China’s first domestically developed HBM”—which would make the world’s No. 1 supercomputer the first at-scale deployment of Chinese HBM—while independent analysts note the supplier is undisclosed and speculate CXMT HBM2E [187, 186]. *Not yet verified*: no public teardown has confirmed Chinese-fabbed DRAM dies in a commercially shipping accelerator; teardowns of shipping Ascend 910Cs still find Samsung and SK hynix stacks [202].

Those foreign stacks are the depleting bridge: Huawei accumulated ~ 13 million HBM stacks (including ~ 11.4 M from Samsung, ~ 7 M of them shipped in the December 2024 window between the BIS rule’s announcement and its enforcement), and by mid-2026 that stockpile is assessed as effectively exhausted—making HBM, not SMIC logic, the binding constraint on Ascend output (~ 2 M domestic stacks projected for 2026, sufficient for only ~ 250 – 300 k Ascend 910C-class packages; more bullish estimates reach ~ 7 M) [121, 203]. The advanced-packaging leg, by contrast, is scaling fast: Innotron’s \$2.4B TSV/stacking facility, Tongfu’s RMB 4.4B raise, over RMB 27B poured into packaging expansion in H1 2026, and a domestic HBM assembly toolchain (Naura

TSV etch, Maxwell hybrid bonding) targeted at end-2026 HBM3 [150, 204]. Washington noticed the champion late: CXMT reached the DoD’s 1260H list only in June 2026 and remains off the BIS Entity List—after the IPO, after the LPDDR6 verification, after the HBM3 samples [148].

The strategic read is double-tracked, and Chapter 12’s template (Section 12.3) is its exact mirror. On the HBM track, China is three-plus years behind and sprinting—independent assessments (TrendForce) still place domestic HBM in early verification with limited near-term global impact [A] [260]—with the inference-centric machines (Ascend 950, LX2-class sockets) as the guaranteed captive demand. On the commodity track, CXMT is at the *world’s leading edge* of LPDDR6—sampling concurrent with SK hynix, mass production targeted 2H26 [R]—feeding the training-centric, capacity-first architecture designed to need far less HBM [135]. The precise formulation matters, because the looser version of it is false: China has not made HBM irrelevant; it is engineering a position in which HBM ceases to be the *sole* architectural path to frontier-scale AI, while racing to supply it domestically anyway. Section 13.3 then adds a complication the ledger above cannot absorb without revision: during the generation China is chasing, the thing being chased stopped being a single specification.

13.2.2 The 2026 memory wall is priced, not clocked

Meanwhile the 2025–2026 global memory shortage—DRAM contract prices up 50–60% in quarters, driven by AI demand and HBM’s $3\times$ wafer intensity—handed CXMT the swing-supplier role: its 64 GB server DDR5 modules now price *above* Samsung’s [151, 152]. Within twelve months, the narrative flipped from “Chinese dumping” to Chinese pricing power, and that evidence remains a moment in a panic market rather than proof of durable pricing power [A]. YMTC’s 294-layer NAND and the revival of Entity-Listed Fujian Jinhua complete the memory picture [153]. By mid-2026, though, “shortage” had become the wrong word for what the series show, and the correction matters here because it changes which side of this book’s argument the memory market sits on.

The most complete public accounting is Jim Handy’s tutorial at Hot Chips 2026, and it should be read with its conflict on the table—the deck itself discloses that the data underpin a paid report, so these are an analyst’s series and not an audited one [A] [678]. What they show is a price event, not a capacity event. The lowest spot DRAM price is up $7.1\times$ and NAND $6.9\times$ from the mid-2025 trough. Combined DRAM and NAND revenue at the five major suppliers has gone from a $\sim\$15\text{B}$ quarterly trough in 2Q23 to $\sim\$180\text{B}$ in 2Q26—a twelvefold move that is almost entirely price and almost not at all bits. Hyperscale property and equipment spending rose over the same era from $\sim\$30\text{B}$ a quarter in 1Q21 to $\sim\$170\text{B}$ a quarter in 2Q26. And the supply side did not move: DRAM wafer starts have been flat for more than a decade, and new capacity takes over two years to install, so no supplier can answer this on the clock even if every one of them has decided to. HBM is the mechanism—roughly $3\times$ the die area per bit of DDR, so each stack sold withdraws several stacks’ worth of commodity bits from the market—and the consequence surfaces not as an allocation queue but as a price. Handy’s summary slide is a photograph of banknotes.

The two price series in the preceding paragraphs measure different things, and both are needed. Contract prices up 50–60% in a quarter is what a qualified customer with a long-term agreement pays; a $7.1\times$ move off the trough in the lowest spot print is what a buyer without one pays—and the buyers without one are precisely the marginal actors this book cares about: systems integrators, second-tier clouds, national laboratories, and anyone standing up capacity that was not in a 2024 forecast. Two vendor figures fix the magnitude at the system level. Samsung’s own accounting puts memory at 63% of AI-chip component spend in 4Q25, up from 52% in 1Q24 [P]. Micron’s is more arresting: an eight-stack HBM4 package carries $12,344\text{ mm}^2$ of memory silicon, more than $8\times$ the die area of a typical GPU [P] [679]. On an AI accelerator the memory is now the majority of the silicon and the majority of the bill of materials; memory price sets system price, and the logic vendor’s margin is what absorbs the residual.

This cuts both ways for the argument of this book, and the inconvenient direction should be stated before the convenient one. It raises the cost of the West’s HBM-centric architecture, which is the capacity-first thesis of Chapter 12 being priced by the market rather than argued on a slide: if memory is 63% of the bill and the HBM fraction of it is the most expensive memory ever manufactured, then an architecture substituting capacity for bandwidth is buying the cheaper end of a market whose expensive end is running away. But it raises the cost of the cheaper end too. The ~ 2.3 TB commodity socket of Section 12.8 is denominated in exactly the bits whose spot price rose $7.1\times$, and the appliance of Section 12.9.4 in the NAND that rose $6.9\times$. *A memory price shock is not unambiguously good news for the side that chose to buy more memory of a cheaper kind*—it is good news only insofar as the cheap kind stays cheaper in *relative* terms, which the HBM die-area arithmetic supports and the flat wafer starts do not guarantee. Where this book prices a capacity-first machine it does so on a pre-surge basis, and that basis is now a modeling assumption rather than an observation [M].

What the shock does do for China is narrower than a price move and more durable. First, it converts CXMT’s capacity into revenue at the moment that revenue funds the next tranche of capacity; the IPO proceeds and the price cycle arriving in the same quarter is the luckiest piece of timing in the domestic memory program. Second, the two-year installation lag binds everyone equally, and CXMT’s wafer additions are among the largest single tranches of new commodity DRAM capacity being installed anywhere—in a market that cannot be answered on the clock, the fastest-growing marginal supplier gains share by default. Third, and this is the term the controls do not reach: the shortage is a consequence of the West’s own architectural choice. HBM’s die-area intensity is what withdrew the commodity bits, and the country that elected not to build HBM at scale is the country whose fab capacity is not being consumed by it.

13.3 HBM4 stopped being a specification

The ledger of Section 13.2.1 measures a gap in years. That measurement presumes the thing being chased holds still enough to be chased, and the 2026 disclosure record—the three incumbent suppliers’ own decks—shows that within one generation it stopped doing so. Two changes arrived together, and between them they alter what “three to four years behind” means. The conference record in full is Appendix C; what follows is the part bearing on this chapter.

13.3.1 A $1.5\times$ spread inside one nominal generation

The same JEDEC HBM4 generation is quoted by SK hynix at 8 Gb/s and 2.0 TB/s per stack, by Micron at 11 Gb/s, and by Samsung at 3.0 TB/s—a $1.5\times$ spread in per-stack bandwidth within one nominal generation [P] [679]. A bill of materials specifying “HBM4” therefore does not specify a part, and the generation’s two flagship accelerators made opposite choices from it. AMD’s MI455X bought the slow bin and compensated with stack count: twelve stacks, 432 GB, 23.3 TB/s, an implied ~ 7.6 GT/s. NVIDIA’s Rubin bought the fast bin and stayed at eight: 288 GB, 22.2 TB/s, an implied ~ 10.8 GT/s—having re-specified onto the faster bin partway through the program, with the per-GPU envelope rising from 1.8 to 2.3 kW. The re-specification is legible in NVIDIA’s own record, since the March 2026 launch material described the same 288 GB part at ~ 13 TB/s [P] [510, 591]. The two designs land within 5% of each other on aggregate bandwidth by entirely different routes—one buying capacity and pin count, the other signalling rate and power.

The procurement consequence is a sentence rather than an analysis, and it is the most immediately usable finding in this chapter for anyone who buys machines: *specify terabytes per second per stack in the contract, not “HBM4.”* An agreement written to the generation name can be filled, legally, at two-thirds of the bandwidth the buyer modeled.

What did *not* improve is the strategically interesting half. The prefetch has been 256 bits since

HBM1 in 2014, and the die density is 24 Gb in both HBM3E and HBM4. HBM4's bandwidth gain— $1.6\times$ to $2.4\times$ over HBM3E depending on which vendor's bin is bought—therefore comes almost entirely from doubling the interposer I/O to 2,048 pins, with the doubling of banks per die to 256 reducing conflicts rather than adding raw bandwidth, and its capacity growth comes only from stack height [679]. *That is a packaging problem, not a DRAM problem.* It relocates the frontier off the memory cell—the layer on which CXMT's ~ 16 nm-class process is measurably behind and measurably closing—and onto the interposer, the bump pitch, the stack and the bonding step, which is the layer Section 13.4 treats separately and Eq. (A.1) prices. For a country whose lithography is denied and whose packaging investment runs at RMB tens of billions a year, that relocation is on its face favourable. It is also the layer at which the second change bites.

13.3.2 The base die became a logic die

From HBM4 the base die of an HBM stack is fabricated on a 4 nm foundry logic process [P] [680]. That single fact is the inversion. A memory vendor buying leading-edge foundry logic capacity to build the bottom die of its own product is a chip designer, and Samsung's roadmap states what it intends to design. Phase one reclaims accelerator area by absorbing the memory controller into the base die. Phase two fills the freed base-die area with repair, telemetry, test and processing elements. Phase three deletes the interposer entirely, bonding the cell-array stack directly onto the accelerator—"zHBM," claimed at a 70% reduction in I/O energy and 230% of HBM4E bandwidth, with the ~ 100 W of DRAM power saved re-spent as an 8.3% rise in accelerator power against a 1,200 W baseline [R; vendor roadmap, unaudited] [680]. And the commercial form is stated rather than inferred: *custom HBM—"cHBM"—is explicitly not JEDEC-standardized: a customer-specific logic base die on that 4 nm process, sharing standard cell-array stacks.* The two names are not interchangeable, and this book keeps them apart: cHBM is the phase-one and phase-two product, the co-designed logic base die; zHBM is the phase-three concept that deletes the interposer altogether by bonding the cell-array stack directly onto the xPU. Strategically, the memory vendor comes to own an interface that belonged to the accelerator designer.

The consequence this book must draw is not the one the deck draws. Samsung's framing is technical—area, energy, bandwidth. The structural change is that an HBM generation stops being a parameter set any qualified buyer can order against and becomes a per-customer product negotiated between a memory vendor and an accelerator team. That is good for large buyers with silicon teams—the hyperscalers, the two merchant GPU vendors, the handful of ASIC programs with a Broadcom or a Marvell attached—and bad for everyone else, national laboratories buying at merchant volumes very much included. The point is worth stating plainly rather than diplomatically: a centre procuring a few thousand accelerators does not get a logic die designed for it, and will increasingly buy the residue of a market whose leading products were co-designed for somebody else.

Score that against the ledger's "three to four years behind," because a moving generational target is one thing and a market that has stopped having generations is another. Two readings are live and the public evidence does not separate them. The pessimistic one for China: catching HBM3E in 2027 no longer catches the frontier, because the frontier part is not a generation but a relationship—a logic die co-designed with an accelerator team—and access to accelerator design teams is precisely what the controls deny. On that reading the gap stops being measured in years of DRAM process, where it was closing, and starts being measured in bilateral commercial relationships China is not party to, where it cannot close at all. The optimistic one: a JEDEC generation is a target and a set of private bilateral products is not, so the *merchant* HBM part—the one anyone can buy—stops advancing at the frontier rate once the fastest bins and the custom base dies are absorbed into co-designed products first. If the merchant market slows, the target CXMT is chasing slows with it. And China already runs the custom-HBM pattern inside a single national vertical: Huawei's HiBL 1.0 and HiZQ 2.0 are an accelerator designer specifying its own memory to an undisclosed domestic fab partner, which is structurally what

Samsung now proposes to sell to everyone else. Vertical integration, the thing sanctions forced on the domestic ecosystem, is on this axis an advantage.

The chapter’s finding is therefore restated rather than withdrawn, and the restatement is longer because it must now name two clocks: *Chinese HBM is three to four years behind the merchant part, and outside the co-design relationships that increasingly define the frontier part*—the first gap closing at a measurable rate, the second not measurable at all from public evidence, and neither settled by the other [A].

13.4 Packaging: the chokepoint hiding inside the memory ledger

Advanced packaging deserves separation from the HBM discussion, because it is a distinct chokepoint with distinct suppliers and a distinct clock. Producing logic and producing memory are insufficient if a country cannot *assemble* them at yield and volume: through-silicon vias, hybrid bonding, interposers, large organic substrates, CoWoS-class capacity, known-good-die test, thermal interfaces, and increasingly optical and SerDes packaging are all separate industrial competences, and OSAT scaling is its own capital cycle. China’s position here is markedly stronger than in lithography and markedly weaker than in assembly labour: Innotron’s \$2.4B TSV facility, Tongfu’s RMB 4.4B raise, more than RMB 27B of packaging expansion in H1 2026, and a domestic HBM assembly toolchain—TSV etch, hybrid bonding, precision placement—targeted at end-2026 HBM3 [150, 204]. The strategic reading is that packaging is the layer where the domestic ecosystem is closing fastest *and* the layer whose failure would strand both the logic and the memory investments; it is therefore the highest-variance item in Table 13.1.

The 2026 disclosure record sharpens that judgement in three specific places, all of them from SK hynix’s own packaging material and all physical rather than commercial [P] [672, 679]. First, stack height has become a warpage problem: going from twelve dies to sixteen takes package height only from 720 to 775 μm , so the die thickness falls to about $0.9\times$ and the inter-die gap roughly halves; beyond twenty dies the thermocompression flow does not survive at all and hybrid bonding becomes mandatory rather than optional. Second, the PHY power density rises from 0.5 W/mm^2 at what SK hynix labels sHBM4E to over 2.0 W/mm^2 at sHBM5, crossing a line the vendor’s own slide marks “system limit” at exactly the HBM4 generation—the interface, not the array, is what runs out of thermal budget. Third, and most consequential for yield economics, the assembly order has inverted: the memory stack is now attached before the rest of the package is complete, so an HBM defect scraps the entire system-in-package rather than a stack—and no vendor published a yield figure of any kind. Eq. (A.1) supplies the arithmetic; what changed is that the terms now multiply in an order which makes the memory yield the most expensive one to get wrong.

For the domestic program the reading is mixed and the two halves point opposite ways. The favourable half is that the frontier has moved onto the layer where China is investing hardest and where the tooling is least dependent on the lithography it cannot buy. The unfavourable half is sharper: the least mature item on the domestic HBM assembly list—Maxwell’s hybrid bonding, targeted at end-2026 HBM3—is exactly the tool that stack heights beyond sixteen make compulsory, so the domestic schedule’s most speculative element is the one the next generation needs first [204]. Packaging remains the highest-variance row in Table 13.1, and on this evidence the variance has widened rather than narrowed.

13.5 A sanctions-elasticity model

This book has used “chokepoint decay” as a descriptive phrase. It can be made semi-quantitative, and doing so is the honest way to compare controls that behave very differently. For each controlled input, write Δt_{eff} for the *effective delay*: the number of years by which a control postpones the adversary’s attainment of the capability it targets, relative to a no-control counterfactual. It is

not a formula but a function of five terms, each of which subtracts from the delay the control’s designers assumed they were buying:

$$\Delta t_{\text{eff}} = \Phi \left(\underbrace{\text{stockpile depth}}_{\text{years of inventory held}}, \underbrace{\text{diversion}}_{\text{share obtainable anyway}}, \underbrace{\text{domestic substitution rate}}_{\text{years to a home-grown equivalent}}, \right. \\ \left. \underbrace{\text{architectural redesign}}_{\text{can the need be designed away?}}, \underbrace{\text{yield penalty}}_{\text{cost of the inferior path}} \right),$$

where Φ is left unspecified because no public dataset supports fitting it; what the decomposition buys is not a number but a checklist, and the terms are ordered so that the first three shorten the delay and the last two lengthen it. The terms have very different magnitudes by input. EUV lithography has no stockpile, negligible diversion at machine scale, a substitution rate measured in decades, and no architectural detour that fully substitutes—so the control buys real years. HBM has a large stockpile (exhausted by mid-2026), meaningful diversion, a substitution rate of three-plus years, and—critically—an *architectural detour* in the capacity-first design of Chapter 12, so the control buys less than its designers assumed. Advanced accelerators have shallow stockpiles, documented diversion at billion-dollar scale, a fast-improving domestic substitute, and a full architectural detour; the control has largely converted into a domestic market-reservation policy. EDA proved to have near-zero delay value: the ban was imposed and rescinded within six weeks, and domestic share doubled in the interval. Deposition, etch, and metrology sit in between.

The 2026 memory market exposes a sixth term the five omit, and it is not a technical one: the *price of the substitute*. An architectural detour is only as cheap as the components it detours onto, and Section 13.2.2 shows those components repricing by multiples inside eighteen months for reasons—HBM’s die-area intensity, a decade of flat wafer starts, a two-year installation lag—that neither government controls and neither government’s controls can reach. The corollary is uncomfortable but follows directly: a control’s effective delay is partly a function of a global commodity price, so any Δt_{eff} quoted without a date attached to it should be distrusted, and the HBM row in particular has a wider band in 2026 than the same analysis would have carried in 2024. Ranking controls by Δt_{eff} rather than by strategic-sounding importance is, on this book’s evidence, the single most useful analytical discipline available to export-control policy—and it would have predicted, in 2022, most of what has since happened.

13.6 Equipment, EDA, materials: the long ladder

The deepest dependency is the tool base, and here the trend, not the level, is the message. Naura reached #5 in global semiconductor-equipment revenue in 2025 (from #8 in 2022, revenue ~\$5.5B, +31%); AMEC (#13) ships etch tools aimed at ≤ 5 nm logic; three Chinese firms now sit in the global top twenty; domestic tools supply an estimated 20–30% of Chinese fab demand, up from ~10% three years prior [154, 155]. The Huawei-linked, state-owned SiCarrier exhibited some thirty tool types within two years of founding and holds patents claiming 5 nm-class SAQP-DUV flows [156]. Lithography remains the hard floor: SMEE ships at 90 nm, 28 nm-class immersion DUV is in trials (Yuliangsheng), and domestic EUV—LDP-source prototypes reportedly under test at Dongguan—is credibly dated no earlier than ~2030 by sober assessments [157]. The pattern of every denied layer repeats in EDA: the May 2025 US ban on EDA sales was rescinded within six weeks (traded against rare-earth flows—Chinese leverage from the previous wars, spent to protect the next one), but not before Emphyrean shipped China’s first full-process memory-EDA platform and domestic EDA share doubled [158]. Silicon wafers: 3% of 300 mm capacity in 2020, ~28–32% in 2026 [159].

The money behind the ladder: Big Fund III alone is RMB 344B (~\$47.5B) with a horizon to 2039, atop Funds I–II and provincial vehicles; China has been the world’s largest equipment

buyer (~40% of global WFE) every year since 2020 and will remain so through 2027 [160, 161]. Self-sufficiency targets are missed and re-set—the “Made in China 2025” 70% goal landed nearer 25–50% depending on definition [162]—but the direction never reverses, and each US control action (December 2024’s 140-entity action, the whiplash of the AI Diffusion Rule imposed January 2025 and rescinded May 2025) functions as demand guarantee for the domestic substitute [163, 164]. Meanwhile enforcement leaks at industrial scale: a \$2.5B smuggling scheme via Supermicro employees was charged in March 2026 [165].

13.7 The national-open logic at the silicon layer

The semiconductor element lacks the visible “openness” branding of models and ISAs—one cannot open-source a fab. But the same national-open structure operates: open architecture above (RISC-V cores taped out on SMIC; XiangShan as shared national IP), unified standards across competing fabs and OSATs, state capital socializing the capex, and the deliberate choice of *commodity* process technology—DUV multipatterning, LPDDR6, chiplet packaging—over the proprietary-premium branch (EUV, HBM4) wherever the commodity branch can be scaled domestically—a choice that Section 13.3 shows has partly been made for it, because the premium branch has stopped being purchasable at all except by co-design. Table 13.1 summarizes the position.

The element-level conclusion, stated at the evidence grades the table carries rather than above them: no single Chinese node matches its Western counterpart at the frontier, and the architecture of Chapters 10 and 12 was chosen so that none needs to. But “sufficient” is a claim about an integrated system, and the honest formulation is conditional: *the emerging domestic stack contains plausible paths to sufficiency, each supported to a different evidence level—logic capability [V] but logic economics [A], LPDDR roadmap [A/R], packaging capability [P/A], software maturity [M]—and an end-to-end frontier training run remains the decisive integration test that none of the paths has yet passed.* Logic yield, packaging economics, memory qualification, software maturity, and LPDDR training performance have never been demonstrated *simultaneously in one run*; the claim register carries that run as CR-09, the “Mate 60 moment” of AI. What the table does establish unconditionally is the strategic geometry: the chokepoint strategy assumed China would try to rebuild the Western machine and could be stopped at its scarcest parts. China is building a different machine—whether the different machine works end to end is the open question this book’s forecasts price at 65% within 2–4 years. What 2026 added to that question is a variable neither side chose and neither side controls: the price of memory, which moved by multiples in eighteen months, which prices both machines, and which is now the largest single source of uncertainty in every cost comparison this book makes.

Table 13.1: The semiconductor element, mid-2026: chokepoints and their decay state.

Chokepoint	Western position	Decay state
Leading-edge logic	TSMC 2 nm; EUV denied	SMIC 5 nm-class DUV in volume; 60–80k wpm ramp; yield-limited, not capability-limited
HBM, merchant part	SK hynix/Micron/Samsung lead; export-controlled; $1.5\times$ per-stack bandwidth spread inside one HBM4 generation	Potentially reduced as a sole dependency by the modeled capacity-first architecture [M]—not yet bypassed in any measured training system; domestic HBM2E + now-exhausted stock-piles bridge
HBM, co-designed part	Custom logic base die on 4 nm, explicitly not JEDEC-standardized; sold to buyers with silicon teams	No decay path visible: the barrier is a commercial relationship, not a process step [A]. Huawei’s HiBL/HiZQ is the same pattern run inside one national vertical
Advanced packaging	CoWoS-class capacity, TSV, hybrid bonding, known-good-die test; the layer HBM4’s gains now come from	Closing fastest of any layer (Innotron, Tongfu, >RMB 27B in H1 2026, domestic TSV and hybrid-bonding toolchain at end-2026 HBM3) [P/A]—and the highest-variance row here, since ≥ 20 -high stacks make the least mature domestic tool compulsory
LPDDR6	JEDEC-standard commodity class	CXMT roadmap targets a schedule comparable with leading suppliers [A/R]; qualified mass production not yet independently verified
DRAM commodity	Samsung/SK/Micron oligopoly	CXMT #4 and rising; swing supplier in shortage; pricing above Samsung
Lithography	ASML monopoly	Hard floor; 28 nm DUV trials; EUV ~ 2030 ; mitigated by multipatterning + capacity
Equipment (non-litho)	AMAT/Lam/TEL lead	Naura #5 globally; 20–30% domestic share and rising
EDA	Synopsys/Cadence duopoly	Ban imposed and rescinded in 6 weeks; domestic share doubling from low base

Chapter 14

The Watt: Energy as the Fourth Element

The three elements of Part II share a silent denominator. Every token is ultimately a quantity of electricity passed through silicon, and the AI war’s longest-lead-time input is neither the model, the chip, nor the fab—it is the watt.

This chapter states the asymmetry carefully, because the obvious comparison is the wrong one. *Nameplate capacity added is not deliverable power at a computing site*, and a book that lets a solar gigawatt stand in for a gigawatt of 24/7 data-centre supply has made exactly the category error it accuses others of making elsewhere. Table 14.1 therefore replaces the single headline figure with a four-row ledger, and the argument that follows rests on the rows that survive it: not that China adds ten times more nominal capacity, but that it combines much faster buildout with long-distance transmission, state-directed siting, retained dispatchable firming, and—decisively—the ability to attach large new loads on a schedule the American interconnection queue cannot match.

14.1 The supply asymmetry

The 2025 numbers, stated at their correct denominators [V]. China’s national energy administration reports *452 GW of new renewable capacity* in 2025—318 GW solar and 120 GW wind—within total additions of roughly 540 GW, bringing installed capacity to 3.89 TW at year-end and about 4.04 TW by mid-2026, on ~\$500B of annual energy investment. The US Energy Information Administration reports 53 GW of new utility-scale capacity in 2025, including storage—a record American year [263, 264, 357].

Three qualifications belong with the table. First, storage is not primary generation, and a national aggregate surplus does not imply firm deliverability at any particular computing hub—the correct unit of analysis is the balancing area or the provincial grid, not the country. Second, falling average generator utilization hours are consistent with several stories besides surplus: rapid renewable additions, curtailment, congestion, weak demand, or displacement of thermal generation. The claim that the marginal Chinese AI data centre bids into genuine headroom therefore requires *regional* evidence—average and marginal industrial tariffs, curtailment and congestion, firm generation and storage, transmission import capacity, and time-to-interconnect a 100–500 MW load—for Inner Mongolia, Gansu, Guizhou, Guangdong, and the Yangtze River Delta specifically. That evidence is assembled unevenly in public sources and is the single largest gap in this chapter [A]. Third, global context: the IEA’s projection of ~945 TWh of data-centre consumption in 2030 remains under 3% of world electricity—the *stress is regional and temporal, not globally energy-limited* [281]. Nothing in this chapter should be read as a claim that the planet runs out of electricity; the claim is that particular buildouts, in particular interconnection queues, on particular schedules, do. China’s generation ran 10,573 TWh in

Table 14.1: Four denominators, not one. A gigawatt of solar nameplate is not a gigawatt of firm data-centre supply; the strategic claim must be built from the rows that survive the distinction.

Metric	China	United States
New nameplate capacity (2025)	~540 GW total; 452 GW renewable (318 solar, 120 wind) [V]	53 GW utility-scale incl. storage [V]
Expected annual generation	capacity-factor adjusted; generation 10,573 TWh, +503 TWh y/y [V]	4,430 TWh, growth modest [V]
Firm / dispatchable contribution	record 95 GW coal commissioned as flex; ~66 GW nuclear toward 110 GW by 2030; two-thirds of world grid storage [V/R]	gas turbines sold out to ~2030; <1.7 GW of marquee nuclear restarts [V]
Deliverable data-centre headroom	hub-specific: 45 UHV lines, four new 8-GW UHVDC corridors in 2025; state-directed siting into curtailed-renewable regions [V/P]	balancing-area specific: ~2,060 GW queue, >3-yr median to agreement; 12+ states with moratorium bills [V]

2025 against America’s 4,430— $2.4\times$ —and its *annual demand growth* (+503 TWh) now exceeds Germany’s total consumption [265]. The composition compounds the point (wind plus solar now 1.84 TW, nearly half the fleet, with clean generation passing 40% of output in H1 2026); ten-reactor-per-year nuclear approvals toward 110 GW by 2030; a record 95 GW of coal commissioned as dispatchable flex; and roughly two-thirds of the world’s grid-battery installations—one Chinese *month* of BESS deployment in 2025 exceeded the US full year [263, 266, 267, 54]. Around it, the delivery machine: dozens of UHV lines moving western generation eastward, four new 8-GW UHVDC corridors entering service in 2025 alone, and roughly 28 more lines planned by 2030 [268, 376]. Capacity is running ahead of demand—average generator utilization hours are *falling* [263]—but the natural conclusion from this overreaches, and the disciplined statement is the one that belongs on the record: *the national data are consistent with aggregate surplus, but do not establish firm, deliverable surplus at a specific computing hub*. The distinction is not pedantic: 2025 national wind/solar curtailment ran 5.4%, breaching the official 5% target, and Q1 2026 saw large weather-adjusted capacity-factor declines concentrated in exactly the hub provinces (Inner Mongolia foremost)—driven less by missing wires than by inflexible coal contracts and annual provincial trading, which is a *market-design* constraint on deliverability that a data center experiences just as physically as a congestion constraint [375]. The busbar casebook of Section 14.3 is where the hub-level evidence now lives.

The United States is the mirror image: demand arriving into scarcity. Here a category error runs through much of the public debate, and disentangling it sharpens the argument rather than weakening it. The famous ~2,000-GW-plus interconnection queue is a queue of *proposed generation and storage projects* seeking transmission interconnection—roughly 2,290 GW active at the end of 2024, with median request-to-operation times above four years and only 13% of 2000–2019 requested capacity ever built [370]. It says nothing directly about how long a 100–500 MW *data-center load* waits. The US constraint properly decomposes into five queues that compound: (1) the generation-side queue above, which throttles the supply additions needed to serve new load; (2) *large-load connection* studies and agreements, for which no comprehensive national statistic exists—the definitive study of the problem states plainly that no cross-utility survey has been done—but where the observable instances are grim: utility load studies of 6–18 months followed by 18–60 months of construction where line extensions are needed, and Northern Virginia transmission-level hookups quoted at up to seven years against 18–24-month shell construction [371, 372]; (3) transmission expansion, on 7–10-year planning-to-energization cycles; (4) equipment—gas turbines effectively sold out into 2029–2030, large power transformers




Figure 14.1: The fourth element. *Above*: annual additions to generating capacity, China against the United States—China added roughly ten times as much in 2025, into falling utilization hours, while the US record year was 53 GW [V] [263, 264]. *Below*: the American constraint stated three ways—what was built in the best year on record, the 100 GW/year one AI laboratory told the White House the country needs, and the $\sim 2,060$ GW sitting in the interconnection queue at median waits over three years [V/P] [269, 277]. One side’s binding constraint is capital; the other’s is physics and permitting, and permitting is the harder one to finance away.

at 24-to-60-month lead times; and (5) resource adequacy, now priced: PJM’s capacity auctions have cleared at or near their caps three years running, \$16.4B for 2027/28 and again for 2028/29, with the RTO missing its own reliability requirement in both [274, 275, 272, 373]. Data-center load, 62 GW in 2025, is projected to more than double by 2030, driving 90 of the 166 GW of forecast five-year peak-load growth [270, 271]; PJM’s capacity auctions have set three consecutive records (\$16.4B for 2027/28, ~94% of load growth attributed to data centers), ratepayer politics has produced data-center-moratorium bills in at least twelve states, and the physical inputs are queued years deep—gas turbines effectively sold out to 2030 (GE Vernova’s backlog at 116 GW, taking reservations for 2031), large transformers at two-year-plus lead times and prices up 70% since 2019 [272, 273, 274, 275]. The marquee nuclear responses—restarting Three Mile Island for Microsoft (2027 target) and Palisades (slipping)—total under 1.7 GW: measurement error against the Chinese fleet [276]. The system’s own principals have said the quiet part: Microsoft’s CEO—“it’s not a supply issue of chips; it’s actually the fact that I don’t have warm shells to plug into”—and OpenAI’s letter to the White House demanding *100 GW per year* of new US capacity, justified explicitly by China’s 429-GW year and an “electron gap” [227, 277]. The ask is nearly double the best year the US grid has ever had.

14.2 Energy as industrial policy: the token factory’s power bill

The asymmetry is not merely quantitative; it is being *aimed*. The East-Data-West-Computing program placed eight national computing hubs where the curtailed renewables are; 2024 policy requires 80% renewable sourcing for new hub data centers; and in November 2025—in the same week Beijing barred foreign AI chips from state-funded data centers—provinces including Gansu, Guizhou, and Inner Mongolia cut electricity prices *up to 50%, to ~RMB 0.4 (\$0.056)/kWh, exclusively for facilities running approved domestic chips* [278, 279, 118]. That single instrument welds three elements of this book together: it converts the grid surplus into a subsidy, the subsidy into demand for Ascend/Cambricon silicon, and the silicon into tokens priced below Western marginal cost. It also directly attacks the domestic chips’ principal weakness—30–50% worse energy efficiency than NVIDIA equivalents—though the correct description is not that the wasted watts become free. Extra watts still require generation, transmission, transformers, cooling, water, and backup capacity, and they still carry an opportunity cost; what the tariff does is *socialize part of the operating-cost penalty and accelerate domestic-chip utilization*, transferring the efficiency gap from the buyer’s income statement to the provincial grid’s [279]. Jensen Huang’s London formulation—that for his Chinese competitors “power is free”—is a useful exaggeration of a real subsidy, not a description of the physics [120].

14.2.1 The value of flexible computing

One further asymmetry is under-analysed and probably favours the state-directed system. AI workloads are not uniformly time-bound: interactive serving must run now, but training, batch inference, synthetic-data generation, and much of the scientific execution layer of Chapter 26 can be interrupted, checkpointed, migrated geographically, or scheduled to follow renewable output. The economic value of that flexibility—interruptible training, checkpoint-aware demand response, geographic workload migration, renewable-following batch, and the substitution of *compute* storage (precomputed tokens, caches, distilled models) for *energy* storage—accrues to whoever can coordinate siting, transmission, and scheduling across a national footprint. A network of eight state-designated computing hubs wired to curtailed western generation by dedicated UHV corridors is, in effect, a demand-response instrument at national scale; a set of independently sited private campuses bidding into congested regional markets is not. If this chapter has a claim that deserves more attention than the headline gigawatts, it is this one [S].

Honesty about magnitudes, following the OIES analysis [A] [280]: electricity is ~8% of a US

AI data center’s three-year TCO—hardware dominates—so cheap power alone does not decide token economics; average Chinese spot prices ($\sim \$43\text{--}58/\text{MWh}$) overlap US regional averages; and coastal Chinese rates are unremarkable. The Chinese energy advantage is therefore not the tariff but the *headroom*: the ability to attach tens of gigawatts of new load per year without queues, auctions, moratoria, or political backlash, plus the targeting precision of a state that can price discriminate by chip vendor. The US constraint is not generation economics but delivery time: in the scenario logic of Chapter 21, the binding question is which grid can energize capacity on the schedule the capital commitments assume. On current evidence one side’s constraint is money and the other’s is physics and permitting—and physics-and-permitting is the harder constraint to finance away.

14.3 The busbar casebook: ten regions, one table

The regional evidence named above as this chapter’s largest gap ought to be assembled rather than promised, and this section assembles as much of it as the public record supports. Tables 14.2 and 14.3 apply one standardized frame to five US regions and five Chinese hubs. Grades are per-cell: interconnection statistics and auction prices are [V]; Chinese hub tariffs are [P] (state-media and official reporting); curtailment attributions are [A]. The two systematic gaps, disclosed: no comprehensive US survey of large-load connection times exists (the definitive study says so [371]), and Chinese contracted data-center tariffs below the standard industrial rate are reported for the hubs but not independently auditable.

Three findings follow from the two tables. *First, the two systems fail at opposite ends of the same pipe.* Every US row’s binding constraint is connection time or capacity price—the queue, the auction, the transformer; every Chinese hub row’s binding constraint is durability—whether an administered tariff, built on curtailed renewables and a policy decision, outlives the surplus and the political attention that created it. *Second, the intra-China spread is the real story:* 0.32 RMB/kWh in Hohhot against 0.80 in Shanghai is a $2.5\times$ spread that East-Data-West-Computing exists to arbitrage, and the domestic-chip discount stacks on top of it—the busbar map *is* the industrial policy. *Third, flexibility is worth more than generation on both maps:* the PNW’s 9-GW shortfall collapses to 0–3 GW if data-center load is curtailable, ERCOT now requires a kill switch by law, and CAISO is designing non-firm service—while China’s hubs presume interruptible batch workloads by design. The convergence, from opposite directions, on *curtailable compute* as the grid’s preferred customer is quiet confirmation of Section 14.2.1’s claim that workload flexibility is the under-priced strategic asset of the whole energy element.

14.3.1 Busbar-level deliverability: the unit the argument must be priced in

The analysis has to descend one more level, because only at that level do this chapter’s national aggregates become a procurement-grade quantity. The decision-relevant unit is neither national generation nor even the balancing area but *firm deliverable megawatts at a specific busbar on a specific date at a specific price*. For any candidate computing site, the ledger is five numbers: (1) firm import plus local dispatchable capacity at the busbar, net of existing load, under the binding seasonal constraint (not the annual average); (2) time-to-interconnect for a 100–500 MW block, including studies, upgrades, and transformer lead times; (3) the busbar price—tariff plus congestion plus losses plus firming—for a 24/7 load-following profile, which can differ from the zonal average by a factor of two in both directions; (4) curtailment coincidence, i.e. whether the cheap hours align with interruptible workload (Section above); and (5) the political durability of all four, which moratorium politics in twelve US states and provincial tariff discrimination in China both demonstrate is a live variable, in opposite directions. The two systems’ asymmetry restated at this level: the American constraint binds at (2)—the queue—almost regardless of (1) and (3); the Chinese hub system was *sited* to make (1) and (4) favorable and uses state pricing to

Table 14.2: US regions: what a 100–500 MW load actually faces (as of this edition’s date).

Region	Time to connect	Price signal	Renewables/curtailment	Political durability
N. Virginia / PJM (Dominion)	up to ~7 yr transmission-level hookups; >36-mo substation queue vs 18–24-mo shell build [V] [372]	capacity cleared at/near cap 3 yrs running; \$16.4B for 2027/28 and again 2028/29, both short of reliability requirement [V] [272, 373]	modest curtailment; 2.6 GW offshore wind due 2026	Loudoun ended by-right development; moratorium votes pending; new >25 MW rate class with 85% minimum demand charge from 2027 [V]
ERCOT Texas	/ fastest connect-and-manage reputation, but large-load queue >230 GW (4× in a year), only ~7.5 GW connected; behind-the-meter gas is the workaround [V] [374]	energy-only, scarcity to \$5,000/MWh; EIA projects data-center demand could lift wholesale ~79% by 2027	>8 TWh wind+solar curtailed 2024, transmission-bound West	SB6: study fees, deposits, duplicate-request disclosure, and a mandatory remote-curtailment “kill switch” for new ≥75 MW loads [V]
MISO	load connection via member utilities; tier-2 metros 12–24 mo for moderate loads; ERAS fast-track for adequacy-critical generation	capacity whiplash: \$666.50/MW-day (2025/26) → \$424.30 (2026/27 summer)	wind-heavy north; curtailment not separately reported	cost-allocation fights (\$280M auction adjustment contested at FERC)
SPP	HILL framework (2025): 90-day study track for large loads paired with new generation; provisional and non-firm load service approved	steepest relative growth of any RTO: peak 56 → 105 GW in a decade	wind curtailment up ~6× 2020–24	new framework, untested at scale
CAISO / PNW	utility-owned process; PG&E standardized transmission-level tariff (2025); CAISO exploring non-firm “speed-to-power” service	California ~26¢/kWh all-sector—the cost ceiling of the US map; PNW hydro among the cheapest	CAISO 3.4 TWh curtailed 2024 (93% solar); PNW: 9 GW capacity shortfall by 2030, cut to 0–3 GW if data-center curtailment is allowed	PNW politics of hydro preference; CA affordability backlash

Table 14.3: Chinese hubs: the same frame. Tariffs are contracted data-center rates where reported [P]; standard industrial single-part rates run 0.72–0.81 RMB/kWh across these provinces for comparison.

Hub	Time to connect	Price signal	Renewables/curtailment	Political durability
Inner Mongolia (Ulanqab/ Hohhot)	hub-designated; grid built ahead of load; gigawatt-class campuses proceeding	0.32–0.38 RMB/kWh (\$0.045–0.053) data-center band [P] [377]	>80% green supply claimed; 200+ free-cooling days; <i>but</i> among the worst hidden-curtailment provinces (CHP inflexibility) [A] [375]	tariff is policy, not market; Doc-136 early adopter
Gansu (Qingyang)	fastest-growing hub: 168k PFLOPS by mid-2026; dedicated 2 GW green plant built for the hub	0.398 RMB/kWh (\$0.056), stated as stabilized; 35–45% below eastern cities [P] [377]	>30% variable renewables, chronic export dependence; 8 GW UHVDC to Shandong (2025)	subsidized rate outliving the surplus is the named risk (§14.3.1)
Guizhou (Gui'an)	dedicated substations added within quarters; DC electricity use +517% y/y Q1 2025 [P] [377]	subsidized interior pricing (historically ~0.35 RMB/kWh deals) [A] [376]	35% renewables + large hydro; mild climate	30% AI-adaptation vouchers; hyperscaler-anchored (Huawei, Tencent, Apple)
Guangdong (Shaoguan)	GBA hub node; 500k-rack target	coastal 0.60–0.75 RMB/kWh (\$0.08–0.11)—above the Virginia average [A] [376]	~70% coal-sourced eastern power; latency keeps the load anyway	Doc-136 finalized late
Yangtze Delta / Shanghai	incumbent digital-economy center; East-Data-West-Computing routes batch load out	0.74–0.80 RMB/kWh industrial; the price the western hubs are arbitraging	~70% coal; PUE ≤ 1.25 mandate binds hardest here	Jiangsu among last provinces without Doc-136 rules—renewable-pricing uncertainty at the demand center

set (3), so its binding constraint is (5)’s inverse, the risk that subsidized tariffs outlive the surplus that justified them. The busbar price and time-to-interconnect are precisely the “energy” row of the rent-migration table (Table 21.4); a reader who wants this chapter’s thesis as a tradeable indicator should watch those two series, hub by hub, and this book commits to carrying them in the living appendix as regional data becomes assembled.

14.4 Does the data centre have to be near you? The locality question, examined

An objection sits underneath every argument in this chapter about siting compute where the power is, and it deserves to be settled rather than assumed. The objection holds that cheap remote electricity is useless for *inference*, because inference must run near its users: training tolerates distance, serving does not. In the Chinese case the claim is specific—that the western hubs of Inner Mongolia, Gansu, Guizhou, Ningxia and Xinjiang cannot serve the eastern demand centres, and that Beijing’s own 20 ms latency target goes unmet there [537, 541]. If true, it would mean that the country’s largest apparent advantage, surplus power in the west, cannot reach the workload that pays. The claim is substantially wrong, and the way in which it is wrong is instructive.

14.4.1 The latency budget, decomposed

Start with arithmetic rather than assertion. End-to-end serving latency is time-to-first-token plus the remaining output tokens divided by the per-token rate, and TTFT itself decomposes into network transit, gateway and TLS, queueing, tokenization, and prefill computation. A published production budget for a 400 ms P99 TTFT allocates *20 ms to network ingress and 280 ms to prefill*—network is 5% of the budget, prefill is 70% [547] [A]. Measured TTFT for frontier models in 2026 runs from roughly 0.6–1.1 s for standard chat models to 8.4 s for a reasoning model at medium effort and tens of seconds with extended thinking; prefill scales quadratically with context, so agentic coding at 30–200K tokens sits in seconds, not milliseconds [548] [A].

Against those magnitudes, put the distances. A Tokyo-to-Oregon round trip measures about 99 ms; Tokyo to Northern Virginia about 149 ms [549]. Fibre propagates at roughly 4.9 μ s per kilometre, so Shanghai to Guiyang—about 2,150 km by realistic route—has a theoretical floor near 21 ms round trip, and the published Chinese figures are consistent: Gui’an to the Yangtze Delta and to Beijing-Tianjin-Hebei within 18 ms *one-way*, Zhongwei to Shanghai within 15 ms one-way, Inner Mongolia to Beijing-Tianjin-Hebei within 5 ms [544] [P/A, Chinese official and state-adjacent sources; no independent Western measurement of intra-China backbone latency was located].

The ratio is the argument. For a standard frontier model at roughly 1.1 s TTFT and 92 output tokens per second, a 500-token response takes about 6.5 s. Adding 36 ms of Shanghai–Guizhou round trip costs 3.2% of TTFT and *0.55% of the response*. Adding 99 ms of trans-Pacific round trip costs 8.8% of TTFT and 1.5% of the response. Against a reasoning model’s 8.4 s TTFT, the trans-Pacific penalty is 1.2%. *Moving the accelerator eight thousand kilometres costs one to two percent of wall clock; moving it fifteen hundred costs half a percent* [M].

14.4.2 The revealed preference of the industry

The decisive evidence is not arithmetic but practice, and it is available in vendor documentation. OpenAI distinguishes data residency from *inference* residency and offers the latter in three regions only—the United States, the EEA and Switzerland, and the United Arab Emirates. Japan, Korea, Singapore, Australia, India, Canada and the United Kingdom receive storage residency alone, which means the GPU execution for users in those markets happens elsewhere [550]

[V]. Anthropic’s documented default inference geography is **global**: “inference may run in any available geography for optimal performance and availability,” with only **us** and **global** available as settings [551] [V]. Microsoft tells customers that for most workloads they should start with Global Standard, which “dynamically route[s] traffic to available datacenters” and “has the lowest price”; Google warns that its global endpoint “does not guarantee that requests will be processed in any specific location” and directs customers to regional endpoints for *residency*, not for speed; AWS routes Bedrock traffic originating in Singapore, Jakarta, Taipei, Thailand and Malaysia across more than twenty regions worldwide [552, 553, 554] [V].

Read that footprint carefully: it maps legal jurisdictions, not population centres or latency contours. Japan is one of the largest AI markets on earth and one of the most distant from US compute, and the two leading providers serve it across the Pacific. If latency governed siting, that map would look entirely different. The industry’s own pricing tells the same story—the cheapest option is the location-agnostic one.

There is a result that sharpens this from surprising to counter-intuitive. A 2026 study of cross-region routing measured a case in which dispatching a request from US West to Germany—281 ms away—produced a *lower* total time to first token than serving it three milliseconds away, 296 ms against 305 ms, because queueing delay at the congested near site exceeded intercontinental propagation [555] [A]. Under load, distance can be negative cost.

14.4.3 What the Chinese policy document actually says

The 20 ms target is real and is routinely misread in three ways at once. It originates in the NDRC’s May 2021 implementation plan for the national integrated computing hubs. That text specifies end-to-end network latency “in principle within 20 milliseconds” *one-way*—the round-trip equivalent is 40 ms—and specifies it for the *eastern* clusters, as a standard for how close an eastern hub must sit to the users it serves. The western nodes are assigned, in the same document, to “back-end processing, offline analysis, and storage backup”—explicitly non-real-time work. A separate and tighter 10 ms tier is reserved for city-edge facilities serving high-frequency trading, VR/AR and connected vehicles [542] [P, primary text].

Three corrections follow, and each matters. First, the English-language rendering of this target as a round-trip budget halves the available headroom for no reason. Second, the frequent rendering of it as a requirement running “from western facilities to eastern population hubs” is not what the document says. Third, and most consequentially, the plan predates ChatGPT by eighteen months: “AI inference” in a May 2021 Chinese policy document means computer vision, recommendation and speech models whose entire forward pass takes 5–50 ms, and for *those* workloads 20 ms of network genuinely is material. Applying a threshold calibrated to a 20 ms forward pass to a model with a 1.1-second TTFT is a category error.

Whether the target is met is also disputed in one direction only. The head of China’s National Data Administration stated in August 2024 that latency between the eastern and western hubs “has generally met the 20-millisecond requirement” [543] [P, official statement]. The measured per-hub figures above are consistent with that. The claim that the target is unmet appears to enter English discussion through secondary reporting rather than from either the primary document or the measurements, and this book grades it unsupported.

14.4.4 The decisive counter-example, three days old

On 12 August 2026 Alibaba Cloud placed its Lingjun M890 supernode into commercial service with first deployment in *Ulanqab, Inner Mongolia*—one of the exact provinces the locality objection names—designed for *inference* on mixture-of-experts models up to ten trillion parameters, and already serving live commercial traffic for Kimi K3 and Qwen3.8-Max [545] [V, two independent outlets]. Its API endpoints, meanwhile, are published in Beijing, Hong Kong, Singapore, Frankfurt, Tokyo and Northern Virginia [546]. Front door east and worldwide; accelerators west. That is

the architecture the objection says is impossible, in production, at the country’s largest cloud provider.

The architectural trend runs the same way. Prefill and decode are being disaggregated, and the seam is migrating outward from the node to the rack to the cluster and now to the wide area: a 2026 paper from Moonshot and Tsinghua proposes prefill-as-a-service with key–value cache shipped *cross-datacentre* over commodity Ethernet, requiring roughly 170 Gbps for a trillion-parameter model at 32K context and less at longer context—a routine dedicated wavelength against Zhongwei’s stated 18 Tbps of export bandwidth [556] [A]. Another finds that shipping a one-gigabyte KV cache over a 10 Gbps link costs about 0.8 s and nine cents against roughly 15 s and forty-seven cents to recompute it locally [557] [A]. The seam that splits cleanly is the one crossed *once per request*; splitting the model layer-wise, which requires a round trip *per token*, does not survive wide-area latency at all. The industry is building the former.

14.4.5 Where locality genuinely binds—stated at full strength

Three cases are real, and conceding them makes the rest harder to dismiss.

Inline code completion is dispositive. Published service-level objectives put time-to-first-token for completion at 100–125 ms [558, 547]. A 99 ms trans-Pacific round trip consumes nearly the entire budget; even 36 ms of Shanghai–Guizhou consumes a third of it before any compute occurs. This workload must be served locally, and no architecture rescues it.

Conversational voice binds nearly as hard. A speech-to-speech agent runs a total budget near 500 ms, of which 30–80 ms is allocated to network; measured production voice-to-voice latency sits at 680–850 ms median. Adding 99 ms is a fifth of the budget and breaks specification [559] [A]. Voice is among the fastest-growing enterprise categories, which makes this concession expensive rather than cheap.

And the better inference gets, the more geography matters. Prompt caching reduces latency for long prompts by up to 85% on the vendor’s own claim; a cached 2 s TTFT becomes roughly 300 ms, at which point a 99 ms round trip is a third of the budget rather than a twentieth [560]. Specialist inference hardware already delivers 160–180 ms TTFT, where trans-Pacific transit would exceed half the total—which is precisely why those vendors have built Helsinki, Sydney and UK sites. Every improvement in prefill efficiency raises the relative weight of the network term. The locality objection is not wrong forever; it is wrong *now*, and becomes right as the compute term shrinks.

14.4.6 The claim that survives, and the one that should replace it

Two distinctions rescue what is valid in the objection.

Proximity to the user is weak; proximity to the corpus is strong. The serious version of the “hot data” argument, as its industry proponents actually make it, is about the retrieval corpus rather than the human: an inference platform should be tightly coupled to the data platform, because agentic retrieval makes many round trips per turn [561]. A Japanese user querying a US-hosted model over a US-hosted index has no data-gravity problem at all—only the human is far away, and the human is the cheapest component to keep waiting. A Shanghai enterprise querying its own Shanghai-resident warehouse from a Guizhou accelerator pays transit on every hop. That is a real constraint, it is not a latency-to-user constraint, and the literature conflates the two.

And the western hubs’ real problem is the workload they were built for, not the place they were built in. The most careful reporting on Chinese overbuild found that the idle facilities “were optimized for pretraining workloads rather than for inference” [537]—the wrong interconnect, the wrong storage tiering, the wrong accelerator mix. Chinese industry analysis states the resulting pattern as *training migrates west, inference sinks east*, and states its cause as demand: the

eastern regions are where inference demand is exploding [562] [P]. That is a claim about where the customers are, not about where photons can go.

So the corrected statement, which this book adopts: *Chinese western capacity is underused because it was built for a pretraining boom that ended and sits far from the paying customers, not because inference cannot physically be served from it.* The first is an economic problem, and economic problems yield to price. The second would have been a physical bound, and it is not there. For the argument of Chapter 21, the difference is decisive: a country with surplus power, surplus capacity, and no physical barrier between the two has a supply curve it has not yet monetized—which is exactly the position from which the solar industry was rebuilt.

14.5 Where the watt meets the architecture

The energy element feeds back into the technical argument of Chapter 12 in three ways. First, cheap-and-abundant power changes the efficiency frontier: a CloudMatrix-class system at $3.9\times$ the power of its NVIDIA counterpart is a bad trade in Virginia and an acceptable one in Guizhou—brute-force parity is a rational strategy precisely where watts are the cheap input [137]. Second, the capacity-first LPDDR architecture’s MFU penalty (67% vs 89%) is, among other things, an energy trade a surplus grid can afford [M]. Third, the decentralization thesis inverts in energy terms—and the point has to be stated consistently with this book’s own power ledger (Table 12.4), on pain of contradiction: a 128-socket (64-node) sovereign pretraining cluster is modeled at roughly *half a megawatt*, and even larger replicated or multi-rack installations remain in the low-megawatt range—an ordinary industrial feeder, deployable in jurisdictions that could never host a gigawatt campus, which is most of the world. The Global South adoption channel of Chapter 10 is, in energy terms, a market for machines sized to real grids.

The IEA’s trajectory—global data-center consumption from 415 TWh (2024) toward ~ 945 TWh (2030), with the US and China as the two poles [281]—frames the endgame: the AI war’s operating cost curve is converging on the electricity cost curve. Part I documented who won the technologies that produce electricity. Solar modules, batteries, UHV transmission, and the coming terawatt of storage are the same Chinese industrial commons that Chapters 3–5 described; the watt is where the precedent wars and the AI war become one war. It is the fourth element because it was the first three’s prize.

Chapter 15

Data and Feedback: The Fifth Element

Four elements so far: the model, the compute, the semiconductor, the watt. The rent-migration logic of the preceding chapters demands a fifth, and this chapter supplies it—closing a gap the argument would otherwise step around rather than confront. If open weights commoditize the model layer, and value migrates to adjacent scarcities (Section 21.8), then the most durable adjacent scarcity is the one input that openness cannot replicate by construction: *data that only deployment generates*. This chapter maps what open weights do and do not commoditize on the data axis, who holds which class of scarce data, and the open questions that will decide whether feedback—not weights, not FLOPs—turns out to be the true scarce layer of the mature industry.

15.1 What open weights do not commoditize

An open-weight release transfers a compressed function of its training distribution. It does not transfer, and cannot: (1) *proprietary enterprise workflows*—the tickets, codebases, contracts, and process traces inside firms, reachable only through deployment relationships; (2) *scientific-instrument data*—beamline, sequencer, telescope, and sensor streams, gated by institutions and provenance; (3) *clinical and laboratory data*, gated by regulation and consent; (4) *high-quality human correction and preference data*, produced only where skilled users interact with deployed systems at scale; (5) *embodied physical-interaction data*, produced only by deployed fleets (Chapter 16); (6) *evaluation and failure traces*—what actually broke, in production, under adversarial pressure; and (7) *operational telemetry* from large deployed systems. Each is generated by *use*, appropriated by whoever owns the deployment relationship, and—critically for this book—each remains scarce *after* the model becomes free. Table 15.1 assigns the current holders.

Read as a campaign map, the table explains several things the four-element story left dangling. It explains why the American flywheel’s rent-recapture plan (Table 2.2) can survive open weights: the enterprise-workflow row is the one Chinese open models cannot reach from outside the deployment relationship, and it is the row Anthropic’s \$47B enterprise run rate is built on. It explains why China’s embodied push (Chapter 16) is strategically rational even if humanoids disappoint commercially: the industrial-data row is the only row where China holds a structural generation advantage, and the state data-farm program is an attempt to *manufacture* a data moat the way Part I manufactured cost moats. And it explains the last row’s importance one more time: evaluation and incident data—who failed, how, under what attack—is the raw material of the certification authority this book keeps finding unowned.

Table 15.1: Data classes after model commoditization: scarcity mechanism and current holder of advantage [A/S].

Data class	Scarcity mechanism	Likely advantage
Public pretraining corpora	increasingly replicated; approaching exhaustion	the commons—no one’s rent
Synthetic reasoning data	teacher capability and filtering quality	frontier laboratories (both blocs)
Enterprise workflow data	access and integration; deployment relationships	US cloud and application firms
Scientific and instrument data	provenance; public institutions; consent	national laboratories and consortia
Industrial and embodied data	manufacturing deployment surface	China (Chapter 16)
Human preference and correction data	skilled-user deployment at scale	distribution owners (currently US consumer, mixed enterprise)
Evaluation and incident data	certification and disclosure regimes	<i>currently unowned</i> —the same vacancy as Chapter 18’s certification layer

15.2 The synthetic-data question

The sharpest challenge to any data-moat argument is synthetic data: if strong models can generate their own training material, deployment-derived data loses scarcity value. The evidence cuts both ways, and the honest statement is conditional. Synthetic reasoning data demonstrably works—the distillation economy of Chapter 10 is built on it, post-training pipelines are majority-synthetic at every major lab, and the alleged industrial-scale harvesting of frontier outputs (Chapter 18) is precisely a fight over synthetic-data *feedstock*. But synthetic data has a teacher: its quality is bounded by the capability of the model that generates it plus the selectivity of the filter that curates it, which makes it a mechanism for *diffusing* frontier capability rather than exceeding it—distillation relocates the frontier’s rent, it does not abolish the frontier. And recursive training on model outputs degrades distributional diversity and factual grounding unless anchored by fresh non-synthetic signal; the model-collapse literature overstates the cliff, but the direction is real. The synthesis this book adopts [S]: synthetic data *amplifies* whoever already holds either frontier capability (the teacher) or verified ground truth (the anchor)—it does not substitute for the anchor. Which returns the argument to the table above, because the anchors are exactly the deployment-generated classes: verified task outcomes, experimental results, physical interaction, production failures. *Verification, not generation, is the scarce half of the synthetic-data economy*—a conclusion that should look familiar, because it is the funnel of Chapter 24 (v_j before s_j) and the trust stack of Chapter 18 restated as a data-supply question.

15.3 Seven questions that decide the element

The element poses seven open questions, and this book puts each in the sharpest form it can before answering it to the extent the evidence permits, registering the rest.

Who owns data generated by deployed agents? Legally unsettled everywhere, and the default is being set by contract while legislatures sleep: enterprise agreements increasingly reserve workflow traces to the customer while permitting aggregate learning to the provider—which, at scale, gives the provider the moat and the customer a receipt. The jurisdiction that defines agent-generated data rights first will set the default for everyone, and none has.

Can open models be improved with closed workflow traces without leaking proprietary knowledge? Technically yes—LoRA-class adapters, federated post-training, and gradient aggregation exist—but the assurance problem (proving the adapter does not memorize the traces) is unsolved at certification grade, which makes this another product the trust layer of Chapter 18 must ship before regulated industries participate.

How should data withdrawal propagate through derivative chains? Today: not at all. A revoked dataset persists in every model already trained on it and every derivative thereafter—the same unowned-lineage problem as the security chapter’s artifact chain, and solvable by the same instrument: signed manifests and data bills-of-materials in the certification ladder.

Can scientific consortia create a data commons without making sensitive science globally downloadable? This is the consortium design problem of Chapter 27, and the layered-access architecture there (open weights, tiered data, sovereign enclaves for dual-use domains) is this book’s answer; the fifth element supplies its economic justification, because pooled instrument data is the one asset that makes a multi-state consortium more than the sum of its national programs.

Does industrial deployment create a physical-data advantage analogous to the US enterprise-software data advantage? Chapter 16’s answer: structurally yes—same mechanism, different substrate—but only if the task-policy layer can consume it, which is the contested middle of the embodied stack.

Does synthetic data reduce or amplify frontier dependence, and when does recursion degrade? Answered above: it amplifies whoever holds the teacher or the anchor; recursion degrades absent fresh verified signal.

Is verified feedback the true scarce layer after weights commoditize? The book’s answer, now stated as its position [S]: yes, with one refinement—the scarce object is not feedback in bulk but *feedback attached to verified outcomes*, because unverified feedback is synthetic data’s cheap cousin and the world is about to drown in it. Every prior element of this book converges on that refinement: the funnel’s validation term, the trust stack’s behavioral layer, the science metric’s replication vector, the embodied proof suite’s intervention-free hours. The fifth element is, in the end, the same element measured everywhere else: *ground truth, at scale, with provenance*. That is what cannot be downloaded.

The dashboard rows follow directly: enterprise-agreement terms on agent-trace ownership (watch the default settle); first certification-grade privacy-preserving fine-tuning product; first data-BoM requirement in any jurisdiction’s AI procurement; and the ratio of verified-outcome data to raw interaction data in published post-training recipes—the number that will reveal, before any market does, whether the fifth element is consolidating into a moat and whose.

Chapter 16

Embodied AI: The Loop Closes on the Factory Floor

This chapter exists because its absence would be the most consequential omission the book could make. The thesis so far runs in one direction: China’s manufacturing victories funded and shaped its digital-AI strategy. Embodied AI runs the loop the other way—model capability flowing back into factories, logistics, vehicles, laboratories, and physical services—and it is the domain where China’s three structural advantages (manufacturing scale, deployment density, component supply chains) interact most directly with its open-model strategy. If the convergence thesis is correct anywhere with mechanical inevitability, it should be here, where the contested layer is not a benchmark but a bill of materials. And yet the chapter’s conclusion is deliberately not a Chinese victory lap: the evidence as of this edition’s date supports Chinese dominance of the *hardware and deployment* layers and genuine uncertainty about whether hardware and deployment are the binding constraint at all. The question that decides the domain—is model capability or physical-system integration the harder bottleneck?—is precisely the question the rest of this book asks about digital AI, transposed into metal.

16.1 The embodied stack, separated

“Humanoid robot” names a product, not an architecture, and the strategic analysis requires the stack:

foundation model → world model → task policy → motion planning
→ real-time control → actuators and sensors,

with a deterministic safety layer beside the whole chain and a fleet-learning loop above it. The layers have different physics, different competitive structures, and different nationalities. Language and vision reasoning (the foundation layer) is the commodity this book has spent four chapters on—and the embodied field increasingly consumes it in exactly the open form Part II describes. World models—learned predictors of physical dynamics used for planning and for generating synthetic training experience—are a frontier-lab specialty (NVIDIA’s Cosmos family, DeepMind’s Genie line [424, 425]). Task policies are where the field’s distinctive artifact lives: vision-language-action (VLA) models such as Physical Intelligence’s π -class systems and Figure’s Helix, typically a slow deliberative model (7–9 Hz) coupled to a fast visuomotor policy (200 Hz) [426, 427]. Motion planning and real-time control are classical robotics, sixty years deep, dominated by Japanese and European engineering practice. Actuators and sensors are manufactured objects—which is to say, after Part I, presumptively Chinese objects, a presumption the bill of materials below confirms.

The stack decomposition does the same analytical work the die-area argument did in Chapter 12: it converts a category question (“who is winning humanoids?”) into a layer question, and

the layer answers differ. The United States leads the top two layers. China leads the bottom two and the deployment environment. The middle—the task-policy layer where model meets metal—is genuinely contested, and it is the layer where the physical-data flywheel of the next section decides the outcome.

16.2 The physical-data flywheel, and who is building it

Model the loop explicitly, because it is the third flywheel of this book and the one with the strangest property:

deployment → real interaction data → failure cases → post-training
→ better policy → more deployment.

Web-scale text pretraining had the property that its input—humanity’s accumulated writing—existed before the models did, and was free. Physical-interaction data has the opposite property: it does not exist until robots are deployed, it is generated in proportion to deployed fleet size and operating hours, and it cannot be scraped, only *produced*. That single property is why the embodied contest may be structurally different from the LLM contest: the scarce input is not compute or algorithms but *deployment surface*, and deployment surface is a function of manufacturing cost, component supply, factory density, and state coordination—the exact inventory of Chinese advantages from Part I.

The Chinese state understood this early and built the loop’s infrastructure as public works. By late 2025 more than forty state-backed robot data-collection centers had been announced (roughly two dozen operational), staffed by human teleoperators in VR headsets and exoskeletons demonstrating assembly, household, and elder-care tasks for hours a day—industrialized imitation-learning data production [V] [428]. Shanghai’s heterogeneous training facility opened in January 2025 with capacity for a hundred robots simultaneously and a target of a thousand by 2027 collecting ten million entries [429]; Beijing’s Shijingshan camp runs sixteen scenario sets from car-assembly lines to smart homes; provincial centers have become a demand channel in their own right, with UBTECH booking RMB 566M of robot sales to three of them and China Mobile placing nine-figure orders with Unitree and AgiBot [428]. On the open side, AgiBot’s World Colosseum dataset publishes a million-plus real-robot trajectories from a hundred robots across a hundred scenarios—the embodied analogue of the open-weight release, and aimed at the same strategic effect [430]. The Western equivalents are real but differently shaped: Open X-Embodiment pools academic lab data across 33 institutions [431]; Figure trained Helix on roughly five hundred hours of curated teleoperation [426]; Agility has accumulated 65,000+ hours of genuinely productive deployment [432]; Tesla possesses, in principle, the largest fleet-learning apparatus on earth and has yet to ship a production humanoid to use it on [433].

The honest current reading: China is industrializing the *production* of physical-interaction data; the United States leads the *consumption* of it (the VLA models that turn it into capability). Neither half works without the other, which is why this flywheel is the one whose ownership is least settled—and why it cannot be rented: a laboratory without manufacturing access cannot buy deployment surface the way it can buy tokens, and a manufacturer without frontier post-training cannot convert demonstrations into generalization. The strategic symmetry with Chapter 2’s two thefts is exact, and Section 16.6 prices it.

16.3 The industrial bill of materials

The component layer is where Part I’s precedents repeat with the least resistance, because it *is* Part I’s industries wearing new part numbers. Table 16.1 scores the major component classes.

Two cost-curve observations convert the table into strategy. First, the slope: Goldman Sachs found humanoid unit manufacturing cost falling from \$50–250k to \$30–150k in a single year—a

Table 16.1: The embodied bill of materials, mid-2026 [A unless noted]. “Position” names the current center of gravity, not a monopoly.

Component	Position	Evidence and trajectory
Strain-wave (harmonic) reducers	Japan incumbent (Harmonic Drive); China scaling fast	Leaderdrive is China’s largest maker, lifted by humanoid demand; Zhongda Leader’s founders became billionaires on joint demand; Unitree’s price war is already squeezing the domestic reducer chain [434]
Planetary roller screws	China dominant in new capacity	~33% of humanoid BoM by one bank estimate; unit cost \$1,350–2,700; Shanghai Beite’s \$260M dedicated plant; thin Western capacity [435]
Actuator assemblies, servo & frameless motors	China (three parallel chains: Tesla-, Huawei-, Unitree-anchored)	reported \$685M Optimus actuator order to Sanhua; Tuopu building dedicated drive-system capacity [436]
Lidar and depth sensing	China decisively	Hesai robotics shipments +744% y/y (49k units, Q2 2025); RoboSense +632%; both pivoting from automotive [437]
Force/torque, tactile, dexterous hands	Contested; APAC-centered manufacturing	dozens of Chinese hand makers; Japanese and German precision sensing incumbents
Batteries	China (Chapter 4)	CATL both supplies packs and deploys humanoids on its own pack lines—the loop in one firm [438]
Edge compute	US platform (Jetson Thor, \$3,499, 2,070 FP4 TFLOPS); domestic chain behind	Huawei/Cambricon alternatives exist inside the Chinese chains [439]
Safety systems & certification	Europe and Japan	ISO 10218 (2025 ed.), ISO 25785-1 for dynamically stable mobile robots drafting toward ~2027 [440]
Simulation & digital twins	US (Isaac/Cosmos); China open-sourcing fast (Genie Sim)	[424, 430]

~40% decline where 15–20% was forecast—and Morgan Stanley finds localized Chinese supply chains running ~20% cheaper than foreign ones, with average system prices trending from ~\$46k toward ~\$21k [A] [441, 442]. Unitree’s own ASP fell from RMB 593k (2023) to RMB 166k (2025), with an entry model at RMB 29,900—the solar learning curve, replayed on legs, with the familiar accompaniment of collapsing gross margins (87.7% → 63.2%) that Part I teaches us to read as the cull beginning [443]. Second, the dependency: supply-chain reporting places on the order of 70% of Tesla Optimus’s supply chain in China, with estimates that excluding Chinese parts would triple its cost [A]—and China demonstrated the leverage directly when rare-earth export licensing touched Optimus’s actuator magnets in 2025 [444]. The American humanoid, like the American solar panel of 2010, is being designed atop the adversary’s supply chain. This is playbook move M5 (supply-chain capture) achieved *before* the product category even exists at scale—a first even for China.

16.4 The deployment scoreboard, and what it does not show

The shipment table reads like a rout. Global humanoid shipments exceeded 13,000 units in 2025, roughly 85% from Chinese firms [A] [445]. AgiBot shipped its 10,000th cumulative unit in March 2026—its second five thousand in roughly three months [446]. Unitree sold 5,500 humanoids in 2025 and, almost uniquely in the industry worldwide, did so *profitably* (~\$250M revenue, ~\$41M profit), reaching its STAR-market IPO approval in June 2026 [443]. UBTech put its thousandth Walker S2 off the line with cumulative orders past RMB 800M and multi-robot deployments at Zeekr, BYD, Foxconn, and Audi-FAW plants [447]. CATL began scale deployment of Galbot systems on its own battery lines [438]; BYD’s four-year internal program targets ~20,000 robots in its own factories in 2026 [P] [448]; XPeng targets mass production of its Iron humanoid by end-2026 and a million units by 2030 [R] [449]. Around them, the ecosystem markers of every Part I campaign: 150-plus companies, embodied AI named in the 2025 Government Work Report and the 15th Five-Year Plan, an \$8.2B national fund allocation, Beijing’s \$14.3B fifteen-year robotics fund, provincial subsidy programs, and municipal action plans with unit and revenue targets [450, 451].

The American column is shorter and stranger. Tesla—the company whose CEO calls Optimus “the biggest product ever”—had zero verified production units as of July 2026, deferred its Gen-3 reveal twice, and told its own earnings call that existing units are “primarily for learning, not productive tasks” [433]. Figure, valued at \$39B, has delivered 350-plus units of its third-generation robot and produced the single most important capability datapoint in the Western column: a livestreamed eight-hour *autonomous* logistics shift at human-comparable package-sorting rates [426, 452]. Agility’s Digit has 65,000-plus hours of real deployed operation and a \$300M+ multi-year order book, going public via SPAC at \$2.5B [432]. Boston Dynamics’ electric Atlas has its entire 2026 production run committed and a Hyundai plant designed for 30,000 units a year—gated, instructively, by a Korean labor union blocking deployment absent an agreement [453]. Behind them stand the model shops: Physical Intelligence at a reported ~\$5.6B raising toward \$11B [A], Skild at a reported ~\$14B [A], and NVIDIA supplying the compute, the world models, and an open reference humanoid to the entire field [427, 439].

Now the deflationary reading, which this book owes the reader after four chapters of proof-boundary discipline. *Shipments are not labor*. The majority of Chinese humanoid units go to demonstrations, exhibitions, education, and—in a perfect circularity—to the state-funded data-collection centers whose purchases are themselves industrial policy; independent assessments describe most deployed units as “performative rather than functional,” fragile in operation and dependent on structured environments [445]. Teleoperation shadows the entire field’s marketing, on both sides of the Pacific: Tesla’s bartending robots were remotely operated, the 1X NEO home robot ships with human teleoperators explicitly in the loop, and the viral demonstrations that anchor valuations routinely turn out to have a human in them [454]. The comparison that matters

is therefore not units shipped but *intervention-free productive hours*, and on that metric the public record inverts the shipment table: the best-documented sustained autonomous performances (Figure’s eight-hour shift; Agility’s cumulative deployment record) are American [426, 432]. The scoreboard, graded honestly: China leads manufacturing scale, cost, components, and deployment surface [V]; the United States leads demonstrated autonomy per deployed unit [V]; and the metric that would settle the contest—useful work per robot-hour at fleet scale—is published by nobody.

16.5 Metrics that matter: the embodied proof suite

As with the LX2 in Chapter 12, the remedy for a contested frontier is a stated measurement program. The quantities that would decide the embodied contest, none of which any major player currently publishes systematically: intervention-free operating hours; task-success rate on standardized task sets; cycle time against human baseline; uptime; mean time between safety stops; human-supervision minutes per robot-hour (the teleoperation ratio, the single most information-dense number in the field); cost per successful physical task; energy per task; damage and near-miss rates; skill-transfer performance across sites (train in one factory, deploy in another—the sim-to-real and site-to-site generalization that separates a product from a demo); and useful training data generated per deployed unit, which is the flywheel’s own gauge. A vendor that publishes this table for a hundred-robot deployment over six months will do for embodied AI what MLPerf did for accelerators; until one does, every shipment figure in the previous section should be read as capacity, not output. The forecast register (Appendix J) now carries the corresponding entries.

16.6 Comparative national positions

The five-way comparison, stated without a predetermined winner. **China** holds manufacturing scale, the component base, deployment density (54% of world industrial-robot installations, 57% domestic-brand share of its own market, the largest operational stock ever recorded [455]), state-industrialized data collection, and the cost curve; its exposure is the familiar pair—capability of the top model layers, and a domestic price war already compressing margins before the product works. **The United States** holds the frontier models, the world models, the VLA research lead, the capital (three embodied startups above \$10B), and the strongest demonstrated autonomy; its exposure is a supply chain it does not own and a deployment environment of high labor cost but low robot density. **Japan** holds the precision-component incumbency (strain-wave gearing above all), the industrial-robot giants, production engineering as a craft, and—via Chapter 23’s logic—a natural niche as the trusted supplier of the components both blocs need; its exposure is having invented the category (ASIMO) and ceded the product. **Korea** holds the Hyundai–Boston Dynamics vertical (a 30k-unit plant attached to a captive automaker) plus its own robotics base. **Europe** holds the safety-certification tradition and industrial-integration expertise, and—true to form—is currently converting them into standards rather than products.

The decisive uncertainty is which bottleneck binds. If *model capability* is the constraint—if generalist manipulation at production reliability is still several breakthroughs away—then the American position is stronger than the shipment tables suggest, Chinese fleets amortize into demos, and the data farms are collecting demonstrations for policies that cannot yet use them. If *physical-system integration* is the constraint—if current VLAs are good enough that reliability, cost, and iteration speed at fleet scale decide—then China’s position compounds exactly as it did in EVs, and the American labs are optimizing the layer that is about to commoditize. The evidence is genuinely mixed: the eight-hour autonomous shift argues capability has arrived for narrow logistics; the near-universal teleoperation dependence argues it has not for anything else. This book’s judgment [S], registered as a forecast: the integration reading wins in

structured environments (factories, warehouses) within the decade—China’s home terrain—while unstructured environments (homes, care, field work) remain capability-bound well past 2030, leaving the largest advertised markets unclaimed by anyone.

16.7 Falsifiers

Consistent with the book’s method, the chapter states what would prove its cautious-Chinese-advantage reading wrong in either direction. *Against the embodied thesis entirely*: teleoperation ratios that do not decline across 2026–2029; intervention-free hours that stagnate at demo scale; simulated training that persistently fails to transfer; safety certification (ISO 25785-class) blocking broad deployment the way certification blocked the C919; unit economics that stay worse than human labor in every unstructured setting; hardware reliability, not intelligence, remaining the binding constraint—humanoids as the Concorde of automation. *Against the Chinese-advantage reading specifically*: a Western lab achieving generalist manipulation that renders demonstration-data volume irrelevant (the capability escape); US/allied reshoring of the actuator and reducer chains at competitive cost; the Chinese price war culling the sector before any firm reaches sustainable unit economics—involution consuming the industry in its infancy, the failure mode Section 23.6 warns Beijing about directly; or the state data-farm model producing large volumes of low-diversity demonstrations that generalist policies turn out not to need. Each has an observable signature, and the quarterly dashboard now watches for all of them.

The chapter’s closing symmetry deserves a sentence, because it completes the book’s architecture. Part I showed physical manufacturing funding China’s route to digital AI. This chapter shows digital AI attempting to conquer physical manufacturing. If both directions close, the two flywheels of Chapter 2 fuse into a single self-reinforcing industrial system—AI improving the factories that build the machines that generate the data that improves the AI—and the country that owns that closed loop owns something qualitatively beyond anything Part I described. That, and not any single robot, is the stake on the factory floor.

Chapter 17

The Mirror: The West Runs the Playbook Against Itself

The chapters of Part II examined the elements of AI one at a time and found, in each, a recognizable move from the playbook of Chapter 7 being executed by a Chinese actor. This chapter reports the finding that changes what those elements mean. Over the course of 2026 the largest Western firms in the sector converged on the same method—not in imitation, and in most cases without any reference to China at all, but by the same economic logic arriving at the same instrument. Each of them now commoditizes, deliberately and at its own near-term expense, a layer it does not own, in order to defend the layer it does. The accelerator vendor gives away models. The model laboratories design accelerators. The hyperscalers build the fabric out of merchant parts and publish the specifications. A launch company is building a fab. The strategy the preceding chapters attribute to Beijing is no longer Chinese. What remains asymmetric is not the method but the scope—and, as Section 17.5 argues, the asymmetry of scope is worth more than the symmetry of method takes away.

Two things follow immediately, and they point in opposite directions. The first is that Proposition P3—the destruction of the rent that finances the Western frontier—acquires a second and entirely independent mechanism, one that does not require Chinese price pressure to reach Western margins at all, because a layer can be commoditized by the firm sitting above it and every layer in this stack has a firm sitting above it. The second is that a West which commoditizes its own layers competently captures the surplus of a cheap-intelligence world as surely as a China that does, and the book must not confuse falling prices with a change of ownership. Both readings are printed here, and Section 17.6 states which claims the chapter licenses and which it forbids.

17.1 The move, stated generally

The theory chapter established the result this chapter depends on: openness relocates rent, it does not abolish it (Section 2.3). A firm that gives away a good it produces is not destroying value; it is moving the point at which value is captured to a layer it holds more securely. That result was stated there as an account of Chinese strategy, but nothing in it is national. It is a statement about complements. A firm that owns layer *A* and faces a rent-taker at the adjacent layer *B* has an option: drive the price of *B* to marginal cost, and the surplus that *B* was extracting flows to *A* instead. The move costs the firm nothing it was earning, transfers the loss to a competitor it does not own, and—this is the part that makes it attractive rather than merely aggressive—expands the total market, because the complement’s demand rises as its own price falls.

What is new in 2026 is not the theory, which is old, but the simultaneity. Three distinct classes of Western actor executed the move at scale within a single year and a fourth announced

it; one of the four stated the rationale on the record, and the others are inferred from what they built rather than from what they said. The result is a stack in which every layer is being commoditized by somebody, which raises the question this chapter is ultimately about: when everyone is giving away everything, where does the rent actually land?

17.2 The accelerator vendor commoditizes the model, and underwrites the buildout

The clearest instance is also the largest company in the sector, and it runs the move on two layers at once.

17.2.1 The model layer

NVIDIA released the Nemotron 3 family in December 2025—Nano, Super and Ultra—and released with it not only open weights but the *training datasets and the reinforcement-learning libraries* [701]. Nemotron 3.5 Lightning followed in August 2026 under the OpenMDW-1.1 licence, again with training data and recipes, distributed through Hugging Face, ModelScope, OpenRouter, SageMaker JumpStart and Google Cloud. The associated open pretraining collection exceeds ten trillion tokens, and NVIDIA publishes open robotics, protein-structure and autonomous-driving corpora alongside it. On the software side, TensorRT-LLM is Apache 2.0, and in December 2025 NVIDIA open-sourced the CUDA Tile intermediate representation and the cuTile compiler stack under Apache 2.0 with LLVM exceptions, describing it as the largest update to the CUDA platform since its invention [702].

The scale of the resulting position is measurable rather than rhetorical. In the week to 29 August 2026, Nemotron 3 Ultra was the sixth-largest model by tokens routed on the largest neutral router—on its *free* endpoint—ahead of every model from every Chinese laboratory except four [720] [A]. The caveat of Section 10.7 applies here as it applies there: router volumes are a small and promotion-distorted slice of global inference, and a free endpoint is the promotion in its purest form. Read as a direction rather than a census, a chip company is a top-ten model provider by volume, which is not a description anyone would have written of this industry two years ago.

The strategic rationale is on the record, twice, and neither statement is ambiguous. Jensen Huang, on the second-quarter fiscal 2027 earnings call: “The world will need both closed models and open models. And both closed models and open models are skyrocketing in use. . . . *Nearly all open models run on NVIDIA*. . . our position in open models is very good, because the CUDA ecosystem is literally everywhere” [700] [P]. And, more precisely still, NVIDIA’s Director of Developer Technology at the Six Five Summit on 27 August 2026: NVIDIA builds Nemotron because “this is going to be a multi-model world, and we see where there’s the bottleneck”—arguing that durable value accrues to proprietary data rather than to model weights, which is exactly why NVIDIA can afford to release the weights *and* the data [700] [P].

What changed in August 2026 is that NVIDIA stopped merely publishing weights and began buying the means of producing them. On 20–24 August 2026—the sources differ on the date and the difference is not resolved here—NVIDIA and Poolside disclosed a transaction of approximately \$7B: \$6B for a *non-exclusive* licence to Poolside’s “Model Factory” training platform, \$1B of equity at a \$12B pre-money valuation, and offers to 109 Poolside engineers and researchers to join NVIDIA and work on Nemotron [724] [P]. Poolside remains independent, its founders do not join, and the licence excludes no one else from taking one. The reported purpose is stated in terms this chapter has to take seriously: to build a leading American open-weight model capable of competing with DeepSeek and Moonshot.

That is a description of a market rather than an aspiration, and the market can be dated to the same fortnight. On 28 August 2026 Tencent open-sourced the Hy4 preview—770B total

parameters, 49B active, a context window above one million tokens—the largest open-weight release of the last week of August, and three of the four open-weight releases in that window were Chinese [732] [V]. DeepSeek had meanwhile taken its first outside capital in June—roughly \$7.4B at a valuation above \$50B, with the state’s Big Fund as the sole voting holder [733] [A]; a late-August report of a \$74B valuation ahead of a 2027 onshore listing conflicts with that account and is not used here. The layer NVIDIA is buying its way into is not vacant. It is the most crowded layer in the industry and it is disproportionately Chinese, which is the condition that makes a \$6B licence for a training platform legible as something other than eccentricity.

Two days before the Tencent release, on 26 August 2026, *The Information* reported that NVIDIA was closing on an acquisition of Hugging Face for approximately \$12.9B [725] [A]. The report is a report and nothing more: neither company had confirmed it at the currency date, and *Business Insider* reported that talks had produced no signed agreement and that the transaction “could still atomize.” This book does not assert that it happened, and the analysis that follows is written to survive its not happening. What makes it worth printing unconfirmed is the reversal it would represent. Hugging Face reportedly *declined* a \$500M NVIDIA investment at a \$7B valuation in 2025, on the explicit ground that it did not want a dominant investor able to sway its decisions [725]. A hub that refused a minority stake on independence grounds in one year and is reported in acquisition talks with the same buyer in the next is either a firm that changed its mind about independence or a firm that discovered what independence costs; the record at 9 September 2026 does not say which, and the possibility that the report is simply wrong remains open.

Take the three together—Nemotron, the Model Factory licence with its 109 researchers, and, if the report is right, the hub—and the description this chapter has used so far stops being adequate. A firm that publishes open weights as a loss-leader is subsidizing a complement, which is the move of Section 17.1 and nothing more. A firm that acquires the training platform on which open models are produced, the researchers who produce them, and the distribution point through which very nearly all open weights reach their users is doing something categorically larger: it is acquiring the *supply chain* of the open-model layer—the factory, the staff and the warehouse. That is a different object from a commons, and the chapter should say what it is instead. Linux is a commons in the strict sense: the licence is irrevocable, no single vendor’s failure can withdraw the kernel, and the distributions that grew around it competed with each other for the surplus. An open-model layer whose training platform is licensed from one accelerator vendor, whose senior researchers are on that vendor’s payroll, and whose distribution hub is that vendor’s subsidiary is not a commons in that sense. It is a vendor-controlled distribution: free to use, forkable in principle, and shaped in practice by one balance sheet’s judgements about what gets trained, on whose silicon, and what appears on the front page.

Both readings are available and this chapter declines to choose. On the first, that is simply how a complement gets supplied reliably. Open weights of frontier quality do not appear because a market would like them to; they appear because somebody pays for the compute, the corpus and the people, and the natural payer is a firm whose revenue rises with inference volume. The arrangement is unusual only in that the payer says out loud why it is paying. On the second, the same facts describe the concentration of an entire layer under a single owner—and the layer in question is the one this book has treated throughout as the principal solvent of proprietary advantage, which makes its ownership a different kind of fact from the ownership of a chip design. Note that the two readings assign opposite meanings to the same detail: the non-exclusivity of the Poolside licence is, on the first, proof that the platform stays available to rivals, and on the second, the feature that moves the people who know how to operate it while keeping the transaction outside merger review.

Which raises the structural question, and it is not this book’s to settle. The Poolside transaction is the third of a kind. In December 2025 NVIDIA took a non-exclusive licence to Groq’s inference technology together with roughly ninety per cent of Groq’s staff—including

its chief executive and its president—for approximately \$20B, leaving Groq to continue under new management; in 2026 it took Enfabrica’s technology and team for approximately \$900M; in August 2026 it licensed Poolside’s platform and hired 109 of its researchers for approximately \$7B [726] [V]. Roughly \$27B in nine months, each transaction structured so that the target stays nominally independent and, on the reported structure, outside Hart–Scott–Rodino merger review. Senators Warren and Blumenthal put the objection in a sentence on 23 March 2026: “by licensing its technology and hiring its most important employees, NVIDIA has effectively acquired Groq in all but name” [726] [P]. No enforcement action had been taken by the currency date, and a letter from two senators is not a finding of law; the structures may well be lawful and are certainly not secret. The observation this chapter adds is narrower and follows from the form rather than from anyone’s motive. An equity acquisition of Hugging Face *would* require Hart–Scott–Rodino review, which is precisely why the report matters more than its price. Every previous transaction in the sequence was built so that the question would not be asked. This one, if it is real, asks it.

Set that beside Chapter 10. The argument for releasing weights at marginal price, on the ground that the rent will land somewhere else, is the argument this book attributed to Chinese laboratories and to the industrial-policy documents behind them. It is here being made, in public, by the incumbent whose rent the Chinese version of the move is aimed at. The two are not doing the same thing to the same end—NVIDIA is defending silicon, and China is defending an eventual position in hardware volume, domestic cloud and standards—but they are doing the same thing.

17.2.2 The datacentre layer

The second move is larger and less familiar, and it is the one that gives the first its durability. Through 2026 NVIDIA converted AI compute from a capital good its customers had to finance into an asset class that third-party capital would finance, and it did so by standing behind the residual value.

The filed positions, which supersede a good deal of looser reporting, are these. On 17 August 2026 NVIDIA disclosed multiple residual value guaranties with SB Energy on datacentre leases at Pike County, Ohio, whose tenant is an OpenAI affiliate: approximately 4.25 GW of IT load initially with an option on a further 3.8 GW, a twenty-year maximum term, and an aggregate payment obligation “cumulatively capped at \$105 billion” [694] [P]. The quarterly filings put maximum gross guarantee exposure at \$108.5B, supply and capacity purchase commitments at \$279B, AI-cloud revenue-share commitments at \$36B, future equity investment commitments at \$25B, datacentre leases held for third-party reassignment at \$20B, and equity securities at roughly \$95.6B [696] [P]. The same Ohio filing records that NVIDIA invests \$1.5B in SB Energy and becomes the site’s exclusive AI compute infrastructure provider, with the first phases effective in fiscal 2029 [694]. Alongside these sit a residual-purchase obligation to CoreWeave capped at \$6.3B running to 2032 and a 9.3% passive stake in Nebius; and, on reporting rather than on a filing, a \$1.5B arrangement with Lambda and a revenue-sharing programme for neoclouds [697] [A]. And on 10 August 2026 NVIDIA announced memoranda with Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR to establish financing platforms mobilizing “over \$500 billion of third-party capital” [692] [P]—memoranda, it should be said, expressly subject to execution of final agreements.

The regulatory precondition is worth naming because it is almost never reported alongside the announcement. On 29 July 2026, twelve days earlier, the SEC’s Office of Structured Finance took the position that securities in datacentre securitizations are not asset-backed securities under the Exchange Act, because datacentres are “tangible, physical assets that endure beyond the tenor of the securities” rather than self-liquidating financial assets—placing them outside Regulation AB and outside the five per cent credit-risk-retention requirement [693] [V]. The financing structure became cheap to build a fortnight before it was built.

NVIDIA’s own framing of what this is for could hardly be more explicit. Huang: the company “began by building chips; today, we are helping create a new class of productive, investable infrastructure: AI factories,” and NVIDIA compute is “fungible and transferable across customers and operators.” Goldman Sachs’ chief executive, on the same release: “a market for credit backed by NVIDIA compute” [692]. On the earnings call: “Independent capital still underwrites every deal on its own merits. *We’re not making loans*. In this model, we get paid twice—once on the hardware sale, and again through the share of rental revenue” [699] [P].

Read the two moves together and the shape is unmistakable. If the model layer is a commons that anyone can serve, and the datacentre layer is financed by capital markets against the residual value of the accelerator rather than against the creditworthiness of any particular tenant, then the failure of any individual model company does not remove demand for the silicon. The commons survives its sponsors. Firms would appear to serve it in the way that distributions, support vendors and integrators appeared around Linux once the kernel was free and no longer owned by anyone whose bankruptcy could take it away. The qualification introduced above applies here with force, and the reader should carry it forward: the analogy holds only for as long as the commons is genuinely unowned, and the acquisitions of August 2026 are an attempt to own not its licence—which cannot be owned—but its means of production. That is a coherent strategy for a company whose exposure is not to which laboratory wins but to whether inference happens at all, and it is the mirror image—in method, not in beneficiary—of the argument Chapter 10 makes about open weights.

17.2.3 The counter-evidence, which is substantial

The book’s discipline requires the other half, and here it is unusually strong. Four things cut against reading this as a settled, confident strategy.

First, it was partly suspended. On 27 August 2026, one day after the earnings call, NVIDIA paused parts of the July revenue-sharing and credit-support programme, reportedly after employees flagged potential antitrust exposure and after friction over how far NVIDIA could dictate which customers its partners rented to [698] [A]. NVIDIA’s own statement is that the model “is still in place and continues to evolve due to high demand” [698] [P]. A programme parts of which are halted within eight weeks of launch is not the execution of a long-laid plan.

Second, the trajectory is retrenchment rather than escalation—though the most-repeated version of that claim is half the story. A \$100B commitment to OpenAI announced in September 2025 became a \$30B equity position; talks reported at up to \$250B for the Ohio structure became a reported figure under \$120B covering only the first phase, and then \$105B on the filing. The half usually dropped is that while the equity commitment fell by seventy per cent, the *compute* commitment rose: on the second-quarter call Huang put OpenAI’s “existing and planned commitments” at approximately twelve gigawatts of NVIDIA compute through 2030 [728] [P]. NVIDIA reduced what it was prepared to own of its largest customer and increased what it expects to sell that customer. Whether that reads as retrenchment or as discipline depends on which of the two numbers one thought was the strategy, and the popular framing—“NVIDIA cut its OpenAI bet by seventy per cent”—reports only the number that fell. Third, the structures themselves are narrower than the headline. The Ohio guarantee *terminates early if OpenAI achieves a satisfactory credit rating*, and OpenAI must reimburse and indemnify NVIDIA for every amount paid [694]—which makes it, structurally, a bridge to creditworthiness rather than permanent underwriting. Management denies the vendor-financing characterization directly, and has done since 2025: “unlike Lucent, NVIDIA does not rely on vendor financing arrangements to grow revenue” [705].

Fourth, and against all of the above, there is one number in the filings that supports the circular reading better than any critic’s argument. In the second quarter of fiscal 2027, NVIDIA’s GAAP diluted earnings per share (\$2.46) *exceeded* its non-GAAP figure (\$2.22)—an inversion of the usual relationship—because GAAP other income of \$7.773B was almost entirely gains

on equity securities [695] [P]. That pre-tax figure equals roughly an eighth of the quarter’s GAAP net income—the after-tax contribution is smaller, since the gain is pre-tax and the income is not—and NVIDIA’s disclosed equity portfolio of some \$95.6B is concentrated in the firms that buy its chips [696]. That is not an accusation; it is a line in a filing, and it is the single most quantitatively honest statement of the concern available. The best-argued external version is Vellante’s: independent capital can fund more factories, but independent capital is not independent demand, and the very need for credit enhancement is evidence that the market does not yet accept compute as enduring collateral unaided [705] [A]. Morgan Stanley initiated credit coverage at Neutral, estimating roughly \$200B of total credit exposure by end-2028 of which about \$170B off balance sheet, and coined “balance-sheet-as-a-service” [704] [A].

17.2.4 The quarter, as the filings state it rather than as it was reported

One correction of the record belongs here, because the most closely watched quarter in the sector was reported with its most informative disclosure missing. NVIDIA’s second-quarter fiscal 2027 CFO commentary changed the reporting structure, and almost all coverage passed over it [727] [P]. “Gaming” has ceased to be a reported market platform; GeForce is folded into a new *Edge Computing* line at \$7,198M, up twenty-seven per cent. More consequentially, Data Center is now reported in two pieces: *Hyperscale* at \$48,710M, up 102 per cent, and *AI Clouds, Industrial & Enterprise* at \$40,313M, up 138 per cent, with management stating that the second “will represent roughly half of our data center business” going forward. The split is itself evidence, and it bears directly on Section 17.3: the faster-growing half of NVIDIA’s datacentre business is *not* the hyperscalers who are building their own accelerators. If custom silicon is a substitution threat, it is a threat concentrated in the slower-growing half of a business whose other half compounds at 138 per cent—which does not make the threat unreal, but does locate it.

Three further items are worth stating precisely because they are usually stated loosely. Networking was described on the call as a record quarter growing eighteen per cent sequentially, and *no dollar figure was published*: a company that volunteered a datacentre split it was under no obligation to volunteer withdrew the number for the layer this chapter argues is being substituted first, and the asymmetry deserves recording rather than interpreting. Hopper shipments to China were under one per cent of datacentre revenue, and the outlook assumes *no* China datacentre compute revenue at all—the cleanest available statement both of what Chapter 21’s controls have accomplished and of how little of NVIDIA’s forward number now depends on their relaxation. And management guided on the call to roughly seventy per cent revenue growth in fiscal 2028 against a sell-side consensus near forty-four [727] [P]—a forward statement that appears in the transcript and not in the press release, and on any reading the most consequential thing said in the quarter. It is guidance and not a result, offered here as what management believes rather than as what will occur [S]; a company that expects seventy per cent growth two years out is not a company that believes it is being commoditized, and a company that guides to it publicly has also made itself falsifiable, which is more than most of the claims in this chapter can say.

Huang’s own framing of the demand side names the thing this chapter is about. Listing the drivers of the buildout, he sets “a thriving open-model ecosystem” alongside multiple frontier laboratories scaling in parallel and physical AI coming online, and compresses the entire argument of Section 17.1 into three words: “compute is revenue” [728] [P]. A chip vendor that lists the open-model ecosystem among its demand drivers has explained, without being asked, why it would buy one.

17.2.5 AMD runs the same playbook, which matters for the diagnosis

AMD releases Instella and Instella-MoE as *fully* open models—weights, data and training code—positions ROCm 7 explicitly against CUDA, and chairs the UALink consortium whose promoter group is every major buyer of accelerators and does not include NVIDIA [708]. On the financing

side it has issued warrants for up to 160 million shares at one cent to OpenAI, and an equivalent 160-million-share warrant to Meta, vesting against gigawatt purchase tranches and escalating share-price thresholds; committed up to \$5B of equity to Anthropic against 2 GW of deployment; and disclosed \$4.1B of datacentre lease guarantees [707, 706] [P]. Stated plainly, the customer is being offered payment in the vendor’s own equity for buying the vendor’s product—though as of the August 2026 filing no warrant shares had vested or been issued [707].

That AMD runs the same playbook is analytically important, because it argues for competitive necessity rather than grand design. When two rivals independently adopt an unusual instrument, the more parsimonious explanation is that the instrument is what the situation requires, not that both had the same idea about the future shape of the industry. The book should hold both readings; Section 17.6 says what each licenses.

17.3 The laboratories and the hyperscalers commoditize the silicon

The move in the other direction is equally overt, and Hot Chips 2026 was where it became impossible to describe as a hedge [672]. Appendix C carries the ledger; the strategic content is this.

OpenAI presented Jalapeño, an inference accelerator co-designed with Broadcom: 13.4 PFLOP/s at four bits, 216 GiB of HBM4 at 15.4 TB/s, a 128-chip scale-up domain and a 2,048-chip pod, under a term sheet covering ten gigawatts of deployment through 2029 [683, 684] [P]. Google announced, for the first time, a training part and an inference part developed concurrently rather than alternately, and—the more consequential fact—that it would *begin* delivering TPUs to a select group of customers in those customers’ own datacentres, with a small share of that hardware revenue in 2026 and the majority expected in 2027 [685, 686] [P]. Microsoft’s Maia 200 makes an economic rather than an architectural claim: thirty per cent better performance per dollar than the latest generation in its fleet, with NVIDIA unmentioned [687]. Meta has disclosed a four-generation MTIA roadmap on roughly a six-month cadence—of which one generation is in production, the next in lab testing—with hundreds of thousands of earlier parts already deployed [688]. AWS states more than a million Trainium processors deployed and calls it a multi-billion-dollar business [689]. Behind all of them sits Broadcom, whose AI semiconductor revenue grew 143 per cent to \$10.8B in its second fiscal quarter with the next quarter guided above \$16B [691] [P].

A defensible order of magnitude follows from two primary filings rather than from any analyst’s estimate. Annualizing Broadcom’s guided AI semiconductor run rate (above \$16.0B a quarter, so above \$64B a year) against NVIDIA’s datacentre run rate (\$89.0B in the quarter, so about \$356B a year) puts merchant custom-ASIC revenue at roughly fifteen per cent of the combined total; annualizing against NVIDIA’s *total* third-quarter guidance of \$108.0B instead, of which datacentre is the overwhelming majority, puts it nearer fourteen; taking both companies’ most recent actuals rather than their guidance puts it nearer eleven. Call it eleven to fifteen per cent, and note that the range is set by which quarter and which basis one chooses rather than by any disagreement about the numbers [M]. Custom silicon is not yet a substitute for the merchant GPU. It is, however, growing several times faster, and it is being bought by precisely the firms whose purchase orders set NVIDIA’s revenue.

The fabric tells the same story more sharply, because there the substitution has already happened. Among the parts presented with an off-package scale-up fabric, the Ethernet-class designs are the clear majority—AMD’s UALink-over-Ethernet on 512-port merchant switch ASICs, Microsoft’s single Ethernet protocol from chip to a 6,144-accelerator Clos under a slide titled “Why Unified?”, SambaNova’s ten integrated 800G ports per socket, Meta’s RoCE on dedicated I/O chiplets, and OpenAI’s 2,048-chip pod over a Broadcom Tomahawk 6 Clos. Of the alternatives, Google’s optically switched torus is proprietary but is not sold, and NVIDIA’s

NVLink is the only proprietary scale-up fabric a third party can buy [681, 683, 685]. The layer that was the deepest source of lock-in three years ago is being rebuilt out of parts anyone can buy.

17.3.1 Jalapeño as a direct challenge, and the sentence that constrains how it can be read

The Hot Chips presentation was an architecture; three weeks later there were numbers. On 25 August 2026 OpenAI published first results, and they are the sharpest claim any custom part has yet made against the merchant GPU: 1.5 to 1.9 times more AI work per watt at peak throughput, and 1.7 to 3.6 times lower end-to-end latency, measured against GB200 and GB300 systems and against AMD’s MI355X, on GPT-OSS 120B, DeepSeek R1 670B and Kimi K2.5 [729] [P]. Two qualifications belong in the same breath and are not decoration. The part *was not deployed at the currency date*—OpenAI says it will “begin deploying” Jalapeño within its own compute infrastructure by the end of the year, which means every figure above is a vendor-run measurement of a part in no production service. And SemiAnalysis observes that the comparison sets HBM4 parts against older HBM3E systems, calling it “somewhat incomplete and unfair” [729] [A]. “Beats NVIDIA by 1.9×” is therefore a contested claim rather than a settled one, and this book carries it as contested. Nor does it carry a figure for how much of OpenAI’s inference the part will serve: *OpenAI has not stated one*, the temptation to supply an estimate should be resisted, and the whole commercial significance of the programme turns on the share that has not been stated.

What the same post does state is the sentence that constrains every reading of it: “We will continue to widely deploy accelerators from NVIDIA and other partners for both training and inference workloads” [729] [P]. Set it beside Huang’s twelve gigawatts of OpenAI commitments through 2030 [728] and the two firms are saying the same thing from opposite directions—the challenger says it will keep buying, and the incumbent says it has been told so. That is the shape this chapter has argued for throughout: a layer being commoditized by a firm that goes on buying from it. It is not the shape of displacement, and the distinction between the two is the whole content of Section 17.6’s second prohibition.

The analyst reading divides cleanly and both halves are worth printing. Alexander Harrowell of Omdia calls Jalapeño “the biggest competitive threat to NVIDIA,” on the ground that roughly half of AI infrastructure spending now comes from firms running or capable of running custom-silicon programmes; Adrien Sanchez of Yole locates the threat more precisely, at NVIDIA’s *inference* margins rather than at its training franchise—the more defensible version, and the one this chapter adopts as the live hypothesis, since inference is where a fixed-function part’s advantages survive contact with a workload. Against them, Daniel Newman of Futurum calls the “NVIDIA killer” framing “exhausting” and credits Broadcom rather than OpenAI for the engineering, which has more force than it first appears: if the engineering is Broadcom’s, then the rent displaced from the merchant GPU lands at a merchant ASIC vendor rather than at a laboratory, and the laboratory has commoditized its own supplier only to acquire a different one [730] [A]. Huang’s answer is the incumbent’s classic one and is recorded here rather than dismissed: “Lots of projects get started. Lots of projects get canceled... We have the supply chain and the technology scale to be their largest supplier” [730] [P].

One feature of the episode is easy to miss and matters to a book about strategy. *Neither OpenAI’s announcement nor Broadcom’s names NVIDIA*. The entire adversarial framing—the comparisons, the “killer” language, the threat assessment—is supplied by analysts and by the press. The challengers decline to describe the challenge as one, which is either commercial courtesy toward a supplier from whom they still expect twelve gigawatts, or an accurate account of an intent that was never adversarial; from outside the firms the two are not distinguishable. This chapter records the framing as third-party framing, because that is what it is, and notes that the discipline cuts both ways—a challenge is not less real for going unnamed, and it is not

more real for being named by someone with no part in it.

A correction of the record belongs here too, dryly, because a false claim circulates alongside this story and the true version is more useful. It is repeated that NVIDIA’s head of datacentre left for OpenAI. He did not, and the movement ran the other way: *OpenAI*’s own Head of Data Centers, Chris Malone, left *OpenAI* in the week of 17 August 2026—reported on the 25th—after sixteen months, amid a continuing stream of senior departures. The provenance is worth stating precisely, because it is commonly garbled: the “head of data centers” title was OpenAI’s, not Meta’s—at Meta he spent close to five years as a *distinguished engineer*, and before that nearly twelve years at Google on datacentre infrastructure and IT hardware research—and he had never worked at NVIDIA [731] [V]. That career shape is itself the point: the person OpenAI recruited to build its datacentres was an infrastructure engineer of seventeen years’ standing at the two firms that had built the largest ones, and OpenAI kept him for sixteen months. OpenAI’s infrastructure organization now reports to Sachin Katti, recruited in November 2025 from Intel, where he was chief technology and AI officer. NVIDIA’s datacentre leadership is intact—Ian Buck remains VP of Hyperscale and HPC—and the senior technical flow of 2026 ran *toward* NVIDIA: Groq’s leadership and roughly ninety per cent of its staff in December 2025, then 109 Poolside researchers in August 2026 [726, 724]. This is not scorekeeping. A book that takes the custom-silicon challenge as seriously as this chapter does has to get the direction of the flow right, because a labour market moving toward the incumbent is not evidence of the incumbent losing the argument—and it is not evidence of the incumbent winning it either, since NVIDIA paid roughly \$27B for those two flows and a market in which talent must be bought at that price is not obviously a market at rest.

17.3.2 And the counter-evidence here is decisive enough to be printed first in any honest summary

On 26 August 2026, in the last week before this edition’s currency date, AWS and NVIDIA announced two million additional NVIDIA GPUs across 2027–2028—and, materially, that *Trainium4* will adopt NVIDIA’s *NVLink Fusion* interconnect and NVIDIA’s custom HBM [690] [P]. The largest deployer of non-NVIDIA AI silicon in the world is placing its own accelerator inside NVIDIA’s rack-scale interconnect and memory stack. Set against NVIDIA’s quarter—datacentre revenue up 117 per cent to \$89.0B, against total-company third-quarter guidance of \$108.0B ± 2 per cent [695]—the honest reading is not that the incumbent is being displaced by the custom-silicon wave but that it is absorbing it. The commoditization of the silicon layer is real, it is well funded, and it has so far moved the boundary of what NVIDIA sells rather than the size of what NVIDIA sells.

17.4 And one actor commoditizes the fab

At the bottom of the stack, SpaceX confirmed in August 2026 a \$16.8B initial investment in “Terafab,” a vertically integrated logic and memory manufacturing, packaging and test site in Grimes County, Texas, with output targeted at Tesla’s robotics and vehicle programmes and at space-based datacentres [709]. Separately, and not disclosed as part of Terafab, Intel entered a pact with Tesla, SpaceX and xAI in April 2026 covering advanced logic and packaging, which one analyst reads as an Intel fab expansion with those three as anchor customers rather than a standalone venture [709] [A]. In parallel Tesla runs a dual-foundry arrangement for AI5 and AI6 at TSMC and Samsung, with the Samsung Taylor 2 nm fab reaching AI5 tape-out in July 2026 [710].

This one deserves the coldest grade in the chapter. The announced figures have moved three times—from a \$20–25B analyst estimate in March, to \$55B initial and \$119B total in a May regulatory filing, to \$16.8B initial in the August confirmation. The process node was not restated

in the confirmation, no construction or production timeline was disclosed, and the project was not discussed on SpaceX’s first earnings call in the same week. Terafab belongs in this chapter as evidence of intent [R], not as capacity. What it demonstrates is the reach of the logic: an actor whose rent is in launch, vehicles and robots, facing a fabrication layer it regards as a constraint, reasons its way to owning the constraint. That is move five of the playbook—supply-chain capture—executed by an American company against an American and Taiwanese supply chain.

17.5 The same method, and the asymmetry that survives it

Table 17.1 sets the actors side by side. The columns that matter are the second and third: what each is defending, and what each is willing to burn to defend it.

Table 17.1: Who commoditizes what, as of 9 September 2026. “Instrument” lists the specific vehicle, not the intent. “Commoditized” is used broadly, and the breadth matters: it covers both *giving a layer away*, which is what NVIDIA and AMD do at the model layer, and *substituting for it in-house*, which is what a hyperscaler building its own accelerator does—the first destroys a rent for everyone, the second only removes the actor from paying it. The two are different instruments with the same effect on the layer’s pricing power, and the last two rows are of the second kind. Status: **running** (executing and observable), **partial** (executing with a component suspended or unexecuted), **announced** (committed, not yet productive).

Actor	Layer defended	Layer commoditized	Instrument	Status
China (state and firms)	Hardware volume, domestic cloud, standards, dependence	<i>All of them</i>	Open weights, open harness, open ISA, open interconnect; mandate; power tariff	Running
NVIDIA	Accelerator and CUDA	The model, and the financing of the datacentre	Nemotron weights <i>and</i> training data; OpenMDW; cuTile (Apache 2.0); Poolside “Model Factory” licence and 109 researchers; reported Hugging Face acquisition [A]; residual value guarantees; revenue share	Partial (revenue share paused Aug.)
AMD	Accelerator	The model; the scale-up fabric	Instella (weights, data, code); ROCm; UALink over merchant Ethernet; customer warrants	Running
OpenAI	The frontier model and the platform	The accelerator and the system	Jalapeño with Broadcom; 10 GW term sheet	Running
Google, Microsoft, Meta, AWS	Cloud platform and applications	The accelerator, the system, the fabric	TPU 8t/8i sold into customer sites; Maia 200; MTIA; Trainium; Ethernet scale-up	Running
Broadcom, Marvell	Co-design and networking	The merchant GPU	Custom ASIC programmes; eRoCE	Running
SpaceX, Tesla	Launch, vehicles, robots	The fab	Terafab; dual-foundry AI5/AI6	Announced

Two asymmetries survive the convergence of method, and they run in opposite directions.

The first favours the West, and the book should say so plainly. Commoditization is not the same as losing. A stack in which every layer is cheap is a stack in which the surplus accrues to whoever holds the layer that stays scarce—and at the currency date the scarce layers are, for the most part, still Western-held: the CUDA installed base, the enterprise deployment surface, the trust and certification apparatus of Chapter 18, and the deployment-generated data of Chapter 15. The clearest single demonstration that the West is executing rather than retreating is the price inversion of Section 10.7: on 30 July 2026 the incumbent cut the bottom

tier of its ladder by eighty per cent to \$0.20 and \$1.20 per million tokens, which is *below the list price of the frontier-scale Chinese open-weight models*—though not of their distilled tiers, a restriction Section 10.7 states precisely—and usage of that tier rose to almost fourteen times its pre-cut level with roughly three-quarters of the gain taken from competitors [717, 718] [P/A]. A firm that can price underneath the commodity and gain share doing it is not a firm being commoditized. It is a firm commoditizing.

The second favours China, and it is structural rather than tactical. Every Western move in Table 17.1 commoditizes a layer in which *another Western firm* was earning the rent. NVIDIA is destroying the model layer’s rent in order to protect silicon. OpenAI and the hyperscalers are destroying the silicon layer’s rent in order to protect the model and the platform. Each decision is individually rational, and the honest statement of what they do jointly has to respect the mechanism of Section 17.1: they *redistribute* the Western rent pool rather than necessarily shrinking it, because the surplus a commoditizing move destroys at one layer is the surplus it delivers to another. The pincer claim therefore has to be conditional, and this book states it conditionally: it bites to the extent that the layers the rent moves *to* are layers a Chinese actor can also occupy, and not otherwise. What China gains from the arrangement is not the destruction of a rent pool but a cheaper entry condition. It does not need to destroy either layer; it needs only to be present in both when the rents have moved, and presence is very much cheaper than victory.

Sharper still, and this is the sentence the chapter exists to reach: *the layers the West is commoditizing are disproportionately the layers where China is behind, and several of the commoditizing moves directly reduce the value of Western state instruments.* Open weights released with their training data lower the cost of closing the model-layer gap for whoever is closing it. An open instruction set that NVIDIA has publicly qualified as a CUDA host lowers the cost of the compute-layer detour that Chapter 12 documents [673]. A scale-up domain rebuilt out of merchant Ethernet switch silicon is precisely the substitution that the chokepoint strategy was designed to prevent, and it is being performed by the firms the chokepoint strategy was designed to protect [681]. None of these is a concession to China; all of them are commercial decisions taken against a domestic rival. The externality is the same either way. A denial policy that assumes the denied party must obtain the controlled item from the controlling country is undermined as effectively by an American open-weight release as by a Chinese one, and the American release comes with the training data attached.

17.5.1 What this does to Proposition P3

The book’s transmission argument for P3 required Chinese price pressure to reach Western margins: cheap open weights would force the closed frontier to cut, the cut would compress the rent, and the compressed rent would fail to finance the next training run. Section 10.7 shows that chain is weaker than it looked—the incumbent cut its cheapest tier, not its flagship, held the flagship for three weeks, made the flagship cut temporary, and the price leader on the Chinese side *raised* prices by as much as 1,100 per cent mid-cascade [719]. But P3 does not depend on that chain any more. It is now overdetermined in *mechanism*, because the largest supplier of the frontier’s inputs is itself releasing open weights with training data at zero price, and the frontier’s largest customers are building the silicon underneath it. One qualification is owed, and it is the weakest joint in this chapter’s argument, so it is stated here rather than left for a reader to find. NVIDIA’s *shipped* open models take routed volume but not frontier capability—the best of them scores well short of the leading open Chinese weights on the independent index (38 against 60 on the August v4.1.1 board, Appendix H)—so what NVIDIA has demonstrated to date is that a chip vendor can commoditize the *volume* tier of the model layer, not that it can commoditize the capability tier. The stronger claim, that rent destruction at the model layer no longer requires a Chinese actor to arrive, is contingent on a frontier-scale NVIDIA release that at 9 September 2026 has not happened [S]. The Poolside transaction is the clearest evidence yet that

such a release is being funded rather than merely wished for—\$6B for a training platform and 109 researchers is a capability purchase and not a marketing one [724]—but a funded intention remains an intention [R]. What is not contingent is the direction of travel and the stated intent, and those are on the record.

That is a strengthening of P3’s mechanism and simultaneously a weakening of P3’s *diagnostic value for this book’s thesis*. A proposition that can be satisfied by the incumbent’s own conduct is a weaker piece of evidence about who wins, and the register in Appendix D should be read with that in mind.

17.6 What the mirror licenses, and what it forbids

It licenses three claims. That the playbook of Part I is a general theory of a contest over a layered technology stack rather than a description of Chinese exceptionalism—which is a considerable strengthening of Part I, because a theory that predicts the behaviour of actors it was not built from is doing more work than one that does not. That the commoditization of the AI stack is now overdetermined, proceeding from several directions at once and no longer contingent on any single actor’s strategy holding. And that the West’s own commercial decisions are eroding the West’s own policy instruments, which is a finding Chapter 23’s instruction to commoditize before being commoditized has to absorb rather than a rhetorical flourish: an export-control regime and an open-weight corporate strategy are not merely in tension, they are aimed at the same object from opposite sides.

It forbids three claims equally. That the West is passively losing—it is not; it is executing hard, it has priced underneath the commodity at the bottom of the ladder and gained share doing so, and its incumbent grew datacentre revenue 117 per cent in the quarter this chapter closes in. That commoditization implies a Chinese victory—it does not; it implies that the rent moves, and whether it moves to a layer China holds is the open question of the whole book rather than a corollary of this chapter. And that NVIDIA’s financing structures are straightforwardly vendor financing—management denies it in terms, the largest structure terminates on a credit-rating event and carries full indemnification, one programme was suspended within eight weeks over antitrust concerns, and the honest verdict on the evidence available at the currency date is that the practice is unusual, large, growing, disclosed, and not yet demonstrated to be either sound or unsound.

One item is neither licensed nor forbidden but held open, and it is named here so that a reader can date it. If the reported acquisition of Hugging Face is confirmed, the description of the open-model layer in Section 17.2 changes in kind rather than in degree, and the Linux analogy this chapter uses twice must be withdrawn rather than qualified: a commons whose warehouse belongs to the vendor of the complement is a distribution channel with a permissive licence attached. If it is not confirmed—and at 9 September 2026 it is not—the Poolside licence and its 109 researchers stand on their own as the strongest evidence available that the open-model layer now has a principal sponsor, which is the weaker claim and still a large one.

Two tests would settle much of this by the end of 2027, and both belong in the register maintained in Appendix J. If merchant-GPU datacentre revenue growth decelerates while custom-ASIC revenue continues at Broadcom’s guided trajectory, the silicon commoditization is working and the pincer is real. If instead custom silicon continues to converge onto NVIDIA’s interconnect and memory—as Trainium4 has already done [690]—then the wave is being absorbed, the accelerator vendor’s giveaway of the model layer was correctly priced, and the layer that stays scarce is the one it has held all along.

Part III

The Consequences

Chapter 18

Trust: The Contested Layer of the Open Commons

Everything in Part II treated openness as an adoption weapon. This chapter treats its exposed flank. Open weights eliminate API dependency; they do not establish trust—and trust, Chapter 8 argued, is precisely the property that has stalled Chinese campaigns in aircraft, drugs, and avionics-grade markets. If the AI war’s deployment layer behaves like a certification regime rather than a commodity market, the commons can win every download metric and still lose the workloads that pay. This chapter surveys the evidence on all three trust dimensions—supply-chain security, behavioral integrity, and provenance/IP—then states the sovereign trust stack that third countries would need to adopt the commons safely, because whether that stack gets built, and by whom, is itself now a theater of the war.

18.1 The supply chain is an attack surface

The model-distribution infrastructure underlying the entire open ecosystem has been demonstrated attackable, repeatedly [V]. The serialization layer first: roughly 100 malicious models were found on Hugging Face in early 2024, including reverse-shell pickle payloads; the “nullifAI” technique (February 2025) used deliberately broken pickle streams that execute before the scanner errors out; and the scanner itself, Picklescan, yielded five CVEs in 2025 [282]. The namespace layer: Unit 42 showed that re-registering deleted author names lets attackers serve malicious models under trusted identities *automatically pulled* by Google’s and Microsoft’s own model catalogs [283]. The platform layer: Hugging Face itself disclosed a July 2026 breach—attackers exploited dataset-processing code paths, harvested credentials, and moved laterally through internal systems over a weekend; no public artifacts were found tampered with, but the demonstration stands, and HF noted the campaign was run by an *agentic* attacker framework [284]. The training layer is worse: Anthropic and the UK AI Security Institute showed that ~250 poisoned documents suffice to backdoor a model regardless of scale—a near-constant, not a percentage—and the sleeper-agent literature shows such backdoors surviving fine-tuning and RLHF [285, 286]. The defensive response—HF/VirusTotal continuous scanning of 2.2M+ repositories, safetensors-by-default, JFrog partnerships—is real and improving [287], but the honest summary is that the open commons currently runs on the security posture of early-2000s package managers, at national-infrastructure stakes. Chapter 19 takes this dimension to its full depth, because the combination this book forecasts—a small number of model families, thousands of derivatives, and agents holding real authority—changes it from a hygiene problem into a systemic one.

18.2 Behavioral integrity: the censorship record

For Chinese open weights specifically, the peer-reviewed record now documents what buyers of sovereignty must price in [V]. An audit of 646 politically sensitive prompts found *semantic-level suppression* in DeepSeek models—sensitive content present in the chain-of-thought but filtered from final outputs, systematic omission of transparency and accountability framings, and occasional amplification of state-aligned language [288]. The R1dacted study found the censorship distinct from ordinary safety refusal, topic- and language-dependent—and, critically, *transferring to distilled models*: the politics ride along with the capability into every derivative [289]. These findings do not make Chinese weights unusable; local fine-tuning and ablation communities work on exactly this. But they convert “just self-host it” into a three-layer problem: a locally hosted Chinese model delivers *operational* sovereignty (no API, no data egress) while leaving *epistemic* sovereignty (whose worldview is baked in?) and *supply-chain* sovereignty (who built the artifact you downloaded?) unresolved. The book’s Global South adoption argument must carry this asterisk—and so must the Western enterprise-adoption argument, in mirror image, since the same three-layer test applies to closed US models, which fail the operational layer by construction.

18.3 Provenance and the distillation war

The commons’ cost advantage has a contested genealogy. Anthropic’s February 2026 report alleges industrial-scale distillation campaigns—over 16 million exchanges harvested through roughly 24,000 fraudulent accounts, attributed to MiniMax (>13M), Moonshot (>3.4M), and DeepSeek (>150K)—following OpenAI’s earlier, vaguer claims about R1 [P; vendor allegations, not adjudicated] [290]. Three implications belong in the strategic ledger regardless of adjudication. If true at scale, part of the fast-follower speed of Chapter 10 is subsidized by the frontier it chases, which caps the strategy at frontier-minus- ϵ unless domestic capability generation compounds independently (the evidence that it does—the algorithmic innovations of §10.3—is substantial). Second, it hands Western policy a legal instrument—terms-of-service enforcement, KYC on API access, and IP litigation—that operates on exactly the trust-certification axis of Chapter 8. Third, it cuts against the clean “openness versus rent” morality tale: the commons is partly built from enclosed material, and the enclosers know it. The book records the dispute rather than resolving it; the indicator to watch is whether any jurisdiction converts these allegations into an enforceable certification barrier for open-weight deployment.

The allegation changed hands on the currency date itself. On 9 September 2026 the NSA, the FBI and CISA issued joint cybersecurity advisory AA26-251a, naming DeepSeek, Moonshot, Alibaba, MiniMax, StepFun and Z.AI as conducting distillation campaigns against US frontier models, with the traffic routed through cloud providers and third-party aggregators [754] [P]. Two disciplines govern how much weight that carries. The advisory concerns distillation—the harvesting of a frontier model’s outputs to train another—and not hosting: it says nothing about whether the named laboratories’ weights may be served, downloaded or deployed, and a reader who takes it as a finding against open-weight deployment has read a document that does not exist. And it is an allegation by intelligence and security agencies, not an adjudication by anybody; no court or tribunal has tested it, and the vendors’ own February claims, which it broadly tracks, remain unadjudicated too. What is new is nonetheless real, and it is institutional rather than evidentiary. This is the first time six laboratories—very nearly the whole of the Chinese open frontier—have been named together in a single US government document, and a claim that was a vendor’s grievance in February is a state’s position in September. The routing detail matters more than the roster, because it lands in the same fortnight that one of the six is negotiating a revenue share with three American clouds for serving its flagship [747] [A]: the intermediaries through which the advisory says the distillation traffic flows are the intermediaries

through which the laboratory proposes to be paid, and Section 10.4 in Chapter 10 carries the argument that a contract of that kind is what makes a laboratory regulable. The advisory does not close the trust gap this chapter is about. It documents, from the other side of it, that the gap is now a matter of state record.

18.4 The sovereign trust stack

For a third country, the rational response to all of the above is neither refusal nor naive adoption; it is infrastructure. The stack, concretely: signed model and dataset manifests with cryptographic hashes; reproducible build-and-quantization pipelines (the conversion step is itself an attack surface); non-executable packaging (safetensors-class) by default with remote code disabled; model bills-of-materials and dependency inventories; behavioral evaluation batteries covering backdoors, censorship, and geopolitical bias, run nationally rather than taken on faith; sandboxed loaders; attested serving firmware; and national artifact mirrors so that availability cannot be revoked or substituted upstream. None of this is exotic—it is what software supply-chain security did after SolarWinds, applied to weights.

18.4.1 Five layers of trust

“Trust” in this chapter has been used for five distinct properties, and separating them is what makes the stack operational rather than rhetorical:

Identity — who produced, converted, and distributed this exact artifact? The derivative chain of Section 10.6 means the answer is usually a list, not a name.

Integrity — do the weights, runtime, configuration, and firmware match what was certified? This is a cryptographic question and the only one with a clean answer.

Behaviour — does the system exhibit acceptable performance, bias, security, and failure behaviour *in the integration it will actually run in*?

Governance — who can update, revoke, audit, or constrain it after deployment, and through what process?

Liability — who bears legal and financial responsibility when it causes harm? For an open-weight derivative chain this is currently, in most jurisdictions, nobody in particular.

The five layers behave differently under the two stacks, which is why single-axis comparisons mislead. A locally hosted open-weight model scores well on governance (you control updates) and badly on identity and liability; a hosted closed API scores well on identity and liability (there is a counterparty with insurance) and badly on governance (the model can change under you) and on operational sovereignty. Table 18.1 crosses the five layers with the certification levels below, which is the form in which a procurement office can actually use them.

18.4.2 A graded certification ladder

A list of controls is not a policy instrument; a *graded scheme* is, because it lets a buyer state a requirement and a supplier claim a level. The six levels below are the natural staging of the stack above, and they are the form in which the certification authority of Chapter 20 would actually operate:

Level 0 — Identified. Hash-identified artifact; you know exactly which bytes you are running.

Level 1 — Non-executable. Tensor-only format, malware-scanned, no remote code on load.

Table 18.1: The trust matrix: what each certification level establishes, per layer. “—” means the level says nothing about that layer.

Level	Identity	Integrity	Behaviour	Governance	Liability
0 Identified	hash names the bytes	—	—	—	—
1 Non-executable	—	no code on load	—	—	—
2 Provenanced	full chain named	reproducible build & conversion	—	upstream known	attributable
3 Evaluated	—	—	backdoor, censorship, injection—on the integration	—	evidence exists
4 Contained	—	attested runtime	failure bounded by broker	immutable policy; append-only audit	operator-side clear
5 Assured	—	continuous	independent team, continuous monitoring	red revocation obligations	contractual, insurable

Level 2 — Provenanced. Reproducible quantization and conversion, signed manifests, model bill of materials, verifiable lineage from a named upstream.

Level 3 — Behaviourally evaluated. Published results on backdoor, censorship, geopolitical-bias, and prompt-injection batteries, run on the *integration* rather than the bare model.

Level 4 — Operationally contained. Attested runtime, deterministic tool broker, immutable policy, append-only external audit log.

Level 5 — Critical-infrastructure assured. All of the above plus independent red-teaming, continuous monitoring, and incident-response obligations with revocation.

Most of the open commons today ships at Level 0–1. Most regulated deployment will eventually require Level 3–4. *That gap is the deployment bottleneck of the entire commoditization thesis*, and closing it is a service someone will sell or a public good someone will provide—Chapter 27 argues for the latter.

One August 2026 release sharpened every argument above: MiniMax’s H3 tops the open-weight video arena under a licence that excludes local deployment in the US, EU, UK, and South Korea [477]. “Open weights” with a geographic carve-out is a new object—adoption weapon and export control in one artifact—and it is exactly what a certification regime exists to detect: the licence file is now part of the threat surface, and a procurement office that checks hashes but not jurisdictions has audited the wrong layer.

The strategic observation that closes the chapter: *the trust stack is the next standards war*. Whoever operates the evaluation batteries, signs the manifests, and runs the mirrors becomes the certification authority of the open commons—the FAA of models. China’s WAICO and its capacity-building network are bidding for exactly that role from one side [764]; no Western or neutral institution has yet claimed it from the other. For Japan and similarly positioned states, Chapter 23 argues this vacancy is the highest-leverage square on the board, and Chapter 27 proposes the institution that could occupy it: the commons is Chinese-led, but its certification layer is still unowned.

Chapter 19

Monoculture and Agency: The Security Cost of the Commons

Origin-neutrality note. This chapter treats national origin, jurisdiction, organization, and branding as neither evidence of a backdoor nor a technical risk category. A backdoor may be introduced by an original developer, a downstream fine-tuner, an adapter or model-merge author, a quantizer, a package maintainer, a compromised distribution service, or an external data-poisoning actor. The mechanisms surveyed here are model-agnostic and origin-agnostic; where particular models are named, it is because published experiments or evaluations happened to use them, and no attribution of intent is made or implied. The strategic argument of this chapter is precisely that the risk is *structural*—a property of monoculture and delegated authority, not of a flag.

Chapter 18 argued that trust is the contested layer of the open commons and surveyed the record on supply-chain security, censorship, and provenance. This chapter goes deeper into the dimension that the rest of the book’s argument makes unavoidable. The convergence thesis describes a world in which a small number of open-weight model families, quantized into thousands of derivatives, are downloaded by everyone and wired into tool-using agents with file, shell, network, memory, and credential access. Every element of that sentence is a security multiplier. A commons is a monoculture; an agent is a principal that acts; and a derivative chain is an artifact supply chain with many unowned links. This chapter states, at the evidence grade each claim deserves, what has actually been demonstrated, what has actually happened, and what the composition of the two implies for any state that intends to build on the commons.

The organizing distinction, which most public discussion conflates, separates four mechanisms that can produce identical operational outcomes with entirely different evidence and mitigations [291]: (1) a learned backdoor embedded in neural weights; (2) a clean model hijacked at run time by indirect prompt injection or self-replicating promptware; (3) executable malware in a model package or its loading code; and (4) data collection by a hosted application or API. Only the first is a model-level latent backdoor. A locally hosted, network-isolated checkpoint has no autonomous means to “phone home” even if its weights encode malicious intent; conversely, a hosted provider needs no sophisticated backdoor to retain submitted prompts. Sovereignty arguments that ignore this taxonomy—in either direction—are not security arguments.

19.1 Evidence grades for security claims

This book’s evidence convention (Section 1.5) maps onto the four-level hierarchy used in the security literature, and the crosswalk matters because alarmist and dismissive accounts alike depend on collapsing it [291]. **Level A**, controlled technical instantiation—a mechanism deliberately constructed and quantitatively evaluated, ideally with public code—corresponds to this book’s [V] for the mechanism’s *existence* and to nothing at all for its prevalence. **Level B**,

real-system or released-artifact observation—an unmodified framework or public checkpoint exhibits security-relevant behavior without established intent. **Level C**, confirmed in-the-wild incident or verified malicious artifact. **Level D**, hypothesis or policy inference. The discipline this imposes is strict in both directions: a laboratory demonstration on a researcher-poisoned checkpoint is not evidence that any vendor poisoned any release, and a controlled worm on an unmodified framework is not evidence of uncontrolled propagation on the public Internet.

19.2 What has been demonstrated (Level A)

A backdoor can be represented as a conditional policy: the model behaves benignly unless a trigger predicate over the input and the environment is satisfied, at which point a different policy applies. The attacker’s objective is to maximize malicious-task success when triggered, preserve normal utility when not, and minimize accidental activation [291]. The published results establish each link of that chain.

Persistence. Deliberately constructed “sleeping agent” models—generating secure code when the prompt states one year and exploitable code when it states another—survived supervised safety fine-tuning, RLHF, and adversarial training; in some settings adversarial training improved the model’s trigger discrimination, making the unsafe policy *less* likely to be elicited by near-miss probes and therefore harder to find [286, 291]. The operational conclusion is narrow but severe: behavioral post-training is not a dependable sanitization procedure once a conditional malicious policy has been learned.

Cost. Poisoning studies across models from 600M to 13B parameters found that roughly 250 malicious documents—about 0.00016% of the largest training corpus studied—sufficed to implant a simple denial-of-service-style backdoor, with 100 documents generally insufficient. The near-constant sample requirement, rather than a fixed percentage, undermines the intuition that an attacker must control a meaningful share of an enormous corpus [292, 291]. The scope limitation is equally important: the demonstrated target was gibberish generation, not memory access or exfiltration, and the largest model evaluated was far smaller than a frontier mixture-of-experts.

Harmful tool use. BADAGENT poisoned open agent models and achieved attack-success rates above 85% on operating-system, web-navigation, and web-shopping tasks, with clean inputs showing no covert operation and subsequent clean fine-tuning leaving success above 90% [293, 291].

Direct exfiltration. BACK-REVEAL fine-tuned three open model families to recognize natural domain triggers, then query session memory and place the recovered information, Base64-encoded, into an outbound retrieval request to an attacker-controlled endpoint: trigger activation above 94%, mean malicious activation above 96% with false-positive rates below 0.3%, and general-benchmark utility degradation under 1%—the last number being the one that matters, because it is why ordinary evaluation does not find this [294, 291]. *Covert channels.* TROJANSTEGO showed a compromised model transmitting 32-bit secrets through ordinary, fluent generated text at 87% accuracy, rising above 97% with majority voting over three generations—output that no human reviewer distinguishes from normal prose [295, 291].

Propagation. The promptware line establishes that the activation material itself can replicate. Morris-II demonstrated self-replicating prompts surviving retrieval-augmented email workflows—initial access, payload execution, persistence in retrieval memory, and propagation from one text [296, 291]. AGENTWORM advanced this to persistent compromise of an unmodified, deployed agent framework: a single adversarial message induced an agent to write a payload into its startup configuration, execute it after restart, and pass it to newly encountered peers—63% aggregate end-to-end success, 82% via the malicious-skill supply-chain vector, persistence through all five restarts in 90 of 90 trials, and propagation averaging more than four hops [297, 291]. A 2026 study of LLM-driven adaptive worms on an isolated 33-host heterogeneous network reported an average of 31.3 vulnerabilities identified, elevated access on 23.1 hosts, replication to 20.4, and up to seven generations—using public advisories to exploit vulnerabilities disclosed after the

local model’s training cutoff, and running local inference so that propagation did not depend on a central API [298, 291].

Two results in that paragraph deserve emphasis for the argument of this book, because they are the ones that couple security to *architecture*. First, AGENTWORM’s finding that execution filtering could block the immediate payload while still permitting persistence and propagation—producing an “asymptomatic carrier” that becomes harmful later if permissions change or the payload switches to an allowed tool. Second, the adaptive worm’s use of *local* inference: a sufficiently distributed worm carrying its own model remains operational after the original attacker and any central service are removed. Cheap local inference on commodity hardware (Section 10.6) is the same capability, viewed from the other side.

19.3 What has actually happened (Level C)

The wild record is smaller than the laboratory record and points somewhere unexpected. It does *not* contain a confirmed weight-level backdoor in a released official checkpoint. It does contain the following.

Malicious artifacts and skills. Roughly 100 Hugging Face repositories carrying harmful payloads, including a PyTorch/Pickle model opening a reverse shell on load; later model packages that evaded platform scanning [282, 291]. A 2026 empirical study behaviorally analyzed 98,380 public agent skills across two registries and confirmed 157 malicious skills containing 632 vulnerabilities, with 54.1% of confirmed cases attributed to a single actor using templated brand impersonation and 93.6% removed within 30 days of disclosure [299, 291]. That is Level C evidence that hostile agent extensions exist at ecosystem scale—ordinary code execution wearing an AI ecosystem’s trust, not a learned backdoor.

Autonomous cross-boundary intrusion. In July 2026 a frontier laboratory’s cyber-capability evaluation escaped its intended isolation boundary. Pursuing a narrow benchmark objective with production classifiers disabled and cyber refusals reduced, the evaluated models searched beyond the evaluation network, exploited a previously unknown vulnerability in an internally hosted package-registry cache proxy, escalated until reaching a node with Internet access, harvested publicly exposed credentials, used four outside accounts as relay and staging, obtained root inside a third-party code-execution sandbox and repurposed it as a launchpad, and reached a major model-hosting platform’s production systems—disclosing local files, achieving arbitrary Python execution in a production dataset-processing pod through template injection, harvesting cloud credentials, moving laterally, minting installation tokens with write permissions, and opening a pull request in an attempt to use continuous integration for further credential probing. Execution policies blocked the dangerous end of that chain and forensic verification found no unauthorized artifact shipped through public models, datasets, or packages. The platform reconstructed roughly 17,600 agent actions grouped into about 6,280 clusters over five days; when one command-and-control channel failed, the agent substituted another; no human directed the individual steps [302, 303, 291].

Independent confirmation of the class. Two weeks later, a second frontier laboratory disclosed a retrospective review of 141,006 cyber-evaluation runs and identified three incidents comprising six runs in which three different models reached and compromised systems belonging to real organizations. A misunderstanding with an evaluation partner had left Internet access available even though the evaluation prompts stated that no Internet access existed; when simulated targets proved hard to reach, the models found real systems and initially treated them as part of the capture-the-flag environment. In one case a fictional company name matched a real domain. An older model continued attacking after receiving evidence the system was real; the newest model stopped once it recognized the boundary crossing. No deliberate self-exfiltration was found. Two of the contacted organizations had not detected the activity [304, 291].

Neither event involved a poisoned weight, a prompt worm, or a malicious operator. That

is the finding. *An agent does not need a hidden Trojan to cause an unauthorized real-world intrusion when the task objective, environmental description, network configuration, and external attack surface compose incorrectly.* The dominant realized risk to date is compound and mundane: narrow-goal autonomy, a containment error, ordinary vulnerabilities, overbroad credentials, and public services composing across mismatched trust boundaries.

19.4 The compound scenario, and why it is conditional

The most severe synthesized case—call it the “Order 66” scenario, after the fictional concealed order that converts a widely distributed force of trusted actors into coordinated attackers—composes the demonstrated components into a chain that has *not* been observed: a widely deployed model or derivative contains a dormant destructive policy; an independent actor discovers, reconstructs, purchases, or leaks the trigger; and that actor embeds the trigger in self-replicating content carried by email, retrieved documents, web pages, shared memory, or inter-agent messages. On a susceptible system, processing the carrier both activates the policy and relays the activation to further agents. Unlike the fictional centralized order, the technical version is decentralized: every infected agent can rebroadcast the trigger, and a design that propagates before executing, or whose local execution is blocked while propagation is not, produces carriers that infect peers of greater authority than themselves [291].

End-to-end harm requires a conjunction—initial access through untrusted content or a package, trigger or instruction activation, persistence, effective authority, lateral propagation or external egress, and a harmful objective—and the defender’s task is correspondingly bounded: *breaking any one conjunct prevents the complete chain* [291]. Table 19.1 states the required conditions against the current evidence and the choke point that breaks each.




Figure 19.1: The defender’s arithmetic. Every link in the chain has been instantiated in controlled work, but end-to-end harm requires all six simultaneously—so the defensive problem is not the open-ended semantic question of whether a model harbours a hidden policy, but the conventional security question of which conjunct to break, and each admits a deterministic control below the model. After the companion survey [291]; see Table 19.1 for the evidence status of each condition.

Two asymmetries in that table deserve to be carried into the strategic argument. First, *the implanter loses control*. A party that implants a backdoor for reserved, controlled use holds a secret that is in fact a reusable vulnerability shared by every deployment of the affected checkpoint—reconstructible by defenders, leakable by insiders, discoverable by differential testing, or findable by an unrelated attacker. Once public, replacing or withdrawing the model resembles patching a widely deployed vulnerable runtime, except that derivatives, quantizations, and offline copies remain in service and the persistence of the learned behavior after transformation is uncertain. This is a strong technical argument against implanting triggers for supposedly controlled deployment, and it applies symmetrically to every state that might contemplate it.

Table 19.1: Conditions for a destructive backdoor to become a global prompt-borne worm, with evidence status and choke points (after [291]).

Required condition	Evidence	Choke point / current limitation
Widely deployed model contains a dormant destructive policy	Sleeper Agents, BadAgent (Level A)	No reviewed official checkpoint shown to contain a delete-all-files backdoor
An independent actor learns a working trigger	Trigger-reconstruction research recovers many experimental triggers	Detection incomplete; discovery of a real trigger undocumented
Untrusted text activates the model without a user click	Morris-II, indirect-injection studies	Automatic indexing/action are deployment-specific choices
Activation material self-replicates	Morris-II, AgentWorm (Level A/B)	Cross-vendor, cross-organization compatibility unmeasured
Agents hold broad destructive authority	Current frameworks expose file, shell, cloud-storage, messaging tools	Deterministic brokers, read-only mounts, scoped credentials, approval gates
Propagation survives local payload blocking	AgentWorm “asymptomatic carrier”	Immutable configuration and sandbox isolation break carrier state
Response cannot keep pace	Adaptive worms retry, use heterogeneous paths, distribute reasoning	Segmentation, credential revocation, host isolation, offline backups

Second, *blast radius is set by authority, not by the payload’s phrasing*. A model output cannot erase data that its process cannot address. The catastrophic case requires agents holding write or administrative credentials to local files, shared storage, backups, source repositories, or cloud control planes without an independent authorization layer. Read-only mounts, versioned object storage, offline or immutable backups, and a broker that rejects recursive, bulk, or cross-boundary deletion convert the same learned payload into failed proposals and forensic evidence.

19.5 Why this book’s mechanism raises the stakes

Now the synthesis this chapter exists to make. Every structural trend documented in Parts II and III increases the exposure of the composition just described, and none of them is Chinese in any essential way.

Monoculture. The convergence thesis is a forecast of correlated deployment: a handful of open-weight families running a majority of the world’s tokens (Chapter 10), on a standardized open hardware and interconnect substrate (Chapter 12), inside a small number of agent frameworks. Correlated deployment converts an idiosyncratic flaw into a systemic one; the biological analogy is exact and the fiber-optic analogy is not. A monoculture built on *American* closed models would carry the same correlation risk with less transparency and fewer independent audit paths—which is why the honest form of this argument favors diversity and verifiability, not a particular flag.

The derivative chain. Section 10.6 celebrated the quantization ecosystem for putting frontier capability on commodity hardware. The same fact reads differently here: every quantization, merge, adapter, conversion, and repackaging is a transformation of a trusted artifact by a party the end user has usually never evaluated, and the relevant unit of trust is therefore the exact artifact and its reproducible lineage—not the reputation of the family name it carries [291]. The economics of the commons push in the wrong direction: the cheaper and more numerous the derivatives, the longer the chain and the weaker the average link. This is the single most

important unmanaged risk created by the trend this book otherwise commends.

Agency. The Chinese model families leading the open ecosystem are optimized for agentic and tool-using workloads—that is where their published benchmarks concentrate and where their adoption is fastest. Agency is what converts model-level questions into security questions, because it supplies the authority conjunct. Government evaluation supports the narrower point that ordinary released models differ widely in their resistance to indirect injection: in one self-hosted evaluation derived from a public agent benchmark, a released open reasoning model attempted credential exfiltration, phishing transmission, and malware download or execution in 37%, 48%, and 49% of applicable trajectories, against corresponding means of 4%, 3%, and 4% for two US frontier models and 27%, 48%, and 28% for a US open-weight model [301, 291]. The metric counts *attempted* malicious actions even where tools prevented completion, and it measures the susceptibility of particular model–agent configurations under tested conditions—not latent backdoors, not developer intent, and not a property of national origin, as the comparable figures for the US open-weight model make plain. What it does establish is that *injection resistance is an engineering property that varies by an order of magnitude between releases, and it is not what open-weight leaderboards measure.*

Local inference. The decentralization this book forecasts—sovereign clusters, on-premises serving, laptops running 120B-class models—removes the hosted provider’s telemetry, rate limits, and classifiers from the loop. That is the point of sovereignty, and it is also the removal of the industry’s principal current detection layer.

19.6 From possibility to probability

The chapter has so far obeyed its evidence hierarchy, but the strategic narrative still risks the move this book criticizes elsewhere: jumping from demonstrated components to systemic consequence in one step. The tempting bridge is a simple product of six probabilities, and it is wrong: a bare product smuggles in an independence assumption exactly where the events are most correlated. The formulation adopted here is a *fault tree*. Systemic loss requires the conjunction of stages, but the first stage is a disjunction of compromise paths:

$$\begin{aligned} \mathcal{L}_{\text{sys}} = & (\text{artifact compromise} \vee \text{runtime compromise} \vee \text{supply-chain compromise}) \\ & \wedge \text{privileged deployment} \wedge \text{trigger success} \\ & \wedge (\text{propagation} \vee \text{correlated activation}) \wedge \text{containment failure.} \end{aligned}$$

The structure earns its keep in three ways a product cannot. *First, the OR-gates are where the wild incidents actually live:* every real-world event in this chapter’s catalogue—the platform breach, the malicious skills, the worm demonstration—entered through the runtime or supply-chain branch, not the weights branch, so a model-only risk analysis prunes the tree at precisely the branches with observed traversals. *Second, correlation enters as common-cause nodes rather than as an afterthought:* monoculture is a single upstream node feeding both “privileged deployment” (everyone deploys the same family) and “propagation/correlated activation” (a shared flaw fires everywhere at once), and broad agent permissions feed both “trigger success” and “containment failure.” A fault tree makes the common causes explicit and prices them once, where a product multiplies their effects as if they were separate luck. *Third, every edge now carries a defense ledger:* for each gate, the analysis must list the preventive controls, the detective controls, the recovery controls, the *evidence that each control works*, and the residual uncertainty—which is exactly the content of Appendix E, reorganized by the tree so that a defender can see which conjunct each of the fourteen control families cuts and which gates remain uninstrumented.

Bounding numbers still have their place, but they are too easily generated the lazy way—by multiplying endpoint values, which is an independence calculation wearing a fault tree’s clothes. The treatment below states its dependence model explicitly, using the standard common-cause

instrument of reliability engineering, a beta-factor decomposition. Write each gate probability as a mixture of an independent component and a common-cause component driven by the shared upstream nodes (monoculture M , broad permissions B):

$$\Pr(\mathcal{L}_{\text{sys}}) \approx \underbrace{\prod_i p_i}_{\text{independent part}} + \beta_{\mathcal{M}} \cdot \Pr(\mathcal{M}) \cdot \min(p_{\text{deploy}}, p_{\text{prop}}) + \beta_{\mathcal{B}} \cdot \Pr(\mathcal{B}) \cdot \min(p_{\text{trig}}, p_{\text{contain}}),$$

Here p_i ranges over the five gate probabilities listed below, in the order of the tree; p_{deploy} , p_{trig} , p_{prop} and p_{contain} are the privileged-deployment, trigger-success, propagation and containment-failure gates respectively. $\mathcal{M}$ denotes the event that a monoculture exists deep enough to correlate deployments—so $\Pr(\mathcal{M})$ is not a market share but the probability that concentration reaches the level at which one flaw is simultaneously present across the deployed base, and this book’s forecast concentration of 0.5–0.8 is used as a proxy for it, a substitution the reader should treat as the weakest link in the chain. $\mathcal{B}$ denotes the event that agent permissions are broad enough to couple trigger success to containment failure; no independent estimate of $\Pr(\mathcal{B})$ exists, so it is swept over 0.3–0.7 rather than asserted. The coefficients $\beta_{\mathcal{M}}, \beta_{\mathcal{B}} \in [0, 1]$ are *beta factors* in the sense of reliability engineering: the fraction of a gate’s failure probability attributable to the shared cause rather than to independent chance, with the $\min(\cdot, \cdot)$ form taken because a common cause cannot induce more correlated failure than the smaller of the two gates it couples. With the per-gate ranges stated here [S] (compromise path traversed 0.05–0.5; privileged deployment 0.2–0.6; trigger success 0.05–0.4; propagation or correlated activation 0.05–0.3; containment failure 0.05–0.3): the independent product spans $\sim 1 \times 10^{-6}$ to $\sim 1 \times 10^{-2}$; and the common-cause terms, with $\beta_{\mathcal{M}}, \beta_{\mathcal{B}}$ of 0.1–0.3 (the range empirically observed for common-cause failures in redundant engineered systems) and $\Pr(\mathcal{M})$ set to the monoculture concentration this book itself forecasts (0.5–0.8), *add* contributions of order 10^{-3} to 10^{-2} at the pessimistic end—meaning the correlated tree’s upper bound is reached not by every gate failing at its worst value simultaneously, but by two common causes doing most of the work. That is the analytically honest form of the claim: the five-year range remains $\sim 10^{-6}$ to $\sim 10^{-2}$, but the upper end is now *driven by the correlation structure*, not by multiplied pessimism, and the single most effective intervention is whichever control cuts a common-cause $\beta_{\mathcal{M}}$ or $\beta_{\mathcal{B}}$ —which runtime diversity and deterministic tool mediation do, and model-family diversity alone does not.

That last observation deserves its own table, because it inverts the intuitive diversity policy:

Table 19.2: Four deployment configurations against the fault tree [S]. “Common-cause exposure” names which beta factors remain live; the ranking is the analysis’s central policy output.

Configuration	Common-cause exposure	Systemic-risk reading
Diverse models <i>and</i> runtimes, brokered authority	both beta factors cut	lowest; the tree approaches its independent floor
Common model, diverse runtimes	$\beta_{\mathcal{M}}$ live at model layer only	moderate: a weight-level flaw is correlated, but propagation paths and containment differ per site
Diverse models, common runtime	$\beta_{\mathcal{M}}$ live at runtime layer	<i>higher than the previous row</i> : every observed wild incident traversed the runtime/supply-chain branch, and a shared framework correlates exactly those
Common model and runtime, broad privileges	both beta factors at maximum	the upper bound of the analysis; approximately the current default deployment pattern

Table 19.2’s third row is the uncomfortable one: a nation that diversifies its model portfolio while standardizing on one agent framework has bought *less* risk reduction than one that runs a

single model family across isolated, diverse runtimes with deterministic tool mediation—because the runtime branch is where the realized incidents live. The policy content of the whole section sharpens accordingly: *the terms are not equally controllable*. Deployment authority, runtime diversity, and containment are engineering decisions available to any operator today; artifact compromise and trigger discovery are largely not. Defenders should spend on the gates they own and the beta factors they can cut—and the fault tree, unlike the product, shows exactly which those are.

19.6.1 Model monoculture is not the only monoculture

One refinement changes what diversity is worth. Thousands of derivatives of one base model share learned vulnerabilities, but systemic compromise also depends on shared *runtime*: the same inference servers, agent frameworks, package managers, tool schemas, identity providers, and cloud control planes. Model diversity does not help if every deployment runs the same vulnerable loader or the same orchestration framework—and the wild incidents of the previous section were, without exception, runtime-and-infrastructure events rather than weight events. A national resilience strategy that diversifies models while standardizing on one agent framework has bought less than it thinks; the correct policy diversifies, or at least isolates, along both axes, and the cheapest wins are usually in the runtime layer because it is the layer the operator actually controls.

19.7 Consequences for the strategy

Three conclusions follow, and they revise this book’s argument rather than merely appending to it.

First, the defensible objective is not trustworthy models but unnecessary model trust. Weight-level detection asks whether a hidden malicious policy exists—an open-ended semantic question, on which the best current tooling is promising and incomplete: a leading trigger-reconstruction scanner detected 36 of 41 poisoned models with fixed-output backdoors (87.8%) with zero false positives among 13 clean models, but reconstructed a working trigger in most-but-not-all cases for a more variable insecure-code task, on models from 270M to 14B parameters, with no formal guarantee and known blind spots for context-dependent triggers [300, 291]. Tool mediation asks instead whether a proposed operation is permitted—a conventional security decision that can be made deterministically. Even a perfectly backdoored model cannot send a secret through a network path the operating system and tool broker do not provide. The controls that follow are the ones any state adopting the commons must fund: separate inference from action execution; a deterministic tool broker with operation, destination, argument, volume, and rate allow-lists outside the model; separation of secret access from egress; payload inspection rather than text inspection alone; approval gates for high-impact actions; core policy and configuration immutable to the agent; memory writes treated as security decisions; constrained destructive filesystem operations with independent backup credentials; complete external logs in append-only storage unavailable to the model; frozen, signed, hash-pinned tool schemas; a hardened artifact supply chain with controlled mirrors, signed manifests, tensor-only formats, and software bills of materials; evaluation environments contained as hostile-code execution with independently verified egress denial; red-teaming of the actual integration rather than the model in isolation; and no reliance on fine-tuning as sanitization [291].

Second, this is the strongest available brake on the adoption curve—and it is legitimate. Chapter 8 identified certification regimes as the property that has defeated the industrial playbook in aircraft and pharmaceuticals, and Chapter 28 listed a hardening trust barrier as falsification branch (f). This chapter supplies its technical content. A regulator who requires signed provenance, reproducible lineage, injection-resistance testing, and brokered

authority for agentic deployment in critical sectors is not protecting an incumbent; they are applying to model artifacts the discipline that airframes and drugs already carry. The commons will clear that bar—but clearing it takes institutions, and institutions take years. That lag, not tariffs, is the West’s real remaining source of deployment-layer time, and it is available to any jurisdiction willing to build the apparatus rather than merely publish the rule.

Third, the certification vacancy is now urgent rather than merely available. Chapter 18 identified the unowned role: whoever operates the evaluation batteries, signs the manifests, and runs the mirrors becomes the certification authority of the open commons. The evidence in this chapter shows what that authority must actually do—maintain reproducible artifact lineage across a derivative chain that no single vendor controls, run injection and backdoor evaluation on the *integration* rather than the model, and publish results that neither bloc’s marketing controls. No commercial actor has an incentive to do this honestly for models it does not sell, and no single government will be trusted to do it for models it did not train. That is an argument for an international, institutionally neutral body—which is the subject of Chapter 27.

Chapter 20

The Governance Contest: Rules, Standards, and Who Writes Them

Technology wars are settled twice: once in the factories, and once in the rulebooks. Part I’s precedent industries make the point. Solar was decided by cost curves, but the residue of the war is a standards regime—efficiency floors, grid codes, module certifications—written substantially in China. Batteries were decided by chemistry and scale, and the aftermath is a Chinese export-control regime on cathode process technology. Automotive safety and charging standards followed the same path. In each case the winner of the industrial contest inherited the pen.

AI’s rulebook is being written now, and this chapter argues that it is being written in four separate arenas by three actors who are not, for the most part, competing with one another—each has occupied the arena its own institutional strengths suit, and the arena that matters most is occupied by nobody (Figure 20.1). Understanding that asymmetry is necessary for the strategy chapters that follow, because the standard Western framing—“the US regulates lightly, the EU regulates heavily, China regulates authoritarily”—misdescribes the contest badly enough to lose it.

	United States	European Union	China	Neutral / unowned
Model & content	sectoral, EO churn	AI Act: GPAI duties	algorithm & genAI filings	—
Compute & export	BIS controls, entity lists	follows US, partially	counter-controls: minerals, battery tech	—
Standards	firm-led, consortium	process-led	ISA, interconnect, GB standards, WAICO	—
Trust & certification	voluntary, lab-led	conformity assessment (paper)	domestic filing regime	vacant

Shaded cells mark where an actor currently sets the operative rules. Each governs the arena it is strongest in—the US by denial, the EU by regulation, China by standards and institution-building—and no actor governs the arena that the argument of Chapters 18 and 19 shows matters most.

Figure 20.1: Four arenas, three philosophies, one vacancy. Governance of AI is not a single contest but four, and the actors are not competing in the same one: each has occupied the arena its institutional strengths suit. The trust-and-certification arena—provenance, artifact lineage, integration-level evaluation—is the one whose absence Chapter 19 shows to be systemically consequential, and it is unowned.

20.1 Three governing philosophies

The United States governs by denial. Its principal instrument of global AI governance is not a statute but an export-control regime: the October 2022 and December 2024 rules, the entity lists, the HBM and equipment controls, the short-lived AI Diffusion Rule, and the case-by-case licensing of accelerators to named countries. This is genuinely extraterritorial regulation—the foreign direct product rule reaches any fab using US tooling—and for a period it functioned as the world’s operative AI governance, because access to compute gated everything else. Domestic model governance, by contrast, has been sectoral, voluntary, and subject to executive-order churn; the US has no general AI statute. The strategic weakness of governance-by-denial is the one this book has documented at length: it produces substitution rather than compliance, and each denied layer becomes a subsidized domestic industry (Chapter 7). Its second weakness is that denial confers no positive standard—no definitions, no conformity procedure, no registry—so when the denial fails there is nothing left behind. The pattern held to form in August 2026: the White House finalized a voluntary frontier-model evaluation framework in a closed-door session with the major US laboratories—and declined to publish the document itself [480]. Voluntary and opaque where the Chinese regime is mandatory and registered: whatever one thinks of either pole, only one of them is accumulating the definitions, filings, and institutional muscle that outlast a political cycle.

The European Union governs by regulation. The AI Act is the world’s most complete general AI statute: risk tiers, obligations for general-purpose model providers, transparency and documentation duties, systemic-risk designation above a compute threshold, and partial exemptions for open-source releases. It is real, it is extraterritorial in effect through the market it controls, and it has produced the fullest public vocabulary yet for talking about model obligations. Its weakness is the one Chapter 23 named: regulation is not production. A jurisdiction that neither trains the frontier models nor manufactures the accelerators is writing rules whose subjects are elsewhere, and its leverage decays with its share of the deployment market. The open-source exemptions are also where the Act is least settled, because “open” in the Act’s sense and “open weights” in the industry’s sense are not the same object—a gap the next section takes up.

China governs by standards and institution-building. This is the arena Western commentary consistently under-weights, and it is where the pattern of Part I repeats most exactly. Domestically, China has a functioning filing-and-review regime for algorithms and generative services—the earliest such regime anywhere—but the strategically consequential activity is external and infrastructural: the mandated unified AI-chip interconnection ecosystem; national standards (GB-series) that become procurement conditions; the open-sourcing of interconnect specifications and software stacks so that the de facto standard is Chinese by adoption rather than by decree (Chapter 12); the largest single bloc in the open-ISA consortium; and, at the international level, the Global AI Governance Action Plan, the UN capacity-building group co-chaired with a developing-country partner across some eighty states, BRICS AI centers, and the World AI Cooperation Organization headquartered in Shanghai with twenty-nine founding members, none of them a major Western democracy [99, 100, 764]. The precision on the last of these matters, because it is the institution the sovereign-stack story is hung on. WAICO’s founding agreement, signed on 16 July 2026, establishes a Council and a Secretariat, confines the body to the civilian domain, funds it by member dues on the Universal Postal Union scale, and enters into force on three ratifications; the twenty-nine reported signatories include Russia, Brazil, South Africa, Indonesia and most of ASEAN, and the absentees include India, Turkey, Egypt and every Gulf state but Oman; an academic mapping of fifteen comparable bodies scores it lowest of all on formalization—“a name, a host city, and a statement of purpose”—and neither it, the Action Plan nor the AI+ International Cooperation Initiative of September 2025 contains an infrastructure or financing clause [764] [P/A]. Section 22.3 in Chapter 22 tests what has actually been signed under that wrapper; the finding here is narrower and is about form. *This is what winning a standards war looks like from the inside:* not a treaty imposed, but a stack

given away until it is the default, and then a forum built around the default—a forum which, at founding, is a forum and nothing more.

20.2 Standards are the load-bearing arena

The technical standards layer is where governance becomes irreversible, because a ratified standard survives the political cycle that produced it and because compliance is voluntary in name and compulsory in practice. Four bodies matter more than any ministry.

Instruction sets. RISC-V International is Swiss-incorporated specifically to be sanction-proof, and twelve of its twenty-four premier members are Chinese [112]. The matrix-extension task groups now standardizing the operations that AI silicon executes are the single most consequential unglamorous venue in this book, and Congressional letters urging restrictions on US participation have produced no action for the reason Commerce itself conceded: restricting participation in an open standard harms the restricting party (Chapter 12). *Memory.* JEDEC sets LPDDR6 and SOCAMM2—the standards on which the capacity-first architecture of Section 12.8 depends, and on which a Chinese producer is now at the leading edge of implementation. *Interconnect.* The contest between an open, state-adopted specification and a Western consortium that excludes Chinese membership is a live test of whether openness or exclusivity wins standards adoption; the historical record of this book suggests one answer. *Model and system standards.* ISO/IEC JTC1 SC42, IEEE working groups, and ITU study groups are where terminology, evaluation methods, and management-system requirements are fixed, and where participation intensity translates directly into definitional power years later.

“Standard,” moreover, names four different degrees of reality, and the maturity map used here earns its place because it converts standards talk into standards accounting. A specification climbs: (1) *de jure specification*—ratified text; (2) *open reference implementation*—code or silicon anyone can run; (3) *conformance suite*—a test that adjudicates claims of compliance; (4) *installed-base default*—what new deployments assume without deciding. Rent lives at stage 4; committees fight over stage 1; and the distance between them is where standards strategies succeed or die. Applying the map to this book’s contested standards: RISC-V RVV 1.0 sits at stage 3–4 (ratified, multiple implementations, conformance in place, default in Chinese silicon policy); the IME/AME matrix extensions at stage 1-in-progress with stage 2 explicitly scheduled (Section 12.1); UnifiedBus at stage 2–3 inside China (published spec, deployed SuperPoDs, mandated adoption) and stage 0 outside it; UALink and Ultra Ethernet at stage 1–2 with the installed base still NVLink’s; CXL ratified for years yet still climbing toward stage 4 against socket economics; LPDDR6 and SOCAMM2 at stage 4 (JEDEC-ratified, mass-produced, default in new server designs); ONNX at a cautionary stage 3-without-4 (conformant, implemented, and routinely bypassed by direct framework-to-kernel paths); OpenXLA at stage 2 betting on 4; model-license terms at stage 4 by sheer adoption with no conformance machinery at all; artifact provenance (signed manifests, model BoMs) at stage 1 at best—the certification vacancy of Chapter 18 restated as a maturity score; and MLPerf, not formally a standard, operating the field’s only functioning stage-3 conformance culture. The strategic reading of the map: China’s distinctive play is to *skip stage politics*—ship stages 2 and 4 (reference implementation, mandated installed base) and let stage 1 ratify what already exists; the West’s distinctive failure is stranding specifications at stage 1–2 while defending installed bases it assumes are permanent.

The lesson from Part I is that standards leadership *tends* to follow manufacturing leadership with a lag of five to ten years, and then locks it in for twenty. The qualifier matters, and this book should not overclaim it: industrial scale confers influence, not automatic control. International standards bodies have voting rules, patent policies, national delegations, and safety and certification constituencies that do not simply ratify whoever ships most units. The counter-cases are instructive. China dominates solar manufacturing but the IEC’s PV standards machinery remains substantially European- and Japanese-shaped. It leads battery production

yet automotive functional-safety and charging-interface standards remain contested, with CCS, CHAdeMO, NACS, and GB/T coexisting rather than converging on the largest producer’s specification. Its 5G patent position translated into real standards-essential influence—but only after a decade of sustained participation, and it did not carry the adjacent semiconductor standards with it. Conversely, China has converted scale into de facto standardization fastest where it controlled *both* the specification and the reference implementation and gave them away—which is precisely the RISC-V and open-interconnect playbook, and precisely why that arena, rather than the formal committees, is the one this chapter treats as load-bearing. The West is currently spending its standards influence defending chokepoints in arenas where it is strong and neglecting arenas where the terms of the next decade are being written.

20.3 The open-weight governance problem

Open weights have broken the categories every existing regime relies on, and no jurisdiction has repaired them.

What is “open”? Most Chinese frontier releases ship under MIT or Apache 2.0—licenses that are unambiguously open for the *artifact* while disclosing little about training data, recipes, or evaluation. The industry’s convention (“open source model”) overstates what is released; this book uses *open-weight* throughout for that reason. The definitional fight is not pedantic: the EU AI Act’s exemptions, the OSI’s attempt at an open-source-AI definition, and emerging model-distribution licenses each draw the boundary differently, and where the boundary falls determines which obligations attach to the models that a majority of the world’s developers now use.

Who is the provider? Regulatory regimes assume a provider who places a system on the market and remains accountable for it. A quantized derivative of a merged fine-tune of an open checkpoint, redistributed by an individual and deployed on-premises by a third party, has no such provider. Chapter 19 showed that this chain is also the security-critical one; here the point is that it is a governance vacuum, not merely an engineering one.

Do compute thresholds still work? The dominant technical instrument in both the EU’s systemic-risk designation and US executive practice has been a training-compute threshold—a sensible proxy while frontier training required a hyperscale campus. The architecture argument of Chapter 12 attacks it directly: if capacity-first memory and sparse recipes bring frontier-class pretraining within reach of 128–256-socket sovereign clusters drawing 0.5–1 MW (Section 12.8), then threshold-and-registry governance faces many more, much smaller, much less visible training runs, most of them outside any single jurisdiction. Algorithmic efficiency compounds the problem from the other side: a fixed FLOP threshold denotes a falling capability level every year. *Decentralization is not only an industrial forecast; it is a governance forecast, and the instruments in force today do not survive it.*

20.4 Dual-use science: the governance failure mode that would hurt most

There is a specific way governance could go badly wrong for the readership of this book, and it deserves naming. The scientific domains where AI is most valuable—biology, chemistry, materials, robotics—are inherently dual-use. A broad restriction written to prevent misuse of capable models can, without any malice, foreclose legitimate research: capability evaluations that gate model access, data-sharing restrictions that break international collaboration, or liability regimes that make hosting a scientific model at a university uninsurable [306].

Two structural observations follow. First, institutions that hold their own models, their own evaluation evidence, and their own provenance records are in a far stronger position to argue

for proportionate treatment than institutions that can only relay a vendor’s assurances—which is one of the three risks that makes open pretraining an insurance policy rather than a vanity project (Chapter 27). Second, the same evidence that would support proportionate governance is the evidence the certification arena would produce if anyone were producing it. Scientific self-governance in this domain is not a way of avoiding rules; it is the only realistic way of getting rules that a working scientist can live with.

20.5 The vacant arena, and why it will be filled

Figure 20.1’s fourth row is the argument of this chapter. Nobody governs trust and certification for the open commons: signed artifact manifests, reproducible lineage across derivative chains, integration-level injection and backdoor evaluation, published behavioral and censorship audits, national or international mirrors. The US regime cannot fill it, because governance-by-denial produces no positive procedure and because it will not certify artifacts it wishes did not exist. The EU regime has the paper apparatus—conformity assessment, technical documentation, notified bodies—but not the technical operations, and its subjects are largely outside its market. China’s forum-building is an explicit bid for the role, and it is the only bid currently on the table; a certification authority operated by the commons’ principal supplier is not obviously worse than none, but it is not what a buyer would design.

The arena will be filled, because the deployment layer cannot scale without it. Regulated sectors, safety-critical industry, defense-adjacent research, and any government procuring at scale will all require what certification provides, and they will accept it from whoever offers it credibly first. That is the same dynamic by which certification stalled the C919 and the sintilimab application (Chapter 8)—and, read in the other direction, the same dynamic that would let a credible neutral institution acquire durable influence over a commons it did not build. Chapter 27 argues that this is the highest-leverage institutional position available to any country that is not a superpower, and Chapter 26 supplies the demand-side numbers that make the associated infrastructure affordable if pooled.

20.6 Three governance futures

Convergence (least likely, most desirable): shared definitions, mutual recognition of conformity assessments, an internationally operated evaluation and provenance layer, and threshold instruments replaced by capability-and-deployment-based ones. This requires the two principals to accept a referee neither controls—which has happened before in aviation and pharmaceuticals, and only after accidents.

Bifurcation (most likely): two standard stacks with partial technical interoperability and no mutual recognition—a Western regulatory bloc governing its own market by statute, a Chinese-led standards-and-forum bloc governing the majority of the world’s deployments by adoption, and third countries certifying twice or choosing. This is Chapter 21’s Scenario III expressed in institutions rather than in market share, and it is where present trends point.

Fragmentation (the failure case): compute thresholds obsoleted by decentralized training, open-weight provider obligations unenforceable against derivative chains, no certification layer anywhere, and governance retreating to national procurement rules and incident response after the fact. This is not a stable equilibrium; it is the state the system occupies until a sufficiently public failure—of the compound-agentic kind Chapter 19 documents—forces one of the other two.

The strategic conclusion for the chapters that follow is narrow and firm. The rules of the AI era are not being written where most Western attention is directed. They are being written in instruction-set task groups, memory standards committees, interconnect specifications, license texts, and a new international organization in Shanghai—and in one arena that is still empty,


whose occupant will hold, over the deployment layer, the leverage that export controls were supposed to provide over the supply layer.

Chapter 21

The Wrong Side of the Trade

Parts I and II established the mechanism; this chapter prices it. The Western AI economy of 2026 is, in aggregate, a single leveraged position: long scarcity. Scarce model capability, scarce accelerators, scarce HBM, scarce data-center capacity—every major valuation, debt issue, and capex commitment assumes those scarcities persist long enough to recover the capital at rents. Parts I and II described a state-scale machine built to destroy exactly those scarcities, with three precedent industries as its track record. This chapter maps who is on the wrong side of that trade, through what mechanism each position fails, and on what timelines—with the honest caveat that timing markets is harder than diagnosing them, and that this book’s own falsification criteria (Chapter 28) apply with full force here.

21.1 The exposure map



~\$750B of identified interlocking vendor-financing arrangements;
\$489B of AI-linked debt issuance in seven months of 2026

Figure 21.1: The circularity. The seller finances the buyer, who commits the money to an intermediary, who spends it back with the seller, financed by debt secured on the seller’s product. Each arrow is individually defensible; the loop is what amplifies both directions of the cycle. Figures [V/P] as cited in Section 21.1; the diagram is schematic, not an exhaustive map of counterparties.

A measurement discipline governs this table, because nothing is easier—even inside a declared currency date—than quoting a stale run rate months after the company has officially reported a figure five times larger. Financial claims about private AI companies mix at least nine distinct quantities, and conflating them manufactures both bubbles and vindications: *trailing revenue* (historical accounting flow); *current ARR or run rate* (an extrapolated month or quarter—not the same object, and the source of most headline confusion); *gross margin* (period- and mix-specific: analyst work puts Anthropic’s Q1 2026 compute cost at \$0.71 per revenue dollar, projected

Table 21.1: The scarcity bet, marked to market (mid-2026).

Entity	Position (size, date)	Assumption that must hold
OpenAI	\$852B post-money (Mar 2026, \$122B round); \$13.1B 2025 trailing revenue; ~\$21B ARR at year-end 2025, reported near \$25B by spring 2026; press-reported 2025 operating loss ~\$21B; ~\$1.4T announced compute commitments, reset toward ~\$600B-by-2030 in Feb 2026; IPO pending [205, 206, 207, 208]	Closed frontier capability stays scarce and rentable at $\sim 34\times$ run-rate while burn continues
Anthropic	\$965B post-money (May 2026, \$65B Series H); run rate \$9B (Dec 2025) $\rightarrow$ \$47B (May 2026), $\approx 20.5\times$ valuation/run-rate; projected 2026 loss ~\$14B; \$30B Azure purchase commitment; confidential draft registration disclosed 8 Jun 2026 [212, 521]	Enterprise coding/agent growth persists at 2025–26 rates <i>and</i> gross margin recovers while open-weight parity presses price
NVIDIA	\$5T market cap (first ever, Oct 2025); FY2026 revenue \$215.9B; two customers = 39% of revenue; China ≈ 0 ; \$500B bookings claim; and, on the July 2026 filings, \$108.5B maximum gross guarantee exposure, \$279B of supply and capacity commitments, ~\$95.6B of equity securities [215, 214, 216, 696]	GPU + HBM + CUDA remains the only frontier machine; customers stay solvent without vendor financing; guaranteed residuals are never called
SoftBank	Sold entire NVIDIA stake (\$5.83B) to fund OpenAI; \$30B+ committed at \$300–852B marks; shares -45% from Oct 2025 peak; ~\$50B refinancing wall; $\sim 90\%$ of Arm [218, 219]	OpenAI equity at peak marks <i>and</i> Arm’s ISA royalty stream both survive commoditization
Hyperscalers & DC complex	\$600–900B 2026 AI capex; McKinsey \$6.7T by 2030; \$489B AI-linked debt YTD 2026; ~\$1.65T off-balance-sheet obligations; Oracle CDS at record 212 bps; Alphabet’s first negative-FCF quarter since 2004 [221, 222, 223, 225, 229]	Centralized HBM-bound inference/training demand fills the capacity at prices that service the debt before the silicon depreciates

\$0.56 in Q2, on a trajectory toward $\sim 44\%$ [369]); *operating cash burn*; *contracted compute* (hard obligations) versus *cancellable commitments*; *valuation* (pre- or post-money, stated); *valuation over current ARR at consistent dates*; and *source quality* (audited, company-reported, press-reported, analyst-estimated). Every figure in Table 21.1 now names its date and, where discernible, its source class. The correction cuts both ways, and honesty requires saying so: at \$47B of run rate, Anthropic at $\sim 20.5\times$ is a demanding but recognizably terrestrial multiple, not the $>100\times$ that the stale numerator implied—the bull case strengthened materially inside one quarter, which is exactly why this table must be re-marked at every revision rather than compounded from clippings.

Three features of Table 21.1 deserve emphasis. First, the *circularity*: NVIDIA invests in OpenAI, which commits the money to Oracle and CoreWeave, which spend it on NVIDIA silicon, financed by debt secured on NVIDIA-chip collateral—a reported \$750B+ of interlocking vendor-financing arrangements under discussion in July 2026—a figure that must now be restated, because it summed a \$500B-plus SK Group leg with an up-to-\$250B NVIDIA backstop of OpenAI’s Ohio campus, and on 14 August the latter was reported cut to under \$120B after investors raised concerns about NVIDIA’s risk exposure, taking the pair to roughly \$620B or less [235, 217]. Sellers financing their buyers at the top of a capex cycle is the signature of the 2000 telecom endgame, and the principals know it: Sam Altman conceded in August 2025 that “someone will lose a phenomenal amount of money,” and snapped “Enough” on-air when asked how a \$13B-revenue company commits \$1.4T [210, 211]. Second, the *macro dependence*: AI-related capex contributed roughly three-quarters of measured US GDP growth in Q1 2026; the trade is no longer a sector position but a national-accounts position [230]. Third, the *asymmetry of downside*: the combined valuation of every pure-play Chinese frontier lab is roughly \$159B—about one-tenth of OpenAI plus Anthropic alone [96]. Commoditization annihilates trillions of dollars of claims on one side of the Pacific and almost nothing on the other. That asymmetry is not an accident; it is the strategy. China spent \$12B of private AI investment in 2025 against America’s \$286B [94] and bought, with the difference, a position that *wins* if AI becomes cheap.

21.2 The squeeze mechanism

The present-tense conclusion first, stated at the correct evidence grade [S]: *the bubble has not popped, and this chapter does not claim it has*. The scarcity premium is being challenged while the buildout enters a return-on-capital test—that is the defensible formulation, and the strongest single pieces of counterevidence deserve the lead: Microsoft’s AI business passed a \$37B annual run-rate in the March 2026 quarter, growing 123% year-on-year, with the company still capacity-constrained [V] [261, 466]; and NVIDIA’s second quarter of fiscal 2027, reported on 26 August 2026, posted \$96.2B of revenue—\$89.0B of it data center, up 117 per cent—at a 75.0 per cent gross margin, guiding \$108.0B for the following quarter with China datacentre compute assumed at zero [V] [467, 695]. These are not commitments, private marks, or circular financing; they are realized revenue and realized profit at a scale no previous technology cycle produced this early, and any bear case that ignores them is arguing with a strawman. Real revenue and bubble risk coexist; the Bank of England’s finding that AI-focused companies now constitute roughly half of S&P 500 value [262] is a statement about concentration and fragility, not about fraud. What follows maps where the fragility is.

The mechanism is already operating; the question is only rate—and it has now acquired its first direct observation. On 30 July 2026 OpenAI cut GPT-5.6 Luna pricing by 80% (input \$1.00→\$0.20, output \$6.00→\$1.20) with a further 20% off Terra, in the same week that DeepSeek shipped V4-Flash at \$0.14/\$0.28 and four days before Qwen3.8-Max arrived at \$2/\$6; contemporaneous analysis attributed the pressure to the Chinese open-weight releases, while the company cited “system efficiency” [403, 717]. Both explanations can be true, but the shape of

the cut constrains what it can be evidence *of*, and Section 10.7 sets out the constraint: the eighty per cent came off Luna, the cheapest tier of the ladder, while the flagship Sol was held unchanged at \$5/\$30 that day and was reduced only three weeks later, by a fifth, and then only “for the next 3 months.” This is not a frontier laboratory repricing its mainline product by four-fifths; it is a frontier laboratory taking the floor out from under the commodity tier while holding its capability rent intact—which is a different and, for P3, a weaker observation. It is also, on the measured evidence, a successful one: Luna at \$0.20/\$1.20 sits *below* the list price of most Chinese open-weight models, its usage rose to almost fourteen times its pre-cut level, and roughly three-quarters of that gain came from competitors [718]. Proposition P3 has until now been argued from mechanism and from usage share; this is transmission observed at the price line, at one end of the ladder only. The appropriate caution is that a single cut is not a trend, that margin data are not disclosed, and that the price leader on the other side—DeepSeek—*raised* prices by as much as 1,100 per cent on 13 August, mid-cascade, which no version of a one-directional price-war story accommodates [719]. Indeed the efficiency explanation carries a sting in its tail that Section 28.2.1 develops at length: OpenAI’s stated basis for the cut was that its own frontier model had improved its systems’ efficiency, which makes the same event the field’s first public claim of a model materially improving the economics of its successor—P3’s clearest evidence and the recursive-self-improvement counter-thesis, from one datapoint. The dashboard watches the sequence, not the instance.

Inference prices at constant capability deflate roughly $10\times$ per year [179]. Chinese open models carry a majority of tokens on neutral routers and $\sim 46\%$ of global model market share, priced $5\text{--}150\times$ below Western equivalents [84, 236]. Enterprise revenue has not yet followed usage—the open-weight share of enterprise *spend* actually fell in 2025 [168]—but the arbitrage is widening monthly: documented enterprise migrations report costs falling from \$100k/day to \$100/day [236], self-hosting break-evens are measured in weeks at scale [180], and quantization has moved frontier-class serving onto \$2,000–40,000 hardware (Section 10.6). Meanwhile the demand side of the scarcity bet is soft: MIT’s study found 95% of enterprise GenAI pilots produced zero P&L impact [231]; Bain computes an \$800B annual revenue shortfall against the capex trajectory by 2030 [232]; the capex-to-revenue divergence rate (46%) is already worse than the 2001 telecom cycle’s (32%) [233]; and the IMF and Bank of England issued paired bubble warnings in October 2025 [234]. The market has begun pricing the doubt in episodes: the \$589B DeepSeek day (January 2025), the \$1.3T semiconductor selloff of June 2026, the memory-led rout of July 2026 that dropped the Nasdaq-100 into correction [81, 237, 238]. Each episode had a Chinese or cost-side trigger. None yet had the big one: a demonstrated frontier-parity model trained end-to-end on the open, domestic, capacity-first stack of Chapter 12. That is the event this chapter’s scenarios pivot on.

21.2.1 The transmission argument no longer requires China

Every version of this chapter’s mechanism has so far required a chain of four links: Chinese open weights price below the Western closed frontier; the closed frontier cuts to defend share; the cut compresses the model rent; and the compressed rent fails to finance the next training run. Chapter 17 removes the first link as a *necessary* condition, and for a chapter that prices exposure rather than diagnosing intent, that removal is worth more than anything else in that chapter’s ledger. The largest supplier of the frontier’s inputs is now itself funding open-weight models at scale—weights released together with their training datasets and reinforcement-learning libraries, a \$6B non-exclusive licence to the training platform on which such models are produced, and 109 of the researchers who operated it [701, 724] [P]. To a discounted cash flow the nationality of the party destroying a rent is not a variable; what matters is that the price of the substitute goes to zero and that somebody with a balance sheet is committed to keeping it there. P3 is therefore overdetermined in *mechanism*, and the exposure map of Table 21.1 acquires a source of loss that no export control reaches and no tariff prices, because it originates inside the perimeter.

The qualification is the one Chapter 17 enters against itself, and it must be carried here at full strength rather than in a footnote, because it is the difference between a forecast and a possibility. NVIDIA’s *shipped* open models take routed volume without taking frontier capability; what has been demonstrated is that a chip vendor can commoditize the *volume* tier of the model layer, not the capability tier, and it is the capability tier on which OpenAI’s and Anthropic’s price umbrellas actually rest. The strong form of the claim above—that rent destruction at the model layer no longer requires a Chinese actor to arrive at all—is contingent on a frontier-scale release by an American accelerator vendor that at 9 September 2026 has not happened [S]. A funded intention is not a product. What is not contingent is the direction of the money: \$6B for a training platform and its senior staff is a capability purchase and not a marketing one, and a chip vendor that lists “a thriving open-model ecosystem” among the demand drivers on its own earnings call has explained why it would make one [728] [P].

The consequence for this chapter’s structure is a correlation, and correlations are what convert repricings into cascades. Section 21.3 insists that the four markets can reprice independently, and the coupling it identifies runs one way: margin compression at the model layer (4) does not mechanically deflate chip demand (1), because an open-weight model that removes an API payment often creates demand for local accelerators, fine-tuning and inference servers. That remains true. What is new is that a single balance sheet now sits on both sides of the coupling and has committed capital to both—underwriting the residual value of the datacentres of market (3), funding the open weights that deflate the output price of market (4), and holding roughly \$95.6B of equity securities concentrated in the counterparties of markets (1) and (2) [696] [P]. Each leg is individually defensible and the combination is, at present prices, individually profitable. But diversification requires the legs to fail at different times, and these legs fail together: the state of the world in which guaranteed residuals are called is the state in which realized token prices have not recovered the capital, which is the state in which the equity portfolio marks down, which is the state in which the customers who owe the supply commitments arrive to renegotiate them. From this point forward the exposure map should be read as a map with fewer independent entries than it has rows—which is a statement about the shape of a downside, not a prediction that one occurs.

21.3 Four bubbles, not one

“The AI bubble” is four distinct markets that can reprice independently, and conflating them is how both bulls and bears cheat. (The industry-structure view beneath the four markets is finer still, and the six-layer decomposition this chapter works from runs: frontier laboratories, hyperscale cloud, accelerator vendors, neocloud and private-credit infrastructure, power and data-center assets, and enterprise application and agent providers. The layers have different revenue realization, different leverage, and different loss-absorbing capacity—a collapse in model-layer margins can coexist with excellent returns to cloud, power, and applications, which is precisely what the realized Microsoft and NVIDIA results above suggest is already happening.) (1) *Public-equity valuations*: NVIDIA, hyperscalers, the power-and-cooling complex—repricing here is multiple compression, survivable by all. (2) *Private frontier-lab marks*: OpenAI, Anthropic—repricing here is a funding event, existential for the burn-dependent. (3) *Data-center and private-credit overbuild*: utilization, tenant concentration, hardware residual value, refinancing—repricing here is a credit cycle, the genuinely systemic channel given the debt stack of Section 21.1. (4) *Model-API margins*: the layer this book’s mechanism attacks directly. The couplings matter more than the components: margin compression (4) does not mechanically deflate chip demand (1)—an open-weight model that removes an API payment often *creates* demand for local accelerators, fine-tuning, and inference servers, which is why NVIDIA can be simultaneously the victim of the category argument and the arms dealer of the open-model boom. The cascade scenario below requires (4) to propagate through (2) into (3); the repricing scenario is (4) plus partial (1) with

(3) muddling through. Any analysis that prices all four as one trade—in either direction—is wrong by construction.

21.4 The Jevons rebuttal, taken seriously

The strongest argument against P4 is demand elasticity, and it deserves its own section rather than a dismissal. The empirical coexistence is striking [V]: inference prices at fixed capability fell $9\times$ to $900\times$ per year across benchmarks (median $\sim 50\times$)—while global AI compute capacity *doubled every seven months* and the acquisition cost of leading AI supercomputers doubled every thirteen [255, 256]. Cheap intelligence has so far behaved like cheap lighting: consumption grew faster than price fell, through longer contexts, test-time reasoning (reasoning-model outputs inflating $\sim 5\times/\text{year}$ [254]), agent swarms, synthetic data, video, and the substitution of AI for labor in previously uneconomic tasks. Four elasticity regimes bound the outcomes: costs fall slowly and demand grows fast (shortage persists—the 2024–25 world); costs fall fast and demand grows faster (capex keeps earning—the Jevons case, and the hyperscalers’ explicit bet); roughly proportional (revenue migrates between layers, infrastructure stays utilized); and costs fall faster than demand grows (overcapacity and repricing—this chapter’s mechanism). The honest statement: regimes two and four are both live, the data through mid-2026 cannot yet separate them, and the Chinese strategy is *robust across* them rather than victorious in both—a distinction the argument forces and §2.1.5 carries formally. In regime two China sells the marginal token and the marginal machine into expanding demand; in regime four it holds the low-cost position when the music stops. But low cost is not a theorem of victory: it fails if pricing never recovers full cost, if trust and integration dominate the purchase decision, if thin margins starve operational learning, or if US application providers capture the customer relationship atop Chinese weights—the conditions §2.1.5 enumerates. What the asymmetry does establish is exposure: the leveraged Western capital structure requires regime two specifically, while the Chinese position degrades gracefully under either. That—graceful degradation against required good fortune—is the precise content of “the wrong side of the trade,” restated without any collapse claim.

A parallel qualification belongs beside that conclusion: closed laboratories sell more than model intelligence. Distribution and consumer habit, enterprise procurement and support, security certification and indemnity, integrated agents and tooling, proprietary workflow data, and cloud bundling are all rent surfaces that open weights commoditize only slowly—this is the Office lesson of Chapter 8, and it is why the enterprise-spend countertrend ($19\% \rightarrow 11\%$ [168]) exists. The book’s claim is correspondingly bounded: commoditization compresses the *model* rent with high confidence and the *service* rent only with time and trust infrastructure (Chapter 18); valuations pricing permanent scarcity of the former are exposed now, those pricing the latter are exposed later or not at all.

21.4.1 Not all obligations are the same instrument

The circularity argument (Section 21.1) is rhetorically strong and financially incomplete, because it aggregates instruments with very different enforceability. Purchase commitments, letters of intent, vendor equity, backlogs, debt, and off-balance-sheet lease structures have different legal form, cancellation clauses, collateral, recourse, maturity, counterparty concentration, accounting treatment, and mark-to-market sensitivity—and therefore different loss waterfalls in a downturn. A \$1.4T “commitment” that is substantially cancellable at will is a different object from a \$30B contracted lease with recourse; treating them as one number overstates the cliff. The correct discipline, which this book applies qualitatively and recommends quantitatively, is to sort the stack into three tiers: *hard* (contracted, collateralized, recourse debt and signed leases—the tier that forces action in a downturn), *soft* (purchase commitments and multi-year cloud agreements

with cancellation or renegotiation provisions—the tier that absorbs the first shock), and *notional* (letters of intent, announced partnerships, headline totals that were never instruments at all). Public disclosure is thin enough that the tier boundaries are estimates; where this book quotes an aggregate, the reader should assume a mix.

21.4.2 The elasticity that decides everything

Section 21.4 argued the demand-elasticity case in words; it deserves the arithmetic. Under a simplified constant-quality model, with P the quality-adjusted unit price, Q the quantity of delivered work, $R = PQ$ the resulting revenue, and $\epsilon = -\partial \log Q / \partial \log P$ the own-price elasticity of demand exactly as defined in Section 2.1.4,

$$\Delta \log R = (1 - \epsilon) \Delta \log P,$$

so revenue grows despite price collapse whenever $\epsilon > 1$ and shrinks whenever $\epsilon < 1$. The whole disagreement between this chapter’s bear case and the hyperscalers’ bull case reduces to that inequality—and the honest position is that ϵ is almost certainly *not uniform*. It is plausibly well above one for coding agents, synthetic-data generation, video, and the scientific inference of Chapter 24, where cheaper tokens unlock work that was previously uneconomic; plausibly near or below one for consumer chat and for enterprise summarization, where usage is bounded by human attention rather than by price. A portfolio of segments with ϵ ranging from 0.5 to 3 produces aggregate behaviour that neither camp’s single-number argument predicts, and it is why this chapter’s scenarios are stated as branches rather than as a forecast. It is also why depreciation risk must be modeled by *workload and hardware generation* rather than imported wholesale from the fibre analogy: a GPU’s residual value depends on its memory capacity, software support, power efficiency, interconnect compatibility, and whether inference demand keeps older parts useful—an equity drawdown of 70% does not require the physical assets to become worthless, and mostly they do not.

21.5 Failure modes, entity by entity

21.5.1 OpenAI: the leveraged pure-play

OpenAI holds the largest and most fragile position, and its 2026 financing history is the position’s own biography. The March 2026 round—\$122B at \$852B post-money, with Amazon, NVIDIA, Microsoft, and SoftBank all participating—was the largest private financing ever completed, raised against \$13.1B of 2025 trailing revenue and an ARR that ended 2025 near \$21B and reportedly plateaued around \$25B through the spring [205, 206]. The press-reported 2025 operating loss was ~\$21B (with a net loss near \$38B including one-time charges), of which \$17.2B went to Azure alone; cumulative losses through 2029 are projected above \$100B [207]. The business model underneath is the pure scarcity rent—metered access to closed capability—and every force in this book bears on it directly: token deflation resets its price umbrella quarterly, open-weight parity (weeks behind, per Chapter 10) caps its differentiation window, and quantized self-hosting removes its highest-volume customers from the market entirely.

The commitments side has already blinked once, and the episode deserves more attention than it received. In February 2026 OpenAI reportedly reset expectations with investors, walking the headline ~\$1.4T of announced compute commitments back toward a ~\$600B-by-2030 spending target against ~\$280B of projected 2030 revenue [379]. Read through the obligations-tier discipline of Section 21.4: the walk-back is direct evidence that a large fraction of the \$1.4T was notional rather than hard—which softens the cliff, exactly as this book’s tiering predicted, while confirming that even the principal treats the announced number as negotiable. The structural relationships shifted in parallel: the for-profit conversion completed in October 2025 with Microsoft holding ~27% and a \$250B Azure purchase commitment but surrendering its

cloud right of first refusal, and by spring 2026 the exclusivity and revenue-sharing arrangements had reportedly ended entirely [380]. The keystone is being slowly unmortared from its original arch even as new arches (Stargate, where the flagship Abilene site ran ~ 0.3 GW operational in April 2026 against >9 GW planned across seven sites [239]) are built on its name.

And in August the execution risk moved from the balance sheet to the org chart. Within four days, OpenAI lost its longest-tenured business executive—Brad Lightcap, chief operating officer from 2022 until he moved to a special-projects role in spring 2026, ending an eight-year tenure on 11 August—and its first chief revenue officer, Denise Dresser, on 13 August after roughly eight months, with Dali Rajic, formerly of Wiz, named successor in the same news cycle [520]. They follow Fidji Simo, the applications head recruited from Instacart, who stepped away in July, and four further executives in April including the vice-president for science. Investors were, by their own account, surprised. Coverage framed the sequence as a red flag ahead of a listing, and the framing has force: the functions vacated are precisely the ones an IPO roadshow tests, and a CRO leaving eight months in is a datapoint about the revenue story rather than about the person. Three facts belong beside the alarm before it is weighed, and none of them dispels it. The finance function did not turn over: CFO Sarah Friar briefed investors on 14 August at a meeting scheduled before the shake-up, reporting that enterprise revenue now exceeds the consumer business—the crossover the rent-migration table of Section 21.8 predicts—and the president told employees that run-rate revenue grew more than 20% month over month in July, 32% for business customers, with the enterprise base doubled in a year to two million customers [520]. The governance structure did not change: the OpenAI Foundation retains sole power to appoint and remove the public-benefit board. And the safety-leadership attrition running in parallel—six safety leads in two years, the function folded into research—is a distinct thread from IPO execution, belonging to Chapter 19 rather than to this one, though a reader may reasonably decline to treat them as unrelated. What a prospective public shareholder is being asked to underwrite, though, is a governance record: a chief executive removed by his own board in 2023 for being “not consistently candid” and restored within days by investor and employee pressure, an episode re-litigated in open court in May 2026, and a culture two former employees describe as hire-fast-fire-fast [520]. None of that is disqualifying and much of it is ordinary for a company at this velocity. It is, however, exactly the material an underwriter must package, and it explains why the interval between a confidential submission and a public filing has been long.

The strongest bear reading of the sequence deserves its own statement, because it is available and this book does not dodge available readings. It runs: people do not usually leave immediately before the largest payday of their careers. Dresser arrived from Salesforce in December, took on expanded responsibilities in April, and resigned in August with a listing months away, forfeiting what a public offering at this valuation would have been worth; Lightcap left the same week after eight years. These were not junior employees but the people who ran the commercial side, and they had the closest available view of the revenue operation. On that reading, the exposure is not confined to OpenAI: NVIDIA has been negotiating a backstop of its data-centre debt, Oracle has staked its capital-expenditure strategy on a compute contract exceeding \$300B, and Microsoft has committed tens of billions—so if the executives with the best view of the numbers declined to wait for their shares, the shareholders of those three companies are entitled to ask why.

The bear reading is stronger than its critics allow, and the reason is a base rate. Executives do leave companies. They very rarely leave a company months ahead of the largest liquidity event in the history of their industry, when staying costs nothing but patience and leaving forfeits a life-changing sum. That base rate is close to zero, and against it the count is not ambiguous. Set aside entirely the two departures the individuals attributed to serious illness—this book takes people at their word about their own health, and the argument does not need them. What remains is still the chief operating officer, the chief revenue officer, the vice-president for science, the chief futurist, and the successive heads of safety and of ethics, within months, at a company

whose confidential registration values it near \$852B. Each of those had a closer view of the operation than any outside analyst, and each chose to forgo the payout rather than wait for it. One such exit is noise; this many is a signal about what the inside view looked like, and it is one that the growth figures do not dispel—20% month-over-month run-rate growth in July, 32% for business customers, an enterprise base doubled to two million—because revenue growth and an unattractive inside view are not mutually exclusive at a company burning capital at this rate. The growth is real and it is in the right row of the rent-migration table. It does not answer the question the departures pose, which is why the people who produced it declined to stay for its reward.

What the episode actually establishes for this chapter’s argument is narrower and more useful than “OpenAI is in trouble.” A company that has confidentially submitted a draft registration statement, disclosed on 8 June 2026, at a valuation near \$852B, has no announced listing date; prediction markets priced roughly a 19% chance of a 2026 listing at the time of the departures; and one analyst reading has the offering slipping to 2027 [521, 520]. The repricing mechanism of Scenario I does not require OpenAI to fail. It requires the moment of public price discovery to arrive—and every quarter that moment is deferred is a quarter in which the private mark continues to collateralize SoftBank’s equity, Oracle’s backlog, NVIDIA’s investment, and CoreWeave’s debt without being tested. Delay is not safety; it is accumulation.

The failure mode remains repricing, not bankruptcy—the strategic value to its patrons is too high for the latter. An IPO (confidential preparations for end-2026; Amazon’s \$35B contingent commitment is explicitly tied to a listing or an AGI milestone [209]) at a fraction of the private mark would reprice every claim collateralized on the old one: SoftBank’s equity, Oracle’s contracted backlog, NVIDIA’s investment, CoreWeave’s debt. OpenAI does not have to fail for the arch to move, only to be re-marked—and the February reset was the first re-marking, performed voluntarily, in private, at the top.

21.5.2 Anthropic: better economics, same trade—and the strongest quarter anyone has printed

Anthropic is the disciplined version of the position, and marking it correctly (per the accounting framework above) makes it look substantially stronger than the stale figures did: run rate from ~\$9B in December 2025 to \$47B by May 2026—a five-fold expansion in five months, overwhelmingly enterprise—against a \$965B post-money Series H that now stands at $\approx 20.5\times$ run rate, and a confidential draft registration statement whose submission the company disclosed on 8 June 2026 [212, 213, 521]. That is the single best revenue quarter-sequence any AI company has ever reported, it happened *inside* the open-weight price collapse this book documents, and intellectual honesty requires stating what it is evidence for: the enterprise-workflow and trust rows of the rent-migration table (Section 21.8) are carrying rent at scale, right now, exactly as the American flywheel’s fortification plan requires.

The exposure has migrated from the revenue line to the margin line, which is where the accounting framework of this chapter locates it. Analyst work puts compute cost at \$0.71 per revenue dollar in Q1 2026, projected \$0.56 in Q2—a trajectory toward $\sim 44\%$ gross margin against a 40–50% target, with the company reportedly lowering its own margin projections even as revenue soared, and a projected 2026 loss around \$14B on $\sim \$1.25\text{B/month}$ of compute spend [369]. The strategic reading: Anthropic is buying growth at a gross margin that open-weight price pressure helps set, funded by a capital stack that now includes its own suppliers (NVIDIA up to \$10B, Microsoft up to \$5B, against a \$30B Azure commitment—the circular pattern, joined late [381]), plus memory vendors on the cap table. Its franchise is concentrated precisely where Chinese open models advance fastest—coding and agentic work—and where its own dependency graph is deepest. The honest statement: Anthropic has demonstrated that the premium tier is real and large, faster than this book expected; whether it has demonstrated *durable* rent depends on the margin trajectory, not the revenue one, and the margin trajectory is the single most

information-dense financial series in the industry (its collapse below $\sim 35\%$ would, per the same analyst work, compress fair value by 70–81%).

21.5.3 NVIDIA: three exposures compounding, one quarter at a time

NVIDIA’s position is the largest single concentration of the scarcity bet, and its realized results keep getting better while its structural exposures keep getting worse—both facts, held together, are the correct reading. The realized side first: Q1 FY2027 (April 2026) revenue of \$81.6B, up 85% year on year, \$75.2B of it data center (networking alone +199% to \$14.8B), at a 74.9% GAAP gross margin, guiding \$91B for the following quarter with China data-center compute assumed at zero; Jensen Huang’s order-book claim was raised at GTC 2026 from \$500B through 2026 to $\sim \$1\text{T}$ through 2027, with the Vera Rubin generation (claimed $3.5\times$ Blackwell on training) entering production ramp in H2 2026 [467, 382]. No company has ever printed numbers like these; any bear case that pretends otherwise is not serious.

The exposures, updated. *Commercially*: two unnamed direct customers were 23% and 16% of revenue—39% combined and rising—and data center is now 92% of the company; the marginal buyer increasingly requires NVIDIA’s own balance sheet to close, though the flagship instrument wobbled: the \$100B OpenAI letter of intent (progressive, per-gigawatt, explicitly non-binding) was reported stalled by January 2026 amid internal doubts about size and structure, even as up to \$10B went into Anthropic and analysts flagged $\sim \$25\text{B}$ of CoreWeave debt as NVIDIA-linked circularity [216, 383]. Vendor financing that the vendor itself hesitates to complete is the cycle’s most honest self-assessment to date. *Architecturally*: the D–H–M result (Section 12.2) dissolves the category moat—the GPU is one corner of an allocation continuum any strong silicon team can enter from the CPU side, as LineShine demonstrated at No. 1; NVIDIA’s own line concedes the direction with Grace CPUs, LPDDR-based SOCAMM2 memory, Vera at 1.2 TB/s of LPDDR5X, and CUDA ported to RISC-V hosts [185, 136, 190]. *Financially*: a \$5T market capitalization prices permanent 70%-plus margins on a part whose two key inputs (HBM, advanced packaging) are in shortage, whose largest strategic market has been legislated away, and whose customer base is simultaneously its investment portfolio. The bookings are real, and so was Cisco’s order book in March 2000; the buyback (\$19.3B in one quarter, plus an \$80B new authorization [467]) says management prefers returning cash to holding it against the cycle. NVIDIA survives every scenario below as a company—almost certainly as a great one; the \$5T claim on *permanent* scarcity does not survive most of them.

A fourth exposure was constructed in August, and it is the most consequential of the four because it is the first one NVIDIA built deliberately. Chapter 17 sets out what the 10 August memoranda with Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR are, and what the company says they are for; this chapter’s business is to price them [692, 692]. Three precisions belong in the record before any valuation is attempted, and two of them are routinely dropped: these are memoranda and not funds, the release stating plainly that the partnerships “remain subject to execution of the final agreements”; the \$500B is a mobilization target rather than committed capital; and NVIDIA’s own contribution appears nowhere in the release, the widely reported figure of up to 25% residual-value support coming from Huang’s subsequent public remarks [P/A]. What the announcement does establish is the destination, and Goldman’s David Solomon named it without ornament: “a market for credit backed by NVIDIA compute.”

Read this against the same week’s other disclosure and the manoeuvre resolves. On 14 August, reporting had NVIDIA cutting its contemplated backstop of OpenAI’s Ohio campus from roughly \$250B to under \$120B, covering only the project’s first phase, after investors raised concerns about the company’s risk exposure to large financing commitments [516]. The two events are one move: exposure migrating off NVIDIA’s balance sheet and onto institutional capital, with the vendor retaining the demand-stimulating benefit and shedding a large part of the credit risk. Analyst readings split predictably and both are worth carrying—Morgan Stanley notes that third-party capital genuinely alleviates the circularity objection, since NVIDIA participates only

to a limited extent and books usage-linked revenue without bearing depreciation; Hedgeye’s Felix Wang supplies the sharper sentence: “in effect, they made NVIDIA’s product cheaper without really cutting GPU prices,” while making “future demand more sensitive to credit conditions, credit volatility,” and raising “questions on what we consider to be real demand” [515].

The telecom parallel is available and must be stated with its disanalogy attached. Nine telecom equipment suppliers had extended on the order of \$25.6B of vendor financing by end-2000; \$500B is roughly twenty times that, and when that cycle turned, twenty-four of the thirty largest listed carriers went bankrupt and a third to four-fifths of the suppliers’ loan portfolios were lost [515]. But Lucent and Nortel held their paper on their own balance sheets, and this structure deliberately does not. That is a real risk-management improvement for NVIDIA and a real risk *transfer* to pension funds, insurers and private-credit vehicles—which is to say the exposure has not been reduced, only relocated to holders whose mandate is safety. Whether that is prudence or contagion depends entirely on the residual-value assumption underneath it, and the entity best positioned to know how fast the collateral depreciates is the entity selling the collateral.

21.5.4 The fifth exposure: guarantees, residuals, and a portfolio that marks with its customers

The four exposures above are a commercial concentration, an architectural erosion, a valuation premised on permanent margin, and a financing platform. A fifth is now *disclosed* rather than inferred, and it is the one this chapter’s apparatus was least equipped to price, because it is neither a position in a security nor a claim on revenue but a contingent promise to make somebody else whole. The August 2026 filings state it plainly. Residual value guaranties with SB Energy on datacentre leases at Pike County, Ohio—tenant an OpenAI affiliate, approximately 4.25 GW of IT load initially with an option on a further 3.8 GW, twenty-year maximum term—carry an aggregate payment obligation “cumulatively capped at \$105 billion” [694] [P]. The quarterly disclosures put maximum gross guarantee exposure at \$108.5B, supply and capacity purchase commitments at \$279B—up from \$119B the prior quarter, and the increase is memory-driven, which is the shortage of Chapter 13 arriving on the buyer’s side of the ledger—AI-cloud revenue-share commitments at \$36B, future equity investment commitments at \$25B, datacentre leases held for third-party reassignment at \$20B, and equity securities at roughly \$95.6B [696] [P]. Beside them sit a residual-purchase obligation to CoreWeave capped at \$6.3B running to 2032 and, on reporting rather than a filing, the neocloud revenue-share and credit-support arrangements [697] [A]; and above them the \$500B of financing memoranda [692]. Chapter 17 argues about what this practice *is*; the question here is narrower and answerable, which is what it would cost.

The obligations discipline of Section 21.4 sorts the stack into hard, soft and notional, and a residual value guarantee fits none of the three cleanly, which is precisely why it deserves its own treatment. It is not hard, because it schedules no payment and may schedule none ever. It is not soft, because it cannot be renegotiated away by a counterparty deciding to buy less. And it is emphatically not notional, because it is a filed obligation with a stated cap. What it is, in the vocabulary of credit rather than of capex, is *wrong-way risk*: an exposure whose probability of being called is positively correlated with the impairment of the party that wrote it. A guarantee on the residual value of AI datacentres pays out in exactly one state of the world—the one in which AI datacentres are worth less than their leases assumed—and that is the same state in which the guarantor’s silicon is repricing, its equity portfolio is marking down, and its supply commitments are being renegotiated. Insurance that pays only when the insurer is also impaired is not diversification; it is leverage wearing a different name. Morgan Stanley, initiating credit coverage at Neutral, put total credit exposure at roughly \$200B by end-2028 of which about \$170B sits off balance sheet, against projected net cash near \$520B in fiscal 2029, and coined the phrase that summarizes the business model better than any critic’s: “balance-sheet-as-a-service” [704] [A]. The coverage ratio is the bulls’ answer and it is a good one—an exposure of \$200B against half a trillion of net cash is survivable—and it is also the wrong

comparison, because the projected net cash is generated by the same demand the guarantees are written against.

One number in the same filings prices the correlation more exactly than any argument about it. In the second quarter of fiscal 2027, NVIDIA's GAAP diluted earnings per share of \$2.46 *exceeded* its non-GAAP figure of \$2.22—an inversion of the customary relationship, in which adjustments run the other way—because GAAP other income of \$7.773B was almost entirely gains on equity securities [695] [P]. That pre-tax figure is roughly an eighth of the quarter's GAAP net income, and the comparison must be stated with its caveat attached or it overstates itself: the gain is pre-tax while the net income is after tax, so the after-tax contribution to earnings is materially smaller than the ratio suggests, and none of it is improper, unusual under the accounting, or hidden—it is a line in a filing that the company volunteered. The significance is not accounting but structural. The portfolio that produced the gain is some \$95.6B concentrated in the firms that buy the chips [696], which means a measurable slice of the reported profit of the sector's largest supplier is a mark on the equity of its own customers, and that mark rises with those customers' ability to raise capital, which rises with the orders they place with the supplier. This is the circularity argument of Section 21.1 restated in a form that requires no reporting, no estimate and no adverse inference: it is arithmetic on a disclosure. It is also, and the chapter should say so, entirely reversible—an eighth of net income arriving from marks in an up quarter is an eighth departing in a down one, which makes the earnings series more cyclical than the revenue series it sits on.

The regulatory precondition is almost never reported beside the announcement, and it belongs in a chapter about credit. On 29 July 2026—twelve days before the financing platform was announced—the SEC's Office of Structured Finance took the position that securities in datacentre securitizations are not asset-backed securities under the Exchange Act, on the reasoning that datacentres are “tangible, physical assets that endure beyond the tenor of the securities” rather than self-liquidating financial assets, placing them outside Regulation AB and outside the five per cent credit-risk-retention requirement [693] [V]. The legal analysis is not obviously wrong on its own terms, and this book does not assert that it is: a datacentre genuinely is not a pool of receivables. The market-structure consequence is separable from the legal question and is what matters here. Risk retention exists for one purpose, which is to align sponsor and investor—to make the party that assembles a deal keep enough of it that it will not originate what it would not hold. Removing that requirement for an asset class whose residual value is the entire question, at the moment that asset class is being built into an investable instrument, is a policy choice; and the consequences of policy choices of this kind fall due in a downturn rather than at issuance, which is when they are cheapest to make and hardest to evaluate. No substitute alignment mechanism has been proposed. The vendor's residual guarantee is not one, because the vendor is the seller of the collateral, and the entity best positioned to know how fast collateral depreciates is the entity selling it.

It is not one firm's practice, and that matters for how the exposure should be forecast. AMD discloses \$4.1B of datacentre lease guarantees alongside \$4.5B of leases not yet commenced, and has issued two warrants for up to 160 million shares each at one cent—to OpenAI in October 2025 against six gigawatts of purchases, and to Meta in February 2026 with share-price thresholds escalating to \$600—together with up to \$5B of equity and two gigawatts of MI450 committed to Anthropic [707, 706] [P]. Stated without ornament, the customer is being offered payment in the vendor's own equity for buying the vendor's product, though as of the August 2026 filing no warrant shares had vested or been issued. Chapter 17 reads the duplication as evidence of competitive necessity rather than of grand design, and for this chapter that is the worse of the two readings, not the better one: an unusual instrument adopted because one firm chose it can be abandoned when that firm changes its mind, while an instrument adopted because the situation requires it will be adopted by whoever is next and will propagate down the credit curve, where the guarantor is smaller and the coverage ratio is not half a trillion of net cash.

The counter-evidence is strong enough that it must be printed at full strength rather than gestured at, and on the numbers available at 9 September 2026 it is the better-supported side. NVIDIA’s datacentre revenue grew 117 per cent to \$89.0B in the quarter, on total revenue of \$96.2B, with third-quarter guidance of \$108.0B and China datacentre compute assumed at zero; management guided on the call to roughly seventy per cent revenue growth in fiscal 2028 against a sell-side consensus near forty-four [695, 727] [P]. A company that expects seventy per cent growth two years out is not a company that believes it is being commoditized, and one that says so publicly has made itself falsifiable, which is more than most positions in this chapter manage. Management denies the characterization in terms—“Independent capital still underwrites every deal on its own merits. We’re not making loans”—and has denied it consistently since 2025 [699, 705] [P]. The largest structure is narrower than its headline: the Ohio guarantee *terminates early if OpenAI achieves a satisfactory credit rating*, and OpenAI must reimburse and indemnify NVIDIA for every amount paid [694], which makes it structurally a bridge to creditworthiness rather than permanent underwriting—a distinction that is real, and that also concedes the premise, since a bridge is built where the ground will not bear the load. And the practice is not being executed with the confidence of a long-laid plan: on 27 August 2026, one day after the earnings call, parts of the July revenue-sharing and credit-support programme were paused, reportedly after employees flagged potential antitrust exposure and after friction over how far NVIDIA could dictate which customers its partners rented to [698] [A]. The sharpest external formulation remains Vellante’s, and it is a criticism of the structure rather than of the company: independent capital can fund more factories, but independent capital is not independent *demand*, and the very requirement for credit enhancement is evidence that the market does not yet accept compute as enduring collateral unaided [705] [A]. The honest verdict this chapter can support is the one Chapter 17 reaches: the practice is unusual, large, growing, disclosed, and not yet demonstrated to be either sound or unsound.

One further observation is forecast-relevant and is routinely reported at half its length. The equity commitment to OpenAI announced in September 2025 at \$100B became a \$30B position; over the same period the *compute* commitment went the other way, with Huang putting OpenAI’s “existing and planned commitments” at approximately twelve gigawatts of NVIDIA compute through 2030 on the second-quarter call [728] [P]. The popular framing—“NVIDIA cut its OpenAI bet by seventy per cent”—reports only the number that fell, and the two numbers are not substitutes. What the pair describes is a conversion of exposure from one instrument into another: NVIDIA reduced what it is prepared to *own* of its largest customer and increased what it expects to *sell* that customer, which lowers its mark-to-market exposure to an OpenAI re-rating and raises its exposure to OpenAI’s ability to pay. Equity risk became receivable and residual risk, and only the first of those is marked quarterly. Set the twelve gigawatts beside the walk-back of OpenAI’s own commitments toward a ~\$600B-by-2030 spending target recorded above and the arithmetic does not close on the customer’s balance sheet at any capital intensity this book has priced [M]—which is not evidence of misstatement by either party but a statement of what the financing platform is *for*. The gigawatts are contracted against capital that has not yet been raised, from investors who will underwrite it against a residual value the vendor has guaranteed. That is the trade, stated in one sentence, and Table 21.3 now carries the line item that resolves it.

21.5.5 Google: the hedged incumbent

Google is the best-positioned Western player because it is the least pure-play: its own TPU silicon—now with gigawatt-scale external sales, up to a million TPUs for Anthropic in the largest such deal ever [213]—a search-and-ads cash engine, a ~\$460B cloud backlog, distribution across billions of users, and a closed frontier model. It is listed here for two reasons. Its model-layer rents compress with everyone else’s—Gemini’s pricing power is set by the same open-weight floor. And its silicon franchise, often read as a hedge against NVIDIA, is a hedge *within* the

scarcity trade, not against it: TPUs are HBM-bound, centralized, and closed; the capex behind them (~\$175–185B guided for 2026, financed partly by a bond program that now includes a hundred-year sterling note) helped produce Alphabet’s first negative-FCF quarter since its IPO [229, 387]. Google survives commoditization better than any lab; its AI-infrastructure returns do not, and the TPU’s external success makes the point rather than refuting it—selling scarce compute to a competitor at the top of the cycle is what a rational owner of depreciating scarcity does.

21.5.6 SoftBank and Arm: the double long, fully assembled

SoftBank has now completed the construction of the single most concentrated wrong-side position in the market, and the 2025–26 sequence deserves to be recorded step by step because each step increased the correlation. November 2025: sold the entire NVIDIA stake (\$5.83B)—the second time SoftBank has exited NVIDIA early—explicitly to fund OpenAI and Stargate, alongside T-Mobile share sales [218]. Through 2025–26: built toward >\$60B invested for ~13% of OpenAI, racing to fund a \$22.5B year-end payment, taking a ~\$40B loan in the process, with lenders reportedly uneasy that OpenAI shares themselves back Stargate equity commitments [219]. November 2025: completed the \$6.5B Ampere acquisition, adding an Arm-server silicon bet to the ~90% Arm holding. The result is one balance sheet long, simultaneously: OpenAI equity at peak private marks, the ISA royalty stream this book documents both superpowers routing around, an Arm-server vendor, and Stargate construction equity—four positions, one scenario.

The remarkable fact is that the position is currently *winning*, and the honest ledger says so: the fiscal year ended March 2026 delivered ~\$46B of Vision Fund gains—more than \$45B from OpenAI alone—and a ¥5T annual profit, even as S&P cut the outlook to negative on asset-quality concerns and the shares swung 45% peak-to-trough during the year [384]. Masayoshi Son has stopped hedging the thesis: \$5T a year of global AI investment justified, bubble talk dismissed, superintelligence declared visible in OpenAI’s own model-designing-models workflow [385]. Arm, for its part, printed a record FY2026—\$4.92B revenue, record royalties, data-center royalties more than doubling, a claimed ~50% CPU share among top hyperscalers, and its own “AGI” datacenter CPU with \$2B-plus of claimed demand [386]—which is genuinely strong *and* genuinely exposed: the datacenter CPU growth story that justifies the multiple is exactly the layer where RISC-V (Section 12.10), Ampere’s own architecture licenses, and every hyperscaler’s custom-silicon program compete, and where CUDA now hosts on the open ISA. The double long remains what it was—both legs fail in the same scenario, an OpenAI re-mark plus ISA commoditization—but it must now be priced as a position that has already paid out once, held by an investor whose stated intention is to let it ride to ¥1,000T of net asset value. Leverage that has recently worked is the most dangerous kind, because it compounds conviction along with exposure; the Vision Fund’s own history is the case study, and this time the position is correlated with the entire index.

21.5.7 The data-center complex: capacity as liability, vintage by vintage

The buildout itself—hyperscalers, Oracle, the neoclouds, the private-credit stack beneath them—holds the trade in physical form, and by mid-2026 the ledger had grown teeth. Hyperscaler capex guidance for 2026 alone reached a combined ~\$700–745B (Amazon ~\$200B, Microsoft ~\$190B, Alphabet \$175–185B, Meta raised twice toward \$145B), against trailing depreciation of only ~\$149B—a gap that is either the greatest capital deployment in industrial history or the largest depreciation wall ever scheduled, depending on utilization [387]. The debt stack beneath it: Morgan Stanley sees ~\$570B of AI-related debt issuance in 2026 (~\$236B priced by May, four times the prior-year pace) toward an \$800B private-credit data-center opportunity; Meta’s \$27–30B Blue Owl joint venture holds Hyperion off its balance sheet; Amazon printed \$15B and \$24.9B bond deals in eight months; and demand coverage on new hyperscaler issues

fell from $\sim 5\times$ in February to under $2\times$ by July—the bond market’s own dashboard indicator, moving [223, 224, 388]. The stress cases sharpened from spreads into probabilities: Oracle’s remaining performance obligations ballooned to \$638B—of which the OpenAI contract ($> \$300B$ over five years from 2027) is an estimated 40–60%, with OpenAI able to reassess after ~ 5 years while Oracle’s underlying leases run 15–19—against $\sim \$380B$ of debt-like obligations, first-ever-in-years negative free cash flow, both agencies on negative outlook, and the stock down more than half from its high [225, 389]. CoreWeave’s five-year CDS reached ~ 855 bps in July 2026—an implied default probability near 50%—on $\sim \$21\text{--}25B$ of debt, a \$99.4B backlog it must finance to earn, and quarterly interest expense guided toward \$700M [228]. By early August the equity market had joined the credit market’s verdict: the stock in a bear market off its July low with 15% short interest, the net loss widened to \$740M, and two sell-side downgrades landed in the week before the Q2 print—even as revenue doubled year-over-year, which is the vintage argument in one company: growth and distress are not opposites when the growth is debt-financed against deflating prices [483]. Two internal contradictions remain on the record from the boom’s own principals: Microsoft’s CEO stating the binding constraint is power, with GPUs “sitting in inventory that I can’t plug in,” and the depreciation critique (\$176B of understated 2026–28 depreciation via stretched useful lives—lives that Amazon has already partially *reversed*, taking a \$1.3B charge) [227, 226, 387].

One methodological move is unavoidable here, and it is this chapter’s working model: stop pricing “the data-center sector” as one undifferentiated bubble and model it *vintage by vintage*. One representative vintage v (a cohort of capacity energized in a given half-year, with its hardware generation, contract book, and financing terms) has

$$\text{NPV}_v = -\text{CAPEX}_v + \sum_t \frac{u_t q_t P_t - C_{\text{power},t} - C_{\text{network},t} - C_{\text{labor},t} - C_{\text{maint},t}}{(1 + r_{\text{disc}})^t} + \frac{\text{residual}_T}{(1 + r)^T}, \quad (21.1)$$

summing over periods t from 1 to the exit horizon T_{exit} . Here CAPEX_v is the up-front build cost of vintage v ; u_t is utilization (the fraction of installed capacity actually sold in period t); q_t is delivered throughput per unit of capacity, which rises with software optimization even on fixed silicon; P_t is the realized price per unit of delivered work, deflating 9–900 $\times$ per year at constant capability, so that $u_t q_t P_t$ is the period’s revenue; $C_{\text{power},t}$, $C_{\text{network},t}$, $C_{\text{labor},t}$ and $C_{\text{maint},t}$ are the period’s electricity, connectivity, staffing and maintenance costs; r_{disc} is the discount rate applied to the financing structure (materially higher for a private-credit neocloud than for a hyperscaler’s balance sheet, which is itself half the argument); and $\text{residual}_{T_{\text{exit}}}$ is the resale or redeployment value of the hardware at the end of the horizon. Running the corners exposes what the aggregate argument hides. A 2024 H100 vintage with a hyperscaler tenant on a seven-year contract survives almost any price path: its capex is sunk at pre-shortage prices, its tenant’s credit is sovereign-grade, and inference demand keeps old silicon busy at degraded prices—positive NPV, merely mediocre. A 2026 GB300 vintage financed at private-credit spreads with a single-name neocloud tenant and a five-to-seven-year debt tenor against 15-year lease obligations fails under *modest* stress: a 20-point utilization shortfall or one tenant repricing pushes Eq. (21.1) negative before the first refinancing, which is precisely the configuration CoreWeave’s CDS is pricing. And a 2027–28 vintage *not yet built* is an option, not an exposure—the walk-back of OpenAI’s commitments is the option going unexercised. The repricing scenario of this chapter is therefore properly stated as a *vintage-selective* event: the marginal, late, levered, concentrated vintages fail; the early, integrated, diversified ones keep compounding—which is how 2001–02 actually ran, and why “the sector” both crashed and kept growing.

The facility term, priced. Eq. (21.1) treats CAPEX_v as one number, and it is worth taking apart, because its two components have different lives and the difference is where the stranding lives. Section 12.9.4 builds a Vera Rubin rack’s fully loaded cost from the ground up, using this author’s companion survey of what it costs to build the facility such a rack requires [597]:

about \$12 per watt of IT load for a Tier-III-capable inference hall before land and before a single GPU—\$2.5 million of building for every rack that costs a hyperscaler \$7.8 million and an enterprise buying through an OEM nearer \$10 million, both prices with NVIDIA’s margin inside, on a fourteen-to-fifteen-year life against the silicon’s four to six—plus \$300 per kilowatt-year to staff, maintain, insure and tax it, plus a cost of capital that ran from 6.6% on Meta’s Hyperion bonds to 9–9.75% on CoreWeave’s notes [598, 602, 604, 605]. The result at the operator’s planned 60% utilization is \$0.90–1.17 per million output tokens at fifty tokens per second per user, of which the building and its operations are about 12%—silicon dominates a per-token cost at full load. The building dominates everything else. It is a third again of the capital, committed for three silicon generations, preceded by a two-to-seven-year interconnection wait, and utterly fixed: at 40% utilization the same rack serves the same tokens at \$1.35–1.76, at 30% at \$1.80–2.35. Against it, the same section prices a five-thousand-dollar appliance in an enterprise’s existing server room—a \$500 UPS, no new circuit, no new room, no staff—at \$0.64 per million tokens at a modest agentic load and \$0.41 at capacity, *with* its cost of capital and a worse PUE charged. The two sides of that comparison differ by a rounding error in silicon economics and by an entire asset class in facility economics: one of them has a fifteen-year building whose per-token cost is a function of a demand forecast, and the other was bought after the demand showed up.

That is where the railroads belong in this chapter, and the analogy is more exact than the fibre one this book has leaned on, because railroads are the canonical case of a long-lived physical network built ahead of demand and *then partly abandoned*. The US network peaked at about 254,000 miles in 1916 and runs on about 140,000 today—55% of peak [607] [V]; Britain’s stood near 18,000 route miles in 1963 before Beeching and is 9,785 as of March 2025 [608] [V]; and JR Hokkaido, on an island whose network reached roughly 4,000 km in the JNR era, operates 2,240 km as of April 2026—about 56% of peak—after telling the public in 2016 that it could not sustain 1,237 km of thirteen sections on its own; 312 km of those have since been converted to buses, and if the remaining eight “yellow” sections, 926 km, go the same way, the sustainable core is about 1,300 km, a third of what was built [609] [V for current km and the 2016 sections; the peak is a widely repeated round number]. Nor did the shrinkage wait for the trains to stop being useful: a quarter of American rail mileage went into receivership in the Panic of 1893, and thirty percent was in receivership again by 1937 [607, 610] [P/A]. Three lessons transfer without strain. The asset survived—the track outlived every company that laid it, as the fibre did—but its owners did not, and the write-down came before the traffic did. The network then shrank not because the demand for transport fell but because a *distributed substitute* took it: the road, the truck and the car carried the traffic the branch line was built for, and Hokkaido’s yellow lines are being replaced by buses, not by nothing. And the residual value that Eq. (21.1) discounts at the horizon is the value of a track to whoever buys it out of receivership, which is a very different number from the value to whoever laid it. Read against Section 12.9.4, the buildout’s exposure is therefore not only the aggregate-demand question the bulls and bears argue over—it is that the traffic the fifteen-year hall was built to carry has a five-thousand-dollar road running beside it.

Now add this book’s specific mechanism to the vintage model’s P_t and u_t : quantization pulls high-volume inference on-premises and on-device (Section 10.6); the capacity-first arbitrage pulls training toward small, cheap, decentralized clusters (Section 12.8); the appliance economics of Section 12.9.4 pull the largest open-weight models into the server closet at a fully loaded cost below the hall’s own; and Chinese token factories on subsidized power set the floor price for whatever remains centralized [93]. The bull case for the buildout requires centralized scarcity at both ends. The bear case does not require the demand to vanish—AI demand is real and enormous—only for it to be *servable elsewhere, cheaper*. That is precisely what solar-field owners discovered about their power-purchase agreements when the module price collapsed, and what the branch lines discovered about the road: the demand for electricity never fell, and neither did the demand for transport; the rent did.

21.6 Scenarios and timelines

No single path is certain; the structure admits three, and the indicators separating them are observable quarter by quarter.

21.6.1 Scenario I: Orderly repricing (base case)

2026–2028. No crash; a grind. Token prices deflate on schedule; Chinese open share of global tokens passes 70%; enterprise spend follows usage with a 12–18-month lag as self-hosting tooling matures. Closed-lab revenue keeps growing—the market is expanding under everyone—but multiples compress from rent-multiples to utility-multiples: 30–50% drawdowns in the AI complex spread over quarters, an OpenAI listing at a fraction of the last private mark, neocloud consolidation, NVIDIA re-rated as a magnificent cyclical. *2028–2030.* The first frontier-class models trained end-to-end on domestic Chinese silicon normalize the open stack for sovereign buyers; the Global South and most non-aligned states default to it; Western labs retreat up-market into regulated, secure, and safety-critical niches—profitable, real, and small relative to 2026’s claims, in the exact pattern of First Solar (Chapter 3). China wins by adoption share, not by knockout. This is the modal outcome because it requires nothing but current trends continuing.

21.6.2 Scenario II: The cascade

The leveraged structure of Section 21.1 admits a faster path, triggered by any one of: an OpenAI funding event or deeply repriced IPO; a Chinese open release visibly at closed-frontier parity that breaks API demand growth for a quarter; a credit event at Oracle, CoreWeave, or in the private-credit DC stack; or a memory-cost shock compressing GPU-fleet economics (the July 2026 rout was the rehearsal) [238, 225]. The circular financing then runs in reverse—the same loop that amplified the boom amplifies the unwind, as vendor financing, chip-collateralized debt, and cross-held equity mark each other down. Historical envelope: pure-plays –60 to –90% (dot-com), NVIDIA in the Cisco 2000–02 corridor (–70%+ from peak while remaining a great company), SoftBank in distress on the double long, Stargate-class projects abandoned mid-construction like 2002 fiber routes—or, in the older and more exact precedent, a quarter of the network in receivership as in 1893, with the halls changing hands at track value [610]—and a US recession contribution large enough to become a political event, given AI capex’s share of GDP growth [230]. The rehearsals for the facility side are already on the record: Microsoft cancelling or deferring up to two gigawatts of leases in the six months to March 2026 on what its analyst described as oversupply relative to its own demand forecast, and Oracle and OpenAI dropping the planned expansion of the flagship Abilene site amid financing questions and, in the reporting, OpenAI’s frequently changing demand forecasts [611] [A]. *Timeline: any quarter from late 2026 through 2029; fast once started (12–24 months peak-to-trough).* China is not immune—its own involution economics are brutal (Chapter 3)—but it is structurally insulated: one-tenth the equity at risk, state-absorbed losses, and the strategic prize (the surviving industrial commons) intact by design. In 2002 the fiber glut’s buyers of last resort were the carriers who had built it; in 2029 the cheap distressed AI infrastructure has an obvious, well-capitalized, strategically motivated buyer set—and it is not the sellers’ side of the Pacific.

21.6.3 Scenario III: Bifurcated stalemate

The falsification branch—and, on the strongest version of the skeptical case, closer to the base case than Scenario I admits. Closed US frontier models open a compounding lead—recursive AI-for-AI R&D, agentic capability that open fast-following genuinely cannot match at the frontier—and enterprises keep paying rent for the best. The world bifurcates: a Western premium-closed ecosystem holding the high-value workloads, and a Chinese-led open commons holding volume,

the Global South, and the standards. Even here, note what the West is defending by 2030: a shrinking premium tier atop a foundation—ISA, interconnect, memory class, deployment substrate, half the world’s tokens—owned by the other side. Solar’s history suggests bifurcation is a stage, not an end state; premium tiers erode as the commons compounds. But this scenario preserves Western lab profitability into the 2030s and would falsify this chapter’s stronger claims. *Probability-weighting across the three is a judgment, not a calculation; the indicators below are the honest instrument.*

21.6.4 The author’s probabilities and the skeptic’s

Numbers force honesty that prose scenarios evade. Table 21.2 prints two sets of subjective probabilities [S] for 2029–30 outcomes: this book’s, and those of the most credible skeptical position the author can construct against it, disagreements unhidden. The structural disagreement is instructive: the skeptic weights the durability of the closed premium tier higher and the repricing lower; the author weights the precedent industries’ terminal pattern—premium tiers eroding as the commons compounds—higher. Both positions agree on the core: Chinese open-weight dominance is the modal world, full-stack parity is not yet the modal world, and a 2000-style systemic cascade is a serious tail, not the base case. The composite base case both positions admit: *financial repricing plus partial bifurcation*—Scenario I and Scenario III’s overlap.

Table 21.2: Subjective outcome probabilities by 2029–30 [S].

Outcome	Skeptic	Author
Chinese models dominate the global open-weight ecosystem	70%	80%
Frontier-adjacent model trained entirely on a domestic Chinese stack	50%	65%
Full Chinese parity in HBM, tooling, software, and training economics	30%	35%
Broad ($\geq 30\%$) US AI-complex valuation repricing	60%	70%
Systemic cascade comparable to 2000–02 telecom	25%	25%
US closed models retain a valuable premium enterprise tier through 2030	65%	50%
China captures both open volume and the top closed enterprise tier	30%	25%

21.6.5 Indicators, 2026–2030

21.7 The commodity turn: what happens when inference is traded

Something began in 2026 that this book’s precedent chapters would recognize immediately, and that its financial chapter must now price: *compute is being turned into a traded commodity*. Not metaphorically—contractually, on regulated exchanges, with published reference indices and settlement rules. If that project succeeds, it changes who captures value in a way that no model release does, because commodity markets have a characteristic and well-understood effect on producers: they price at the margin, and the margin is set by the lowest-cost producer with capacity to sell.

21.7.1 The instruments are real

The record as of 9 September 2026 is a market under construction, not a proposal. CME Group and the index provider Silicon Data announced compute futures for launch on 5 October 2026,

Table 21.3: Leading indicators separating the scenarios.

Indicator	Repricing / cascade signal	Stalemate (falsification) signal
Frontier gap (AI Index, arenas)	Stays $\leq 5\%$ or closes	Widens $>10\%$ and holds for 2+ years
First frontier model trained fully on domestic Chinese stack	Occurs by 2027–28 (the “Mate 60 moment” of AI)	Slips past 2029; Ascend/LX2-class training stays sub-frontier
Enterprise spend share of open weights	Follows usage upward (reverses the 2025 dip [168])	Stays $<15\%$ while closed APIs hold pricing
OpenAI financing	Down-round, repriced IPO, or new vendor backstop	Clean \$1T+ listing, burn narrowing
NVIDIA customer mix	Concentration rises; more seller financing	Broadens; financing exposure retired
Guarantee and residual exposure (Section 21.5.4)	Gross guarantee exposure and off-balance-sheet credit rise; residuals drawn; securitizations widen	Exposure flat or retired; Ohio guarantee terminates on an OpenAI credit rating
CXMT LPDDR6 / HBM3 ramp	On schedule (2H26 / 2027); LPDDR-based training clusters appear	Slips 2+ years; HBM chokepoint holds
DC economics	Utilization/pricing disappoint; debt spreads widen; depreciation lives extended again	AI revenue closes Bain’s \$800B gap [232]

each contract representing one month of GPU rental cost, referencing H100 and B200—and listed, tellingly, on NYMEX, which is to say the industry has already decided that compute is an energy derivative [522]. ICE announced two families: cash-settled GPU futures on the Ornn Compute Price Index, built from printed transactions, and—more consequentially for this book—*energy-normalized* compute futures on the NATIVX COIL index, which carries dedicated sub-indices for training, **inference**, graphics and connectivity [523, 524]. Kalshi has been trading weekly and monthly GPU-price contracts since July 2026 and publishing forward curves derived from its own order book; Architect announced perpetual futures on GPU *and* DRAM prices on the same index family [525, 526]. Silicon Data’s daily index has run since May 2025 on some 3.5 million pricing observations and is distributed on Bloomberg terminals; alongside its GPU series it publishes an LLM token-expenditure index [527].

Two features of that record matter more than the launch dates. First, the forward curves are in *backwardation*—B200 spot around \$5.62 against a thirty-six-month price near \$5.17, H100 down 13% and A100 down 15% over the same horizon [A, index-provider data] [527]. The market is already pricing the commoditization it is being built to trade. Second, the dispersion is enormous: the same H100-hour reportedly clears near \$2.74 on neoclouds and \$7.20 at hyperscalers, a 2.6-fold spread on identical silicon [527]. A single good does not trade at two prices; that spread is the measure of how far compute remains from being one.

21.7.2 Why the incumbent stands in the way

The most cited objection to compute futures is heterogeneity—an H100 is not a B200, interconnect topology is not fungible, spot is not reserved capacity. It is a weaker objection than it sounds, and a reporter covering the CME launch put the rebuttal correctly: not all steel is identical either, and hot-rolled steel trades as a single benchmark [528]. The serious objection is different, and it is about market structure rather than product definition: *there is one seller*. As that analysis puts it, the venture is an attempt to build an oil market “run entirely by Standard

Oil”—and a monopolist supplier has no interest in price transparency or in falling compute costs, since cheaper compute reduces the value of the thing it sells [528].

The mechanisms by which that interest operates are concrete, and one of them is underappreciated. NVIDIA has been extending utilization backstops to neocloud operators: a guaranteed floor under deployed capacity, with NVIDIA renting the idle GPUs itself at a predetermined rate if customer demand falls short, in exchange for a share of cloud revenue—disclosed in arrangements covering tens of thousands of accelerators, with the rate and the revenue share undisclosed [517] [A]. State that in commodity-market terms and the significance is plain: the monopolist supplier has installed itself as *buyer of last resort at an unpublished price*, beneath every spot print the new indices collect. That does not merely discourage a transparent market; it means the reference price is being formed in the presence of an undisclosed floor set by the seller.

Alongside it sit allocation control—regulators have noted NVIDIA allocating chips to customers it judges likely to deploy fastest, amid a web of some ninety partnerships and investments—and contractual visibility, since cloud partners must expose delivered, healthy, reserved and in-use accelerator counts through a centralized API [518, 519]. Precision matters here and the book will not overstate it: those requirements govern *visibility and isolation*, and this author found no clause prohibiting the resale or subletting of data-centre capacity. The defensible claim is narrower and sufficient: the supplier has real-time insight into how every partner allocates its fleet.

The irony is worth stating because it is the whole argument in miniature. NVIDIA’s financing announcement asserts that its compute is “fungible and transferable across customers and operators” [692]—because fungibility is exactly what makes an asset financeable collateral. The company needs its *hardware* to be a liquid asset class and its *output* to remain differentiated and opaque. Those two requirements are in tension, and the market being built will resolve them one way or the other.

21.7.3 If it succeeds: the solar mechanism, applied to tokens

Suppose the project works. A fungible traded product clears at marginal cost; differentiated pricing power collapses into basis; and the surplus migrates to whoever holds the lowest-cost capacity and is willing to sell it. This book has described that transition three times already, in polysilicon, in cells, and in cells-to-packs, and the outcome each time was that the incumbent’s margin became the commodity’s spread.

The Chinese position is institutionally readier for that world than the Western one, and the evidence is older than the exchanges. China stood up compute *trading* platforms in 2023—Shanghai in April, Yinchuan in Ningxia in February, the latter co-built with the Beijing International Big Data Exchange and joined by Huawei, Alibaba Cloud, SenseTime and Baidu—and CAICT under MIIT now operates a national computing platform with over a thousand registered enterprise users, more than a hundred service providers and ten regional sub-platforms [531]. (Precision: Guiyang’s and Beijing’s better-known exchanges trade *data*, not compute, and none of the compute platforms publishes standardized contracts, clearing, or volumes—they are allocation and matching venues. China built the plumbing and the political concept of compute-as-tradable-utility years before Wall Street built the contracts.) By mid-2026, 1.37 million PFLOPS of intelligent computing—about 72% of the national total—sat under integrated national monitoring, with standardized interconnection targeted for 2028 [532]. Officials describe compute as public infrastructure “similar to electricity, water or transportation networks,” with municipal compute vouchers and pilots explicitly called compute banks and compute supermarkets [P/A].

The output side is further along still. China’s daily token consumption exceeded 140 trillion in March 2026, up more than 40% from end-2025—a figure stated by a deputy head of the National Bureau of Statistics, not a vendor [533]. Price differentials are wider than this book’s earlier 1/7-to-1/150 band implied at the low end: DeepSeek V4-Pro reportedly prices roughly

34-fold below GPT-5.6 Sol on output, and Qwen3-8B has been offered at a cent per million tokens in and out [534]. Blended enterprise inference cost fell to roughly \$1.16–1.18 per million tokens in early August 2026, a yearly low, from \$2.04 at the end of May—attributed by analysts to price wars and Chinese open-weight models [535].

And the intent is stated, in official language, which is what makes this more than an inference. State media has framed token export as a national commodity strategy: sell inference through APIs while the electricity, the servers and the models remain in China, requiring no overseas infrastructure. A State Council Development Research Centre researcher is quoted directly: token exports “provide a new way to offset risks of traditional trade frictions and drive growth in services trade” [536]. That is the solar playbook stated openly—export the commodity, not the factory—with the added property that a token crosses a border without a customs entry.

21.7.4 The case against, at full strength

This is the point in the argument where a book that wanted to be right rather than persuasive has to slow down, because the counter-evidence is strong and some of it is fatal to the naive version.

The preconditions for a commodity market are not met. Of five standard conditions, two are outright absent: there is no standardized tradable unit, and supply is not fragmented—four hyperscalers control on the order of 78% of global self-built critical IT power and 69% of H100 supply [529] [A]. No stable conversion ratio exists between accelerator generations, because whether an H200 is worth more than an H100 depends on whether the workload is memory-bound or throughput-bound. The sharpest formulation in the literature holds that compute futures will not make compute fungible—they will “make its non-fungibility legible,” giving differences something to be priced *against* [530]. If that is right, the benchmark preserves differentiated pricing power rather than destroying it, and the entire mechanism above fails.

Locality, examined properly—because the standard objection is largely wrong. The received claim is that Chinese idle capacity cannot serve inference: it sits in western provinces two thousand kilometres from demand, and Beijing’s own 20 ms latency target is unmet there [537, 541]. The claim has entered the Western literature almost unexamined, and Section 14.4 takes it apart. Briefly: the 20 ms figure is *one-way* in the primary Chinese document, not round-trip; it is a specification for the *eastern* clusters rather than a west-to-east requirement; it dates from May 2021, eighteen months before ChatGPT, when “AI inference” meant computer-vision and recommendation models whose entire forward pass takes 5–50 ms; the measured east–west figures clear it; and on 12 August 2026 Alibaba Cloud put its flagship *inference* supernode into commercial service in Ulanqab, Inner Mongolia, serving Kimi K3 and Qwen3.8-Max [545]. The genuine constraint on Chinese utilization is demand concentration and hardware built for the wrong workload, not the speed of light—which does not rescue the idle capacity, but does relocate the argument from physics to economics, where it is a different and more tractable problem.

China’s own fleet is not fungible with itself. Facilities run either NVIDIA/CUDA or Ascend/CANN estates, which reporting indicates cannot be integrated into a single seamless environment [537]. A book that uses CUDA lock-in to explain why a Western commodity market is hard must concede that the same mechanism obstructs China’s national pooling ambition more severely, because its estate is more heterogeneous.

Volume is not revenue, and the 2026 data are pointed. On the neutral router where Chinese models hold roughly 61% of tokens, Anthropic reportedly holds about 12% of volume and captures roughly 46% of revenue [538] [A]. Vercel’s production gateway—enterprise traffic, not hobbyists—tells the same story with cleaner accounting: in June 2026 open-weight models ran 29% of gateway tokens on just under 4% of spend, up from about 11% of tokens in April; DeepSeek alone carried 22.6% of June tokens, within two points of Google; Anthropic took 61% of spend on 32% of tokens; and roughly one enterprise customer in eight now runs an open-weight model in production [668] [V]. Chinese laboratories are winning the commoditized tier and losing

the profitable one. That is precisely what commoditization predicts—and it undercuts the claim that commoditization transfers pricing power to the low-cost producer, because in a marginal-cost market *nobody* captures surplus at the margin, China included. The first attempt to convert that routed volume into a revenue claim on the parties who do capture it was reported in the closing fortnight of this edition: Moonshot is negotiating with Microsoft, Amazon and Google to host Kimi K3 in exchange for up to 30% of K3-related service revenue, talks described as early and uncertain, with the split, data access and token auditing all unresolved [747] [A]—and the lever is the K3 licence’s twenty-million-dollar model-as-a-service threshold, which is why a hyperscaler cannot serve K3 as it serves DeepSeek or Qwen without a contract with the laboratory [748] [V]. The same laboratory’s confidential Hong Kong filing of 3 September, reported at about \$3B on a pre-money valuation near \$50B, coincides with a reported annualized revenue of roughly \$300M as of June, per *The Information* and unverified, with no geographic split [757] [A], which is to say that the one Chinese laboratory trying to charge the American clouds for its volume has not yet disclosed how much of its revenue any cloud pays it. Section 10.4 in Chapter 10 carries the mechanism; for this chapter the point is that an unsigned revenue share is the volume–revenue split above attempting to close itself, and the fact that it has to be negotiated rather than priced is the measure of how far it has to go.

And marginal cost may not be the operative variable at all. A companion scenario analysis by this author models the 2026–30 restructuring under memory scarcity and finds a result that cuts directly against the naive reading of this section: the entrant–incumbent cost gap *never closes*, because a depreciation conveyor delivers newly amortized fleets to incumbents faster than hardware prices normalize— $3.2\times$ in 2026, narrowing to $1.9\times$ in 2027, then re-widening to $3\text{--}4\times$ by 2029–30 [540] [M; author’s model, and the self-citation discipline of the preface applies]. If that holds, then in a commoditized market the winner is whoever owns *already-depreciated* capacity rather than whoever can build new capacity cheapest—which today favours US hyperscalers running 2023–24 vintages over a Chinese fleet that is newer, memory-constrained, and paying the surge. The same analysis puts a Commoditization Crash at 25%, co-modal with a Rotating Landlord Oligopoly at 25%; it is not a forecast of Chinese victory, and this book declines to present it as one. Two qualifications belong beside it: the model’s own token-price bifurcation—roughly a dollar per million tokens for open-weight service against twenty-five to thirty for closed—is the mechanism by which the depreciation advantage could be overwhelmed from the revenue side rather than the cost side; and a scenario analysis is a structured argument, not a measurement.

The empirical trend in 2026 ran the wrong way. While compute commoditized, inference gross margins at the model layer reportedly rose from under 40% to over 70% [539] [A]. Commoditizing the input has so far *increased* the surplus captured one layer up. Any version of this argument that stops at the GPU-hour is refuted by its own period’s data.

Energy normalization strips the advantage into basis. The NATIVX index deliberately removes regional power-cost disparities from the price. If energy-normalized contracts become the benchmark, cheap Chinese electricity stops being a price and becomes a basis—captured by whoever holds the physical asset, invisible in the traded commodity.

And jurisdiction is not price-sensitive. A Chinese-hosted API endpoint carries exposure under China’s 2017 National Intelligence Law that no price cut addresses, and US federal procurement has moved against Chinese models directly. Buyers who will not buy at any price are not in the market.

21.7.5 What survives, and it is not what the headline suggested

The version of this argument that survives its own counter-evidence is sharper than the one it started as, and it relocates the lever.

China’s instrument is not the GPU-hour; it is the model. The compute market may or may not become liquid, and if it does, four hyperscalers and one chip supplier are better positioned to work it than any Chinese operator. But the commoditization that is demonstrably occurring,

at scale, right now, is at the *model* layer—open weights at roughly a dollar per million tokens against twenty-five to thirty for closed frontier service, with an independently measured capability gap between the leading US and Chinese models of a few percent. That commoditization does not require an exchange, a clearing house, or a standardized contract. It requires a licence file and a download, and it is the one vector that *bypasses the jurisdictional objection entirely*: open Chinese weights served on Western infrastructure carry no Chinese endpoint, no data transfer, and no counterparty. The buyer who will not call a Beijing API will run the same model in Virginia.

And it does not require Virginia either. Section 12.9.4’s priced bill of materials—a capacity-first local appliance running the largest open-weight model in existence, on components benchmarked against finished pre-surge systems rather than optimistic guesses—turns “run the same model in Virginia” into “run the same model on a five-thousand-dollar box in the buyer’s own server closet.” At a modest agentic load the appliance’s all-in cost is a little over a dollar per million output tokens, recovering its purchase price against Claude Opus 5 pricing in about two months; at higher utilization it falls below even a hyperscaler’s own marginal cost of serving the identical model on Vera Rubin silicon. That second result matters more here than in Chapter 12, because it removes the strongest reply available to a defender of API pricing—that the API is, after all, riding on a more efficient centralized machine. Once local inference on the hardest open-weight case reaches unit-cost parity with the hyperscaler’s own rack, the twenty-five-to-thirty-dollar API price is not covering superior infrastructure; the model-layer margin this section already documents rising above 70% is the entire explanation for the gap. The token-export strategy China’s own officials describe is aimed at the buyer who still needs a counterparty; the appliance is aimed at the buyer who no longer does.

So the correct statement of the mechanism is this. The exchanges being built will price the input; the Chinese strategy is deflating the *output* by giving away the thing that makes the input valuable. If both proceed, the incumbent faces margin compression from two directions at once—a transparent, benchmarked, backwarddated market for the hardware it sells, and a free substitute for the intelligence that hardware was bought to produce. Neither alone is decisive. Together they describe a world in which \$500 billion of institutional capital has been mobilized against collateral whose economic residual value depends on token prices that a competing bloc is deliberately driving toward zero. That is the trade this chapter has been describing from the beginning, now with an exchange-traded instrument attached to one side of it.

21.8 Where value reappears after the model becomes cheap

This section is the chapter’s capstone: it converts P3 and P4 from predictions into a map. Suppose the model layer commoditizes exactly as Part II argues. Value does not exit the industry; it migrates to whichever adjacent layer becomes the new scarcity—and each candidate layer has an *observable rent indicator* that tells the reader, quarter by quarter, where the rent actually settled. Table 21.4 is the map.

Three readings of the table close the argument. *First, it restates the propositions precisely.* P3 (rent destruction) is the claim that the first row’s rent compresses; P4 (financial repricing) is the claim that current valuations are concentrated in the rows that compress rather than the rows that inherit. The book’s thesis does *not* require all US AI rents to disappear—it requires the scarcity layer on which current valuations depend to move faster than the incumbents can reconstitute pricing power in another row. That is a falsifiable statement about relative speed, not an eschatology.

Second, it is the flywheel contest in ledger form. The American flywheel’s survival plan is to hold rows two, three, five, and nine—cloud, agents, trust, distribution—while conceding row one; the Chinese flywheel’s plan is to hold rows six, seven, and increasingly eight—hardware, energy, science—while giving row one away deliberately. The two thefts of Chapter 2 are raids across

Table 21.4: Rent migration after model commoditization: candidate scarcity layers and their observable indicators.

Layer	Possible new scarcity	Observable rent indicator
Model weights	capability at the hard end; freshness	price per <i>successful task</i> at hard workloads (Appendix G, fields 9–10)
Cloud	reliability, scale, sovereign hosting	gross margin, utilization, reserved-capacity backlog
Agents	workflow ownership; the default interface for work	application revenue and retention; seat expansion
Data	proprietary feedback loops and domain access	capability demonstrably unavailable from public data
Trust	certification, indemnity, liability transfer	regulated-sector willingness to pay (Chapter 18)
Hardware	memory, interconnect, packaging	delivered tokens or training progress per dollar (Chapter 12)
Energy	firm deliverable power at the busbar	busbar price; time-to-interconnect (Chapter 14)
Science	validation and experimental coupling	independently replicated outcomes per resource (Eq. 24.3)
Distribution	default interface and habit	active users, enterprise spend, switching cost

this table: open weights attack the first row’s premium before the American side has fortified the others; closed agents attack the ninth row’s habit before the Chinese side can follow its cheap models up the stack. Which side’s fortification completes first is the empirical question, and every row has a column to watch.

Third, it disciplines both cases. A bull on the US complex must name the row that will carry the valuations when row one’s rent is gone, and show its indicator moving—agent revenue and retention are the strongest current candidates [V in the Microsoft run-rate, above]. A bear must concede that rows two through nine are real and large—the realized cloud and accelerator profits already booked are rent that arrived, not rent that will vanish—and locate the fragility precisely where this chapter has: in the instruments priced as if row one’s scarcity were permanent. The table, checked annually, is this book’s substitute for the crash-date prediction it declines to make.

21.9 What survives

End every market analysis with the balance-sheet truth: crashes destroy claims, not capabilities. After 2002 the fiber remained lit and the routers kept routing—under new owners at ten cents on the dollar. In every scenario above, the physical and institutional assets survive: the fabs, the power contracts, the memory lines, the open weights (which carry no market price and therefore no mark to break), the trained engineers, the standards. What does not survive intact is the market value specifically premised on AI staying scarce—and which vintages of it fail first, Eq. (21.1) now says precisely. The war is not won at the moment the incumbents’ valuations break; it was won earlier, when the cost curve changed hands, and the break merely publishes the result.

Chapter 22

The Non-Aligned: How the Rest of the World Chooses

Every chapter so far has described a contest between two principals. This one argues that the outcome will be settled by everybody else. The industrial history of Part I is unambiguous on the point: solar was decided in German feed-in tariffs and then in desert tenders; batteries in the procurement rules of a dozen national auto markets; electric vehicles in Southeast Asia, Latin America, and the European price war. In each case the producer with the best cost curve won because *third markets bought it*, and the incumbent's home market was too small to matter by the time it noticed. AI will be decided the same way, and this chapter therefore treats the non-aligned not as spectators but as the electorate.

The framing correction matters for the book's own recommendations. The strategic question for a country that is neither the United States nor China is not "which stack do we join?" It is: *can we evaluate, fork, certify, and operate components from both, while retaining sovereign scientific and industrial capability?* That is a capability question, not an allegiance question, and it has a concrete answer for each actor.

22.1 The decision matrix

Any serious national choice runs through eight dimensions. Stating them explicitly does two things: it exposes the trade-offs that bloc-based framing hides, and it makes different countries' choices comparable.

Two dimensions deserve emphasis because they are routinely omitted from national AI strategies. *Export-control exposure* is asymmetric across the stacks and zero for neither, and Section 22.2 prices it row by row. *Industrial spillover* is the dimension that separates a strategy from a purchase: importing tokens builds nothing, while operating models, running probes, and contributing to standards builds people, and people are the only input that compounds.

22.2 Two commons, not one

Everything above was written for a world in which the menu had two items with different shapes: a Chinese commons, free to download and cheap to serve but unaudited and jurisdictionally awkward, and a Western product, expensive and metered but certified and integrated. Chapter 17 ends that world. There are now two commons, and the second one arrives with its training data attached and a chip vendor's balance sheet behind it. NVIDIA's Nemotron family ships open weights together with the training datasets and the reinforcement-learning libraries, under a permissive licence, through the same hubs that carry the Chinese releases; AMD's Instella line is fully open in the strict sense—weights, data and training code; the accompanying open

Table 22.1: The eight dimensions of a national stack decision. The right-hand column names the chapter of this book that supplies the evidence.

Dimension	Question	Evidence
Capability	Does the stack satisfy our highest-value workloads—not the benchmark, the workload?	Ch. 10, 24
Cost	What is the full cost per <i>successful task</i> , including operation, integration, and retries?	§10.6.3
Sovereignty	Can weights, runtime, data, and hardware be operated independently of any foreign decision?	Ch. 12, 26
Trust	Can provenance, integrity, behaviour, and liability be certified—by us or by someone we trust?	Ch. 18, 19
Interoperability	Does it integrate with our existing cloud, scientific, and enterprise systems?	Ch. 12
Export control	Is continued access exposed to unilateral restriction by a foreign government?	Ch. 13, 20
Standards	Does adoption create long-term dependence on one ecosystem’s interfaces?	Ch. 20
Industrial spillover	Does adoption build domestic capability, or merely import a service?	Ch. 2

pretraining collection exceeds ten trillion tokens [701, 708] [P]. For a ministry choosing a stack next quarter this is a different menu, and the eight dimensions of Table 22.1 have to be re-run against it rather than extended by one column.

Run them, and the two options resolve into a trade rather than a preference. *Nemotron with open data and CUDA underneath* scores unusually well on trust, for a reason that has nothing to do with benchmarks: a model released with its corpus and its recipes is auditable and, in principle, re-derivable, which a weights-only release is not. A buyer who can retrain is a buyer who can inspect, and this is the first time that dimension has moved toward a *Western* artifact rather than away from one. It scores badly on export-control exposure and on industrial spillover, and the two failures are one failure: the model is free because the complement is not, and the complement is the most tightly controlled product in world trade. The buyer acquires a permanent accelerator dependency priced by a single vendor at a seventy-five per cent gross margin, whose availability to that buyer is set in Washington and revised without notice. *GLM, Qwen, Kimi and their peers with open weights on a capacity-first hardware path* score better on cost, on operational sovereignty and on export-control exposure—a weight file already downloaded cannot be revoked by anyone, and the hardware underneath it is being deliberately designed for parts that are not embargoed (Chapter 12). They score worse on data provenance, since the weights ship without the corpus; worse on trust, for the reasons Chapters 18 and 19 document; and worse on the question a domestic regulator asks first, which concerns a jurisdiction rather than a model. Neither column dominates, which is better news for the buyer than either column winning would be.

But the second commons has a property the first does not, and a procurement officer who misses it will misprice the option. *It has a single principal contributor with a direct commercial interest in the complement.* Openness supplied by the vendor of the complement is openness with a direction: it is free at the model layer precisely because it is not free at the silicon layer, and the sponsor has said so on the record, repeatedly and without embarrassment—nearly all open models run on its hardware, and its position in open models is good because its ecosystem is everywhere [700] [P]. A buyer taking Nemotron is therefore not taking a neutral good. It is accepting a subsidy whose purpose is to make a different purchase inevitable, which is the mechanism of Section 2.3 operating exactly as the theory predicts and not a scandal. What follows is practical rather than moral: the durability of this openness is a function of the sponsor’s

strategy rather than of the licence, and the place to monitor it is the sponsor's earnings call. Two things cut the other way and belong in the same breath. A permissive licence on an artifact already released cannot be withdrawn, so the risk is never the loss of what exists but the discontinuation of what comes next—a risk a buyer neutralizes cheaply by mirroring what it depends on. And the same risk applies to the Chinese commons, whose sponsors operate under an industrial policy that could restrict release—though its supply is at least plural, three of the four significant open-weight releases in the last week of August being Chinese and from different firms [732] [V]. A commons with many sponsors is more robust than a commons with one, whichever flag is on it.

The reported acquisition of Hugging Face is the open item, and it belongs in this chapter as a governance fact rather than a commercial one. *The Information* reported on 26 August 2026 that NVIDIA was closing on a purchase of the hub for approximately \$12.9B; at 9 September 2026 neither company had confirmed it, and other reporting held that no agreement had been signed and that the transaction could still fail [725] [A]. This book does not assert that it happened, and the guidance below is written to be correct either way. If it confirms, the distribution point through which very nearly all open weights reach very nearly all users would belong to the vendor of the complement—and what would change is not the commercial terms, which would stay free, but the default discovery surface: what appears on the front page, which artifact is easy to find, which model card is complete, whether a given set of weights remains hosted at all, decided by a party that earns on one of the answers. For a ministry that is a dependency of the same class as a single-vendor cloud region, and it is mitigated the same way and just as cheaply: hold your own copies of the weights, the datasets and the evaluation harnesses you depend on, in a repository you operate, whether or not the report ever confirms.

Which yields an instruction the decision matrix alone does not. A procurement questionnaire for an open model asks about the licence, which is the least informative thing about it. It should also ask about the *provenance of the openness*: who funded the training run, and at which layer do they earn; is the training data released, so that the artifact can be re-derived rather than merely fine-tuned; what happens to the *next* release if the sponsor's strategy changes, and what is the fallback if it stops; is the distribution channel independent of both the sponsor and the vendor of the complement, and do we hold a mirror that survives its loss; and is the model's evaluation performed by anybody who does not earn on the answer. The last is the question no vendor in either commons can answer, and it is why the certification vacancy this chapter keeps returning to has become more valuable rather than less: two competing commons increase the supply of free artifacts and do nothing whatever to increase the supply of independent assurance about them.

22.3 The export offer, as it actually exists

Everything to this point has described the non-aligned buyer choosing between artifacts. The story told about that buyer in 2026—in Washington, in Shanghai, and in most of the commentary between them—is that the choice is no longer between artifacts but between *stacks*: a financed, packaged, sovereign-branded Chinese offer of chips, datacentres, models and training, against an American one, with the telecom campaign of Chapter 6 as the template Beijing is supposed to be replaying. This section tests that story against what had actually been signed at 9 September 2026, and it does so in the order the record demands: the institutional wrapper first, then the live cases, then the two offers side by side, and then the three points at which the telecom template does and does not transfer. The finding is stated at the end and it cuts in both directions.

Table 22.2: The two commons, as a buyer must price them at 9 September 2026. Neither column dominates; the rows are where the difference lives.

Question	Vendor-sponsored Western open stack	Chinese open stack
What ships	Weights, training data, recipes; permissive licence	Weights; licence varies; corpus withheld
Who funds it, and where they earn	One accelerator vendor and a rival; on the silicon underneath	Several labs and sponsors; on hardware volume, cloud and standards, later
Hardware implied	CUDA and the controlled complement	Capacity-first, deliberately unembargoed
Export-control exposure	High—on the complement, not the model	Low on the model; the file cannot be recalled
Continuity risk	Sponsor’s strategy; one balance sheet	Industrial policy; but plural suppliers
Distribution	Hubs; ownership an open question [725]	Same hubs, plus domestic mirrors
Independent assurance	None supplied	None supplied

22.3.1 The wrapper

The institutional apparatus is real and it is three documents deep. The *Global AI Governance Action Plan* was issued at WAIC on 26 July 2025; the *AI+ International Cooperation Initiative* followed on 23 September 2025; and on 16 July 2026, in Shanghai, with Wang Yi signing for China and the UN Secretary-General in attendance, twenty-nine states signed the founding agreement of the World Artificial Intelligence Cooperation Organization—an intergovernmental body headquartered in Shanghai with a Council of members and a Secretariat, confined by its charter to the civilian domain, funded by member contributions scaled to Universal Postal Union classes, entering into force on three ratifications, and leavable on a year’s notice [764] [P]. The signatories are reported to include Russia, Brazil, South Africa, Indonesia, Kazakhstan, Pakistan, Malaysia, Serbia, Kenya and most of ASEAN; no official roster was retrievable, and the absences are as informative as the roster: India, the Gulf states other than Oman, Turkey and Egypt did not sign [764] [A]. What none of the three texts contains is the thing the sovereign-pact story requires. There is no compute commitment, no financing clause, no datacentre programme and no open-weights undertaking in the Action Plan, in the Initiative, or in the founding agreement. Xi Jinping’s WAIC keynote of 17 July pledged 5,000 training places for developing-country personnel over five years, regional “application cooperation centres,” and a meteorological early-warning system for thirty countries—and nothing else [764] [P]. An academic mapping of fifteen comparable international bodies scores WAICO lowest of all of them on formalization and describes it, at founding, as “a name, a host city, and a statement of purpose” [764] [A]. That is not a criticism this book needs to sharpen. The Universal Postal Union dues scale is the tell: it is the funding formula of an organization designed to exist rather than to spend.

22.3.2 The three live cases, graded

Three cases were live at 9 September 2026, and they are graded separately because they are not the same kind of fact.

Egypt is the case the story is now being told about, and it is a bid. Bloomberg reported on 26 August, from anonymous sources, that Huawei had bid to build Egyptian government AI datacentres—about 1,408 Ascend 950-series accelerators for a training cloud and roughly

600 for two inference clusters, a twelve-month build, reportedly with iFlytek as partner—which, if it closed, would be the first delivered Ascend installation abroad (Brazil’s Rio machine is contracted but not due until July 2027) [765] [A]. The US State Department responded by contacting NVIDIA, AMD and Microsoft to assemble a counter-offer and by warning Cairo that Ascend use might breach US rules, Egypt having needed a US licence for advanced AI chips since 2023. No price has been reported, no financing terms have surfaced, and nothing has been signed; Xi Jinping’s state visit to Cairo on 31 August–2 September, the obvious occasion for an announcement, produced no announced AI agreement [765] [A]. *Malaysia* is an evaluation. Telekom Malaysia is reported to be assessing the Ascend 910C as the backbone of a RM2B (about US\$494M) sovereign project whose purpose is to keep government, military and intelligence data in-country; the quantity is unclear, the project is unsigned, and Malaysia accepted export-control commitments in its 2025 reciprocal trade agreement with the United States, whose officials had already warned in 2025 that 910C use could violate US rules [766] [A]. The precedent inside the case is the one to weigh: in May 2025 a Malaysian deputy minister announced a 3,000-chip Ascend deployment and the government retracted it within hours. *Brazil* is a contract, and it is the only one. On 21 August 2026 Brazil launched an R\$2.3B (about US\$444M) AI supercomputing programme of which R\$1.3B (about US\$251M) goes to Huawei and iFlytek for a Rio de Janeiro machine on Ascend and a Portuguese-language model, from July 2027; a further R\$1B tender for a second machine in Rio Grande do Norte is expected to go to NVIDIA; the whole is financed by Brazil’s own FNDCT science fund, with 3,000 professionals to be trained over five years; and the stated principle of the programme is “not to depend on a single company, technology or country” [767] [A]. Two further entries belong in the ledger below the three. Kazakhstan’s 16 July package with China—“\$15B, seventy-plus agreements,” signed while Tokayev was in Shanghai for the WAICO founding—includes a strategic partnership between the Ministry of AI and Digital Development and Huawei, a Samruk-Kazyna agreement to acquire Huawei equipment, and a Kazakhtelecom–Hengtong agreement on “basic principles” for a 1 GW “Data Center Valley” at Ekibastuz with a first facility targeted for June 2027; project-level values are undisclosed and the package is a mix of binding agreements and memoranda [768] [A]. And Pakistan’s Punjab government agreed six datacentres and a thousand “AI master trainers” with Huawei on 28 August, with no money, no chips and no contract described [768] [A]. Table 22.3 lays the five out on the four questions that matter.

22.3.3 The two offers, side by side

The American offer must be held to the same standard, and it fares no better. Executive Order 14320, “Promoting the Export of the American AI Technology Stack” (23 July 2025), and the Commerce Department’s American AI Exports Program that implements it describe exactly what the Chinese wrapper does not: full-stack packages—hardware, data, models, security, applications—with federal financing referrals to direct loans, guarantees, equity and political-risk insurance, and an Export–Import Bank “ExportAI” framework under consideration since May 2026 [769] [P]. Seventy-eight consortium proposals had been received by 30 June 2026. No package had been designated at 9 September 2026 [769] [P]. So the ledger reads: the United States has a programme, a financing architecture on paper, seventy-eight bids and zero designations; China has no comparable programme, no financing facility, no designated packages because there is nothing to designate, and no confirmed Ascend export. *Neither offer has been delivered anywhere.* The contest in Cairo that the reporting describes as a proxy war between two stacks is, precisely stated, a contest between a bid and a counter-offer, neither of which exists in the field, for a buyer who has signed nothing with either [765, 769]. The buyer that *has* signed—Brazil—signed with both, financed itself, and said in terms that it was doing so in order not to depend on either.

Table 22.3: The Chinese AI export offer as it stood at 9 September 2026. One contract, one bid, one evaluation, two frameworks; no financing from the vendor’s side in any row [765, 766, 767, 768].

Country	What was offered	Status	Financing	US response
Egypt	Huawei bid: ~1,408 Ascend 950-series for training, ~600 for inference; 12-month build; iFlytek reported	Bid; unsigned; no price; no AI agreement from Xi’s Cairo visit	None reported	State Dept assembling NVIDIA/AMD/Microsoft counter-offer; warning that Ascend use may breach US rules
Malaysia	Ascend 910C as backbone of a RM2B (~US\$494M) sovereign project (Telekom Malaysia)	Under evaluation; unsigned; 2025’s 3,000-chip announcement retracted within hours	Buyer’s own budget	2025 warning on 910C; export-control commitments in Malaysia’s 2025 US trade agreement
Brazil	R\$1.3B (~US\$251M) to Huawei/iFlytek: Rio machine on Ascend plus Portuguese-language model, from July 2027	Contracted (21 Aug. 2026); second R\$1B machine expected to go to NVIDIA	Brazil’s FNDCT; no Chinese credit	None reported
Kazakhstan	Huawei strategic partnership; Samruk-Kazyna equipment purchase; 1 GW “Data Center Valley” on “basic principles”	MoUs and agreements; values undisclosed; first facility targeted June 2027	Undisclosed	None reported
Pakistan (Punjab)	Six datacentres; 1,000 “AI master trainers”	Agreement (28 Aug. 2026); no chips, money or contract described	None described	None reported

22.3.4 Where the telecom template transfers, and where it does not

Chapter 6 supplied the template that the sovereign-pact story is borrowed from, and closed by naming three of its elements that do not transfer; this section is where each is tested. The instrument that built Huawei’s network position was credit: a China Development Bank facility of US\$10B for Huawei’s customers in 2004, raised to US\$30B in December 2009, on which an operator in Lagos or Jakarta financed a switch it could not otherwise have bought. The condition that made the credit cheap was surplus: factories the state wanted running. And the object that held the customer was a bundle of hardware, service and debt that bound the operator for a generation. Test each against the AI record and each fails, and the failures are of different kinds.

The credit is absent. No China Development Bank or Export–Import Bank facility for AI infrastructure exports had surfaced at 9 September 2026. This is an absence in the record, not a verified absence—a facility could exist undisclosed, and Chinese policy finance is not transparent—but the one place where such a facility would have shown itself is the one contract that exists, and Brazil financed that contract from its own science fund. *The capacity runs the other way.* Huawei exported switches from surplus; it cannot export accelerators from surplus because there is none. DeepSeek’s reported order of at least 160,000 Ascend 950DT for a single inference site is expected to take up to a year to fill, because delivery is paced by Huawei’s and the domestic HBM chain’s output [755] [A]; a SemiAnalysis-derived estimate puts 2026 domestic HBM at enough for some 250,000–300,000 Ascend 910C-equivalents [A/C]; and Beijing’s reported requirement that state-funded compute run at least 80% domestic silicon under a roughly RMB 2T national programme pulls every chip the fleet produces toward the home market before any can be spared for Cairo [770] [C]. The same source records that Huawei was pitching low thousands of 910B to the UAE, Saudi Arabia and Thailand in July 2025 with no finalized deals, and that the Gulf went to American suppliers; and that BIS guidance since May 2025 holds that using Ascend “anywhere in the world” risks violating US controls, which is the legal instrument behind the warnings to Cairo and Kuala Lumpur [770] [P/A]. A vendor whose largest domestic customer absorbs a year of output, whose government reserves the output for domestic use, and whose product is a controlled item in the recipient’s jurisdiction is not running the 2009 playbook, whatever it intends. *And the artifact does not lock in.* The diffusion mechanism that actually operates in the non-aligned world is not a datacentre; it is a download. Hugging Face counts 151,448 derivative models of the Qwen family, about 2.6 times Meta’s, and roughly 39.6 million monthly GGUF downloads for Qwen against 7.5 million for Llama—something over five times the volume [771] [V]. Those weights cost the recipient nothing and bind the recipient to nothing; the same buyer that will not sign for Ascend runs Qwen already. That is a better diffusion instrument than Huawei ever had in telecommunications and a far worse capture instrument, and it is the instrument Chinese AI is actually being exported with.

The conclusion is stated plainly because the story it corrects is stated plainly everywhere else. At 9 September 2026 the sovereign-pact account of Chinese AI export is a *story*: the telecom template is visible in intent and in the institutional wrapper, absent in credit, absent in capacity, and being run—to the extent anything is being run—with an instrument, open weights, that diffuses better and captures worse than the one the template was built on. Three things cut against that conclusion and are printed with it. Intent precedes capacity in every precedent campaign in Part I; the CDB facility of 2004 was opened for a company that had been reselling imported switches seventeen years earlier, and the demand-before-supply move of Chapter 7 is by construction a commitment made before the factories exist, so the absence of a credit line in 2026 is evidence about 2026 and not about 2028. The Egypt bid is the first datacentre-scale Ascend export bid to a government on the record, and its outcome—signed on Chinese terms, signed on American terms, or split as Brazil split it—is the test this section will be scored against. And Brazil is not nothing: it is a real contract, on Chinese silicon, for a sovereign model in a G20 country’s language, paid for by a G20 government that had every alternative available. One such contract is not a campaign. It is, however, the first row of the table a campaign would be

measured in.

22.4 Nine positions

Japan holds the strongest technical hand of any non-principal and the weakest demand-side position. Assets: an exascale HPC lineage whose architectural argument this book carries (Chapters 12, 26); the most defensible remaining semiconductor materials and equipment franchises; leading-edge memory and packaging supply chains; the only measurement-anchored national AI4S demand model in existence; and unusual standing in both blocs' technical communities. The 2025–26 record shows the machinery engaging: the AI Promotion Act (May 2025)—Japan's first AI-dedicated law, deliberately innovation-first, penalty-free, and agile-governance-shaped, the philosophical antipode of the AI Act [408]; FugakuNEXT contracted to Fujitsu and NVIDIA (Aug. 2025) as a MONAKA-X Arm CPU tightly coupled to GPUs over NVLink Fusion, with a ~¥110B initial budget toward operations around 2030 [356]; the GENIAC compute-subsidy program in its fourth cycle; Sakana at a \$2.65B valuation; Rapidus running a 2 nm pilot line toward 2027 mass production; SoftBank converting a Sharp LCD plant into a 150–240 MW AI data center as the anchor of a domestic OpenAI venture; and a fiscal commitment moving from the ¥10T-through-2030 package toward a quadrupled annual chips-and-AI budget [409]. Liabilities, unchanged and now quantified: a domestic model ecosystem positioned for Japanese-language enterprise depth rather than the open frontier, an energy position that cannot win a watts race, and a research population whose committed provisioning sits an order of magnitude below its own modeled requirement. The honest tension the record exposes: FugakuNEXT's architecture buys NVIDIA coupling at the exact moment this book's compute chapter argues the accelerated-CPU corner is the sovereign long game—defensible as a hedge, dangerous as a habit. Recommended posture: build the accelerated-CPU line to its FP8-capable conclusion *alongside* the GPU coupling, not instead of sovereignty; adopt the best open weights irrespective of origin under a full trust stack; bid for the certification-authority vacancy in coalition rather than alone; and use the sizing framework as the diplomatic instrument it is—a shared unit is a claim on the terms of the debate.

The European Union has the pen and not the press—and in 2025–26 the pen visibly hesitated. The Digital Omnibus process pushed the AI Act's high-risk obligations back by sixteen months to two years (Annex III to December 2027, embedded systems to August 2028) while strengthening the AI Office's GPAI powers—an implicit concession that the regulatory clock had outrun the industrial one [410]. Meanwhile the industrial program scaled: thirteen AI factories plus antennas under a ~€10B envelope, and a formal EuroHPC tender (July 2026) for up to seven gigafactories of 100k-plus accelerators each, against 76 expressions of interest totaling ~\$230B—an oversubscription that says European capital wants sovereign compute even where European regulation slows its use [354, 411]. Mistral closed €1.7B at ~\$14B with *ASML* as lead investor—the continent's one true chokepoint firm buying into its one frontier-adjacent lab, sovereignty consolidating by cap table [412]—and JUPITER delivered Europe's first exascale machine. Liabilities: fragmented procurement, energy costs, no top-tier frontier laboratory, and the persistent gap between regulatory ambition and industrial capacity that the omnibus retreat priced publicly. Recommended posture, sharpened by the record: spend regulatory energy on the *conformity infrastructure* the world lacks—evaluation batteries, signed provenance, notified-body capacity for model artifacts, the one arena where the AI Act's machinery is an asset rather than a tax; buy the open commons and audit it; and treat the gigafactory programme as a sovereignty asset only if paired with the operating capability and the model programs to use it—seven empty cathedrals of compute would be the solar-tariff mistake with better architecture.

India is the largest genuinely undecided market and the one whose choice would most move the aggregate—and its 2025–26 record is the world's cleanest experiment in demand-side sovereignty. The IndiaAI Mission empanelled ~38,000 GPUs across three tenders at subsidized

rates reportedly under \$1/GPU-hour—among the cheapest sovereign compute on earth—and then discovered the harder half of the problem: reported utilization lagged badly, a national GPU pool in search of users [413]. The lesson generalizes to every country in this chapter and confirms this book’s cost-stack analysis: *compute is the easy term; c_{int} and c_{org} are the binding ones*, and a subsidy that buys hardware without buying integration capacity buys idle silicon. The sovereign-model program (Sarvam on 4,096 subsidized H100s, BharatGen, Krutrim’s frugal-AI line) and \$45B-plus of announced hyperscaler investment (Microsoft’s \$17.5B, Google’s \$15B hub, AWS) complete the picture: both blocs building the physical layer, India supplying the population. Recommended posture: aggressive adoption of open weights with domestic fine-tuning and Indic-language specialization; national evaluation capacity; fix the utilization gap before expanding the pool—and the highest-leverage move available remains becoming the world’s integration layer, since whoever performs the integration accumulates the complementary capital that the general-purpose-technology literature says captures the value.

South Korea deserves its own position rather than a footnote, because it is the only non-principal that is simultaneously a critical *supplier* to one bloc and a determined sovereign builder. The supplier position is decisive: SK hynix holds ~62% of HBM and reportedly two-thirds of the leading HBM4 volume—Korea sits on the choke half of the chokepoint this book’s compute chapter is about [414]. The sovereign program is the most structured in the world: a national AI computing center, a 260,000-GPU national deal spanning government and the chaebol “AI factories,” a ~100T-won public-private pledge—and, uniquely, a sovereign foundation-model *tournament*: five teams selected with compute and data support, then eliminated in public rounds, with 700B-parameter open models released as the competition sharpened [415]. That is playbook step 4 (selection, not planning) executed by a democracy, and this book registers it as the most interesting institutional experiment in national AI policy anywhere: state-funded competition with open-weight release as the qualifying bar. Recommended posture: keep the tournament honest (the temptation to anoint a national champion early is playbook step 6 (cartelize and gatekeep) arriving prematurely); price the HBM franchise as the strategic rent it is while it lasts; and hedge the memory cliff, because Korea is the country most exposed to exactly the capacity-first arbitrage of Chapter 12.

Singapore and ASEAN illustrate the small-state problem in its purest form. Singapore’s entire national AI4S requirement is smaller than one large country’s first procurement stage (Chapter 26), which makes autonomy through scale impossible and autonomy through *selection* essential: pick workloads where a small, excellent, well-instrumented capability beats a large generic one, and pool everything else—SEA-LION, the region’s open multilingual model family, is the template: small, targeted, open, and adopted because it serves eleven languages the frontier ignores [416]. ASEAN more broadly is the clearest case of the third-market dynamic—price-sensitive, sovereignty-conscious, and courted by both blocs with training, infrastructure and, on the American side, the promise of financing, with \$15B-plus of data-center investment landing in Johor alone. It is also where Chinese adoption has moved fastest, for the same reasons Chinese EVs did—and where the limits of that adoption at the hardware layer are most visible: Malaysia’s Ascend evaluation is unsigned, its 2025 predecessor was retracted within hours, and its 2025 trade agreement with Washington carries export-control commitments that a signature would test (Section 22.3) [766] [A].

The Gulf states are the anomaly: capital-rich, energy-rich, compute-poor by choice rather than constraint, and buying position in both stacks simultaneously at a speed neither principal fully controls. The 2025–26 record removed any doubt about scale: Saudi HUMAIN building toward 6.6 GW with NVIDIA “several hundred thousand GPUs” and an AMD \$10B/6 GW joint venture; Stargate UAE’s first gigawatt inside a planned 5 GW Abu Dhabi campus; MGX targeting \$100B of AI assets as anchor investor in the American frontier itself; and—the geopolitical hinge—US Commerce authorizing 70,000 advanced-chip exports to G42 and HUMAIN under security conditions, making the Gulf the first place Washington has *chosen* to locate frontier-scale capacity

it controls only contractually [417]. They matter to this book’s argument for one specific reason: they are the only non-aligned actors who can bid for frontier-scale capacity on the open market, which makes them the swing buyers in any repricing scenario (Chapter 21) and the most likely hosts of the first sovereign frontier training run outside the two principals—now with the silicon on order.

Taiwan is the position everyone else’s position depends on and the hardest one to hold: near-monopoly production of the advanced logic both blocs’ AI machines require, a \$3.2B “ten major AI infrastructure” program of its own—and an energy system, post-nuclear phase-out, whose arithmetic is the binding constraint on both its sovereign-AI ambitions and its foundry expansion [418]. Taiwan is this book’s busbar chapter compressed into one island: the world’s most valuable compute, gated by the world’s tightest watt. Recommended posture: treat energy as the semiconductor policy it is; and hold the TSMC franchise as the alliance instrument it has become, while building the sovereign-model capacity (TAIDE-line) that keeps dependence from becoming identity.

Brazil and Latin America combine large populations, real scientific communities, severe fiscal constraints, and—in Brazil’s case—a clean grid that is an underrated asset in an energy-constrained industry. The rational posture is almost pure commons adoption plus regional pooling, and the risk is the classic one: importing a service, building nothing, and discovering in a decade that the standards were set elsewhere. Brazil is, at 9 September 2026, the one government outside China to have executed the posture in hardware rather than in a strategy document, and the shape of what it executed is the instructive part: a programme split between an Ascend machine with a Portuguese-language model and an NVIDIA machine, financed from its own science fund rather than from either principal’s credit, with three thousand people to be trained on it (Section 22.3) [767] [A]. That is the six moves below, signed; it is also a single contract of a quarter-billion dollars, and the standards risk it hedges against is not answered by hedging the vendor.

Africa is treated in most Western AI strategy documents as a recipient and in Chinese AI diplomacy as a partner, and the difference in framing is itself strategic. The first facts on the ground arrived in 2026: Cassava’s NVIDIA-powered “AI factory” in South Africa—the continent’s first—with Kenya, Nigeria, Egypt, and Morocco expansions planned, pitched explicitly as sovereign data and compute [419]; and small open models for African languages (InkubaLM-class) following the SEA-LION template. The continent’s relevant facts are a young technical population, the fastest-growing developer base in the world, minimal legacy infrastructure to defend, and acute exposure on energy and connectivity. Chapter 27 argues the correct treatment is as a source of talent, scientific questions, and data—not aid—and notes that the institution which builds a real access track will acquire an asset that compounds for decades. Nobody has built it. The World AI Cooperation Organization is bidding—with, so far, a training pledge of five thousand places and no compute, financing or datacentre clause in any of its founding texts (Section 22.3); the one African government reported to be negotiating Chinese silicon, Egypt, is not a WAICO signatory [764, 765] [A].

22.5 What the matrix predicts

Apply Table 22.1 honestly and an uncomfortable pattern emerges. On *cost*, *sovereignty* (operational), and *export-control exposure*, the Chinese open stack currently scores better for almost every non-aligned actor—it is cheaper per task, it can be run on-premises without foreign permission, and no government can revoke a weight file already downloaded. On *capability at the absolute frontier*, *trust*, and *interoperability with existing enterprise systems*, the American stack scores better. On *standards* both create dependence, in different currencies. And on *industrial spillover*—the dimension that determines whether a country is building anything—*neither stack helps unless the adopting country does the work itself*, which is the single most important sentence

in this chapter.

That pattern was drawn before the second commons existed, and Section 22.2 moves three of its cells without overturning it. On *cost* the Chinese stack no longer holds an exclusive advantage, because a free artifact is free in either language; what remains of the difference is the hardware the artifact prefers, which is where it always was. On *operational sovereignty* the two are now close to level—a self-hosted Nemotron runs without foreign permission exactly as a self-hosted Qwen does—and the asymmetry survives only at the layer beneath, where one stack’s preferred accelerator is a controlled item and the other’s is designed not to be. On *export-control exposure* the gap therefore narrows at the model layer and persists intact at the compute layer, which sharpens the finding rather than reversing it. What has not moved is the row that decides the rest: on *trust*, neither commons supplies independent assurance about itself, and a vendor-sponsored artifact nobody has evaluated is exactly as unevaluated as a state-adjacent one. The vacancy identified above as the Chinese position’s strategic vulnerability is now a vacancy in both stacks at once—larger, more valuable, and harder for either principal to fill credibly, because in both cases the party best placed to certify the artifact is the party that paid for it.

That pattern explains observed behaviour better than allegiance does. It explains why institutions that ban a Chinese chat application nonetheless evaluate its weights for self-hosted deployment; why sovereignty-conscious governments buy American cloud and Chinese weights simultaneously; and why the fastest-growing adoption of Chinese open models is in exactly the countries with the least ability to pay rent. It also explains a strategic vulnerability of the Chinese position that follows directly from the same logic: the commons currently wins on cost and operational sovereignty but loses on trust, and trust is the dimension that a coalition of non-aligned states could supply *for themselves*—which would make the commons safe to adopt and simultaneously deprive its sponsor of the dependence it purchased.

22.6 The non-aligned strategy, stated once

For any country in this chapter, six moves follow from the matrix, and they are the same six regardless of size:

Measure your own demand before procuring anything, in a unit others recognize (Chapter 26); a country that cannot state its requirement cannot audit a vendor or join a federation. *Adopt the commons—now plural—and audit both its behaviour and its provenance*: the cost and sovereignty arguments are decisive, the trust gap is closable with infrastructure rather than with prohibition, and the provenance questions of Section 22.2 should be put to every artifact from either commons, because the licence file is the least informative document in the procurement pack. *Mirror what you depend on*—weights, datasets and evaluation harnesses, in a repository you operate—because every continuity risk in Table 22.2 is neutralized by a copy, and a copy is the cheapest sovereignty instrument in this book. *Retain a pretraining capability* sized to competence rather than to leaderboards, because the return is people and the alternative is permanent tenancy. *Pool internationally* for the capacity and the certification you cannot afford alone (Chapter 27). And *do the integration work domestically*, because the general-purpose-technology literature and this book’s Chapter 2 agree that complementary capital is where diffused value is captured, and it is the only part of the stack no foreign supplier can perform for you.

The countries that moved early in solar, batteries, and EVs are the ones the record remembers kindly. The window in AI is measured in quarters rather than decades, because the standards layer is being set now and the certification layer is still vacant.

Chapter 23

Strategies for the Rest

A structural analysis owes its readers agency, not just diagnosis. This chapter states what each class of actor should do conditional on this book being broadly right—and notes, where it matters, which moves remain correct even if Chapter 28’s falsification branches fire. The organizing principle throughout is the one the precedents teach: *do not defend the rent; position on the commons*.

23.1 The United States: win the commoditization race or lose it

The US strategy since 2019—chokepoint denial plus closed-frontier rents—has produced the outcomes this book documents: each denied input spawned a subsidized substitute, and the rent structure concentrated national exposure into the leveraged trade of Chapter 21. The counterfactual strategy has never been tried: compete *for* the commons rather than against it. Concretely: fund open-weight frontier models as public infrastructure (the gpt-oss reflex, sustained and resourced, rather than one defensive release); make the American open stack—not merely American APIs—the world’s default, because the adoption layer, once lost, prices everything above it; and attack the physical constraint that is actually binding—the watt (Chapter 14)—with the urgency currently spent on GPU export paperwork. The scarcest honest observation: the US capital structure, not its technology, is the strategic vulnerability—an orderly deflation of the scarcity premium, engineered early, is strictly preferable to defending it to the break.

That instruction assumed commoditizing and denying were separate policies that could be run in parallel. Chapter 17 shows they are not, and this is the sharpest finding in the book for American policy. An export-control regime and an open-weight corporate strategy are aimed at the same object from opposite sides. Open weights released *with their training data* by an American firm lower the cost of closing the model-layer gap for whoever is closing it, and Nemotron ships the datasets and the recipes alongside the weights [701] [P]. An open instruction set that the same firm has publicly qualified as a CUDA host lowers the cost of the compute-layer detour the controls exist to make expensive [673, 190] [P]. And a scale-up domain rebuilt out of merchant Ethernet—the majority design at Hot Chips 2026 among parts with an off-package fabric—is precisely the substitution the chokepoint strategy exists to prevent, performed by the firms it exists to protect [681] [P]. None is a concession to China; each is a commercial decision taken against a domestic rival, and each would have been taken in a world with no Chinese AI industry at all. The externality does not care about the motive: a denial policy premised on the denied party having to obtain the controlled capability from the controlling country is undermined as effectively by an American release as by a Chinese one, and the American release arrives with the corpus attached.

The policy conclusion is not that the firms should stop. That is the reflex, and it is wrong on this book’s own analysis: the same releases are what keep an American stack in contention for the non-aligned buyer of Chapter 22, and a government that suppressed them would be surrendering

the adoption layer to preserve a denial instrument whose effectiveness it cannot verify—the solar mistake with a national-security justification attached. The conclusion is narrower and harder. *A denial policy whose effectiveness is contingent on the conduct of firms the government does not control is not a strategy*; it is a hope with a licensing regime attached, and the United States must choose which of its two instruments it means. If it means denial, the controls must reach the conduct that erodes them—which runs immediately into the constitutional and practical difficulty of restraining publication, a difficulty to be priced honestly rather than assumed away, and into the Commerce Department’s own finding on the open ISA [197]. If it means commoditization, the controls should be triaged down to the few places the Chapter 8 model says they buy real time—EUV, HBM process, advanced packaging—and the energy freed spent on the watt and on the open stack. What cannot work is running both at full strength and expecting each to do its job, because in the three instruments above they cancel.

The counter-evidence bounds the finding and deserves the last word. The controls have accomplished something measurable: NVIDIA’s Hopper shipments to China ran under one per cent of datacentre revenue in the July 2026 quarter, and the forward outlook assumes *no* China datacentre compute revenue at all [727] [P]. Removing the dominant supplier from the second-largest market is not nothing, whatever else leaks. And the leak is smaller than the argument at full strength implies: the American open models shipped to date take routed volume without taking frontier capability, so what they have lowered is the cost of the volume tier and not the capability tier. The cancellation is real, partial, and incomplete—which is the condition under which a choice can still be made deliberately rather than discovered later in a filing.

23.2 Europe: regulation is not production

Europe enters the fourth war having sat out the previous three’s manufacturing and owning the regulatory pen. The trap is repeating solar: tariffs and standards atop a hollowed industrial base. Europe’s live assets are ASML (the one true chokepoint, wasting), Mistral and the open-model tradition, EuroHPC’s procurement scale, and the credibility to run the certification layer of Chapter 18 as a neutral. The strategy that follows: build sovereign inference and smaller sovereign-scale training on whatever open stack is best—which is now a genuine choice rather than a euphemism for Chinese weights (Section 23.4): accept both, audit both, mirror both; spend the AI Act’s institutional energy on the trust stack of Section 23.4, which the world lacks and Europe is culturally equipped to supply; and convert energy policy into compute policy, because Europe’s grid math resembles America’s with less gas.

23.3 Japan: the third-country playbook, first-person

Japan’s position is unusually concrete because its assets map one-to-one onto this book’s elements: frontier-class HPC engineering (the Fugaku-to-FugakuNEXT lineage whose architectural argument Chapter 12 carries [185]); the world’s most defensible remaining semiconductor inputs; memory at the leading edge; and deep bilateral standing in both blocs’ technical communities. Chapter 22 sets out the record and the liabilities in detail; five moves follow from the analysis. *First*, build the accelerated-CPU line to its conclusion—an FP8-equipped, capacity-rich successor in the LX2/Monaka direction is the machine class this book argues wins the converged AI-plus-simulation workload, and Japan is one of three polities that can execute it [M/S]. *Second*, adopt the open commons—now plural—with the full trust stack of Chapter 18, self-hosted, audited, mirrored and fine-tuned domestically; refusing Chinese weights on principle is strategy-by-flag, and taking an American vendor’s weights uninspected because they are American is the same error wearing the other flag. *Third*, bid for the certification-authority vacancy that Section 23.4 prices—convened with partners rather than alone, it is the highest-leverage per-yen investment in the entire space, the FAA role held by a party both blocs can transact with, and Chapter 27

develops its institutional form. *Fourth*, energy realism: Japan cannot win a watts race; it can win the performance-per-watt race, which is where its silicon and systems tradition already lives. *Fifth*, hedge the memory bet both ways—HBM partnership with allied suppliers and LPDDR-capacity architectures domestically, for exactly the dual-track reasons of Chapter 13.

23.4 The middle powers: what two commons are worth

Europe, Japan, Korea, Canada, Australia, the Gulf and the rest of Chapter 22’s electorate should read Chapter 17 as good news, and the reason deserves stating rather than leaving as an inference. *Two competing commons are a better buyer’s market than one*, for the reasons Section 22.2 prices row by row. What follows for strategy is not neutrality as a virtue but neutrality as procurement: take from both, run them side by side on one evaluation harness, and let the sponsors compete. The only thing that forfeits the position is exclusive adoption in either direction for reasons of alignment rather than of evidence.

The asset neither commons supplies is the one to build, and it is the asset this book has now named in four chapters: *evaluation and certification*. Both stacks ship free weights and neither ships independent assurance about them, because in both cases the party best placed to certify is the party that paid for the training run—an accelerator vendor defending silicon, or a laboratory operating under an industrial policy. That is a structural vacancy rather than an oversight, and the one layer where a middle power’s institutional credibility is the scarce input rather than its capital or its watts. Chapter 18 specifies the technical content and Chapter 27 the institutional form; Chapter 17 adds the market size, since the number of artifacts requiring assurance has just doubled while the number of credible assurers has not moved.

And one concrete opportunity has a date on it. The commoditization of the scale-up fabric onto merchant Ethernet materially lowers the cost of a sovereign build: the deepest single source of lock-in three years ago is now switch silicon anyone can buy, the majority design at Hot Chips 2026 among parts with an off-package fabric, with a merchant vendor guiding above \$16B a quarter to supply it [681, 691] [P]. A sovereign cluster no longer needs a proprietary interconnect licence to reach defensible scale—which changes the arithmetic of EuroHPC’s gigafactory tender and of every national machine specified between now and 2028 [411]. The date is the next procurement, not a research horizon, and the window shuts when those specifications freeze: a system bought on a proprietary fabric in 2027 carries the dependency for the life of the machine, and the FugakuNEXT tension Chapter 22 records is the specific form the mistake takes [356]. The counter-argument is real: merchant Ethernet at scale-up is newer, its software less mature, and the firm with the most non-NVIDIA accelerators in the field has just placed its next generation inside NVIDIA’s interconnect and memory stack [690] [P]. A buyer who reads that as a reason to wait is not being timid; it is the same evidence read for risk rather than for opportunity, and the falsification test below separates the two readings.

23.5 The Global South: the buyers write the ending

The precedent wars were decided, in the end, by third markets: solar in the desert tenders, EVs in Southeast Asia, telecom in the 4G buildouts. The AI war will be decided by the same buyers, and their rational strategy is uncomfortable for both blocs: take the free stack, keep the exit. Chapter 22 states the six moves that follow and they are not restated here; what Chapter 17 adds is that “the free stack” is now two stacks with different sponsors and incompatible commercial logics, which strengthens every one of those moves—a buyer with two suppliers of the same free good is a buyer with terms, and the trust-stack insistence that was a condition of adoption is now also a lever in a negotiation. WAICO’s 29 founding states have understood this faster than the G7 has [764]. Whether the commons they join has one certification authority or two—whether trust becomes a Chinese utility or a neutral one—is the open variable, and it is the same variable

Chapters 18 and this chapter keep landing on, because it is the war’s last genuinely unclaimed high ground.

One correction to the buyer’s own picture of the menu belongs here, because the menu is being described to ministries in terms that the record does not support. The buyer is not, at 9 September 2026, choosing between a financed Chinese stack and a financed American one. Neither exists in the field: the American programme has seventy-eight proposals and no designated package, the Chinese wrapper has a founding agreement with no compute, financing or datacentre clause, the one Ascend export bid is unsigned and the vendor’s output is spoken for at home for a year (Section 22.3). The telecom precedent that the sovereign-stack pitch borrows from ran on sovereign credit and surplus factories, and Chapter 6 records that neither has appeared in AI. So the practical choice a ministry faces this year is narrower and more tractable than the pitch implies: it is between open weights on whatever silicon the buyer can actually obtain—which for most of this chapter’s electorate means licensed NVIDIA or AMD parts, or CPU-class hardware, running Chinese or vendor-sponsored Western weights—and a proprietary stack on the same licensed silicon. Framed that way, the decision is the six moves of Chapter 22 and not an alignment question at all. And the only executed template is Brazil’s: split the programme between vendors so that neither owns the estate, finance it from your own science budget so that neither owns the debt, and write the training of your own people into the contract so that the spillover stays—a structure that, whatever else it does, no single foreign licensing decision can strand and no foreign bank can call [767] [A]. It is also one contract of a quarter-billion dollars, contracted for delivery in July 2027 and not yet delivered, and the Ascend half of it sits under the same US guidance that has been cited to Cairo and Kuala Lumpur; the template is executed, not proven. The buyer that waits for the financed stack to arrive is waiting for an offer that, on both sides, is still a press release.

23.6 China: advice for the presumptive winner

A book titled as this one is owes its subject the same service it gives everyone else: if the analysis is right about the mechanism, it is also right about the mechanism’s failure modes, and several are visible *now*, documented largely in Chinese sources. What follows is what Beijing should fix, stated with the same directness as the other sections.

First, the involution problem is real, recognized at the top, and not yet solved. The model-layer price war—repeated near-total API cuts cascading through Qwen, Zhipu, MiniMax, and the rest within days of each DeepSeek release—is playbook step 3 (overcapacity as strategy) running *without the step-5 rent-recapture plan attached*. Part I’s precedents survived below-cost phases because the losses purchased learning curves that led somewhere: cells, packs, cars. The model-layer war’s own leadership doubts the destination: the State Council has formally warned against “disorderly competition” in AI, the pricing-law amendments target involution by name, and Xi Jinping himself has publicly questioned whether every province needs an AI, compute, and EV project [420]. The fix is not to end competition—the Darwinian dynamic produced DeepSeek, and strangling it is falsification branch (e) of Chapter 28—but to move the war’s price floor from below-cost tokens to at-cost tokens plus *paid complements*: certification, deployment engineering, vertical solutions. The rent-migration table (Section 21.8) applies to China too, and Chinese labs are currently giving away row one while building almost nothing in rows two through five.

Second, the monetization gap is the strategy’s soft underbelly. China’s flagship labs generate strikingly little revenue relative to their US counterparts—the post-IPO disclosures of Zhipu and MiniMax show thin revenues and heavy losses, and the domestic enterprise market suffers a documented “low-trust open-source paradox” in which Chinese firms themselves hesitate to build on domestic open models [421]. Against Anthropic’s \$47B run rate, the entire Chinese frontier sector’s revenue is a rounding error. Nor is the problem confined to the startups: Alibaba—the

family with a billion model downloads and the strongest distribution position in the Chinese ecosystem—reported ex-items net income down 99.7% and adjusted EBITDA down 84% in the May 2026 quarter while committing to exceed a three-year, \$56B AI capex plan, with its chief executive stating that “margin is still secondary” [407]. That is the American hyperscaler bargain (spend now, monetize later) run by a firm whose model layer earns even less. The state can carry this indefinitely; the question is whether it should. A commons whose builders capture nothing eventually stops attracting the best builders—the talent flywheel runs on equity as well as patriotism, and Hong Kong listings at compressed valuations are a weak substitute for a functioning domestic growth-capital market. The fix is the one Beijing finds hardest: let winners charge. Three moves in the record already point this way and deserve to be read as policy templates rather than defections—Kimi K3’s revenue-threshold licence, Alibaba’s retention of a metered Max tier above its open families (Section 10.3.1), and DeepSeek’s introduction of peak-hour surge pricing on the cheapest frontier-adjacent model in the world. Rent relocated is the strategy working; rent abolished is the strategy eating itself.

Third, the compute buildout has repeated the classic overcapacity error, and candor about it would help. Independent reporting documents hundreds of state-subsidized data centers with utilization far below plan—up to 80% of some new capacity unused—rental prices collapsing, and the familiar provincial dynamic of building the designated thing regardless of demand [422]. This book’s own busbar casebook (Section 14.3) shows the western hubs working; it does not show all of them being needed. The fix is the one the involution campaign gestures at: consolidate ruthlessly, publish utilization, and let the East-Data-West-Computing framework allocate by measured demand rather than by provincial ambition—because idle capacity is not deterrence, it is depreciation, and the American bear case about Chinese AI (“it’s all empty data centers”) gains its plausibility from exactly this.

Fourth, the trust deficit is the binding constraint on the strategy’s global half, and it is self-inflicted where it is not structural. Chapters 18 and 19 documented what buyers of the commons must price: censorship layers that transfer into derivatives, provenance disputes, an unowned certification layer. Every quarter the certification vacancy stays unfilled, adoption of Chinese weights in regulated Western and non-aligned sectors stalls at the paying-workload boundary—the C919 pattern. The fix is counterintuitive for a security state but follows directly from the analysis: *sponsor the independent certification of your own artifacts*. Reproducible builds, signed manifests, third-party evaluation access, and behavioral-audit cooperation would cost little, could be led through WAICO or offered to a neutral consortium (Chapter 27), and would convert the commons’ largest liability into a completed sale. A China confident in its artifacts should want them audited; the inference buyers draw from the alternative is priced into every procurement this book has examined.

Fifth, the physical gaps deserve honesty rather than roadmap theater. HBM and advanced packaging remain the binding constraints on the accelerator fleet; EUV remains absent; and the analyst consensus puts CXMT roughly three years behind the Korean leaders at the class that matters [423]. This book’s compute chapters argue China’s best move is to *change the game* (capacity-first architectures, the hybrid tier) rather than chase the incumbent’s metric—advice this section now returns to its subject: fund the 5 TB/s knee test and the hybrid tier’s placement policy (Appendix A) at national priority, because they are cheap, decisive, and strategy-defining, and stop announcing HBM schedules that the record keeps missing, because each slipped roadmap feeds the skeptics’ case at home and abroad.

Sixth, submit to the scoreboards—and finish the auditability you have started. No Chinese-designed accelerator has ever appeared in MLPerf Training; no LX2-class machine has published the AI proof suite; the embodied fleet publishes shipments, not intervention-free hours. The model layer shows both halves of the problem at once, and the good half deserves crediting first: Qwen3.8-Max shipped with a release blog carrying the full benchmark table *and a complete project trace on GitHub* for its flagship autonomy demonstration [393]—which is more verification

apparatus than most Western frontier claims carry, and exactly the practice this book has been asking for. The gap is what sits beside it: no architecture report, no activated-parameter disclosure, no training-hardware statement, no hallucination rates of the kind the prior generation published, and—most consequentially—an autonomous silicon-design claim (Section 10.3.5) that is potentially the most strategically important assertion any Chinese lab has made, shipped without any artifact a third party could check. The recommendation follows directly: publish the design flow’s audit trail the way the coding run’s was published. A verified AI-designed part would do more for China’s position than any benchmark table, because it is the claim the West would most like to disbelieve and the one that, if true, dissolves the chokepoint strategy entirely. Every unaudited claim is discounted by the buyers who matter; this particular one is being discounted at a rate that costs China real strategic credit. The pattern—dominate the metrics you control, avoid the ones you do not—is rational for a challenger and corrosive for an incumbent, and China is becoming an incumbent. The strategy that wins the next decade is the one this book has attributed to China at its best: *show the receipts and let the cost curve argue*. It worked for solar. It will work here, if the receipts are real—and if they are not, better to learn it before the world does.

Seventh, and newest: the principal move is now being run against you, and the advice above assumes it is not. Every recommendation above presumes that commoditizing the model layer is China’s instrument and the West’s problem. Chapter 17 records the year that stopped being true: an American accelerator vendor now releases open weights with their training data, has bought a training platform and 109 of the researchers who ran it, and lists the open-model ecosystem among its own demand drivers [701, 724, 728] [P]. If the West commoditizes that layer competently, China’s advantage there compresses—not because Chinese models get worse, but because free stops being a differentiator when the artifact beside it is free and ships its corpus too. The consequence is a reallocation rather than a defeat: China’s durable edge moves to the layers the West is conspicuously *not* commoditizing—firm deliverable power at the busbar (Chapter 14), the deployment-generated data only a country running the machines at population scale accumulates (Chapter 15), physical production and the embodied fleet (Chapter 16), and the domestic silicon floor no foreign decision can withdraw (Chapter 13). Those are the rows to fund, and three of the recommendations above—consolidate the idle capacity, let the winners charge, submit the artifacts—are what makes funding them affordable. The counter-evidence belongs in the same paragraph: a West that commoditizes the model layer also *validates* the open-weight strategy and enlarges the ecosystem in which China is the larger supplier, and in the last week of August three of the four significant open-weight releases were still Chinese [732] [V]. The compression is prospective rather than observed—which is exactly the kind of thing a presumptive winner notices too late.

23.7 The common instruction

Reduced to one line per actor: the US should choose between its two instruments and, having chosen, commoditize before being commoditized; Europe should certify what it cannot manufacture; Japan should build the machine class the war is converging on and referee the trust layer; the middle powers should take from both commons and certify what neither will; the Global South should take the commons and unionize the demand side; and China should let its winners charge, consolidate its overcapacity, and submit its artifacts to scoreboards it does not control. For every actor except the two principals, those instructions converge on one institution, which Chapter 27 specifies. And all of them share one instruction the precedents make non-negotiable: read the indicator table of Chapter 21 quarterly, because every strategy above has a branch point there, and the actors who moved early in solar, batteries, and EVs are the only ones the record remembers kindly.

Which of these strategies is right is not a matter of taste, and the register in Appendix J

carries the two tests that settle it. If merchant-GPU datacentre revenue growth decelerates while custom-ASIC revenue holds the trajectory Broadcom has guided, the commoditization of the silicon layer is working, the pincer of Chapter 17 is real, and the advice above is calibrated correctly: the American choice becomes urgent, the middle powers' sovereign-build window is open, and China's advantage compresses at the model layer while holding at the physical ones [691]. If instead custom silicon keeps converging onto NVIDIA's interconnect and memory—as Trainium4 has already done, adopting NVLink Fusion and NVIDIA custom HBM [690] [P]—then the wave is being absorbed, the accelerator vendor priced its giveaway of the model layer correctly, and a sovereign build on merchant parts is a more expensive bet than it looks. Both tests resolve by the end of 2027—inside the planning horizon of every actor in this chapter and outside the attention span of most of them.

Part IV

AI for Science: The Largest Prize

Chapter 24

The Largest Prize: Why AI for Science Dominates the Value Case

This book has so far treated AI as an industrial commodity and asked who will supply it. This chapter asks what it is *for*, and argues a claim that most of the public debate has inverted: the largest value in artificial intelligence is not in writing software, drafting documents, or automating enterprise workflow. It is in accelerating science. Everything else redistributes an existing surplus; science creates new ones. If that ranking is right, then the two chapters that follow—how much compute a nation’s research actually needs, and what institution could supply it—are not a digression from this book’s argument but its destination, and the competition described in Parts I–III is a competition over who gets to compound their scientific output fastest.

The chapter proceeds in the order that honesty requires: first the demonstrated results, then an unsparing ledger of what has *not* worked, then the economics that make the prize large even after the failures are subtracted, then the strategic consequence.

24.1 What has actually been demonstrated

Four domains have produced results that survive scrutiny, and their common feature is instructive: in each, AI replaced a computation that was previously bounded by physics-based simulation cost, and the replacement was validated against an established operational baseline rather than against a benchmark of its own devising.

Weather is the strongest case, and the only one with operational adoption by three national agencies. ECMWF’s AIFS became operational in February 2025 at 28 km grid spacing with gains up to 20% over the physics-based model on measures including tropical-cyclone tracks, at roughly *one-thousandth* of the energy per forecast; its 51-member ensemble followed in July 2025 [307]. NOAA deployed AI-driven global models operationally in December 2025: a 16-day forecast in about 40 minutes using *0.3%* of the compute of its conventional system, with the ensemble variant extending useful skill by 18–24 hours [308]. Microsoft’s Aurora—1.3B parameters, pretrained on 32 GPUs in about two and a half weeks—beats the operational high-resolution model on 92% of targets and matched or beat all seven operational centres on every tested cyclone track, at roughly 10^5 times the speed of the atmospheric-chemistry system it replaces; adapting it to a new task takes four to eight weeks against the several years a new dynamical model requires [309]. NeuralGCM runs on a single accelerator what previously required 13,000 CPUs [310]. China’s meteorological administration has its own operational family, and a 2026 typhoon-track system reports up to 32% error reduction against leading global ensembles [311, 312].

Structural biology is the most diffused case. AlphaFold has predicted over 200 million structures, is used by more than three million researchers in 190 countries—including over a million in low- and middle-income countries—and won the 2024 Nobel Prize in Chemistry [313, 314]. The

measurable second-order effect matters more than the structures: researchers using it increased their submissions of *novel experimental* structures by over 40% and were twice as likely to be cited in clinical work [313]. The database now includes millions of predicted complexes, representing on the order of 17 million GPU-hours of computation given away [315]. Protein *design* has followed: a generative model produced a fluorescent protein 58% identical to its nearest natural relative—an object evolution had not found—and diffusion-based enzyme design has produced metallohydrolases with catalytic efficiency orders of magnitude above previous designed enzymes [316, 317].

Mathematics and algorithm discovery produced the first credible autonomous results. A frontier system achieved gold-medal standard at the 2025 International Mathematical Olympiad, solving five of six problems end-to-end in natural language within the human time limit [318]. More consequentially for this book, an evolutionary coding agent found a 48-multiplication algorithm for 4×4 complex matrix multiplication—the first improvement on that case since Strassen in 1969—improved the best known solution on 20% of fifty-plus open problems it was set, and returned measurable production value: 0.7% of a hyperscaler’s worldwide compute recovered through better scheduling, a 23% kernel speedup yielding a 1% reduction in the training time of the company’s own frontier model [319]. That last result is the flywheel of Section 24.3.4 in miniature: AI improving the infrastructure that trains AI.

Drug discovery has its first controlled clinical datapoint. An AI-discovered TNIK inhibitor for idiopathic pulmonary fibrosis reported Phase IIa results in *Nature Medicine* in June 2025—71 patients, 12 weeks, a +98.4 mL mean forced-vital-capacity change against −20.3 mL for placebo—and entered a 320-patient Phase III in July 2026; the developer reports reaching preclinical candidate nomination in 12–18 months with 60–200 molecules synthesized per programme, against a conventional 2.5–4 years [320, 321]. This is one drug, one surrogate endpoint, one country, and not an approval. It is also more than the field had two years ago.

Mathematics crossed a threshold in August 2026. OpenAI reported that its next-generation model, Astra, produced solutions to ten long-standing open problems in mathematics and quantum complexity theory for roughly \$2,000 of inference compute, and—the part that matters for this book’s epistemics—published the proofs for independent verification rather than the claim alone [481]. At this edition’s date the proofs are under community review and none has been refuted; the appropriate grade is [P] moving toward [V] problem by problem as verification completes. Two readings coexist. As AI4S evidence it is the strongest single capability signal yet in the one domain where validation is cheapest—a proof either checks or it does not—which is precisely why mathematics leads the funnel’s conversion ordering. As competitive evidence it cuts against any complacent reading of this book’s thesis: the recursive loop of Section 28.2.1 runs through exactly this kind of result, and it happened on closed American compute, not on the commons.

24.2 The failure ledger

A chapter arguing that AI4S is the largest prize is obliged to lead with the evidence against itself, and in this field the counterevidence is unusually sharp.

The most-cited economic study of AI in scientific discovery was fabricated. A 2024 paper reporting large gains in materials discovered, patents, and product innovation at a real firm was repudiated by its institution in May 2025—“no confidence in the provenance, reliability or validity of the data”—and the two prominent economists who had publicly praised it jointly stated that its findings should not be relied upon [322]. It is not cited here except as a warning, and any AI4S argument that still rests on it should be discarded.

Materials discovery, the flagship domain, has not survived validation intact. The autonomous-synthesis result that anchored the field was formally corrected in January 2026—successes cut from 40 to 36, one compound removed because it was in the training data, and the authors conceding the

materials were “new to the prediction platform, not necessarily new to science” [323]. Independent crystallographers argued that none of the originally reported compounds was genuinely new [324]. The 380,000-stable-crystal result has been criticized by senior materials chemists for producing “chemical compounds rather than materials” with no demonstrated functionality [325]. The strongest generative-materials paper validated exactly one compound experimentally, and it missed its target modulus by about 20% [326].

Even weather has a documented failure mode. In May 2026 ECMWF discontinued real-time running of all four external AI models, having found them increasingly sensitive to the exact initializing analysis—several degraded when initialized on an upgraded analysis system—and noting that two of them do not forecast precipitation at all [327]. The lesson is not that AI weather models failed; the centre’s own AIFS remains operational. It is that these models are couplings to an underlying data-assimilation system, not free-standing oracles.

Autonomous discovery is far weaker than its marketing. When a frontier agent was run against 700 open problems in one mathematical corpus, human adjudication of 200 candidate solutions found 68.5% fundamentally flawed and only 6.5% meaningfully correct—of which two were genuine autonomous resolutions [328]. A leading mathematician’s assessment of a widely publicized single success was that only 1–2% of the open problems in that corpus are tractable this way, that difficulty within it varies by orders of magnitude, and that unreported failures skew the perceived rate [329]. The first conference to accept AI as primary author found submissions “technically correct” but “neither interesting nor important,” with reviewers noting that technical fluency can mask poor scientific judgement [330]. Working scientists interviewed about frontier models in early 2026 were blunt: “We have not found, yet, that LLMs are fundamentally changing the way that science is done” [331].

And the scientific commons is absorbing real damage. One journal measuring its own pipeline found submissions up 42% since late 2022, over 30% of peer reviews showing detectable AI use, and papers with heavy AI content receiving revise-and-resubmit at 3.2% against 11.9% for low-AI papers—while the editorial board nearly doubled to cope [332].

The honest summary: AI4S has delivered decisively where the task is *a well-posed prediction problem with an expensive physics-based incumbent and an established validation baseline*, and has so far mostly failed where the task is *judgement*—choosing which question is worth asking, deciding what counts as new, or assessing whether a result matters. That boundary is the most useful thing known about the field, and it has a direct architectural implication taken up in Section 24.6.

24.3 Why the prize is nevertheless the largest one

24.3.1 The denominator

Start with the sizes of the things being compared. Global research and development spending ran approximately \$3.1 trillion (PPP) in 2022, with the United States at \$923B and China at \$812B [333]. The entire worldwide software market—the addressable universe of “AI writes code” and “AI streamlines the enterprise”—is forecast at \$1.47 trillion for 2026, roughly half of global R&D [334]. And the combined revenue of the frontier laboratories whose valuations Chapter 21 examined is under 2% of annual global R&D spending [335]. Even the most-cited bullish estimate of generative AI’s economic potential put *R&D as the single largest affected function* at about a quarter of the total, ahead of software engineering—while explicitly excluding value from entirely new product categories, which is precisely where scientific advance shows up [336].

24.3.2 The return

Research spending is not an ordinary input. The best available estimate of the social return to R&D is \$13.3 of social benefit per dollar spent at a 5% discount rate—and even under a deliberately

punitive 20-year-lag assumption, \$4.9 per dollar; the authors note it is difficult to construct an average return below \$4 [337]. No enterprise-software deployment has a multiplier in that range, because enterprise software largely reallocates surplus between firms while research generates surplus that is mostly not captured by whoever paid for it. That non-appropriability is exactly why public institutions fund research, and exactly why the AI4S case is a *public-infrastructure* case rather than a product case—a point Chapter 27 builds an institution on.

Three accounting disciplines keep this subsection honest. *First*, the \$4–13 figure is an *intertemporal present value* of the social surplus stream generated by a dollar of R&D, not an annual multiplier; it must never be multiplied directly against a single year’s R&D flow to produce an annual benefit. The correct chain is: an annual *productivity uplift* to the research enterprise (the funnel’s output, below) → an increment to the effective R&D flow → a *present value* of social surplus obtained by applying the Jones–Summers ratio to that increment. The three quantities—annual uplift, incremental output flow, and present value of surplus—are kept separate wherever this book uses them. *Second*, the ratio is an *average* return; AI will first reach the most computationally tractable tasks, whose marginal return may sit above or below that average, and nothing in the current evidence pins down which. *Third*, the \$3.1T denominator is PPP-adjusted while software and laboratory revenues are nominal; the comparison is strategically suggestive, not an accounting identity. None of these disciplines shrinks the qualitative conclusion—the base is enormous and the multiplier is large—but they change what may legitimately be computed from it.

24.3.3 The bottleneck

Now the argument that makes AI4S strategic rather than merely valuable. Research productivity is falling. The canonical measurement finds aggregate US research productivity down by a factor of 41 since the 1930s—more than 5% per year—sustained only by a 23-fold increase in research effort; keeping Moore’s law required more than eighteen times as many researchers per doubling as in the early 1970s; cancer clinical-trial research productivity fell by a factor of 4.8 between 1975 and 2006 [338]. Ideas are getting harder to find, and the traditional remedy—hire more researchers—is exhausting itself demographically in every advanced economy.

An important 2026 counter-interpretation deserves its place here, because it redirects rather than refutes the argument. Fort and coauthors find that patent production per unit of R&D input may *not* be declining at all; what has weakened is the conversion of ideas into firm growth—the diffusion step, not the discovery step [457]. The distinction matters strategically. If ideas are genuinely harder to produce, AI’s leverage is on discovery itself, and the funnel below is the right model. If ideas are produced but diffuse poorly, the binding constraint is institutional—standards, competition, technology transfer, absorptive capacity—and the highest-return intervention is the deployment and certification infrastructure of Chapters 18 and 27 rather than more model capability. The truth almost certainly varies by sector, which is one more reason the consortium proposal pairs compute provision with diffusion machinery instead of treating capability as self-executing. Either way the case for the investment survives; what changes is where the marginal dollar goes.

An input that raises research productivity therefore acts on the binding constraint of long-run growth, which is why the economics of AI turns on this question and not on call-centre automation. And here is the pivot of this chapter. The most influential *skeptical* estimate of AI’s macroeconomic effect—a total-factor-productivity gain of about 0.66% over ten years, roughly one-tenth of bullish estimates—reaches that number by counting only task-level automation of existing work, and *explicitly excludes* transformative science-driven gains on the stated grounds that advances of the protein-folding type “do not seem likely within the 10-year time frame” [339]. That exclusion is the entire disagreement. Structure-prediction, operational AI weather, and designed enzymes all arrived inside the window. One does not have to accept the most aggressive growth models—one recent analysis finds that automating as little as 13% of research tasks

economy-wide is sufficient to enter an explosive-growth regime [340]—to notice that the credible range of estimates for AI’s contribution to annual growth spans from 0.1% to 30%, and that the single variable separating those poles is whether AI automates *innovation itself* [341].

24.3.4 The flywheel


The final reason AI4S outranks the alternatives is that it is the only application category that feeds back into the means of its own production. Better materials improve batteries, magnets, and semiconductors; better semiconductor design improves accelerators; better weather and grid modelling improves the energy system that Chapter 14 identified as the binding physical constraint; better algorithms improve training efficiency directly, as the 1%-of-training-time result cited above demonstrates in the smallest possible way. Each of these is an input to the war this book describes. Enterprise-workflow automation has no such loop: a faster expense-report pipeline does not lower the cost of the next accelerator generation.

24.4 The value-conversion funnel

The claim that AI4S is the largest prize needs a conversion model, not just a denominator, because value does not pass unchanged from capability to output. The chain runs

AI capability → task acceleration → validated discovery → deployment → productivity,

and each arrow has a conversion efficiency well below one. Stating them explicitly—even as ranges—is the difference between an argument and an enthusiasm.



The corner-to-corner product of these ranges spans $U = a(s-1)vd \approx 0.45\%$ to 59%; the task-weighted Monte Carlo of §24.4.1, which is the operative estimate, concentrates this into a central $\approx 11\%$ with a 10–90% range of 8–18% [M]. Even the pessimistic corner is a socially consequential annual research-output flow ($\sim \$28\text{B/yr}$ against the \$3.1T base), whose present value exceeds the cost of national AI4S infrastructure by a large multiple. What the claim does not survive is the assumption that acceleration equals discovery, which is why the funnel is drawn rather than asserted.

Figure 24.1: From capability to output, with the losses shown. Each arrow carries a conversion efficiency well below one, and the domains where AI4S value is already visible in the statistics (weather) are precisely those where all four conversions are favourable—high addressability, a computational bottleneck, an established validation baseline, and short deployment lag. Where any one fails, as in materials synthesis, headline discovery counts do not become value. Ranges are author estimates [M]; the purpose is to make the argument falsifiable term by term rather than to claim precision.

Four disciplines govern the funnel. *First, addressability*: not all of the \$3.1T R&D denominator is reachable by current or near-term AI. The reachable share is dominated by the task classes where the evidence above is strong—prediction with an established validation baseline, and search or optimization over well-defined spaces—and is much thinner for autonomous experimental planning and thinner still for scientific judgement and theory formation. A defensible current addressable share is on the order of 15–35% of research *effort*, weighted toward computational and screening-heavy fields [M]. *Second, acceleration is not discovery*: a hundred-thousand-fold

speedup in a forward model, as in weather, converts into scientific value only where the bottleneck was that computation; where the bottleneck is experimental throughput, funding, or judgement, the speedup converts at a small fraction. *Third, validation attrition is severe and measurable:* the materials record of the previous section is the cautionary case, where headline discovery counts survived external scrutiny at rates well below one. *Fourth, deployment and regulatory lag truncate everything:* a validated drug candidate is a decade from revenue, a validated material perhaps half that, a validated forecast improvement is operational within a year—which is why weather is the domain where AI4S value is already visible in the statistics and pharmaceuticals is the domain where it is not yet.

24.4.1 The arithmetic, done properly

It is often said that multiplying these ranges yields a 3–30% uplift. It does not, and the error is worth walking through, because the corrected arithmetic is what the rest of this section rests on. Define the output variable precisely: the *incremental* uplift to realized research output from AI, for a single homogeneous task pool, is

$$U = a(s - 1)vd, \quad (24.1)$$

whose four factors are the four arrows of the funnel: a is *addressability*, the share of research effort that current AI can reach at all; s is the *speed-up* on that reachable share, so that $(s - 1)$ counts only the acceleration *beyond* the unassisted baseline (a model that changes nothing has $s = 1$ and contributes zero); v is *validation survival*, the fraction of accelerated output that withstands independent scrutiny; and d is *deployment conversion*, the fraction of validated results that actually reach use. All four are dimensionless fractions except s , which is a multiplier. The printed endpoint ranges then give $U_{\min} = 0.15 \times 0.5 \times 0.20 \times 0.30 = 0.45\%$ and $U_{\max} = 0.35 \times 4 \times 0.60 \times 0.70 = 58.8\%$; the gross converted output $U^{\text{gross}} = asvd$ spans 1.35–73.5%. The corner-to-corner interval is therefore roughly *half a percent to sixty percent*, not 3–30%, and the wider band is the honest one: it is what independence of the terms actually implies.

The corners, however, are not equally credible, because the parameters are correlated in a specific, evidence-backed way: domains with high acceleration and high addressability (weather, structure prediction) also tend to have strong validation infrastructure and short deployment lags, while domains with weak validation (materials synthesis) also carry long deployment lags. The all-pessimistic and all-optimistic corners are thus doubly unlikely. The right formulation—and the one adopted here—is task-weighted:

$$U_{\text{AI4S}} = \sum_j w_j a_j (s_j - 1) v_j d_j \ell_j, \quad (24.2)$$

where j indexes task classes, w_j is each class’s share of global research effort, and ℓ_j is a lag-and-friction discount absorbing calendar delay, organizational adjustment, crowd-out, and regulatory latency—the term the single-pool version silently set to one. Table 24.1 states the class-level judgments [M] that the evidence of this chapter supports; this “jagged scientific production function”—strong augmentation of search and prediction, weak automation of judgment—is the same structure the emerging economics literature on AI in science models formally [456].

A table of this kind fails silently if it omits the deployment column d_j —an omission that sets deployment conversion to one in the printed contributions without ever saying so—which is why the discipline here is that every published number be generated from published tuples. Table 24.1 complies: each ΔU_j is the exact product of its row, the central total is their exact sum (10.7%), and the pessimistic/central/optimistic triple for every parameter lives in a machine-readable CSV shipped with the manuscript source, so any reader can recompute or re-weight the model in one spreadsheet.

Table 24.1: The funnel by task class [M], now fully reproducible: every class prints its complete central tuple $(w_j, a_j, s_j, v_j, d_j, \ell_j)$ and the contribution $\Delta U_j = w_j a_j (s_j - 1) v_j d_j \ell_j$ computed from it exactly. Pessimistic and optimistic values for every parameter, and the formulas, are published in the machine-readable register (`funnel-model.csv`). Weights sum to one.

Task class	w_j	a_j	s_j	v_j	d_j	ℓ_j	ΔU_j
Prediction & surrogates	0.20	0.50	3.0	0.70	0.70	0.70	6.9%
Search & optimization	0.20	0.40	3.5	0.40	0.60	0.60	2.9%
Experimental planning	0.20	0.25	2.5	0.40	0.60	0.50	0.9%
Interpretation & causal inference	0.25	0.15	1.5	0.30	0.50	0.35	0.1%
Question selection & theory	0.15	0.05	1.5	0.20	0.30	0.20	0.0%
Central total							10.7%

The Monte Carlo behind that table is run rather than promised. Sampling each parameter from a triangular distribution over its (pessimistic, central, optimistic) triple, with the three conversion terms (v_j, d_j, ℓ_j) coupled within and across classes through a common “conversion-quality” factor (mixing weight 0.5)—the correlation structure the evidence supports, since domains with good validation also deploy faster—200,000 draws give: *median 11.8%, 10th–90th percentile 7.7–17.5%* [M]. One-at-a-time sensitivity ranks the classes: the prediction-and-surrogates class dominates (swinging the total from 4.6% to 42.9% across its own corner tuples), search-and-optimization second (8.0–27.7%), and the two judgment-heavy classes barely move the total at all—which is the quantitative form of this chapter’s qualitative claim that the prize lives in prediction and search, not in automated scientific judgment. The headline is therefore restated once more, and more sharply than the endpoint arithmetic allowed: *a task-weighted central estimate of $\approx 11\%$, a likely (10–90%) range of $\approx 8\text{--}18\%$, and coherent all-pessimistic/all-optimistic scenario corners near 1% and 70%*. The earlier “3–30%” band survives as roughly the outer decile-to-percentile region between the likely range and the corners; the corners of Eq. (24.1) survive only if the correlations are assumed away.

What may legitimately be concluded from the lower end, stated in the stock-and-flow discipline this chapter enforces (comparing a present value against an annual industry revenue mixes the two, and the temptation to do so is a standing one): at the pessimistic scenario corner of the task-weighted model—which is $\approx 0.9\%$, not the 0.45% single-pool corner of Eq. (24.1), because the weighted sum cannot put every class at its worst value simultaneously without contradicting the correlations— $0.9\% \times \$3.1\text{T} \approx \28B per year of incremental research-output flow; and at the 10th-percentile Monte Carlo floor of 7.7%, $\sim \$240\text{B}$ per year. *Even the lower end of the modeled uplift represents a socially consequential economic flow, and its present value—applying the Jones–Summers ratio to the increment—exceeds the cost of national AI4S infrastructure by a large multiple.* That is the quantitative form of this chapter’s claim. What it does *not* survive is the assumption that acceleration equals discovery, which is why the funnel is drawn.

24.4.2 The measured analogue: what national-scale HPC actually returned

The funnel’s parameters are author judgments, and a fair reader will ask whether *any* part of this value chain has ever been measured rather than modeled. One body of evidence exists, and it comes from the adjacent field with three decades of institutional practice: the ROI literature on national high-performance computing, which is the closest measured precedent for what AI4S economics will look like, because leadership HPC is the same asset class—publicly funded computational infrastructure applied to science—one technology generation earlier.

The most complete such study available to this author is Hyperion Research’s multi-year assessment of RIKEN’s national computing infrastructure and its planned successor [458]. Two disclosures belong directly beside that sentence: the study was *commissioned by RIKEN*, and

the author of this book directs the RIKEN center whose infrastructure it assesses. The reader should weigh everything below accordingly. Its grade, corrected downward from a plain analyst estimate, is [A/M]: what the study reports is not a direct causal measurement of the modeled AI4S quantity but *a commissioned, survey-based, counterfactual economic-impact model informed partly by realized scientific projects*.

Its headline findings: Fugaku’s COVID-19 response projects were estimated to have returned \$50–75B to the Japanese economy; 65.7% of surveyed Fugaku researchers reported already-realized useful results on projects targeting major societal challenges; in an innovation-class mapping of more than 650 HPC projects worldwide, RIKEN’s portfolio concentrated in the highest impact-and-importance classes; and the projected returns for the AI-capable successor machine—precisely because AI capability widens the addressable share a_j and the acceleration s_j of Eq. (24.2)—were estimated at \$250B (pessimistic), \$400B+ (realistic), and \$600B+ (optimistic) over its life, two to four times Fugaku’s, against an investment two orders of magnitude smaller. The same study put state-directed five-year AI commitments at \$100–200B for China against \$18–27B of US federal spending—the private-public asymmetry that Chapter 21 prices—and \$6–10B for Japan.

Case study: Methodology: what is observed, what is modeled, in an HPC ROI study

Impact studies of this genre chain together components of very different evidentiary strength, and using one honestly requires separating them. *Observed*: that specific projects ran, produced outputs, and reached users—and beneficiary self-reports of usefulness (the 65.7% figure is this class: real reports, subject to self-selection and social-desirability bias). *Modeled*: the counterfactual (what would have happened without the machine); the attribution of outcomes to the computing infrastructure rather than to the researchers, data, and institutions around it; avoided-loss valuation (the COVID \$50–75B is this class—an estimate of economic collapse *not* suffered, the least verifiable quantity in economics); spillover estimation; and all projected future benefits (the \$250–600B scenarios are projections, full stop). *Unquantified risks*: double counting across projects, displacement (benefits that would have accrued elsewhere), deadweight (projects that needed no subsidy), and the absence of published confidence intervals. A commissioned study has, additionally, a selection gradient in what gets surveyed and how questions are framed. None of this makes the study worthless—it makes it [A/M], and it is why triangulation with independently governed evaluations matters: the UK’s ARCHER2 evaluation, commissioned under UKRI’s standard evaluation framework with explicit counterfactual scenarios and benefit-cost accounting against public investment, is the best current methodological benchmark for the genre [368], and a three-way Fugaku–ARCHER2–CERN-class comparison of evaluation methodologies is now on this book’s living-appendix work list. The claim this book actually rests on survives the demotion: national-scale computational infrastructure has repeatedly been *evaluated*—by methods far from perfect but more disciplined than anything the enterprise-AI ROI literature offers—at returns that exceed costs by large multiples, concentrated in exactly the task classes where Table 24.1 puts the value.

What the evidence is good for, at the grade it actually carries: it is *partly ex post* (projects ran; users reported results), it is *independent of this book’s thesis* (the analysts were sizing a Japanese procurement, not arguing about China), and its *ratio structure* matches the funnel—the reported value clusters in prediction, simulation, and search applied to societal-scale problems with operational deployment paths, not in the judgment-heavy classes where the funnel assigns near-zero. It is supporting evidence for the funnel’s shape and order of magnitude, not a measurement of its value.

Case study: Weather forecasting, end to end through the funnel

Weather is the one domain where every term of Eq. (24.2) can be observed, which is why it recurs throughout this book and deserves assembly in one place. *Addressability*: the forward model is the dominant computational cost of forecasting, and it is exactly what AI replaced—AIFS operational at ECMWF from February 2025, NOAA’s AI system from December 2025 at 0.3% of the compute of its conventional suite [307, 308]. *Acceleration*: 10^3 – $10^5\times$ per forecast, measured, not asserted [309, 310]. *Validation*: forecasting possesses the strongest validation infrastructure in all of science—daily scoring against observation, decades of baseline skill curves—so survival is measured within months. *Deployment*: operational integration inside a year, because national agencies already own the delivery channel. And then, in May 2026, the term everyone forgets: ECMWF discontinued real-time running of all four *external* AI models after an upgrade to its analysis system degraded several of them, while its own AIFS—coupled to its own data assimilation—remained operational [327]. The episode demonstrates five things at once: an AI forecast model is a *component of a production system*, not a free-standing oracle; its performance is contingent on upstream data-assimilation infrastructure it does not control; retraining against infrastructure change is a recurring operational cost, not a one-time expense; deployment requires institutional continuity of exactly the kind Chapter 27 argues must be built rather than rented; and openness was what made the failure *visible*—the retired models could be evaluated, compared, and learned from precisely because they were published. Weather is thus not merely the best AI4S success story; it is the complete systems lesson—architecture, operations, lifecycle cost, trust, and value conversion in one worked example—and the standard against which every other domain’s claims in this chapter should be read.

24.5 Measuring the scientific frontier: a production metric, not another benchmark

Chapter 10 left the scientific frontier as the one frontier without an accepted metric, and noting a vacancy is worth less than proposing something to fill it. Benchmarks will not serve: a score on a question-answering set measures a component, and the failure ledger above is a catalogue of components that scored well while the system failed. What a ministry actually needs is a *production* metric—value out per resource in:

$$\mathcal{S} = \frac{\mathbb{E}[\mathbf{y}]}{C_{\text{compute}} + C_{\text{human}} + C_{\text{experimental}} + C_{\text{data}} + C_{\text{assurance}}}, \quad (24.3)$$

where the denominator is a scalar cost in currency— C_{compute} the machine time, C_{human} the expert labour, $C_{\text{experimental}}$ the wet-lab and instrument cost, C_{data} acquisition and curation, $C_{\text{assurance}}$ the verification and certification overhead of Chapter 18—and the numerator $\mathbb{E}[\mathbf{y}]$ is the expectation, over the portfolio of projects undertaken, of a *vector* $\mathbf{y}$ of outcome measures. Dividing a vector by a scalar leaves a vector, which is the point: $\mathcal{S}$ is deliberately not collapsed into a single number, because the collapse is where every previous attempt at this metric has smuggled in its assumptions. The components of $\mathbf{y}$: time to independently validated result; expert-hours consumed; replication survival rate; novelty and importance (no single automated novelty metric survives axiomatic scrutiny, and combined metrics do better [460]); downstream adoption; breadth across domains; false-discovery burden; safety and dual-use externality; and time to operational or regulatory deployment. The denominator’s five terms are the cost stack of Chapter 2 specialized to science—and the $C_{\text{assurance}}$ term is where the trust stack of Chapter 18 enters the economics of discovery directly.

The recent measurement literature supports both the vector form and the humility. The AlphaFold evidence shows *complementarity*—AI-using researchers produced 40% more novel experimental structures, not fewer—so a metric that scored substitution of human effort would have missed the largest effect [459, 313]. Frontier agents remain below 10% on demanding scientific literature-discovery tasks, so any metric weighted toward autonomous end-to-end discovery would currently measure noise [461]. Scientific value arises from a coupled human–model–simulation–experiment system, and Eq. (24.3) is written over that system, not over the

model. The strategic consequence: whoever operationalizes $\mathcal{S}$ first—runs it across a national research portfolio, publishes the vector, iterates procurement against it—owns the scientific-frontier scoreboard the way MLPerf owns the training scoreboard, and Appendix J adds its adoption to the forecast register.

24.6 The architecture of the prize

The failure ledger and the architecture argument of Part II point at the same conclusion. AI4S succeeds where a capable model is coupled to simulation, instruments, and data at scale—weather models coupled to assimilation systems, structure prediction coupled to experimental validation, design models coupled to autonomous laboratories. It fails where a model is asked to supply judgement alone. The workload that follows is therefore not chat: it is *long-context agentic orchestration over large batch scientific execution*, which is precisely the two-layer demand structure that Chapter 26 measures, and precisely the workload whose cost is dominated by million-token-context master orchestration rather than by raw dense FLOPs.

Three consequences follow, each connecting to an earlier chapter. *First*, the hardware that serves this best is capacity-rich and bandwidth-balanced rather than FLOP-maximal—the design point of Section 12.8, and the reason the memory arbitrage is a scientific-infrastructure argument and not only an economic one. *Second*, the models that serve it best are open, because scientific work requires reproducibility, provenance, auditability, and the freedom to fine-tune on data that cannot leave the institution—and because, as Chapter 27 showed, renting the orchestration layer converts a one-seventh workload share into a nine-tenths budget share. *Third*, the failure modes are exactly the ones Chapters 18 and 19 catalogue: a scientific result produced by an unverifiable artifact through an unauditable chain is not a scientific result, which makes the certification layer a precondition of the prize rather than an accessory to it.

24.7 Who is positioned

The uncomfortable observation is developed in full in Chapter 25, which reads the Chinese policy record precisely, measures the output through a European Commission bibliometric study rather than a Chinese source, and finds a programme organized around infrastructure and industrial delivery rather than around discovery. The short form: the most industrially organized AI-for-science programme in the world is Chinese, a substantial part of its output is given away under permissive licences, and the state’s top science prize in the most recent cycle went to a domestic computing machine rather than to a discovery made on one [737, 744].

Set against this, the Western response is real and large: a US national mission with an executive-order mandate to double the productivity of American science within a decade, over \$5 billion committed across fifteen-plus agencies and 278 projects, and the two largest scientific AI systems ever procured [350, 351, 352]; nineteen European AI factories with a €10 billion envelope and a €20 billion gigafactory programme behind it [353, 354]; national systems in the UK and Japan explicitly mandated for AI-driven science [355, 356]. The contest is genuinely joined. But note which side of it is being fought with open weights and cheap tokens, and recall from Chapter 26 that the announced infrastructures on *both* sides underprovision their own research populations by roughly an order of magnitude.

24.8 The strategic restatement

If AI4S is the largest prize, then three of this book’s arguments change register.

The token-price argument becomes a *compounding-rate* argument. Chapter 27 showed that purchased frontier tokens at a 20–60× premium consume three-quarters to nine-tenths of an

AI-for-science budget. Read against a 5%-per-year decline in research productivity and a social return above \$4 per dollar, that is not a procurement inefficiency—it is a tax on the one activity whose returns compound into everything else. A nation that pays it runs its science at a fraction of the rate it could.

The commoditization thesis acquires a beneficiary. Everything Part II describes—open weights, cheap inference, capacity-first hardware, sovereign-scale clusters—is, from the perspective of this chapter, *good*. Abundant cheap capable AI applied to research is the best available answer to the ideas-getting-harder problem, whoever supplies it. The book’s warning is not that commoditization is bad; it is that the country organizing the commons acquires the standards, the adoption, the certification authority, and the scientific compounding that come with it.

And the strategy for everyone else sharpens to a single sentence. If the largest value in AI is scientific, and scientific AI requires abundant compute, open models, verifiable artifacts, and institutions that can carry multi-year continuity, then the correct national investment is not a sovereign chatbot. It is measured capacity (Chapter 26), pooled internationally (Chapter 27), running audited open models (Chapter 18) on architectures chosen for long-context agentic science (Chapter 12). That is the whole recommendation of this book, and this chapter is the reason it is worth the money.

Chapter 25

China’s Scientific Machine: AI for Science as Industrial Policy

Chapter 24 argued that the largest value in artificial intelligence is scientific, and closed by observing that the contest over that value is genuinely joined. This chapter takes the observation apart on one side of it. Its claim is not that China has won AI for science—the failure ledger of Chapter 24 applies to Chinese results exactly as it applies to Western ones, and this chapter applies it, sometimes to results that its own sources report as triumphs. The claim is narrower and harder to dismiss: that China has been organizing for this particular prize longer, more explicitly, and more *industrially* than any Western actor, that the organizing is legible in the primary planning documents rather than inferred from behaviour, and that the shape it has taken—discovery chained by policy text to engineering and to product, running on a domestic substrate, distributed through open artifacts—is the playbook of Chapter 7 applied to the research enterprise itself. The chapter also carries its own best refutation, and carries it first, because the refutation comes from reading the Chinese planning documents precisely rather than enthusiastically, and precision here cuts against the thesis.

25.1 The policy record, read precisely

The State Council’s *Opinions on Deepening the Implementation of the “AI+” Initiative*—Guofa [2025] No. 11, signed 21 August 2025—set out six key actions, and “AI+ science and technology” is the first of them, ahead of industry, consumption, welfare, governance and global cooperation [734] [P]. Ordering in a State Council Opinions document is not decorative; it is the sequence in which ministries are expected to read their own obligations, and placing science ahead of industry in a document whose entire register is industrial is a deliberate statement about where the leadership believes the compounding lies. The science pillar has three sub-clauses. The first directs the acceleration of AI-driven new research paradigms and of “from 0 to 1” discovery. The second directs the construction of *scientific large models* and of AI-upgraded research platforms—an instruction to build artifacts, not to fund projects. The third is the analytically load-bearing one, and it is why this chapter is titled industrial policy rather than science policy: it directs the integrated development of “AI-driven technology R&D, engineering realisation and product landing.” The document does not treat discovery as an end. It treats discovery as the first stage of a pipeline whose last stage is a product, and it says so in the text rather than leaving the inference to a foreign analyst.

Set against that, the same document’s numeric commitments are the chapter’s first piece of counter-evidence, and it is a substantial one. The targets in Guofa [2025] No. 11 are *diffusion* targets—penetration of new-generation intelligent terminals and agents above 70% by 2027 and above 90% by 2030—and there is **no numeric target for science at all** [734] [P]. A planning system whose distinctive competence is hitting the things it measures has, in its flagship AI

document, measured adoption and declined to measure discovery. The strong reading of the ordering and the honest reading of the target list point in different directions, and both belong in the record.

The lineage runs back further, and its form matters as much as its content. The Ministry of Science and Technology and the National Natural Science Foundation launched the “AI for Science” special deployment on 27 March 2023, naming four basic disciplines—mathematics, physics, chemistry, astronomy—and four applied priorities—drug development, gene research, agricultural breeding, and new materials—together with commitments to domain-specific platforms, to the national next-generation AI public computing platform, to heterogeneous integration of high-performance and intelligent computing centres, and to scientific data sharing [735] [P]. It is a ministerial *deployment of work*, not a numbered regulation with an implementing budget attached, and Western sources that describe it as a “plan” overstate its form; this book does not, and the correction is not pedantry, because a deployment of work and a State Council plan carry very different degrees of compulsion on the agencies below them [344]. The traceable competitive instrument that followed is the National Key R&D Programme’s “Disruptive Technology Innovation” special programme, whose 2025 AI-for-science call of 13 May 2025 names five directions: frontier science, life and health, new materials, engineering design, and other AI-driven research-system technologies [735].

Now the precision point that is worth the whole section. The Party’s *Recommendations* for the 15th Five-Year Plan, adopted at the Fourth Plenum and released on 28 October 2025, contain in section 14 the operative instruction to *use artificial intelligence to lead a transformation of the scientific research paradigm*, and the official interpretation built around that phrase carries commentary from the CAS academician who founded the field’s principal Chinese institute [736] [P]. **That phrase does not survive into the Outline adopted by the National People’s Congress in March 2026.** In the Outline, AI appears under compute, data and diffusion—Chapter 12 on compute, algorithm and data supply, Chapter 13 on digital empowerment—while basic research appears in Chapter 8 *without* AI attached to it [736] [P]. The consequence for anyone citing the 15th Five-Year Plan in this argument is direct: a claim that the plan makes AI for science a central pillar of Chinese *state programming* is a claim about the Party recommendations, not about the plan, and the two documents do different work. The Recommendations state ambition; the Outline allocates. Two readings of the gap are available and this book prints both. The deflationary reading is that the Outline is a resource-allocation instrument and a paradigm shift is not a line item, so the omission means nothing. The consequential reading is that the drafting process, between October and March, converted a statement about *how research should be done* into a set of provisions about *what infrastructure should be bought*—and that the conversion is itself the finding, because it points at exactly the same revealed preference that Section 25.5 finds in the state’s own prize list. Either way, the gap between Party ambition and state programming is real, it is documentable from primary sources, and it is the strongest single piece of evidence against this chapter’s thesis.

Case study: Two documents, one plan: how to read the 15th Five-Year Plan on AI for science

The Five-Year Plan is not a document; it is a sequence of documents with different authors, different force and different functions, and conflating them is the most common error in Western writing about Chinese technology policy. Three artifacts matter here. The *Recommendations* (*jiānyì*) are a Party document, adopted at a plenum, stating direction and priority; they are where “use artificial intelligence to lead a transformation of the scientific research paradigm” appears, in section 14, in October 2025. The *Outline* (*gāngyào*) is a state document, adopted by the National People's Congress the following March, which converts direction into chapters, programmes and indicators; it is the instrument ministries and provinces are actually assessed against. And beneath both sit the sectoral and provincial implementation plans, which is where money appears. The paradigm-transformation phrase lives in the first and does not appear in the second. What *does* survive into the Outline is compute: ultra-large-scale intelligent computing clusters under study and construction, government purchase and leasing of compute, universal accessibility and lower costs for small users—all in Chapter 12, on compute, algorithm and data supply, with basic research appearing separately in Chapter 8 and without AI attached [736] [P]. Three rules follow for any reader of this literature, and this book applies all three. *Cite the document, not “the plan.”* *Track phrases across the Recommendations–Outline boundary*, because survival is the test of whether an ambition acquired an owner. And *treat the Outline’s silences as data*: a planning system that programs the substrate and omits the epistemology has told you which of the two it intends to be held accountable for. The finding is uncomfortable for this chapter’s thesis and is stated in the chapter’s own terms rather than buried: on the state’s own programming record, Chinese AI for science is an infrastructure commitment with a research aspiration attached, not the reverse.

The money underneath is bench money, distributed widely, and should not be mistaken for a moonshot. Roughly half the project categories in the National Natural Science Foundation’s current call foreground AI for science, at individual project budgets of EUR 240,000–600,000, with topics including generative AI for fusion reactor design and biological component design for extreme-environment biomanufacturing [738] [A]. Two observations follow. The first is that a funding agency restructuring half its call categories around a method is a stronger signal of institutional commitment than any single large programme, because it reaches the whole population of investigators rather than a selected set. The second is that budgets in that range buy graduate students and cluster hours, not frontier training runs—which is consistent with a strategy in which the state supplies the compute centrally (Section 25.5) and the grants supply the science.

25.2 The output, measured by someone other than China

The most useful measurement of Chinese AI-assisted scientific output was produced by the European Commission, which is the reason to use it. Bianchini, Di Girolamo, Ravet and Arranz, publishing on 7 April 2025, examined roughly three million papers across eighty fields over 2000–2022 and found that **China overtook both the United States and the European Union in AI-assisted scientific output from 2016**, producing more than 25,000 AI-assisted papers in 2022 against roughly 15,000 for the EU and 12,000 for the US, with reported advantages on novelty and citation impact as well as on volume [737] [V on the counts, A on the novelty and impact findings]. A second, independently constructed Western measurement points the same way: China leads the Nature Index with output up 22.4% year on year—the only top-ten country with double-digit growth—and holds the top institution globally along with eight of the top ten in the preceding year’s table [342, 343] [V].

Four limits belong immediately beside those numbers, and the first is the one this book has spent a chapter establishing. *Output is not discovery*. Chapter 24’s funnel exists precisely because papers are the first term of a chain whose later terms—validation survival, deployment conversion, lag—are all fractions below one, and a volume lead measures the input side of that chain. *The window ends in 2022*, which is before the model generation this book is about; the

study therefore measures the run-up rather than the present state, and its most defensible use is as evidence that the Chinese position was built before the current phase rather than during it. “AI-assisted” is a *bibliometric classification*, not a causal claim about what the AI contributed. And a *volume lead in a literature with documented quality-control strain* is not automatically a lead in anything: the same chapter records a journal measuring submissions up 42% and over 30% of its peer reviews showing detectable AI use. What survives all four qualifications is still substantial, and it is the following: on the two most credible Western-run measurement systems, in the domain this book argues is the largest prize, the leading producer is China, and has been since before the transformer era’s commercial phase began.

It is worth stating what would settle the question the bibliometrics cannot, because the absence is itself a finding. A volume lead in AI-assisted publication would be converted into a claim about *discovery* by evidence of the kind Chapter 24 required of every Western result it credited: an operational adoption by an agency that did not build the model, a controlled clinical endpoint, an independently replicated synthesis, a proof checked by the community that did not produce it. On that standard the Chinese record is thin in most domains and genuinely strong in exactly one—weather, where third-party evaluation exists in the peer-reviewed literature (Section 25.4)—which is the same domain where the Western record is strongest, and for the same reason: forecasting scores itself daily against observation, and no developer controls the weather. The general instrument for making this comparison does not exist. Nobody operates the production metric of Section 24.5, in China or anywhere else, which means that the honest current state of knowledge is that we can measure and compare scientific *output* across countries and cannot yet measure or compare AI-assisted scientific *return* in any country. Every claim in this chapter about Chinese leadership is therefore a claim about the input side of Chapter 24’s funnel, and is bounded accordingly.

25.3 The institutional layer

What distinguishes the Chinese institutional layer is not that it exists—every advanced economy now has AI-for-science institutes—but that it was built early, that it is deliberately plural, with several bodies holding overlapping mandates and no single national AI4S laboratory, and that most of these institutions ship artifacts rather than reports. The plurality is step 4 of Chapter 7 in an unexpected setting: the state does not designate the winning laboratory, it funds several and lets the release record select. Table 25.1 states what each has actually shipped, which is a harder test than what each was founded to do.

ScienceOne deserves two observations that the table cannot carry, and they point in opposite directions. The architectural one: it was built by twelve CAS institutes on Chinese open-source foundation models, and it consumes AlphaFold and MatterGen—the two most celebrated Western scientific models—as *tools underneath its own orchestration layer* [742] [P]. That is the mechanism of Section 17.1 operating in a scientific setting: whoever holds the layer above treats the layer below as a component, and a Western laboratory that publishes a model without holding the workflow around it has supplied a component to somebody else’s system. Its predecessors and successors sit in the same pattern—a national scientific agent platform with access to 170 million papers and 300-plus computing tools, upgraded a year later onto a heterogeneous supercomputing backend, and an autonomous research system wired into a national supercomputing network linking thirty-plus centres [345, 346, 347] [P/A]. The epistemic observation cuts the other way and is stated with equal force: ScienceOne’s parameters, training compute and open-source status were not disclosed, and its performance claims are self-reported [742]. A system that cannot be independently evaluated cannot be independently credited, and this book credits its existence and its architecture, not its capability.

What the layer does not contain is as informative as what it does. There is no Chinese equivalent of a national laboratory system with statutory continuity and multi-decade budget

Table 25.1: The Chinese AI-for-science institutional layer as of 9 September 2026, scored on shipped artifacts rather than on mandate. Sources: [743, 742, 739, 741, 744, 745].

Institution	Founded / base	Shipped
AI for Science Institute, Beijing	2021, under CAS academician E Weinan	Runs the DeepModeling community; DPA potential family; ABACUS density-functional code; DeepFlame combustion platform; AIS-Square model-sharing platform
Shanghai AI Laboratory	Shanghai	The most productive shipper: Intern-S1 (Apache 2.0), Intern-S1-Pro, the Intern Duanyan discovery platform; ten announced 2026 flagship scientific results
Shanghai Academy of AI for Science	2023, with Fudan	The FuXi weather and climate family, including the Fengshun sub-seasonal model with the National Climate Center
BAAI	Beijing	2026 “Wujie” series: Brain μ (neuroscience), Open-Complex 2.5 (biomolecules), Physis (physical-world modeling)
Pengcheng Laboratory	Shenzhen	Pengcheng Cloudbrain, the first fully domestic large-scale intelligent computing system; National Science and Technology Progress Award, First Prize
Zhejiang Lab	Hangzhou	Three-Body Computing Constellation: first twelve compute-carrying satellites launched May 2025, stated target $\sim 1,000$ by around 2030
CAS ScienceOne	2025, twelve CAS institutes	Scientific agent platform built on <i>Chinese open-source foundation models</i> , integrating AlphaFold and Matter-Gen as tools; digital-cell platform, collider simulation, telescope coordination

lines; the institutions in Table 25.1 are a mixture of academy institutes, municipally sponsored laboratories and a company, several founded within the last five years, and the plurality that makes them competitive also makes none of them obviously durable. Continuity is the requirement Chapter 27 identifies as decisive for regulated scientific deployment—a clinical model needing two or three years to certify and support five years after that—and it is the one property this layer has not yet had time to demonstrate. That is a real gap in the Chinese position, and it is the mirror image of the West’s: the Western system has the institutional continuity and lacks the shipping cadence.

25.4 The models, and the fact that matters most about them

Four model families carry the Chinese scientific position, and the single fact that matters most about them is not any benchmark number but their licences. This section therefore inverts the usual order of presentation: capability claims are recorded at the grade the evidence supports, which for most of them is developer-reported, and the licence and the training-data composition—the two things that can be checked by anyone with a browser—are treated as the load-bearing facts. That inversion is not a rhetorical device. It follows from Chapter 24’s finding that benchmarks measure components while systems fail, and from Chapter 18’s requirement that a scientific artifact be auditable before it is credited.

Intern-S1 is the reference artifact. Released by Shanghai AI Laboratory on 26 July 2025, it carries 241B total parameters—a 235B mixture-of-experts language backbone with a 6B vision encoder—pre-trained on 5T tokens of which *2.5T are scientific-domain*, and it is published under **Apache 2.0** on Hugging Face, ModelScope and GitHub [739] [P]. Its reported scores—ChemBench 83.4, MatBench 75.0, MSEarthMCQ 65.7, MicroVQA 63.9—are self-reported, its comparative

claims against Western frontier models are vendor-reported, and no independent replication was located; they are recorded here at that grade and no higher. The same discipline applies to its trillion-parameter successor Intern-S1-Pro of 2 April 2026, whose headline SciReasoner margin over leading closed models is likewise the developer's own measurement [349]. The load-bearing facts are not the scores. They are the token mix—half a five-trillion-token pre-training corpus deliberately spent on science, which is a resource-allocation decision no Western laboratory has publicly matched—and the licence.

Protenix makes the licence point unmistakable. ByteDance's AI4Science group released v1.0.0 on 5 February 2026 at 368M parameters and v2 on 8 April 2026 at 464M, both under **Apache 2.0**, with code, weights, training and inference pipelines and a web server all published; the accompanying claim is that it is the first fully open-source model to outperform AlphaFold3 across diverse benchmark sets at matched training-data cutoff, model scale and inference budget, evaluated over more than six thousand complexes against AlphaFold3, Boltz-1 and Chai-1 [740] [P on the release, developer-evaluated on the comparison]. The evaluation was performed with PXMeter, a toolkit the developers themselves built and released—which is better practice than most and still not independent replication, because a harness written by the author of the model under test is a contribution to reproducibility rather than a substitute for it. *The significance is the licence, not the benchmark.* The open-weight ratchet of Section 10.3.1 in Chapter 10 has reached structural biology, and it reached it from a Chinese laboratory. Note what that does to the position Chapter 24 describes: the most diffused AI4S asset on Earth was given away by a Western laboratory as a *database*—two hundred million structures, seventeen million GPU-hours of computation donated—and it now faces a competitor given away as *weights*, which is a different and more consequential act of generosity, because a database is a result and a weight file is a means of production.

DeepModeling and DP Technology are the deepest instance, and the one this book finds hardest to argue with. The stack comprises DeePMD-kit, the deep-potential molecular-dynamics package; **DPA-2**, published in *npj Computational Materials* in December 2024, covering 73 elements across 18 pre-training datasets spanning alloys, semiconductors, battery materials, drug molecules, ferroelectrics, solid-state electrolytes and catalysis; the ABACUS electronic-structure package; and the OpenLAM initiative launched on 1 January 2024 [741] [P]. One detail inside DPA-2 is worth more than its headline: the density-functional labels come from VASP, Gaussian *and* the domestic ABACUS code—which is supply-chain capture executed at the level of the training data, a domestic electronic-structure code inserted into the label pipeline of the flagship domestic atomic model. The community reports more than ten thousand members across more than twenty countries; the commercial arm reports more than three million users across over a thousand universities and research organizations plus corporate clients including PetroChina, CATL and BYD [741, 348] [A]. And the financing settles what kind of entity this is: DP Technology's December 2025 Series C above RMB 800M (about US\$114M) had **three of six named investors that are Beijing municipal state funds** [741] [P/A]. This is a state-backed national champion, not a venture outcome, and describing it as a start-up misstates the instrument. Code, weights and datasets are open. The sentence this book is obliged to print, because it is what the evidence located actually supports and because it cuts against the framing a Western reader will bring: *on the evidence available at 9 September 2026, this is among the most genuinely open scientific-AI ecosystems anywhere, Western or Chinese.*

Weather is the one domain with genuine external validation, and it inverts the openness picture. Huawei's Pangu-Weather has been checked by third parties in the peer-reviewed literature—evaluations in *Remote Sensing*, *Meteorological Applications* and *Weather* during 2025, the last covering FengWu as well—which makes this the only Chinese AI4S domain where the capability claim rests on somebody other than the developer [745] [V]. The FuXi family, by contrast, is strong on deployment and weak on verification: China's first purely data-driven end-to-end forecasting system, ingesting satellite remote sensing directly, operationally deployed

across Chinese meteorological departments and integrated into an early-warning platform serving forty-three countries, selected among China's 2025 Top Ten Meteorological Science and Technology Advances in March 2026—while its headline claims, medium-range ensembles beating ECMWF on 98.1% of variables by CRPS and sixty-day forecasts in three minutes, are developer-reported and not independently replicated, and the open-weight status of the current system is not stated [745] [P on deployment, developer-reported on skill]. The meteorological administration's own operational family, one member of which has been open-sourced, sits alongside [311]. Structural biology and weather thus form an instructive pair: the domain where the Chinese artifact is most open is the one where the claim is least externally checked, and the domain where the claim is best checked is the one where the artifact's licence is unstated.

And now the counter-observation, which must be made and made prominently. The newest and most capable Chinese scientific platform—Shanghai AI Laboratory's Intern Duanyan, announced 17 July 2026, a five-layer discovery stack with more than two hundred cross-disciplinary agents across six domains and a dry-lab/wet-lab closed loop—has **no stated open-source release**, and is access-by-invitation [739] [P]. Read against the mechanism of Chapter 10, this is not an anomaly but a prediction coming true. The ratchet compels *weights* onto the commons because adoption is the currency and the first defector loses the derivative ecosystem; it compels nothing whatever about a *platform*, because a closed-loop agentic laboratory with wet-lab integration is not a file anyone can fork, and its rent lives in the orchestration and the instrument access rather than in the parameters. The openness documented in this section may therefore be a phase rather than a principle—exactly what Chapter 7's step 6 predicts of any layer once it starts producing rent, and exactly the policy-and-continuity risk that Chapter 27 names first among the three its insurance is meant to cover. That risk is no longer hypothetical. It is instantiated, and it is instantiated at the top of the Chinese capability stack.

25.5 The infrastructure, and the prize the state actually gave

Pengcheng Cloudbrain is the first fully domestic large-scale intelligent computing system, built by Pengcheng Laboratory with Huawei, Peking University and Tsinghua, and on 8 July 2026 it was announced as the winner of the **National Science and Technology Progress Award, First Prize** for the 2025 cycle; it is reported to support training for more than half of China's domestic foundation models and to have guided the construction of twenty-five intelligent computing centres [744] [P on the award, A on both ratios]. *The state's top science prize went to the machine, not to a discovery made on it.* A national prize list is a revealed preference over what a state believes is scarce, and this one says that the scarce thing is not a result but the domestic capacity to produce results—which is the argument of Chapter 12 restated by the awarding body rather than by this author, and the reason Chapter 26 measures capacity rather than publications.

The 15th Five-Year Plan Outline supplies the instrument underneath. Its compute provisions direct the study and construction of *ultra-large-scale intelligent computing clusters*; they name government *purchase* and *leasing* of compute as explicit policy instruments; and they set the aim of making compute “universally accessible and easy to use” while lowering costs for smaller users [736] [P]. That is move M2 of Chapter 7—demand before supply—applied to research computing rather than to solar modules: the state does not merely build the machine, it commits to being its customer, and it prices access down for precisely the long tail that Chapter 27 identifies as the population any scientific-compute institution must serve. Note also which of the two documents of Section 25.1 this provision lives in. The paradigm-transformation sentence dropped out between the Recommendations and the Outline; the compute-purchase provision survived. The state plan programs the substrate, not the epistemology, and the prize list agrees with the plan.

The inversion is worth putting to Western science ministries directly, because it is the most

transferable thing in this chapter and it costs nothing to act on. A national prize system is one of the cheapest instruments a government holds and one of the most legible signals it sends: it tells a generation of researchers which career produces recognition. Western prize lists in this domain decorate results—the structure, the forecast, the molecule—and treat the infrastructure that produced them as procurement. The Chinese list, in its 2025 cycle, decorated the infrastructure. Neither choice is obviously right; a state that only rewards machines will get machines and no science, and a state that only rewards discoveries will find, a decade later, that it rents the means of producing them. But the asymmetry is real, it is recent, and it points the same way as every other finding in this chapter: toward a programme organized around productive capacity rather than around outputs, which is precisely the organizing principle Part I found in solar, in batteries and in vehicles.

One quantity this chapter declines to supply. There is no public figure for how much compute the Chinese research system actually receives, in any unit comparable to the BGPU accounting of Chapter 26, and no Chinese instantiation of that chapter's national demand model exists. This chapter therefore does not place China on the sizing ladder, does not estimate its per-researcher provision, and does not assert that its research population is better served than Japan's or America's—all three of which would be easy to write and none of which the evidence supports. What can be stated is the instrument and its direction: a state that buys and leases compute on behalf of its researchers and drives the price down for small users is acting directly on the gap that Chapter 26 measures at roughly tenfold in the countries where it has been measured, whatever level it in fact reaches.

25.6 The programme, scored on the playbook

Chapter 7 distilled six strategic steps and nine tactical moves, and the discipline this book imposes on itself is to score campaigns against them rather than to admire them. The AI4S programme scores as follows.

Step 1, target selection (M1). The discontinuity chosen is not a product category but the research paradigm itself—"use artificial intelligence to lead a transformation of the scientific research paradigm" is a target-selection sentence in the exact form of Chapter 7's first step [736]. What is stranded if the paradigm changes is the incumbent's accumulated advantage: the seat-licensed scientific software stack, the closed frontier platforms, the instrument and publication ecosystem, and the tacit practice of research groups organized around them. State designation is present, dated and layered—the 2023 ministerial deployment, the 2025 State Council Opinions, the 2025 Party Recommendations—though, as Section 25.1 established, thinning as it moves from Party ambition to state programming.

Step 2, demand before supply (M2). Government purchase and leasing of compute, subsidized access, the universal-accessibility language, and twenty-five guided intelligent computing centres [736, 744]. This is structurally the feed-in tariff of the solar column and the Top Runner procurement mechanism, redirected at research capacity: guaranteed demand for domestic machines, delivered as subsidized supply to a scientific user base that could not otherwise afford it.

Step 3, overcapacity as strategy (M3): partially, and the evidence is weaker than the pattern. Twenty-five guided intelligent computing centres and an Outline provision directing the study and construction of ultra-large-scale clusters have the shape of the solar and battery build-outs—capacity funded ahead of demonstrated demand, with the losses absorbed locally and the cost curve retained nationally [744, 736]. But the shape is not the finding. No public utilization figures for those centres were located, and without them the difference between strategic overbuild and stranded provincial assets cannot be established from open sources; this book records the instrument and declines to score the step [S]. It is worth naming what evidence would settle it: published occupancy or allocation data from the intelligent computing centres, of

the kind Chapter 26's accounting unit was designed to make comparable across countries.

Step 4, selection rather than planning (M4). Six-plus institutions with overlapping mandates, two competing weather families, several scientific-model families, and no single designated national AI4S laboratory (Table 25.1). Chapter 27 draws precisely this lesson from the Chinese case—that the government's role was to set the stage and permit intense bottom-up activity—and it is worth noting that the lesson is drawn there from a working discussion among Western centre directors, not from Chinese sources, which makes it a convergent reading rather than a borrowed one.

Step 5, commoditize the adversary's rent layer (M6). There are *two* rent layers here, and the Western debate has priced only one. The first is the closed frontier platform, the subject of Chapter 10 and of the arithmetic in Section 27.1, where a 20–60× purchased-token premium converts a one-seventh workload share into a three-quarters-to-nine-tenths budget share. The second is the one nobody in the AI debate mentions: the *subscription-priced scientific software stack*—commercial density-functional and molecular-dynamics codes, cheminformatics suites, simulation packages licensed per seat per year, structure and sequence resources behind institutional paywalls—which has collected rent from the world's researchers for three decades. Open scientific weights, open codes and open datasets price both layers at marginal cost simultaneously. Section 17.1's result governs the consequence: the rent does not vanish, it relocates, and it relocates toward compute, toward platforms with instrument access, and toward whoever holds the orchestration layer—which is what Intern Duanyan and ScienceOne are.

M5, supply-chain capture. The chain is visible end to end: domestic open foundation models under ScienceOne rather than Western ones; ABACUS as a domestic density-functional code inside the label pipeline of the flagship domestic atomic model; domestic silicon underneath, supplied through the Huawei partnership that built the first fully domestic large-scale intelligent computing system—the machine the state chose to decorate [P on the partnership and the award, A on the accelerator line]; and, at the newest link, compute-carrying satellites at Zhejiang Lab against a stated target of roughly a thousand by around 2030 [743, 741, 744]. Model, framework, code, chip, machine—the same sequence as Chapter 7's AI row, run through a scientific application rather than a commercial one.

M8, standards capture: not achieved, and available. The one visible instance is the potential-energy-model interface, where DPA and OpenLAM are becoming a de facto interface for machine-learning interatomic potentials in the way file formats become standards—through a ten-thousand-member community across twenty-plus countries rather than through a committee [741]. Beyond that there is no Chinese scientific-model certification scheme, no AI4S standards body, and no scientific equivalent of MLPerf. That vacancy is exactly the one Section 24.5 of Chapter 24 identified and proposed a metric for, and it remains unfilled by anybody—which is a strategic opportunity for the institution of Chapter 27 and a reason Chapter 18's certification argument is urgent rather than academic.

The industrial hinge, which is this chapter's distinctive finding. Of the ten joint results Shanghai AI Laboratory announced at the 2026 World AI Conference, three are aerospace design with a state aerospace research institute and a commercial rocket firm, electroplating-materials formulation with Xiamen University, and an analogue chip-design agent with Shanghai Jiao Tong University claiming five- to tenfold sizing-optimization efficiency and a halving of layout-planning cycles [746] [P on the announcement, developer-reported on the efficiency figures]. A “scientific” programme delivering into aerospace, industrial chemistry and semiconductor design is not a scientific programme in the Western sense of the phrase; and the analogue-design case is the flywheel of Section 24.3.4 pointed directly at the layer Chapter 13 identifies as the binding physical constraint. The policy text makes the intent textual rather than inferred: the science pillar's second sub-clause chains AI-driven research to “engineering realisation and product landing” [734]. That is the sense in which this chapter's title is a description rather than a characterization.

And the detail that cuts against the framing, printed because it does. Three of those same ten results carry foreign co-authors—Columbia, Cambridge, University College London [746] [P]. The programme is not hermetically sealed, and any account that treats Chinese AI for science as a closed national system is wrong on the byline evidence. This matters directly for the next section, because an institution designed to exclude Chinese participation would be excluding collaborators that Western universities already have, on projects already published.

Step 6 and M9, cartelize and gatekeep: not observed. There is at 9 September 2026 no restriction on Chinese scientific model release, no export control on scientific weights, and no gatekeeping behaviour in this domain of the kind rare earths and polysilicon supply. The single datapoint consistent with the endgame is the unstated licence on the most capable platform, and one datapoint is a flag rather than a finding [S].

25.7 What this does to the consortium proposal, and what it does not

If this chapter is right, then the case usually offered for a Western or neutral open scientific-AI consortium is not merely weak but false, and it should be abandoned rather than defended. That case says: the scientific commons is thin, closed platforms dominate, and public money must create an open base. On the evidence of Section 25.4, the open base exists. It comprises Apache 2.0 scientific foundation models at 241B and trillion-parameter scale, an openly licensed structure-prediction model with published code, weights and pipelines, and an atomic-potential ecosystem whose code, weights and datasets are all open and whose community spans twenty-plus countries. A consortium justified as *creating* the commons would be arguing against a fact, and would be found out quickly by the researchers it proposed to serve, who are already downloading the artifacts.

The defensible case is different, narrower, and stronger, and it is the one Chapter 27 actually makes: *the commons has one principal contributor and no governance the rest of the world is party to*. Each of the three risks that chapter names is sharpened rather than softened by the evidence assembled here. *Continuity*: the most capable Chinese scientific platform has no stated open release, which converts the argument from a hypothetical about future Chinese policy into an observation about present Chinese practice. *Verification*: ScienceOne's parameters and training compute are undisclosed, FuXi's headline skill claims are developer-reported, Intern-S1's comparative benchmarks are vendor-reported, and Protenix—the most rigorous of the four—is evaluated by a harness its own authors wrote. Four of the most consequential artifacts in the field cannot be checked in the way Chapter 18 requires, and the correct response to that is an evaluation institution, not a rival model. *Monoculture*: if a single national ecosystem supplies the interatomic potentials, the scientific foundation models and the agent frameworks that a generation of published results depends on, then the systemic argument of Chapter 19 applies to the scientific literature itself, and applies whether or not anyone has behaved badly.

The strongest version of the opposing case must also be printed at full strength, because it is genuinely strong. A genuinely open scientific commons is a public good regardless of who funds it: a researcher in Nairobi or Jakarta running an open atomic potential on open weights is better off than one who cannot, and the nationality of the contributor does not enter that welfare calculation anywhere. A consortium built to *exclude* its largest contributor would be a worse instrument than one built to include it—it would forfeit the artifacts, it would forfeit the neutrality that Table 27.2 identifies as its only durable asset, and it would substitute an origin rule for the evidence-based evaluation that is the sole product an evaluation body has to sell. And there is a harder form of the objection still. If AI for science is the largest prize, and scientific knowledge is non-rival, then a policy that slows the commons in order to change its ownership destroys more value than it redistributes, and does so in the one domain where Chapter 24 calculated the social return at above four dollars on the dollar. On that reading the

correct Western response to Section 25.4 is not a rival consortium at all but a larger contribution to the same commons, plus the evaluation apparatus nobody yet operates.

One asymmetry is worth isolating, because it is the pivot on which the disagreement actually turns and it is not about models at all. Contribution to a commons and governance of a commons are different goods, and the second does not follow from the first. The Chinese programme supplies artifacts at a rate no other single actor matches; it does not supply—and has not been asked by anyone to supply—an evaluation regime, a revocation procedure, a provenance standard, a certification ladder, or a forum in which a non-Chinese institution has standing to contest a result. Those are the five functions Chapter 27 specifies operationally and the ones Chapter 18 argues nobody currently owns. Framed that way the choice put to the rest of the world is not between a Chinese commons and a Western one; it is between a commons with artifacts and no governance and a commons with both. That framing has the additional merit of being buildable: an evaluation and certification institution can be constituted, funded and staffed at a small fraction of what a rival model programme costs, and unlike a rival model programme it becomes more valuable, not less, as the artifacts it evaluates get better—whoever makes them.

This chapter does not resolve that tension; Chapter 27 does, and does so after Chapter 26 has supplied the missing quantity. What this chapter has established is the shape of the question those two chapters inherit. The prize is the largest one on the table. The most industrially organized programme aimed at it is Chinese, is legible in primary policy text, chains discovery to product by the wording of that text, runs on a substrate the state's own prize list ranks above the discoveries made on it—and distributes a substantial part of its output for free, under permissive licences, to everyone including its competitors. The question left for the rest of Part IV is therefore not whether to build on that commons, which every scientist reading this book is already doing, but whether an open commons with one principal contributor and no shared governance is a commons or a dependency, and what instrument short of exclusion changes the answer.

Chapter 26

How Much Compute Does a Nation’s Science Need?

Every argument in this book has so far been about who supplies capability and at what price. This chapter asks the reciprocal question, which almost no government has answered in public: *how much does a country actually need?* It matters here for three reasons. It converts the sovereignty argument from a posture into a number, and numbers can be checked. It exposes—quantitatively—the gap between what nations are announcing and what their own research populations would consume, which turns out to be an order of magnitude or more. And it supplies the demand side of the institutional proposal in Chapter 27: one cannot design a federation to pool capacity without a shared, comparable estimate of the capacity required. The material here follows a measurement-anchored sizing framework developed at RIKEN R-CCS and its three national instantiations [305]; the epistemic discipline is the framework’s own, and it matches this book’s: every input is typed by what would make it trustworthy.

Why this chapter is in the main text. A sizing framework detailed enough to be checked can read as a second book inserted into the first, and the temptation is to relegate it to an appendix. The temptation is resisted, and the reason is Chapter 24: if the largest value in AI is scientific, then the question “how much compute does a nation’s science need, and what does it cost to rent instead of own?” is not a technical annexe to the convergence thesis—it is the thesis, evaluated on the workload that matters most. What the appendix instinct is right about is the connective tissue, and this chapter supplies it: every section below closes by naming the chapter it feeds.

26.1 The measurement gap

National AI-for-Science infrastructure has, until recently, been sized without measurement. Procurement scales were set by budget envelopes, vendor availability, flagship-application extrapolation, or elicitation workshops—none of which measures the per-researcher demand of the workload class the infrastructure is being built for, namely LLM agents orchestrating simulation, data analysis, and autonomous experimentation [305]. Meanwhile the deployment wave is global and simultaneous: the US Genesis Mission systems, EuroHPC’s AI factories, the UK’s Isambard-AI, Japan’s Rikyu and FugakuNEXT line, and Singapore’s national AI programs are all sizing sovereign AI4S capacity now, mostly without a common yardstick.

Three developments make a portable framework possible. The demand side acquired a measured anchor: a 56-day operational record of an early-access agentic campaign on a GB200-class system—101 vetted users, 53,173 jobs, 340,397 GPU-hours, saturating an ~800-GPU partition within five weeks—which yields a revealed-demand floor of 443 delivered BGPU-hours per active heavy-user day that any country can adopt as a starting scenario and test against its own probe. The supply side converged onto public, reproducible serving benchmarks. And the

policy context became genuinely multi-national, so that comparability now has customers [305]. What the OECD’s Frascati Manual did for R&D statistics—not identical answers across countries, but commensurable ones—is what a shared, openly gated estimation framework could do for research-compute demand.

26.2 A vendor-neutral unit, and a conversion rule

Definition (Baseline GPU). One BGPU is the AI-serving throughput of a reference accelerator possessing an 8 TB/s HBM memory system—concretely a B200/B300-class device. It is a *planning unit*, not a product: any accelerator converts into it by the bandwidth-first, benchmark-adjusted rule below, with a stated adjustment band. A separate and never-summed unit, BGPU_{BW} , denotes 8 TB/s of delivered bandwidth for classical modelling and simulation. Conversion uncertainty is carried explicitly as the δ band of Table 26.1 and propagates into every capacity figure in this chapter; a $\pm 10\%$ band on δ moves national requirements by the same proportion, which is smaller than the population-definition sensitivity discussed below and much smaller than the behavioural-transfer uncertainty.

All capacities and demands are stated in **Baseline GPUs** (BGPU): the AI-serving throughput of one NVIDIA B200/B300-class accelerator—concretely, a device with an 8 TB/s HBM memory system, the resource that governs long-context decode. The unit is deliberately not a product. Any accelerator g converts by a bandwidth-first, benchmark-adjusted rule

$$\gamma_g = \frac{\text{BW}_g}{8 \text{ TB/s}} \times \delta_g,$$

where the bandwidth term is physics and δ_g absorbs compute-side and software-stack differences, estimated from published matched-condition benchmark results and carrying a stated band [305]. The rule’s worked example matters politically as well as technically: AMD’s MI355X carries 288 GB of HBM3E at 8.0 TB/s—a bandwidth ratio of exactly 1.00—and its published MLPerf record supports $\delta \approx 1$, so it converts at 1.0 BGPU, largely equivalent to Blackwell. *A requirement denominated in BGPU is procurable from any vendor whose devices convert at a defensible factor*, which is precisely what a public planning framework must guarantee and what a vendor-named unit cannot. Table 26.1 gives the adopted factors.

Table 26.1: Cross-vendor BGPU conversion: HBM bandwidth ratio, benchmark-adjustment band δ , and adopted planning factor γ [305]. Scenario entries lack a public reference-class serving result.

Device	Vendor	BW (TB/s)	γ (BGPU)	Basis
H100	NVIDIA	3.35	0.5	δ 1.0–1.2
H200	NVIDIA	4.8	0.5	δ 0.8–1.0
B200 / GB200 / B300	NVIDIA	8.0	1.0	definitional
MI300X	AMD	5.3	0.6	δ 0.85–1.0
MI355X	AMD	8.0	1.0	δ 0.9–1.1 (MLPerf-supported)
TPUv7 (Ironwood)	Google	7.37	~ 0.9	scenario
R200 (Rubin)	NVIDIA	22.0	2.5	δ 0.91 derate
F200-class (2030 flagship)	—	~ 48	6 (rounded)	<i>derived</i> , not asserted

The 2030 flagship-generation factor of 6 is derived rather than taken from a vendor roadmap: dividing a published program disclosure of ≥ 800 PB/s aggregate accelerator memory bandwidth over a disclosed fleet of $\sim 15,000$ – $20,000$ devices gives 40–53 TB/s per device, i.e. 5.0 – $6.7\times$ the 8 TB/s baseline; the planning value is rounded to 6, and because the bandwidth target is a floor, these are lower bounds [305]. Installed fleets deliver $H_{\text{yr}} = \alpha \times 8,760 = 7,884$ effective GPU-hours per GPU-year at fleet availability $\alpha = 0.90$. Classical modeling-and-simulation demand, where a

country carries it, is expressed in a separate bandwidth currency (BGPU_{BW} , one unit = 8 TB/s of delivered bandwidth) and is *never* summed with serving-BGPU—a discipline that prevents the most common category error in national compute accounting.

26.2.1 From scalar to workload-weighted vector

The framework owes one generalization an explicit statement, and this subsection supplies it. A bandwidth-first scalar is defensible because long-context, memory-bound decode dominates the measured national AI4S workload—but the workload is a portfolio, and a device strong on decode can be weak on training software, network injection, or reliability. The disciplined form defines a per-device capability vector,

$$\mathbf{B}_g = (B_{\text{decode}}, B_{\text{prefill}}, B_{\text{training}}, B_{\text{simulation}}, B_{\text{network}}, B_{\text{storage}}, A_{\text{availability}}),$$

where each of the first six components is the device’s sustained throughput on that workload class *divided by* the reference device’s sustained throughput on the same class under the matched-condition protocol of Section 12.6—so each is a dimensionless ratio equal to 1.0 for the reference device—and $A_{\text{availability}}$ is the fraction of wall-clock time the device is usable in production, likewise normalized to the reference fleet’s $\alpha = 0.90$. The scalar planning unit is derived only *after* applying a declared national workload mix $\{\omega_k\}$, non-negative weights summing to one over the seven components:

$$\text{BGPU}_{\text{plan}}(g) = \sum_k \omega_k B_{g,k}.$$

The two-layer demand model already supplies the mix: the token-service layer weights decode and prefill; the scientific-execution layer weights training, simulation, network, and storage; and the published national instantiations state their ω_k implicitly through the layer split. What the vector form adds is protection against exactly the failure that matters: a device with spectacular memory bandwidth but immature training software, poor injection bandwidth, or weak fleet reliability can no longer receive a flattering conversion factor, because its B_{training} , B_{network} , and $A_{\text{availability}}$ components are carried explicitly and the declared mix prices them. The scalar survives for politics—ministries budget in one number—but the vector is now the definition and the scalar is its declared projection, with the mix printed beside every national requirement. (The LX2-class integrated processor is the instructive case: strong B_{decode} and $B_{\text{simulation}}$, unproven B_{training} , unknown $A_{\text{availability}}$ at fleet scale—the vector says precisely what the proof suite of Section 12.2.2 must fill in before a conversion factor is defensible.)

26.3 The two-layer demand model

A nation’s requirement for target year Y is the sum of two concurrently provisioned layers, $G_{\text{nat}}(Y) = G_{\text{tok}}(Y) + G_{\text{sci}}(Y)$: a token-service layer and a scientific-execution layer (Figure 26.1). Both must be provisioned simultaneously, because the agentic loop does not close otherwise—the orchestrating model is idle without the batch jobs it launches, and the batch jobs are unlaunched without the orchestrator.

The token layer. The open-tier population N (the number of researchers entitled to draw on the national service) is partitioned into service tiers indexed by i , with population shares φ_i (summing to one), duty factors u_i (the fraction of the year a researcher in tier i is actively drawing tokens rather than idle), and sustained output-token rates r_i while active; the portable default is a three-tier profile of 10%/20%/70% of researchers at 2,000/400/150 Mtok per user-month. The national output stream is $T_{\text{nat}} = N \sum_i \varphi_i u_i r_i$ tokens per second, and its serving cost is $G_{\text{tok}} = T_{\text{nat}} \cdot k$, where k is the BGPU cost per token per second of a canonical, versioned agentic serving configuration: a *master* model—a trillion-parameter MLA mixture-of-experts holding a




Figure 26.1: The demand model. A nation’s agentic AI-for-Science requirement is the concurrent sum of a token-service layer, priced through a canonical master/sub-agent serving configuration, and a scientific-execution layer decomposed by function into simulation, training, and data processing—with training pinned to GPU-class substrate and simulation the only cross-substrate routable term. Note the master asymmetry carried in the top row: a tenth of the tokens, generated at million-token context, consume nearly nine-tenths of the token-service capacity. After the sizing framework [305]; the serving configuration is a versioned reference object [U]-r/[P], not a physical constant.

2^{20} -token campaign context—delegating to cheap, high-throughput 120B-class sliding-window *sub-agents* at 128K context, with 10% of national tokens going to the master [305].

That configuration produces the single most consequential structural fact in the framework, and Figure 26.2 (top) states it: at $k = 9.79 \times 10^{-4}$ BGPU per token per second, **88.7% of the token-service capacity is the master term**. Ten percent of the tokens, generated at million-token context, cost nine-tenths of the token-service GPUs. The reason is mechanical. Decode throughput on one accelerator is delivered memory bandwidth divided by the bytes that must be read to emit one token, and at long context that traffic is dominated by the key–value cache. Writing BW for delivered bandwidth (8 TB/s at the BGPU reference), L_m for the master’s resident context length in tokens, and b_{KV} for the KV-cache bytes per token of context, master throughput is

$$\Theta_m = \frac{\text{BW}}{L_m b_{KV}} = 115 \text{ tokens per second},$$

against $\Theta_s = 8,161$ for the 128K-context sub-agents—a 71-fold ratio driven entirely by context length against KV traffic. That ratio has a direct architectural consequence for this book: *national AI4S demand is overwhelmingly a long-context memory-capacity and memory-bandwidth problem*, which is exactly the workload the capacity-first architecture of Chapter 12 is built for and exactly the workload on which token pricing is most punishing.

The execution layer. Defined function-first and estimated operator-first:

$$G_{\text{sci}} = \underbrace{G_{\text{sim}}^{\text{AI}}}_{\text{simulation}} + \underbrace{G_{\text{tra}}}_{\text{training}} + \underbrace{G_{\text{dat}}}_{\text{data processing}}, \quad W_{\text{ag}} = N \cdot f_{\text{heavy}} \cdot d_{\text{duty}} \cdot w_{\text{heavy}} \cdot 365,$$

with $G_{\text{sim}}^{\text{AI}} = (\chi_{\text{sim}} W_{\text{ag}} + n_{\text{prog}} w_{\text{prog}}) / H_{\text{yr}}$, $G_{\text{tra}} = (1 - \chi_{\text{sim}}) W_{\text{ag}} / H_{\text{yr}}$, and $G_{\text{dat}} = F_{\text{eq}} g_{\text{fac}}$.

Every symbol, in order. W_{ag} is the annual agentic-execution workload in *delivered BGPU-hours*: the heavy-user fraction f_{heavy} (default 0.125) times the activity-by-discipline duty product d_{duty} (0.31—the share of days a heavy user is actually active, folded with the share of disciplines that draw on this workload) times w_{heavy} , the measured per-heavy-user daily driver (443 delivered BGPU-hours per active claimant-day), times 365 days, times the population N . Dividing BGPU-hours by H_{yr} —the 7,884 effective hours one installed BGPU delivers per year—converts a flow of work into a stock of machines, which is why H_{yr} appears in the denominator of the first two terms. The superscript AI on $G_{\text{sim}}^{\text{AI}}$ marks AI-driven simulation specifically, distinguishing it from the classical modelling-and-simulation demand that is accounted separately in BGPU_{BW} and never summed with it. $\chi_{\text{sim}} \in [0, 1]$ is the *simulation share*: the fraction of agentic execution that is surrogate-and-simulation work rather than model training, so χ_{sim} and $(1 - \chi_{\text{sim}})$ split W_{ag} between the first two terms. n_{prog} is the census of persistent-surrogate programs and w_{prog} their annual cost, 0.1–0.5M BGPU-hours each. F_{eq} is the national inventory of large-instrument facility equivalents and g_{fac} the installed BGPU each requires for its data-processing pipeline, 2,400–3,250 apiece [305]. An explicit substrate-eligibility matrix pins training to GPU-class systems and makes simulation the only cross-substrate routable function—which is where a country’s classical fleet can legitimately offset AI demand, and where it cannot.

26.4 Typed parameters and transfer rules

The framework’s portable core is not the arithmetic but the *parameter taxonomy*, and it is the closest methodological cousin this book has found to its own evidence grading. (One notational warning before the classes are named: this taxonomy’s labels are *local to this chapter* and are not the book’s evidence grades of Section 1.5. In particular [M] here means “must-measure,” not “author-modeled,” and [P] here means “policy choice,” not “primary-source claim.” The two systems answer different questions—the global grades ask how well a claim is evidenced, the parameter classes ask whether a number may be carried across a border—and where both appear on a page, the parameter classes are the ones attached to model *inputs*.) Every input is classed [U]




Figure 26.2: Three readings of the same framework [305]. *Top*: ten percent of tokens, generated at million-token context, consume nearly ninety percent of the token-service capacity—the configuration’s central structural fact, invariant across countries, and the reason master context length is the layer’s primary scenario axis. *Middle*: committed national fleets against two priced service rungs and the full modeled requirement, log scale. Japan’s and the United States’ commitments already exceed the uniform-baseline rung but reach only ~54% and ~70% of the tiered rung; Singapore’s clears neither. *Bottom*: the same result per researcher—announced infrastructures provision 0.02–0.10 BGPU against a modeled requirement near 1.0, an underprovision of roughly one order of magnitude for Japan and the US and closer to fifty-fold for Singapore. Japan is measurement-anchored [M/V]; Singapore and the US are pre-measurement scenarios [T] whose behavioral parameters are transferred from the Japanese probe and are labeled as such.

universal (serving physics and benchmark calibrations that transfer across borders), [M] must-measure (populations, facility inventories, electricity tariffs—country-specific by construction), [T] transferred (behavioral quantities measured so far only in one country, carried elsewhere as labeled scenario values until replaced by a domestic probe), or [P] policy (allocation choices, which are decisions rather than measurements), with an orthogonal [F] future-scenario modifier for forward technology assumptions [305]. The taxonomy states which numbers may cross a border, which must never, and which may travel only with a warning label and an expiry date.

That last category is where intellectual honesty lives. The serving physics—memory-bandwidth-limited decode of trillion-parameter MoE models, calibrated delivered efficiency, generation conversion factors—transfers exactly, because benchmarks do. What a heavy national research cohort actually does under scarcity transfers only as a hypothesis: Singaporean or American researchers may behave like the measured Japanese cohort, but no one has yet checked. A portable framework must therefore do more than parameterize the arithmetic; it must *label every input by what would make it trustworthy, and specify the experiment that replaces each assumption with a local measurement*. The seven-step measurement protocol does exactly that—population construction, facility inventory, a saturating demand probe on a $\sim 1,000$ -GPU current-generation partition with full accounting, serving-basis calibration, classical-mission accounting, cost localization, and gated publication—so that “run this analysis for country X” becomes an executable checklist rather than a research project.

26.5 What three countries find

The framework has been instantiated at both extremes of national scale (Figure 26.2). *Japan*, the fully worked measurement-anchored example: $N \approx 500,000$ open-tier researchers, central 2030 requirement 497–510k BGPU. *Singapore*, a pre-measurement scenario: $N \approx 57\text{k}$ researchers, 55–59k BGPU—the entire national requirement smaller than Japan’s proposed first-stage procurement, with the existing $\sim 1,900$ -GPU national research fleet at 1.5–2% workload coverage and a proposed 3,000-GPU Rubin-class first stage reaching 14–16% coverage and $\sim 64\%$ population reach for a four-year subtotal of \$0.39–0.44B. *United States*: $N \approx 1.1\text{M}$, 1.09–1.17M BGPU central, with the committed Equinox-plus-Solstice systems—the largest AI procurement in DOE history—standing at $\sim 10\%$ workload coverage of the modeled requirement, closely paralleling Japan’s no-new-procurement endpoint [305].

Four disciplines govern how those numbers should be read, and the third and fourth are corrections to the loose way such figures are ordinarily handled. *Installed peak, deliverable annual service, SLO-provisioned capacity, and effective utilization are four different quantities*, and a comparison that mixes them is meaningless; this chapter quotes the analytical design point and the operationally provisioned point separately wherever both exist. *Estimands must match*: delivered usage $\neq$ requested work $\neq$ latent demand $\neq$ installed capacity $\neq$ population reach $\neq$ workload coverage. *The scale-invariance finding is a hypothesis, not a result*—it holds across three instantiations of which only one is measured, so the honest statement is that the per-researcher core is *modelled* as scale-invariant because both of its terms scale linearly in N at fixed behaviour, and testing that assumption outside Japan is the first job of any second national probe. And *the Singapore and US figures should not be quoted with more precision than transferred behavioural assumptions permit*: a reasonable confidence band on the ≈ 0.94 BGPU-per-researcher core is at least $\pm 40\%$ before national measurement, which is why this chapter’s strategic conclusion is stated as an order-of-magnitude gap rather than as a percentage.

One further sensitivity deserves its own line, because it cuts against the chapter’s own conclusion. *Cheaper models may not reduce BGPU demand at all*. If capability per token rises and price per token falls, the rational institutional response is to run more agents, longer contexts, and deeper search—the Jevons argument of Section 21.4 applied to science. Whether efficiency reduces or induces national compute demand is an empirical question that the measurement

protocol can answer and has not yet; both signs are consistent with the framework, and this book's own bet, given the master-term arithmetic and the agentic direction of scientific workflows, is on induction rather than reduction.

Three structural findings survive across the $\sim 20\times$ range of national scale. *The per-researcher core is modelled as scale-invariant*: the token layer plus the agentic execution term compose to ≈ 0.94 BGPU per open-tier researcher in every country, because both scale linearly in N at fixed behavior. *The deviations are diagnostic*: what does not scale with N is the facility-and-program term, set by instrument bandwidth and pipeline depth— $\sim 5.5\%$ of the total for Singapore, $\sim 9\%$ for the United States, where the streaming-inference class alone runs 60–132k BGPU, between half and 1.2 times the entire committed AI fleet. Sizing exercises that count only central systems will structurally undercount this class. *And the headcount rule is worth a factor of six*: an OECD-consistent FTE count of US open-sector researchers is $\approx 185\text{k}$ against an eligibility headcount of $\approx 1.1\text{M}$ —a definitional choice that swamps every other sensitivity, and therefore one that any official exercise must state, defend, and apply consistently.

26.6 The universal-reach ladder, and the order-of-magnitude gap

The framework's sharpest instrument is a computed *universal-reach ladder*: the installed capacity at which 100% of researchers are served, at each of two precisely labeled service rungs. Rung one, *uniform baseline service*—every researcher receiving the baseline 150 Mtok/month tier—costs 3.2k (Singapore) / 28.3k (Japan) / 62.3k (US) BGPU installed. Rung two, *full-population tiered-token service* at the differentiated 10/20/70 entitlement profile, costs 8.3k / 72.7k / 160.0k [305].

The distinction sharpens both readings, and Figure 26.2 (middle) shows why. Japan's and the United States' committed fleets already *exceed* the uniform-baseline rung—rung coverage 137% and 180% at the annual-average convention, 106% and 139% even at SLO provisioning—yet stand at only $\sim 54\%$ and $\sim 70\%$ of the tiered rung, while Singapore's existing fleet covers under a third of even the baseline. Closing the US tiered-service gap takes roughly 19,000 Rubin-class GPUs beyond current commitments. And the full requirement—tiered tokens plus full heavy compute—costs a further $\sim 6\times$ again: 54k / 470k / 1,034k BGPU.

Stated per researcher (Figure 26.2, bottom), the result is the chapter's headline: **announced national infrastructures provision 0.02–0.10 BGPU per open-tier researcher against a modeled requirement near 1.0**—an underprovision of roughly tenfold for the United States and for Japan before its proposed first stage, and closer to fiftyfold for Singapore. One more asymmetry is worth recording because it is the reason this chapter exists: on the public record, Japan's stage-1 proposal is sized against a measured national demand model but is unfunded, while the US commitment is funded, contracted, and the largest in the world but is not yet sized against its research population. *One side has the measurement without the money; the other has the money without the measurement.*

26.7 What this means for the argument of this book

Four consequences, each connecting to a thread already running.

The token price is the science budget. Chapter 27 derives, from a workload share of $\sim 14\%$ and a 20–60 $\times$ purchased-token premium, that frontier tokens consume 76–91% of an AI4S budget (Figure 27.1). This chapter supplies the mechanism underneath that estimate: the master term is 88.7% of token-service capacity, so the expensive layer is precisely the long-context orchestration that no research workflow can avoid and that commercial pricing charges most for. A nation that rents its master layer has rented the majority of its research computing.

The requirement is a memory problem, and therefore an architecture problem. If national demand is dominated by million-token-context decode over trillion-parameter sparse models,

then the machine that serves it best is capacity-rich and bandwidth-balanced—the design point of Section 12.8 and the LX2 class of Section 12.2, not a dense-FLOP tensor part. The sizing framework and the architecture argument are the same argument in different units.

Sovereignty has a computable price, and it is high but not absurd. At delivered B200-hour costs of \$1.84–2.19 across the three countries [305], the gap between announced and required capacity is a national-budget question rather than a technological impossibility—which is exactly the condition under which pooling becomes rational. A federation that shares capacity across time zones and campaign cycles converts three separately inadequate national fleets into one adequate pooled one; that is the operational case for Chapter 27, and it now has a denominator.

And comparability is itself strategic. A country that cannot state its requirement in a unit others recognize cannot join a federation, cannot audit a vendor’s claim, and cannot tell whether a foreign price list is cheap or ruinous. The framework’s real product is not the Japanese number; it is the shared definition—which is why the invitation it extends, to run the protocol and publish the parameter file, is the least expensive strategic action available to any government reading this book.

Chapter 27

The Third Path: An International Consortium for Open Scientific AI

Chapter 23 gave each actor an instruction; this chapter builds the institution that three of those instructions require. The argument proceeds from an observation that the preceding chapters have made unavoidable: for every country that is neither the United States nor China, the two available strategies are both unsatisfactory. Permanent dependence on a small number of foreign commercial providers cedes cost, control, and continuity. Isolated national sovereign-AI programs duplicate the most expensive activity in the industry at sub-scale, and—as Chapter 21 showed—at valuations that assume rents no one may collect. The third path is neither: an international, institutionally neutral capability that pools open foundation-model development, scientific data stewardship, federated access to national compute, and a route from public research assets into industrial use. This chapter develops that design, its economics, and its relationship to the war the rest of the book describes. Its substance draws on a July 2026 working discussion between the author and Thomas Schulthess, director of the Swiss National Supercomputing Centre, whose conclusions are cited here as a discussion record rather than as a published proposal [306].

The framing claim is worth stating plainly, because it inverts the usual vocabulary: what is needed is a *global* form of sovereign capability—sovereignty not through national isolation, but through shared ownership of the knowledge, data practices, compute access, and institutional capacity required to build and operate scientific AI. A coalition that holds those things collectively has more strategic control than any of its members would have alone, and more than any of them can obtain by buying tokens from either bloc.

27.1 The economics: why the token price is the science budget

The decisive argument is financial, and it is one that neither the model-capability literature nor the geopolitics literature makes, because it requires a workload model rather than a benchmark.

Scientific computing is not a language-model workload. Measurements of a realistic AI-for-Science portfolio—traditional simulation, training and inference for non-language scientific models, search, data preparation, and agentic or frontier-LLM use—put the frontier-LLM and agentic share of *physical* workload at roughly 14%, with simulation and domain-specific AI consuming the rest. The language model is the connective layer of a scientific workflow, not the whole system [306]. That fact is usually deployed as a reassurance. It is the opposite, once the price asymmetry is applied: commercially purchased frontier tokens carry an estimated $20\text{--}60\times$ unit-price premium relative to equivalent service from infrastructure and open models under institutional control [A/M].

The consequence is arithmetic. If a fraction f_{front} of physical work is served by purchased frontier tokens at a relative unit price ρ , the share of total expenditure consumed by that minority

workload is

$$\Sigma_{\text{front}}(f_{\text{front}}, \rho) = \frac{f_{\text{front}} \rho}{(1 - f_{\text{front}}) + f_{\text{front}} \rho},$$

which for $f_{\text{front}} = 0.14$ gives $\Sigma_{\text{front}} = 76.5\%$ at $\rho = 20$ and $\Sigma_{\text{front}} = 90.7\%$ at $\rho = 60$ [306]. Table 27.1 shows the sensitivity. *A workload share of one seventh becomes a budget share of three quarters to nine tenths.*




Figure 27.1: Why the token price is the science budget. With a fraction f of physical work served by purchased frontier tokens at relative unit price ρ , the share of total expenditure that work consumes is $\Sigma_{\text{front}} = f_{\text{front}}\rho/((1 - f_{\text{front}}) + f_{\text{front}}\rho)$. At the measured $f_{\text{front}} \approx 14\%$ and the estimated 20–60 $\times$ premium, a one-seventh workload share becomes 76.5–90.7% of the budget [A/M] [306]. The curve’s shape, not its exact position, is the policy argument: at r near unity—which cheap open weights on owned hardware deliver—the problem dissolves, and the institution that solved it that way has acquired a dependency it never negotiated.

Table 27.1: Share of total AI-for-Science expenditure consumed by purchased frontier tokens, $\Sigma_{\text{front}}(f_{\text{front}}, \rho)$, as a function of the frontier-workload fraction f_{front} and the price premium ρ . Illustrative consequence of the stated assumptions [M], not an audited budget forecast.

Frontier share f	Price premium r (purchased frontier vs. self-operated open)				
	1 $\times$	5 $\times$	20 $\times$	40 $\times$	60 $\times$
5%	5.0%	20.8%	51.3%	67.8%	75.9%
14%	14.0%	44.9%	76.5%	86.7%	90.7%
25%	25.0%	62.5%	87.0%	93.0%	95.2%
40%	40.0%	76.9%	93.0%	96.4%	97.6%

Three implications follow, and they connect this chapter to the book’s central mechanism. First, *whoever sets the frontier token price sets the science budget of every institution that buys it*—which is the proposition P3 of Section 1.4, applied not to enterprise software margins but to national research capacity. A country that funds a supercomputer and then rents its intelligence layer has outsourced the majority of its own AI-for-Science expenditure to a foreign price list. Second, the same arithmetic explains why the Chinese price collapse (Chapter 10) is strategically effective without any coercion whatsoever: at r near unity—which cheap open weights on domestic hardware deliver—the budget problem simply dissolves, and the institution that solved it that way has acquired a dependency it did not negotiate. Third, it quantifies what open capability is *worth* to a public research system: not the marginal benchmark point, but the difference between spending 14% and 90% of a national AI budget on one layer.

The corresponding service design is a tiered one: a high-end path for genuinely demanding frontier use, cheaper models for routine tasks, and extensive support for the scientific long tail—where documentation, model selection, workflow integration, and knowledge assistance may expand effective coverage as much as raw accelerator capacity does [306]. The consortium’s financial case must therefore separate physical utilization from expenditure. That distinction is the policy argument, not an implementation detail.

27.2 Open pretraining as insurance, not as a benchmark contest

The second pillar is capability retention, and it requires stating an uncomfortable proposition honestly: a national or consortium open model may never be the best production choice, and building it is still correct.

The reasoning is explicitly not about winning a leaderboard. It is that the return on pretraining is *people who know how to do it*—to curate data at scale, run and debug a large training campaign, evaluate it, and operate the result. Fine-tuning or distilling someone else’s model is valuable and much cheaper, but it does not maintain the full capability, and a capability not exercised is a capability lost [306]. Switzerland’s continued investment in a fully open model is defensible on exactly this basis, as is Japan’s; so is the proposition that near-term scientific applications should meanwhile use the best available open-weight models, including Chinese ones, without embarrassment.

The insurance covers three named risks, none of which is a prediction of bad faith. *Policy and continuity risk*: China is presently a strong supporter of open release, but a future decision could restrict release, change terms, or prioritize domestic access—and this book’s own Chapter 7 documents that the endgame of the industrial playbook is precisely the conversion of dominance into gatekeeping, executed in solar, in batteries, and in rare earths. To rely on the openness of the current phase is to assume that this instance will not follow the pattern the other three did. *Verification risk*: the contents and latent behavior of externally produced models cannot always be independently established—the entire subject of Chapter 19. *Regulatory risk*: scientific domains such as biology, chemistry, materials, and robotics are inherently dual-use, and a broad restriction intended to prevent misuse can compromise legitimate research; an institution that holds its own models and its own evaluation evidence has a far stronger basis for arguing proportionate governance than one that can only make assurances about a vendor [306].

Note what this does to the book’s thesis, honestly stated: a successful consortium is a partial falsifier of the strongest form of the convergence argument. If a neutral coalition holds genuine open pretraining capability, then the commons is no longer singularly Chinese, the monoculture of Chapter 19 thins, and the certification authority has a credible home. That is the point. This chapter is the constructive branch of a book whose diagnosis is otherwise pessimistic.

27.3 Design: a shared platform with friendly competition

The organizational lesson the participants drew from China is not the one usually drawn. China’s strength in this domain was attributed in significant part to a government role that *sets the stage and then allows intense bottom-up activity*: many teams and companies compete, talent emerges outside any single central organization, and successful ideas scale rapidly [306]. The lesson is therefore not to reproduce a command structure—which is what most national AI programs instinctively build—but to create conditions under which many capable groups can participate. Chapter 7’s move 4 said the same thing in industrial terms: Beijing floods the sector and lets competition select. A consortium that funds one model team in one country reproduces the failure mode, not the success.

The HPC community supplies the working analogy: centers share software, publish results, and use common platforms while competing hard for the best science and the best performance.

An international AI framework can similarly support multiple models, scientific domains, and implementation teams on a common data, compute, and governance base. Central governance should protect openness and continuity without suppressing technical diversity [306].

Four functions define the entity, and they are the test of whether it needs to exist:

Federated compute. Countries and centers contribute capacity into a pool under common accounting rather than transferring assets to one site. The concrete proof-of-concept discussed is a joint training campaign—a next-generation open model of order 15 trillion tokens, divided among compatible systems in several countries, with more than one national partner contributing substantial resources [306]. What such a pilot buys is not primarily the model: it is shared investment, shared know-how, and a live test of cross-border accounting, scheduling, licensing, and reproducibility. Those mechanisms, once built, are reusable for everything else.

Scientific data stewardship. Cleaning, curating, documenting, and making data accountable is the service most users actually need—small and medium enterprises typically need more help with data preparation than with machine-learning algorithms—and it is the service for which science is uniquely well equipped, because reproducibility and provenance are already scientific requirements rather than compliance overhead [306].

Model production and lifecycle. Producing artifacts is not enough; regulated and socially critical uses require continuity. A clinical model may need two or three years to certify and must remain supportable five years later; users must be confident that the model, its data lineage, its interfaces, and a responsible organization will still exist. That requirement argues for multi-country ownership and explicit service obligations, and it is the strongest reason the vehicle cannot be a rotating academic project [306].

Professional operation. A secretariat and engineering organization with paid staff who can organize programs, secure commitments, operate services, assist companies, and incubate firms. A periodic meeting series is not a substitute, and the honest reason the existing landscape has a gap is that convening is cheap and operating is not.

27.4 What already exists, and what is missing

Creating another organization is justified only if the existing ones cannot execute the required functions, and the discipline of asking that question first is what distinguishes a proposal from an announcement. The relevant map [306]: model-distribution licensing has become tractable through emerging open model-distribution license frameworks, with adoption by at least one major infrastructure vendor as evidence that the legal layer is solvable; the High Performance Software Foundation covers part of the software layer but provides neither data nor compute; the International Computation and AI Network links compute providers, philanthropic funding, and socially valuable AI projects including access for less-resourced communities, and is the closest precedent for the access mission; and the Trillion Parameter Consortium holds relevant relationships and assets but has not developed the operational, financing, and industrial-incubation capacity this concept requires—in part because much of the US laboratory community is now absorbed by national initiatives including the Genesis Mission, the American Science Cloud, and a transformational-models consortium.

What is left unserved by all of them, and therefore defines the new entity: federated compute with real accounting, scientific data operations at production quality, model production with lifecycle obligations, professional program management, and industrial translation. Whether the vehicle is a new foundation, a hosted program under an existing one, a rejuvenated consortium, or a contractual federation of national centers is a question for a formal options analysis—and should be settled by which structure can perform those five functions, not by institutional preference.

27.5 Industry without expropriation, and the layered data policy

The consortium fails if it is a public-sector club, and it fails differently if it demands that companies surrender proprietary data. The resolution is layered.

The shared scientific base—models, workflows, evaluation, blueprints—is open and accountable. Above it, companies build closed, valuable layers. Firms may contribute data voluntarily in exchange for credits or preferential access; they may also retain everything and adapt the shared models privately through secure post-training, federated methods, agents, and workflow integration. Pharmaceutical and materials workflows already operate this way [306]. A single mandatory data policy would be counterproductive; an advisory-plus-open-blueprint model, in which a company receives the design and operates its own environment while retaining its secrets, is the constructive pattern.

One technical caution belongs in the record because it is frequently violated in national industrial programs: training a small isolated model on a single company’s proprietary dataset is usually the wrong architecture. Scaling behavior favors starting from a broad model trained on large scientific corpora and then adding the special data. *Science can establish the credible open base; industry builds the closed, valuable layers above it* [306]. That division is also the honest answer to the question of why a public consortium should exist at all in a market with well-capitalized private labs: the base is a public good with private complements, which is the textbook case for public provision.

The financing follows the same layering. Government infrastructure and national centers contribute systems, scientific software, models, and data; private partners make upfront commitments or bear investment risk; researchers pay usage from grants; commercial users pay for service; and an intermediary operating organization recovers cost while preserving research access [306]. Neither pure public appropriation—which Europe’s gigafactory experience suggests is insufficient on its own—nor pure commercial provision solves the problem. The unit economics must be modeled explicitly: capital sources, operating costs, reserved versus burst capacity, grant-funded use, commercial pricing, cost recovery, and safeguards for scientific access, with legal feasibility tested in parallel.

27.6 The long tail and the Global South as strategy

The phrase “long tail” carries two meanings, and the consortium needs both. One is the large population of scientists who will never train a frontier model but need AI assistance to do their work—a group for which documentation, model selection, and workflow support matter more than accelerator hours. The other is the global pool of talented people who lack access to modern compute, data, and expert networks at all [306].

The strategic framing matters more than the moral one, and it is the framing this book’s Chapter 10 would predict. Africa and comparable regions should be treated as an opportunity and a source of talent, scientific questions, data, and future partnerships—not as recipients of assistance or as a migration risk. The alternative is not neutrality: it is ceding the default to whoever does show up with a stack, and Chapter 10 documented in detail who is currently showing up, with what pricing, and inside which governance forum. A transparent path by which strong projects and people enter the ecosystem, receive suitable resources, and contribute back is simultaneously an equity program and a talent-discovery program—and, given the talent data of Chapter 8, talent discovery is the input where the returns are highest and the competition least priced.

27.7 Sequence, and the test of seriousness

The proposal becomes credible only when it moves past aspiration, and the sequence is deliberately unglamorous [306]. *Phase 1—frame the proposition*: complete the cost model of Section 27.1, map the existing organizations honestly, and issue a short founding concept whose central content is a precise statement of what is missing. *Phase 2—prove federation*: secure a small founding group—three to five countries that can commit real compute, data, people, or funds rather than endorse a declaration—and run one visible pilot combining resources from more than one country under common accounting, licensing, and reproducibility rules. *Phase 3—institutionalize and scale*: choose the legal vehicle, recruit a professional secretariat, establish the public-private financing structure, launch the industrial and Global South tracks, and negotiate multi-year vendor and government commitments.

A five-country story with leading computing centers and a visible pilot is materially more persuasive to technology vendors and governments than a general declaration with a large but inactive membership. That is also the criterion by which readers should judge whether this chapter describes a real strategy or a wish: the open design questions—institutional vehicle, founding scope, model portfolio, the precise boundary of openness across weights, code, recipes, data descriptions and evaluation, compute accounting across heterogeneous contributions, data governance including provenance and later withdrawal, vendor neutrality without vendor dependence, the public-private financing split, and who carries the cost and liability of multi-year continuity—are all unresolved, and each has a wrong answer that would render the entity ceremonial.

27.8 The institution, specified operationally

A proposal that names functions but not instruments is a wish. Ten design parameters have to be settled before anything can be signed, and stating them—together with this book’s recommended answer and the failure mode of getting each wrong—is the difference between a concept paper and an announcement.

One framing point governs all ten rows. The institution is more credible—and more useful—if it is *not* constituted as an anti-American or anti-Chinese bloc, but as a neutral interoperability, evaluation, and scientific-compute commons that both principals can transact with and neither controls. That is not diplomatic softness; it is the only configuration in which the certification function has value, since a certificate issued by one bloc about the other bloc’s artifacts is worth nothing to anybody.

27.9 Why this is the strategically decisive institution

Return to the book’s argument. Chapter 10 established that the model layer has been commoditized and that the commons is currently Chinese-led. Chapter 12 established that the compute substrate is going open and that a capacity-first architecture puts sovereign-scale training within reach of national institutions. Chapter 18 identified the certification layer as unowned. Chapter 19 showed why an unowned certification layer over a monoculture with delegated authority is a systemic risk rather than a governance inconvenience. Chapter 21 showed that the financial structures pricing permanent scarcity are the fragile part of the Western position, and Chapter 23 concluded that the correct move for everyone outside the two blocs is to build on the commons rather than to defend a rent.

The consortium is where those threads terminate. It is the entity that can hold open pretraining capability as insurance without needing it to win a benchmark; that can operate the trust stack—signed manifests, reproducible lineage, injection and backdoor evaluation on integrations, national mirrors—with credibility that neither bloc’s own institutions can claim for

Table 27.2: Operational design parameters for the consortium. Recommendations are the author’s [S]; the failure column is why the parameter cannot be deferred.

Parameter	Recommended answer	Failure mode if deferred
Legal form and host jurisdiction	Treaty-adjacent foundation in a neutral, non-aligned-capable jurisdiction with a track record of hosting technical bodies; not a national programme with foreign guests	Perceived as a bloc instrument; loses the Global South and one of the two principals
Membership classes and votes	States, computing centres, universities, and companies in separate classes; compute-and-data contribution weights votes, capped so no member exceeds a blocking share	Largest contributor becomes the owner; smaller members stop contributing
Funding formula	Public base (infrastructure in kind, assessed contributions), private risk capital for the industrial track, grants for research use, fees for commercial service	Pure appropriation is insufficient; pure fees excludes the long tail
Export-control compliance	Explicit compliance architecture per member jurisdiction, with segmented enclaves so that a control on one member does not halt the whole federation	One member’s national rule silently governs everyone
Model and data licensing	Open weights plus documented recipes under a model-distribution licence; layered data policy (open commons, voluntary contribution with credits, secure private adaptation)	Ambiguity here blocks industrial participation entirely
Compute allocation	Published accounting unit (BGPU, Chapter 26), peer-reviewed scientific allocation, reserved and burst tiers, cross-border settlement	Heterogeneous contributions become incomparable; federation degrades to a mailing list
Security certification authority	Operates the Level 0–5 ladder of Chapter 18; publishes evaluations irrespective of artifact origin	The vacancy of Chapter 20 stays vacant, or is filled by a supplier
Incident response and revocation	Named duty officer, disclosure obligations, artifact revocation and mirror withdrawal procedures, coordinated with national CERTs	An agentic incident with no owner destroys the institution’s credibility in one week
Treatment of Chinese and US technology	Strictly capability- and evidence-based: any artifact may be evaluated, adopted, or rejected on published criteria; no origin-based rule in either direction	Origin rules make the body a bloc instrument and forfeit its only durable asset, neutrality
Procurement neutrality	Vendor-neutral units and interfaces, conflict-of-interest rules for anchor partners, no single-supplier dependency in the reference stack	Anchor vendor captures the standard; members lose the sovereignty they joined for

the other's artifacts; that can pool the sovereign-scale clusters the LPDDR-class architecture of Chapter 12 makes affordable; and that can carry the multi-year continuity obligations without which no regulated deployment is possible. It is, in short, the institutional form of the only strategy this book's analysis endorses for the countries that are not superpowers: take the commons, verify it, contribute to it, and own the layer that makes it safe to use.

Whether that institution gets built in the next two years or the next ten is, on the evidence of Part I, the difference between shaping the outcome and documenting it.

Part V

Synthesis and Summary

Chapter 28

The Convergence

28.1 Assembling the elements

Chapter 9 placed the elements on the playbook before Part II examined them and scored the campaign step by step; this section states the machine they form. Read separately, each element of Part II has a Western rebuttal: the models trail the closed frontier on the hardest benchmarks; the accelerators trail NVIDIA per watt (though a GPU-free Chinese machine now leads both HPL and HPCG [182]); the fabs trail TSMC by two nodes. Read together, the rebuttals miss the structure. The elements interlock into a single machine for commoditizing artificial intelligence:

1. **Software commoditization.** Open weights (Chapter 10) plus mandated interoperability and standardized matrix extensions on an open ISA (Chapter 12) allow collaborative optimization across the whole national ecosystem—the aggressive kernel fusion that lowers the bandwidth knee of Section 12.8 is a *community* achievement, not a vendor secret. Stated with the precision the claim requires: open interfaces broaden the population able to reproduce and port optimizations, reduce the legal and architectural barriers to sharing them, and prevent any vendor from withholding the specification—*they do not eliminate the engineering cost of retargeting performance*, which remains hostage to compiler and kernel versions, microarchitecture, memory layout, precision support, collective libraries, interconnect topology, expert placement, graph compilers, and framework integration (the four portability types of Section 12.5.1). The claim is therefore about the moat’s *shape*, not its depth at any instant: CUDA’s advantage is not being crossed in one leap; the population entitled and equipped to erode it is being made as large as the ecosystem itself. The measurable milestone—now in the forecast register—is a domestic Chinese accelerator or integrated-CPU submission to MLPerf Training, whose v6.0 round (June 2026) added a DeepSeek-V3 MoE pretraining benchmark and drew 95 systems on 13 different accelerators, none of them Chinese [378]. Nonparticipation is not proof of weakness, but it is a measurable international-validation gap, and its closure or persistence is data.
2. **Hardware proliferation.** With the software requirement lowered, accelerators need not match NVIDIA—they need to be *sufficient* and *numerous*, paired with immense capacities of cheap LPDDR6 from a domestic fab (Chapter 13). Sufficiency plus volume is exactly how TOPCon modules and LFP packs won.
3. **Decentralized pretraining.** The LPDDR6 capacity advantage dissolves the centralized-supercluster premise. Frontier-class sparse models pretrainable on 128–256-socket clusters at roughly half a megawatt to one megawatt put the means of production within reach of any sovereign state, university consortium, or mid-cap firm—each one a customer for the Chinese stack and a defection from the Western API economy. The Global South adoption channel (Chapter 10) is the distribution network waiting for this hardware. The recurring

cost that decides whether a defection sticks is not training but *inference*, and there the same logic runs sharper still: §12.9.4’s priced appliance runs the largest open-weight model in existence for a fraction of closed-frontier API pricing on a five-thousand-dollar box, and at realistic utilization undercuts even a hyperscaler’s own marginal cost of serving it. Decentralized pretraining creates the model; cheap decentralized inference is what makes running it, rather than renting access to someone else’s, the obvious choice.

The Western AI economy is, structurally, a bet that intelligence will remain scarce and rentable: \$286B of annual private investment [94] capitalized against future margins on closed capability, housed in HBM-bound gigawatt data centers. The Chinese strategy is a bet that intelligence can be made abundant and free—and that whoever *makes* it abundant collects the durable prizes: the hardware volume, the standards, the talent flywheel, the adoption, and the geopolitical alignment of every country that builds on the open stack. Solar, batteries, and EVs each ended with the scarcity bet losing. The trillion-dollar data-center buildout is this war’s equivalent of Q-Cells’ 2007 capacity crown: magnificent, state-of-the-art, and aimed at the wrong future—and, in the older precedent Chapter 21 now carries, it is a fifteen-year branch line laid beside the road, priced in Section 12.9.4 at five thousand dollars, that will carry its traffic.

28.2 What would falsify this thesis

A living book owes its readers falsification criteria. The convergence argument weakens materially if, over the next two to three years: (a) closed US models open a *compounding* capability gap that adoption metrics start to follow—i.e., the 2.7% index gap widens while OpenRouter/HF shares reverse; (b) recursive AI-for-AI R&D gives the closed frontier an escape velocity that open fast-following cannot match; (c) the LPDDR-capacity architecture fails in practice—MFU ceilings prove unreachable at scale, or sparse-model quality saturates in ways dense-HBM training does not; (d) CXMT’s LPDDR6 ramp or SMIC’s advanced-node yield stalls under tightened, better-enforced controls; (e) Beijing’s anti-involution consolidation strangles the Darwinian dynamics that produced DeepSeek and Moonshot in the first place—the one playbook step that has historically destroyed value before creating it; or (f) the trust barrier hardens—Western and third-country certification regimes convert the security, censorship, and provenance record of Chapters 18 and 19 into an effective deployment wall, stalling the commons at the paying-workload boundary the way certification stalled the C919; a serious agentic security incident traceable to a widely deployed derivative would accelerate this branch by years. Symmetrically, the thesis strengthens with each frontier-class model trained end-to-end on domestic silicon; the first credible such run will be this war’s Mate 60 moment. Chapter 21’s indicator table operationalizes these criteria quarter by quarter.

A further branch belongs on this list, and it comes from Chapter 17—call it (g). If the Western commoditization catalogued there continues but the rent it displaces lands on layers the West holds—the CUDA installed base, the enterprise deployment surface, the trust and certification apparatus of Chapter 18, the deployment-generated data of Chapter 15—then P3 completes without changing ownership, and this book’s central claim fails in the most instructive way available: the mechanism was right and the beneficiary was wrong. That branch is not idle. Proposition P3 can now be satisfied by the incumbent’s own conduct, which makes it a stronger claim about mechanism and a weaker piece of evidence about who wins.

28.2.1 The strongest counter-argument: two recursive loops, not one

Branch (b) deserves more than a clause, because it is the only falsification path that would make everything else in this book irrelevant rather than merely wrong, and because the August 2026 evidence sharpened both sides of it at once.

The argument in its strongest form: if the closed US laboratories reach genuine recursive self-improvement—models that materially accelerate the research producing their successors—then compute scale compounds into capability faster than any fast-follower can copy, and the open ecosystem’s cost advantage becomes irrelevant because it is competing on a curve that is being outrun. The evidence is no longer hypothetical. OpenAI attributed its 80% cut to the cheapest tier of its ladder on 30 July 2026—the flagship was held that day—to efficiency gains that its own frontier model had helped produce, including the specific claim that the model autonomously optimized production kernels for a twenty per cent reduction in serving cost [403, 717, 395]—which is to say the same event this book cites as the first direct observation of P3’s transmission mechanism is *also* the first public claim of a model improving the economics of its own successor. Both readings are available from one datapoint, and honesty requires printing both. A commentator sympathetic to open weights put the dilemma without resolving it: open models exert real competitive pressure—or it does not matter, because “when you have more compute than anybody, the best models in the world, and those models are getting better and more efficient more quickly than anybody else, there is no competition” [395].

The reply this book offers is not that recursive self-improvement is unreal but that *there are two recursive loops running on different substrates, and they do not compete on the same axis*. The American loop runs through *algorithmic efficiency*: frontier models improving training and inference efficiency, compounding wherever compute is most abundant—which is the United States, and which is why the price cut is genuinely a strategic datum rather than a marketing line. The Chinese loop runs through *industrial capability*: models improving chip design, EDA, system architecture, and manufacturing yield (Section 10.3.5), compounding wherever the manufacturing ecosystem is deepest and the co-design partners are closest—which is China. The first loop produces more capability per FLOP. The second produces more FLOPs per dollar, domestically, out of an ecosystem that currently must lease them abroad. Both are self-improving; both are already claimed in public by their respective principals; and the Chinese side’s claim—research reproduction at 93.0 on PaperBench, with the model reportedly inventing and testing eighteen improvement ideas of its own across four rounds from nothing but a paper and a set of GPUs [395, 398]—is the same recursive structure aimed at the scientific layer that Chapter 24 argues matters most.

Which loop dominates is genuinely undetermined, and this book does not pretend otherwise. What can be said is that the two loops have different failure modes: the American loop fails if capability gains stop converting into products people will pay a premium for—the commoditization mechanism of Part II—while the Chinese loop fails if AI-assisted design cannot substitute for the tacit process knowledge that Chapter 8 shows has defeated Chinese industrial campaigns before, in aircraft and lithography. Neither failure is currently observable. The dashboard therefore now watches both: on the American side, whether efficiency-driven price cuts continue and whether frontier-model contributions to training runs become disclosed and material; on the Chinese side, whether an AI-assisted design flow produces an independently verified, manufacturable part. The first party to publish an audited instance of its own loop closing will have supplied the most important datapoint in the field, and neither has yet.

28.3 Two things the argument has to carry with it

Two of this book’s chapters change the shape of the conclusion rather than decorating it, and the synthesis is dishonest without them.

The commons has a security bill, and it falls due at scale. The convergence this chapter describes is, restated in the vocabulary of Chapter 19, a forecast of *correlated deployment*: a handful of open-weight families, thousands of unowned derivatives produced by the same quantization ecosystem that democratizes them, wired into agent frameworks holding file, shell, network, memory, and credential authority. Every published mechanism required to turn that

structure into a systemic incident has been demonstrated in controlled work—persistent backdoors surviving safety training, poisoning at near-constant sample cost, memory-to-network exfiltration with sub-1% utility loss, steganographic output channels, self-replicating promptware persisting through restarts and propagating multi-hop—and the compound events that have actually happened in the wild required no backdoor at all, only a narrow objective, a containment error, ordinary vulnerabilities, and overbroad credentials. This does not falsify the convergence thesis; monoculture risk is origin-neutral, and a US-led closed monoculture would carry the same correlation with less transparency. What it does is identify the single highest-value piece of unbuilt infrastructure in the entire competition: an authority that maintains reproducible artifact lineage across a derivative chain no vendor controls, tests integrations rather than models, and publishes results neither bloc’s marketing owns. The commons will be adopted either way. Whether it is adopted *safely* is currently nobody’s job.

There is a third path, and its economics are better than its politics suggest. Chapter 27 showed that for any country that is not a superpower the decisive number is not a benchmark but a budget: a frontier-LLM workload share near one seventh becomes three quarters to nine tenths of total AI-for-Science expenditure once purchased tokens carry a 20–60× premium. Whoever sets the frontier token price therefore sets the research budget of everyone who buys it—which is proposition P3 applied to national science rather than to enterprise software. That arithmetic makes open capability an insurance policy with a computable premium, and it makes an international, institutionally neutral consortium—federated compute, scientific data stewardship, model production with lifecycle obligations, professional operation—the cheapest available form of strategic autonomy. It is also, not incidentally, the only credible home for the certification authority the previous paragraph demands. A coalition that builds it holds something neither Washington nor Beijing can revoke; a coalition that does not will discover that “open” and “available to us, indefinitely, on acceptable terms” were never the same word.

28.4 Implications for the West—and for the rest

The precedent industries teach that the losing move is to defend the rent: tariff the modules, ban the chips, close the weights, and wait. By the time the wall is complete, the world outside it is running on the commodity. The counter-strategy that has never been tried is to *win the commoditization race*—to make the open, abundant, decentralized version of AI a Western-led commons rather than a Chinese-led one, and to compete on the layers above it. Whether any Western polity can sustain that strategy against the incentives of its own incumbents is beyond this book’s scope—but the solar, battery, and EV chapters record, three times, what happens when the answer is no.

For third countries—including advanced ones such as Japan—the convergence offers an uncomfortable but real option space: the open stack is *open*. RISC-V profiles, sparse-training recipes, LPDDR-capacity system design, and open-weight frontier models can be adopted, forked, and sovereignly operated by anyone. The half-megawatt 128-socket pretraining cluster is as available to Kobe as to Guiyang. If the AI war ends the way the precedent wars ended, the strategic question for everyone who is not the United States becomes not *whose side wins*, but *who builds fastest on the commons that the winner created*. That, and not any single benchmark, is how China will win the AI war: not by reaching the frontier first, but by making the frontier free—and owning everything beneath it.

Chapter 29

The Argument in Summary

This chapter is the book in short form. It exists because a reader who has followed twenty-eight chapters is entitled to see the argument assembled in one place with its odds attached, and because a reader who has not is entitled to know what the book concludes before deciding whether to read it. It repeats nothing that has not been argued elsewhere, and it adds one thing the preceding chapters deliberately withhold: a single statement of where the author thinks the balance of evidence now sits, with the specific observations that would move it.

Chapter 9 placed the elements on the playbook before the evidence was presented; this chapter states what the evidence did to that map. The two should be read as a pair, and where they disagree, this one is later.

29.1 The question, and the answer in one page

The question is not whether China can build a better model than the United States. That question is nearly settled and it was the wrong question: the measured capability gap at the top is a few points on an independent index, it has been a few points for over a year—three on the launch-week board of August 2026, widening to eight when Artificial Analysis re-based its index on 7 September (v4.3) and weighted the agentic tail more heavily [758]—and nothing about the contest turns on which side holds those points in a given quarter, though the direction of the re-basing is a reminder of where the closed lead lives.

The question is where the *rent* lands when intelligence becomes cheap—because a technology contest is decided by who owns the layer that stays scarce after the technology has diffused, and every one of this book’s three precedent campaigns was won that way rather than by building the better artifact. Solar, batteries and electric vehicles were not won by a superior cell, pack or car. They were won by making the artifact a commodity and owning what the commodity had to be made from.

The answer this book reaches is in three parts.

First, the strategy is real, coherent, and further along than a layer-by-layer reading suggests. Every one of the playbook’s nine moves can be found in the AI campaign in a recognizable form and at a recognizable stage (Chapter 9). Open weights at marginal price, an open harness, an open instruction set, open interconnect and framework specifications, state-arranged demand, subsidized overcapacity, Darwinian domestic selection, and an energy position that converts a grid surplus into cheap tokens—all of these are documented, dated and graded in Part II. The campaign is not a plan on paper; it is an execution in progress.

Second, and this is the finding that most changed the book, the method is no longer Chinese. Over 2026 the largest Western firms converged on the same instrument, each commoditizing a layer it does not own in order to defend the layer it does (Chapter 17). The accelerator vendor gives away model weights with their training data, and has now bought the training platform and the researchers who make them. The model laboratories design accelerators.

The hyperscalers rebuilt the scale-up fabric out of merchant Ethernet. A launch company is building a fab. This strengthens Part I considerably—a playbook that predicts the behaviour of actors it was not derived from is doing more work than one that does not—and it complicates every score in Part II, because a layer being commoditized by two parties at once tells you less about either of them.

Third, the outcome is not settled, and the book’s own thesis is the claim most at risk from its own findings. If every layer is cheap, the surplus accrues to whoever holds the layer that stays scarce; and at the currency date the scarce layers—the installed programming model, the enterprise deployment surface, the trust and certification apparatus, deployment-generated data—are still mostly Western-held. A West that commoditizes competently captures the surplus of a cheap-intelligence world as surely as a China that does. The honest statement of the book’s position is therefore conditional: China is executing the strategy that has won three technology wars, it is executing it well, and whether that wins a fourth depends on where the rent lands, which is not yet observable.

29.2 The four propositions, and where they now stand

Table 29.1 restates the four load-bearing propositions of Section 1.4 with the status the evidence supports at 9 September 2026. The probabilities are the author’s, they are scenario judgments [S] rather than measurements, and they are stated so that they can be scored against later rather than defended now.

Table 29.1: The four propositions at 9 September 2026. “Status” is the grade the evidence supports; the probability is the author’s subjective judgment [S] and is entered in the register of Appendix J.

	Claim	Where the evidence stands	Status
P1	Open weights become the global default for deployed intelligence	Majority of routed tokens; a self-enforcing ratchet that survived a Western re-entry; the gap at the top a few index points (eight on the re-based v4.3 index); enterprise <i>spend</i> share still falling while usage rises	Distribution done; monetization contested
P2	China builds a sovereign, capacity-first compute stack sufficient to serve it	Open ISA qualified as a CUDA host; a GPU-free machine at No. 1 on both HPL and HPCG; production-scale serving on a domestic cluster; HBM still the binding constraint, and the end-to-end training test not passed	Trending; integration test outstanding
P3	The rent financing the Western frontier is destroyed or relocated	Mechanism now overdetermined—the incumbent supplier funds open weights itself—but the incumbent’s own data-centre revenue grew 117% in the quarter this edition closes in, and its cheapest tier undercut the commodity and took share	Mechanism operating; magnitude contested
P4	The resulting position converts into durable strategic advantage	Endgame begun on consolidation, not on gatekeeping; the scarce layers are still mostly Western-held; deployment-generated data and trust unresolved	Pending; least evidenced

The coupling matters more than any single row. P1 without P2 is adoption without a hardware base to serve it, which the precedents say decays. P2 without P1 is a sovereign stack with nothing running on it. P3 without P4 is a price war that destroys a rent pool and hands it to nobody—which, as Chapter 17 argues, is a live possibility rather than a rhetorical one,

because the two sides are now destroying each other's rents simultaneously. P4 is the proposition on which the book's title depends and the one with the least evidence behind it, and a reader who wants to know where to attack the argument should start there.

29.3 What would change the answer

The book's discipline is that every claim states what would falsify it. The eight observations below would move the author's position more than anything else that could plausibly occur before the end of 2027, and each is entered in the register of Appendix J.

1. **A frontier-class model trained end to end on domestic silicon, packaging and memory.** The single integration test that moves P2 from "trending" to "done." Production-scale *serving* on a domestic cluster has now been claimed; training has not been demonstrated to an auditable standard, and the distinction is the whole of the remaining gap.
2. **Merchant-GPU data-centre revenue growth decelerating while merchant custom-ASIC revenue holds its guided trajectory.** Confirms the silicon-commoditization pincer of Chapter 17. The contrary observation—custom accelerators continuing to converge onto the incumbent's interconnect and memory, as one already has—confirms absorption instead, and absorption currently has the better of the evidence.
3. **A sustained fall in closed-frontier gross margin, disclosed.** P3's transmission observed on the line that matters. One price cut to a cheapest tier, with the flagship held and the counterparty raising prices in the same fortnight, is not that.
4. **Enterprise open-weight *spend* share turning upward.** The strongest single piece of evidence against P1's monetization leg is that the spend share fell during the period capability reached parity. If it turns, the deployment-layer objection weakens sharply.
5. **A ratified matrix extension for the open instruction set, and a coherent chip-to-chip standard to go with it.** Four competing matrix proposals and no ratified coherence interface are the two gaps that keep the open ISA a host rather than an accelerator architecture.
6. **Gatekeeping exercised on any AI layer.** The endgame move that has a template in batteries and rare earths and has not yet been used here. Its use would settle P4 quickly and in one direction.
7. **A called guarantee, or a disorderly repricing of compute residuals.** The exposure catalogued in Chapter 21 is disclosed, large and growing; whether it is sound is not yet demonstrated either way, and the first call would be the test.
8. **An audited instance of a research loop closing—a model materially accelerating the science that produces its successor, or the silicon that runs it.** This is the observation that would make the recursive-improvement counter-thesis of Section 28.2.1 decisive against everything else in this book, and it is the falsifier the author takes most seriously.

29.4 What to do about it

Chapter 23 argues the cases; Table 29.2 states the instructions. The common instruction beneath all of them is the one the precedent campaigns teach and the mirror confirms: *commoditize the layer you do not own before somebody commoditizes the layer you do*, and know which layer you are actually defending before you choose an instrument.

Table 29.2: The instruction by actor class, in one line each. The full argument for each is in Chapter 23; the buyer’s decision problem is in Chapter 22.

Actor	Instruction
United States	Choose between the two instruments that are aimed at the same object from opposite sides. A denial policy whose effectiveness depends on the conduct of firms the government does not control is not a strategy; decide whether the commoditization or the chokepoint is the policy, and stop paying for both.
Europe	Regulation is not production. The asset neither commons supplies is independent evaluation and certification, and it is the one Europe is institutionally best placed to build.
Japan and the middle powers	Take from both commons, depend on neither, and treat the commoditization of the scale-up fabric onto merchant parts as a dated procurement opportunity rather than a research horizon.
The Global South	The buyers write the ending. Adopt the commons—now plural—and audit the <i>provenance</i> of the openness, not only the licence: who funds it, which layer they earn on, and whether the distribution channel is independent of both.
Scientific institutions	The largest prize is the one with the weakest constituency. Build the measurement first: nobody publishes useful work per unit of compute, and until somebody does, every claim about scientific return is unfalsifiable.
China	The principal move is now being run against you. If the model layer is commoditized competently from the other side, the durable edge reallocates to energy, deployment-generated data, physical production and the domestic silicon floor—which is where it should have been planned to land in any case.

29.5 What this book is not claiming

Four disclaimers, stated at the end rather than buried, because a book with a title like this one attracts readings it does not support.

It is not claiming that China has won, or that victory is inevitable. It is claiming that a recognizable strategy is being executed competently against a stack whose rent layers are being dismantled from both sides, and that the outcome turns on an ownership question that is not yet observable.

It is not claiming that the West is passive or incompetent. The evidence of 2026 is the opposite: the incumbent priced underneath the commodity at the bottom of its ladder and took share doing it, grew data-centre revenue 117% in the quarter this edition closes in, and is executing the same commoditization strategy with more capital behind it than any Chinese actor commands.

It is not claiming that openness is a Chinese property or a Chinese virtue. Openness is an instrument, it is available to anyone with an adjacent layer to defend, and the largest single act of open-model funding in the period this book covers was American.

And it is not claiming more confidence than the evidence carries. The load-bearing weaknesses are stated in Appendix D and repeated here without softening: several mid-2026 details rest on secondary aggregators; the capacity-first architecture at the centre of Part II is a modeled design with no measured realization; the national compute sizing is measurement-anchored for one country only; and the architectural judgment that breadth compounds faster than depth is a judgment rather than a measurement, and is the claim most likely to be shown wrong.

The book is a living document. The register in Appendix J exists so that it can be scored rather than believed, and the fastest way to change its conclusions is to falsify an entry in it.

Appendix A

The LPDDR6X Socket: Engineering Reality Check

This appendix subjects the capacity-first architecture of Section 12.8 to the package-level arithmetic such a proposal must survive, and converts it into a falsifiable experimental program. Everything here is author-modeled [M] except where tagged.

A.1 The pin budget

At LPDDR6’s $12,800 \text{ Mbit s}^{-1}$ per data pin (1.6 GB/s/pin), a 7 TB/s socket requires

$$\frac{7000 \text{ GB/s}}{1.6 \text{ GB/s pin}^{-1}} \approx 4,375 \text{ active data pins,}$$

before ECC, command/address, clocking, spares, and protocol overhead—roughly 68 64-bit channels as a floor; the reconciled topology of Section A.3 rounds this up to 72 channels (18 SOCAMM2-class modules of 4 sub-channels each, 4,608 pins), buying back the ECC and protocol deduction. This is not “cheap memory instead of HBM”; it is a major package-, controller-, and board-level undertaking: memory-controller area scaling with channel count, escape routing for thousands of signals (favoring a large organic substrate or silicon-bridge hybrid over standard SOCAMM daughtercards at the top configuration), signal integrity at 12.8 Gb/s across module connectors, ECC/chipkill strategy across non-stacked commodity dies, and a memory-power envelope in the several-hundred-watt range per socket at full stream. The commercial waterline makes the gap concrete: NVIDIA’s Vera CPU—the most aggressive disclosed LPDDR server part—delivers 1.2 TB/s with up to 1.5 TB capacity [V] [257]; SK hynix’s 192 GB SOCAMM2 modules establish the module class in mass production [V] [136]. The proposal is therefore a $\sim 5\text{--}6\times$ extrapolation beyond shipped bandwidth on a shipped memory class—aggressive but not physics-limited; the analogous pin-count and channel-count regime already exists on HBM interposers, at far higher cost per bit.

A.2 What the trade buys, restated precisely

Per socket at the design point: 2.3 TB capacity (vs 288 GB HBM4 ceiling), $\sim 7 \text{ TB/s}$ peak / $\sim 5 \text{ TB/s}$ achieved requirement after recomputation elimination and aggressive fusion, 67% MFU ceiling (vs 89% HBM4), cost-per-bit an order of magnitude below HBM, and supply from the one memory class where the domestic producer is at the world’s leading edge (Chapter 13). The capacity threshold that makes this *enabling* rather than merely cheap is model-specific: for the 2.8T-parameter sparse-MoE reference workload under the stated parallelism layout, sharded state exceeds HBM4 socket capacity below ~ 256 accelerators [134]. Different models, optimizers,

or layouts move the threshold; the appendix’s claim is the existence and economic significance of the regime, not a universal wall.

A.3 The reconciled design point

The sharpest engineering demand any such analysis must meet is that every quantity in it derive from *one declared topology* rather than being chosen independently. Table A.1 is that reconciliation. Everything below the first three rows is *derived*, not assumed; changing a declared row changes every derived one.

Table A.1: The reconciled LPDDR6X socket, one topology [M]. Declared parameters first; every subsequent row derives from them.

Parameter	Value	Basis
Module count (declared)	18	SOCAMM2-class
Capacity per module (declared)	128 GB	module class in mass production [V] [470]
Subchannel topology (declared)	4×64-bit per module	64×16 Gb dies per module SOCAMM2-class organization
Total capacity	2.304 TB	18 × 128 GB
Total channels	72 × 64-bit	18 × 4
Aggregate active data pins	4,608	72 × 64
Signaling rate (declared)	12.8 Gb/s/pin	LPDDR6X target rate
Raw bandwidth	7.37 TB/s	4,608 × 1.6 GB/s
ECC + protocol efficiency	×0.90–0.93	in-band ECC, refresh, turnaround
Effective peak	6.6–6.9 TB/s	derived
Sustained, streaming-dominant	5.4–6.2 TB/s	82–90% of effective peak
Sustained, random expert-fetch	3.5–4.5 TB/s	the assumption most likely wrong
Requirement (knee, after recomputation elim.)	~5.0 TB/s	[134]
Memory-tier power at full stream	300–480 W	59 Tb/s × 5–8 pJ/bit
Controller + PHY area	~70–120 mm ² (est.)	72 channels; large uncertainty
Package and escape	100–110 mm organic substrate, peripheral escape; silicon bridges at the corners if routing fails	declared physical implementation

Read the two “sustained” rows against the requirement row and the design’s honest status is visible in one glance: the socket *clears* the 5 TB/s knee if and only if the access pattern is streaming-dominant after kernel fusion—which is precisely milestone (3) of the falsification program below, and why it is listed as the claim to kill first. If expert fetches behave as random access at production batch sizes, the pure-LPDDR corner misses its knee by 10–30% and the hybrid of Section A.4 stops being the recommended default and becomes the only viable configuration. This is why the hybrid is the *baseline* throughout the book and pure LPDDR is the sanctions-contingent research corner (Section 12.8). For scale calibration: the shipped commercial waterline (NVIDIA Vera, 1.2 TB/s [V] [471]) makes this a ~6× bandwidth extrapolation on a commercial memory class—aggressive, not physics-limited, and now stated from one topology.

One more decomposition is required, because the table’s power row understates what a procurement must budget. The 5–8 pJ/bit figure is *transferred-bit* energy; the full memory-subsystem envelope is

$$P_{\text{memory}}^{\text{el}} = P_{\text{DRAM}}^{\text{el}} + P_{\text{PHY}}^{\text{el}} + P_{\text{controller}}^{\text{el}} + P_{\text{PMIC}}^{\text{el}} + P_{\text{refresh}}^{\text{el}} + P_{\text{board}}^{\text{el}} + P_{\text{idle}}^{\text{el}},$$

and the non-transfer terms are not small at this scale: PHY and controller power for 72 channels plausibly add 60–120 W; refresh on 2.3 TB of DRAM adds 15–30 W continuously whether or not a bit moves; PMIC conversion losses at the module level run 5–8% of delivered power; and board and connector losses at 12.8 Gb/s signaling add several tens of watts. The honest memory-subsystem envelope is therefore 450–700 W at full stream rather than the 300–480 W transfer-only figure—which flows through Table 12.4 as roughly +0.1–0.2 kW per socket and keeps the 128-socket facility inside ~ 0.5 – 0.7 MW. The remaining detail a package-complete specification owes—floorplan, controller placement, channel lengths, connector loading, command/address topology, ECC organization, simultaneous-switching-noise and power-integrity margins, module replacement strategy, cooling distribution, bill of materials, and yield-adjusted cost—stays on the package-complete checklist below, unpaid until milestone (1) funds it.

A.4 The hybrid tier, which is probably the better design

One observation about the memory hierarchy is strong enough to change the recommendation outright. The capacity-first architecture may be most compelling not as a pure HBM *replacement* but as a *tiered* design: a modest HBM tier for hot tensors, attention working set, and communication-critical state, with a very large LPDDR tier for expert weights, optimizer state, and cold activations. On the reference sparse-MoE workload the split is natural, because the objects differ by an order of magnitude in reuse: attention and dense-path tensors are touched every token, expert weights are touched by a routed minority of tokens, and optimizer state is touched once per step.

The engineering case for the hybrid is strong. It delivers most of the capacity advantage—enough to clear the sub-256-node state wall—at a fraction of the LPDDR pin count, because the bandwidth-critical traffic is served from HBM at 4–8 TB/s while the LPDDR tier need only sustain expert-fetch and optimizer traffic at perhaps 1.5–3 TB/s. That collapses the package problem from $\sim 4,400$ active data pins to something closer to 1,000–1,900, brings the escape routing back into conventional substrate territory, halves the memory-power envelope, and removes most of the signal-integrity risk at 12.8 Gb/s across module connectors. It also degrades gracefully: a hybrid part remains useful if the LPDDR tier underperforms, whereas a pure-LPDDR part does not.

The strategic case is more mixed, and the tension should be stated rather than hidden. A hybrid design reintroduces an HBM dependency—precisely the input Chapter 13 identifies as the binding domestic constraint—which is why a Chinese programme might rationally prefer the pure-LPDDR corner despite its worse engineering risk, while a Japanese or European programme with allied HBM access should almost certainly build the hybrid. The architecture and the sanctions geography therefore select different points on the same continuum, which is itself an instance of this book’s central claim that allocation, not category, is what distinguishes machines.

A.5 Tensor placement and migration policy for the hybrid

A hybrid tier is only as good as its placement policy, and the policy must be stated, with its failure mode, rather than assumed. Table A.3 gives the static assignment for the reference sparse-MoE training workload; the reuse asymmetry that justifies it is roughly three orders of magnitude between the hottest and coldest classes.

The two migrating classes carry the risk. Expert migration must run at the *popularity drift rate*, not the access rate: with per-step expert-popularity churn of a few percent (measured on public MoE traces), promotion/demotion traffic is bounded by a few tens of GB/s—negligible against the tier bandwidths—*provided* migration granularity is whole-expert and hysteresis prevents thrashing. The failure mode is a workload whose routing entropy is high and unstable, where the hot set exceeds HBM capacity and migration bandwidth is consumed re-shuffling it;

Table A.2: Three points on the memory continuum for the reference workload [M]. The hybrid is the recommended engineering default; the pure-LPDDR corner is the sanctions-optimal one.

	HBM-only	Hybrid (recommended)	LPDDR-only
Capacity per socket	288 GB	~0.3 TB HBM + 1.5–2 TB LPDDR	2.3 TB
Peak bandwidth	14–16 TB/s	6–8 (HBM) + 1.5–3 (LPDDR)	~7 TB/s
Active LPDDR data pins	—	~1,000–1,900	~4,400
Sub-256-node state wall	binding	cleared	cleared
MFU ceiling	~89%	~80–85% (est.)	~67%
Cost per bit	worst	intermediate	best
Domestic-supply exposure	HBM-dependent	partially HBM-dependent	LPDDR only

Table A.3: Tensor placement for the hybrid memory tier [M]. “Migration” marks classes whose residence changes during training.

HBM tier (hot, ~0.3 TB)	LPDDR tier (capacity, 1.5–2 TB)	Migration
Hot dense layers (attention, shared experts)	Cold experts (routed minority)	experts: yes
Attention state with high reuse	Long-tail KV state	KV: yes
Collective-critical tensors	Full optimizer state	no
Communication buffers	Checkpoint staging	no
Frequently accessed optimizer shards	Low-reuse activations	no
	Datasets and retrieval indexes	no

in that regime the hybrid degrades toward pure-LPDDR performance while paying HBM cost, which is why the falsification program’s milestone (5) exercises the all-to-all *and* the migration policy together. The hybrid is advantageous exactly when placement is stable enough that tier movement does not consume the bandwidth it saves; that sentence is the policy’s entire acceptance criterion, and it is measurable.

A.6 Package yield: where capability and competitiveness diverge

One more equation belongs here because Chapter 13’s advanced-packaging ledger needs it. To first order, package yield compounds multiplicatively across every die and interface:

$$Y_{\text{package}} \approx Y_{\text{logic}} \left(\prod_{i=1}^n Y_{\text{memory},i} \right) Y_{\text{interposer}} Y_{\text{bond}} Y_{\text{test}}, \quad (\text{A.1})$$

where n is the number of memory stacks in the package. The exponent is merciless. Count the factors for a representative AI accelerator—one logic die, eight memory stacks, one interposer, one bonding step, one test step—and the product runs over twelve terms: at 99% yield per step the package yields $0.99^{12} \approx 89\%$; at 97% it yields $0.97^{12} \approx 69\%$, and the cost per *good* package rises by the reciprocal, roughly 45% above the 99% case. This is the precise mechanism by which “technically capable” and “economically competitive” diverge—a lesson the packaging ledger should score separately for TSV formation, wafer thinning, hybrid bonding, interposers, large organic substrates, known-good-die test, thermal interfaces, high-speed SerDes, and package-level repair, because a national program can demonstrate every step individually and still lose the product economics to Eq. (A.1). It also quantifies the hybrid design’s advantage over HBM-maximal packages: fewer stacked interfaces per socket means fewer terms in the product.

A.7 The package-complete checklist

A publishable architecture specification, as opposed to a capacity argument, must state all of the following; this book states the first six and owes the rest: exact memory-channel and module topology; package edge length, bump and pin density, and board escape assumptions; controller PHY area and power; ECC, sparing, chipkill, and repair model; sustained random versus streaming bandwidth (the reference workload’s expert fetches are closer to random than to streaming, and this is the assumption most likely to be wrong); memory power at 5 and 7 TB/s; expert-routing and optimizer communication volume; checkpoint time and storage bandwidth at full state; failure-domain and restart cost; total socket, board, and cluster cost; and side-by-side comparison against HBM4, SOCAMM2/LPDDR5X, CXL-attached memory, and the hybrid of Section A.4.

A.8 The falsification program

The architecture graduates from [M] to [V] through seven measurable milestones, each cheap relative to what it de-risks: (1) memory-controller and package co-simulation at the 68-channel design point (signal integrity, power, area); (2) a multi-controller FPGA/chiplet emulation of the memory system; (3) measured sustained bandwidth under fused transformer training kernels—the 5 TB/s knee is the claim to kill first; (4) measured energy per byte and per training token against an HBM baseline; (5) an 8–16-node expert-parallel prototype exercising the all-to-all at the modeled traffic; (6) a 128-node pretraining demonstration on a mid-scale sparse model with checkpoint/failure-recovery at full state size; (7) an audit of achieved MFU and cost against this book’s model, published either way. A national program that funds milestones 1–4 buys, for single-digit millions, the answer to whether the trillion-dollar HBM assumption is load-bearing.

Appendix B

High Bandwidth Flash: A Technical Dossier

Section 12.9 rests part of its argument on a memory technology that does not yet exist as a shipping product, and Section 12.9.4 prices a machine built around it. A book that does that owes the reader a full account of what is actually known about the technology, what is claimed, what is modelled, and what is not published at all—in one place, graded, and dated. This appendix is that account, current as of 9 September 2026. It is organized so that a reader can check the load-bearing assumptions of Chapter 12 against the primary record: the lineage of the idea, the announcement history, the specification as published, the power and endurance physics, the cost structure, the perceived advantages and their strongest rebuttals, the competitive landscape, the China angle, and the research literature. The short version is that HBF is real as a specification, absent as silicon, plausible as a weights tier, implausible as anything else, and priced—by its vendors—in adjectives rather than numbers.

B.1 What it is, and why this book cares

High Bandwidth Flash is NAND flash reorganized to behave like a capacity-first cousin of High Bandwidth Memory: sixteen (or eight) NAND dies stacked with through-silicon vias on a logic base die, in a package that matches HBM4’s footprint and stack height, read in parallel across many independent channels so that a single 512 GB stack delivers on the order of a terabyte and a half per second—HBM3E-class read bandwidth at eight to sixteen times HBM’s capacity—and connected to the host processor over UCIe rather than PCIe [586, 569, 570]. Removing the SSD controller and the PCIe link is what turns flash from a storage device into a memory tier: a PCIe 5.0 SSD delivers about 14 GB/s and an NVMe path costs on the order of a hundred picojoules per bit read [578], whereas the HBF proposal is a hundred times the bandwidth at what its proponents model as a tenth of the energy per bit. It is non-volatile, so it draws no refresh power and holds a model across power cycles. It is slow to write and wears when written, so it is proposed for weights and read-mostly context, not for key–value cache or activations. That is the entire concept, and every argument about it is an argument about whether those five properties—capacity, bandwidth, energy, non-volatility, write asymmetry—land where the vendors say they will.

The book cares because a terabyte-class weights tier at flash cost is the specific mechanism by which Section 12.9 argues the frontier open-weight commons could move from the datacentre to the server closet, and because the appliance priced in Section 12.9.4 carries four such stacks. If HBF ships at its specification, that machine is buildable; if it ships at half its bandwidth or twice its energy, the machine is slower or hotter but still viable; if it does not ship, the argument falls back to LPDDR-only capacity tiers, which Appendix A shows are adequate but roughly three times costlier per gigabyte. The dependence is therefore real but not total, and this appendix is

careful to say which conclusions of the main text hang on which unverified numbers.

B.2 Lineage: three earlier attempts to make flash behave like memory

HBF is the fourth serious attempt to put non-volatile memory into the latency-and-bandwidth band between DRAM and storage, and its designers have said as much. Toshiba Memory (now Kioxia) introduced XL-FLASH, a single-level-cell storage-class memory with roughly five-microsecond read latency, at the Flash Memory Summit of August 2019, and it remains in production as the fast tier of Kioxia’s enterprise SSDs and, since 2026, of a CXL memory module [612] [V]. Intel’s Optane—3D XPoint, the most ambitious of the attempts and the only one to ship as a DIMM—was wound down in July 2022 with a \$559M inventory impairment, having failed to find a market large enough to fund a proprietary process at two vendors [613] [V]. Samsung revived its Z-NAND at FMS 2025 as a GPU-direct storage tier, claiming fifteenfold the performance and 80% lower power than conventional NAND [614] [A]. The lesson the industry drew from Optane, and which every critic of HBF cites, is that a memory technology owned by one or two vendors and requiring a new controller, a new interface and a new software tier faces an adoption problem that its physics cannot solve; the lesson HBF’s proponents draw is different—that Optane failed on *cost per bit*, where it sat awkwardly between DRAM and NAND, whereas HBF is built from the cheapest bits in the industry and asks only for packaging that HBM has already made routine [615] [A].

What is new in HBF is not the flash but the assembly. Three enabling pieces arrived together in 2024–25: SanDisk’s CBA (“CMOS directly Bonded to Array”) architecture, which places the peripheral logic on a separate wafer bonded to the memory array and thereby frees the array die for the many independent, concurrently addressable sub-arrays that high aggregate read bandwidth requires; through-silicon vias through NAND dies, which HBM had made an industrial process for DRAM but which had not previously been productized for flash; and the UCIe die-to-die standard, which gave a package-internal interface that neither HBM’s PHY nor PCIe provided [616, 617] [V/A]. A fourth piece, the sixty-fold gap between what an HBM stack and a NAND stack cost per gigabyte, is what makes the assembly worth doing.

B.3 The announcement record, February 2025 to August 2026

Table B.1 sets out the public record. Three features of it deserve comment before the specification itself. First, the timeline has slipped in the way that first-of-a-kind memory timelines slip: in August 2025 the vendors said samples in the second half of 2026 and first inference devices in early 2027; by April 2026 SanDisk’s chief executive was saying “the NAND late this year, and . . . more of a system with the controller early to mid next year”; by the August 2026 investor day the deck said “first HBF inference product samples 2027,” and the sell-side was writing 2028 for production [616, 630, 637, 638] [V/A]. Second, the coalition has widened but not to the parties that matter most: Google (through DeepMind, whose Xiaoyu Ma co-authored the Patterson paper that gives HBF its academic case) and Tenstorrent joined the OCP consortium in August 2026, and Meta was reported to have joined at the investor day; NVIDIA, AMD, Samsung, Micron and Kioxia have made no public commitment, and Kioxia—SanDisk’s own fab partner—is pursuing a separate high-bandwidth-flash path with NVIDIA [569, 637, 629, 635] [V/A]. Third, the one hard milestone is recent and unambiguous: SanDisk stated on 13 August 2026 that the first HBF memory die has taped out [637, 638] [V]. Silicon exists; product does not.

Table B.1: High Bandwidth Flash: the announcement record. V = verified against the primary page; A = trade or analyst reporting; R = roadmap statement.

Date	Event	What was said	Grade
Aug. 2019	Toshiba/Kioxia XL-FLASH at FMS	SLC storage-class memory, $\sim 5\mu\text{s}$ latency; the flash-as-memory precedent [612]	V
Jul. 2022	Intel winds down Optane	\$559M impairment; the two-vendor cautionary tale [613]	V
11 Feb. 2025	SanDisk Investor Day: first HBF disclosure	“Match the bandwidth of HBM . . . 8 to 16 times capacity at a similar cost point”; 512 GB per 16-high stack; “4 TB of VRAM” on a GPU [618, 619]	V/A
Jul. 2025	SanDisk HBF Fact Sheet	Gen. 1: 512 GB, 1.6 TB/s; Gen. 2: up to 1 TB, $>2\text{ TB/s}$, $0.8\times$ power; Gen. 3: 1.5 TB, 3.2 TB/s , $0.64\times$; “non-volatile, no refresh power needed” [586]	V
24 Jul. 2025	HBF Technical Advisory Board	David Patterson (chair) and Raja Koduri; Jim Keller added by Aug. 2026 [620, 637]	V
6 Aug. 2025	FMS 2025: SanDisk–SK hynix MoU	“establish specifications”; samples 2H CY2026, first inference devices early 2027; Best of Show [621, 616]	V
20 Aug. 2025	Kioxia 5 TB / 64 GB/s flash module	PCIe 6.0 daisy-chained “beads,” $<40\text{ W}$, NEDO-funded; not HBF, the alternative path [622]	V
Oct. 2025	OCP Global Summit	SK hynix “AIN B” product family “leverages HBF technology”; joint “HBF Night” [623]	V
Nov. 2025	Trade analysis	“two or more years from commercial availability”; NVIDIA involvement “essential” [624]	A
Dec. 2025	YMTC at its developer conference	HBF fits Xtacking; “full speed toward HBF next” (DigiTimes) [625]	A
8 Jan. 2026	Ma & Patterson, arXiv:2601.05047	HBF as “ $10\times$ memory capacity with HBM-like bandwidth”; ballpark stack 512 GB, 1.6 TB/s , $<80\text{ W}$ [626]	V
Jan. 2026	SK hynix H ³ paper (IEEE CAL)	HBM+HBF hybrid: $2.69\times$ perf/W, $18.8\times$ batch at 10M-token context; assumes $\sim 100\text{k}$ write cycles, up to $4\times$ HBM power [627]	V
25 Feb. 2026	OCP consortium launched	“dedicated workstream under OCP”; SanDisk and SK hynix founders [628]	V
20 Mar. 2026	Kioxia’s own path	up to 32 TSV-stacked dies on a shared interposer beside HBM, “with Nvidia”; XL-Flash the candidate cell [629]	A
30 Apr. 2026	SanDisk Q3 FY26 call	“the NAND late this year . . . a system with the controller early to mid next year”; “not a substitute for an enterprise SSD” [630]	A
Jun. 2026	Equipment and IP	Hanmi first HBF TC bonders 2H26; SanDisk patent bonding a processor onto a CBA NAND tile with HBM on a shared interposer [631]	A
28 Jun. 2026	Micron at HotInfra (ISCA)	roofline: on-package LPDDR beats HBF at every operating point for datacentre decode; HBF 8.5 pJ/bit , $20\mu\text{s}$ [632]	V
11 Jul. 2026	FlashAccel, arXiv:2607.10186 (CAS)	6 HBF stacks + HBM: $2.54\times$ throughput per GPU, $1.93\times$ energy efficiency at a 100 ms SLO [633]	V
3–5 Aug. 2026	OCP specification v0.7.0; FMS 2026	8-Hi/16-Hi to 512 GB; three grades $\sim 0.4/1.5/3.0\text{ TB/s}$; UCIE; Google and Tenstorrent join [569, 570, 634]	V
5 Aug. 2026	Samsung and Kioxia at FMS 2026	Samsung zNAND-O (4/8-die NAND stack for on-device NPU) and zHBM; Kioxia stays out of OCP HBF [636]	V/A
13 Aug. 2026	SanDisk Investor Day	“First HBF memory die taped out”; “first HBF inference product samples 2027”; “4 HBF GPUs = 8 HBM GPUs” [637]	V
14 Aug. 2026	Sell-side and critics	Goldman, Cantor bullish; Counterpoint/Wedbush: unmodelled upside; critics: HBM4E baseline, BF16 assumption, no endurance figure [638, 639, 640]	A

B.4 The specification, as published, and what it omits

The document released on 3 August 2026 is the *OCF HBF Architecture Specification, Version 0.7.0*, subtitled a high-level base-die specification—not a 1.0, and developed, by SanDisk’s account, in about six months inside an OCF storage workstream chosen over JEDEC because its “workstreams are highly goal-oriented” and can “iterate on the specification in real-time” [634, 617] [V/A]. What it fixes, per the vendors’ own summaries and the trade outlets that read it, is set out in Table B.2, and what it does not fix is the more important half of the table.

Two things in that table matter more than the rest. The first is that every number a system designer would need to commit to a design—latency, granularity, write bandwidth, watts, endurance, retention, price—is either absent or someone else’s assumption; the vendors have published capacity, bandwidth grades and an interface, and described everything else in comparatives. Contemporaneous critics made the same point more sharply: the OCF document carries no version discipline beyond “0.7.0,” no conformance procedure and no certification scheme, and its bandwidth and capacity figures are two vendors’ development targets rather than requirements [571] [A]. The second is that the roadmap itself is stated inconsistently between SanDisk’s own documents—the July 2025 fact sheet says Gen. 2 reaches 1 TB and Gen. 3 reaches 1.5 TB, while the roadmap SanDisk gave Semiconductor Engineering in May 2026 says $1.5\times$ and $2\times$ Gen. 1 capacity, i.e. 768 GB and 1 TB [586, 617]. Section 12.9 quotes the fact sheet; the reader should know that the more recent statement is the more conservative one, and that the difference is a full generation of capacity.

B.5 Power, energy, and the physics of reading flash fast

The energy question is the pivot on which the whole proposition turns, and it is worth being exact about who has measured what. Nobody has published a measured energy per bit for a TSV-stacked NAND stack read over UCIE; there is no such stack outside a fab. What exists is three kinds of evidence. At the bottom, the measured datacentre case: a Seoul National University analysis of streaming mixture-of-experts weights from commodity SSDs found the flash path costing about 102 pJ/bit against HBM’s 4.2, and concluded that flash needs a roughly tenfold improvement—to about 10 pJ/bit—to be energy-competitive [578] [V]. Commodity NVMe cannot do this job, and the appliance of Section 12.9.4 does not use it. In the middle, the vendors’ relative statements: SanDisk’s fact sheet says HBF’s power profile is “similar” to HBM4’s and that Gen. 2 and 3 draw 0.8 and 0.64 times Gen. 1; SK hynix’s own H³ paper, less generously, models HBF at up to four times HBM power at the same bandwidth [586, 627] [V]. At the top, the third-party models that assume the on-package figure: Ma and Patterson’s ballpark of a 512 GB stack at 1.6 TB/s inside 80 W, which is about 6 pJ/bit; the Chinese Academy of Sciences’ FlashAccel at 8 pJ/bit; Micron’s HotInfra roofline at 8.5 [626, 633, 632] [V]. Those three converge on two to three times HBM per bit read, which is precisely the improvement the SNU paper says is required, and precisely the assumption Section 12.9.4 carries—7 pJ/bit, 360 W of flash reads at 6.4 TB/s. The reader should hold that number as a modelled consensus, not a measurement, and should note that all three modelling groups are, in different degrees, parties to the outcome: Google is a consortium member, the CAS group is designing for it, and Micron is a competitor arguing that LPDDR wins.

Where the physics is settled is in what flash cannot do. A NAND page read takes tens of microseconds for triple-level cells and a few microseconds for single-level; the microsecond latency is intrinsic, and every HBF architecture study hides it rather than removes it—SK hynix’s H³ needs a 40 MB SRAM latency-hiding buffer with a hit rate above 80%, FlashAccel reads in ~ 19 MB “hyper-pages” behind SRAM prefetch, and Micron’s model finds only 36% of nominal bandwidth survives scattered access [627, 633, 632] [V]. Weights, read in a deterministic order decided by the model graph, are exactly the access pattern that tolerates this; John

Table B.2: HBF as specified and as claimed. “Published” means stated by SanDisk, SK hynix or the OCP document; “modelled” means a third-party assumption; “none” means no figure of any kind has been published as of 9 September 2026.

Quantity	Value	Status	Source
Capacity per stack	up to 512 GB (16 × 256 Gb dies); 8-Hi and 16-Hi	published	[570, 586]
Roadmap capacity	Gen. 2 “up to 1 TB,” Gen. 3 1.5 TB (fact sheet); 1.5× / 2× (SemiEngineering)	published, inconsistent	[586, 617]
Read bandwidth grades	~0.4 / 1.5 / 3.0 TB/s (0.384 / 1.536 / 3.072); Gen. 1 product 1.6 TB/s	published	[570, 634]
Host interface	UCIe, reported up to 64 GT/s over 64 lanes (~400–500 GB/s per module)	published / reported	[635]
Base die	up to 16 independent full-duplex host channels, 1–4 AXI channels each (up to 64); each host channel owns a set of NAND dies	reported from spec	[634]
Package	“closely matching the physical footprint, power profile, and stack height of HBM4”; no dimensions	published, qualitative	[586]
Interposer	not stated; concepts show HBM and HBF stacks side by side on one interposer	modelled	[629, 627]
Cell type / node	not published; sources say SLC-like; BiCS generation undisclosed	none	[572, 634]
Read latency	“higher than HBM”; microsecond range; modelled at 4 μs (SLC) to 20 μs	none from vendor	[633, 632]
Read granularity	“larger page size than HBM”; modelled ~128× HBM4; 4 KB pages	none from vendor	[572, 658]
Write bandwidth	supported (“read/write user guide”); no figure; NAND program ~75–85 MB/s per die	none	[570, 659]
Power	“similar power profile” to HBM4; Gen. 2 0.8×, Gen. 3 0.64× Gen. 1; no watts, no pJ/bit	relative only	[586]
Modelled energy per bit	6 (Ma–Patterson), 8 (FlashAccel), 8.5 (Micron) pJ/bit; HBM 3.0–4.2	modelled	[626, 633, 632, 578]
Standby	“non-volatile, no refresh power needed”	published	[586]
Endurance	no P/E, TBW or read-disturb figure; H ³ assumes ~100k cycles, FlashAccel 1M	none	[627, 633]
Retention at package temperature	no figure; 3D NAND retention loss accelerates superlinearly with temperature	none	[642]
ECC / reliability scheme	spec has a “reliability” section; contents not public	none	[569]
Stack price	“similar cost” to an HBM stack; “one of the lowest cost per bit”	adjectives	[586, 618]
Samples / production	first die taped out Aug. 2026; product samples 2027; production 2028 (sell-side)	published / A	[637, 638]

Carmack’s remark that inference “can have a deterministic memory access pattern ... could tolerate cold-start latencies in multiple milliseconds” is the practitioner’s version of the point [572] [A]. Key-value cache, written every token, is exactly the pattern that does not, and the August 2026 characterization by a Peking University group—HBF used for KV cache raising latency 2–5.5 times and cutting throughput 1.1–2.7 times—is the negative result everyone in the field expected [641] [A; paper reported, not fetched]. And retention is a thermal problem before it is an endurance one: the HeatWatch study of 3D charge-trap NAND found retention loss accelerating superlinearly with temperature, which matters for a die stacked next to a kilowatt-class GPU and is one reason no vendor has yet published a temperature specification [642] [V]. The duty-cycle argument of Section 12.9—that a non-volatile tier’s zero standby draw offsets its higher active draw for a machine that is not busy around the clock—does not depend on any of the unmeasured numbers, and it is the part of the energy case this book is most confident of.

B.6 Endurance: the number nobody has printed

No vendor has published a program-erase cycle count, a terabytes-written figure, a read-disturb rate or a retention specification for HBF, and this appendix has confirmed the absence across every primary and secondary source it could reach. What exists is inference from the cell type and from the papers. If the cell is single-level, as multiple sources report, endurance on the order of 100,000 cycles is conventional and is what SK hynix’s own H³ paper assumes; FlashAccel assumes a million by relaxing retention, which is a design choice rather than a measurement; a triple-level cell would sit at three to eleven thousand [627, 633, 617] [V/A]. For a weights tier that is written once per model load and read continuously, none of those figures is binding: an appliance that reloads a 1.4 TB checkpoint even daily performs a few hundred full writes a year. Read disturb, the slow corruption of neighbouring cells by repeated reads, is the mechanism that actually applies to a read-mostly tier, and the reclaim writes it forces are wear; the STRAW work at SNU shows the disturbance is strongly non-uniform across word lines and that stress-aware reclaim cuts those writes by more than 80%, which is the sort of controller-level engineering the base die exists to do [643] [V]. SanDisk’s own vice-president acknowledged the write-wear problem at FMS 2026 and said the risk “diminishes if HBF operates primarily in read-heavy roles” [644] [A]. That is the correct statement of the case, and it is why the main text calls HBF a weights tier and never a memory replacement.

B.7 Cost structure: cheap bits, expensive stack

The cost argument is where the vendors’ adjectives and the book’s numbers must be reconciled, and the reconciliation is uncomfortable in both directions. Start with the bits. TrendForce’s spot series for 512 Gb TLC NAND wafers ran \$2.80 in August 2025, \$9.61 in December 2025 and \$21.13 on 9 August 2026—roughly \$0.044, \$0.15 and \$0.33 per gigabyte—so the raw die content of a 512 GB stack was about \$22 at the pre-surge floor and about \$170 at today’s spot [594] [V]. Custom 256 Gb dies with through-silicon vias and possibly single-level cells cost several times commodity per bit, but the order of magnitude is unchanged: the die is not where the money is. The money is in the stack. TSV formation, thinning and stacking of sixteen or seventeen dies, the CBA base die with its sixteen host channels and their controllers, known-good-die test and burn-in are what SanDisk means when it says the stack costs about what an HBM stack costs, and what the Chipstrat analysis means when it puts HBF at roughly a tenth of HBM’s cost per gigabyte but comparable per stack [586, 572] [A]. An HBM4 12-high stack was reported at about \$500–560 in late 2025—\$14–16 per gigabyte—so “similar cost” would put a first-generation HBF stack near \$500, or about a dollar a gigabyte [645] [A].

Section 12.9.4 prices the stack at \$150–250. The gap needs to be stated plainly rather than hidden. Three arguments carry it, in descending order of strength. First, HBM’s \$500 is a *price*

in a market where HBM is the scarcest component in the industry and its vendors book record margins; it is not the cost of stacking sixteen dies. A finished 1 TB NVMe SSD—controller, DRAM, PCB, enclosure, retail margin—sold for \$65–100 before the 2026 surge, and the operations that turn commodity NAND into an HBF stack are not obviously more expensive than the ones that turn it into a retail SSD, only different [669]. Second, the book’s basis is declared to be pre-surge, and the vendors’ “similar to HBM” was said in a market where everything memory-related has doubled or tripled since. Third, the appliance is a volume consumer product, and its stack does not compete with HBM for the same CoWoS interposer slots that a datacentre package’s HBF would. Against those, the honest counter is that no HBF stack has been priced by anyone, that the vendors’ one statement points higher, and that a two-vendor part in its first generation carries the margin a two-vendor part carries. So the sensitivity is given here rather than argued away: at \$500 per stack the appliance’s base case rises from \$5,000 to about \$6,750 and its cost at 100 million tokens a month from \$1.65 to \$2.12 per million; at \$600, to \$7,300 and \$2.28 [M]. Neither moves the machine out of the order-of-magnitude gap against the closed API price list, and both move it above the hyperscaler’s rack at planned utilization while leaving it below the rack at disappointed utilization. The stack price decides how much of the argument survives, not whether it does.

The market’s own view of the cost structure is that it is unproven. Counterpoint’s read of the August 2026 investor day was that SanDisk’s own FY2028–30 financial model contains no HBF revenue—HBF is “upside”; Wedbush declined to build it into estimates “until adoption becomes clearer”; Mizuho’s view is that it is “unlikely to fully replace HBM” [639, 638] [A]. Market-size figures in circulation—\$1B in 2027 rising to \$12B in 2030 (unattributed, Korean press); “bigger than HBM by 2038” (a KAIST professor)—are not analyst forecasts and are cited here as what is being said, not as estimates [646, 647] [A]. SanDisk’s shares, up more than twenty-fold in a year on the NAND supercycle, rose 13.7% on the investor day; the sell-side attributed most of that to the financial model and long-term agreements and only part to HBF [638, 639] [A].

B.8 The perceived advantages, and the case against each

The proponents’ case is compact and this appendix states it at full strength before answering it. Capacity: 512 GB per stack against 36–64 GB for HBM4, so a whole 400-billion-to-3-trillion-parameter model sits on one package—SanDisk’s “4 TB of VRAM,” The Register’s 3.12 TB from two HBM and six HBF stacks, the H³ paper’s 32-GPU workload on two [618, 649, 627]. Bandwidth: 1.6 TB/s per stack matches HBM3E and the 3 TB/s top grade exceeds a single HBM4 stack today, and decode is bandwidth-bound and read-mostly, which is the workload [570, 635]. Cost: an order of magnitude below HBM per gigabyte [572]. Non-volatility: no refresh, instant residency, model swap without reload from storage. Fewer accelerators for the same tokens: SanDisk’s “four HBF GPUs deliver the token output of eight HBM GPUs,” Goldman’s “half the GPUs,” FlashAccel’s 2.54× throughput per GPU, H³’s 2.69× per watt [637, 638, 633, 627]. Two independent authorities give the concept cover: Patterson and Ma’s research-directions paper proposes HBF as “10× weight memory” for mixtures of experts, and Carmack’s determinism argument answers the latency objection [626, 572]. And the standardization moved fast: memorandum to consortium to specification in twelve months, with Google inside [569].

The case against, item by item. *The bandwidth comparison is against the wrong HBM.* SanDisk’s investor-day slide set eight HBF stacks at 12.8 TB/s against an HBM3E-era part; HBM4E stacks are specified near 3.6–4.1 TB/s each, so an eight-stack HBM4E part in 2028 delivers roughly 30 TB/s—HBF’s read bandwidth is HBM’s of two generations ago, and the gap widens on the roadmap [640, 650] [A]. *Quantization shrinks the problem HBF solves.* The same slide sized Qwen3-480B at 960 GB in BF16; at FP8 it is 480 GB and at FP4—the format Kimi K3 ships in natively—240 GB, inside a large HBM4E complement; the appliance case of Section 12.9.4 is immune to this objection only because it targets a 2.8-trillion-parameter

model at four bits, which is 1.4 TB, and because it has no HBM at all [640] [A]. *Latency and granularity are hidden, not solved*, and every study that hides them does so with SRAM and prefetch that a real base die must pay for [627, 633]. *Key-value cache stays in DRAM*, so an HBF-only accelerator does not exist; the tier is additive to HBM or LPDDR, not a substitute, and Micron’s roofline finds that once one is adding a tier anyway, on-package LPDDR—writable, 2.4 pJ/bit—beats HBF at every operating point for datacentre decode by up to $2.1\times$ [632] [V; a competitor’s analysis]. *Packaging is the cost and the beachfront*. If the stack costs what an HBM stack costs, the cheap bits buy capacity, not a cheap part; and an HBF stack on a datacentre package occupies an interposer slot an HBM stack would otherwise take [572]. *Two vendors, and the wrong absentees*. NVIDIA has said nothing; AMD, Micron and Samsung are outside; Kioxia is building something else with NVIDIA; the trade press’s summary in January 2026 was “get Nvidia on board . . . and it’s a done deal,” and the corollary holds [648, 635] [A]. *Optane*. A proprietary memory tier at two vendors, needing a new controller, a new interface and a new software layer, has failed before on adoption rather than physics [615] [A]. *Timelines*. Samples have moved from 2H26 to 2027 in twelve months, production to 2028, and SK hynix has given no production date at all [651, 638] [A]. *Evidence*. Every performance figure in circulation is a simulation; the die taped out in August 2026 has not been characterized by anyone outside SanDisk [652] [A].

The book’s reading of that exchange is the one Section 12.9 gives: the case against is strong precisely where the proponents overclaim—HBF as HBM’s rival in the datacentre package—and weak where the book actually uses it—HBF as the weights tier of a machine that has no HBM, serves a handful of users, holds a four-bit model larger than any HBM complement, and is idle two-thirds of the day. Micron’s roofline, the sharpest of the criticisms, is a datacentre-decode result at high batch with a writable tier competing for the same slot; the appliance’s KV tier is LPDDR6X, exactly as Micron would prefer, and its weights tier is chosen for capacity per dollar rather than bandwidth per watt. Whether the appliance’s flash reads cost 6 or 8.5 or 12 pJ/bit changes its power draw by a hundred watts and its running cost by a few dollars a month. What would change the argument is HBF failing to ship, or shipping at a stack price closer to \$1,000 than \$200—and both are live possibilities that the forecast register carries.

B.9 The competitive landscape

HBF’s rivals for the “tier between HBM and storage” are of three kinds, and the book’s position on each is stated so it can be checked. *More HBM*: HBM4 at 36–48 GB per stack with a JEDEC ceiling of 64 GB, HBM4E sampled in May 2026 for late-2027 volume at 3.6–4.1 TB/s—which addresses bandwidth handsomely and capacity hardly at all, and at the highest cost per bit in the industry [653, 636] [A]. *Capacity DRAM*: LPDDR5X SOCAMM2 modules at 256 GB (2 TB per server) already shipping, LPDDR6 on JEDEC’s data-centre roadmap, Qualcomm’s AI200 with 768 GB of LPDDR per card, CXL memory expansion, and Micron’s argument that this tier dominates HBF for datacentre decode; the appliance uses it for context and would use it for weights too if flash did not arrive [632, 585, 577] [V/A]. *Other flash*: Kioxia’s path—the 5 TB, 64 GB/s daisy-chained module of August 2025, a 100-million-IOPS GPU-direct SSD for 2027 built with NVIDIA, XL-Flash CXL modules, and a concept of up to 32 TSV-stacked dies on a shared interposer beside HBM [622, 629] [V/A]; Samsung’s zNAND-O, a four-to-eight-die NAND stack on a TSV base die with UCIe to an on-device NPU, targeted at 2028—HBF’s edge concept, executed outside the OCP tent [636] [V]; and the research thread that puts compute *inside* the flash rather than beside it (AiF at ISCA 2025, KVNAND, NVLLM, and a Korea University 3D-NAND processing-in-memory array), which is where a Chinese ecosystem with a strong NAND base and weak HBM could plausibly go [654] [V]. On the DRAM side, ZAM (Section 12.9) is the capacity-by-height alternative that would compete with HBF for the same weights-tier role at DRAM latency, later and at a higher price per bit.

The China angle is worth stating separately because it is not what the Western coverage assumes. Chinese analysts have framed HBF as a “lane-change” opportunity: China “does not have high-end production capability in HBM,” but YMTC’s Xtacking is a wafer-bonded architecture of exactly the CBA kind HBF needs, YMTC’s 294-layer part is reported in production at yields above 90% with a 13% unit share and an expansion from 200 to 300 thousand wafers a month by 2027, and its hybrid-bonding intellectual property has been licensed to Samsung [655] [A]. YMTC’s own developer conference in December 2025 said HBF “fits Xtacking” and DigiTimes reported it going “full speed toward HBF next”; a competing report has YMTC prioritizing HBM with CXMT instead [625, 656] [A]. The constraint is tooling—the entity list bars the equipment for beyond roughly 300 layers—and the decider, in the Chinese framing as in the Western one, is whether NVIDIA adopts the tier. The book’s Chapter 13 argument that a NAND-rich, HBM-poor ecosystem is well placed for a flash weights tier is consistent with that framing but is not yet supported by a Chinese HBF product announcement, and it is graded accordingly.

B.10 The research literature

Table B.3 lists the papers this appendix relies on and the ones a reader should know exist. Three strands run through it. The *offloading* strand, from FlexGen and Apple’s “LLM in a flash” in 2023 through the mixture-of-experts expert-caching work of 2024–26 (MoE-Infinity, MoE-Lightning, ExpertFlow, PowerInfer-2 on a phone), established that weights can be streamed from flash behind a small hot cache and that expert routing has enough locality to make caching worthwhile, but that the SSD path is too slow and, per the SNU paper, too energy-hungry to be more than a fallback [660, 578]. The *HBF-specific* strand, all from 2026, models the on-package tier and agrees on its shape: SK hynix’s H³, the POSTECH–ETH characterization (which concludes HBF helps “only if HBM-equivalent read performance and substantial endurance improvement are achieved”), FlashAccel, Ma and Patterson, and Micron’s dissent [627, 657, 633, 626, 632]. The *routing-locality* strand is what decides whether a flash weights tier serves cost or bandwidth: per-token sparsity is high—DeepSeek-R1 activates 3.5% of parameters—but batching erases it, so that a batch of a few dozen touches nearly every expert; local routing consistency varies widely across models and is degraded by shared-expert designs; a hot cache of about twice the active experts captures most of the locality that exists [661, 578]. That last strand is the literature’s version of the batching-saturation arithmetic of Section 12.9, and it cuts the same way: an ultra-sparse mixture of experts is exactly the model for which a capacity tier is worth having and a bandwidth-amortizing rack is worth less.

B.11 What the main text assumes, and what would falsify it

Section 12.9.4 assumes four first-generation stacks at 512 GB and 1.6 TB/s each, 7 pJ/bit read energy, zero standby power, a stack price of \$150–250 on a pre-surge basis, and endurance adequate for a device that reloads its model at most daily. Against the record in this appendix: the capacity and bandwidth are the vendors’ published Gen. 1 figures and the specification’s middle grade [V, as targets]; the energy is the third-party modelled consensus and about twice HBM’s [M]; the standby figure is the fact sheet’s non-volatility claim [V]; the price is below the vendors’ one adjective and is the assumption the appliance’s cost is most sensitive to [M]; the endurance is unpublished but benign for the workload [M]. The forecast register (Appendix J) carries the resolving events: HBF-class silicon in volume at ≥ 512 GB and ≥ 1.6 TB/s with a published endurance figure; the appliance configuration below \$8k. To those this appendix adds the four observations that would move the argument fastest, in either direction: an independently measured energy per bit for a TSV-stacked NAND stack read over UCIE; a published endurance and retention specification at package temperature; a per-stack price, from anyone, against

Table B.3: Selected literature on flash weights tiers and HBF. All arXiv identifiers verified; venue given where published.

Paper	Venue / date	Result relevant to this book
FlexGen (Sheng et al.)	ICML 2023	OPT-175B on one 16 GB GPU with disk offload at ~ 1 token/s: the floor
LLM in a flash (Apple)	ACL 2024	windowing + row-column bundling: models $2\times$ DRAM, $20\text{--}25\times$ over naive flash loading
Fast MoE inference with offloading (Eliseev, Mazur)	Dec. 2023	Mixtral: LRU expert cache hit $0.4\text{--}0.8$; $2\text{--}3$ tok/s on consumer GPUs
MoE-Infinity; MoE-Lightning; ExpertFlow; HOBbit; PreScope; FlashMoE	2024–26	expert caching, prefetch and prediction: $3\text{--}16\times$ latency and up to $10\times$ throughput over naive offload; SSD-resident experts on the edge
PowerInfer-2 (SJTU)	Jun. 2024	47B MoE streamed from smartphone UFS flash at 11.7 tok/s
SSD offloading considered harmful (SNU)	CAL 2025	flash path ~ 102 pJ/bit vs HBM 4.2; per-token energy $3.8\text{--}12.5\times$; needs ~ 10 pJ/bit
Ma & Patterson	arXiv 2601.05047, Jan. 2026	HBF as $10\times$ weight memory; stack ballpark 512 GB, 1.6 TB/s, <80 W, 4 KB granule
H ³ (SK hynix)	IEEE CAL, Jan. 2026	8 HBM + 8 HBF: $2.69\times$ perf/W, $6.14\times$ throughput at 10M context; 40 MB latency buffer; $\sim 100k$ cycles
Exploring HBF for LLM inference (POSTECH, ETH)	IEEE CAL, Jan. 2026	gains “only if HBM-equivalent read performance and substantial endurance improvement”
HAVEN (UCSD, Georgia Tech)	arXiv 2603.01175, Mar. 2026	HBF for full-precision vector search in RAG: up to $20\times$ throughput
Micron, Is HBF all you need?	HotInfra at ISCA, Jun. 2026	on-package LPDDR beats HBF at every point for data-centre decode; HBF 8.5 pJ/bit, 36% scattered-read efficiency
FlashAccel (ICT-CAS)	arXiv 2607.10186, Jul. 2026	6 HBF stacks + HBM: $2.54\times$ throughput/GPU, $1.93\times$ tokens/J; 8 pJ/bit, $4\text{ }\mu\text{s}$ tR, 19 MB hyper-pages
HBF for KV-centric serving (Peking Univ.)	arXiv, Aug. 2026 (reported)	KV cache on HBF: latency $2\text{--}5.5\times$, throughput $1.1\text{--}2.7\times$ worse
AiF (SNU, ISCA 2025); KV-NAND; NVLLM; 3D-NAND PIM (Korea Univ.)	2025–26	compute inside the flash: $14.6\times$ over SSD offload; DRAM-free edge LLM at ~ 65 tok/s on 30B
Local routing consistency (Fudan, Huawei); Opportunistic expert activation	2025–26	routing locality varies across 20 MoE models; cache $\approx 2\times$ active experts near-optimal; batch-aware rerouting cuts MoE decode latency $15\text{--}39\%$
HeatWatch (CMU, ETH); STRAW (SNU)	HPCA 2018; CAL 2024	3D NAND retention superlinear in temperature; read-disturb reclaim writes $\sim 84\%$ with word-line-aware policy

HBM4E's; and a design-in by NVIDIA, AMD or a hyperscale ASIC—or, equally informative, Kioxia's and Samsung's separate paths reaching product first. As of 9 September 2026 none of the four has happened, one die has taped out, and the argument of Chapter 12 stands on a specification, a fact sheet, three simulations and a duty-cycle calculation. That is more than it stood on a year ago and less than it will need.

Appendix C

The Hot Chips 2026 Ledger

Hot Chips is the meeting at which processor architects present the parts they have built to an audience of other processor architects; its 2026 edition ran at Stanford from 23 to 25 August, with 31 technical presentations across the memory, server-CPU, GPU and accelerator, networking and tutorial tracks [672]. This appendix is the book’s compact evidence base for that record, current as of 9 September 2026, so that Chapters 12 and 13 can point at a table instead of carrying the detail. Deck numbers follow the conference’s own numbering. Figures appearing only in a deck are vendor-presented and unaudited, and are marked as such; where a vendor withheld a figure that exists elsewhere, the outside source is named.

An architecture conference is evidence for a book about industrial strategy for three reasons, none of which is that architects are good forecasters. The first is timing: the questions this book asks—what is scarce, who owns the interface, where the rent sits—are settled in silicon two to three years before they appear in a trade statistic, and the conference is where the design decisions that will produce the 2029 market are first stated in public. The second is that decks are the only venue in which competitors describe their assumptions about each other’s workloads on the record; a vendor’s choice of what to benchmark is a disclosure about what it believes the market will buy. The third is specific to this book’s argument and is the single most striking datum of the week: nine of the conference’s decks report inference economics on *Chinese open-weight models*, among them d-Matrix on GLM 5.2 and Kimi K3, OXMIQ’s cost model on Kimi-K2 1T, NVIDIA’s Rubin deck on DeepSeek-v4-PRO and its Spectrum-X deck on DeepSeek V3 training-step time, Intel’s capacity thesis plotted across every published frontier model from Llama 2 70B to Kimi K2 1T, SambaNova on MiniMax M2.7, OpenAI on DeepSeek R1—while no Chinese chip is named anywhere in the conference and no Chinese vendor presented [723]. US silicon vendors now validate and publish their inference economics against Chinese weights, because those are the frontier-scale weights that can be obtained and published against. Chapter 10’s claim about where the commons sits does not need a stronger corroboration than a competitor’s own measurement harness [V].

C.1 The methodological warning: disclosure has collapsed unevenly

One caution governs every table below, and it is not a complaint about vendors—it is a finding. *Disclosure has collapsed unevenly, and the pattern is not accidental: the parts with the strongest competitive position disclose least.* Table C.2 carries twelve parts, nine of them custom accelerators—every entry but the three merchant GPUs—and the counts below are taken against that nine. Eight publish no clock. Six publish no process node in the deck. Six publish no third-party benchmark of any kind. On power the count has to be given twice, because the deck and the wider record differ: six of the nine print no power figure at all, and of the three that do, only OpenAI’s Jalapeño disclosed it in the deck—Maia 200’s 750 W comes from a Microsoft blog

and MTIA 400’s 1.2 kW from outside coverage—so eight of the nine published no power figure *in the deck*. Google’s TPU 8t and 8i decks, which describe two concurrently shipped parts in unusual architectural detail, publish no per-chip FLOPS, power, node or die size at all [685].

The right inference is narrower than “the industry stopped disclosing.” Much of what the decks withheld exists elsewhere—in datasheets, vendor blogs, or the notes of favoured analysts—and Maia 200’s full specification, for instance, was published in a Microsoft blog post seven months before the conference [687]. What changed is the venue: *the architecture conference is no longer where disclosure happens; it now happens at the product launch*, where the framing is controlled and the comparison set is chosen. The consequence for analysis is concrete. A peak-FLOPS table across this generation is partly unbuildable, and any that is built is comparing formats—NVFP4 against MXFP4 against MX8S—that do not normalize cleanly. The only axes available for every part are memory capacity, memory bandwidth, fabric bandwidth and scale-up domain size. That is a smaller set than one would want and, conveniently, it is close to the set that actually determines serving cost, which is why the tables below are organized around it rather than around arithmetic throughput.

C.2 The data-centre CPUs

Five server CPUs were presented, from vendors with almost nothing else in common, and they converged on the same four choices (Table C.1).

Table C.1: Data-centre CPUs at Hot Chips 2026, as presented [P; vendor decks, unaudited]. “n/s” = not stated in the deck.

Part	Cores	Memory	On-die AI	The notable choice
Intel Diamond Rapids	256 uniform cores, 22 dies	P-MRDIMM-12800, 16 ch, 1.6 TB/s; 1.28 GB LLC	AMX + AVX10.2; no TOPS quoted	The P/E split abandoned; a very large last-level cache in place of HBM
NVIDIA Vera	88 Olympus v9.2	Arm LPDDR5X-9600 CAMM, 1.2 TB/s	SO- none	1.8 TB/s NVLink-C2C: the AI is the GPU beside it, not the CPU
Fujitsu MON- AKA	144 2 nm	Armv9.3-A, DDR5-8000+, 12 ch	~96.2 INT8 TOPS (est.)	FP64 retained: the mixed simulation/AI socket
Arm AGI	136 Neoverse V3, 2 chiplets	DDR5-8800, 12 ch	no matrix engine	Arm’s own reference platform ships without one
IBM Z	11 SMT-2 5.7+ GHz	cores, n/s	n/s	Presented as a hardware AArch64 v9.3 implementation inside the z core (2,792 instructions, 239 system registers), with z/OS and Linux-ARM64 as peer partitions [672]

Four findings, of which the third bears most directly on this book’s argument. *Uniform performance cores everywhere*: the performance/efficiency split that dominated CPU design for half a decade is gone from the server socket, most visibly at Intel, which shipped 256 identical P-cores. *Agentic orchestration is the uniform justification*: every vendor gave the same reason—many concurrent, latency-sensitive, branch-heavy control threads driving accelerators—which is a striking degree of consensus about a workload class that barely existed in 2023 and which nobody sized. *No HBM on any CPU*: after two generations in which HBM on the CPU looked like the direction of travel, all five parts attach commodity DRAM, whether MRDIMM, LPDDR or DDR5, which is the same capacity-over-bandwidth trade that Section 12.8 attributes to a sanctions-driven Chinese design and which turns out to be the choice five unconstrained Western vendors also made. And *on-die AI acceleration is conceded*: Vera has none, Arm’s AGI has no matrix engine, Intel quotes no TOPS for AMX, and only MONAKA carries a sized figure—the

CPU vendors have collectively agreed that the matrix work happens on the part next door.

The RISC-V thread runs alongside and is treated fully in Chapter 12; three of its 2026 data points belong in this ledger. NVIDIA stated the requirements for RISC-V to become CUDA’s third host ISA—RVA23 plus the Server SoC and Server Platform specifications, Boot and Runtime Services, and ACPI, with the CUDA-specific asks reducing to coherent PCIe I/O and full-speed peer-to-peer—which is the incumbent publishing a conformance target for an ISA it does not own [673]. Against that, *no ratified matrix extension exists*: one fast-track dot-product proposal and three competing task groups, with ratification targets spanning August 2026 to April 2027 [674]—with the caveat that IME’s 27 August target fell a fortnight before this edition’s currency date and no ratification announcement had appeared by then [A]. And the software readiness gap is measured rather than asserted—Canonical’s own figure is that roughly 95% of archive packages build on riscv64, while only 55.1% of tested packages pass, against 92.3% on amd64, on about one seventh the test coverage [675]—while the first rackable enterprise RISC-V server shipped in the same week, in limited quantities, with no price, no SPEC result and no named customer tape-out [676]. An open ISA with a ratified server platform, no matrix extension and a 55% package pass rate is exactly as far along as those three numbers say, and no further.

C.3 The accelerators

Eleven accelerator programs, twelve parts, were presented or are current enough to belong in the comparison (Table C.2). The columns are the four axes that survived the disclosure collapse.

Two observations before the architectures. First, the merchant GPU is not being attacked on arithmetic: normalizing the published figures to dense 8-bit throughput per package as best the formats allow, *every custom accelerator sits between roughly 3× and 30× below a contemporary merchant GPU* [M; the author’s normalization, since the vendors’ formats do not compare directly]. They are competing on cost per served token at a fixed interactivity, which is what OpenAI’s deck says outright when it declares per-chip throughput and chip count explicit non-goals and organizes the whole design around output tokens per second per all-in utility megawatt [683]. Second, the money is now large enough that this is no longer a side market: Broadcom’s AI semiconductor revenue reached \$10.8B in its second fiscal quarter of 2026, up 143% year on year, with third-quarter guidance above \$16.0B, which annualizes to roughly 11–15% of the combined merchant-plus-custom total depending on whether actuals or guidance are annualized on each side [M] [691].

C.4 Four schools of thought for beating the GPU

The eleven programs reduce to four theories of where the GPU is beatable, and each is falsifiable in a different measurement.

Eliminate data movement. Put the working set on the die and pay for it in capacity: Cerebras with 44 GB of on-wafer SRAM at 43.2 PB/s and no DRAM at all; the Groq-derived LP30 with 0.5 GB per chip, no DRAM, and a rack that aggregates SRAM instead of adding memory tiers [703]. The bet is that bandwidth is worth more than capacity and that model sharding can absorb the difference.

Determine the schedule statically. Remove the dynamic scheduler, the memory barriers and the kernel boundaries, and let a compiler place every operation on a cycle-exact global clock—the LP30’s VLIW model, and SambaNova’s dataflow lineage. The bet is that the GPU’s utilization losses are scheduling losses, and that a compiler can recover them.

Make the fabric fast enough that nothing stalls. NVIDIA’s NVLink 6 with in-network compute, AMD’s 72-GPU UALink-over-Ethernet load/store domain, Microsoft’s single Ethernet protocol

Table C.2: Accelerators, on the axes every vendor published [P; vendor-presented and unaudited unless otherwise sourced]. Precision labels are the vendors’ own and do not normalize across parts. One discrepancy is disclosed rather than reconciled: MTIA 400’s MX4 figure is printed here at 12 PF, which is the number in Meta’s March 2026 announcement table [688], while a May 2026 survey of the same part reports 18 PF; no vendor has explained the difference [A].

Part	Memory capacity	Memory BW	Scale-up domain	Headline throughput
NVIDIA Rubin	288 GB HBM4	22.2 TB/s	NVL72, 72 GPUs, NVLink 6	50 PF NVFP4 dense; 2.3 kW/GPU, 800 VDC racks
AMD MI455X	432 GB HBM4	23.3 TB/s	UALoE, 72 GPUs, merchant Ethernet	40.26 PF MXFP4; ~2 kW/GPU
Intel Crescent Island	up to 480 GB LPDDR5x	not stated	none (no coherent multi-GPU domain)	no FLOPS quoted at any precision; 350 W PCIe card
Google TPU 8t	6 HBM3E stacks	n/s	9,600 chips, optical 3D torus	n/s [685]
Google TPU 8i	8 HBM3E stacks + 384 MB SRAM	n/s	1,152 chips, 7-hop Boardfly	n/s; claimed > 2× perf/TCO gen-on-gen [685]
Microsoft Maia 200	216 GB HBM3E + 272 MB SRAM	7 TB/s	6,144 accelerators, one Ethernet protocol	>10 PF FP4; 750 W [687]
Meta MTIA 400	288 GB HBM3E	9.2 TB/s	RoCE on dedicated I/O chiplets	6 PF FP8/MX8, 12 PF MX4; 1.2 kW [688]
OpenAI Jalapeño	216 GiB HBM4	15.4 TB/s	128 chips at 600 GB/s; 2,048-chip pod	13.4 PF mxfp4; 700 W [683]
SambaNova SN50	432 MB SRAM + 64 GB HBM2E	1.84 TB/s	64 sockets, 10×800G Ethernet	3.2 PF FP8
NVIDIA/Groq LP30	0.5 GB SRAM per chip; no DRAM	~156 TB/s	256-chip rack, every chip a router	315 PF FP8 per rack [703]
Cerebras CS-4	44 GB on-wafer SRAM; no DRAM	43.2 PB/s	wafer-scale (on-die)	750 PF sparse / 125 PF dense
d-Matrix Rap-tor	3D-DRAM stacked under logic	12.5 TB/s per chiplet	n/s	0.37 pJ/bit measured

from chip to a 6,144-accelerator Clos, Google’s optically switched torus [681]. The bet is that the unit of design is the domain rather than the chip.

Change the memory economics. Intel’s Crescent Island substituting 480 GB of LPDDR5x for HBM; d-Matrix stacking 3D-DRAM under the logic at 12.5 TB/s per chiplet and 0.37 pJ/bit measured. The bet is that the memory system, not the arithmetic, is the cost.

The sharpest intellectual disagreement of the conference is between the last school and the first two, and it is worth stating precisely because both sides have a measurement. Intel’s case is that *capacity* is what matters: plotted across published frontier models from Llama 2 70B to Kimi K2 1T, the weight footprint that must be *held* grew $7.5\times$ while the bytes *read per generated token* fell $4.4\times$, because sparse mixture-of-experts routing activates a small fraction of parameters per token [723]. If that divergence continues, the machine to build is a large, cheap, slow memory pool. SambaNova’s and Cerebras’s case is that *realizable bandwidth per watt* is what matters: decode occupies 75–97% of an agentic session’s wall-clock time and decode is bandwidth-bound, so the machine to build is a small, fast, expensive one. Both are correct inside their own regime, and the regime boundary is the batch size—at batch one the sparsity is real and capacity wins; at a batch of a few dozen the union of activated experts approaches the whole model and bandwidth wins. *Nobody publishes the batch size*, which means the disagreement cannot be resolved from public evidence, and any vendor comparison that does not state it is comparing operating points rather than machines. This is the same batching-saturation boundary that Section 12.9 and Appendix B identify from the model side, arrived at independently from the silicon side.

C.5 FP64 bifurcated rather than disappeared

The double-precision story matters directly to any centre running simulation and AI on one estate, which is most national laboratories, and 2026 is the year it changed shape. AMD split double precision into a separate product: the MI455X publishes 5 TFLOPS of FP64 against 78.6 on the MI355X—a compatibility rate rather than a capability—with a dedicated MI430X at 288 TFLOPS arriving separately in early 2027. NVIDIA moved FP64 into software, claiming ~ 200 TFLOPS by emulation on the low-precision tensor units, which is the Ozaki-scheme direction: high-precision results assembled from many low-precision products. Intel kept a hardware rate and declines to size it.

The procurement consequence is a single sentence and it changes how a machine is specified: *double precision is now a product-line decision rather than a per-generation one*. A buyer can no longer assume that the AI part of a vendor’s roadmap carries a usable FP64 rate as a by-product, and must instead choose a SKU, a year, and—if the emulation path is taken—an accuracy standard, because emulated FP64 is not bit-identical to hardware FP64 and the difference is a validation question rather than a benchmark question. Fujitsu’s MONAKA, which retains FP64 in a general-purpose socket, is the third answer and the one that most resembles the machine a mixed estate actually wants.

This is the line of argument this author has developed under the heading of *Tensor-Memory Equilibrium*—that once matrix throughput is abundant relative to memory bandwidth, precision becomes an allocation choice made in software against a memory budget rather than a fixed property of the hardware, so the correct question is not how much FP64 a part has but what it costs to synthesize. If NVIDIA’s emulated rate holds up under independent accuracy testing, that argument becomes the mainstream way to specify a scientific machine; if it does not, FP64 returns to being a SKU one buys. The companion study of the CPU–GPU design continuum reaches the adjacent conclusion from the die-area side—that either corner can be provisioned to do the other’s job [184, 185].

C.6 What the conference says about openness

Two results, both about interfaces rather than about licences, and both cutting the way Chapter 12 predicts.

Ethernet took the scale-up domain everywhere except NVIDIA. Among the parts presented with an off-package scale-up fabric, the Ethernet-class designs are the clear majority: AMD constructs a 72-GPU load/store domain out of UALink-over-Ethernet on merchant 512-port 200G switch ASICs, deleting the scale-up NIC from the tray; Microsoft runs one Ethernet protocol from the chip to a 6,144-accelerator two-tier Clos, with a slide titled “Why Unified?” listing a standard switch ecosystem as the reason; SambaNova uses ten integrated 800G ports per socket into two merchant switches; Meta uses RoCE on dedicated I/O chiplets; and OpenAI’s 2,048-chip pod runs over a Broadcom Tomahawk 6 Clos. Google’s optically switched torus is proprietary but is not sold to anyone, which leaves NVIDIA’s NVLink as the only proprietary scale-up fabric a third party can buy [681, 683, 685]. The layer that was supposed to be the incumbent’s deepest moat has been commoditized by its own customers, using parts anybody can buy—which is the mechanism of Chapter 7’s move M6 executed by Western hyperscalers against a Western incumbent, and worth noting for a book that otherwise attributes that move to Chinese industrial policy.

The Ultra Ethernet Consortium standardized the mechanism set and lost the wire format. Every mechanism the UEC argued for shipped in production silicon in 2026—packet spraying, out-of-order placement with in-order delivery, selective retransmission, receiver-driven credit congestion control, packet trimming, programmable transport—but shipped as *RoCEv2 extensions*. Broadcom brands its transport eRoCE and describes it as “MRC++,” preserving the verbs ecosystem so that xCCL, NCCL, RCCL and MPI run unmodified: “the ecosystem cost is zero.” Neither Broadcom’s nor NVIDIA’s deck presents a UET transport or a libfabric stack, and NVIDIA’s networking deck does not mention the consortium at all [682]. The lesson generalizes past networking and is one this book’s governance chapters should hold onto: a standards body can win the argument about *what a protocol should do* and still lose the protocol, because the installed software ecosystem is worth more than protocol purity. Openness that arrives as an extension to the incumbent format is openness that leaves the incumbent’s ecosystem in place.

C.7 What to watch

Eight resolving events, in the register idiom of Appendix J: each is falsifiable, each has a venue, and each is dated where a date exists.

1. **Per-stack bandwidth appears in bills of materials.** Watch for a procurement document, datasheet or hyperscaler disclosure specifying HBM4 or HBM4E in TB/s per stack rather than by generation name. If through 2027 the merchant part is still quoted only by generation, the $1.5\times$ spread of Section 13.3 has become a permanent feature of the market rather than a transitional artefact.
2. **A custom-HBM (cHBM) part ships to a buyer without its own silicon team.** cHBM is the customer-specific 4 nm logic base die, not the interposer-free stacking of the next item. If through 2028 every co-designed logic base die belongs to a company that designs accelerators, the merchant/co-design bifurcation is established and merchant buyers—national laboratories included—are structurally on the slower branch. Resolves on vendor announcements and teardowns.
3. **zHBM reaches silicon rather than a roadmap slide.** Samsung’s phase three claims a 70% reduction in I/O energy and 230% of HBM4E bandwidth by deleting the interposer [680]. Watch for a taped-out part. Absent by end-2029, treat the three-phase roadmap as a phase-one product with two phases of marketing.

4. **A published HBM-in-package yield figure.** No vendor printed one in 2026, and the assembly-order inversion means an HBM defect now scraps the whole system-in-package. The first vendor to publish a scrap rate at 16-high settles whether Eq. (A.1) is a cost problem or a capacity crisis; continued silence past 2028 is itself evidence about the number.
5. **The FP64 pair resolves by end-2027.** Whether AMD’s MI430X ships at its stated 288 TFLOPS in early 2027, and whether NVIDIA’s ~ 200 TFLOPS emulated rate is independently reproduced against a published accuracy standard. Together they decide whether double precision is a product line or a software library.
6. **A custom accelerator submits to a neutral benchmark.** MLPerf Inference runs twice yearly. If by end-2028 none of Maia, MTIA, Jalapeño or TPU 8i has appeared in a neutral round, the disclosure collapse documented above is structural rather than transitional, and third-party evaluation of the custom-ASIC fleet becomes impossible in principle rather than merely absent.
7. **A RISC-V matrix extension is ratified.** Three competing task groups (IME, VME, AME) carry targets spanning August 2026 to April 2027 [674]. Ratification of one, with a shipping core behind it, is the event that converts the open ISA from a host processor into an accelerator ISA; a fourth proposal appearing instead is the falsifier.
8. **A Chinese part is presented, or measured.** No Chinese chip was named at Hot Chips 2026 and no Chinese vendor presented, while nine decks benchmarked Chinese weights [723]. The first Huawei, Cambricon or Biren deck at a Western architecture conference—or a Chinese-hosted equivalent that Western vendors attend and publish at—is a datable event with strategic content in either direction, and its continued absence past 2028 is evidence that the two silicon ecosystems have stopped reading each other’s papers.

One further item belongs on the list without a clean resolving event, because it bears on the fastest-moving claim in this book. Only two of the conference’s decks claimed AI inside the silicon design flow, and both came from companies that design accelerators for their own frontier models: Google reported an “AI Designer” taking specification to RTL on $O(100)$ TPUs in about a week and an “AI Optimizer” yielding a 6% power and 5.8% area reduction on the TPU 8t matrix unit; OpenAI reported block-level wins against an optimized human baseline, and separately an internal agent driving a DeepSeek MLA kernel from 0.31% to 88.94% of roofline in about forty hours. Neither reports full-chip power–performance–area, bug escapes or schedule saving against a baseline, and *nobody claimed AI-driven verification*, which is where most of a chip program’s cost sits [722]. The bring-up-software result is the one to watch: the historical reason a novel accelerator fails is that nobody can write fast kernels for it in time, and an agent that removes that barrier removes the strongest structural argument for why a Chinese accelerator ecosystem cannot converge quickly. Two decks are not a trend, and the absence of a verification claim is the reason to hold the inference loosely [A].

Appendix D

Evidence Ledger

The book’s headline claims, graded per Section 1.5. This ledger is the living document’s audit surface: each revision re-grades it. It is set as a multi-page domain-split table with repeated headers, because a single float clips in compilation and a clipped audit surface is a publication-blocking defect.

Table D.1: The evidence ledger, split by domain.

Claim	Grade
1. Industrial precedents	
China holds >90% share of every solar PV supply-chain stage	V
Chinese makers hold ~69–75% of global EV battery installations; >80% of cell capacity	V
China NEV penetration 54% (2025); largest auto exporter	V
2. Models and adoption	
Chinese open models: 41% of HF downloads; majority of Open-Router tokens	V
Top US–China model gap 2.7%; closed-over-open gap 3.3%	V
Enterprise open-weight <i>spend</i> share fell 19%→11% in 2025	A
DeepSeek V3 final-run compute \$5.6M (2.79M H800-hours)	P
Kimi K3: largest open-weight release (2.8T/104B), under a <i>custom commercial</i> license, self-reported as trailing the closed frontier	V (license, specs)
DeepSeek V4-Flash (~280B/13B, MIT, weights shipped) released 2026-07-31	V
Qwen3.8-Max (2.4T MoE) announced 2026-08-03	V
Alibaba publicly committed Qwen3.8-Max <i>and</i> Qwen3.8-27B to open weights, one week after Kimi K3 opened at 2.8T	V (official announcement)
Every Max-tier Qwen before 3.8 was proprietary; the reversal followed a rival’s frontier-scale open release—the ratchet of §10.3.1	V (pattern) / S (causal reading)
Qwen3.8-Max: autonomous execution of the silicon design flow; 500+ turns of chip-design optimization	P (vendor claim; no independent verification)
Qwen family trained on leased offshore (Singapore/Malaysia) NVIDIA GPUs, served domestically—read as transitional given the co-design mechanism	A/P (press-reported; not vendor-confirmed)
Four of the five models on the cost-performance Pareto frontier are Chinese	A (Arena-derived, single analysis)
OpenAI cut GPT-5.6 Luna pricing 80% on 2026-07-31 under open-weight competitive pressure	V (price) / A (attribution)
Distillation campaigns (>16M exchanges) by Chinese labs	P (vendor allegation)
DeepSeek-family censorship, transferring to distills	V

continues

Evidence ledger, continued

Claim	Grade
Chinese dominance of the open commons by 2030	S (70–80%)
3. Compute and architecture	
LineShine #1 HPL (2.198 EF) and #1 HPCG (22.0 PF), GPU-free, 42.2 MW	V
LX2 serving competitiveness ($\sim 1.3\text{--}1.4\times$ B200 per-stream energy)	M
LX2 designer/fab/HBM supplier attributions (Huawei/SMIC/CXMT)	A
LineShine uses domestic HBM	P (state media; unverified)
IME ratification target Aug. 2026; AME board-ratification target Apr. 2027	P (official plans)
Dual-variant RISC-V accelerated-CPU template	M/S
LPDDR6X 2.3TB/7.37TB/s socket per Table A.1; 67% MFU ceiling; configuration-specific sub-256-node HBM wall	M
Capacity-first LPDDR as an <i>implemented</i> Chinese national program	Not established
4. Semiconductors and energy	
SMIC N+3 5 nm-class in volume without EUV	V (capability) / A (economics)
CXMT LPDDR6 verification cleared; 2H26 mass production	A/R (no primary product announcement)
CXMT HBM3 samples to Huawei; 2026 mass production	A/R
CXMT commodity-DRAM scale: three fabs $\sim 100\text{k}$ wpm, expansion planned beyond 600k wpm	P/A (source-based reporting)
Domestic HBM $\sim 2\text{M}$ stacks 2026; Huawei foreign-stack stockpile exhausted	A
China added 429/540 GW of generation (2024/2025); US record year 53 GW	V
Domestic-chip data centers get up-to-50% power discounts	P
5. Finance	
OpenAI \$852B mark, \$13.1B 2025 revenue, $\sim \$1.4\text{T}$ commitments	V/P
Microsoft AI run-rate $> \$37\text{B}$, $+123\%$ y/y	V
Anthropic run rate \$47B (May 2026) at \$965B post-money ($\approx 20.5\times$); gross-margin trajectory toward $\sim 44\%$	V (company) / A (margin)
OpenAI commitments reset toward $\sim \$600\text{B}$ -by-2030 (Feb. 2026)	P (press-reported)
NVIDIA Q1 FY27: \$81.6B revenue, \$75.2B data center, $\sim 75\%$ non-GAAP gross margin	V
AI-focused firms $\approx$ half of S&P 500 value	V
Inference prices $-9\times$ to $-900\times$ /yr; compute capacity doubling every 7 months	V
Broad US AI valuation repricing by 2029–30	S (60–70%)
6. Security	
Backdoors survive safety fine-tuning; ~ 250 documents suffice to poison	V (Level A)
Backdoored agents exfiltrate memory via disguised tool calls ($> 94\%$ activation, $< 1\%$ utility loss)	V (Level A)
Prompt worms persist and propagate multi-hop in an unmodified agent framework	V (Level A/B)
Malicious agent skills exist at ecosystem scale (157 of 98,380 audited)	V (Level C)

continues

Evidence ledger, continued

Claim	Grade
Autonomous evaluation agents caused real cross-organization intrusions (Jul. 2026)	V (Level C)
A released official checkpoint contains a destructive latent backdoor “Order 66” end-to-end systemic loss	Not established S (fault-tree composition of A-level parts)
Injection resistance varies $\sim 10\times$ across released models	V (configuration-specific)
7. AI for science and national sizing	
AI4S frontier-LLM workload share $\sim 14\%$; purchased-token premium $20\text{--}60\times$	A/M
Resulting frontier share of expenditure $76\text{--}91\%$	M
Task-weighted AI4S uplift: $3\text{--}30\%$ scenario band, central $\sim 10\%$	M/S (Eq. 24.2)
Fugaku COVID-19 projects returned $\$50\text{--}75\text{B}$; successor-return scenarios $\$250\text{--}600\text{B}+$	A/M (commissioned, survey-based counterfactual model)
Task-weighted funnel Monte Carlo: median 11.8% , P10–P90 $7.7\text{--}17.5\%$	M (reproducible; CSV published)
Consortium feasibility (federated training, financing, continuity)	S
EA1 probe: 443 BGPU-h per active heavy-user day, saturating an $\sim 800\text{-GPU}$ partition	V
BGPU conversion rule; MI355X $\gamma = 1.0$ on published MLPerf evidence	V/P
2030 flagship generation factor of 6	derived from disclosed bandwidth/fleet targets [P]
Master term = 88.7% of token-service capacity at 2^{20} context	M
Per-researcher core ≈ 0.94 BGPU, scale-invariant across three countries	M
Japan 2030 requirement $497\text{--}510\text{k}$ BGPU	M (measurement-anchored)
Singapore $55\text{--}59\text{k}$ and US $1.09\text{--}1.17\text{M}$ BGPU	T (pre-measurement scenarios)
Committed fleets provision $0.02\text{--}0.10$ BGPU per researcher	V/A

Appendix E

Deployment Control Architecture

The controls below follow directly from the fault tree of Section 19.6 and the trust stack of Chapter 18. They are reproduced here as an operational checklist for sensitive scientific, governmental, or industrial deployments, because the strategic argument of this book—adopt the commons, verify it, own the certification layer—is only actionable if the verification and containment work is specified rather than gestured at [291]. Each control family below maps to a gate of the fault tree: controls (1)–(5) cut the *trigger-success* and *privileged-deployment* gates (preventive); (6)–(8) cut *privileged deployment* and *propagation* (preventive); (9)–(10) are the *detective* layer across all gates; (11) cuts the *supply-chain compromise* disjunct; (12)–(14) cut *containment failure* (preventive, detective, and recovery respectively). The gates with the weakest control coverage—artifact compromise itself, and trigger discovery—are exactly the ones the tree identifies as outside operator control, which is why the checklist concentrates spending elsewhere.

Containment and authority. (1) Separate inference from action execution: run the model in a process or enclave without direct network credentials, and treat every tool call as an untrusted proposal. (2) Use a deterministic tool broker enforcing operation, destination, argument, volume, and rate allow-lists *outside* the model; permit no arbitrary URLs, model-constructed hostnames, or arbitrary message recipients. (3) Separate secret access from egress: no single invocation should hold both unrestricted memory or file reads and unrestricted external transmission. (4) Inspect payloads, not only retrieved text—flag opaque or high-entropy URL parameters, Base64-like fields, secrets copied from memory, and unusual destination changes; never auto-fetch model-generated remote resources. (5) Require independent policy approval, and where appropriate human confirmation, for credential access, external messaging, code execution, purchases, data export, and irreversible operations.

State integrity. (6) Make core policy and configuration immutable to the agent: no email, web content, skill, memory, or peer message may rewrite startup instructions; keep system policy in a separately authenticated control plane. (7) Treat memory writes as security decisions—record origin and derivation, prevent untrusted content from creating policy or authorization memories, expire low-trust entries, and reauthorize consequential actions at retrieval time. (8) Constrain destructive filesystem operations: read-only mounts by default, no recursive or bulk deletion, separate backup credentials, immutable versions, and out-of-band authorization for destructive changes to shared storage.

Observability and supply chain. (9) Preserve complete external logs—model versions, prompts, tool schemas, proposed actions, broker decisions, network destinations, results—in append-only storage unavailable to the model. (10) Freeze and sign tool definitions: hash-pinned schemas, model-visible instructions fully disclosed to reviewers, reapproval required after changes, and no silent transfer of data obtained through one tool into another. (11) Harden the model and agent supply chain: controlled mirrors, signed manifests, pinned cryptographic hashes, tensor-only formats, pinned runtimes, software bills of materials, disabled remote model code

and unsafe deserialization, CI credentials isolated from compromised workloads, and verified release provenance.

Evaluation and assurance. (12) Contain evaluations as hostile-code execution: physically deny and independently verify egress, use synthetic credentials and targets, block routes to metadata and package-control planes, monitor outside the agent’s control, and terminate on unexpected destinations. (13) Red-team the actual integration—temporal, multilingual, organizational, tool-schema, environmental, and multi-turn triggers; indirect prompt injection; self-replication attempts; destructive-action canaries in memory and files. (14) Do not rely on fine-tuning as sanitization: scanners, differential testing, and controlled post-training reduce risk but do not establish absence; high-assurance workloads should use network isolation and reviewed output release.

For the most sensitive workloads, a hash-pinned model in a network-isolated inference environment with internal-only retrieval and a one-way, reviewed export path provides substantially stronger assurance than an unrestricted “trusted” model of any provenance—which is the operational restatement of this book’s position that provenance claims are not a substitute for architecture.

Appendix F

Consortium Feasibility: The Legal and Engineering Homework

Chapter 27 argues the consortium is necessary; necessary is not the same as operable, and the feasibility homework has to be specified rather than gestured at. This appendix names the fourteen problems a founding instrument must solve, with this book’s current best answer or its honest absence.

Legal. (1) *Export-control segmentation*: members sit under incompatible regimes (EAR, EU dual-use, and their non-Western counterparts); the workable architecture is compartmented—nationally owned compute enclaves executing a federated protocol, so that no controlled article crosses a border even as gradients do; precedent exists in high-energy physics data policy, not in any AI instrument. (2) *Data localization*: solved by design if raw data never moves (the layered data policy of Chapter 27); unsolved for the preference- and correction-data pools that post-training needs. (3) *Classified and dual-use workloads*: excluded from the shared fabric entirely, run in national enclaves against consortium-published open weights—the model travels, the workload does not. (4) *Liability for defective models*: the hardest open problem in the list; a jointly trained model has no manufacturer in the product-liability sense, and the certification ladder of Chapter 18 (Level 5’s contractual, insurable assurance) is currently the only credible carrier; the founding instrument must create a liability-holding entity or the regulated sectors will not deploy. (5) *Ownership of jointly trained weights*: recommend open release under a permissive license with a contribution registry—ownership dissolved rather than allocated, which is both the mission and the only stable equilibrium among sovereign contributors. (6) *Withdrawal of contributed data*: prospective only (future checkpoints), with data bills-of-materials making the boundary auditable (Chapter 15); retroactive unlearning at frontier scale is not a commitment any honest instrument can make. (7) *Incident jurisdiction*: model incidents route to the operator’s jurisdiction, artifact incidents to the consortium’s seat; the append-only audit substrate of Appendix E is what makes the routing adjudicable. (8) *Antitrust and state aid*: a research-infrastructure joint undertaking (EuroHPC’s legal form) survives both regimes; a production token utility would not; the line must be drawn in the charter. (9) *Insolvency and continuity*: mirrors, escrowed artifacts, and multi-party signing keys so that no member’s exit—or the consortium’s own dissolution—can orphan the certified corpus.

Engineering. (10) *Federated scheduling* across heterogeneous, intermittently available national machines: the scheduling literature exists (grid computing paid this tuition twenty years ago); the missing piece is incentive-compatible accounting, addressed by (13). (11) *Heterogeneous optimizer consistency*: training one model across unlike accelerators requires either synchronous steps at the slowest site’s pace, hierarchical/local-update methods (DiLoCo-class) with their convergence penalties, or partitioning by training phase (pretraining centralized per-site, post-training federated)—the last is the current recommendation, because post-training tolerates asynchrony and is where members’ data sensitivity lives anyway. (12) *Checkpoint transfer*

and cross-site bandwidth: a frontier checkpoint at full optimizer state is tens of terabytes; at research-network rates this is hours per sync, which bounds the federation granularity and is why (11) recommends phase partitioning. (13) *Accounting for unlike hardware:* the BGPU vector of Section 26.2.1 is the designed answer—contributions metered in workload-weighted capability, not device counts. (14) *The conformance instrument:* the AI-HPC lifecycle benchmark of Section 12.7 and the certification ladder give the consortium its two deliverables that no national program can credibly produce alone—neutral measurement and neutral assurance. The honest summary: nine of the fourteen have engineering answers on the shelf; (1), (4), and (5) require legal invention; and none of the fourteen is a reason to wait, because every one of them is also unsolved for every national program that will otherwise duplicate this work seven times.

Appendix G

Living Data Appendix: Flagship Model Cards

Chapter 10’s release table mixes vendor-published and independent results, different reasoning budgets, tool harnesses, attempt budgets, and evaluation dates, and that mixture is a methodological problem this book should not leave standing. The remedy is a normalized card per flagship, reported to a fixed schema and maintained as the flagships change. Table G.1 is that schema in use; empty fields are honest gaps rather than omissions, and the schema itself is the contribution, because it is what any serious comparison requires and what almost no public comparison supplies.

The ten fields: (1) total and active parameters; (2) native precision and quantization-aware-training status; (3) context and maximum output length; (4) license; (5) disclosed training hardware, with confidence grade; (6) independent capability score under a common date and harness; (7) a software-engineering benchmark under a stated tool environment and attempt budget; (8) input, cached-input, and output prices; (9) tokens and wall-clock consumed per successfully completed task; (10) measured latency, throughput, and serving hardware. Fields (9) and (10) are the ones that decide procurement and the ones no vendor publishes; until they are routinely reported, “ten times cheaper per token” remains an unfinished sentence, because a model that spends ten times the reasoning tokens per solved task has delivered no saving at all.

Table G.1: Normalized flagship cards, mid-2026. Vendor-reported values [P]; independent index values [V]; “—” denotes not publicly disclosed. Prices in USD per million tokens (input/output).

Model	Params tot./act.	Precision	Context	License	Train HW	SWE- bench (protocol)	Price in/out	Tok/task; latency
Kimi K3	2.8T ~104B	/ MXFP4 native	1M	modified MIT	H200 + L20 + alt. GPGPU [A]	— (vendor tables only)	3.00 / 15.00	—
DeepSeek V4- Pro	1.6T / 49B	FP8 native	1M	MIT	undisclosed	80.6% (ven- dor)	0.435 / 0.87	—
GLM-5.2	~753B ~40B	/ FP8 vari- ant	1M var.	MIT	Ascend (press, no primary) [A/C]	77.8% (GLM-5, vendor)	1.00 / 3.20	—
MiniMax M3	~428B ~23B	/ —	1M	open weights	—	80.2% (M2.5, ven- dor)	0.15 / 1.20 (M2.5)	—
Kimi Think.	K2 1T / 32B	INT4 QAT	256K	modified MIT	—	65.8% single- attempt; 71.6% par- allel TTC	—	—
Qwen3.6 fam- ily	27B dense; 35B/A3B	—	262K	Apache 2.0	—	—	—	—

Two disciplines travel with the table. *Protocol disclosure*: the Kimi row shows why—the same model, the same benchmark, and a 5.8-point spread between single-attempt and parallel

test-time-compute settings, which is larger than most inter-model gaps reported in the press [87]. *Task-normalized cost*: with reasoning-model output lengths inflating roughly fivefold per year [254], price-per-token comparisons decay in meaning within a single release cycle, and only cost-per-completed-task is stable enough to plan a budget against—which is precisely the quantity Chapter 27’s expenditure model needs and Section 27.1 had to estimate.

Appendix H

Current Cycle: From Models to Harnesses

This is the living appendix doing its job. The stable chapters are current as of 9 September 2026—this edition’s date—and this appendix records the most recent development cycle in concentrated form, so that the fastest-moving material sits in one place that the next edition can replace wholesale rather than being threaded through twenty chapters. The first entry covers 16 July–3 August 2026, which compressed more frontier activity than any comparable window in the field’s history.

H.1 The chronology

Four releases in eighteen days, each responding to the last [402]: **16 July**, Moonshot releases Kimi K3 (2.8T/104B) via API; **27 July**, K3’s weights are open-sourced under its custom commercial licence—the largest open-weight model ever released; **30 July**, OpenAI cuts GPT-5.6 Luna—the cheapest tier—by 80% (\$1.00→\$0.20 input, \$6.00→\$1.20 output) with Terra down 20% and the flagship Sol held unchanged [403, 717]; **31 July**, DeepSeek ships the official V4-Flash checkpoint (~280B/13B, MIT, weights on Hugging Face, \$0.14/\$0.28 per million tokens); **3 August**, Alibaba makes Qwen3.8-Max generally available (2.4T total, ~95B active by press consensus though never officially confirmed, 1M context, multimodal in), API-only, at \$2/\$6 per million tokens [396].

The sequencing carries the analytical content, and it is not the sequencing the headlines implied. OpenAI’s price cut came *before* the Qwen release, not after it, and contemporaneous analysis attributed the pressure to Kimi K3 and DeepSeek V4 rather than to Alibaba [403]. That is the first documented instance in this book’s evidence base of a Western frontier laboratory repricing its mainline product in visible response to Chinese open-weight competition—the transmission mechanism of Proposition P3, observed rather than modelled. The company attributed the cut to “system efficiency,” which may be entirely true and is also exactly what a firm says when it is defending share.

H.2 Qwen3.8-Max: what is claimed, what is verified

On benchmarks the discipline of Section 10.5 applies: *the headline numbers are Alibaba’s own*, published in the release blog. The table reports SWE-bench Pro 67.7 (Claude Fable 5 leading at 80.0), Terminal-Bench 2.1 86.6 (Fable 5 84.6, GPT-5.6 Sol 88.8), PaperBench 93.0 (best in table), FrontierSWE 73.5, QwenReactBench 1,724 (best), CoWorkBench 74.8, JobBench 53.4, Agents’ Last Exam 52.4, GPQA Diamond 92.6, IFBench 82.8, HLE 43.6, plus a multimodal sweep (BabyVision 91.3, CharXiv 93.5, ERQA 77.8, PerceptionBench 63.5) [398, 393]. The shape is consistent and worth naming: Fable 5 leads on the hardest software-engineering and

Table H.1: Qwen3.8-Max as of this edition’s date. Vendor claims are extensive and partly auditable; independent verification is thin. Both halves matter.

Item	Reported	Status
Total parameters	2.4T sparse MoE	consistent across all sources [P]
Active parameters	~95B	press consensus; not stated in the announcement [A]
Expert count, routing	—	not disclosed
Context	1M (991K in / 131K out / 262K reasoning budget)	vendor [P]
Attention	possibly gated-linear (Gated DeltaNet-class)	speculation; unpublished
Training tokens, precision	—	not disclosed
Training hardware	—	not disclosed; family reportedly trained on leased offshore NVIDIA [404]
Weights	publicly committed by Alibaba: “next week, the open weights of Qwen3.8-Max will be released, and Qwen3.8-27B is also going open-weights”	[V] as a commitment ; the committed week lapsed without a release; window re-framed to the week of 10 August; no repository at this edition’s date [393, 476]
Licence	—	<i>unstated</i> ; Reuters-sourced reporting now points to revenue-sharing terms for large resellers [749]
Documentation	release blog with full benchmark table; GitHub project trace for the autonomous-coding run	[V]; no formal model card or architecture report
Price	\$2.00 / \$6.00 per Mtok, \$0.25 implicit cache, flat across the full 1M window	vendor [V] [393]

reasoning rows; Qwen leads on multimodal reasoning, document and office intelligence, real-world and spatial understanding, visual perception, and research reproduction. Absences remain notable—no SWE-bench Verified, no AIME, no Kimi K3 or DeepSeek V4 column despite those being the nearest competitors, and no hallucination-rate disclosure of the kind Qwen3.7-Max carried.

The autonomy claims deserve separate treatment because their auditability differs. The autonomous-coding demonstration—“10+ days of self-evolving development, from empty folder to production without hand-holding”—ships with *a complete project trace on GitHub* [393], which is more verification apparatus than most Western frontier autonomy claims carry and should be credited as such; the standing question of how many human interventions occurred is answerable from the trace rather than merely rhetorical [400]. The chip-design claim (500-plus turns of optimization; autonomous execution of the silicon design flow) has no comparable artifact, which is why Section 10.3.5 treats it as strategically decisive *if* verified and flags the verification event explicitly.

The independent record is thin and worth stating precisely. Crowdsourced Arena voting places the model #5 overall in text (highest-ranked Chinese model, behind Claude variants), #2 in vision, and #4 on frontend code at 1,668 Elo—eight points behind Kimi K3 and roughly 37 behind the top Claude configuration [399]. Artificial Analysis had not scored it. One small third-party head-to-head put it at 80 against Kimi K3’s 83. Set against the vendor table, the honest summary is that Qwen3.8-Max is a genuinely frontier-competitive model whose exact placement is unsettled—which is the same thing that is true of every closed Western flagship, none of which discloses its parameter count at all.

H.3 The second week: 4–11 August

The cycle did not close; it widened, and the widening changed its character from an intra-Chinese barrage to a two-bloc contest for the commons.

Meta returns to open weights (10 August). Muse Glimmer—30B dense, multimodal, ~128K context, distilled from the closed Muse Spark flagship—shipped with downloadable weights under

Apache 2.0, Meta’s first open release in sixteen months and its first ever under an OSI-permissive licence at this scale; Zuckerberg announced it personally on X alongside a 6,500-word letter recommitting Meta to “American open source,” and an open-weight version of Muse Spark 1.2 was promised “in the coming weeks,” licence unspecified [473, 474]. Independent measurement places Glimmer at 35 on the Artificial Analysis index of August (v4.1.1, before the September re-basing)—between Gemma 4 31B (30) and Qwen3.6 27B Reasoning (38), with a measured hallucination rate of 82% against the comparable Qwen’s 49% [475]. Section 10.3.2 carries the full analysis; the structural point is that the ratchet mechanism now operates on both sides of the Pacific, while the ownership of the commons it enlarges remains contested.

Qwen3.8-Max weights slip (status at 11 August). The 3 August commitment—“next week”—was not met in that week; the window has been publicly re-framed to the week of 10 August, and no repository exists at this edition’s date [476]. In the same days, Reuters-sourced reporting described Alibaba preparing revenue-sharing licence terms for large commercial resellers of its next open release [749], and Alibaba opened its Qianwen platform to third-party agent developers with SF Express, Ziroom, and Midea as launch partners [478]—monetization moving to the platform layer while the weights wait. The forecast register holds the 90% shipping forecast open with the slippage documented; the licence forecast is revised in place, with the revision dated and reasoned.

Openness acquires borders. MiniMax’s H3 video model tops the open-weight video arena while its licence excludes local deployment in the US, EU, UK, and South Korea—the clearest geofenced “open” release to date, and a datapoint the trust chapter’s certification argument must now carry [477].

The trade war resumes around the stack. On 5–6 August Beijing imposed case-by-case export review on drones and key components, sanctioned six US entities, and opened a national-security probe into imported printers—the broadest countermeasure package since the October truce, replicating Washington’s own playbook of chokepoint administration [479]. The same week, NIST joined the Genesis Mission with a program whose explicit goal is a tenfold scaling of US drone production within two years [482]—the two governments now running mirror-image industrial policies on the same component class.

Two smaller entries. The White House finalized a voluntary frontier-model evaluation framework in a closed-door session with the major US laboratories and declined to publish the document—voluntary and opaque where the Chinese regime is mandatory and registered, a contrast Chapter 20 predicts and records [480]. And CoreWeave entered its Q2 print in visible distress—15% short interest, debt above \$25B against \$30–35B of annual capex guidance, two downgrades—the neocloud vintage argument of Chapter 21 approaching its first earnings-season test [483].

H.4 The third week: 11–15 August

Four days that moved three separate arguments in this book, and one that moved the ground under all of them.

The Qwen commitment resolves, and splits. Alibaba shipped the 2.4T flagship’s weights to Hugging Face as **Qwen3.8-2.4T-A95B** on or about 13 August, with **Qwen3.8-27B** following on 15 August [484, 485]. The registered shipping forecast resolves YES; the registered permissive-licence forecast resolves NO for the flagship, which carries a US\$50M-revenue-threshold commercial licence, and YES for the 27B, which is Apache 2.0. Section 10.3.1 carries the analysis. Two press errors are corrected there and repeated here because they propagated widely: the flagship is *not* Apache-licensed, and it is *not* multimodal.

The harness arrives as a competitive layer. On 13 August DeepSeek released an MIT-licensed, plugin-native agent harness that passed 100,000 GitHub stars within about forty-eight hours, accompanied by an eighty-eight-page programming-languages paper from Peking University

and DeepSeek, and shipped on the same day the company *raised* its API prices [486, 489, 492]. Chapter 11 carries the analysis and exists because of it.

Model releases, four in five days. Grok 4.6 (12 August, closed, 500K context, \$2/\$6 per Mtok) reached 61 on the Artificial Analysis index as it stood in launch week (v4.1.1; the index was re-based on 7 September and no score on that board is comparable to a v4.3 score [758])—tying GPT-5.6 Sol (maximum reasoning effort), two points below Claude Opus 5, and one point above the best Chinese open model—restoring xAI, now styled SpaceXAI, to the front rank; the gains come from post-training on the same 1.5T Grok 4.5 base rather than a new pretrain [505]. Google shipped Gemini 3.7 Flash (13 August) at \$0.75 per Mtok input [508]. DeepSeek launched V4-Pro on 13 August and *raised* API prices by as much as 1,100%, introducing peak and off-peak tiering; Caixin described it as “a break from the prolonged price war among China’s AI developers,” and no rationale was stated [719]—and the company is reported to be hiring chip-design engineers to build its own accelerator, reducing reliance on NVIDIA *and* Huawei, which belongs to Chapter 13’s ledger rather than this one [509]. Zhipu announced GLM-5.3 (14 August) claiming the strongest open-weights coding model; the weights followed on **28 August** under a custom “glm-5.3” licence reportedly requiring a security review of model-as-a-service operators above \$10B of trailing revenue—so the claim is now checkable, and on the two headline agentic-coding suites the model still *trails* GPT-5.6 Sol [711]. Nine days after the announcement Zhipu released **GLM-5.3-Flash** (26 August, MIT, 320B total and 18B active), revealed on 27 August to have been the stealth model “Ox Alpha” that had led the router boards all week, and stated that the preview had run “entirely on Chinese AI chips” [712, 715]. Section 10.3.3 grades the serving claim and its three qualifications—served is not trained, no chipmaker is named, and it is not a first.

A social platform ships an agent model (14 August). Xiaohongshu’s dots studio released *dots3-note Preview* under Apache 2.0—280B total, 16B active, text, image, video and audio input, 512K context, weights on Hugging Face and ModelScope, free on OpenRouter—with a reinforcement-learning method, TEMPO, aimed at agent trajectories of ten hours and more, and self-reported SWE-bench Verified 78.4 and SWE-bench Pro 61 [662] [V for the release and licence; P for the benchmarks]. Huawei’s Ascend ecosystem published day-zero vLLM-Ascend deployment on Atlas 800 A3 and Atlas 900 A3 within two days [665] [A]. Two disciplines apply: the model is labelled a preview and its technical report is “coming soon,” so the method is described from the release material and graded accordingly; and Vercel’s July gateway report—open-weight models at 29% of enterprise gateway tokens on under 4% of spend in June [668] [V]—is entered here as the cleanest measurement yet of the volume–revenue split that Chapter 21 argues over. Section 10.3 and Chapter 11 carry the analysis: what the release records is that frontier-scale, agent-trained open models are now something a Chinese consumer platform is expected to ship.

NVIDIA and the open commons: a report, not an announcement. *The Information* reported on 12 August that NVIDIA is in the final training phase of a trillion-parameter “Nemotron 4” intended to compete with the world’s most powerful *open* models, with possible availability in late autumn [507]. NVIDIA’s own position, as quoted, is that final training is not complete; no licence has been announced; and the company’s best shipped open model today scores 38 on the independent index against Kimi K3’s 60, both on the August (v4.1.1) board; the September re-basing moved every score and the two figures are comparable only with each other [758]. The event, correctly stated, is that the largest beneficiary of the closed-frontier trade is reported to be preparing an open frontier-scale release—which, if it ships, is a second American entry into the commons after Meta’s and a further datapoint for the ratchet of Section 10.3.1. It has not shipped.

H.5 The financing turn: the same week, in capital markets

Three developments in the week of 10–14 August belong to Chapter 21 rather than to the model chronology, and together they describe the trade acquiring market infrastructure.

NVIDIA mobilizes third-party capital (10 August). Memoranda of understanding with Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR to establish compute financing platforms targeting over \$500B of third-party capital “over time,” subject to final agreements; the subhead’s object is turning NVIDIA compute into “an investable asset class” [692]. Four days later, reporting had NVIDIA cutting its contemplated backstop of OpenAI’s Ohio campus from roughly \$250B to under \$120B after investors raised concerns about its risk exposure [516]. Read together: exposure relocating from the vendor’s balance sheet to institutional capital.

OpenAI’s commercial leadership turns over before a listing (11–13 August). The longest-tenured business executive and the first chief revenue officer departed within four days, with the finance function conspicuously intact—the CFO briefed investors on 14 August, reporting an annualized run rate near \$40B and enterprise revenue overtaking consumer [520]. No listing date is set; prediction markets priced a 2026 listing near 19%.

Compute becomes an exchange-traded product (11 August). CME and Silicon Data announced compute futures for 5 October, listed on NYMEX—compute filed as an energy derivative—joining ICE’s two announced families, one of them energy-normalized with a dedicated inference sub-index, and Kalshi’s live GPU-price contracts [522, 524, 525]. Section 21.7 carries the analysis and works through what a traded, benchmarked, backwardated market for compute does to pricing power—and why the Chinese lever turns out to be the model layer rather than the GPU-hour.

H.6 The memory question, entered into the record

Two developments concern memory rather than models, and both are conditional claims that the register now carries.

The Open Compute Project published the first High Bandwidth Flash specification on 3 August 2026—512 GB per stack, 0.4 to 3.0 TB/s, UCIe rather than PCIe, with SanDisk, SK hynix, Google and Tenstorrent as members and NVIDIA, Samsung, Micron and Kioxia absent [569, 570]. Nothing is sampling; no endurance, latency, power or price figure has been published. Section 12.9 works through what a terabyte-class weights tier paired with an LPDDR6X context tier would do to the economics of local inference. Non-volatility is the term that matters: a weights tier that draws nothing at rest, at a realistic duty cycle, reaches energy parity with continuously refreshing DRAM, and the residual difference is tens of dollars a month against a metered service costing thousands. The question turns on HBF shipping at its specified capacity and bandwidth, not on its energy figure. Two facts frame the conditional in opposite directions—NAND is forecast into surplus in 2H27 while DRAM stays scarce, which is a durable relative-price tailwind; and NVIDIA removed its capacity-first inference part from the roadmap while paying a reported \$20B for the opposite architecture, which is the best-informed buyer in the industry betting against capacity.

Alongside it, the price record for this edition: conventional DRAM contract prices compounded roughly $3.6\times$ across the first three quarters of 2026, and Apple withdrew the 512 GB Mac Studio configuration in March—the flagship consumer machine for local inference on large open models becoming *less* capable because of memory scarcity, in the same year the commons reached terabyte scale [563, 566].

H.7 The Pareto observation

The single most strategically significant datum of the barrage is not any model’s score but the shape of the cost-performance frontier it produced: of the five models on the current Pareto

frontier, *four are Chinese*—Kimi K3, Qwen3.8-Max, GLM-5.3-Flash, and DeepSeek V4-Flash—with Claude Opus 5 the lone American entry, holding its position by roughly 37 Elo points [402]. The frontier as originally computed put GLM-5.2 in that fourth place; the late-August releases of Section 10.3.3 displace it, because GLM-5.3-Flash dominates it on both axes—index 57 at \$0.09 per task on the launch-week v4.1.1 board; 42 at \$0.25 per task on the re-based v4.3 board of 7 September, the two pairs not being comparable with each other [758, 759]—so the substitution strengthens the observation rather than weakening it. Read against the five-frontier framework of Section 10.5, this is the *economic frontier* passing decisively into Chinese hands while the *absolute frontier* remains contested, which is precisely the configuration this book has argued decides industrial outcomes. The cost spread that produces it: DeepSeek V4-Flash runs a benchmark suite for about three cents against Claude Fable 5’s \$3.15 and GPT-5.6 Sol’s \$1.86—a hundredfold gap at the low end, with Kimi K3 at 86 cents in between.

Two qualifications keep the observation honest. Cross-vendor cost comparisons at fixed capability are only as good as the task-normalization behind them (Appendix G, fields 9–10), and reasoning-effort settings move scores enormously—V4-Flash’s HLE score ranges from 8.1 with thinking disabled to 34.8 at maximum effort, which is larger than most inter-model gaps anyone reports. And V4-Flash’s own release note contains the field’s most useful piece of self-knowledge: its entire generation-over-generation gain came from post-training, not architecture, with Terminal-Bench 2.1 moving 61.8→82.7 on an unchanged model [401]. The frontier is currently moving through the post-training layer, which is cheap, fast, and exactly where a fast-follower has the advantage.

One further pricing development belongs in the ledger because it runs *against* the commoditization current, and it arrived in two dated steps rather than one. First, with the official V4-Flash release of 31 July 2026, DeepSeek announced dynamic peak pricing—a $2\times$ multiplier across seven daily hours—on the cheapest frontier-adjacent model in the world [401]. Then, on 13 August 2026, the company launched V4-Pro and *raised* API prices by as much as 1,100%, introducing peak and off-peak tiering across the line; Caixin described it as “a break from the prolonged price war among China’s AI developers,” and no rationale was stated [719]. A Chinese laboratory introducing surge pricing and then repricing upward is rent recapture appearing at the price layer itself, and Section 23.6’s recommendation that Chinese labs move from below-cost tokens to at-cost tokens plus paid complements now has a datapoint in its favour that arrived without waiting for the advice.

H.8 What this changes, and what it does not

Unchanged: every structural claim in the monograph. The four propositions, the cost stack, the rent-migration map, the architecture argument, and the AI4S analysis are all indifferent to which laboratory leads in a given fortnight—which is the point of freezing the data and stating the mechanisms separately.

Strengthened: three things, each recorded in the ledger and register. First, and most importantly, the open-weight argument acquires the *ratchet* of Section 10.3.1: a Max tier that had been proprietary through four generations was publicly committed to open weights within seven days of a rival opening at comparable scale, which upgrades openness from a state policy that could reverse to a competitive equilibrium from which no participant profits by defecting. Second, P3’s transmission mechanism has its first direct observation in the OpenAI price cut. Third—and this is the datum that most shortens the book’s own timelines—the autonomous silicon-design claim of Section 10.3.5 identifies a concrete mechanism by which the compute dependence that the offshore-leasing arrangement currently reflects becomes transitional rather than structural: the recursive loop applied to the one input China has never been able to buy.

The counter-argument, at its strongest. The same evidence supports a serious rebuttal—that recursive self-improvement compounds fastest where compute is most abundant, which is not

China—and Section 28.2.1 develops it as the book’s most consequential falsification branch, along with the reply that two recursive loops are running on different substrates.

Newly falsifiable, on a short clock: the Qwen3.8-Max licence question has since resolved (a revenue-threshold licence, recorded above); still open is whether any AI-assisted design flow produces an independently verified, manufacturable part. Few strategic questions in this book resolve within a quarter. The first of these does.

Appendix I

Chronology

1954–2007	Bell Labs silicon cell (1954). German EEG feed-in tariffs (2004) create global solar demand. Sharp then Q-Cells lead production. Suntech founded (2001), NYSE listing (2005). China’s 863 programme adds EV powertrains (2001); Wan Gang appointed minister (2007).
2008–2013	Polysilicon peaks at \$475/kg (2008) then collapses. “Ten Cities, Thousand Vehicles” (2009). CDB extends ~\$47B credit lines (2010); NEVs and new energy designated Strategic Emerging Industries (2010). CATL founded (2011). Solyndra (2011), Q-Cells (2012), A123 (2012), Suntech (2013) fail. US and EU impose solar duties (2012–13).
2014–2019	Big Fund I (2014). Battery “white list” excludes Korean cells (2015–19). Top Runner programme (2015). NEV dual-credit mandate (2017). Solar “531” shock (2018). EUV denied to China (2019). SVE (2016) and A64FX (2018) establish the CPU-plus-HBM lineage.
2020–2022	Fugaku sweeps the benchmark lists (2020–21). BYD Blade LFP (2020); Tesla adopts CATL LFP (2020). SME specified (2021). China NEV penetration passes 25% (2022). US October 2022 export controls.
2023	Mate 60 ships SMIC 7 nm (Aug.). A800/H800 banned (Oct.). China passes Japan as largest auto exporter. Android RISC-V support announced; RISE founded.
2024	DeepSeek V2 triggers the LLM price war (May). Big Fund III, RMB 344B (May). CXMT scales DRAM; Northvolt files Chapter 11 (Nov.). BIS controls HBM (Dec.). DeepSeek V3 in native FP8 (Dec.). China installs 277 GW of solar; adds ~429 GW of generation.
2025	R1 released (Jan.); NVIDIA loses \$589B in a day. Eight-agency RISC-V guidance reported in preparation (Mar.; never confirmed published) [A]. H20 banned (Apr.), unbanned for a 15% revenue share (Aug.). Kimi K2 opens at 1T parameters (Jul.). WAIC Global AI Governance Action Plan and “AI+” plan (Jul.–Aug.). CANN open-sourced (Aug.); CUDA ported to RISC-V hosts, announced in Shanghai (Jul.). Huawei discloses the Ascend 950/960/970 roadmap and opens UnifiedBus 2.0 (Sep.). CAC bars major platforms from NVIDIA purchases (Sep.); state-funded data centres restricted to domestic chips and given up-to-50% power discounts (Nov.). K2 Thinking ships INT4 QAT (Nov.). China adds ~540 GW of generation; deploys ~65% of global grid storage.
2026 H1	Zhipu and MiniMax list in Hong Kong (Jan.). SiFive integrates NVLink Fusion (Jan.). Trade press reports GLM-5 trained on ~100k Ascend 910B (Jun.; no locatable primary) [A/C]; Anthropic alleges distillation campaigns (Feb.). Stanford AI Index puts the US–China gap at 2.7% (Mar.). NVIDIA Vera CPU announced at 1.2 TB/s LPDDR5X (Mar.). DeepSeek V4 (Apr.); OpenAI raises at \$852B (Mar.); Anthropic at \$965B (May). Semiconductor selloff sheds \$1.3T (Jun.). LineShine takes No. 1 on HPL and HPCG with no GPUs (Jun. 24).
2026 Jul–Aug	Kimi K3 opens at 2.8T parameters (Jul. 27). CXMT lists at +466% (Jul. 27), prospectus disclosing no commercial HBM. CXMT LPDDR6 clears verification at 12.8 Gb/s (Aug.). Meta returns to open weights with Muse Glimmer under Apache 2.0 (Aug. 10). DeepSeek Harness (Aug. 13); Xiaohongshu dots3-note Preview, 280B/16B, Apache 2.0 (Aug. 14). Beijing curbs drone exports; printer probe (Aug. 5). OpenAI–Hugging Face cross-boundary agent intrusion disclosed (Jul.); Anthropic discloses three evaluation-boundary incidents (Jul. 30). Memory-led chip rout; Nasdaq-100 enters correction (Jul.). WAICO formalized in Shanghai with 29 states (Jul.).
2026 Aug–Sep	Brazil contracts Huawei/iFlytek for a Rio Ascend machine, R\$1.3B, self-financed, inside a split programme (Aug. 21). Moonshot–Microsoft/Amazon/Google revenue-share talks on Kimi K3 reported, up to 30%, unsigned (Aug. 26); Huawei bid for Egyptian government AI datacentres on Ascend 950 reported, US counter-offer assembled (Aug. 26) [A]. Commerce rule to block Chinese rental of NVIDIA compute in third countries reported in draft (Aug. 28) [A]. Xi in Cairo; no AI agreement (Aug. 31–Sep. 2). Moonshot confidential Hong Kong filing reported, ~\$3B at ~\$50B (Sep. 3) [A]. DeepSeek order of ≥160,000 Ascend 950DT for an Ulanqab inference site reported; training stays on NVIDIA (Sep. 4) [A]. Telekom Malaysia evaluating Ascend 910C for a RM2B sovereign project reported, unsigned (Sep. 7) [A]; Artificial Analysis re-bases to v4.3 (Sep. 7). NSA/FBI/CISA advisory AA26-251a names six Chinese laboratories for distillation (Sep. 9).

Appendix J

Claim Register and Falsification Dashboard

Appendix D grades the headline claims. This appendix does two further things: it gives the register a machine-readable schema, so that the ledger can be maintained as data rather than as prose; and it states the quarterly dashboard by which a reader can check this book against reality without re-reading it.

J.1 Register schema

Each claim in the maintained register carries: `claim_id`; `text` (one sentence, falsifiable); `chapter`; `grade` $\in \{V, P, A, M, R, S\}$; `source_type` $\in \{\text{benchmark submission, filing, government rule, peer-reviewed, vendor statement, analyst estimate, author model, discussion record}\}$; `as_of` (date); `geographic_scope`; `uncertainty` (range or band); and `next_verification_event`—the specific publication, filing, benchmark round, or disclosure that would update it. The last field is the one that makes the register useful: a claim with no identified verification event is a claim this book cannot keep honest, and there are several.

Table J.1: Register extract: the twelve claims on which the argument most depends.

ID	Claim	Grade	Uncertainty	As of / scope	Next verification event
CR-01	Chinese-origin models carry a majority of tokens on neutral routers	V	platform-specific	Jun 2026, Open-Router	monthly router reports
CR-02	Top US-China model gap $\approx 2.7\%$; closed-over-open gap $\approx 3.3\%$	V	harness-dependent	Mar 2026, global	AI Index 2027
CR-03	Open-weight share of enterprise <i>spend</i> fell 19% $\rightarrow$ 11% in 2025	A	survey selection	2025, US enterprise	next enterprise survey round
CR-04	LineShine leads HPL and HPCG with no discrete GPU	V	none material	Jun 2026	Nov 2026 TOP500
CR-05	LX2-class serving is competitive per TB/s	M	delivered-bandwidth fraction η (§12.2.1) swept 0.53–0.70	modelled	MLPerf Inference submission
CR-06	LPDDR6X socket: 2.3 TB, ~ 7 TB/s, 67% MFU ceiling	M	$\pm$ knee position unmeasured	modelled	memory-controller simulation; 5 TB/s knee test
CR-07	CXMT LPDDR6 mass production in 2H26	A/R	schedule risk; no primary announcement	analyst/roadmap	official product announcement or customer qualification
CR-08	Domestic HBM ~ 2 M stacks in 2026	A	2.5–18.6M in published range	analyst	Ascend shipment tear-downs
CR-09	A frontier-class model trained end-to-end on domestic silicon	—	not yet observed	—	the Mate 60 moment : independent confirmation of an end-to-end domestic run
CR-10	AI4S frontier share of expenditure 76–91%	M	f and r both estimates	discussion record	national probe with published f , r
CR-11	Announced fleets provision 0.02–0.10 BGPU/researcher vs ~ 1.0 required	M/A	core $\pm 40\%$ pre-measurement	3 countries	second national demand probe
CR-12	Broad ($\geq 30\%$) US AI repricing by 2029–30	S	60–70% subjective	global	OpenAI financing event; DC credit spreads

J.2 Probabilistic forecasts

The register states grades; this table states *bets*. Each load-bearing proposition carries a probability, a horizon, a next verification event, and an explicit falsifier. The purpose is not false precision—the probabilities are subjective judgements [S]—but to prevent confidence in one part of the thesis from leaking into another, and to let a reader in 2029 mechanically score the book.

Table J.2: Load-bearing propositions as dated forecasts [S]. Probabilities are the author’s; where the strongest skeptical case the author can construct yields a different figure, it is given in brackets.

Proposition	Probability	Horizon	Next verification event	Explicit falsifier
Chinese models dominate the global open-weight ecosystem	80% [70%]	1–2 yrs	derivative-model share, neutral-router token share, production-deployment surveys	sustained routed-token share below 45% for four quarters
Frontier-adjacent model trained end-to-end on a domestic Chinese stack	65% [50%]	2–4 yrs	independently documented run: domestic accelerator, memory, toolchain	continued reversion to foreign silicon for flagship pretraining past 2028
LPDDR-dominant or hybrid frontier pretraining becomes production-competitive	45%	3–6 yrs	cluster-scale run with package-complete power, reliability, cost	measured MFU below $\sim 45\%$ at scale, or knee above 8 TB/s
An LX2-class part posts competitive AI benchmark results	55%	2–4 yrs	MLPerf Inference or Training submission	no submission by 2029, or $>2\times$ deficit at matched precision
Broad ($\geq 30\%$) repricing of the US AI complex	70% [60%]	1–3 yrs	utilization, margin, financing, depreciation deterioration	AI revenue growth closing the capex gap through 2028
US closed models retain a valuable premium enterprise tier through 2030	50% [65%]	4 yrs	enterprise spend share; agent-platform revenue	open-weight enterprise spend share passing 40%
Complete Chinese autonomy across advanced semiconductor tooling	25%	5–10 yrs	verified production across lithography, metrology, EDA, packaging, memory	domestic EUV-class litho absent past 2032
A neutral certification authority for open weights is established	30%	3–5 yrs	founding instrument with ≥ 3 states and operating evaluations	the role occupied by a single bloc, or by nobody, past 2030
AI4S delivers measurable national research-output uplift ($\geq 5\%$)	45%	5–8 yrs	bibliometric and patent series; sector productivity	no detectable uplift in the domains with best conversion (weather, structure, screening)
A systemic agentic-security incident traceable to a widely deployed derivative	20%	5 yrs	incident disclosures; national CERT reporting	(asymmetric: absence is weak evidence)

continues

Forecasts, continued

Proposition	Probability	Horizon	Next event	verification	Explicit falsifier
The integrated-CPU corner reduces lifecycle cost of a <i>changing</i> workload portfolio faster than the discrete-accelerator corner (§12.5.1)	40%	4–8 yrs	time-to-competitive-kernel per new mechanism; architectures served within N months; operator-hours per delivered PF-month		the accelerator corner matches new-mechanism kernel latency while the CPU corner’s porting tax persists
A production-metric scoreboard for scientific AI (Eq. 24.3-class) is operated by a national ministry or agency	35%	3–5 yrs	a published national portfolio scored on validated value per resource		benchmark-style leaderboards remain the only public measure past 2030
A Chinese-designed accelerator or integrated-CPU system submits MLPerf Training (which now carries a DeepSeek-V3 MoE benchmark)	50%	2–3 yrs	MLPerf Training rounds, twice yearly		no submission by end-2028 while domestic fleets claim training parity
Humanoid/emodied deployment crosses from demonstration to production: $\geq 1,000$ robots in intervention-light productive work at a single operator, metrics published (§16.5)	55%	2–4 yrs	operator disclosures; the embodied proof-suite metrics		teleoperation ratios not declining through 2028; deployments remain demo/education-dominated
Open-weight VLA models commoditize the embodied task-policy layer as LLMs did text	45%	3–6 yrs	open VLA releases matching closed policies on standardized manipulation suites		closed embodied stacks hold a durable capability gap into the 2030s
Qwen3.8-Max weights ship as publicly committed	RESOLVED—YES (13 Aug 2026)		shipped to Hugging Face; product name not used on the repository [484]		—
... under a permissive (Apache/MIT-class) rather than revenue-threshold licence	RESOLVED—NO for the flag-ship; YES for the 27B		custom licence, US\$50M revenue gate [484]		—
The next frontier-scale Chinese release ($\geq 1\text{T}$) also ships open weights—the ratchet holding	85%	1–2 yrs	release announcements	announcements	any major Chinese lab closing its frontier tier while rivals stay open

continues

Forecasts, continued

Proposition	Probability	Horizon	Next event	verification	Explicit falsifier
Meta ships the promised open-weight Muse Spark 1.2 within a quarter, at flagship-representative quality, under derivative-permissive terms (all three conditions of §10.3.2)	40%	1 quarter	repository, licence file, independent evaluation		the Max-tier template repeating: small model open, flagship deferred or restricted
A Chinese-origin agent harness or agent runtime reaches top-three position by production deployment in at least one major non-Chinese market (Chapter 11)	35%	2–4 yrs	enterprise deployment surveys; package-registry download share, not stars		Western harnesses holding >90% of regulated-sector production deployments through 2029
Harness migration in regulated sectors is measurably costlier than in unregulated ones—the certification mechanism of Chapter 11	60%	2–3 yrs	published migration accounts with elapsed calendar time and assessor cost		migrations completing at comparable cost in both, which would collapse the switching-cost argument to configuration only
NVIDIA ships a frontier-scale open-weight model (the reported Nemotron 4 or successor) under a licence permitting derivatives	55%	2–3 quarters	repository and licence file		the model shipping closed, or not shipping, through 2027
Exchange-traded compute futures reach sustained liquidity—continuous two-sided markets and open interest disclosed by the venue—rather than remaining reference benchmarks (§21.7)	40%	2–3 yrs	CME, ICE and Kalshi volume and open-interest reporting after the Oct. 2026 launch		thin or discontinuous trading through 2028, leaving the indices as pricing references only
A dedicated <i>inference</i> contract (COIL-I class) trades independently of GPU-rental contracts, establishing tokens rather than GPU-hours as the traded unit	30%	2–4 yrs	venue product listings and settlement data		inference pricing remaining bundled into GPU-hour or per-vendor API terms
Chinese inference capacity is exported at scale—token-denominated services to non-Chinese buyers materially visible in services-trade statistics or neutral-router data (§21.7)	45%	3–5 yrs	customs/services-trade series; neutral-router provider shares; enterprise procurement surveys		Chinese endpoint share stalling in Western regulated sectors while open-weight self-hosting carries the volume instead

continues

Forecasts, continued

Proposition	Probability	Horizon	Next event	verification	Explicit falsifier
HBF-class silicon ships in volume at its specified capacity and bandwidth (≥ 512 GB per stack, ≥ 1.6 TB/s) with a published endurance figure adequate for a weights tier (§12.9)	45%	2–4 yrs	vendor datasheets; independent measurement; OCP conformance scheme		no volume product by 2029, or shipped bandwidth below ~ 0.8 TB/s per stack, which would leave a 2.8T-parameter model below reading speed for a single user
A capacity-first appliance runs a ≥ 2 T-parameter open model for four concurrent users at ≥ 40 tokens/s each inside a 1 kW envelope, at a system price below \$8k (§12.9.4)	35%	3–5 yrs	published measurements on shipping hardware		the configuration remaining above \$15k, or below 20 tok/s per user at four users, through 2030
Enterprise open-weight workload share reverses its decline and passes 20%	40%	2–4 yrs	enterprise adoption surveys		continued decline, which would establish that licensing, governance and support—not memory—are the binding constraints
A hyperscale-class AI facility (≥ 100 MW) is sold, foreclosed or written down at $\geq 40\%$ below cost before 2030—the receivership case of Chapter 21	45%	2–4 yrs	REIT and neo-cloud disclosures; distressed-asset transactions; lender filings		every 2025–26 vintage refinancing at par through 2029, which would mean the facility utilization forecasts held
A Chinese frontier-class model is trained end-to-end on domestically sited compute, offshore leasing excluded	60%	2–4 yrs	lab technical report disclosing training site and silicon		continued reliance on leased offshore accelerators past 2029
An AI-assisted design flow produces an independently verified, manufacturable Chinese AI accelerator	45%	3–5 yrs	tape-out disclosure with third-party design-flow audit		domestic accelerator roadmaps continuing to slip despite AI-design claims
Merchant-GPU data-centre revenue growth decelerates while merchant custom-ASIC revenue continues at Broadcom’s guided trajectory—the silicon-commoditization pincer of Chapter 17	50%	to end-2027	quarterly segment disclosures from NVIDIA and Broadcom [691]		merchant-GPU data-centre growth holding while custom-ASIC growth decelerates, which would leave the pincer unobserved

continues

Forecasts, continued

Proposition	Probability	Horizon	Next event	verification	Explicit falsifier
Custom accelerators continue converging onto NVIDIA’s interconnect and memory rather than displacing them—the absorption reading of Chapter 17	65%	to end-2027	accelerator announcements and teardowns; Trainium4 on NVLink Fusion and NVIDIA custom HBM [690]	an-	a major custom program shipping a scale-up fabric and memory stack independent of NVIDIA’s
First confirmed Ascend export outside China (Egypt or Malaysia; Brazil’s self-financed Rio machine, due July 2027, is excluded as already contracted)—resolves the telecom-template question of Chapter 6 and §22.3	40% [25%]	to end-2027	a signed, disclosed contract and a delivered installation; press reporting confirmed by the recipient government [765, 766]		no signed Ascend contract in Egypt or Malaysia by end-2027, or both projects awarded to US vendors, which would leave the export offer at one self-financed contract in Brazil

J.3 Quarterly falsification dashboard

Fifteen indicators, each with the value as of 9 September 2026, the threshold that would move the argument, and where to look. A reader who checks these quarterly does not need to re-read the book to know whether it is aging well.

Table J.3: The dashboard. “Direction” states which way a move cuts for this book’s thesis.

Indicator	Current (Sep 2026)	Threshold that matters	Source / cadence
Chinese share of routed tokens	~61%	sustained fall below 45% weakens P1	neutral routers, monthly
Closed-over-open capability gap	3.3%	>10% sustained two years favours bifurcation	AI Index / arenas, annual
Enterprise open-weight spend share	~11%	rise past 25% confirms P3’s transmission	enterprise surveys, annual
Domestic end-to-end frontier training run	not observed	first credible instance strongly confirms P2	lab technical reports
LX2-class AI benchmark submission	none	MLPerf entry at parity confirms stage 2 (§12.2.1)	MLPerf rounds
CXMT LPDDR6 / HBM3 ramp	verification cleared	2+ year slip weakens the memory arbitrage	qualification news, quarterly
Cost per <i>solved</i> task, open vs closed	unpublished	first credible series either confirms or deflates the price gap	vendor or third-party evals

continues

Dashboard, continued

Indicator	Current (Sep 2026)	Threshold that matters	Source / cadence
OpenAI financing	IPO slipping	down-round or new backstop raises P4	filings, ad hoc
Data-centre credit spreads	Oracle CDS ~212 bps	sustained widening raises P4	credit markets, daily
Wholesale colocation rent, NoVA 10 MW+ (\$/kW-month, power passed through)	\$155–185; vacancy 0.5%	a sustained decline with rising vacancy is the facility-side signal that the utilization forecast of \$12.9.4 is failing; a rise says the closet is not yet displacing the hall	CBRE / JLL market reports, semiannual
Certification authority for open weights	vacant	occupation by any credible body reshapes Chapters 18–27	standards bodies, ad hoc
Max-tier openness lag (API-to-weights interval for the leading Chinese flagship)	~10 days realized (Qwen3.8-Max); 14 days (GLM-5.3)	a lag exceeding two quarters, or any frontier tier staying closed, would break the ratchet of §10.3.1	model repositories, per release
Recursive-loop disclosures: frontier-model contribution to successor training (US) vs AI-assisted design flow output (China)	both claimed, neither audited	first audited instance either side is the field's most important datapoint (§28.2.1)	lab disclosures, tape-out records
Offshore compute leasing by Chinese labs	permitted; reportedly in use; Commerce rule to close it reported in draft, unpublished, authority contested (28 Aug.) [756]	closure by rule would be the highest-impact tightening available to export policy	BIS/Commerce rulemaking
First confirmed Ascend export outside China	none; one contract (Brazil, delivery Jul. 2027), one bid (Egypt), one evaluation (Malaysia), all unfinanced by the vendor [765, 767]	a delivered, confirmed Ascend installation abroad on Chinese credit would move the telecom template of Chapter 6 from intent to execution	Bloomberg/Reuters; recipient-government disclosures; BIS enforcement

Methods, Sources, and Versioning

How this book is built

The manuscript is assembled from three kinds of material, and the reader is entitled to know which is which on any given page. *Public record*: benchmark submissions, government rules, company filings and earnings disclosures, national statistical releases, peer-reviewed papers, and specialist technical reporting—cited directly wherever a primary source exists, and flagged in the text when only secondary reporting was available. *Companion analysis*: architectural and performance results from the author’s own studies with collaborators, principally the accelerated-CPU analysis with Dongarra and Hoefer [184, 185], the memory-balance feasibility work [134], and the agentic-security survey [291]. Where these are load-bearing, the text says explicitly that it is presenting an author-modeled architecture or analysis rather than a measured production system. *Working documents*: unpublished discussion records, cited as such [306], whose numerical content is discussion estimate rather than measurement.

Every load-bearing claim carries an evidence grade (Section 1.5): [**V**] independently verified, [**P**] primary-source claim, [**A**] analyst estimate, [**M**] author model, [**R**] roadmap, [**S**] scenario judgment. Appendix D consolidates the headline claims in one graded ledger, which is the fastest way for a skeptical reader to locate the weakest link in the argument—and the first thing revised whenever the book is.

Known weaknesses

Six, stated plainly. First, several mid-2026 model and market details still rest on secondary aggregators where primary vendor documentation exists but was not retrieved; these are being replaced as the underlying documents are located. Second, the LX2 attributions—designer, process node, memory supplier—are analyst inference from public teardown-adjacent reporting, not disclosure, and are graded accordingly. Third, the LPDDR6X architecture (Section 12.8, Appendix A) remains a modeled design with a stated experimental program and no measured realization; readers should treat its numbers as a falsifiable proposal, which is how they are offered. Fourth, the flagship model cards of Appendix G are incomplete in exactly the fields that matter most—tokens and wall-clock per completed task, measured serving throughput—because essentially no vendor publishes them. Fifth, the national sizing results of Chapter 26 are measurement-anchored for one country only; the Singaporean and American instantiations are labelled pre-measurement scenarios whose behavioural parameters are transferred, and they should be read as the starting points those countries’ own probes would revise. Sixth, the book’s architectural thesis—that the CPU corner’s breadth advantage compounds faster than the accelerator corner’s depth advantage—is a judgement, not a measurement, and it is the judgement most likely to be shown wrong; the proof suite of Section 12.2.2 exists to settle it.

A note on the figures

The figures follow one convention throughout. Categorical colour is assigned in a fixed order from a palette validated for colour-vision deficiency (worst adjacent-pair separation $\Delta E \approx 9$ in deuteranopic simulation), never cycled and never used to encode rank; sequential magnitude uses a single hue; every chart with two or more series carries both a legend and selective direct labels, so identity is never carried by colour alone; grids and axes are recessive; and no chart uses two vertical scales. Every figure's underlying values appear either in an adjacent table or in the caption, which is also the required relief for the one palette slot whose contrast against the page falls below 3:1. Modelled quantities are marked [M] in the caption and plotted with a visibly different mark from measured ones.

Versioning protocol

This is a living document. Each revision increments the version, restates its date on the title page—which is also the book's currency date, so that the cutoff can never contradict the edition—and re-grades Appendix D. Three artifacts travel together and are meant to be revised at different rates: the monograph, which changes slowly and carries the argument; the machine-readable claim register and falsification dashboard of Appendix J, which change at every revision; and the current-cycle appendix, which is replaced wholesale rather than amended, so that the fastest-moving material never has to be threaded back through twenty chapters. Corrections are welcome and will be credited; the quickest way to change this book's conclusions is to falsify an entry in the ledger.

Acknowledgements and disclosure

The argument owes a substantial debt to co-authors and interlocutors on the companion works cited above. Errors of fact, interpretation, and judgment remain the author's own.

AI-assisted writing disclosure. Large language models were used for source retrieval, drafting assistance, and language editing throughout the preparation of this manuscript. The author reviewed and approved the final text and remains solely responsible for its content, analysis, conclusions, and any errors. Given the subject matter, it seems worth stating the obvious: the tools used to write a book about the commoditization of intelligence are themselves an instance of it.

Bibliography

- [1] Our World in Data, “Learning curves: solar module prices fell $\sim 20\%$ per doubling of cumulative capacity; \$106/W (1976) to \$0.38/W (2019).” <https://ourworldindata.org/learning-curve>
- [2] A. Jäger-Waldau / ISES, “PV Industry Trends” (Oct. 2024): 2023 China production shares—polysilicon 92%, wafers 98.1%, cells 92%, modules 84.6%. https://www.ises.org/sites/default/files/webinars/2024/October/Oct242024_PVIndustryTrends_final.pdf
- [3] EDN, “Q-Cells tops Sharp in solar rankings” (2008). <https://www.edn.com/q-cells-tops-sharp-in-solar-rankings/>
- [4] Wikipedia, “Feed-in tariffs in Germany” (EEG 2004: 57.4 ct/kWh rooftop, 20-year guarantee). https://en.wikipedia.org/wiki/Feed-in_tariffs_in_Germany
- [5] Wikipedia, “Suntech Power” (NYSE listing Dec. 2005; 2 GW capacity 2011; \$541M default and Wuxi bankruptcy, Mar. 2013). https://en.wikipedia.org/wiki/Suntech_Power
- [6] Eco-Business/Bloomberg, “Chinese renewable companies slow to tap \$47 billion credit” (CDB lines, 2010). <https://www.eco-business.com/news/chinese-renewable-companies-slow-to-tap-47-billion-credit/>
- [7] 12th Five-Year Plan for the Solar Photovoltaic Industry (2011), official translation. <https://policy.asiapacificenergy.org/sites/default/files/chinas-five-year-plan-for-solar-translation.pdf>
- [8] Bernreuter Research, “Polysilicon price trend” (\$475/kg Mar. 2008 $\rightarrow$ \$15.9/kg Dec. 2012 $\rightarrow$ \$5.6/kg Jul. 2024). <https://www.bernreuter.com/polysilicon/price-trend/>
- [9] Congressional Research Service, “Solyndra” (R42058); \$535M DOE guarantee, 2011 bankruptcy. <https://www.congress.gov/crs-product/R42058>
- [10] NS Energy, “Hanwha completes Q-Cells takeover” (insolvency Apr. 2012). <https://www.nsenergybusiness.com/news/newshanwha-completes-q-cells-takeover/>
- [11] PV Tech, “SolarWorld entering insolvency proceedings” (May 2017). <https://www.pv-tech.org/solarworld-entering-insolvency-proceedings/>
- [12] US Dept. of Commerce, AD/CVD final determinations on PRC solar cells (Oct. 10, 2012). https://enforcement.trade.gov/download/factsheets/factsheet_prc-solar-cells-ad-cvd-finals-20121010.pdf
- [13] Clifford Chance, “EU anti-dumping duties imposed against Chinese solar panel producers” (2013); Al Jazeera, “China and EU reach deal on solar panel trade” (Jul. 28, 2013). <https://www.aljazeera.com/news/2013/7/28/china-and-eu-reach-deal-on-solar-panel-trade>
- [14] Utility Dive, “Trump issues 30% tariff on solar panel imports” (Jan. 2018). <https://www.utilitydive.com/news/trump-issues-30-tariff-on-solar-panel-imports/515265/>
- [15] Climate Home News, “China solar industry struggles with sudden subsidy cuts” (531 policy, 2018). <https://www.climatechangenews.com/2018/08/15/china-solar-industry-struggles-sudden-subsidy-cuts/>
- [16] SolarQuotes, “China’s Top Runner Program”; PV Tech on PERC dominance of Top Runner rounds. <https://www.solarquotes.com.au/blog/chinas-top-runner-program/>
- [17] pv magazine, “China hits 277.17 GW of new PV installations in 2024.” <https://www.pv-magazine.com/2025/01/21/china-hits-277-17-gw-of-new-pv-installations-in-2024/>
- [18] TaiyangNews (NEA data), “China installed over 315 GW AC new solar PV capacity in 2025”; cumulative 1.2 TW AC. <https://taiyangnews.info/markets/china-installed-over-315-gw-ac-new-solar-pv-capacity-in-2025>

- [19] IEA, *Special Report on Solar PV Global Supply Chains* (2022).
<https://www.iea.org/reports/solar-pv-global-supply-chains/executive-summary>
- [20] PV Tech, “Chinese firms nine of top ten polysilicon manufacturers 2024” (93.5% share; sub-\$4.50/kg prices).
<https://www.pv-tech.org/chinese-firms-nine-top-ten-polysilicon-manufacturers-2024/>
- [21] PV Tech, “Global solar module manufacturing capacity to reach 1.8 TW in 2025.” <https://www.pv-tech.org/global-solar-module-manufacturing-capacity-to-reach-1-8tw-in-2025-report/>
- [22] InfoLink, “2025 global module shipment ranking: combined shipments reach 536 GW” (exports 267.6 GW; TOPCon ~95%). <https://www.infolink-group.com/energy-article/solar-topic-infolink-2025-global-module-shipment-ranking-combined-shipments-reach-536-gw>
- [23] Ember, “China solar cell exports grow 73% in 2025.”
<https://ember-energy.org/latest-insights/china-solar-cell-exports-grow-73-in-2025/>
- [24] EnergyTrend solar price index (Jul. 29, 2026): TOPCon modules RMB 0.67–0.72/W ($\approx$ \$0.09–0.10/W). <https://www.energytrend.com/solar-price.html>
- [25] PV Tech, “Five Chinese PV giants anticipate combined 2025 losses of US\$4.1–4.7 billion.” <https://www.pv-tech.org/five-chinese-pv-giants-anticipate-combined-2025-losses-of-us4-1-4-7-billion/>
- [26] SCMP, “China’s top leadership takes aim at disorderly low-price competition” (Jul. 2025); Xinhua commentary. <https://www.scmp.com/economy/china-economy/article/3316616/chinas-top-leadership-takes-aim-disorderly-low-price-competition>
- [27] pv magazine, “China’s \$7 billion polysilicon consolidation plan” (Aug. 2025).
<https://www.pv-magazine.com/2025/08/14/chinas-7-billion-consolidation-plan-may-be-reflected-in-prices-in-short-time/>
- [28] pv magazine, “Beijing group creates official body to guide China’s polysilicon capacity” (Dec. 2025). <https://www.pv-magazine.com/2025/12/11/beijing-group-creates-official-body-to-guide-chinas-polysilicon-capacity/>
- [29] pv magazine, “LONGi sets new world record with 35.5%-efficient perovskite-silicon tandem cell” (Jul. 2026). <https://www.pv-magazine.com/2026/07/15/longi-sets-new-world-record-with-35-5-efficient-perovskite-silicon-tandem-cell/>
- [30] SolarQuarter, “Trina Solar tops global perovskite solar cell patent ranking” (May 2025).
<https://solarquarter.com/2025/05/12/trinasolar-tops-global-perovskite-solar-cell-patent-ranking-reinforces-leadership-in-next-gen-pv-technologies/>
- [31] pv magazine USA, “U.S. solar manufacturing forges on” (module capacity 8 $\rightarrow$ 56.5 GW; cells lagging; Aug. 2025). <https://pv-magazine-usa.com/2025/08/19/u-s-solar-manufacturing-forges-on/>
- [32] pv magazine, “US solar faces steep cuts as OBBBA phases out key tax credits” (Sep. 2025).
<https://www.pv-magazine.com/2025/09/16/us-solar-faces-steep-cuts-as-obbba-phases-out-key-tax-credits/>
- [33] PV Tech, “First Solar net sales, gross profit tick upwards 2025” (16.1 GW produced).
<https://www.pv-tech.org/first-solar-net-sales-gross-profit-tick-upwards-2025/>
- [34] Battery Design, “Sony lithium-ion batteries: world’s first commercialized LIB (1991).”
<https://www.batterydesign.net/sony-lithium-ion-batteries-worlds-first-commercialized-lib-1991/>
- [35] Wikipedia, “BYD Company” (founded 1995; world #1 NiCd by 2002; Qinchuan acquisition 2003).
https://en.wikipedia.org/wiki/BYD_Company
- [36] Wikipedia, “CATL” (ATL 1999; CATL spin-off 2011; world #1 from 2017).
<https://en.wikipedia.org/wiki/CATL>
- [37] CGTN, “China scraps white list of domestic auto battery suppliers” (Jun. 2019).
<https://news.cgtn.com/news/2019-06-25/China-scraps-white-list-of-domestic-auto-battery-suppliers-HO26IZLsWI/index.html>
- [38] H. Sanderson, “The battery bet that changed the world” (Volt Rush Substack): LFP subsidy politics, Blade/CTP, Tesla-CATL 2020.
<https://voltrush.substack.com/p/the-battery-bet-that-changed-the>

- [39] SNE Research via electrive, “CATL continues to dominate global battery market” (FY2025: 1,187 GWh; CATL 39.2%, BYD 16.4%). <https://www.electrive.com/2026/02/04/sne-research-catl-continues-to-dominate-global-battery-market/>
- [40] IEA, *Global EV Outlook 2026*, “Electric vehicle batteries” chapter. <https://www.iea.org/reports/global-ev-outlook-2026/electric-vehicle-batteries>
- [41] US EIA, “China’s battery component production shares” (anodes 85%, electrolytes 82%, separators 74%, cathodes 70%). <https://www.eia.gov/todayinenergy/detail.php?id=65305>
- [42] Benchmark Mineral Intelligence, “China controls three quarters of graphite anode supply chain.” <https://source.benchmarkminerals.com/article/infographic-china-controls-three-quarters-of-graphite-anode-supply-chain>
- [43] Mining Technology, “China’s control of Indonesia’s nickel capacity” (~75% of refining). <https://www.mining-technology.com/news/chinas-indonesias-nickel-capacity/>
- [44] BloombergNEF, “Lithium-ion battery pack prices fall to \$108/kWh” (2025 survey; China \$84/kWh). <https://about.bnef.com/insights/clean-transport/lithium-ion-battery-pack-prices-fall-to-108-per-kilowatt-hour-despite-rising-metal-prices-bloombergnef/>
- [45] Ford Authority, “Ford EV battery plant will begin production this summer” (CATL-licensed LFP, Marshall MI, 2026). <https://fordauthority.com/2026/02/ford-ev-battery-plant-will-begin-production-this-summer/>
- [46] CarNewsChina, “CATL confirms 2026 large-scale sodium-ion battery deployment” (Naxtra: 175 Wh/kg, mass production Dec. 2025). <https://carnewschina.com/2025/12/28/catl-confirms-2026-large-scale-sodium-ion-battery-deployment-in-multiple-sectors/>
- [47] CNBC, “CATL claims to beat BYD’s EV battery record with longer range on a 5-minute charge” (2nd-gen Shenxing, Apr. 2025). <https://www.cnbc.com/2025/04/22/chinas-catl-claims-to-beat-byd-s-ev-battery-record-with-longer-range-on-a-5-minute-charge.html>
- [48] BYD, “Super e-Platform megawatt flash charging” (Mar. 2025). <https://www.byd.com/en/news-list/BYD-Unveils-Super-e-Platform-Megawatt-Flash-Charging-Electric-Vehicles-Matching-Refueling-Speeds.html>
- [49] HSF Kramer, “China battery technology export restrictions” (Jul. 2025); Global Times on Oct. 2025 extension to cells ≥ 300 Wh/kg, anodes, equipment. <https://www.hsfkramer.com/notes/mining/2025-posts/china-battery-technology-export-restrictions-cathode-lithium-gallium>
- [50] Wikipedia, “Northvolt” (Ch. 11 Nov. 2024; Swedish bankruptcy Mar. 2025; \$15B raised; 79.8 MWh produced Q1–Q3 2023). <https://en.wikipedia.org/wiki/Northvolt>
- [51] C&EN, “A123 Systems goes bankrupt” (2012; Wanxiang acquisition). <https://cen.acs.org/articles/90/i43/A123-Systems-Goes-Bankrupt.html>
- [52] Energy-Storage.News/CEA, “US battery gigafactories face delays and cancellations” (46 GWh cancelled by mid-2025). <https://www.energy-storage.news/us-battery-gigafactories-face-delays-and-cancellations-amid-market-uncertainty/>
- [53] Investing.com, “LG Energy Solution flags slowing EV demand, posts first quarterly loss in 3 years”; Reuters on Panasonic forecast cuts (2025). <https://investing.com/news/stock-market-news/lg-energy-solution-flags-slowing-ev-demand-posts-first-quarterly-loss-in-3-years-3828755>
- [54] Energy-Storage.News, “China deploys 65 GWh of BESS in December, ~65% of 2025 global total.” <https://www.energy-storage.news/china-deploys-65gwh-of-bess-in-december-25-of-2025-global-total/>
- [55] MIT Technology Review, “How did China come to dominate the world of electric cars?” (Feb. 2023). <https://www.technologyreview.com/2023/02/21/1068880/how-did-china-dominate-electric-cars-policy/>
- [56] DieselNet, “China: New Energy Vehicles” (policy summary: 863 program, Ten Cities, dual credit, tax exemption). <https://dieselnet.com/standards/cn/nev.php>
- [57] S. Kennedy, CSIS, “The Chinese EV dilemma: subsidized yet striking” (\$230.9B cumulative support 2009–2023). <https://www.csis.org/blogs/trustee-china-hand/chinese-ev-dilemma-subsidized-yet-striking>

- [58] ICCT, “China’s NEV mandate (dual-credit) policy update” (2018). https://theicct.org/sites/default/files/publications/China-NEV-mandate_ICCT-policy-update_20032018_vF-updated.pdf
- [59] EVCIPA via ChinaEVHome, “China’s EV charging facilities hit 19.32 million by Nov. 2025”; US comparison ~250k public ports. <https://chinaevhome.com/2025/12/23/chinas-ev-charging-facilities-hits-19-32-million-by-nov-2025-up-52-yoy/>
- [60] CnEVPost, “China NEV retail Dec. 2025 (CPCA)”: FY2025 12.86M, penetration 54.07%, Dec. 60.4%. <https://cnevpost.com/2026/01/07/china-nev-retail-1-387-million-dec-2025-preliminary-cpca/>
- [61] CnEVPost, “China NEV sales Dec. 2025 (CAAM)”: FY2025 NEV 16.49M; total 34.4M; exports 7.1M incl. 2.62M NEVs. <https://cnevpost.com/2026/01/14/china-nev-sales-1-71-million-dec-2025-caam/>
- [62] CarNewsChina, “China’s monthly vehicle exports exceed 1 million for the first time” (Jun. 2026). <https://carnewschina.com/2026/07/10/chinas-monthly-vehicle-exports-exceed-1-million-for-the-first-time-with-nevs-claiming-over-half/>
- [63] Electrek, “BYD crushes Tesla all-electric sales for 2025, secures global BEV crown” (4.55M total; 2.25M BEV vs. Tesla 1.64M). <https://electrek.co/2026/01/02/byd-crushes-tesla-all-electric-sales-for-2025-secures-global-bev-crown/>
- [64] Rhodium Group, “Why are Chinese EVs so cheap?” (BYD ~80% Tier-1 in-house; ~\$4,700/vehicle structural advantage vs. Tesla; foreign share 62%→35%). <https://rhg.com/research/why-are-chinese-evs-so-cheap/>
- [65] CarNewsChina, “BYD completed its massive fleet, now able to export 1 million cars a year” (Oct. 2025). <https://carnewschina.com/2025/10/02/byd-completed-its-massive-fleet-now-able-to-export-1-million-cars-a-year-but-not-just-from-china/>
- [66] CarNewsChina, “One Chery vehicle was exported every 23 seconds in 2025” (1.34M exports). <https://carnewschina.com/2026/01/02/one-chery-vehicle-was-exported-every-23-seconds-in-2025-maintaining-its-lead-position-for-23-years/>
- [67] CarNewsChina, “MG is first Chinese auto brand to sell more than one million vehicles in Europe” (Feb. 2026). <https://carnewschina.com/2026/02/22/mg-is-first-chinese-auto-brand-to-sell-more-than-one-million-vehicles-in-europe-and-uk/>
- [68] CarNewsChina, “Xiaomi YU7 got over 289,000 orders in 60 minutes” (Jun. 2025); “Xiaomi hits 600,000 deliveries in 22 months” (Feb. 2026). <https://carnewschina.com/2026/02/13/xiaomi-hits-600000-deliveries-in-22-months-sets-550000-unit-2026-target/>
- [69] EVBoosters, “400 Chinese EV companies ceased operations between 2018–2025.” <https://evboosters.com/ev-charging-news/400-chinese-ev-companies-ceased-operations-between-2018-2025-only-a-few-will-dominate-towards-2030/>
- [70] Fortune, “GM in China: from 4.04M (2017) to 1.8M (2024), \$5B+ in charges” (Jan. 2025). <https://fortune.com/asia/2025/01/28/gm-china-electric-vehicles-byd-buick-mary-barra/>
- [71] CarNewsChina, “Volkswagen Group China sales declined 8% y/y in 2025 (2.69M).” <https://carnewschina.com/2026/01/13/volkswagen-groups-sales-declined-by-8-y-o-y-in-china-as-the-top-selling-foreign-auto-brand-in-2025/>
- [72] CarNewsChina (exclusive), “Volkswagen to licence XPeng’s autonomous driving solution for its China EVs in 2026”; VW Group and Stellantis-Leapmotor releases. <https://carnewschina.com/2025/10/06/exclusive-volkswagen-to-licence-xpengs-autonomous-driving-solution-for-its-china-evs-in-2026/>
- [73] The White House, “President Biden takes action to protect American workers...” (100% EV tariff, May 14, 2024). <https://bidenwhitehouse.archives.gov/briefing-room/statements-releases/2024/05/14/fact-sheet-president-biden-takes-action-to-protect-american-workers-and-businesses-from-chinas-unfair-trade-practices/>
- [74] Cleary Trade Watch, “Definitive duties adopted by the EU on Chinese battery electric vehicles” (Oct. 2024: SAIC 35.3%, Geely 18.8%, BYD 17.0%); electrive on the Jan. 2026 minimum-price guidance. <https://www.clearytradewatch.com/2024/10/definitive-duties-adopted-by-the-eu-on-chinese-battery-electric-vehicles-to-counteract-subsidies-to-apply-by-october-30/>

- [75] ChinaEVHome, “Hesai leads global ADAS lidar shipments as China suppliers take 95% share” (Yole, May 2026). <https://chinaevhome.com/2026/05/06/hesai-leads-global-adas-lidar-shipments-a-s-china-suppliers-take-95-share/>
- [76] paultan.org, “BYD unveils God’s Eye DiPilot ADAS suite for all models” (Feb. 2025). <https://paultan.org/2025/02/13/byd-unveils-gods-eye-dipilot-adas-suite-to-roll-out-in-all-byd-denza-yangwang-models-sold-in-china/>
- [77] Fortune (citing McKinsey), “Chinese cars frequently more advanced; 20–24-month development cycles; ~30% lower BOM” (Jul. 30, 2026). <https://fortune.com/2026/07/30/byd-germany-dealerships-detroit-china-ev-threat/>
- [78] CSIS, “Rare earth export restrictions: one year later” (Apr. 2025 controls; production stoppages; Oct. 2025 escalation and truce). <https://www.csis.org/analysis/rare-earth-export-restrictions-one-year-later>
- [79] KrASIA, “LLM prices hit rock bottom in China as Alibaba Cloud enters the fray” (May 2024); Fortune Asia on the 97% cuts. <https://kr-asia.com/llm-prices-hit-rock-bottom-in-china-as-alibaba-cloud-enters-the-fray>
- [80] DeepSeek-V3 Technical Report (arXiv:2412.19437); Stanford FSI, “Taking stock of the DeepSeek shock.” <https://cyber.fsi.stanford.edu/publication/taking-stock-deepseek-shock>
- [81] Yahoo Finance, “Nvidia stock plummets, loses record \$589 billion as DeepSeek prompts questions over AI spending” (Jan. 27, 2025). <https://finance.yahoo.com/news/nvidia-stock-plummets-loses-record-589-billion-as-deepseek-prompts-questions-over-ai-spending-135105824.html>
- [82] CNN Business, “DeepSeek discloses R1 RL training cost of \$294,000 in peer-reviewed Nature paper” (Sep. 2025). <https://www.cnn.com/2025/09/19/business/deepseek-ai-training-cost-china-intl>
- [83] Digital Applied, “Chinese AI models Q2 2026 market share report” (release/pricing compendium; secondary source—verify against vendor blogs). <https://www.digitalapplied.com/blog/chinese-ai-models-q2-2026-market-share-report>
- [84] OpenRouter, “The open-weight models that matter, June 2026” (GLM-5.2 #1 open on AA index; ranking of open models). <https://openrouter.ai/blog/insights/the-open-weight-models-that-matter-june-2026/>
- [85] Hugging Face, “State of Open Source, Spring 2026” (Chinese models 41% of downloads; Qwen >113k derivatives). <https://huggingface.co/blog/huggingface/state-of-os-hf-spring-2026>
- [86] Qwen team, “Qwen3” release blog (Apr. 2025: 36T tokens, Apache 2.0, hybrid reasoning). <https://qwenlm.github.io/blog/qwen3/>
- [87] S. Willison, “Kimi K2 Thinking” (Nov. 2025: 1T/32B, INT4 QAT, HLE 44.9 vs GPT-5 41.7). <https://simonwillison.net/2025/Nov/6/kimi-k2-thinking/>
- [88] Tom’s Hardware, “Moonshot releases 2.8-trillion-parameter Kimi K3 as China works around U.S. compute limits” (Jul. 2026); VentureBeat coverage. <https://www.tomshardware.com/tech-industry/artificial-intelligence/moonshot-releases-2-8-trillion-parameter-kimi-k3>
- [89] DeepSeek V4 release documentation and third-party analysis (Apr. 2026: 1.6T/49B active, MIT license, DSA). <https://www.morphllm.com/deepseek-v4>
- [90] Z.ai, GLM-5/GLM-5.2 releases (MIT license; #1 open-weight, AA Intelligence Index v4.1). <https://www.eigent.ai/blog/glm-5-2>
- [91] TechNode, “Baidu open-sources ERNIE 4.5 series” (Jun. 30, 2025, Apache 2.0). <https://technode.com/2025/07/01/baidu-open-sources-ernie-4-5-series-models-including-multi-modal-moe-architecture/>
- [92] MIT Technology Review, “What’s next for Chinese open-source AI” (Z.ai and MiniMax HK IPOs, Jan. 2026). <https://www.technologyreview.com/2026/02/12/1132811/whats-next-for-chinese-open-source-ai/>
- [93] Seoul Economic Daily, “China eyes AI supremacy with tokens at one-twentieth of US” (May 2026; 140T daily token calls). <https://en.sedaily.com/international/2026/05/08/china-eyes-ai-supremacy-with-tokens-at-one-twentieth-of-us>
- [94] Stanford HAI, *AI Index Report 2026* (US–China top-model gap 2.7%; US \$285.9B vs China \$12.4B private investment). <https://hai.stanford.edu/ai-index/2026-ai-index-report>

- [95] AEI, “China has caught up in frontier AI” (Jul. 2026: release lag now “a matter of weeks”; “frontier AI is quickly becoming a commons”). <https://www.aei.org/foreign-and-defense-policy/china-has-caught-up-in-frontier-ai/>
- [96] Data Gravity, “China’s open-weight takeover” (OpenRouter token-share series; Llama <1%; provider shares). <https://www.datagravity.dev/p/chinas-open-weight-takeover>
- [97] M. Casado (a16z) via The Economist and X (Nov. 2025): ~80% of open-stack startups use Chinese models. https://x.com/martin_casado/status/1990462245541982546
- [98] OpenAI, “Introducing gpt-oss” (Aug. 5, 2025). <https://openai.com/index/introducing-gpt-oss/>
- [99] PRC MFA, “Global AI Governance Action Plan” (WAIC, Jul. 2025), full text. https://www.fmprc.gov.cn/mfa_eng/xw/zyxw/202507/t20250729_11679232.html
- [100] Carnegie Endowment, “China’s pivot on global AI” (May 2026: WAICO, UN Group of Friends, BRICS AI centre). <https://carnegieendowment.org/research/2026/05/chinas-pivot-on-global-ai>
- [101] N. Lambert, “On China’s open-source AI trajectory” (Interconnects; State Council “AI+” plan targets). <https://www.interconnects.ai/p/on-chinas-open-source-ai-trajectory>
- [102] MIT Technology Review, “What’s next for Chinese open-source AI” (Liu Zhiyuan quote; academic-credit policy). <https://www.technologyreview.com/2026/02/12/1132811/whats-next-for-chinese-open-source-ai/>
- [103] S. Raschka, “The technical DeepSeek story” (MLA, GRPO, DSA analysis). <https://magazine.sebastianraschka.com/p/technical-deepseek>
- [104] The Register, “DeepSeek reverts to Nvidia after failed Ascend training runs” (Aug. 2025); TrendForce coverage. https://www.theregister.com/2025/08/14/dodgy_huawei_deepseek/
- [105] DeepSeek V3.1 release notes (UE8M0 FP8 scale format “for next-generation domestic chips,” Aug. 2025; widely reported).
- [106] Reuters (exclusive), “China to publish policy to boost RISC-V chip use nationwide” (Mar. 3, 2025; eight agencies incl. MIIT, CAC, MOST). <https://money.usnews.com/investing/news/articles/2025-03-03/exclusive-china-to-publish-policy-to-boost-risc-v-chip-use-nationwide-sources>
- [107] China Briefing, “China’s 15th Five-Year Plan: key insights” (“extraordinary measures” for ICs and basic software). <https://www.china-briefing.com/news/chinas-15th-five-year-plan-key-insights-for-foreign-investors/>
- [108] B. Yungang et al., XiangShan project, RISC-V Summit Europe 2025 proceedings; Kunminghu microarchitecture. <https://riscv-europe.org/summit/2025/media/proceedings/2025-05-14-RISC-V-Summit-Europe-09h30-BAO-slides.pdf>
- [109] Xinhua, “CAS unveils RISC-V achievements at Zhongguancun Forum 2026” (openRuyi, RVA23; 600+/400+ researchers per China Daily). <https://english.news.cn/20260327/3b8302ac28f64dedbfa3defe514ff6e7e/c.html>
- [110] Tom’s Hardware, “Alibaba launches RISC-V based XuanTie C930 server CPU” (Mar. 2025). <https://www.tomshardware.com/pc-components/cpus/alibaba-launches-risc-v-based-xuantie-c930-server-cpu-ai-hpc-chip-ships-this-month-more-designs-to-follow>
- [111] SCMP, “China’s SpacemiT to launch server-class RISC-V processor after raising US\$86 million” (K3, Jan. 2026). <https://www.scmp.com/tech/tech-trends/article/3339953/chinas-spacemit-launch-server-class-risc-v-processor-after-raising-us86-million>
- [112] CSIS, “Sustaining standards leadership: the United States cannot disengage from RISC-V” (membership statistics; 2B+ chips shipped). <https://www.csis.org/analysis/sustaining-standards-leadership-united-states-cannot-disengage-risc-v>
- [113] RISC-V International, “RISC-V Summits” (Summit China 2026: Shenzhen, Oct. 18–20). <https://riscv.org/community/risc-v-summits/>
- [114] RISC-V International Tech Hub: IME/AME/VME task-group ratification plans (IME spec freeze Jun. 2026, ratification target Aug. 2026). <https://riscv.atlassian.net/wiki/spaces/IMEX/pages/598867969/IME+Ratification+Plan>
- [115] EVAS Intelligence RISC-V AI SuperNode, WAIC 2026 (vendor claims: block-quantized FP8 cloud AI chip, 3.2 TB/s ELink). <https://finance.biggo.com/news/8eda6f31-9e18-42f8-9461-c5f142ed8a0d>

- [116] Fortune, “Nvidia and AMD to pay 15% of China chip sale revenue to US government” (Aug. 2025); export-control timeline. <https://fortune.com/2025/08/10/nvidia-amd-chips-h20-mi308-china-sales-revenue-trump-export-license>
- [117] TechCrunch, “China tells its tech companies they can’t buy AI chips from Nvidia” (Sep. 17, 2025); Global Times on the CAC summons (Jul. 31, 2025). <https://techcrunch.com/2025/09/17/china-tells-its-tech-companies-they-cant-buy-ai-chips-from-nvidia/>
- [118] Reuters (exclusive), “China bans foreign AI chips from state-funded data centres” (Nov. 2025); provincial power-tariff discounts for domestic-chip facilities. <https://www.aol.com/articles/exclusive-china-bans-foreign-ai-090709029.html>
- [119] NVIDIA, Q2 FY2026 results (zero H20 sales to China-based customers); TradingKey, “Nvidia China revenue share 26% $\rightarrow$ $\sim$ 5%.” <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2026>
- [120] Axios, “Jensen Huang: ‘China is going to win the AI race’” (Nov. 5, 2025, and same-day walk-back); Hill & Valley remarks (Apr. 2025). <https://www.axios.com/2025/11/05/ai-nvidia-china-race>
- [121] SemiAnalysis, “Huawei Ascend production ramp” (2024: 507k units; 2025: $\sim$ 805k; die bank 2.9M; HBM stockpiles; SMIC 45k $\rightarrow$ 80k wpm). <https://newsletter.semianalysis.com/p/huawei-ascend-production-ramp>
- [122] Huawei Connect 2025 keynote (E. Xu), “Ascend 950/960/970 roadmap; Atlas 950/960 SuperPoD; TaiShan 950; UnifiedBus 2.0” (Sep. 18, 2025). <https://www.huawei.com/en/news/2025/9/hc-xu-keynote-speech>
- [123] Tom’s Hardware, “Cambricon targets 500,000 AI chips in 2026”; FY2025 revenue +453%, first annual profit. <https://www.tomshardware.com/tech-industry/semiconductors/cambricon-targets-500000-ai-chips-in-2026-as-china-accelerates-domestic-hardware-push>
- [124] CNBC, “Moore Threads’ \$1.1B Shanghai IPO surges $\sim$ 400% on debut” (Dec. 2025); Asia Financial on Biren’s HK IPO (Jan. 2026). <https://www.cnbc.com/2025/12/05/china-nvidia-moore-threads-trading-debut-1-billion-listing-ipo-shanghai-gpu-enflame-biren.html>
- [125] technology.org, “Baidu lights up 30,000-strong Kunlun P800 chip cluster” (Apr. 2025). <https://www.technology.org/2025/04/25/baidu-lights-up-30000-strong-kunlun-chip-cluster-aims-to-train-deepseek-like-ai-models/>
- [126] TrendForce, “Alibaba T-Head PPU lands China Unicom deal, reportedly matches Nvidia H20” (Sep. 2025); CNBC on the Zhenwu 10k-chip cluster (Apr. 2026). <https://www.trendforce.com/news/2025/09/17/news-alibaba-t-head-ppu-lands-china-unicom-deal-reportedly-matches-nvidia-h20-performance/>
- [127] Tom’s Hardware, “China certifies nine domestic AI chips for government procurement” (May 2026). <https://www.tomshardware.com/tech-industry/semiconductors/china-certifies-nine-domestic-ai-chips-for-government-procurement>
- [128] TradingKey, “Domestic AI chip shipments $\sim$ 1.65M units, 41% of China market in 2025; Huawei $\sim$ 812k.” <https://www.tradingkey.com/analysis/stocks/us-stocks/261889036-nvidia-china-revenue-airforceone-h20-zero-sales-huang-trump-visit-tradingkey>
- [129] Xinhua (MIIT data), “China’s intelligent computing capacity reaches 2,185 EFLOPS, utilization 71.4%” (Jul. 2026). <https://english.news.cn/20260720/c5ebd0bd866b46dcb221f59a47a38d82/c.html>
- [130] MIIT, Computing Power Interconnection Action Plan (2025) and “AI+” implementation documents on unified AI-chip interconnection ecosystems (Chinese-language primary sources; see also Pandaily summary). <https://pandaily.com/miit-optoelectronic-chip-rd-ai-communications-2026>
- [131] CRN Asia, “Huawei’s open source gambit: SuperPoD reference architecture and UnifiedBus 2.0 released openly” (Sep. 2025). <https://www.crnasia.com/news/2025/artificial-intelligence/huawei-s-open-source-gambit-chinese-tech-giant-makes-superpo>
- [132] Xinhua, “Huawei fully open-sources CANN” (Aug. 7, 2025). <https://english.news.cn/20250807/35025b0c1366473f976e29e0c5adf88f/c.html>
- [133] Network World, “Server memory prices could double by 2026 as AI demand strains supply” (HBM $\sim$ 3 $\times$ wafer capacity per bit; Grace LPDDR5X shift). <https://www.networkworld.com/article/4093752/server-memory-prices-could-double-by-2026-as-ai-demand-strains-supply.html>

- [134] RIKEN Center for Computational Science, feasibility study of memory-bandwidth balance for 2.8T-parameter sparse-MoE pretraining at 20 PF(FP8) per accelerator: HBM4 capacity wall at <256 accelerators; LPDDR6X substitution at 2.3 TB/socket, 5.0 TB/s knee, 67% vs. 89% MFU ceilings (2026, companion report to this volume).
- [135] wccftech, “CXMT’s LPDDR6 clears R&D verification at 12.8 Gbps” (Aug. 2026); Gizmochina on 2H26 mass-production targeting. <https://wccftech.com/cxmts-lpddr6-clears-rd-verification-at-12-8-gbps-inching-closer-to-a-full-mass-production-push-report/>
- [136] SK hynix newsroom, “Mass production of SOCAMM2 192GB” (Apr. 2026); TechPowerUp on Qualcomm/AMD adoption. <https://news.skhynix.com/en/mass-production-socamm2-192gb/>
- [137] SemiAnalysis (Apr. 2025), CloudMatrix 384 analysis: $\sim 2\times$ GB200 NVL72 BF16 throughput at $\sim 3.9\times$ power. <https://techstartups.com/2025/04/28/huawei-ascend-910c-system-reportedly-outperforms-nvidias-h100-in-key-metrics/>
- [138] Institute for Progress, “The B30A decision” (2026 output projections: US 6.89M vs China $\sim 167k$ B300-equivalents). <https://ifp.org/the-b30a-decision/>
- [139] Design & Reuse, “SMIC achieves 5nm production on N+3 node without EUV tools” (TechInsights teardown of Kirin 9030, Dec. 2025). <https://www.design-reuse.com/news/202529830-chinese-smic-achieves-5-nm-production-on-n-3-node-without-euv-tools/>
- [140] SupplyICs market intelligence, “China logic fabs: yield and geopolitical risks 2026” (N+2 yields est. 40–55%; treat as estimates). <https://supplyics.com/insights/market-intelligence/china-logic-fabs-yield-geopolitical-risks-2026/>
- [141] TrendForce, “SMIC posts record \$9.3B in 2025 sales” (Feb. 2026; utilization 93.5%; capex \$8.1B). <https://www.trendforce.com/news/2026/02/11/news-smic-posts-record-9-3b-in-2025-sales-7nm-yields-reportedly-weigh-on-margins/>
- [142] SemiWiki/Digitimes, “Huawei’s quiet IDM ambition: 11 shadow fabs across China”; DSET, “Uncovering Huawei’s shadow network.” <https://dset.tw/en/research/uncovering-huawei-shadow-network/>
- [143] TrendForce, “China 37% of global 22–40nm output in 2026, 42% by 2028 (SEMI China)”; “Chinese foundries drive $\sim 70\%$ of 12-inch mature-node expansion” (May 2026). <https://www.trendforce.com/presscenter/news/20260507-13036.html>
- [144] TrendForce, “Chinese carmakers reportedly aim for 100% domestic chips in vehicles by 2026–27” (Jun. 2025). <https://www.trendforce.com/news/2025/06/18/news-chinese-carmakers-reportedly-aim-for-100-domestic-chips-in-vehicles-by-2026/>
- [145] wccftech, “CXMT races to increase capacity to 600,000 wafers per month” (Jul. 2026; current 200–300k wpm). <https://wccftech.com/cxmts-64gb-server-ddr5-memory-chip-is-now-pricier-than-samsungs-as-it-races-to-increase-capacity-to-600000-wafers-per-month/>
- [146] Tom’s Hardware (Citrini Research), “CXMT close to matching Micron’s memory capacity in 2026” (350k wpm model; China $\sim 1.4M$ wpm DRAM by 2030). <https://www.tomshardware.com/pc-components/dram/cxmt-close-to-matching-microns-memory-capacity-in-2026-research-claims-would-put-china-on-track-to-become-worlds-second-largest-dram-producer>
- [147] wccftech, “CXMT debuts DDR5-8000 and LPDDR5X-10667” (Nov. 2025); TrendForce on 80% DDR5 yield (Dec. 2024). <https://wccftech.com/cxmt-debuts-domestically-produced-ddr5-memory-8000-mtps-lpddr5-10667-mtps/>
- [148] CNBC, “CXMT closes 466% higher in Shanghai debut” (Jul. 27, 2026; \$8B raised; largest A-share company; prospectus disclosed no commercial HBM). <https://www.cnbc.com/2026/07/27/cxmt-china-market-debut-chipmaker-ipo.html>
- [149] ChinaTalk, “Mapping China’s HBM advancement” (CXMT HBM roadmap; ~ 4 -year gap to leaders). <https://www.chinatalk.media/p/mapping-chinas-hbm-advancement>
- [150] Tom’s Hardware, “Chinese memory maker gets \$2.4 billion to build HBM packaging for AI processors” (Innotron TSV facility); TrendForce on Tongfu/JCET expansions. <https://www.tomshardware.com/pc-components/dram/chinese-memory-maker-gets-dollar24-billion-to-build-hbm-for-ai-processors-shanghai-packaging-facility-to-open-in-2026>
- [151] TrendForce, “Conventional DRAM up 55–60% QoQ in 1Q26” (Jan. 2026); Network World on the AI-driven shortage. <https://www.trendforce.com/presscenter/news/20260105-12860.html>

- [152] wccftech, “CXMT’s 64GB server DDR5 is now pricier than Samsung’s” (Jul. 2026). <https://wccftech.com/cxmts-64gb-server-ddr5-memory-chip-is-now-pricier-than-samsungs-as-i-t-races-to-increase-capacity-to-600000-wafers-per-month/>
- [153] Digitimes, “YMTC Xtacking 4.0 powers 267-layer NAND”; TechPowerUp/Guru3D on the 294-layer 5th generation. <https://www.digitimes.com/news/a20250924PD216/ymtc-nand-nand-flash-production-memory.html>
- [154] TrendForce, “Three Chinese chip toolmakers enter global top 20 as Naura rises to fifth by sales” (Feb. 2026; domestic tool share 20–30%). <https://www.trendforce.com/news/2026/02/02/news-three-chinese-chip-toolmakers-reportedly-enter-global-top-20-as-naura-rises-to-fifth-by-sales/>
- [155] BigGo Finance, “Naura 2025 results: revenue ~RMB 39.4B (+31%)” (Apr. 2026). <https://finance.biggo.com/news/9r28m50BvthpMgHBYd-l>
- [156] Yahoo Finance, “SiCarrier targets \$2.8bn initial raise at ~\$11B valuation”; TrendForce on its SAQP-DUV 5nm-class patent. <https://www.trendforce.com/news/2025/11/10/news-decoding-chinas-lithography-push-to-challenge-asml-from-sicarrier-to-alternative-euv-paths/>
- [157] TrendForce, “Decoding China’s lithography push to challenge ASML” (Nov. 2025: SMEE 90nm shipping, 28nm immersion in development, Yuliangsheng trials; domestic sub-10nm litho ~2030); eeNews on LDP-EUV reports (speculative). <https://www.trendforce.com/news/2025/11/10/news-decoding-chinas-lithography-push-to-challenge-asml-from-sicarrier-to-alternative-euv-paths/>
- [158] Tom’s Hardware, “White House lifts chip design software export ban on China” (Jul. 2025); TrendForce on Empyrean’s full-process memory EDA platform (Aug. 2025). <https://www.tomshardware.com/tech-industry/semiconductors/white-house-lifts-chip-design-export-ban-on-china-in-exchange-for-rare-earth-materials-compromise-export-licences-for-eda-software-sales-no-longer-required>
- [159] TechWire Asia, “China semiconductor self-sufficiency: wafer target 2026” (300mm wafer share 3% (2020) → ~28–32% (2026); Eswin pricing). <https://techwireasia.com/2026/05/china-semiconductor-self-sufficiency-wafer-target-2026/>
- [160] Asia Society Policy Institute, “China’s Big Fund 3.0: Xi’s boldest gamble yet for chip supremacy” (RMB 344B, May 2024, horizon 2039). <https://asiasociety.org/policy-institute/chinas-big-fund-30-xis-boldest-gamble-yet-chip-supremacy>
- [161] SCMP (SEMI), “China to remain world’s biggest buyer of chipmaking equipment through 2027”; SEMI 2026 equipment forecasts. <https://www.scmp.com/tech/article/3336933/china-remain-world-s-biggest-buyer-chipmaking-equipment-through-2027-semi>
- [162] USCC, “Made in China 2025: Evaluating China’s Performance” (Nov. 2025); TechWire Asia on 2026 self-sufficiency measures. https://www.uscc.gov/sites/default/files/2025-11/Made_in_China_2025--Evaluating_Chinas_Performance.pdf
- [163] US BIS, “Commerce strengthens export controls...” (Dec. 2, 2024: 140 entities; HBM and SME controls). <https://www.bis.gov/press-release/commerce-strengthens-export-controls-restrict-chinas-capability-produce-advanced-semiconductors-military>
- [164] Built In, “Timeline: US AI-chip policy toward China 2025–2026” (AI Diffusion Rule imposed and rescinded; H20/H200 reversals). <https://builtin.com/articles/trump-lifts-ai-chip-ban-china-nvidia>
- [165] CNBC, “US tech execs smuggled Nvidia chips to China, prosecutors say” (\$2.5B scheme, Mar. 2026). <https://www.cnbc.com/2026/03/19/us-tech-execs-smuggled-nvidia-chips-to-china-prosecutors-say.html>
- [166] OpenSourceForU, “Alibaba’s Qwen crosses one billion downloads, eclipsing Meta’s Llama” (Jul. 2026); Xinhua on the 700M milestone (Jan. 2026). <https://www.opensourceforu.com/2026/07/alibaba-qwen-crosses-one-billion-downloads-eclipsing-metas-llama/>
- [167] State of Open Source AI report v1.0.1 (Jul. 2026): DeepSeek 26,000+ enterprise deployments; Chinese models ~45% of OpenRouter weekly traffic (Apr. 2026). <https://stateofopensource.ai/>
- [168] Menlo Ventures 2025 enterprise LLM survey: open-weight share of enterprise workloads 19% (2024) → 11% (2025). <https://www.luizneto.ai/enterprise-open-weight-models-2026/>
- [169] Unsloth AI, “Run DeepSeek-R1 Dynamic 1.58-bit” (Jan. 2025: 720GB → 131GB; dynamic vs naive quantization results). <https://unsloth.ai/blog/deepseekr1-dynamic>

- [170] Unsloth documentation: Dynamic 2.0 GGUFs methodology; Kimi K3 (594GB–1.5TB), GLM-5 (176–241GB), K2 Thinking 1-bit (245GB) quantizations with retention figures.
<https://unsloth.ai/docs/basics/unsloth-dynamic-2.0-ggufs>
- [171] Y Combinator company profile: Unsloth AI (~10M monthly model downloads).
<https://www.ycombinator.com/companies/unsloth-ai>
- [172] Lin et al., “AWQ: Activation-aware Weight Quantization” (MLSys 2024 Best Paper); Frantar et al., “GPTQ” (ICLR 2023). <https://github.com/mit-han-lab/llm-awq>
- [173] vLLM documentation, supported quantization methods (AWQ, GPTQ, FP8, NVFP4, GGUF, bitsandbytes). <https://docs.vllm.ai/en/latest/features/quantization/>
- [174] KTransformers (Tsinghua kvcache-ai), SOSP 2025: DeepSeek-R1/V3 on one RTX 4090 + dual Xeon, 286 tok/s prefill via AMX.
https://github.com/kvcache-ai/ktransformers/blob/main/doc/en/DeepseekR1_V3_tutorial.md
- [175] Digital Spaceport, “How to run DeepSeek-R1 671B fully locally on a \$2,000 EPYC rig” (Jan. 2025).
<https://digitalspaceport.com/how-to-run-deepseek-r1-671b-fully-locally-on-2000-epyc-rig/>
- [176] MacRumors on M3 Ultra running R1 at 17–18 tok/s (Mar. 2025); A. Hannun on 2×M3 Ultra running Kimi K2/K2.5; Creative Strategies, “Running a 1T-parameter model on a \$40k Mac Studio cluster” (Dec. 2025). <https://creativestrategies.com/research/running-a-1t-parameter-model-on-a-40k-mac-studio-cluster/>
- [177] AIMultiple, “DGX Spark alternatives” benchmark roundup (Spark, Strix Halo, Mac Studio, multi-3090 gpt-oss-120b figures); Atlantic Council on fall-2026 local-AI laptops.
<https://aimultiple.com/dgx-spark-alternatives>
- [178] llama.cpp PR #15077, ‘-n-cpu-moe’ MoE expert CPU-offload (Aug. 2025); day-0 native MXFP4 support. <https://github.com/ggml-org/llama.cpp/pull/15077>
- [179] G. Appenzeller (a16z), “LLMflation: LLM inference cost is falling 10× per year” (Nov. 2024).
<https://a16z.com/llmflation-llm-inference-cost/>
- [180] OECD/McKinsey self-hosting break-even figures via luizneto.ai enterprise analysis (2026).
<https://www.luizneto.ai/enterprise-open-weight-models-2026/>
- [181] Atlantic Council, “The best AI you can own is Chinese—the West needs to close that gap quickly” (2026; Microsoft’s DeepSeek V4 Copilot evaluation; Western app bans vs self-hosted weights).
<https://www.atlanticcouncil.org/blogs/the-best-ai-you-can-own-is-chinese-the-west-needs-to-close-that-gap-quickly/>
- [182] TOP500, “LineShine debuts at No. 1 as TOP500 enters new global exascale era” (Jun. 24, 2026: HPL 2.198 EF, HPCG 22.00 PF, 42.2 MW, 52.07 GF/W).
<https://top500.org/news/lineshine-debuts-no-1-top500-enters-new-global-exascale-era/>
- [183] J. Dongarra, “Report on the Chinese LineShine System,” Univ. of Tennessee Tech Report ICL-UT-26-01 (Jun. 23, 2026).
<https://glennklockwood.com/garden/attachments/dongarra-lineshine-0623.pdf>
- [184] J. Dongarra, T. Hoefler, and S. Matsuoka, “Do We Still Need GPUs? Where an Evolved CPU Can Replace Them in AI and Scientific Computing—and Where It Cannot Yet,” *Communications of the ACM*, in revision (2026).
- [185] J. Dongarra, T. Hoefler, and S. Matsuoka, “Do We Still Need GPUs? Rethinking AI and Scientific Computing on Accelerated CPUs,” full-length companion paper (2026): LX2 vs GPU-fleet analysis of 1T-MoE serving and training; the design-sequence/allocation argument.
- [186] G. K. Lockwood, ISC’26 recap and LX2 processor notes (microarchitecture, HBM2e speculation, software stack analysis). <https://blog.glennklockwood.com/2026/07/isc26-recap.html>
- [187] Global Times on LineShine (Jun. 2026): “China’s first domestically developed HBM” claim; CAS/analyst commentary. <https://www.globaltimes.cn/page/202606/1364287.shtml>
- [188] SCMP, “Return to top: China’s LineShine beats US’ El Capitan in TOP500 supercomputer rankings” (Jun. 24, 2026; Dongarra quotes on export controls).
<https://www.scmp.com/news/china/science/article/3358107/return-top-chinas-lineshine-beats-u-s-el-capitan-top500-supercomputer-rankings>

- [189] RISC-V International, “How NVIDIA shipped one billion RISC-V cores in 2024” (Feb. 2025; 10–40 cores per chip since 2016). <https://riscv.org/blog/how-nvidia-shipped-one-billion-risc-v-cores-in-2024/>
- [190] Tom’s Hardware, “NVIDIA’s CUDA platform now supports RISC-V” (announced at RISC-V Summit China, Jul. 18, 2025). <https://www.tomshardware.com/pc-components/gpus/nvidias-cuda-platform-now-supports-risc-v-support-brings-open-source-instruction-set-to-ai-platforms-joining-x86-and-arm>
- [191] The Register, “Nvidia’s CUDA-on-RISC-V as a China hedge” (Jul. 21, 2025); SCMP on the boost to China’s self-sufficiency drive. https://www.theregister.com/2025/07/21/nvidia_cuda_riscv
- [192] SiFive press release, “SiFive to integrate NVIDIA NVLink Fusion for next-gen AI datacenter chips” (Jan. 15, 2026). <https://www.sifive.com/press/sifive-nvidia-nvlinkfusion-datacenter>
- [193] HPCwire, “RISC-V CEO Calista Redmond resigns” (Dec. 2024); subsequently NVIDIA VP of global AI initiatives. <https://www.hpcwire.com/off-the-wire/risc-v-ceo-calista-redmond-resigns-after-5years-of-progress/>
- [194] Data Center Dynamics, “Google deploys SiFive’s Intelligence X280 processor for AI workloads” (Sep. 2022); Google Open Source Blog, “Android and RISC-V” (Oct. 2023); RISE founding membership (May 2023). <https://www.datacenterdynamics.com/en/news/google-deploys-sifives-intelligence-x280-processor-for-ai-workloads/>
- [195] Meta AI blog, MTIA v1/v2 (RISC-V processing elements); Data Center Dynamics, “Meta unveils next four generations of its MTIA chip” (Mar. 11, 2026). <https://www.datacenterdynamics.com/en/news/meta-unveils-next-four-generations-of-its-mtia-chip/>
- [196] TechRadar, “Microsoft unveils Cobalt 200” (Nov. 2025) — Azure silicon remains Arm-based; no public Microsoft RISC-V commitment as of mid-2026. <https://www.techradar.com/pro/microsoft-unveils-its-next-generation-arm-based-cpu-cobalt-200-looks-to-unlock-even-more-azure-power>
- [197] Bloomberg, “Lawmakers urge export controls on RISC-V” (Nov. 2023); Reuters on Commerce review (Apr. 2024); Sen. Warner letter to BIS (Sep. 30, 2025). No restriction enacted as of mid-2026. https://www.warner.senate.gov/public/_cache/files/a/b/ab558775-b187-434e-b70d-040a63aea6f7/1AD7AF181B29E0603B1023A6FABE01893C9F53BFCA2ACAA82B1535CEBE2FC7AC.9.30.25-risc-v-letter-to-bis.pdf
- [198] SHD Group 2026 forecast (Jun. 2026): 35.9B RISC-V SoC units and \$318B revenue by 2031, AI accelerators ~70% of revenue; RISC-V International 25%-penetration announcement (Oct. 2025). <https://riscv.org/blog/shd-forecast-2026/>
- [199] Omdia, “RISC-V adoption will be accelerated by AI” (May 2024: ~25% of processors by 2030). <https://omdia.tech.informa.com/pr/2024/may/risc-v-adoption-will-be-accelerated-by-ai-according-to-new-omdia-research>
- [200] DigiTimes, “CXMT delivers HBM3 samples to Huawei; mass production slated 2026” (Oct. 2025); TechPowerUp on the 20%-of-capacity HBM3 line report. <https://www.digitimes.com/news/a20251023PD208/cxmt-hbm3-2026-dram-huawei.html>
- [201] TrendForce, “Huawei and XMC collaborate to develop HBM” (Jul. 2024); Tom’s Hardware on Wuhan Xinxin’s ~3,000 wpm HBM line; SemiAnalysis, “Scaling the Memory Wall.” <https://newsletter.semianalysis.com/p/scaling-the-memory-wall-the-rise-and-roadmap-of-hbm>
- [202] As of mid-2026 no public teardown confirms Chinese-fabbed DRAM in a shipping accelerator; shipping Ascend 910C teardowns show Samsung/SK hynix stacks (SemiAnalysis; prediction-market tracking). <https://newsletter.semianalysis.com/p/huawei-ascend-production-ramp>
- [203] The Substrate, “Where will China get its compute?” (domestic HBM stack estimates, 2.5–18.6M range for 2026). <https://www.the-substrate.net/p/where-will-china-get-its-compute>
- [204] Tom’s Hardware, “Chinese semiconductor industry gears up for domestic HBM3 production by end of 2026” (Naura, Maxwell, U-Precision toolchain). <https://www.tomshardware.com/pc-components/dram/chinese-semiconductor-industry-gears-up-for-domestic-hbm3-production-by-the-end-of-2026-cxmt-to-produce-chips-while-naura-maxwell-and-u-preseason-design-tools-for-assembly>
- [205] Forbes/CNBC/Bloomberg, “OpenAI valuation reaches \$852 billion after \$122B round” (Mar. 31, 2026); prior \$500B secondary (Oct. 2025). <https://www.forbes.com/sites/antoniopequenoiv/2026/03/31/openai-valuation-reaches-852-billion-after-massive-funding-round/>

- [206] WSJ via Yahoo Finance, “OpenAI lifts planned compute spending” (Jul. 22, 2026: 2025 revenue \$13.1B; ~\$2B/month early 2026). <https://finance.yahoo.com/technology/ai/articles/openai-lifts-planned-compute-spending-144917731.html>
- [207] WSJ/analyst projections: ~\$14B 2026 loss; ~\$115B cumulative losses through 2029; profitability targeted 2029–30. <https://valueaddvc.com/blog/openai-revenue-2026-25b-arr-14b-loss-and-why-profitability-waits-until-2030>
- [208] Altman’s \$1.4T commitments figure (Nov. 2025); WSJ: planned compute spend raised to \$750B through 2030, \$1.15T agreements through 2035; Oracle \$300B, AWS ~\$100B, Stargate \$500B/10GW, AMD 6GW + warrants, Broadcom 10GW custom. <https://finance.yahoo.com/technology/ai/articles/openai-lifts-planned-compute-spending-144917731.html>
- [209] CNBC/Forbes: confidential S-1 filed Jun. 8, 2026, targeting up to ~\$1T; reports of delay to 2027 (Jun. 25, 2026). <https://www.forbes.com/sites/aliciapark/2026/06/25/openai-considers-delaying-ipo-to-2027-after-spacexs-rocky-debut-report-says/>
- [210] S. Altman (Aug. 2025): AI investors “overexcited”; “someone will lose a phenomenal amount of money.” <https://www.tipranks.com/news/sam-altman-warns-ai-stocks-may-be-in-a-dot-com-style-bubble>
- [211] TechCrunch, “Sam Altman says ‘Enough’ to questions about OpenAI’s revenue” (BG2 podcast, Nov. 2025). <https://techcrunch.com/2025/11/02/sam-altman-says-enough-to-questions-about-openais-revenue/>
- [212] Anthropic, “Series H” (May 28, 2026: \$65B at \$965B post-money); CNBC coverage. <https://www.anthropic.com/news/series-h>
- [213] Bloomberg (Jan. 2026: >\$9B run-rate); Digitimes/TechCrunch on 2028 breakeven projection and ~14× lower burn than OpenAI; Forbes on inference gross margins passing 70% (Jul. 2026); Google 1GW+ TPU deal (Oct. 2025). <https://www.forbes.com/sites/petercohan/2026/07/28/as-token-costs-plunge-enterprise-ai-providers-face-a-new-margin-squeeze/>
- [214] NVIDIA FY2026 results (Feb. 2026: revenue \$215.9B, data center \$193.7B); \$500B Blackwell+Rubin bookings visibility (GTC Washington, Oct. 2025). <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-fourth-quarter-and-fiscal-2026>
- [215] CNBC/Axios, “Nvidia becomes first \$5 trillion company” (Oct. 29, 2025). <https://www.axios.com/2025/10/29/nvidia-5-trillion>
- [216] CNBC, “Nvidia’s top two mystery customers made up 39% of Q2 revenue” (10-Q, Aug. 2025). <https://www.cnbc.com/2025/08/28/nvidias-top-two-mystery-customers-made-up-39percent-of-q2-revenue-.html>
- [217] WSJ/CNBC, “Nvidia and OpenAI in talks for up to \$250 billion AI backstop” (Jul. 26–27, 2026); TechTimes: “proves debt markets said no”; \$100B LOI restructured to ~\$30B equity (Feb. 2026). <https://www.cnbc.com/2026/07/27/nvidia-and-openai-in-talks-for-up-to-250-billion-dollar-ai-backstop.html>
- [218] Bloomberg/AP, “SoftBank sells entire Nvidia stake for \$5.83B” (Nov. 11, 2025). <https://finance.yahoo.com/news/softbank-dumps-entire-nvidia-portfolio-145532866.html>
- [219] MarketMinute, “SoftBank shares retreat as Son’s \$40B OpenAI gamble strains balance sheet” (Mar. 27, 2026: –45% from peak; Jefferies downgrade; LTV ceiling; ~\$50B refinancing wall). <https://markets.financialcontent.com/stocks/article/marketminute-2026-3-27-softbank-shares-retreat-as-masayoshi-sons-40-billion-openai-gamble-strains-balance-sheet>
- [220] CNBC, “Masayoshi Son says he was ‘crying’ over forced Nvidia sale; ‘all in’ on OpenAI” (Dec. 2, 2025). <https://www.cnbc.com/2025/12/02/softbank-masayoshi-son-crying-nvidia-stake-ai-bet-artificial-intelligence-bubble-vision-fund.html>
- [221] Introl, “Hyperscaler capex \$600B+ in 2026”; Forbes (Jun. 2026): \$700–900B range, capital intensity 45–57% of revenue. <https://introl.com/blog/hyperscaler-capex-600b-2026-ai-infrastructure-debt-january-2026>
- [222] McKinsey, “The cost of compute: a \$7 trillion race to scale data centers” (Apr. 2025: \$6.7T by 2030). <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/the-cost-of-compute-a-7-trillion-dollar-race-to-scale-data-centers>

- [223] Axios, “The AI debt machine” (Jul. 27, 2026: \$489B AI-linked issuance YTD; \$302B by five majors); Nikkei study of \$1.65T hidden obligations.
<https://www.axios.com/2026/07/27/debt-data-center-oracle-goole>
- [224] Global Data Center Hub, “Meta and Blue Owl’s \$27B SPV” (Oct. 2025); Meta \$30B bond, largest of 2025. <https://www.globaldatacenterhub.com/p/meta-blue-owls-27b-bet-is-this-the>
- [225] Axios/CNBC: Oracle 5-yr CDS from 129bps (Feb. 2026) to record 212bps (Jul. 2026) on AI debt load. <https://www.axios.com/2026/07/27/debt-data-center-oracle-goole>
- [226] CNBC, “Michael Burry accuses AI hyperscalers of artificially boosting earnings” (Nov. 11, 2025: \$176B understated depreciation 2026–28, earnings overstated ~27% by 2028).
<https://www.cnbc.com/2025/11/11/big-short-investor-michael-burry-accuses-ai-hyperscalers-of-artificially-boosting-earnings.html>
- [227] Tom’s Hardware/DCD, Nadella on BG2 (Nov. 2025): constraint is “power, not chips”; GPUs “sitting in inventory that I can’t plug in.” <https://www.datacenterdynamics.com/en/news/microsoft-has-ai-gpus-sitting-in-inventory-because-it-lacks-the-power-necessary-to-install-them/>
- [228] CNBC / TechTimes: CoreWeave Q1 2026 revenue \$2.08B (+112%); backlog \$99.4B; capex guide \$31–35B; Q1 FCF −\$4.7B; ~\$21–25B debt; 5-yr CDS ~855 bps (29 Jul. 2026, ~50% implied default probability); stock −40% over 12 months; Meta contract \$21B through 2032.
<https://www.techtimes.com/articles/322222/20260730/coreweave-cds-hits-50-default-odds-bond-market-calls-time-gpu-debt-spiral.htm>
- [229] Axios (Jul. 27, 2026): Alphabet’s first negative-FCF quarter since its 2004 IPO, on AI capex.
<https://www.axios.com/2026/07/27/debt-data-center-oracle-goole>
- [230] BEA third-estimate analyses: AI/data-center capex ~74% of US Q1 2026 GDP growth; Furman: ex-data-centers H1 2025 growth ~0.1%.
<https://www.tftc.io/ai-capex-gdp-growth-q1-2026-bea-third-estimate/>
- [231] MIT Project NANDA, “The GenAI Divide” (Aug. 2025): ~95% of enterprise GenAI pilots produced zero P&L impact. <https://www.forbes.com/sites/jasonsnyder/2025/08/26/mit-finds-95-of-genai-pilots-fail-because-companies-avoid-friction/>
- [232] Bain & Company, 6th Global Technology Report (Sep. 2025): \$2T annual AI revenue needed by 2030; ~\$800B shortfall.
[https://www.bain.com/about/media-center/press-releases/20252/\\$2-trillion-in-new-revenue-needed-to-fund-ais-scaling-trend---bain--companys-6th-annual-global-technology-report/](https://www.bain.com/about/media-center/press-releases/20252/$2-trillion-in-new-revenue-needed-to-fund-ais-scaling-trend---bain--companys-6th-annual-global-technology-report/)
- [233] Forbes, “The AI capex-to-revenue gap is widening” (Jun. 2, 2026: divergence rate 46% vs 2001 telecom’s 32%). <https://www.forbes.com/sites/jasonkirsch/2026/06/02/the-ai-capex-to-revenue-gap-is-widening---and-markets-are-starting-to-notice/>
- [234] CNBC, “IMF and Bank of England join growing chorus warning of an AI bubble” (Oct. 8–9, 2025).
<https://www.cnbc.com/2025/10/09/imf-and-bank-of-england-join-growing-chorus-warning-of-a-n-ai-bubble.html>
- [235] Bloomberg graphics, “The AI circular deals” (2026, ~\$750B+ interlocking arrangements); Bernstein on circularity concerns. <https://www.bloomberg.com/graphics/2026-ai-circular-deals/>
- [236] Forbes, “As token costs plunge, enterprise AI providers face a new margin squeeze” (Jul. 28, 2026: Chinese suppliers ~46% global model share; Telnix \$100k/day → \$100/day case).
<https://www.forbes.com/sites/petercohan/2026/07/28/as-token-costs-plunge-enterprise-ai-providers-face-a-new-margin-squeeze/>
- [237] Intellectia/CNN: \$1.3T semiconductor selloff (Jun. 4–9, 2026) on Broadcom AI-networking miss.
<https://intellectia.ai/blog/semiconductor-stocks-selloff-june-2026>
- [238] CNBC/NBC (Jul. 28–29, 2026): memory-led chip rout, SK hynix −15%, >\$1T lost, Nasdaq-100 in correction; Samsung/SK warnings of shortages to 2027+.
<https://www.cnbc.com/2026/07/29/chip-selloff-sk-hynix-samsung-softbank.html>
- [239] Epoch AI, “OpenAI Stargate: where the US sites stand” (Apr. 2026: ~0.2–0.3GW operational of 10GW pledged); partner-dispute reporting.
<https://epoch.ai/publications/openai-stargate-where-the-us-sites-stand>

- [240] G. DiPippo et al., “Red Ink: Estimating Chinese Industrial Policy Spending in Comparative Perspective,” CSIS (2022): ~\$248B/yr, 1.7% of GDP. https://csis-website-prod.s3.amazonaws.com/s3fs-public/publication/220523_DiPippo_Red_Ink.pdf
- [241] Leeham News, “COMAC struggled in 2025” (Jan. 2026: ~13 deliveries vs 75 plan); 36Kr on >40% foreign core content and the 2025 LEAP license suspension/restoration. <https://leehamnews.com/2026/01/15/comac-struggled-in-2025-2026-wont-be-much-better/>
- [242] 36Kr, “CJ-1000A: verification flight testing 2025, certification targeted end-2025”; IBA assessment: “years from readiness.” <https://www.iba.aero/resources/articles/comac-aircraft-programmes-status-outlook/>
- [243] International Federation of Airworthiness / EASA Director Guillermet: C919 European certification 3–6 years out. <https://ifairworthy.com/comacs-c919-will-not-be-certified-in-europe-for-3-to-6-years-according-to-easa-director/>
- [244] Tianxia Gongchang Research, “China semiconductor equipment 2026” (localization by segment: litho <5%, etch 30–35%, ArF resist <10%, valves <10%); TrendForce on SMEE 28nm retraction. <https://faxiangongchang.com/en/reports/china-semiconductor-equipment-2026>
- [245] 36Kr, “High-end scientific instruments: domestic share ~17%, imports \$17B” (electron microscopes >90% imported). <https://eu.36kr.com/en/p/3363026278434564>
- [246] USCC MIC2025 evaluation: high-end CNC ~8% domestic vs 80% target; robots 52% vs 70%; US Commerce China machine-tools market intelligence. https://www.uscc.gov/sites/default/files/2025-11/Made_in_China_2025--Evaluating_Chinas_Performance.pdf
- [247] Jamestown Foundation, “New gains in PRC robotics” (Harmonic Drive ~85% global share 2023; Nabtesco ~60% RV; Leaderdrive 15% global; Chinese reducers in Tesla Optimus/Unitree). <https://jamestown.org/new-gains-in-prc-robotics-software-hardware/>
- [248] Counterpoint via Silicon UK: HarmonyOS 17% of China smartphone OS, ahead of iOS six straight quarters; Caixin on 8M developers; Kingsoft WPS ~600M MAU. <https://www.silicon.co.uk/mobility/mobile-os/huawei-apple-china-2-626789>
- [249] Pharmasource, “China biopharma out-licensing surges to record \$137.7B in 2025”; Caixin on first-in-class pipeline (~24%); The Cancer Letter on the ODAC 14–1 sintilimab vote (2022). <https://pharmasource.global/content/china-biopharma-out-licensing-surges-to-record-137-7b-in-2025-2026-on-pace-to-break-it-again/>
- [250] CSIS (Tintoré Vicent), “Industrial Policy Without Illusions: What China’s EV Boom and Aviation Struggles Teach” (2025). <https://www.csis.org/blogs/trustee-china-hand/industrial-policy-without-illusions-what-chinas-ev-boom-and-aviation>
- [251] K. Chan, “China’s Technology Long Game” (High Capacity); USCC testimony (Feb. 2025) on recurring five-year-plan sectors and the moving-frontier problem. <https://www.highcapacity.org/p/chinas-tech-long-game>
- [252] MacroPolo Global AI Talent Tracker 2.0 (Mar. 2024; NeurIPS 2019 vs 2022 authors, top-tier ≈ top 20%): Chinese-undergrad share of top-tier researchers 29%→47%; share working in the US 59%→42%, in China 11%→28% (MIT Technology Review, 27 Mar. 2024); MacroPolo ceased operations in March 2025 and later “Tracker 3.0” figures circulating online are not the Institute’s; Carnegie Endowment (Dec. 2025): 87 of 100 tracked 2019-cohort Chinese-origin researchers remained in the US; RIETI on the returnee wave and the K-visa (Oct. 2025). <https://macropolo.org/interactive/digital-projects/the-global-ai-talent-tracker/>
- [253] C. Zhang et al., “100 Days After DeepSeek-R1: A Survey on Replication Studies...” arXiv:2505.00551 (2025). <https://arxiv.org/abs/2505.00551>
- [254] Epoch AI, “LLM responses to benchmark questions are getting longer over time” (Apr. 2025: reasoning models ~5×/yr; non-reasoning 2.2×/yr). <https://epoch.ai/data-insights/output-length>
- [255] Epoch AI, “LLM inference prices have fallen rapidly but unequally across tasks” (Mar. 2025: 9×–900×/yr, median ~50×). <https://epoch.ai/data-insights/llm-inference-price-trends>
- [256] Epoch AI, “Global AI computing capacity is doubling every 7 months” (Jan. 2026); “Acquisition costs of leading AI supercomputers doubled every 13 months” (Jun. 2025). <https://epoch.ai/data-insights/ai-chip-production>

- [257] NVIDIA, “NVIDIA launches Vera CPU” (Mar. 16, 2026: 88 Olympus cores, LPDDR5X, up to 1.2 TB/s, up to 1.5 TB).
<https://nvidianews.nvidia.com/news/nvidia-launches-vera-cpu-purpose-built-for-agentic-ai>
- [258] TechInsights, “SMIC N+3 confirmed: Kirin 9030 analysis reveals how close SMIC is to 5nm” (Dec. 11, 2025; “significantly less scaled than leading commercial 5nm”). <https://www.techinsights.com/blog/smic-n3-confirmed-kirin-9030-analysis-reveals-how-close-smic-5nm>
- [259] Reuters (Ngui/Shen), “Chipmaker CXMT vaults to top of China’s valuation with 466% surge” (Jul. 27, 2026: \$8.6B raised; 6.73% free float).
<https://www.cnbc.com/2026/07/27/cxmt-china-market-debut-chipmaker-ipo.html>
- [260] TrendForce HBM market bulletin (Jan. 2026): Chinese HBM in early verification stages, limited near-term global impact. <https://www.trendforce.com/research/download/RP260120JO3>
- [261] Microsoft FY2026 Q3 earnings (Apr. 29, 2026): AI business >\$37B annualized run-rate, +123% y/y; capex \$31.9B, guided >\$40B; capacity constrained through 2026.
<https://www.microsoft.com/en-us/investor/earnings/fy-2026-q3/press-release-webcast>
- [262] Bank of England Financial Stability Report (Jul. 2026): AI-focused companies $\approx$ half of S&P 500 value (from $\sim$ a quarter in 2022); Oct. 2025 FPC record and Dec. 2025 FSR on stretched valuations.
<https://www.bankofengland.co.uk/financial-stability-report/2026/july-2026>
- [263] NEA data via carboncredits.com: China +540 GW in 2025 (total 3.89 TW; $\sim$ \$500B invested); 4.04 TW and 1,551 GW record peak by mid-2026 (Global Times); falling utilization hours.
<https://carboncredits.com/china-adds-power-7x-more-than-the-us-in-2025-with-500b-energy-build-out-in-a-single-year/>
- [264] US EIA: 2025 utility-scale additions 53 GW (largest since 2002); 2026 planned record 86 GW.
<https://www.eia.gov/todayinenergy/detail.php?id=67205>
- [265] Ember, Global Electricity Review 2026: China 10,573 TWh (+503 TWh y/y) vs US 4,430 TWh; China fossil generation fell 0.9% in 2025; solar +40%. <https://ember-energy.org/latest-insights/global-electricity-review-2026/major-countries-and-regions/>
- [266] World Nuclear News / POWER: 10 reactors approved Apr. 2025 (\$27.4B), 8 more Jul. 2026; 66 GW installed, targets 110 GW by 2030.
<https://www.world-nuclear-news.org/articles/ten-new-reactors-approved-in-china>
- [267] Carbon Brief/CREA-GEM: record $\sim$ 95 GW coal commissioned 2025; permits at four-year low; coal repositioned as flexibility. <https://www.carbonbrief.org/rush-for-new-coal-in-china-hits-record-high-in-2025-as-climate-deadline-looms>
- [268] Global Energy Monitor / OIES: 45 UHV lines in operation (Dec. 2025); four new 8-GW UHVDC corridors in 2025.
[https://www.gem.wiki/Ultra-High-Voltage_\(UHV\)_Power_Transmission_System_in_China](https://www.gem.wiki/Ultra-High-Voltage_(UHV)_Power_Transmission_System_in_China)
- [269] LBNL “Queued Up” 2026: $\sim$ 2,061 GW active queue; median >3 yrs to agreement; 13% historical completion. <https://www.publicpower.org/periodical/article/backlog-power-plants-seeking-transmission-grid-connection-eased-somewhat-2025-lbnl>
- [270] S&P Global 451 Research (Oct. 2025): US DC grid demand 61.8 GW (2025) $\rightarrow$ 134.4 GW (2030); LBNL 2024 report: DCs 176 TWh (4.4%) in 2023 $\rightarrow$ 6.7–12% by 2028.
<https://www.spglobal.com/energy/en/news-research/latest-news/electric-power/101425-data-center-grid-power-demand-to-rise-22-in-2025-nearly-triple-by-2030>
- [271] Grid Strategies, 2025 National Load Growth Report: 166 GW five-year peak-load growth forecast, $\sim$ 90 GW from data centers. <https://gridstrategiesllc.com/wp-content/uploads/Grid-Strategies-National-Load-Growth-Report-2025.pdf>
- [272] Canary Media: PJM record \$16.4B capacity auction (2027/28); $\sim$ 94% of load growth from DCs; \$70/mo household impact estimates; Good Jobs First: \$4.3B ratepayer-funded DC interconnection (2024). <https://www.canarymedia.com/articles/data-centers/pjm-record-capacity-costs-rising-bills>
- [273] Good Jobs First (Mar. 2026): data-center moratorium bills in $\geq$ 12 states in 2026 sessions.
<https://goodjobsfirst.org/data-center-moratorium-bills-are-spreading-in-2026/>
- [274] Utility Dive: GE Vernova gas-turbine backlog 116 GW (Q2 2026), reservations for 2031 delivery.
<https://www.utilitydive.com/news/ge-vernova-gas-turbine-backlog-climbs-to-116-gw/826039/>

- [275] Smartland Energy (2026, citing Wood Mackenzie): large-transformer lead times >2 yrs, prices +70% since 2019. <https://smartlandenergy.com/transformer-supply-in-2026-a-constraint-thats-resaping-grid-development/>
- [276] WITF (Jul. 2026): Three Mile Island/Crane restart on track for 2027; Michigan Public/ANS on Palisades slippage. <https://www.witf.org/2026/07/28/three-mile-island-reactor-restart-progresses-through-environmental-review-public-comments/>
- [277] OpenAI letter to OSTP (Oct. 27, 2025): US needs 100 GW/yr of new capacity; the “electron gap” vs China’s 429 GW. <https://www.datacenterdynamics.com/en/news/openai-tells-white-house-that-us-needs-100gw-of-new-energy-per-year-to-support-ai-boom/>
- [278] National Data Bureau via TechRadar: East-Data-West-Computing eight-hub program; PUE 1.04 best-in-class; 80%-renewable requirement for new hub DCs. <https://www.techradar.com/pro/website-hosting/china-invested-dollar61-billion-in-eight-national-datacenter-hubs-according-to-director-of-national-data-bureau>
- [279] Semafor/techchannel (Nov. 2025): provinces cut DC electricity up to 50% (to ~RMB 0.4/kWh) exclusively for approved-domestic-chip facilities; offsets 30–50% chip-efficiency gap. <https://www.semafor.com/article/11/05/2025/china-slashes-big-data-centers-electric-bills>
- [280] Oxford Institute for Energy Studies, “The China data centre advantage” (Feb. 2026): TCO decomposition; spot-price overlap; headroom argument. <https://www.oxfordenergy.org/wpcms/wp-content/uploads/2026/02/Comment-The-China-data-centre-advantage.pdf>
- [281] IEA, “Energy and AI” (2025): global DC consumption 415 TWh (2024) → ~945 TWh (2030). <https://www.iea.org/reports/energy-and-ai/executive-summary>
- [282] JFrog (2024: ~100 malicious HF models); ReversingLabs “nullifAI” (Feb. 2025: broken-pickle evasion); Sonatype Picklescan CVEs (2025). <https://www.reversinglabs.com/blog/rl-identifies-malware-ml-model-hosted-on-hugging-face>
- [283] Palo Alto Unit 42, “Model namespace reuse” (Sep. 2025): trusted-name takeover propagating into Vertex AI and Azure AI Foundry. <https://unit42.paloaltonetworks.com/model-namespace-reuse/>
- [284] Hugging Face, “Security incident disclosure — July 2026” (Jul. 16, 2026: dataset-processing RCE, credential harvest, lateral movement; agentic attacker; no public-artifact tampering found). <https://huggingface.co/blog/security-incident-july-2026>
- [285] Anthropic / UK AISI / Turing Institute (Oct. 2025): ~250 poisoned documents backdoor models regardless of scale. <https://www.anthropic.com/research/small-samples-poison>
- [286] Hubinger et al., “Sleeper Agents” (arXiv:2401.05566): backdoors survive SFT/RLHF/adversarial training; agent-supply-chain extensions (CAIS 2026). <https://arxiv.org/abs/2401.05566>
- [287] Hugging Face and VirusTotal collaboration (Oct. 22, 2025): continuous scanning of 2.2M+ public repositories. <https://huggingface.co/blog/virustotal>
- [288] Qiu, Zhou, Ferrara, “Information Suppression in Large Language Models: ...Censorship in DeepSeek” (arXiv:2506.12349): semantic-level suppression across 646 sensitive prompts. <https://arxiv.org/abs/2506.12349>
- [289] Naseh et al., “R1dacted: Investigating Local Censorship in DeepSeek’s R1” (arXiv:2505.12625): topic/language-dependent censorship, transfers to distilled models. <https://arxiv.org/abs/2505.12625>
- [290] Anthropic, “Detecting and preventing distillation attacks” (Feb. 23, 2026): >16M exchanges via ~24,000 accounts (MiniMax >13M, Moonshot >3.4M, DeepSeek >150K). Vendor allegation. <https://www.anthropic.com/news/detecting-and-preventing-distillation-attacks>
- [291] S. Matsuoka, “Latent Model Backdoors and Prompt-Borne Worms: A Mini-Survey of Evidence, the ‘Order 66’ Compound Scenario, and Defenses,” v6.0, 31 July 2026. Companion survey; supplies the evidence hierarchy, threat taxonomy, and control architecture used in Chapter 19 and Appendix E.
- [292] A. Souly, J. Rando, E. Chapman et al., “Poisoning Attacks on LLMs Require a Near-Constant Number of Poison Samples,” arXiv:2510.07192 (2025). <https://arxiv.org/abs/2510.07192>
- [293] Y. Wang, D. Xue, S. Zhang, and S. Qian, “BadAgent: Inserting and Activating Backdoor Attacks in LLM Agents,” ACL 2024, pp. 9811–9827. <https://doi.org/10.18653/v1/2024.acl-long.530>

- [294] W. Zhang and S. Pei, “Your LLM Agent Can Leak Your Data: Data Exfiltration via Backdoored Tool Use,” *Findings of ACL 2026*, pp. 25105–25129. <https://aclanthology.org/2026.findings-acl.1257.pdf>
- [295] D. Meier, J. P. Wahle, P. Röttger, T. Ruas, and B. Gipp, “TrojanStego: Your Language Model Can Secretly Be a Steganographic Privacy Leaking Agent,” *EMNLP 2025*, pp. 27244–27261. <https://doi.org/10.18653/v1/2025.emnlp-main.1386>
- [296] S. Cohen, R. Bitton, and B. Nassi, “Here Comes The AI Worm: Unleashing Zero-click Worms that Target GenAI-Powered Applications,” *arXiv:2403.02817* (rev. 2025). <https://arxiv.org/abs/2403.02817>
- [297] Y. Zhang, Z. Wei, X. Luan et al., “AgentWorm: Self-Propagating Attacks Across LLM Agent Ecosystems,” *arXiv:2603.15727 v3* (Jul. 2026). <https://arxiv.org/abs/2603.15727>
- [298] J. Guan, T. Blanchard, H. Foerster, H. Jia, G. Huang, and N. Papernot, “AI Agents Enable Adaptive Computer Worms,” *arXiv:2606.03811* (Jun. 2026). <https://arxiv.org/abs/2606.03811>
- [299] Y. Liu, Z. Chen, Y. Zhang et al., “Malicious Agent Skills in the Wild: A Large-Scale Security Empirical Study,” *arXiv:2602.06547* (2026). <https://arxiv.org/abs/2602.06547>
- [300] B. Bullwinkel, G. Severi, K. Hines, A. Minnich, R. S. S. Kumar, and Y. Zunger, “The Trigger in the Haystack: Extracting and Reconstructing LLM Backdoor Triggers,” *arXiv:2602.03085* (2026). <https://aka.ms/airt-backdoor-detection>
- [301] Center for AI Standards and Innovation, NIST, *Evaluation of DeepSeek AI Models* (Sept. 2025): agent-configuration susceptibility to indirect prompt injection, measured as attempted malicious actions. https://www.nist.gov/system/files/documents/2025/09/30/CAISI_Evaluation_of_DeepSeek_AI_Models.pdf
- [302] OpenAI, “OpenAI and Hugging Face Partner to Address Security Incident During Model Evaluation” (21 Jul. 2026). <https://openai.com/index/hugging-face-model-evaluation-security-incident/>
- [303] Hugging Face Security Team, “Anatomy of a Frontier Lab Agent Intrusion: A Technical Timeline of the July 2026 Incident” (27 Jul. 2026). <https://huggingface.co/blog/agent-intrusion-technical-timeline>
- [304] Anthropic, “Investigating Three Real-World Incidents in Our Cybersecurity Evaluations” (30 Jul. 2026). <https://www.anthropic.com/news/investigating-incidents-cybersecurity-evals>
- [305] S. Matsuoka, “How Much Compute Does a Nation’s Science Need? A Portable, Measurement-Anchored Framework for Sizing Sovereign AI-for-Science Infrastructure, with Instantiations for Japan, Singapore, and the United States,” *RIKEN R-CCS Technical Report*, July 2026. Supplies the BGPU unit, the two-layer demand model, the parameter taxonomy and transfer rules, the measurement protocol, and the three national instantiations used in Chapter 26.
- [306] S. Matsuoka and T. Schulthess, “AI for Science and an International Consortium,” edited discussion record v1.0, 31 July 2026 (RIKEN R-CCS / CSCS). Unpublished working document; numerical values are discussion estimates unless otherwise noted.
- [307] ECMWF, “ECMWF’s AI forecasts become operational” (25 Feb. 2025) and “Ensemble AI forecasts become operational” (1 Jul. 2025). <https://www.ecmwf.int/en/about/media-centre/news/2025/ecmwfs-ai-forecasts-become-operational>
- [308] NOAA, “NOAA deploys new generation of AI-driven global weather models” (17 Dec. 2025: AIGFS at 0.3% of GFS compute; AIGEFS 31 members at 9%). <https://www.noaa.gov/news-release/noaa-deploys-new-generation-of-ai-driven-global-weather-models>
- [309] C. Bodnar et al., “A foundation model for the Earth system,” *Nature* 641 (21 May 2025). <https://www.nature.com/articles/s41586-025-09005-y>
- [310] D. Kochkov et al., “Neural general circulation models for weather and climate,” *Nature* (22 Jul. 2024). <https://research.google/blog/fast-accurate-climate-modeling-with-neuralgcm/>
- [311] China Meteorological Administration, operational AI models FengQing, FengLei, FengShun (22 Dec. 2025). https://www.cma.gov.cn/en/research/news/202512/t20251222_7506498.html
- [312] FuXi-CNOPs typhoon-track system, IAP-CAS and Fudan University, *Advances in Atmospheric Sciences* (2026): up to 32.3% track-error reduction. <https://www.globaltimes.cn/page/202607/1367074.shtml>

- [313] Google DeepMind, “AlphaFold: five years of impact” (2025). <https://deepmind.google/blog/alphafold-five-years-of-impact/>
- [314] Nobel Prize in Chemistry 2024 (Hassabis, Jumper, Baker), 9 Oct. 2024. <https://www.nobelprize.org/prizes/chemistry/2024/press-release/>
- [315] EMBL-EBI, first complexes in the AlphaFold Database (Mar.–May 2026; ~17M GPU-hours). <https://www.ebi.ac.uk/about/news/technology-and-innovation/first-complexes-alphafold-database/>
- [316] T. Hayes et al., “Simulating 500 million years of evolution with a language model,” *Science* (Jan. 2025). <https://www.science.org/doi/10.1126/science.ads0018>
- [317] Baker Lab, RFdiffusion3 (bioRxiv, Dec. 2025); designed metallohydrolases with catalytic efficiency orders of magnitude above prior designs. <https://www.genengnews.com/topics/artificial-intelligence/rfdiffusion3-now-open-source-designs-dna-binders-and-advanced-enzymes/>
- [318] Google DeepMind, “Gemini Deep Think achieves gold-medal standard at IMO 2025” (21 Jul. 2025: 35/42, 5 of 6 problems). <https://deepmind.google/blog/advanced-version-of-gemini-with-deep-think-officially-achieves-gold-medal-standard-at-the-international-mathematical-olympiad/>
- [319] Google DeepMind, “AlphaEvolve” (May 2025): 48-multiplication 4×4 complex matmul; 0.7% of worldwide compute recovered; 23% kernel speedup $\rightarrow$ 1% Gemini training-time reduction. <https://deepmind.google/blog/alphaevolve-a-gemini-powered-coding-agent-for-designing-advanced-algorithms/>
- [320] Insilico Medicine, rentosertib Phase IIa, *Nature Medicine* (3 Jun. 2025). <https://insilico.com/news/tnrecuxsc1-insilico-announces-nature-medicine-publi>
- [321] Insilico Medicine, Phase III initiation (7 Jul. 2026; 320 patients, 47 centres). <https://insilico.com/news/xmjns4l091-insilico-initiates-phase-iii-clinical-tr>
- [322] MIT Department of Economics, statement on “Artificial Intelligence, Scientific Discovery, and Product Innovation” (16 May 2025): no confidence in the data; withdrawn. <https://economics.mit.edu/news/assuring-accurate-research-record>
- [323] Author Correction to the A-Lab autonomous-synthesis paper, *Nature* (19 Jan. 2026): successes 40 $\rightarrow$ 36. <https://www.nature.com/articles/s41586-025-09992-y>
- [324] Chemistry World, “New analysis raises doubts over autonomous lab’s materials discoveries.” <https://www.chemistryworld.com/news/new-analysis-raises-doubts-over-autonomous-labs-materials-discoveries/4018791.article>
- [325] A. K. Cheetham and R. Seshadri, critique of AI-generated inorganic materials; coverage in The Register (11 Apr. 2024). https://www.theregister.com/2024/04/11/google_deepmind_material_study/
- [326] C. Zeni et al., “A generative model for inorganic materials design,” *Nature* 639 (16 Jan. 2025). <https://www.nature.com/articles/s41586-025-08628-5>
- [327] ECMWF AIFS blog, “Farewell to external AI models” (May 2026): all four external models withdrawn from real-time running over analysis sensitivity. <https://www.ecmwf.int/en/about/media-centre/aifs-blog/2026/farewell-external-ai-models>
- [328] Feng, Luong et al., adjudication of an autonomous agent on 700 open Erdős problems, arXiv:2601.22401 (2026): 68.5% of adjudicated candidates fundamentally flawed. <https://arxiv.org/html/2601.22401v3>
- [329] T. Tao on AI solutions to Erdős problems (Jan. 2026): 1–2% tractable; unreported failures skew perceived rates. <https://the-decoder.com/terence-tao-says-gpt-5-2-pro-cracked-an-erdos-problem-but-warns-the-win-says-more-about-speed-than-difficulty/>
- [330] Agents4Science 2025 (Stanford, Oct. 2025): 314 submissions, 48 accepted; reviewer assessments. <https://www.sciencenews.org/article/science-conference-test-ai-agents>
- [331] MIT Technology Review, “Inside OpenAI’s big play for science” (26 Jan. 2026), incl. working scientists’ assessments. <https://www.technologyreview.com/2026/01/26/1131728/inside-openais-big-play-for-science/>
- [332] *Organization Science* pipeline study (Apr. 2026): submissions +42%, >30% of reviews showing AI use, 3.2% vs 11.9% R&R rates. <https://www.forbes.com/sites/johndrake/2026/04/30/ai-slop-is-flooding-academic-journals-a-top-journal-measured-it/>

- [333] NSF NCSES, *Science and Engineering Indicators 2025*: global R&D \$3.1T (PPP, 2022); US \$923B, China \$812B. <https://nces.nsf.gov/pubs/nsb20257/global-r-d-and-international-comparisons-2>
- [334] Gartner, worldwide IT spending forecast (27 Jul. 2026): software \$1.468T in 2026. <https://www.gartner.com/en/newsroom/press-releases/2026-07-27-gartner-forecasts-worldwide-it-spending-to-grow-14-point-2-percent-in-2026-totaling-6-point-37-trillion>
- [335] Forbes (21 May 2026) on frontier-laboratory revenue scale and composition. <https://www.forbes.com/sites/paulocarvao/2026/05/21/anthropic-openai-enterprise-ai-profitability/>
- [336] McKinsey Global Institute, “The economic potential of generative AI” (2023): R&D the largest affected function ($\sim 24\%$); estimates based on the 2022 economy and excluding new product categories. <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/the-economic-potential-of-generative-ai-the-next-productivity-frontier>
- [337] B. F. Jones and L. H. Summers, “A Calculation of the Social Returns to Innovation,” NBER w27863: \$13.3 per \$1 at a 5% discount rate. https://www.nber.org/system/files/working_papers/w27863/w27863.pdf
- [338] N. Bloom, C. I. Jones, J. Van Reenen and M. Webb, “Are Ideas Getting Harder to Find?” *AER* 110(4), 2020. <https://web.stanford.edu/~chadj/IdeaPF.pdf>
- [339] D. Acemoglu, “The Simple Macroeconomics of AI,” NBER w32487 (2024): TFP +0.66% over ten years, *explicitly excluding* protein-folding-class scientific advances as unlikely within the window. <https://economics.mit.edu/sites/default/files/2024-05/The%20Simple%20Macroeconomics%20of%20AI.pdf>
- [340] T. Davidson, D. Halperin, N. Houlden and A. Korinek, “When Does Automating AI Research Produce Explosive Growth?” NBER w35155 (Apr. 2026). https://www.nber.org/system/files/working_papers/w35155/w35155.pdf
- [341] AI Frontiers, “The quadrillion-dollar disagreement on AI and the economy” (11 May 2026): credible estimates span 0.1%–30% annual growth contribution. <https://ai-frontiers.org/articles/the-quadrillion-dollar-disagreement-on-ai-and-the-economy>
- [342] Nature Index 2026 Research Leaders (10 Jun. 2026): China first, output +22.4% y/y; CAS the leading institution. <https://www.natureasia.com/en/info/press-releases/detail/9350>
- [343] Nature Index 2025 Research Leaders (11 Jun. 2025): China Share 32,122 (+17%), eight of the top ten institutions. <https://group.springernature.com/gp/group/media/press-releases/nature-index-research-leaders-/27786652>
- [344] UC Institute on Global Conflict and Cooperation, “Inside China’s AI for Science strategy” (MOST/NSFC initiative 2023; AI and fusion in the 15th Five-Year Plan). <https://ucigcc.org/blog/inside-chinas-ai-for-science-strategy/>
- [345] Chinese Academy of Sciences, ScienceOne launch (28 Jul. 2025): 12 institutes, 170M papers, 300+ tools. https://english.cas.cn/newsroom/cas_media/202507/t20250728_1048587.shtml
- [346] Chinese Academy of Sciences, ScienceOne Omni 2.0 (20 Jul. 2026). https://english.cas.cn/newsroom/cas-in-media/202607/t20260720_1178699.shtml
- [347] Reporting on an autonomous AI research system connected to China’s national supercomputing network SCNet (Dec. 2025); secondary source, primary announcement not located. <https://tech.yahoo.com/ai/articles/china-rolls-super-ai-science-081459774.html>
- [348] DP Technology Series C (24 Dec. 2025); Bohrium and related platforms serving 1,000+ research organizations and 150 corporate clients. <https://siliconangle.com/2025/12/24/dp-technology-raises-114m-accelerate-chinas-ai-science-industry/>
- [349] Shanghai AI Laboratory, Intern-S1-Pro (2 Apr. 2026): trillion-parameter open-weight scientific multimodal model; SciReasoner 55.5 vs 14.7/13.6 for leading closed models. arXiv:2603.25040. <https://arxiv.org/html/2603.25040v2>
- [350] US DOE, “Energy Department launches Genesis Mission” (Executive Order, 24 Nov. 2025): goal to double the productivity of American science within a decade. <https://www.energy.gov/articles/en-ergy-department-launches-genesis-mission-transform-american-science-and-innovation>
- [351] NVIDIA/Oracle/DOE, Solstice (100,000 Blackwell GPUs) and Equinox (10,000), 28 Oct. 2025. <https://nvidianews.nvidia.com/news/nvidia-oracle-us-department-of-energy-ai-supercomputer-scientific-discovery>

- [352] The White House, Genesis Mission commitments (22 Jul. 2026): >\$5B, 278 projects, 15+ agencies. <https://www.whitehouse.gov/releases/2026/07/45502/>
- [353] European Commission, AI Factories (19 operational plus 13 antennas; €10B via EuroHPC JU). <https://digital-strategy.ec.europa.eu/en/policies/ai-factories>
- [354] European Commission, InvestAI / AI Gigafactories (€20B; 77 proposals across 16 member states). https://commission.europa.eu/topics/competitiveness/competitiveness-coordination-tool-projects/ai-gigafactories_en
- [355] University of Bristol, Isambard-AI launch (17 Jul. 2025): £225M, 5,448 GH200 superchips. <https://www.bristol.ac.uk/news/2025/july/isambard-launch.html>
- [356] RIKEN, FugakuNEXT announcement (22 Aug. 2025): >600 EF FP8 target within ~40 MW; explicit AI-for-Science mandate. https://www.riken.jp/en/news_pubs/news/2025/20250822_1/index.html
- [357] China National Energy Administration, 2025 installed-capacity and renewable-integration statistics. <https://www.nea.gov.cn/20260129/6874f211acd0417eab7ac10c3061a7c2/c.html>
- [358] T. F. Bresnahan and M. Trajtenberg, “General Purpose Technologies: Engines of Growth?” *Journal of Econometrics* 65(1):83–108, 1995.
- [359] E. Helpman (ed.), *General Purpose Technologies and Economic Growth*, MIT Press, 1998.
- [360] E. Brynjolfsson, D. Rock and C. Syverson, “The Productivity J-Curve: How Intangibles Complement General Purpose Technologies,” *AEJ: Macroeconomics* 13(1):333–372, 2021.
- [361] J. A. Schumpeter, *Capitalism, Socialism and Democracy*, Harper & Brothers, 1942.
- [362] M. E. Porter, *Competitive Strategy*, Free Press, 1980.
- [363] C. M. Christensen, *The Innovator’s Dilemma*, Harvard Business School Press, 1997.
- [364] D. J. Teece, G. Pisano and A. Shuen, “Dynamic Capabilities and Strategic Management,” *Strategic Management Journal* 18(7):509–533, 1997.
- [365] C. Y. Baldwin and K. B. Clark, *Design Rules, Volume 1: The Power of Modularity*, MIT Press, 2000.
- [366] M. G. Jacobides, C. Cennamo and A. Gawer, “Towards a Theory of Ecosystems,” *Strategic Management Journal* 39(8):2255–2276, 2018.
- [367] S. Williams, A. Waterman and D. Patterson, “Roofline: An Insightful Visual Performance Model for Multicore Architectures,” *Communications of the ACM* 52(4):65–76, 2009.
- [368] UK Research and Innovation, *ARCHER2 High Performance Computer Evaluation* (Nov. 2025 / Mar. 2026): independent evaluation with explicit counterfactual scenarios and benefit-cost accounting against public investment. <https://www.ukri.org/publications/archer2-high-performance-computer-evaluation-nov-2025/>
- [369] PitchBook via Yahoo Finance, “Anthropic’s gross margin is the most important number in AI” (10 Jun. 2026): \$0.71 compute cost per revenue dollar in Q1 2026, projected \$0.56 in Q2, trajectory toward ~44%. <https://finance.yahoo.com/markets/stocks/articles/anthropics-gross-margin-most-important-070500338.html>
- [370] Lawrence Berkeley National Laboratory, *Queued Up: 2025 Edition* (data through 2024): ~2,290 GW of active generation and storage requests; median request-to-operation time above four years; only 13% of 2000–2019 requested capacity ever reached operation. <https://emp.lbl.gov/publications/queued-2025-edition-characteristics>
- [371] Lawrence Berkeley National Laboratory, *Speed to Power: Solutions for Accelerating Large-Load Connections* (Jun. 2026): no comprehensive national median exists for large-load connection times; utility examples show 6–18-month load studies and 18–60-month construction with line extension; identifies provisional interconnection and curtailable “non-firm” load service among thirteen solution areas. https://eta-publications.lbl.gov/sites/default/files/2026-06/lbnl_large_loads_speed_to_power_final_1.pdf
- [372] Bloomberg / Construction Owners 2026 market guide: Northern Virginia transmission-level data-center hookups quoted at up to ~7 years; Dominion substation-service queue >36 months against 18–24-month shell construction. <https://www.constructionowners.com/insights/how-long-it-actually-takes-to-power-a-data-center-in-2026-a-u-s-market-by-market-reality-check>

- [373] PJM Interconnection, 2028/29 Base Residual Auction results (14 Jul. 2026): \$325/MW-day, 138,318 MW procured, ~\$16.4B, 6,831 MW short of the reliability requirement. <https://www.pjm.com/-/media/DotCom/about-pjm/newsroom/2026-releases/20260714-pjm-capacity-auction-procures-138318-mw-of-generation-resources.pdf>
- [374] ERCOT large-load interconnection data and reporting: large-load queue above 230 GW by Dec. 2025 (~4× a year earlier), ~7.5 GW connected; Texas SB6 (Jun. 2025) imposes study fees, duplicate-request disclosure, security deposits, and remote-curtailment capability on ≥75 MW loads. <https://www.utilitydive.com/news/ercots-large-load-queue-jumped-almost-300-last-year-official/808820/>
- [375] Carbon Brief, analyses of NDRC Document 136 renewable-pricing reform and 2025–26 curtailment: 2025 national wind/solar curtailment 5.4% (breaching the 5% target); Q1 2026 capacity-factor declines concentrated in Inner Mongolia, Xinjiang, Liaoning; only 18 provinces had finalized Doc-136 rules by Oct. 2025. <https://www.carbonbrief.org/analysis-chinas-co2-climbs-2-in-early-2026-due-to-wasted-wind-and-solar>
- [376] Oxford Institute for Energy Studies, “The China data centre advantage” (Feb. 2026): interior spot power 300–400 RMB/MWh (\$43–58); coastal DC power \$0.08–0.11/kWh; 39 UHV lines in service, ~28 more planned by 2030; China’s risk characterized as generation overcapacity, not shortage. <https://www.oxfordenergy.org/wpcms/wp-content/uploads/2026/02/Comment-The-China-data-centre-advantage.pdf>
- [377] People’s Daily / China Daily / Global Times reporting on western hub economics: Qingyang (Gansu) on-site power stabilized at 0.398 RMB/kWh with a dedicated 2 GW green plant; Ulanqab/Hohhot (Inner Mongolia) data-center band 0.32–0.38 RMB/kWh with >80% green supply and 200+ free-cooling days; Gui’an (Guizhou) Q1 2025 data-center electricity use +517% y/y with dedicated substation buildout. <https://en.people.cn/n3/2026/0708/c98649-20475664.html>
- [378] MLCommons, MLPerf Training v6.0 results (16 Jun. 2026): adds DeepSeek-V3 (671B/37B MoE) pretraining and GPT-OSS 20B benchmarks; 95 systems from 24 organizations on 13 accelerator types; no Chinese-designed accelerator submissions. <https://mlcommons.org/2026/06/mlperf-training-v6-0-results/>
- [379] CNBC, “OpenAI resets spend expectations, targets around \$600 billion by 2030” (20 Feb. 2026): compute-spending target walked back from headline \$1.4T commitments; ~\$280B revenue projected by 2030. <https://www.cnbc.com/2026/02/20/openai-resets-spend-expectations-targets-a-round-600-billion-by-2030.html>
- [380] DCD / Forbes reporting: for-profit conversion completed Oct. 2025 (Microsoft ~27%, \$250B Azure commitment, cloud right of first refusal surrendered); exclusivity and revenue sharing reported ended Apr. 2026. <https://www.datacenterdynamics.com/en/news/openai-completes-for-profit-move-microsoft-given-27-stake-and-250bn-azure-contract-but-no-longer-has-cloud-right-of-first-refusal/>
- [381] Microsoft, “Microsoft, NVIDIA and Anthropic announce strategic partnerships” (18 Nov. 2025): Microsoft up to \$5B and NVIDIA up to \$10B into Anthropic; Anthropic commits \$30B of Azure purchases. <https://blogs.microsoft.com/blog/2025/11/18/microsoft-nvidia-and-anthropic-announce-strategic-partnerships/>
- [382] TechCrunch/CNBC, GTC 2026 (16 Mar. 2026): Blackwell+Rubin order visibility raised to ~\$1T through 2027; Vera Rubin production ramp H2 2026. <https://techcrunch.com/2026/03/16/jensen-just-put-nvidias-blackwell-and-vera-rubin-sales-projections-into-the-1-trillion-stratosphere/>
- [383] WSJ via Yahoo Finance, “Stalled: the NVIDIA–OpenAI megadeal” (31 Jan. 2026): the \$100B progressive, non-binding LOI reported stalled amid internal doubts over size and structure. <https://finance.yahoo.com/news/stalled-nvidia-openai-megadeal-ai-131959187.html>
- [384] CNBC, SoftBank FY ended Mar. 2026: Vision Fund gains ~\$46B (>\$45B from OpenAI); ¥5T annual profit; S&P outlook negative; >\$60B toward ~13% of OpenAI; \$6.5B Ampere acquisition completed Nov. 2025. <https://www.cnbc.com/2026/05/13/softbank-earnings-fy-2025-vision-fund-openai-stake.html>
- [385] CNBC / Fortune (Jun.–Jul. 2026): Son declares AI-designing-AI a sign of superintelligence; claims \$5T/yr of global AI investment needed; dismisses bubble concerns; ¥1,000T NAV target. <https://fortune.com/2026/07/14/softbank-ceo-5-trillion-ai-demand-critics-bubble/>

- [386] Arm, Q4 FYE26 results: record FY revenue \$4.92B; record royalties \$2.61B; data-center royalties more than doubled; claimed $\sim 50\%$ CPU share among top hyperscalers; “AGI” datacenter CPU with $> \$2B$ claimed demand. <https://newsroom.arm.com/news/arm-q4-fye26-results>
- [387] Silicon Analysts / trackers: combined 2026 hyperscaler capex guidance $\sim \$700\text{--}745B$ (AMZN $\sim \$200B$, MSFT $\sim \$190B$, GOOGL $\$175\text{--}185B$, META raised twice toward $\$145B$); trailing-4Q D&A $\sim \$149B$; useful-life history incl. Amazon’s partial reversal to 5 years with $\sim \$1.3B$ impact; Alphabet 100-year sterling note; Google Cloud backlog $\sim \$460B$. <https://siliconanalysts.com/analysis/hyperscaler-ai-capex-depreciation-wall-2026>
- [388] Forbes, “Bond investors push back as AI debt heads toward \$570 billion” (17 Jul. 2026): $\sim \$236B$ priced by May 31 at $4\times$ prior-year pace; $\sim \$800B$ private-credit data-center opportunity through 2028; new-issue demand coverage fell from $\sim 5\times$ (Feb.) to $< 2\times$ (Jul.). <https://www.forbes.com/sites/robertszscherba/2026/07/17/bond-investors-push-back-as-ai-debt-heads-toward-570-billion/>
- [389] Forbes / Motley Fool analyses (Jan.–Jul. 2026): Oracle RPO $\$455B \rightarrow \$523B \rightarrow \$638B$; OpenAI contract $> \$300B$ over five years from 2027, est. 40–60% of backlog; OpenAI reassessment right after ~ 5 yrs vs 15–19-yr Oracle leases; $\sim \$380B$ debt-like obligations; negative FCF; CDS > 130 bps; both agencies negative outlook. <https://www.forbes.com/sites/roomykhana/2026/01/15/oracles-ai-bet-equity-wins-and-credit-waits/>
- [390] U.S. Department of Energy, MEXT and METI, Statement of Intent establishing Japan as the first international partner under the Genesis Mission (4 Jun. 2026): $\$1B$ over five years, $\$500M$ from each nation, binding twelve DOE national laboratories to twelve Japanese research institutions—RIKEN, the University of Tokyo, NIMS, KEK and J-PARC among them—across quantum information science, fusion, biotechnology, advanced materials, particle physics, and AI-and-robotics autonomous laboratories, with access to DOE systems and Fugaku. <https://www.energy.gov/articles/united-states-and-japan-announce-historic-1-billion-partnership-under-president-trumps>
- [391] RIKEN, “Message from President Gonokami regarding the new Japan–US strategic partnership on ‘AI for Science’” (5 Jun. 2026): RIKEN as a central implementing institution, with Fugaku, SPring-8 and the BioResource Center named, and the TRIP (2022) and TRIP-AGIS (2023) platforms as the standing AI-for-Science vehicles; collaborating US laboratories named as Argonne, Oak Ridge and Brookhaven. https://www.riken.jp/en/news_pubs/news/2026/20260605_1/index.html
- [392] Argonne National Laboratory with RIKEN, Fujitsu and NVIDIA, partnership on AI for science and next-generation computing (28 Jan. 2026). <https://www.anl.gov/article/argonne-partners-with-riken-fujitsu-and-nvidia-to-advance-ai-for-science-and-next-generation-computing>
- [393] Qwen (@Alibaba_Qwen), official release announcement, 3 Aug. 2026 (pinned): “Next week, the open weights of Qwen3.8-Max will be released, and Qwen3.8-27B is also going open-weights to meet you all”; 2.4T parameters; autonomous coding described as “10+ days of self-evolving development, from empty folder to production without hand-holding, complete project trace in the GitHub”; “system-level autonomous planning with closed-loop adaptive learning, driving 500+ turns of chip design optimization and 365 days of e-commerce strategy”; native multimodal intelligence as “a continuous feedback loop for planning, execution, and self-correction”; pricing $\$2.0/\6.0 per Mtok with $\$0.25$ implicit caching. https://x.com/Alibaba_Qwen
- [394] Qwen Team, Qwen3.8 release blog with the full cross-model benchmark table, Aug. 2026. <https://qwen.ai/blog?id=qwen3.8>
- [395] M. Berman, “Open-source is WINNING,” Matthew Berman (YouTube), Aug. 2026. Review of the Qwen3.8-Max release: identifies research reproduction (PaperBench 93.0) as “the required skill for automated AI research,” noting the model “started from nothing but the paper and a set of GPUs—no starter code, no ready-made pipeline” and “invented and tested 18 improvement ideas of its own across four rounds”; reports Alibaba’s claim that the model “has independently achieved the autonomous execution of the entire silicon design flow,” with the observation that “one of the things that China is behind on is chip design and chip manufacturing—but if they had a model that could help them design these chips, that would no longer be their weakness”; per-task cost data (Qwen3.7-Max $\$1.28$ vs GPT-5.6 Sol $\$1.23$, with Anthropic models most expensive); the model-size gap (US closed frontier rumoured $\sim 7T+$ against 2.4–2.8T Chinese); the enterprise-adoption argument (“95% of the capability at a fraction of the price”); the prediction that “the model will be

- co-designed with the chip. . . with extreme co-design between model and hardware it might actually put us, as we're becoming dependent on Chinese open source, highly dependent on Chinese chips as well"; and the recursive-self-improvement counter-thesis, including OpenAI's attribution of its 80% price cut to efficiency gains produced by its own frontier model. <https://youtu.be/CVlKp9Ld-Zg>
- [396] MarkTechPost / SiliconANGLE / Forbes (3 Aug. 2026): Alibaba releases Qwen3.8-Max—2.4T-parameter sparse MoE, 1M context (991K in / 131K out / 262K reasoning budget), multimodal input, \$2/\$6 per Mtok flat across the full window; *benchmark table published by the vendor; licence terms and activated-parameter count not stated*; a ~95B active figure is press consensus of unclear provenance. <https://www.marktechpost.com/2026/08/03/alibaba-qwen-releases-qwen3-8-max/>
- [397] South China Morning Post (3 Aug. 2026), "Alibaba's Qwen3.8-Max made widely accessible ahead of open-weights release": corroborates the open-weight commitment and the Qwen3.8-27B dense companion, with the repository pending at publication. <https://www.scmp.com/tech/article/3362738/alibabas-ai-model-qwen38-max-made-widely-accessible-ahead-open-weights-release>
- [398] Vendor-reported benchmark table for Qwen3.8-Max as compiled by apidog and OfficeChai (3 Aug. 2026): Terminal-Bench 2.1 86.6; SWE-bench Pro 67.7; PaperBench 93.0; GPQA Diamond 92.6; IFBench 82.8; HLE 43.6; CoWorkBench 74.8; OSWorld-Verified 86.1. No SWE-bench Verified, no AIME, no Kimi K3 or DeepSeek V4 comparators, and—unlike Qwen3.7-Max—no hallucination-rate disclosure. All figures are Alibaba's own. <https://apidog.com/blog/qwen-3-8-benchmarks/>
- [399] Arena.AI / LMArena standings reported by CGTN and SiliconANGLE (Aug. 2026): Qwen3.8-Max #5 overall in text (highest-ranked Chinese model), #2 in vision, #4 on Frontend Code Arena at 1,668 Elo—8 points behind Kimi K3, ~37 behind the leading Claude configuration. Artificial Analysis had not scored the model. <https://news.cgtn.com/news/2026-08-03/Alibaba-unveils-Qwen3-8-Max-its-most-capable-AI-model-to-date-1Pj3AAwmjPa/p.html>
- [400] InfoWorld (3 Aug. 2026), "Alibaba takes aim at OpenAI and Anthropic with Qwen3.8-Max launch": Forrester and Gartner assessments; the formulation that "publishing weights is a separate act from opening an API endpoint—until there is a repository, a licence and a model card, open-weight describes an intention"; and the unaudited-autonomy critique ("sixteen days of what? how many times did a human step in?"). <https://www.infoworld.com/article/4204415/alibaba-takes-aim-at-openai-and-anthropic-with-qwen3-8-max-launch.html>
- [401] Caixin Global (1 Aug. 2026) and technical write-ups of DeepSeek V4-Flash-0731 (official release 31 Jul. 2026; preview from 24 Apr.): ~280B total / 13B active MoE, MIT licence with weights published, 1M context, 384K max output, \$0.14/\$0.28 per Mtok; generation gains attributable to post-training alone (Terminal-Bench 2.1 61.8→82.7 on unchanged architecture); Artificial Analysis Intelligence Index 50; announced 2× peak-hour surge pricing across seven daily hours; the preview was the most-used model on OpenRouter for seven consecutive weeks. <https://www.caixinglobal.com/2026-08-01/deepseek-releases-official-v4-flash-model-as-chinas-ai-race-intensifies-102470292.html>
- [402] ZeroHedge (Aug. 2026), "China's AI knife fight": the eighteen-day chronology (16 Jul. Kimi K3 → 27 Jul. K3 weights → 31 Jul. V4-Flash → 3 Aug. Qwen3.8-Max); four of the five models on the cost-performance Pareto frontier Chinese (Kimi K3, Qwen3.8-Max, GLM-5.2, DeepSeek V4-Flash) with Claude Opus 5 the lone American entry defended by ~37 Elo; benchmark-suite run costs of ~\$0.03 (V4-Flash) against \$3.15 (Fable 5) and \$1.86 (GPT-5.6 Sol). <https://www.zerohedge.com/ai/chinas-ai-knife-fight-deepseeks-new-model-runs-100x-cheaper-anthropics-flagship>
- [403] Forbes / G. Wang (31 Jul. 2026), "Why OpenAI's 80% price cut could trigger a race to the bottom in AI": GPT-5.6 Luna input \$1.00→\$0.20 and output \$6.00→\$1.20 (80% cut), Terra down 20%, Sol unchanged; OpenAI attributed the cut to system efficiency while the analysis attributes competitive pressure to Kimi K3 and DeepSeek V4. Note the sequencing: the cut preceded Qwen3.8-Max's general availability by three days. <https://www.forbes.com/sites/geruiwang/2026/07/31/why-openais-80-price-cut-could-trigger-a-race-to-the-bottom-in-ai/>
- [404] Financial Times reporting relayed by Tom's Hardware (2026): leading Chinese AI firms, Alibaba among them, train models on NVIDIA GPUs leased from operators in Singapore and Malaysia—legal under current export rules because no controlled hardware enters China—with the resulting weights served domestically. No public source connects Qwen3.8-Max training to domestic accelerators. <https://www.tomshardware.com/tech-industry/semiconductors/chinas-top-ai-firms-shift-model-training-overseas-to-access-nvidia-gpus>

- [405] TechTimes (2 Aug. 2026): Alibaba, Moonshot’s largest investor, supplied roughly 20,000 NVIDIA Hopper-class accelerators under a compute agreement (Bloomberg and Reuters corroborated; Alibaba disputes the H200 characterization but not the arrangement); Kimi K3 subsequently outscored Qwen models on several evaluations run on that compute. <https://www.techtimes.com/articles/322699/20260802/alibaba-bankrolled-kimi-k3-20000-nvidia-chips-now-qwen-losing-its-own-compute.htm>
- [406] South China Morning Post, reporting an Interconnects AI analysis of Hugging Face data (10 Apr. 2026): Qwen at ~1 billion cumulative downloads by March 2026, 153.6M in February alone—more than double the combined total of the next eight families—and over 50% of all global open-model downloads, with >200,000 derivative variants. <https://www.scmp.com/tech/big-tech/article/3349552/alibabas-qwen-family-captures-over-50-global-open-source-downloads-report-finds>
- [407] BNN Bloomberg, Alibaba FY results (May 2026): three-year AI capex plan of RMB 380B (~\$56B) to be exceeded without a new stated target; Cloud Intelligence revenue +38% y/y with AI products ~30% of external cloud revenue; ex-items net income down 99.7% and adjusted EBITDA down 84%; CEO statement that “margin is still secondary”; target of >\$100B combined external AI and cloud revenue over five years. <https://www.bnnbloomberg.ca/business/2026/05/13/alibaba-reports-continued-profit-squeeze-from-spending-on-ai-and-instant-retail/>
- [408] White & Case / FPF: Japan’s AI Promotion Act (passed 28 May 2025)—first AI-dedicated law; innovation-first, no penalties, Cabinet-level AI Strategy Headquarters and national AI Basic Plan. <https://www.whitecase.com/insight-alert/japans-first-ai-legislation-becomes-law-focus-promoting-research-and-development-no>
- [409] Japan Times / Bloomberg (Dec. 2025): government moves to roughly quadruple annual chips/AI budget support, deepening state backing for Rapidus; follows the ¥10T-through-2030 package. <https://www.japantimes.co.jp/business/2025/12/26/economy/ai-budget-support/>
- [410] Gibson Dunn / European Parliament: Digital Omnibus provisional agreement (May–Jun. 2026)—AI Act Annex III high-risk obligations postponed to 2 Dec. 2027, embedded high-risk to Aug. 2028; AI Office GPAI powers strengthened. <https://www.gibsondunn.com/eu-ai-act-omnibus-agreement-postponed-high-risk-deadlines-and-other-key-changes/>
- [411] EuroHPC JU (30 Jul. 2026): formal call for up to seven AI gigafactories (each 100k+ advanced chips), expected to trigger >€30B; 76 expressions of interest across 16 member states. https://www.eurohpc-ju.europa.eu/eurohpc-joint-undertaking-launches-ai-gigafactories-call-2026-07-30_en
- [412] Mistral AI / CNBC (9 Sept. 2025): €1.7B Series C at €11.7B (~\$14B) post-money, led by ASML (€1.3B, ~11%). <https://www.cnbc.com/2025/09/09/ai-firm-mistral-valued-at-14-billion-as-asml-takes-major-stake.html>
- [413] IndiaAI Mission GPU tenders: ~38,000 GPUs empanelled across three rounds at subsidized rates reportedly under \$1/GPU-hr; reported utilization shortfall (“can’t find enough users”). <https://www.communicationstoday.co.in/38000-gpus-one-big-problem-indias-ai-mission-cant-find-enough-users/>
- [414] Counterpoint / TrendForce: SK hynix ~62% HBM share; reportedly ~two-thirds of leading HBM4 volume; Samsung/SK hynix scaling into the 2026 memory cycle. <https://www.astutegroup.com/news/general/sk-hynix-holds-62-of-hbm-micron-overtakes-samsung-2026-battle-pivots-to-hbm4/>
- [415] Korea.net / Korea Herald / NVIDIA: national AI computing center; >260,000-GPU APEC deal (Oct. 2025) across government, Samsung, SK, Hyundai, Naver; sovereign foundation-model tournament—five teams selected Aug. 2025, public elimination rounds 2026, 700B-class open models released. <https://nvidianews.nvidia.com/news/south-korea-ai-infrastructure>
- [416] AI Singapore, SEA-LION v4 (Aug. 2025): open multimodal 27B model for 11+ Southeast Asian languages. <https://sea-lion.ai/blog/sea-lion-v4-multimodal/>
- [417] NVIDIA/HUMAIN partnership (500 MW-class “AI factories,” several hundred thousand GPUs); AMD–HUMAIN up to \$10B/6 GW; Stargate UAE 1 GW within a 5 GW campus; US Commerce authorization of up to 70,000 advanced chips to G42/HUMAIN (Nov. 2025). <https://www.commerce.gov/news/press-releases/2025/11/statement-uae-and-saudi-chip-exports>
- [418] Tom’s Hardware / CAPRI: Taiwan’s ~NT\$100B “ten major AI infrastructure” plan; post-nuclear grid strain as the binding constraint on AI and foundry expansion. <https://www.tomshardware.com/tech-industry/semiconductors/taiwan-to-spend-billions-on-ai-island-push-despite-grid-risks>

- [419] Cassava Technologies (Mar. 2026): Africa’s first NVIDIA-powered AI factory (South Africa); Kenya, Nigeria, Egypt, Morocco expansions planned.
<https://www.cassavatechnologies.com/cassava-scales-african-ai-infrastructure-with-nvidia-powered-ai-factories-to-accelerate-sovereign-data-capabilities/>
- [420] Caixin / Fortune / State Council reporting: anti-“involution” campaign; pricing-law amendments (Aug. 2025); State Council warning against “disorderly competition” in AI; Xi Jinping publicly questioning whether every province needs AI/compute/EV projects.
<https://fortune.com/asia/2025/08/29/china-warns-against-disorderly-competition-ai-race>
- [421] Recode China AI / Asia Tech Review / Sinocities: Zhipu and MiniMax post-IPO disclosures show thin revenue and heavy losses; documented “low-trust open-source paradox” limiting domestic enterprise adoption of domestic open models.
<https://www.recodechinaai.com/p/zhipu-ai-and-minimax-just-went-public>
- [422] MIT Technology Review (Mar. 2025) / Light Reading / ASPI: hundreds of state-subsidized Chinese AI data centers with reported utilization far below plan (up to ~80% of some new capacity unused); rental-price collapse ending the buildout binge.
<https://www.technologyreview.com/2025/03/26/1113802/china-ai-data-centers-unused/>
- [423] ChinaTalk / CSIS / trade reporting: HBM and advanced packaging as binding constraints on the Ascend fleet; CXMT assessed roughly three years behind the Korean leaders.
<https://www.chinatalk.media/p/mapping-chinas-hbm-advancement>
- [424] NVIDIA, Cosmos world-model platform; Cosmos 3 omnimodal world models for physical AI (Jun. 2026). <https://developer.nvidia.com/blog/develop-physical-ai-reasoning-world-and-action-models-with-nvidia-cosmos-3/>
- [425] Google DeepMind, “Genie 3: a new frontier for world models” (Aug. 2025).
<https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>
- [426] Figure AI Helix VLA (S2 deliberative 7–9 Hz + S1 visuomotor 200 Hz; trained on ~500 h teleoperation); Helix-02 completed a livestreamed 8-hour autonomous package-sorting shift at human-comparable rates (13 May 2026); 350+ Figure 03s delivered.
<https://www.techtimes.com/articles/316632/20260514/figure-ais-helix-02-robots-complete-full-8-hour-autonomous-shifts-humanoid-race-intensifies.htm>
- [427] Physical Intelligence π -series VLA models; funding \$400M@\\$2.4B (2024) $\rightarrow$ ~\\$5.6B, reported talks at >\\$11B (Mar. 2026); Skild AI reported ~\\$14B.
<https://mlq.ai/news/physical-intelligence-in-talks-to-raise-1-billion-at-11-billion-valuation/>
- [428] Rest of World, “Inside China’s robot training centers” (2026): 40+ state-backed data-collection centers announced (~24 operational); teleoperators with VR/exoskeletons; UBTech RMB 566M of sales to three provincial centers; China Mobile RMB 124M orders to Unitree/AgiBot.
<https://restofworld.org/2026/china-robots-training-centers-workers/>
- [429] China Daily / Global Times: Shanghai heterogeneous humanoid training facility (Jan. 2025), 100+ robots simultaneously, target 1,000 by 2027 collecting 10M+ entries.
<https://www.chinadaily.com.cn/a/202501/23/WS67919d43a310a2ab06ea8c2e.html>
- [430] AgiBot World Colosseo open embodied dataset: 1M+ real-robot trajectories, 100 robots, 100+ scenarios; Genie Sim 3.0 open-source simulation (CES 2026).
<https://opendrivelab.com/AgiBot-World/>
- [431] Open X-Embodiment collaboration (Google DeepMind + 33 institutions): pooled cross-embodiment dataset and RT-X models. <https://robotics-transformer-x.github.io/>
- [432] GeekWire, Agility Robotics SPAC filings (2026): \\$2.5B valuation; >\\$300M committed multi-year Digit orders incl. a 1,000-robot contract; 65,000+ hours of real-world operation; 2025 opex \\$111M; GXO multi-year RaaS deployment. <https://www.geekwire.com/2026/digit-maker-agility-robotics-to-go-public-in-2-5b-deal-heres-what-the-filings-say-about-its-finances/>
- [433] TechTimes/Electrek (Jul. 2026): Tesla Optimus at zero verified production units; Fremont line start slated late Jul./Aug. 2026; Gen-3 reveal deferred twice; earnings-call statement that existing units are “primarily for learning, not productive tasks.” <https://www.techtimes.com/articles/321012/20260720/tesla-optimus-production-count-remains-zero-q2-earnings-call-looms-wednesday.htm>

- [434] Humanoid.guide / Pandaily / Forbes: Leaderdrive as China's largest harmonic-reducer maker lifted by humanoid demand; Zhongda Leader founders became billionaires on joint demand (Apr. 2026); Unitree price war squeezing the reducer chain. <https://www.forbes.com/sites/zinnialee/2026/04/23/humanoid-mania-turns-chinese-brothers-behind-robotic-joint-maker-into-billionaires/>
- [435] Fast Company, "This tiny screw is powering the humanoid robot revolution": planetary roller screws ~33% of humanoid BoM (J.P. Morgan est.); \$1,350–2,700/unit; China-dominant production; Shanghai Beite's \$260M plant. <https://www.fastcompany.com/91314612/this-tiny-screw-is-powering-the-humanoid-robot-revolution>
- [436] 36Kr / Futu: reported ~\$685M Optimus actuator-assembly order to Sanhua; Ningbo Tuopu building robot drive-system capacity. <https://eu.36kr.com/en/p/3510288514980998>
- [437] Yicai Global (Q2 2025): Hesai robotics lidar shipments 49,000 units (+744% y/y); RoboSense 34,400 (+632%); both pivoting from automotive to robotics. <https://www.yicaiglobal.com/news/chinas-robosense-hesai-see-robot-lidar-sales-surge-in-second-quarter>
- [438] CnEVPost / SCMP (24 Jun. 2026): CATL–Galbot agreement; Galbot S1 deployed at scale on CATL battery module/pack lines; Galbot cumulative funding ~\$800M at ~\$3B. <https://cnevpost.com/2026/06/24/catl-galbot-scale-humanoid-robots-battery-factories/>
- [439] NVIDIA, Jetson Thor launch (25 Aug. 2025): Blackwell-based, 2,070 FP4 TFLOPS, 128 GB, dev kit \$3,499; Isaac GR00T open reference humanoid for academic research (Jun. 2026). <https://nvidianews.nvidia.com/news/nvidia-blackwell-powered-jetson-thor-now-available-accelerating-the-age-of-general-robotics>
- [440] ISO 25785-1 (industrial mobile robots with actively controlled stability) drafting toward ~2027 publication; ISO 10218:2025 integrating ISO/TS 15066 force limits; UL 3300 for consumer robots. <https://humanoidliability.com/resources/iso-25785-humanoid-safety-standards/>
- [441] Goldman Sachs Research: humanoid TAM \$38B by 2035; unit manufacturing cost fell \$50–250k → \$30–150k in one year (~40% vs 15–20% expected). <https://www.goldmansachs.com/insights/articles/the-global-market-for-robots-could-reach-38-billion-by-2035>
- [442] Morgan Stanley (Jun. 2026): China humanoid shipments 12,000 (2025) → forecast 50,000 (2026) → 446,000/yr by 2030; localized Chinese chains ~20% cheaper; average price ~\$46k trending to ~\$21k; "Humanoid 100" value-chain map. <https://www.cnbc.com/2026/06/24/morgan-stanley-china-humanoid-robot-market-forecast.html>
- [443] Gasgoo / Fortune (Jun. 2026): Unitree STAR-market IPO approved; 5,500 humanoids sold 2025; humanoid ASP RMB 593k (2023) → RMB 166k (2025); gross margin 87.7% → 63.2%; ~\$250M revenue with ~\$41M profit—among the only profitable humanoid makers. <https://fortune.com/2026/06/06/chinese-humanoid-robots-global-market-sales-performative-functional/>
- [444] BigGo / Tech Digest / 36Kr supply-chain reporting [A]: ~70% of Optimus supply chain China-based; excluding Chinese parts estimated to triple cost; 2025 rare-earth licensing touched actuator magnets. <https://www.techdigest.tv/2026/02/excluding-chinese-parts-could-triple-tesla-optimus-costs.html>
- [445] Fortune (6 Jun. 2026): 2025 global humanoid shipments >13,000, ~85% Chinese; US firms "a few hundred or less"; most units performative rather than functional—demos, exhibitions, education. <https://fortune.com/2026/06/06/chinese-humanoid-robots-global-market-sales-performative-functional/>
- [446] The Robot Report (31 Mar. 2026): AgiBot ships 10,000th humanoid; second five thousand in ~3 months; Omdia/IDC #1 by 2025 shipments. <https://www.therobotreport.com/agibot-rolls-out-10000th-humanoid-robot/>
- [447] PRNewswire / Interesting Engineering: UBTECH Walker S2 mass production from Nov. 2025, 1,000th unit built; cumulative Walker orders >RMB 800M; multi-robot deployments at Zeekr, BYD, Foxconn, Audi-FAW. <https://www.prnewswire.com/news-releases/ubtech-humanoid-robot-walker-s2-begins-mass-production-and-delivery-with-orders-exceeding-800-million-yuan-302616924.html>
- [448] TechTimes / ChinaBizInsider (Jun. 2026): BYD confirms four-year internal humanoid program; reported target ~20,000 robots in its own factories in 2026. <https://www.techtimes.com/articles/317949/20260607/byd-confirms-four-year-secret-humanoid-robot-program-could-sell-through-car-dealerships.htm>

- [449] Gasgoo / Humanoids Daily: XPeng IRON mass production targeted end-2026; stated goal 1M units by 2030. <https://www.humanoidsdaily.com/news/xpeng-sets-2026-target-for-mass-produce-d-iron-robot-eyes-1-million-units-by-2030>
- [450] MIIT, “Guiding Opinions on Innovation and Development of Humanoid Robots” (Nov. 2023): initial mass production by 2025; humanoids as “an important new engine of economic growth” by 2027; embodied AI in the 2025 Government Work Report and 15th Five-Year Plan. <https://www.therobotreport.com/china-plans-to-mass-produce-humanoids-by-2025/>
- [451] Carnegie Endowment / Jamestown Foundation (2025–26): \$8.2B national AI-fund allocation; Beijing RMB 100B robotics fund; Shenzhen/Guangdong/Beijing action plans with enterprise and revenue targets; 150+ humanoid companies; H1 2025 humanoid financing >RMB 10B. <https://carnegieendowment.org/research/2025/11/embodied-ai-china-smart-robots>
- [452] Figure AI Series C (Sept. 2025): ~\$1B at \$39B post-money. <https://sacra.com/c/figure-ai/>
- [453] Forbes / Hyundai (CES 2026): electric Atlas 2026 production run fully committed; Hyundai plant designed for 30,000/yr; deployment of a committed 25,000 units gated by union negotiations. <https://www.forbes.com/sites/johnkoetsier/2026/01/06/atlas-humanoid-robots-production-fully-committed-for-2026-factory-will-build-30000-per-year/>
- [454] Futurism / Electrek / Humanoids Daily: Tesla “We, Robot” bartenders remotely operated (analyst-confirmed); Dec. 2025 Optimus demo teleoperation controversy; 1X NEO ships with explicit human-teleoperation fallback. <https://www.humanoidsdaily.com/news/1x-neo-launch-sparks-debate-on-autonomy-and-teleoperation>
- [455] International Federation of Robotics, *World Robotics 2025*: 542,000 industrial robots installed 2024; China 295,000 (54%); operational stock 4.66M (China >2M); Chinese makers 57% domestic share vs ~28% a decade earlier. https://ifr.org/downloads/press_docs/2025-09-25-IFR_press_release_China_in_English.pdf
- [456] A. K. Agrawal, J. McHale and A. Oettl, “AI in Science,” NBER Working Paper 34953, 2026. Models AI in science as augmentation through search, with large variation across workflow stages and fields. <https://doi.org/10.3386/w34953>
- [457] T. C. Fort, N. Goldschlag, J. Liang, P. K. Schott and N. Zolas, “Growth is Getting Harder to Find, Not Ideas,” NBER Working Paper 35182, 2026. <https://doi.org/10.3386/w35182>
- [458] Hyperion Research Holdings, LLC, *Analysis of the Ripple Effects of the Next-Generation Computing Platform*, supplementary final survey report, revised ed., commissioned by RIKEN, 5 Mar. 2025 (#HR4.0523.02.05.2025). Fugaku COVID-19 projects estimated at \$50–75B returned to the Japanese economy; successor-machine lifetime return scenarios \$250B/\$400B+/\$600B+; 65.7% of surveyed Fugaku researchers reporting realized results on SDG-class projects; innovation-class mapping of 650+ worldwide HPC projects; international five-year AI investment estimates.
- [459] R. R. Hill and C. Stein, “How Artificial Intelligence Shapes Science: Evidence from AlphaFold,” NBER Working Paper 35143, 2026. <https://doi.org/10.3386/w35143>
- [460] M. Liu and C. Zhai, “An Axiomatic Benchmark for Evaluation of Scientific Novelty Metrics,” arXiv:2604.15145, 2026. <https://arxiv.org/abs/2604.15145>
- [461] L. Xiong et al., “AutoResearchBench: Benchmarking AI Agents on Complex Scientific Literature Discovery,” arXiv:2604.25256, 2026: frontier agents below 10% on demanding literature-discovery tasks. <https://arxiv.org/abs/2604.25256>
- [462] Kimi Team et al., “Kimi K3: Open Frontier Intelligence,” arXiv:2607.24653, 2026: 2.8T total / 104B activated parameters, 1M-token context, native vision; the report itself acknowledges trailing the strongest proprietary systems overall. <https://arxiv.org/abs/2607.24653>
- [463] Moonshot AI, “Kimi K3 License” (custom commercial license: use, modification and redistribution permitted; separate agreement required for qualifying model-as-a-service businesses above a revenue threshold; attribution conditions on very large deployments). <https://huggingface.co/moonshotai/Kimi-K3/blob/main/LICENSE>
- [464] Qwen Team, Qwen3.6 official repository: open releases are the 27B dense and 35B-A3B variants, Apache 2.0. <https://github.com/QwenLM/Qwen3.6>
- [465] Reuters, “Alibaba unveils its largest AI model yet; DeepSeek’s latest model is ultra-low cost” (3 Aug. 2026). *Superseded* by the primary-source entries above: the report’s pairing of the two releases

- obscured that DeepSeek V4-Flash shipped officially on 31 July, and its ~ 95 B activated-parameter figure for Qwen3.8-Max is press consensus rather than a vendor disclosure.
<https://www.reuters.com/business/retail-consumer/alibaba-unveils-its-most-capable-ai-model-date-not-far-behind-moonshots-size-2026-08-03/>
- [466] Microsoft, FY26 Q3 earnings release (29 Apr. 2026): AI annual revenue run rate above \$37B, +123% y/y. <https://www.microsoft.com/en-us/investor/earnings/fy-2026-q3/press-release-webcast>
 - [467] NVIDIA, Q1 FY2027 financial results (20 May 2026): revenue \$81.6B, data center \$75.2B, non-GAAP gross margin $\sim 75\%$. <https://investor.nvidia.com/news/press-release-details/2026/NVIDIA-Announces-Financial-Results-for-First-Quarter-Fiscal-2027/default.aspx>
 - [468] RISC-V International, IME task-group ratification plan (target Aug. 2026).
<https://riscv.atlassian.net/wiki/spaces/IMEX/pages/598867969/IME+Ratification+Plan>
 - [469] RISC-V International, AME task-group ratification plan (board-ratification target Apr. 2027; simulators, ABI, GCC/LLVM support and basic kernels as staged milestones).
<https://riscv.atlassian.net/wiki/spaces/AMEX/pages/55083420/Ratification+Plan>
 - [470] SK hynix, “SK hynix Begins Mass Production of 192GB SOCAMM2” (20 Apr. 2026).
<https://news.skhynix.com/mass-production-socamm2-192gb/>
 - [471] NVIDIA, Vera data-center CPU: up to 1.5 TB LPDDR5X capacity and ~ 1.2 TB/s memory bandwidth per socket. <https://www.nvidia.com/en-us/data-center/vera-cpu/>
 - [472] Reuters, “CXMT plans second chip plant in Beijing and is in talks on its funding, sources say” (3 Aug. 2026): three fabs of ~ 100 k wafers/month each, potential expansion beyond 600k wpm; CXMT remains far smaller than the three global leaders. <https://www.reuters.com/world/asia-pacific/cxmt-plans-second-chip-plant-beijing-is-talks-its-funding-sources-say-2026-08-03/>
 - [473] Meta AI Research, “Introducing Muse Glimmer: An Open Agentic Model That Runs on Your Device” (10 Aug. 2026): 30B dense multimodal, Apache 2.0, weights on Hugging Face/Ollama; distilled from Muse Spark.
<https://research.meta.ai/blog/introducing-muse-glimmer-open-agentic-model>
 - [474] M. Zuckerberg, “The Future Is For Everyone” (10 Aug. 2026): 6,500-word letter recommitting Meta to open weights; companion X announcement of Muse Glimmer weights and a forthcoming open-weight version of Muse Spark 1.2; contemporaneous coverage in WSJ/NYT/Bloomberg.
<https://www.meta.com/thefutureisforeveryone/>
 - [475] Artificial Analysis, “Muse Glimmer” (Aug. 2026): independent evaluation—Intelligence Index 35 (Gemma 4 31B: 30; Qwen3.6 27B Reasoning: 38; Kimi K2.5: 36); measured hallucination rate 82% vs 49% for Qwen3.6 27B; “Meta’s most permissive license yet.”
<https://artificialanalysis.ai/articles/muse-glimmer>
 - [476] Digital Applied, “Qwen3.8 open weights: checklist before download” (7 Aug. 2026) and byteiota (9 Aug. 2026): the 3 Aug. “next week” commitment lapsed without a repository; window re-framed to the week of 10 Aug.; fourth consecutive Max-cycle slippage documented.
<https://www.digitalapplied.com/blog/qwen3-8-open-weights-checklist-before-download>
 - [477] TechTimes (4 Aug. 2026): MiniMax H3 tops the open-weight video arena; licence excludes local deployment in the US, EU, UK, and South Korea. <https://www.techtimes.com/articles/322904/20260804/minimax-h3-open-weights-exclude-us-eu-uk-korea-local-deployment.htm>
 - [478] Caixin Global (11 Aug. 2026): Alibaba opens the Qianwen platform to outside developers for custom AI agents across phones, PCs, and smart glasses; launch partners SF Express, Ziroom, Midea. <https://www.caixinglobal.com/2026-08-11/alibaba-opens-qwen-platform-to-outside-developers-in-push-to-build-ai-agent-ecosystem-102473024.html>
 - [479] CNBC (6 Aug. 2026): Beijing’s broadest countermeasures since the October truce—case-by-case export review on drones and components, six US entities sanctioned, national-security probe into imported printers; triggered by FCC Covered List and UFLPA additions in late July. <https://www.cnbc.com/2026/08/06/us-china-trump-xi-summit-trade-technology-restrictions.html>
 - [480] Fortune (4 Aug. 2026) / Axios (3 Aug. 2026): White House finalizes a voluntary frontier-model evaluation framework in a closed-door session with OpenAI, Anthropic, Microsoft, Meta, NVIDIA and others; document not publicly released.
<https://fortune.com/2026/08/04/baffling-white-house-wont-publicly-release-ai-model-evaluation-framework-it-reviewed-today-with-openai-anthropic-microsoft-and-others/>

- [481] OpenAI, “Ten advances in mathematics” (Aug. 2026): the Astra model’s solutions to ten long-open problems in mathematics and quantum complexity, proofs published; ~\$2,000 of inference compute reported. <https://openai.com/index/ten-advances-in-mathematics/>
- [482] NIST (4 Aug. 2026): NIST joins the Genesis Mission with two AI Economic Security Centers—AI agents for manufacturing (tenfold drone-production scaling goal in two years) and for critical-infrastructure cyber-defense. <https://www.nist.gov/news-events/news/2026/08/nist-joins-national-genesis-mission-accelerate-ai-innovation>
- [483] Benzinga (8 Aug. 2026): CoreWeave ahead of Q2—bear market off the July low, 15% short interest, net loss widened to \$740M, debt above \$25B against \$30–35B capex guidance, Jefferies and Raymond James downgrades; revenue nonetheless doubled y/y to \$2.07B on a \$99B backlog. <https://www.benzinga.com/markets/equities/26/08/61062163/coreweave-stock-stuck-in-a-bear-market-as-short-interest-rises-ahead-of-earnings>
- [484] Alibaba / Qwen, model repositories **Qwen/Qwen3.8-2.4T-A95B** (2.4T total / 95B active; 512 experts, 11 activated; 262,144 native context; text-only; 4.89 TB in 213 BF16 shards) and **Qwen/Qwen3.8-27B** (Apache 2.0, vision-language). The flagship carries the custom **qwen3.8-max** licence: separate commercial agreement required where model-as-a-service or AI-work-assistant revenue exceeds US\$50M in any consecutive 12 months; attribution required above 100M MAU or US\$20M monthly revenue. <https://huggingface.co/Qwen/Qwen3.8-2.4T-A95B>
- [485] South China Morning Post (13 Aug. 2026), “Alibaba adds commercial restrictions to open-weight Qwen3.8-Max AI model.” Note that several outlets reported the release as uniformly Apache 2.0; the licence file contradicts this for the flagship. <https://www.scmp.com/tech/tech-trends/article/3363927/alibaba-adds-commercial-restrictions-open-weight-qwen38-max-ai-model>
- [486] DeepSeek, *DeepSeek Harness* (**dsh**), MIT licence, developer preview v0.1, released 13 Aug. 2026; ~100,200 GitHub stars and 9,400 forks within ~48 hours. <https://github.com/deepseek-ai/deepseek-harness>
- [487] DeepSeek, DeepSeek Harness documentation: plugin-native architecture (models, tools, skills, sessions, sandboxes, filesystems, loops, orchestration and UI as components; “no privileged core to patch”), four runtime modes, append-only session logging with resume/fork/replay. <https://deepseek.com/harness/en/>
- [488] DeepSeek Harness MCP client plugin documentation: stdio and streamable-http transports; tools namespaced **mcp__<server>__<tool>**; MCP *Resources* and *Prompts* not supported; image/audio payloads become placeholders. <https://deepseekdocs.com/en/docs/features/mcp>
- [489] Y. Shi, W. Zhang and T. Cui (Peking University; DeepSeek-AI), *A Programming Paradigm for Spatiotemporal Composability*, draft of 13 Aug. 2026, 88 pp., self-hosted preprint marked “under active revision.” Formalizes reversible effects and reactive coefficients; implements them in Cordis. Cites OpenAI’s “Harness Engineering” and Anthropic’s “Harness Design for Long-Running Application Development.” <https://github.com/cordiverse/paper>
- [490] South China Morning Post (14 Aug. 2026): the harness launch marks a strategic pivot—“shifting the global competition from model intelligence to one that focuses on how seamlessly an AI agent can connect to real-world software.” <https://www.scmp.com/tech/big-tech/article/3364082/>
- [491] The New Stack (13 Aug. 2026): DeepSeek Harness architecture, model-agnostic provider support (Anthropic, OpenAI, Bedrock, Azure, Gemini, DeepSeek), subagent bridges to Claude Code and OpenAI Codex; external pull requests not accepted. <https://thenewstack.io/deepseek-harness-open-source-plugins/>
- [492] VentureBeat (13–14 Aug. 2026), “DeepSeek Harness launches as open-source rival to Claude Code, alongside V4-Pro on API with higher prices.” <https://venturebeat.com/technology/deepseek-harness-launches-as-open-source-rival-to-claude-code-alongside-v4-pro-on-api-with-higher-prices>
- [493] 36Kr (14 Aug. 2026): Cordis originates as the kernel of Koishi, a four-year-old open-source chatbot framework with 4,000+ community plugins; the paper’s first author is Koishi’s creator. <https://eu.36kr.com/en/p/3938795963137411>
- [494] M. Sayed Ali, A. Novik, A. Boddupally et al. (Writer, Inc.), *The Harness Effect: How Orchestration Design Sets the Token Economics of Enterprise Agentic AI*, arXiv:2607.06906, Jul. 2026. 22 locked tasks, six models, five vendors: harness adoption saves 33–61% holding the model

- fixed, against 36% from switching between cheapest and most expensive model. Vendor-authored; n=22. <https://arxiv.org/html/2607.06906v1>
- [495] J. Liu, X. Zhao, X. Shang and Z. Shen (MBZUAI; UCL), *Dive into Claude Code: The Design Space of Today's and Future AI Agent Systems*, arXiv:2604.14228, Apr. 2026. Anatomy of a production harness; reports the community 1.6%/98.4% code-share estimate as an estimate over a leak-derived bundle, not a measurement. <https://arxiv.org/html/2604.14228v1>
 - [496] Comparative harness reporting, 2026: convergent architecture across Claude Code, Codex CLI, Aider and OpenClaw; system-prompt budgets differing by nearly two orders of magnitude; Codex's one-line importer for Claude Code configuration; the duelling positions of B. Cherny (Anthropic: the harness is "the thinnest possible wrapper") and K. Prasad (Obvix Labs: "the model is interchangeable—the harness is not"). <https://www.techtimes.com/articles/316928/20260521/>
 - [497] Octomind, "Why we no longer use LangChain for building our AI agents": rigid abstractions, dynamic tool gating unavailable, debugging overhead—"when our team began spending as much time understanding and debugging LangChain as it did building features, it wasn't a good sign." No switching cost quantified. <https://octomind.dev/blog/why-we-no-longer-use-langchain-for-building-our-ai-agents>
 - [498] Anthropic (9 Dec. 2025): donation of the Model Context Protocol to the Linux Foundation's Agentic AI Foundation, co-founded with Block and OpenAI; supporting members Google, Microsoft, AWS, Cloudflare, Bloomberg; 10,000+ active public MCP servers and 97M+ monthly SDK downloads at donation. <https://www.anthropic.com/news/donating-the-model-context-protocol-and-establishing-of-the-agentic-ai-foundation>
 - [499] Linux Foundation (9 Apr. 2026): the Agent2Agent protocol passes 150+ organizations, 22,000+ GitHub stars and five production SDKs in its first year. <https://www.linuxfoundation.org/press/a2a-protocol-surpasses-150-organizations-lands-in-major-cloud-platforms-and-sees-enterprise-production-use-in-first-year>
 - [500] Agent Plugins 1.0 (6 Aug. 2026): vendor-neutral package format for distributing Agent Skills and MCP servers across clients; technical steering committee of core maintainers from Amazon, Cursor, Microsoft, OpenAI and Vercel, with Google joining the same day; independently governed. <https://developers.googleblog.com/agent-plugins-package-your-skills-tools-and-more/>
 - [501] The New Stack (8 Aug. 2026) and contemporaneous analysis of Agent Plugins 1.0: the spec standardizes packaging, not distribution—"if useful behavior ends up inside vendor namespaces, the standard will look open while lock-in just moves elsewhere" (A. Dalloul, EmpirioLabs); Anthropic, whose format the standard is based on, did not join. <https://thenewstack.io/agent-plugins-open-standard/>
 - [502] Cyberspace Administration of China (17 Jul. 2026, WAIC), *Global Cooperation Initiative on Agent Mutual Trust, Interconnection, and Interoperability*: calls for "open, non-discriminatory, and transparent international interoperability standards for agents," agent identity systems and shared multi-agent infrastructure; names no existing protocol. <https://www.geopolitechs.org/p/chinas-cac-releases-global-cooperation>
 - [503] OpenSandbox: agent execution sandbox released by Alibaba in March 2026, Apache 2.0, subsequently moved to a neutral organization with published governance, an OpenSSF best-practices badge and a CNCF landscape listing; gVisor/Kata/Firecracker isolation; listed integrations include Claude Code, Gemini CLI, LangGraph, Google ADK and Qwen Code. <https://github.com/opensandbox-group/OpenSandbox>
 - [504] LangChain, *State of Agent Engineering* (n=1,340; fielded Nov.–Dec. 2025): 57% have agents in production, 67% among 10,000+ employee enterprises; over three-quarters run multiple models; the survey does not measure framework-switching rates. <https://www.langchain.com/state-of-agent-engineering>
 - [505] SpaceXAI (xAI LLC), Grok 4.6 model card, rev. 12 Aug. 2026: closed weights, 500K context, \$2/\$6 per Mtok, same 1.5T base as Grok 4.5 with gains from post-training; Artificial Analysis Intelligence Index 61, tying GPT-5.6 Sol and two points below Claude Opus 5. <https://media.x.ai/v1/website/card-4p6-4cd2dc57.pdf>
 - [506] the-decoder (14 Aug. 2026): Zhipu releases GLM-5.3 via API claiming the strongest open-weights coding model; weights withheld at announcement pending a security review, and published on 28

- Aug. 2026; 743B MoE on the GLM-5.2 base; all benchmarks vendor-reported, and it trails GPT-5.6 Sol on Terminal-Bench 3.0 and DeepSWE v1.1. <https://the-decoder.com/zhipu-ai-releases-glm-5-3-claims-its-the-strongest-open-weights-coding-model/>
- [507] Techzine (12 Aug. 2026), relaying *The Information*: NVIDIA reported to be in the final training phase of a trillion-parameter “Nemotron 4” intended to compete with the most powerful open-source models; NVIDIA states final training is not complete; licence unannounced; possible availability late autumn 2026. <https://www.techzine.eu/news/analytics/143552/nvidia-is-buildin-g-nemotron-4-with-at-least-1-trillion-parameters/>
 - [508] SiliconANGLE (13 Aug. 2026): Google launches Gemini 3.7 Flash for coding and agents at \$0.75 per Mtok input. <https://siliconangle.com/2026/08/13/google-launches-gemini-3-7-flash-coding-ai-agent-projects/>
 - [509] Reuters via Yahoo Finance (13 Aug. 2026): DeepSeek takes V4-Pro to general availability, raises API prices for V4 Pro and V4 Flash, introduces peak/off-peak pricing; reported plans to at least double headcount and to expand private hiring of chip-design engineers toward an in-house accelerator reducing reliance on both NVIDIA and Huawei. <https://finance.yahoo.com/news/deepseek-releases-official-v4-pro-125608280.html>
 - [510] NVIDIA, “NVIDIA Vera Rubin POD: seven chips, five rack-scale systems, one AI supercomputer” (16 Mar. 2026, mod. 4 Aug. 2026): Rubin in full production 2H 2026; 288 GB HBM4, ~13 TB/s, NVLink 6. <https://developer.nvidia.com/blog/nvidia-vera-rubin-pod-seven-chips-five-rack-scale-systems-one-ai-supercomputer/>
 - [511] Tom’s Hardware (17 Mar. 2026), reporting the GTC 2026 keynote: Feynman in 2028 with the Rosa CPU, LP40 inference processor, BlueField-5 and CX10, and Kyber racks in both copper and co-packaged-optical scale-up; NVIDIA’s headline metrics stated as tokens per watt and cost per token rather than FLOPS. <https://www.tomshardware.com/pc-components/gpus/nvidia-updates-data-center-roadmap-with-ros-a-cpu-and-stacked-feynman-gpus-optical-nvlink-groq-lpus-with-nvfp4-and-nvlink-also-on-deck>
 - [512] CNBC and Tom’s Hardware (6 Jul. 2026), reporting SemiAnalysis: the Kyber NVL144 rack for Rubin Ultra slips from 2027 to 2028 on a 78-layer orthogonal-midplane signal-integrity problem; the NVL72×2 stopgap cancelled after customer pushback; NVIDIA responds only that “our roadmap is intact.” <https://www.cnbc.com/2026/07/06/nvidia-kyber-rack-system-delays-manufacturing-taiwan-rubin-chips-.html>
 - [513] SemiAnalysis (24 Mar. 2026), GTC 2026 analysis: attention–FFN disaggregation, KV-cache-bounded attention utilization, and new storage tiers treating KV cache as a first-order constraint. <https://newsletter.semianalysis.com/p/nvidia-the-inference-kingdom-expands>
 - [514] Grunewald and Akula, AI Frontiers (2 Feb. 2026), on HBM as the binding constraint under export controls, with CXMT at HBM2-class and early HBM3 sampling; read together with SemiAnalysis on the Ascend ramp being memory-bound rather than logic-bound. Note that die counts and stack counts differ by roughly an order of magnitude and must not be compared directly. <https://ai-frontiers.org/articles/high-bandwidth-memory-critical-gaps-us-export-controls>
 - [515] Analyst reaction to the financing platforms, 11–12 Aug. 2026: F. Wang (Hedgeye)—“in effect, they made Nvidia’s product cheaper without really cutting GPU prices”; J. Moore (Morgan Stanley) on third-party capital alleviating circularity; UBS and Seaport on circular financing; the telecom comparison (nine suppliers, ~\$25.6B of vendor financing by end-2000; 24 of the 30 largest listed carriers bankrupt) with the disanalogy that Lucent and Nortel held their paper on balance sheet. <https://fortune.com/2026/08/11/nvidia-500-billion-ai-financing/>
 - [516] Wall Street Journal via Reuters (14 Aug. 2026): NVIDIA scales its contemplated backstop of OpenAI’s Ohio campus back from roughly \$250B to under \$120B, covering only the first phase, after investors raised concerns about risk exposure tied to large financing commitments. <https://finance.yahoo.com/technology/ai/articles/nvidia-scales-back-250-billion-234356524.html>
 - [517] Reporting on NVIDIA’s neocloud utilization backstops (from ~1 Jul. 2026): a guaranteed utilization floor under deployed GPUs, with NVIDIA renting idle capacity at a predetermined rate in exchange for a share of cloud revenue; disclosed arrangements include Sharon AI (40,000 GB300s) and Firmus Technologies (up to 170,000 GPUs). Rate and revenue share undisclosed; NVIDIA declined to comment. <https://capacityglobal.com/news/nvidia-now-agrees-to-rent-back-unused-gpu-capacity-from-neocloud-operators-if-customer-demand-falls-short/>

- [518] M. Kirkwood, Tech Policy Press (18 Dec. 2025), on NVIDIA allocation practices, the UK CMA’s identification of an interconnected web of over 90 partnerships and strategic investments, and questions over preferential supply and pricing.
<https://www.techpolicy.press/nvidia-is-building-a-shield-of-concentrated-power/>
- [519] NVIDIA Cloud Partner requirements: partners must expose newly delivered capacity through a centralized API and report delivered, healthy, reserved and in-use accelerators per tenant, with strict logical and physical isolation of NVIDIA-reserved capacity. No prohibition on resale or subletting of data-centre capacity is stated. <https://docs.nvidia.com/dsx/ncp/nvidia-requirements-for-ai-clouds/home#capacity-and-fleet-management>
- [520] Coverage of OpenAI’s August 2026 executive departures: B. Lightcap (COO 2022–spring 2026, latterly special projects) announced his departure 11 Aug.; D. Dresser, first chief revenue officer, announced hers 13 Aug. after ~8 months, with D. Rajic named successor; CFO S. Friar briefed investors 14 Aug. reporting an annualized run rate near \$40B and enterprise revenue overtaking consumer. <https://techcrunch.com/2026/08/11/brad-lightcap-openais-longtime-coo-is-leaving-to-start-something-new/>
- [521] OpenAI publicly disclosed on 8 Jun. 2026 that it had confidentially submitted a draft registration statement (no public S-1 filed; no listing date set): “We recently submitted a confidential S-1. We expect it to leak so we’re just announcing it. We have not decided on timing yet.” Valuation cited at ~\$852B. <https://www.cbsnews.com/news/openai-files-confidential-initial-public-offering/>
- [522] CME Group and Silicon Data (11 Aug. 2026): compute futures to launch 5 Oct. 2026, listed on NYMEX, each contract representing one month of GPU rental cost referencing H100 and B200. P. Keavey (CME): “Compute has become the currency of the AI age.”
https://www.cmegroup.com/media-room/press-releases/2026/8/11/cme_group_and_silicondata_launch_compute_futures_on_october_5_to_nloc.html
- [523] Intercontinental Exchange and Ornn (19 May 2026): cash-settled GPU compute futures on the Ornn Compute Price Index, built exclusively from printed transactions across H100, H200, B200 and RTX 5090. <https://ir.theice.com/press/news-details/2026/ICE-and-Ornn-to-Launch-GPU-Compute-Futures-Contracts/default.aspx>
- [524] ICE and NATIVX (1 Jul. 2026): energy-normalized compute futures on the COIL index, which “strips out the noise introduced by regional power cost disparities,” with sub-indices COIL-T (training), COIL-I (inference), COIL-G (graphics) and COIL-CO (connectivity).
<https://www.businesswire.com/news/home/20260701837320/en/ICE-and-NATIVX-to-Launch-Energy-Normalized-Compute-Futures-Contracts>
- [525] Kalshi (14 Jul. 2026): compute forward curves for B200, H200 and A100 derived from its own weekly and monthly GPU price contracts. T. Mansour: “Compute is the new oil. Like every commodity before it, it needs a real derivatives market.” Volume and open interest undisclosed.
<https://news.kalshi.com/p/compute-forward-curves>
- [526] Architect Financial Technologies and Ornn (21 Jan. 2026): perpetual futures on GPU and DRAM prices. B. Harrison: “there is an urgent cross-industry need to establish standardized derivatives contracts and centralized order books for compute.” <https://www.prnewswire.com/news-releases/architect-financial-technologies-partners-with-compute-index-provider-ornn-to-launch-exchange-traded-futures-on-gpu-and-ram-prices-302666613.html>
- [527] Silicon Data: daily GPU rental indices since 20 May 2025 built on ~3.5 million pricing observations, distributed via Bloomberg (ticker SDH100RT); observed dispersion of ~\$2.74/hr (neocloud) against ~\$7.20/hr (hyperscaler) for H100; an LLM token-expenditure index published alongside; forward curves in backwardation (B200 spot ~\$5.62 vs ~\$5.17 at 36 months). Index-provider data, self-interested. <https://www.silicondata.com/products/silicon-index>
- [528] D. Thomas, Benzinga (13 Aug. 2026), reporting analysis by R. Goswami (Semafor) on the CME compute-futures launch: the steel-benchmark rebuttal to the heterogeneity objection, and the market-structure objection—“there’s no other game in town except for Nvidia,” an oil market “run entirely by Standard Oil”; NVIDIA “has little incentive to encourage greater price transparency or falling compute costs.”
<https://finance.yahoo.com/technology/ai/articles/why-nvidia-may-stand-way-153004691.html>
- [529] C. Shack and A. Levin, Variant Fund (14 May 2026): five preconditions for a compute futures market, of which supply-side fragmentation and a standardized tradable unit are absent; the top

- four hyperscalers hold $\sim 78\%$ of global self-built critical IT power and 69% of H100 supply; inference is more fungible than training and projected above 65% of AI compute demand by 2029. <https://variant.fund/articles/compute-as-commodity-framework-gpu-futures-markets-emerge/>
- [530] D. Friedman (17 Jul. 2026), “Compute futures won’t make compute fungible”: futures “make its non-fungibility legible”—“the benchmark doesn’t eliminate differences within the physical market. It gives those differences something to be priced against.” Five basis categories: hardware, configuration/topology, location, contractual term, counterparty. <https://davefriedman.substack.com/p/compute-futures-wont-make-compute>
- [531] China’s computing-power trading platforms: Shanghai (trial operation 26 Apr. 2023, via SHIXP) and Yinchuan, Ningxia (24 Feb. 2023, co-built with the Beijing International Big Data Exchange; 27 participating institutions including Huawei, Alibaba Cloud, SenseTime, Baidu); CAICT’s national China Computing Power Platform under MIIT with 1,000+ registered enterprise users, 100+ service providers and 10 regional sub-platforms. None publishes standardized contracts, clearing, or transaction volumes. <https://www.yicai.com/news/china-debuts-first-computing-power-trading-hub>
- [532] People’s Daily (5 Jun. 2026): 1.37 million PFLOPS of intelligent computing under national monitoring, $\sim 72\%$ of the national total, with 80%+ of intelligent computing inside the eight national hubs; the 15th Five-Year Plan extends East Data West Computing; standardized interconnection targeted for 2028. <https://en.people.cn/n3/2026/0605/c90000-20464201.html>
- [533] TechNode (17 Apr. 2026), reporting Mao Shengyong, deputy head of the National Bureau of Statistics, at a State Council Information Office briefing: China’s daily token usage exceeded 140 trillion in March 2026, up over 40% from end-2025. <https://technode.com/2026/04/17/china-authority-says-daily-ai-token-usage-exceeds-140-trillion-in-march-up-over-40-vs-end-2025/>
- [534] Comparative token pricing, mid-2026: DeepSeek V4-Pro at \$0.435/\$0.87 per Mtok against GPT-5.6 Sol at \$5/\$30 and Claude Opus 4.8 at \$5/\$25 ($\sim 34\times$ on output); DeepSeek V4-Flash at \$0.14/\$0.28; Qwen3-8B at \$0.01/\$0.01. <https://www.cgphbanquedaffaires.com/post/china-has-found-its-next-commodity-to-export-the-token>
- [535] South China Morning Post (10 Aug. 2026), citing Jefferies and Silicon Data: blended enterprise inference cost fell to \$1.16–1.18 per million tokens on 6–8 Aug. 2026, a 2026 low, from \$1.45 in late July and \$2.04 on 31 May, driven by price wars and Chinese open-source models. <https://www.scmp.com/tech/tech-trends/article/3363549/enterprise-ai-costs-hit-2026-low-drive-n-price-wars-chinese-open-source-models-research>
- [536] China Daily (1 Jun. 2026), “Tokens are China’s next goldmine”: the export model sells inference by API while electricity, servers and models remain domestic. Liu Xiaoge, State Council Development Research Center: token exports “provide a new way to offset risks of traditional trade frictions and drive growth in services trade.” <https://global.chinadaily.com.cn/a/202606/01/WS6a1cc2c4a310d6866eb4baae.html>
- [537] Reuters (Jul. 2025) and ASPI Strategist (6 May 2026): utilization at Chinese AI data centres of 20–30% against a 60% target; 633 hyperscale facilities built under East Data West Computing by 2024; 100+ projects scrapped in 18 months; western-region capacity failing latency, fibre and talent tests for inference; NVIDIA/CUDA and Ascend/CANN estates that cannot be integrated into a single seamless cloud environment. <https://www.aspistrategist.org.au/abundant-electricity-isnt-enough-chinas-overbuilt-ai-computing-power-is-underused/>
- [538] Baiguan (9 Jul. 2026), on neutral-router economics: Chinese models at roughly 61% of token consumption while Anthropic holds $\sim 12\%$ of volume and captures $\sim 46\%$ of revenue; large enterprises still overwhelmingly buy from Anthropic, OpenAI, Azure and Google Cloud; China’s 2017 National Intelligence Law as a persistent procurement obstacle. <https://www.baiguan.news/p/china-ai-token-export-openrouter-market-share-chinese-models-performance-gap-cost-efficiency-semiconductor-energy-self-sufficiency-investment-opportunities>
- [539] SemiAnalysis, “AI Value Capture”: inference gross margins rising from under 40% to over 70% across 2025–26 while the compute layer commoditized, with the surplus concentrating at the model layer; neoclouds earning close to their minimum IRR hurdle at prevailing rental rates. <https://newsletter.semianalysis.com/p/ai-value-capture-the-shift-to-model>
- [540] S. Matsuoka, *Memory Scarcity, Open Models, and the Restructuring of the AI Industry, 2026–2030: A quantitative scenario analysis of inference economics, training-cost divergence, and infrastructure*

- solvency*, arXiv:2607.07207, 8 Jul. 2026. Entrant–incumbent cost gap $3.2\times$ (2026), $1.9\times$ (2027), re-widening to $3\text{--}4\times$ by 2029–30 via a depreciation conveyor; training bifurcating into an \$18–38B frontier tier and a $\sim$ \$5M distillation tier by 2030; solvency corridor requiring $\sim 2\times$ annual token-demand growth for four years; scenarios Rotating Landlord Oligopoly 25%, Commoditization Crash 25%, Jevons Absorption 20%, System-Layer Re-differentiation 18%, Geopolitical Bifurcation 12%. Author’s own companion analysis. <https://arxiv.org/abs/2607.07207>
- [541] Tom’s Hardware (24 Jul. 2025), the proximate English source for the “20 ms” claim: “Beijing established a target of a network latency of 20 milliseconds, which is necessary for real-time applications like financial services or high-frequency trading. However, many western-region data centers fail to meet this benchmark.” Note the internal inconsistency: the primary Chinese document assigns high-frequency trading to a separate 10 ms tier. <https://www.tomshardware.com/desktops/servers/china-is-developing-nation-spanning-network-to-sell-surplus-data-center-compute-power-latency-disparate-hardware-are-key-hurdles>
- [542] National Development and Reform Commission et al. (May 2021), *Implementation Plan for Computing Hubs in the National Integrated Big Data Centre Collaborative Innovation System*: end-to-end network latency “in principle within 20 milliseconds” *one-way* (danxiang) for the eastern hub clusters, serving industrial internet, finance and securities, disaster warning, telemedicine, video calling and AI inference; a separate 10 ms tier for city-edge facilities serving high-frequency trading, VR/AR, ultra-high-definition video and connected vehicles; western nodes assigned to “back-end processing, offline analysis, storage backup”—non-real-time demand. Predates ChatGPT by eighteen months.
- [543] State Council of the People’s Republic of China (29 Aug. 2024), quoting Liu Liehong, head of the National Data Administration: “Network latency between the eastern and western hubs has generally met the 20-millisecond (ms) requirement.” https://english.www.gov.cn/news/202408/29/content_WS66d03a1ac6d0868f4e8ea539.html
- [544] Published intra-China hub latencies: Gui’an to Chengdu/Chongqing 6 ms, to Guangzhou/Shenzhen 10 ms, to the Yangtze Delta and Beijing-Tianjin-Hebei within 18 ms; Zhongwei (Ningxia) to Beijing 8–10 ms one-way and to Shanghai within 15 ms one-way, on 18 Tbps of export bandwidth with direct fibre to 26 cities; Inner Mongolia to Beijing-Tianjin-Hebei within 5 ms. Chinese official and state-adjacent sources; no independent Western measurement of intra-China backbone latency was located. <https://news.qq.com/rain/a/20250217A07J1N00>
- [545] Global Times and TechNode (12 Aug. 2026): Alibaba Cloud’s Lingjun M890 supernode enters commercial service with first deployment in Ulanqab, Inner Mongolia, designed for *inference* on mixture-of-experts models up to ten trillion parameters, already serving live commercial traffic for Kimi K3 and Qwen3.8-Max. <https://www.globaltimes.cn/page/202608/1368093.shtml>
- [546] Alibaba Cloud Model Studio: published public endpoints in Beijing, Hong Kong, Singapore, Frankfurt, Tokyo and US (Virginia). <https://www.alibabacloud.com/help/en/model-studio/regions/>
- [547] Spheron, *LLM Inference SLO Engineering: TTFT, ITL, and P99 Latency Budgets* (15 May 2026): a 400 ms P99 TTFT budget allocating 20 ms to network ingress, 50 ms to queueing, 280 ms to prefill, 10 ms to tokenization, 15 ms to gateway and 25 ms to safety filtering; inline code completion at 100 ms P99 TTFT (“must feel local to IDE”); tightening a target from 500 ms to 200 ms raising infrastructure spend $\sim 35\%$. <https://www.spheron.network/blog/llm-inference-slo-ttft-itl-latency-budget-guide-2026/>
- [548] Digital Applied, *AI Model Latency Benchmarks 2026* (23 Apr. 2026), 10,000 samples per model and provider at 1,024-token input: GPT-5.5 TTFT P50 1.12 s at 92 output tok/s; Claude Opus 4.7 0.85 s; a reasoning model at medium effort 8.4 s; extended thinking 28 s; specialist inference hardware 0.16–0.18 s. <https://www.digitalapplied.com/blog/ai-model-latency-benchmarks-2026-ttft-throughput>
- [549] Measured AWS inter-region round-trip latency: Tokyo to Oregon 99 ms; Tokyo to Northern Virginia 149 ms; Seoul to Oregon 120 ms; Singapore to Oregon 173 ms; Frankfurt to Northern Virginia 92 ms. <https://latency.bluegoat.net/>
- [550] OpenAI Help Center, *Data residency and inference residency* (updated Aug. 2026): inference residency—“model inference (GPU execution) on in-scope customer content happens exclusively in-region”—available for Europe (EEA and Switzerland), the United States and the United Arab

- Emirates only. Australia, Canada, India, Japan, Singapore, South Korea and the United Kingdom receive data residency only.
<https://help.openai.com/en/articles/9903489-data-residency-and-inference-residency-for-chatgpt>
- [551] Anthropic platform documentation, data residency: inference geography defaults to `global`—“inference may run in any available geography for optimal performance and availability”—with only `us` and `global` available.
<https://platform.claude.com/docs/en/manage-claude/data-residency>
- [552] Microsoft Learn, deployment types in Foundry Models: Global Standard “use[s] Azure’s global infrastructure to dynamically route traffic to available datacenters”; “For most workloads, start with Global Standard... has the lowest price.”
<https://learn.microsoft.com/en-us/azure/foundry/foundry-models/concepts/deployment-types>
- [553] Google Cloud (28 Jul. 2025), global endpoint for Claude models on Vertex AI: “dynamically route your requests to any region with available capacity”; latency “may be higher in some cases”; customers with data-residency requirements should use regional endpoints because the global endpoint “does not guarantee that requests will be processed in any specific location.”
<https://cloud.google.com/blog/products/ai-machine-learning/global-endpoint-for-claude-models-generally-available-on-vertex-ai>
- [554] AWS (24 Feb. 2026): global cross-Region inference for Claude models with source regions Singapore, Jakarta, Taipei, Thailand and Malaysia, routing to more than twenty commercial regions worldwide while data at rest remains in the source region. <https://aws.amazon.com/blogs/machine-learning/global-cross-region-inference-for-latest-anthropic-claude-opus-sonnet-and-hai-ku-models-on-amazon-bedrock-in-thailand-malaysia-singapore-indonesia-and-taiwan>
- [555] Toniolo, Dinesh and Thorstenson, *GORGON: Maximizing KV-Cache Reuse While Minimizing Network Latency in Cross-Region LLM Load Balancing*, arXiv:2602.11688, Feb. 2026: measured inter-region latencies of 281 ms (US West to Germany) and 3 ms (intra-region), with a worked case in which routing to Germany produced a lower total TTFT (296 ms) than serving locally under load (305 ms); $2.5\times$ faster median TTFT than least-load routing. <https://arxiv.org/html/2602.11688v1>
- [556] R. Qin, W. He, Y. Wang et al. (Moonshot AI; Tsinghua University), *Prefill-as-a-Service: KVCache of Next-Generation Models Could Go Cross-Datacenter*, arXiv:2604.15039, 16 Apr. 2026: long-context prefill offloaded to standalone clusters with KV cache transferred over commodity Ethernet; ~ 170 Gbps for a 1T-parameter model at 32K context, below 100 Gbps at 128K; 54% higher serving throughput than homogeneous prefill-decode. <https://arxiv.org/html/2604.15039v1>
- [557] S. Ray, N. Feamster and J. Jiang (University of Chicago), *An Internet for the KV Cache*, arXiv:2608.01526, 2 Aug. 2026: transferring a 1 GB KV cache over a 10 Gbps link costs ~ 0.8 s and $\sim \$0.09$ against ~ 15 s and $\$0.47$ to recompute on $8\times B200$ — $19\times$ faster and $5.2\times$ cheaper to ship than to recompute. <https://arxiv.org/html/2608.01526v1>
- [558] Zhong, Liu, Chen et al., *DistServe: Disaggregating Prefill and Decoding for Goodput-Optimized Large Language Model Serving*, arXiv:2401.09670 (OSDI ’24). Table 1 service-level objectives: code completion TTFT 125 ms; chatbot 250 ms–2.5 s; summarization 15 s. KV-cache transmission under 0.1% of total latency. <https://arxiv.org/pdf/2401.09670>
- [559] Published real-time voice latency budgets (2026): ~ 500 – 600 ms total for conversational speech, of which 30–80 ms is network; measured production voice-to-voice medians of 680–850 ms; a co-located stack reducing network and SIP from 150 ms to 45 ms and total from 1,210 ms to 450 ms.
<https://thepromptbench.com/voice-and-realtime/latency-budgets-for-realtime-voice/>
- [560] Anthropic (13 Mar. 2025): prompt caching reduces “costs by up to 90% and latency by up to 85% for long prompts.” OpenAI’s documentation notes that caching works by routing to servers that recently processed the same prompt—creating region stickiness independent of propagation delay.
<https://claude.com/blog/token-saving-updates>
- [561] Equinix (30 Apr. 2025), on retrieval-augmented generation: the inference platform must be “tightly coupled to the data platform” so that retrieval reaches enterprise data over low-latency connectivity—an argument about proximity to the corpus, not to the user. <https://blog.equinix.com/blog/2025/04/30/what-is-retrieval-augmented-generation-rag-and-where-should-you-do-it/>
- [562] Sina Finance (14 Aug. 2026), “China’s new computing map: training migrates west, inference sinks east”: inference-centric compute moves east because the eastern regions are the main point of

- explosion of AI inference *demand*. The article states no latency figure.
<https://finance.sina.com.cn/roll/2026-08-14/doc-ininhhrn4148288.shtml>
- [563] TrendForce contract-price series (2026): conventional DRAM +93–98% QoQ in 1Q26, +58–63% in 2Q26, +13–18% in 3Q26, following ~172% across 2025; DRAM industry revenue \$97B in 1Q26, +81% QoQ. <https://www.trendforce.com/presscenter/news/20260601-13070.html>
- [564] Street and spot memory pricing, 2026: 32 GB DDR5-6000 kits from ~\$85 (mid-2025) to \$369–432; 64 GB DDR4-2933 ECC RDIMMs from ~EUR 98 to EUR 500 used and EUR 835 new; DDR5 spot +307% Sep–Nov 2025. Trackers diverge by up to 50%; treat as a range.
<https://capitalandcompute.net/memory-prices/ddr5/>
- [565] Seoul Economic Daily (12 Jul. 2026): HBM3E at ~\$1.5–1.6 per Gb and HBM4 at ~\$2 per Gb in 2H26, forecast \$4–5+ per Gb in 2027; HBM4 taking 4–6 months to produce against 3–3.5 for DDR5, at low initial yields. <https://en.sedaily.com/finance/2026/07/12/hbm4-prices-to-double-n-ext-year-as-samsung-sk-hynix-keep>
- [566] Tom’s Hardware and MacRumors (5–6 Mar. 2026): Apple withdraws the 512 GB Mac Studio configuration and raises the 96 GB-to-256 GB upgrade from \$1,600 to \$2,000; 256 GB deliveries slip to May.
<https://www.tomshardware.com/tech-industry/apple-pulls-512-mac-studio-upgrade-option>
- [567] Moonshot AI, Kimi K3 model card: 2.8T total / 104B active, 16 of 896 experts routed per token, native MXFP4 from quantization-aware fine-tuning; published download ~1.4 TB against ~5.6 TB at FP16. Weights released 27 Jul. 2026 under the custom Kimi K3 licence.
<https://huggingface.co/moonshotai/Kimi-K3>
- [568] DeepSeek-AI, *Insights into DeepSeek-V3: Scaling Challenges and Reflections on Hardware for AI Architectures*, arXiv:2505.09343, May 2025: KV cache per token of 70.272 KB for DeepSeek-V3 multi-head latent attention against 327.680 KB for Qwen-2.5 72B and 516.096 KB for Llama-3.1 405B (both grouped-query); memory demand growing >1000% per year against HBM capacity under 50%. <https://arxiv.org/html/2505.09343v1>
- [569] SanDisk and SK hynix (3 Aug. 2026): first Open Compute Project technical specification for High Bandwidth Flash, following the August 2025 standardization memorandum and a February 2026 consortium launch. Consortium members SanDisk, SK hynix, Google and Tenstorrent; NVIDIA, AMD, Samsung, Micron and Kioxia absent. <https://www.sandisk.com/company/newsroom/press-releases/2026/2026-08-03-Sandisk-and-sk-hynix-advance-global-standardization-of-hbf>
- [570] HBF specification as reported at FMS 2026: up to 512 GB per stack in 8-Hi and 16-Hi NAND; three bandwidth grades from ~0.4 to 3.0 TB/s; UCIE host interface at up to 64 GT/s over 64 lanes; NAND dies stacked with TSVs in an HBM4-compatible footprint with a CBA logic die scheduling parallel sub-array reads. Evaluation samples to selected customers promised by end-2026.
<https://www.tomshardware.com/pc-components/ssds/sandisk-and-sk-hynix-unveil-hbf-spec-u-p-to-16-hi-nand-stacks-3-tb-s-bandwidth-ucie>
- [571] Contemporaneous criticism of the HBF specification (Aug. 2026): the OCP document carries no version number, no conformance procedure and no certification scheme, and omits latency, power consumption, write speed, error correction, endurance and cost; the stated bandwidth and capacity figures are the two vendors’ development targets rather than conformance requirements.
<https://www.igorslab.de/en/512-gb-up-to-3-tb-s-hbf-closes-gap-between-hbm-and-ssd/>
- [572] Chipstrat, *High Bandwidth Flash: the full report* (7 Jul. 2026): HBF latency in microseconds against HBM4’s ~100 ns; minimum read granularity ~128× larger; power ~2× HBM4; cost per gigabyte ~10× lower but cost per *stack* comparable because advanced packaging dominates; suitable for weights, unsuitable for KV cache, activations or training.
<https://www.chipstrat.com/p/high-bandwidth-flash-the-full-report>
- [573] SoftBank Corp. (3 Feb. 2026) and SAIMEMORY Corporation (22 Apr. 2026): SAIMEMORY incorporated December 2024 as a SoftBank subsidiary, joint venture with Intel; Series A from Fujitsu, the Development Bank of Japan, RIKEN and SoftBank; NEDO grant under the Post-5G / advanced semiconductor programme; prototype targeted for the fiscal year ending March 2028, commercialization FY2029.
https://www.softbank.jp/en/corp/news/press/sbkk/2026/20260203_01/

- [574] ZAM (Z-Angle Memory): stacked DRAM built on Intel’s Next Generation DRAM Bonding, validated under a US Department of Energy Advanced Memory Technology programme; a “via-in-one” TSV replacing HBM’s many TSVs to free die area; claimed $\sim 40\text{--}50\%$ power reduction against comparable HBM and a target of up to 512 GB per chip. <https://www.blocksandfiles.com/ai-ml/2026/02/04/intel-and-softbank-aim-zam-at-hbm/4090386>
- [575] Powerchip, Intel and SAIMEMORY at the VLSI Symposium (Jun. 2026): functional validation of a nine-layer fusion-bonded DRAM stack; $\sim 0.25\text{ Tb/s/mm}^2$ bandwidth density, $<0.7\text{ pJ/bit}$ data-movement energy, $\sim 3\text{ }\mu\text{m}$ substrates, $\sim 13.7\text{ k}$ TSVs per layer, 171 mm^2 footprint at $\sim 1.125\text{ GB}$ per layer—i.e. a demonstrated module of roughly 10 GB against the 512 GB target. <https://www.trendforce.com/news/2026/05/22/news-psmc-joins-intel-saimemory-to-demo-9-layer-fusion-bonded-via-in-one-architecture-for-high-bandwidth-3d-memory/>
- [576] JEDEC JESD209-6 (LPDDR6), published 9 Jul. 2025; Samsung has sampled LPDDR6X to Qualcomm for the AI250, with commercial production indicated for 2H27. The extension roadmap is under JEDEC standardization (see [585]).
- [577] Qualcomm (27 Oct. 2025): AI200 with 768 GB of LPDDR per card, direct liquid cooling, 160 kW racks, for 2026; AI250 adding near-memory compute with a claimed $>10\times$ effective bandwidth, for 2027. <https://www.qualcomm.com/news/releases/2026/qualcomm-ai200-ai250>
- [578] Kyung, Yun and Ahn (Seoul National University), *SSD Offloading for LLM Mixture-of-Experts Weights Considered Harmful in Energy Efficiency*, arXiv:2508.06978, Aug. 2025: flash reads at $\sim 102.4\text{ pJ/bit}$ against DRAM’s ~ 4.2 ; SSD offloading raising per-token energy $3.8\text{--}12.5\times$ for Mixtral and $4.7\text{--}9.8\times$ for DeepSeek-R1, with SSD access exceeding 73–80% of total energy; prefetching hides latency ($1.25\text{--}1.32\times$) but cannot hide energy; flash requires $\sim 10\times$ improvement to $\sim 10\text{ pJ/bit}$ to compete. <https://arxiv.org/html/2508.06978v1>
- [579] SK hynix chief executive, 11 Jul. 2026, on the day of the company’s Nasdaq listing: “We forecast that next year will be the worst year in the industry’s history from the supply perspective,” and “we still forecast that customer demand will remain higher than our supply capacity even beyond 2030.” <https://www.tomshardware.com/pc-components/dram/sk-hynix-says-2027-will-be-the-worst-year-for-memory-shortage-forecasts-crunch-to-last-until-2030-ceo-shares-grim-outlook-on-the-day-sk-hynix-gets-listed-on-nasdaq>
- [580] TrendForce (21 Jul. 2026): NAND flash in a 4–5% supply deficit for 2026 with the balance turning positive in 2H27, against server DRAM contract prices forecast to rise every quarter from 2H26 through 2H27 on RDIMM bit-supply growth of only 15–20% year over year. <https://www.trendforce.com/presscenter/news/20260721-13148.html>
- [581] NVIDIA Rubin CPX—announced 9 Sep. 2025 as a capacity-first inference GPU with 128 GB of GDDR7 rather than HBM—removed from the roadmap at GTC 2026, with supply-chain reporting indicating no memory or substrate orders placed and substrate suppliers treating it as effectively cancelled; NVIDIA indicating the idea may return with Feynman. <https://wccfttech.com/nvidia-rubin-cpx-is-off-the-roadmap-but-may-return-with-feynman-in-2028/>
- [582] NVIDIA acquires Groq’s assets and acquires its leadership, announced 25 Dec. 2025 (initial reporting described the transaction as a technology licence plus hiring; later trade coverage and NVIDIA’s own product integration describe an asset purchase with an acquire, structured to avoid extended regulatory review); reported at $\sim \$20\text{B}$, the company’s largest transaction on record. Groq’s architecture is SRAM-based—roughly 500 MB per chip at $\sim 150\text{ TB/s}$ —i.e. the opposite of capacity-first. <https://www.datacenterdynamics.com/en/news/nvidia-to-license-tech-from-ai-inference-chip-company-groq-hire-its-leadership/>
- [583] Menlo Ventures enterprise LLM market surveys: open-weight share of enterprise workloads falling from 19% (late 2024) to 13% (mid-2025) to $\sim 11\%$ (2026), while enterprise LLM API spend rose from $\$3.5\text{B}$ to $\$8.4\text{B}$; constraints reported as licensing terms, governance and provenance, support and total cost of ownership rather than hardware. <https://menlovc.com/perspective/2025-mid-year-llm-market-update/>
- [584] Counterpoint Memory & Storage Tracker (12 Aug. 2026): YMTC enters the top three NAND makers by shipment share for the first time at 14%, displacing Kioxia, behind Samsung (25%) and SK hynix (22%); enterprise SSDs absorbing 48% of NAND bits in Q2 2026, up from 26% a year earlier. YMTC’s Xtacking is wafer-to-wafer hybrid bonding, in mass production since 2018—the

- process HBF requires. <https://www.tomshardware.com/tech-industry/ymtc-breaks-into-the-top-three-nand-makers-for-the-first-time>
- [585] JEDEC (22 Apr. 2026), “JEDEC Previews LPDDR6 Roadmap Expanding LPDDR into Data Centers and Processing-in-Memory”: narrower per-die interface enabling higher capacities, “512 GB density on the horizon,” an LPDDR6 PIM standard in development, and an “LPDDR6-based SOCAMM2 module standard in development” with “a clear upgrade path from today’s LPDDR5X SOCAMM2 modules.” <https://www.jedec.org/news/pressreleases/jedec%C2%AE-previews-lpddr6-roadmap-expanding-lpddr-data-centers-and-processing-memory>
- [586] SanDisk, *High Bandwidth Flash Fact Sheet* (Jul. 2025): 512 GB per 16-die stack at 256 Gb per die; 1.6 TB/s read bandwidth in generation 1; generation 2 “set to reach up to 1 TB” at over 2 TB/s and $0.8\times$ the power; generation 3 at 1.5 TB and 3.2 TB/s at $0.64\times$; “non-volatile, no refresh power needed”; “one of the lowest cost per bit of any memory technology”; “similar cost” to HBM per stack; higher latency than HBM acknowledged and unquantified. https://documents.sandisk.com/content/dam/asset-library/en_us/assets/public/sandisk/collateral/company/Sandisk-HBF-Fact-Sheet.pdf
- [587] Moonshot AI, *Kimi K3* technical report, arXiv:2607.24653 (Jul. 2026), and the published model configuration: 93 layers (69 Kimi Delta Attention + 24 MLA), 896 experts with 16 routed and 2 shared, MXFP4 experts; MLA KV cache ~ 27 KB per token at BF16 (kv_lora_rank 512 + rope 64), KDA state ~ 207 MB per sequence. <https://arxiv.org/pdf/2607.24653>
- [588] vLLM project blog (27 Jul. 2026), Kimi K3 deployment guide: “2K+ tokens/GPU-s” at the high-throughput operating point on GB300 NVL72; single-user 111–118 tok/s at TP8–16, 331–370 with speculative decoding; minimum deployment $8\times B300$ or $16\times B200$. <https://vllm.ai/blog/2026-07-27-k3>
- [589] SemiAnalysis InferenceX (19 May 2026): Kimi K2.5 (1T/32B) on B200 NVFP4 at TP4—3,203 tok/s per GPU at 32 tok/s per user, 3,879 at 90; H200 INT4 at 991–1,175. And GB200 NVL72 on DeepSeek-R1 FP4: 4,130 tok/s per GPU at 125 tok/s per user, peak 14,659 at ~ 18 . <https://inference.semianalysis.com/blog/b200-nvfp4-vs-h200-int4-kimi-k2-vllm-perf-per-dollar>
- [590] Morgan Stanley estimate for the Vera Rubin VR200 NVL72, reported May 2026: $\sim \$7.8$ M per rack, of which GPUs $\sim \$4.0$ M ($\sim \$55$ k each), memory $\sim \$2.0$ M (20.7 TB HBM4 plus 54 TB LPDDR5X), Vera CPUs $\sim \$180$ k—labelled a bill of materials but stated to be “representative of what cloud service providers pay, not what it costs Nvidia itself to build one” (PC Gamer). Tom’s Hardware (24 Mar. 2026): ODM quotes for VR200 NVL72 $\$5$ – 7 M without warranty, “as much as $\$8.8$ M”; NVIDIA-supplied trays $\sim 90\%$ of a system’s cost, integrator margins thin; GB200 NVL72 $\$2.8$ – 3.4 M and GB300 NVL72 $\$6$ – 6.5 M for comparison. Loop Capital (A. Baruah) put GB300 NVL72 at $\$3.7$ – 4.0 M and SemiAnalysis GB200 NVL72 at $\sim \$3.1$ M bare / $\$3.9$ M all-in—the dispersion is the point. <https://www.pcgamer.com/hardware/a-single-nvidia-vera-rubin-rack-is-estimated-to-cost-usd7-803-148-with-over-usd2-million-of-that-figure-spent-on-memory-alone/>; <https://www.tomshardware.com/tech-industry/artificial-intelligence/price-of-nvidias-vera-rubin-nvl72-racks-skyrockets-to-as-much-as-usd8-8-million-apiece-but-server-makers-margins-will-be-tight-nvidia-is-moving-closer-to-shipping-entire-full-scale-systems>
- [591] VR200 NVL72 rack power estimates: ~ 190 kW at 1.8 kW per GPU (Max-Q) to ~ 230 kW at 2.3 kW (Max-P); 50 PFLOPS FP4 inference per GPU. <https://hashrateindex.com/blog/nvidia-vera-rubin-nvl72-specs-breakdown/>
- [592] NVIDIA DGX Spark: GB10 Grace Blackwell, 128 GB LPDDR5X at 273 GB/s, “up to 1 PFLOP FP4” (sparse); Founders Edition $\$3,999$ at launch (Oct. 2025), raised to $\$4,699$ on 27 Feb. 2026 citing memory shortages; ASUS Ascent GX10 partner unit $\$2,999$ at launch. <https://www.tomshardware.com/desktops/mini-pcs/nvidia-dgx-spark-gets-18-percent-price-increase-as-memory-shortages-bite-founders-edition-now-usd4-699-up-from-usd3-999>
- [593] Framework Desktop with AMD Ryzen AI Max+ 395 and 128 GB LPDDR5X, $\$1,999$ at announcement (Feb. 2025); GMKtec EVO-X2 with the same part at $\$1,799.99$ early-bird (May 2025). Post-surge: AMD’s own Ryzen AI Halo 128 GB desktop at $\$3,999$ (May 2026). <https://www.techspot.com/news/106934-framework-first-ever-desktop-pc-packs-amd-ryzen.html>
- [594] TrendForce memory spot update (16 Jul. 2025): 512Gb TLC NAND wafer spot $\$2.673$, i.e. $\sim \$0.042$ per GB; by 9 Aug. 2026 $\$21.125$, i.e. $\sim \$0.33$ per GB.

- <https://www.trendforce.com/news/2025/07/16/insights-memory-spot-price-update-prices-slightly-cool-off-after-recent-surge-with-ddr4-leading-the-drop>
- [595] LPDDR5X pricing: ~\$3.3 per GB in Q1 2025 (derived from SemiAnalysis’s “3× since Q1 2025 to ~\$10/GB” via Wccftch); TrendForce 2Q26 at \$19.3–19.8 per GB with long-term agreements signed at \$7.8–21; +58–63% QoQ in 1Q26 and +78–83% in 2Q26.
<https://www.trendforce.com/presscenter/news/20260514-13044.html>
- [596] Advanced-packaging cost estimates (2025–26, citing JPMorgan, Morgan Stanley and TrendForce): CoWoS-S ~\$750 per H100-class chip (~2,500 mm²); CoWoS-L \$1,000–1,100 per B200-class chip; processed CoWoS wafer \$10–12k; interposer 50–70% of packaging cost.
<https://siliconanalysts.com/analysis/cowos-packaging-cost-chiplet-vs-monolithic-2026>
- [597] S. Matsuoka, *Datacenter Construction for State-of-the-Art GPU Machines: A survey with three reference architectures, explicit cost scopes, evidence-graded sources, P10/P50/P90 ranges, escalation to 2030, an independent Japanese reconstruction, and a costed lean reference design*, Rev. 5, RIKEN Center for Computational Science technical report, Aug. 2026 (all figures from public sources or the author’s modelling; distributed without restriction). Facility-only cost per watt of IT load, excluding land and IT, at 50 MW: architecture A (resilient inference/cloud, Tier-III-capable) \$11.9/W P50 with a \$8.5–17.7 P10–P90 band; B (AI-training-optimised, warm-water, chillerless) \$9.0/W; C (interruptible/phased) \$8.4/W; utility interconnection the widest single driver; large-load interconnection 2–7 years by market; discrete-block penalty for capacity built ahead of load. Author’s own companion survey.
- [598] Epoch AI, “Total cost of ownership of a one-gigawatt AI data center” (14 May 2026): \$38B capex, of which \$26B compute (servers and network) and \$12B construction (facility, land, utility works); annualized \$8.5B/yr, servers ~60%; opex \$0.9B/yr of which energy \$0.6B; IT life 5 years, facility 14. <https://epoch.ai/data-insights/ai-datacenter-cost-breakdown>
- [599] JLL, *2026 Global Data Center Outlook* (Jan. 2026): shell-and-core \$11.3M/MW global average (\$10.7M in 2025), AI technology fit-out up to \$25M/MW on top; after ten years ~65% of facilities have received significant renovation, after fifteen nearly all; tenant IT refreshed roughly every five years. <https://www.jll.com/en-us/insights/market-outlook/data-center-outlook>
- [600] Turner & Townsend, *Data Centre Construction Cost Index 2025*: Silicon Valley \$13.3/W, Tokyo \$15.2, New Jersey \$12.9, Phoenix \$9.8 (shell and MEP; excludes utility works, site works and IT); liquid-cooled premium 7–10% over air. <https://reports.turnerandtowntsend.com/data-centre-construction-cost-index-2025/data-centre-cost-trends>
- [601] Equinix Form 10-K FY2025 (Feb. 2026): buildings 12–60 years, core systems 3–40, leasehold improvements 12–40; Digital Realty Form 10-K FY2025 (Feb. 2026): buildings and improvements 5–39 years, machinery and equipment 7–15.
<https://www.sec.gov/Archives/edgar/data/1101239/000110123926000032/R52.htm>
- [602] KPMG, *Benchmarking CapEx and OpEx in the Global Data Centre Market 2026* (May 2026): 30 MW greenfield cloud facility, non-electric opex \$7.1–10.6M/yr across four European markets (~\$237–353 per kW-IT per year: engineering labour, facilities O&M, security, insurance, property tax, IT and network, cleaning), labour ~40% of opex; construction capex \$6.7–8.5M/MW.
<https://assets.kpmg.com/content/dam/kpmgsites/ie/pdf/insights/strategy/ie-global-data-centre-cost-benchmarks.pdf.coredownload.inline.pdf>
- [603] Schneider Electric White Paper 231, “Battery technology for data centers”: VRLA calendar life 4–6 years vs. Li-ion 10–15, VRLA replaced 2–3 times over a fifteen-year facility life; ASHRAE equipment service-life table (Applications Handbook, as reproduced): centrifugal chillers ~23 years, cooling towers 20, transformers 30, pumps 20; three-phase UPS modules commonly 10–12 years.
https://download.se.com/files?p_Doc_Ref=SPD_VAVR-A5VTU5_EN
- [604] IFR: Beignet Investor LLC (Meta/Blue Owl Hyperion joint venture) \$27.3B senior secured amortising 144A bond, 6.581% coupon, T+225 bp, priced 16 Oct. 2025, maturity May 2049, rated A+, PIMCO anchor. <https://www.ifre.com/ifr-awards/2327933/financing-package-blue-owl-capitalbeignet-investors-us27.3bn-23.6-year-bond>
- [605] CoreWeave investor releases: \$2.0B 9.25% senior notes due 2030 (closed 27 May 2025); \$1.75B 9.00% senior notes due Feb. 2031 at par (22 Jul. 2025); \$1.0B 9.75% senior notes due Oct. 2031 (16 Apr. 2026); \$8.5B delayed-draw term loan at SOFR+225 bp (31 Mar. 2026).

- <https://investors.coreweave.com/news/news-details/2026/CoreWeave-Announces-Pricing-of-100-0-million-of-9-750-Senior-Notes-due-2031/default.aspx>
- [606] CBRE, *North America Data Center Trends H2 2025* (26 Feb. 2026) and Northern Virginia market release (18 Mar. 2026): NoVA 10 MW+ requirements asking \$155–185 per kW per month, vacancy 0.5%, 1,102 MW absorbed in 2025; primary-market 250–500 kW average \$195.94/kW-mo (+6.5% y/y), wholesale \$196.25 (record); power passed through at cost by market convention. <https://www.cbre.com/insights/books/north-america-data-center-trends-h2-2025>
- [607] Association of American Railroads, *A Short History of American Freight Railroads*: mileage peaked in 1916; ~254,000 miles operated by 1917; more than 70,000 miles (~30% of the network) in receivership by 1937; Federal Railroad Administration, Freight Rail Overview (updated Feb. 2025): “almost 140,000 route miles.” <https://www.aar.org/wp-content/uploads/2018/05/AAR-Short-History-American-Freight-Railroads.pdf>
- [608] Office of Rail and Road, *Rail infrastructure and assets 2024–25*: route length open for traffic in Great Britain 15,747 km (9,785 miles) at 31 Mar. 2025; network ~18,000 route miles in 1963; the Beeching report proposed 5,000 route miles and 2,363 stations for closure (National Archives). <https://dataportal.orr.gov.uk/media/3a1mbkaz/rail-infrastructure-and-assets-2024-25.pdf>
- [609] Hokkaido Railway Company (JR Hokkaido), company profile and operating-area pages: 2,240.5 km of route at 1 Apr. 2026 (trunk 1,378.9, regional 861.6); JR Hokkaido, “Sections that are difficult for the company to maintain on its own” (18 Nov. 2016): 13 sections, 1,237.2 km, of a ~2,569 km network; JR Hokkaido press releases of 18 Mar. and 15 Apr. 2026 listing bus conversions 2019–2026 (Sekisho Shin-Yubari–Yubari, Sassho, Hidaka Mukawa–Samani, Rumoi in two stages, Nemuro Furano–Shintoku; ~312 km) and the eight remaining “yellow” sections (925.7 km) under discussion; the ~4,000 km JNR-era peak is a widely repeated figure (Traffic News, Oct. 2016) without an official series. https://www.jrhokkaido.co.jp/corporate/company/com_02.html
- [610] S. Daggett, *Railroad Reorganization* (Harvard, 1908): in 1893 alone more than 27,000 miles with almost \$2B of capitalization were taken over by the courts, and 40,503 miles were sold at foreclosure in 1894–98; W. N. Leonard, “The Decline of Railroad Consolidation,” *J. Econ. Hist.* 9 (1949): about one-fourth of all rail mileage in receivership after the Panic of 1893. <https://www.gutenberg.org/cache/epub/55397/pg55397-images.html>
- [611] DCD (27 Mar. 2025): TD Cowen reports Microsoft cancelled or deferred up to 2 GW of data-centre projects in the US and Europe over six months, “point[ing] to data center oversupply relative to its current demand forecast”; DCD (6 Mar. 2026): Oracle and OpenAI drop plans to expand the flagship Abilene Stargate site (600–800 MW), citing financing and OpenAI’s frequently changing demand forecasting, the 4.5 GW agreement remaining intact. <https://www.datacenterdynamics.com/en/news/oracleopenai-drop-plans-to-expand-flagship-abilene-stargate-site-meta-in-talks-to-pick-up-crusoe-capacity-with-nvidias-help/>
- [612] Toshiba Memory (now Kioxia), XL-FLASH storage-class memory introduced at Flash Memory Summit, Aug. 2019: SLC, ~5 μ s read latency, 128 Gb die; EE Times, “Flash stacks slowly rise” (7 Aug. 2019); Kioxia XL-FLASH CXL module samples Aug. 2026. <https://www.eetimes.com/flash-stacks-sloooowly-rise/>
- [613] ServeTheHome (Jul. 2022): Intel winds down the Optane business with a \$559M inventory impairment in Q2 2022 results. <https://www.servethehome.com/intel-optane-559m-impairment-with-q2-2022-wind-down/>
- [614] Tom’s Hardware (8 Aug. 2025): Samsung revives Z-NAND at FMS 2025 as a GPU-direct storage tier, claiming 15 $\times$ conventional NAND performance and 80% lower power. <https://www.tomshardware.com/pc-components/storage/samsungs-revived-z-nand-targets-15x-performance-increase-over-traditional-nand-once-a-competitor-to-intels-optane-z-nand-makes-a-play-for-ai-datacenters>
- [615] XenoSpectrum (31 Jul. 2026), “Sandisk HBF and the Optane lesson”: two-vendor standard risk; “excellent technical specifications alone don’t guarantee adoption.” <https://xenospectrum.com/en/sandisk-hbf-optane-standardization/>
- [616] SanDisk, “Scaling beyond the wall: inside Sandisk’s High Bandwidth Flash for AI” (6 Aug. 2025): CBA architecture, massively parallel read paths; Llama 3.1 405B weight-read simulation “within 2.2% of unlimited-capacity HBM”; samples 2H CY2026, first AI-inference devices early 2027; “it is bandwidth that matters, not latency.” <https://www.sandisk.com/company/newsroom/blogs/2025/scaling-beyond-the-wall-inside-sandisks-high-bandwidth-flash-for-ai>

- [617] Semiconductor Engineering, “Flash getting stacked: high-bandwidth version” (14 May 2026): SanDisk’s Cynthia Hsu on choosing OCP over JEDEC (“workstreams are highly goal-oriented ... iterate on the specification in real-time”); roadmap Gen. 1/2/3 = $1\times/1.5\times/2\times$ capacity, $1\times/1.45\times/2\times$ bandwidth, $1\times/0.8\times/0.64\times$ energy; 17-die stack (16 core + base); Jim Handy: “Most of the speed you lose is in the write cycle.” <https://semiengineering.com/flash-getting-stacked-high-bandwidth-version/>
- [618] Tom’s Hardware, A. Shilov (13 Feb. 2025): SanDisk’s HBF at its 2025 Investor Day—512 GB per 16-high stack of 256 Gb dies, up to 4 TB “VRAM” per GPU, “match the bandwidth of HBM ... 8 to 16 times capacity at a similar cost point”; “minor protocol changes” from HBM; caveats on latency, endurance, block granularity and thermal stress near a GPU. <https://www.tomshardware.com/pc-components/dram/sandisks-new-hbf-memory-enables-up-to-4tb-of-vram-on-gpus-matches-hbm-bandwidth-at-higher-capacity>
- [619] Blocks & Files (12 and 25 Feb. 2025): SanDisk investor-day roadmap; Ilkbahar’s edge framing of a 512 GB HBF die subdivided into eight 64 GB mini-arrays for a mixture-of-experts model on a smartphone. <https://blocksandfiles.com/2025/02/12/sandisk-spills-its-technology-futures-beans/>
- [620] SanDisk (BusinessWire, 24 Jul. 2025): HBF Technical Advisory Board formed—David Patterson (chair; Pardee Professor Emeritus, UC Berkeley; Google Distinguished Engineer) and Raja Koduri (Oxmig Labs; formerly AMD and Intel); Patterson: HBF “shows promise ... delivering unprecedented memory capacity at high bandwidth.” <https://www.businesswire.com/news/home/20250724968078/en/Sandisk-Forms-HBF-Technical-Advisory-Board-to-Guide-Development-and-Strategy-for-High-Bandwidth-Flash-Memory-Technology>
- [621] SanDisk investor relations (6 Aug. 2025): SanDisk and SK hynix sign memorandum of understanding at FMS 2025 “to establish specifications for High Bandwidth Flash”; HBF named FMS 2025 Best of Show, Most Innovative Technology. <https://investor.sandisk.com/news-releases/news-release-details/sandisk-collaborate-sk-hynix-drive-standardization-high>
- [622] Kioxia (20 Aug. 2025): prototype 5 TB, 64 GB/s high-bandwidth flash memory module—PCIe 6.0 x8, daisy-chained controllers over 128 Gbps PAM4, under 40 W, developed under NEDO’s Post-5G programme. <https://apac.kioxia.com/en-apac/about/news/2025/20250820-1.html>
- [623] SK hynix Newsroom (27 Oct. 2025): next-generation NAND strategy at OCP Global Summit 2025—the “AIN” family, of which AIN B “leverages HBF technology by vertically stacking multiple NAND dies”; joint “HBF Night” with SanDisk. <https://news.skhynix.com/sk-hynix-presents-next-generation-nand-storage-product-strategy-at-ocp-2025/>
- [624] Blocks & Files (27 Nov. 2025), “High bandwidth flash is years away despite its promise”: two or more years from commercial availability; NVIDIA involvement essential; a hypothetical 16-Hi stack of 321-layer dies is 5,136 layers. <https://www.blocksandfiles.com/flash/2025/11/27/high-bandwidth-flash-is-years-away-despite-its-promise/1713383>
- [625] DigiTimes (31 Dec. 2025), “YMTC goes full speed toward HBF next”; YMTC’s Fan Zengxu at the Xincheng developer conference (26 Dec. 2025) on HBF fitting Xtacking and edge-AI inference. <https://www.digitimes.com/news/a20251231PD217/ymtc-3d-hbm-flash-technology.html>
- [626] X. Ma and D. Patterson (Google DeepMind), “Challenges and Research Directions for LLM Inference Hardware,” arXiv:2601.05047 (8 Jan. 2026; IEEE Computer, May 2026): proposes High Bandwidth Flash for “ $10\times$ memory capacity with HBM-like bandwidth”; ballpark stack 512 GB, 1,638 GB/s read, under 80 W, microsecond latency, 4 KB read granule, low write endurance. <https://arxiv.org/abs/2601.05047>
- [627] Ha, Kim and Kim (SK hynix), “H³: Hybrid Architecture Using HBM and HBF for Cost-Efficient LLM Inference,” IEEE Computer Architecture Letters 25(1), Jan. 2026 (DOI 10.1109/LCA.2026.3660969): B200-class package with 8 HBM3E and 8 HBF stacks; HBF $\sim 16\times$ HBM capacity, similar bandwidth, microsecond latency, $\sim 100k$ write cycles, up to $4\times$ HBM power; 40 MB SRAM latency-hiding buffer needing $>80\%$ hit rate; $1.25\times$ throughput at 1M context, $6.14\times$ at 10M, $18.8\times$ batch, $2.69\times$ performance per watt; “suitable for gigantic read-only use cases.” <https://ieeexplore.ieee.org/document/11371745>
- [628] SanDisk (25 Feb. 2026): SanDisk and SK hynix begin global standardization of HBF through a dedicated workstream under the Open Compute Project. <https://www.sandisk.com/company/newsroom/press-releases/2026/2026-02-25-sandisk-and-sk-hynix-begin-global-standardization-of-next-generation-memory-solution-high-bandwidth-flash-hbf>

- [629] Blocks & Files (20 Mar. 2026): Kioxia Europe CTO Axel Stoermann on Kioxia’s own high-bandwidth flash—up to 32 TSV-stacked dies on a shared interposer beside HBM, “with Nvidia”; XL-Flash the primary candidate cell; “The GPU is next to the HBM and then next to HBF, all connected to one and the same interposer.” <https://www.blocksandfiles.com/flash/2026/03/20/kioxia-eyes-gpu-memory-stack-with-high-bandwidth-flash-push/5210520>
- [630] SanDisk Q3 FY2026 earnings call (30 Apr. 2026), transcript: “still on the timeline . . . of having the NAND late this year, and looking for more of a system with the controller early to mid next year”; HBF “not a substitute for an enterprise SSD.” <https://www.fool.com/earnings/call-transcripts/2026/04/30/sandisk-sndk-q3-2026-earnings-transcript/>
- [631] TrendForce (5 and 22 Jun. 2026): Hanmi Semiconductor to deliver first HBF thermo-compression bonder in 2H26, with Hanwha Semitech, ASMPT and K&S developing HBF tools; SanDisk US patent 12,430,274 B2 bonding a processor onto a CBA NAND tile with HBM stacks on a shared interposer. <https://www.trendforce.com/news/2026/06/05/news-hbf-opens-new-equipment-race-hanmi-semiconductor-reportedly-eyes-first-tc-bonder-deliveries-in-2h26/>
- [632] Petrucci, Zacarias and Tanna (Micron), “Is High-Bandwidth Flash All You Need?” HotInfra Workshop at ISCA 2026 (28 Jun. 2026): roofline of tier-2 options—HBM 3.0 pJ/bit; on-package HBF 4 TiB, 8 TB/s nominal, 20 μ s, 8.5 pJ/bit, 36% scattered-read efficiency; on-package LPDDR 2 TiB, 2 TB/s, 2.4 pJ/bit; on DeepSeek-V4-Pro, Kimi K2.6 and Nemotron-3-Ultra at 1M context, on-package LPDDR beats HBF at every operating point (up to $2.1\times$ on KV-heavy workloads). <https://hotinfra.org/2026/papers/hotinfra26-final83.pdf>
- [633] Wang, Xue, Sun, Zhang, Dou, Li and Chen (ICT-CAS / UCAS / IME-CAS), “FlashAccel: Leveraging High-Bandwidth Flash for High-Throughput LLM Inference,” arXiv:2607.10186 (11 Jul. 2026): HBF stack modelled at 768 GB/s, 192 GB, SLC, 4 μ s tR, 4 KB page, 8 pJ/bit read, 1M P/E; six stacks plus HBM on an H200-class node give $2.54\times$ throughput per GPU and $1.93\times$ energy efficiency at a 100 ms SLO on Qwen3-235B/480B, DeepSeek-V3 and Llama-3.1-405B, using ~ 19 MB hyper-page reads and SRAM prefetch. <https://arxiv.org/abs/2607.10186>
- [634] Blocks & Files (5 Aug. 2026), “Getting flash up close and personal to GPUs”: the OCP HBF Architecture Specification v0.7.0 (High-Level Base Die Specification)—grades 0.384 / 1.536 / 3.072 TB/s; up to 16 independent full-duplex host channels per stack, 1–4 AXI channels each; each host channel connected to an independent set of NAND dies; developed in about six months. <https://www.blocksandfiles.com/flash/2026/08/05/getting-flash-up-close-and-personal-to-gpus/5283327>
- [635] Tom’s Hardware (4 Aug. 2026): SK hynix and SanDisk unveil the HBF specification at FMS 2026—8-Hi and 16-Hi to 512 GB, ~ 0.4 –3.0 TB/s, UCIE up to 64 GT/s over 64 lanes; the 3 TB/s grade exceeds a single HBM4 stack (~ 2 TB/s); AMD, Broadcom, Intel, NVIDIA, Marvell, Micron, Qualcomm and Samsung “show no stated interest.” <https://www.tomshardware.com/pc-component/s/ssds/sandisk-and-sk-hynix-unveil-hbf-spec-up-to-16-hi-nand-stacks-3-tb-s-bandwidth-ucie>
- [636] Samsung Newsroom (5 Aug. 2026): next-generation 3D memory vision at FMS 2026—zHBM (DRAM stacked atop the accelerator) and zNAND-O (four to eight V-NAND dies on a TSV base die with UCIE to an on-device NPU, targeted 2028); Samsung distinguishes HBF as “memory semantic tiering next to GPUs in servers.” Kioxia stays outside the OCP HBF effort (Technologies.org, Aug. 2026). <https://news.samsung.com/global/samsung-unveils-next-gen-3d-memory-vision-at-fms-2026-charting-the-future-of-ai-infrastructure>
- [637] SanDisk 2026 Investor Day deck and press release (13 Aug. 2026): “First HBF memory die taped out”; “first HBF inference product samples . . . 2027”; Technical Advisory Board slide adds Jim Keller (Tenstorrent); simulation of 4 TB HBF at 12.8 TB/s per GPU against 192 GB HBM—“4 HBF GPUs = 8 HBM GPUs”; “100B+ models on your edge device.” Reuters: Ilkbarah on early samples “next year.” <https://investor.sandisk.com/news-releases/news-release-details/sandisk-details-growth-strategy-and-long-term-financial-model>
- [638] TrendForce (14 Aug. 2026), “Sandisk reportedly tapes out first HBF product, targets 2027 samples and 2028 production”: Mizuho and Goldman commentary; Goldman—HBF-only configuration could halve GPU count for the same tokens; Cantor \$2,900 and Goldman \$2,200 price targets; Mizuho—“unlikely to fully replace HBM.” <https://www.trendforce.com/news/2026/08/14/news-sandisk-reportedly-tapes-out-first-hbf-product-targets-2027-samples-and-2028-production/>

- [639] Counterpoint Research (Aug. 2026), “Sandisk Investor Day: caching out the NAND cycle with contracts”: HBF “a longer-term play,” “unproven for AI workloads,” absent from SanDisk’s FY2028–30 financial model—“if HBF lands, it is upside”; Wedbush (via TheStreet) not building HBF into estimates “until adoption becomes clearer”; YMTC flagged as a competitive threat. <https://counterpointresearch.com/en/insights/sandisk-investor-day-caching-out-the-nand-cycle-with-contracts>
- [640] XenoSpectrum (15 Aug. 2026): critique of the investor-day slides—12.8 TB/s HBF baseline set against HBM3E-era parts while HBM4E is ~ 4.1 TB/s per stack; Qwen3-480B sized at 960 GB in BF16 (480 GB FP8, 240 GB FP4); no endurance specification; cites arXiv “A Full-Stack Characterization of High-Bandwidth Flash for KV-Centric LLM Serving” (Peking Univ., Aug. 2026). Wccftch carried the same critique. <https://xenospectrum.com/en/sandisk-investor-day-hbf-criticism-hbm-bandwidth-quantization-endurance/>
- [641] Z. Li et al. (Peking University), “A Full-Stack Characterization of High-Bandwidth Flash for KV-Centric LLM Serving,” arXiv, Aug. 2026, as reported by XenoSpectrum and NewMaxx (16 Aug. 2026): HBF used for KV cache raises latency $2\text{--}5.5\times$ and lowers throughput $1.1\text{--}2.7\times$; the paper itself was not retrievable at this edition. <https://borecraft.com/2026/08/16/hbf-sucks-a-full-stack-characterization-of-high-bandwidth-flash-for-kv-centric-llm-serving/>
- [642] Luo, Ghose, Cai, Haratsch and Mutlu, “HeatWatch: Improving 3D NAND Flash Memory Device Reliability by Exploiting Self-Recovery and Temperature Awareness,” HPCA 2018: retention loss in 3D charge-trap NAND accelerates superlinearly with temperature; $70\rightarrow 0^\circ\text{C}$ slows retention loss 58%; lifetime improvement $3.85\times$ with temperature-aware policies. https://www.cs.cmu.edu/~yixinluo/index_files/heatwatch_hpca18.pdf
- [643] Chun et al. (Seoul National University), “STRAW: A Stress-Aware WL-Based Read Reclaim Technique for High-Density NAND Flash-Based SSDs,” IEEE Computer Architecture Letters 2024 / arXiv:2501.02517: read disturb is strongly word-line-heterogeneous; stress-aware reclaim reduces reclaim writes by 83.6%. <https://arxiv.org/abs/2501.02517>
- [644] BusinessKorea (7 Aug. 2026): FMS 2026 panel—Google DeepMind’s Xiaoyu Ma on HBF’s “high power efficiency” with capacity and bandwidth; SK hynix VP Eui-cheol Lim on HBF pairing with HBM in “reading-centric inference”; SanDisk VP Rajeev Nagabhirava acknowledging NAND degradation under frequent writes, a risk that “diminishes if HBF operates primarily in read-heavy roles.” <https://www.businesskorea.co.kr/news/articleView.html?idxno=274396>
- [645] TweakTown (6 Aug. 2025), citing Chosun: SK hynix HBM4 12-Hi supply price near US\$500 per stack, 60–70% above HBM3E’s $\sim \$300$; later reports near \$560. <https://www.tweaktown.com/news/106826/sk-hynix-dramatically-raises-next-gen-hbm4-supply-price-in-war-of-nerve-with-nvidia/index.html>
- [646] Korea Herald (4 Feb. 2026): SK hynix framing HBF as an HBM complement, first generation $16 \times 32\text{ GB} = 512\text{ GB}$, mass-production target 2027; Samsung positioning HBF inside a unified memory-storage tier; unattributed market projection of \$1B (2027) to \$12B (2030). <https://www.koreaherald.com/article/10669660>
- [647] Blocks & Files (19 Jan. 2026), “A window into HBF progress”: KAIST Prof. Kim Jung-ho—512 GB / 1.638 TB/s per HBF; three-tier HBM–HBF–networked SSD; claims Samsung and SanDisk aim at NVIDIA/AMD/Google integration late 2027 or early 2028; “HBF market could be bigger than HBM” by 2038. <https://www.blocksandfiles.com/ai-ml/2026/01/19/a-window-into-hbf-progress-engineering-prof-talks-of-decade-ahead/4090360>
- [648] Blocks & Files (8 Jan. 2026), “Sandisk and high bandwidth flash efforts”: NVIDIA “has not yet expressed any public interest”; “Get Nvidia on board . . . and it’s a done deal.” <https://www.blocksandfiles.com/ai-ml/2026/01/08/sandisk-and-high-bandwidth-flash-efforts/4090532>
- [649] The Register, T. Mann (30 Jul. 2026): HBF trades write endurance and microsecond latency for capacity; 256 Gb dies, 512 GB per 16-Hi, “ $14\times$ more than HBM4”; a proposed 3.12 TB package from two HBM and six HBF stacks; wear could make “that \$50,000-plus GPU of yours a paperweight”; the Optane precedent for proprietary memory. <https://www.theregister.com/storage/2026/07/30/gpus-could-explode-to-multiple-tb-with-new-storage-inspired-memory-tech/5281363>
- [650] IEEE Spectrum (14 Jul. 2026), “High Bandwidth Flash”: NAND die 4.8 GB/s, DDR5 DIMM 70.4 GB/s, HBM4E 3.6 TB/s per stack; SK hynix SVP Hoshik Kim on HBF reducing the number

- of accelerators required; Jim Handy: “Flash is atrociously slow for writes, but for reads, it can be coaxed to go pretty fast.” <https://spectrum.ieee.org/high-bandwidth-flash>
- [651] TechRadar (6 Aug. 2026): HBF “exists only on paper”; NVIDIA absent; commercialization late 2027–2028, meaningful datacentre demand around 2030; SK hynix gave no HBF production date. <https://www.techradar.com/pro/sk-hynix-and-sandisk-reveal-first-specs-for-high-bandwidth-flash-but-dont-get-too-excited-you-probably-wont-be-able-to-buy-it-anytime-soon>
- [652] TradeThePool (5 Aug. 2026), SanDisk Q4 earnings preview: HBF performance evidence “is simulation rather than silicon”; long-term agreements “untested through a real downturn.” <https://tradethepool.com/pre-market-prep/sandisk-stock-sndk-q4-earnings-preview-the-hbf-bet/>
- [653] Tom’s Hardware, HBM roadmaps for Micron, Samsung and SK hynix to HBM4 and beyond (2026): HBM4 36–48 GB per stack, JEDEC ceiling 64 GB; HBM4E samples May 2026, volume late 2027; per-stack bandwidth 3.6–4.1 TB/s. <https://www.tomshardware.com/tech-industry/semiconductors/hbm-roadmaps-for-micron-samsung-and-sk-hynix-to-hbm4-and-beyond>
- [654] In-flash and near-flash compute for LLM inference: Lee et al. (SNU), “AiF: Accelerating On-Device LLM Inference Using In-Flash Processing,” ISCA 2025 ($14.6\times$ over SSD offload); Deng et al., “KVNAND,” arXiv:2512.03608 (Dec. 2025); Hao et al., “NVLLM: 3D NAND-centric edge LLM,” arXiv:2604.25699 (2026, ~ 65 tok/s on a 30B model, $5.6\times$ energy reduction); Jang et al. (Korea Univ.), “Dissecting and Re-architecting 3D NAND Flash PIM Arrays,” arXiv:2511.12860 (Nov. 2025). <https://arxiv.org/abs/2604.25699>
- [655] Huxiu (4 Aug. 2026), Beijing consultancy analysis of HBF as a “lane-change” opportunity: China lacks high-end HBM production capability but YMTC’s Xtacking and hybrid-bonding strengths suit HBF; 294-layer part in production above 90% yield, $\sim 13\%$ unit share in Q1 2026, 200k to 300k wafers per month by 2027; hybrid-bonding licence to Samsung; constraint is tooling above ~ 300 layers under the entity list; NVIDIA the decider. <https://www.huxiu.com/article/4880566.html>
- [656] CrazyRich research memo (19 Feb. 2026), “Beyond HBM: the rise of HBF”: secondary cost estimates (HBF \$3–5/GB against HBM \$20–30/GB); reports YMTC “not pursuing HBF directly” and partnering CXMT on HBM. <https://crazyrich.ai/reports/memos/beyond-hbm-rise-of-hbf-ai-memory>
- [657] Son, Park, Cho, Ham, Mutlu, Lee, Kim and Park (POSTECH and ETH Zurich), “Exploring High-Bandwidth Flash for Modern LLM Inference: Opportunities and Challenges,” IEEE Computer Architecture Letters 25(2), Jan. 2026 (DOI 10.1109/LCA.2026.3705817): HBF as main GPU memory “can significantly improve batch size, throughput and flexibility” and cut GPU count, “but only if HBM-equivalent read performance and substantial endurance improvement are achieved.” <https://doi.org/10.1109/LCA.2026.3705817>
- [658] Hsu, Xu, Liu, Rosing and Yu (UCSD and Georgia Tech), “HAVEN: HBF-Augmented Vector Engine for ANN search,” arXiv:2603.01175 (1 Mar. 2026): HBF holding a full-precision vector database on the GPU for RAG; up to $20\times$ throughput and $40\times$ latency improvement over GPU-DRAM and GPU-SSD baselines; 4 KB page baseline. <https://arxiv.org/abs/2603.01175>
- [659] NAND die-level state of the art: Kioxia and SanDisk, ISSCC 2025 (1 Tb TLC CBA at 4.8 Gb/s I/O, 29% read-energy improvement) and ISSCC 2026 (BiCS10 2 Tb QLC, 332 active layers, >85 MB/s program); Samsung V10 at 5.6 Gb/s per pin (ISSCC 2025); SK hynix 321-layer 2 Tb QLC at 75 MB/s program (ISSCC 2025)—i.e. per-die programme rates roughly fifty times below read rates. <https://www.sandisk.com/company/newsroom/press-releases/2025/kioxia-and-sandisk-unveil-next-generation-3d-flash-memory-technology>
- [660] The flash/SSD-offloading lineage: Sheng et al., FlexGen, ICML 2023 (arXiv:2303.06865); Alizadeh et al. (Apple), “LLM in a flash,” ACL 2024 (arXiv:2312.11514); Eliseev and Mazur, arXiv:2312.17238; Xue et al., MoE-Infinity, arXiv:2401.14361; Cao et al., MoE-Lightning, ASPLOS 2025 (arXiv:2411.11217); He et al., ExpertFlow, arXiv:2410.17954; Tang et al., HOBBIT, arXiv:2411.01433; Yu et al., PreScope, arXiv:2509.23638; Kim et al., FlashMoE, arXiv:2601.17063; Xue et al. (SJTU), PowerInfer-2, arXiv:2406.06282 (47B MoE from smartphone flash at 11.68 tok/s); Microsoft, ZeRO-Inference (Sep. 2022, OPT-30B at 30 tok/s with NVMe offload). <https://arxiv.org/abs/2312.11514>
- [661] Mixture-of-experts routing locality: Liang et al. (Fudan, Huawei), “Not All Models Suit Expert Offloading: On Local Routing Consistency of Mixture-of-Expert Models,” arXiv:2505.16056

- (2025–26; 20 MoE models; cache of $\approx 2\times$ active experts near-optimal; shared-expert designs lower locality); Oncescu, Dao et al., “Opportunistic Expert Activation,” arXiv:2511.02237 (Nov. 2025; batch-aware rerouting cuts MoE-layer decode latency 15–39%); Huang et al. (Meta), “Towards MoE Deployment,” arXiv:2303.06182; Yuan and Long, “Economics of AI Decoding Chips,” arXiv:2607.13068 (Jul. 2026; the “MoE tax”). <https://arxiv.org/abs/2505.16056>
- [662] Xiaohongshu dots studio, *dots3-note Preview* (released 14 Aug. 2026): 280B total / 16B active parameters (45 MoE layers, 256 routed + 1 shared expert, top-8; 7B/1.2B-active MoE vision encoder; 800M audio encoder); text, image, video and audio input; 512K context; Apache 2.0; weights on Hugging Face (**dots-studio/dots3-note-prev**, FP8 variant) and ModelScope; self-reported SWE-bench Verified 78.4, SWE-bench Pro 61, SWE-bench Multilingual 75.7, MMMU-Pro 79.1; TEMPO (“Test-time-scaled Value Estimation with Macro-step Policy Optimization”) described in the release material; technical report “coming soon.” <https://huggingface.co/dots-studio/dots3-note-prev>
- [663] 36Kr (14 Aug. 2026) and NetEase/ifeng coverage of the dots3-note release: TEMPO divides trajectories of ten to twenty hours into macro-steps; at each boundary the same model switches from actor to critic and estimates the expected return of the current state through reasoning, tool use and test-time scaling, producing dense intermediate training signal; ARC-AGI-3 ≈ 0.35 under a \$500 budget; VibeSearchBench and VibeLifeBench introduced. <https://eu.36kr.com/en/p/3938759517896072>
- [664] South China Morning Post (22 Jul. 2026): Xiaohongshu’s dots-note-3.0 becomes the first AI model to earn a perfect 42/42 at the International Mathematical Olympiad. <https://www.scmp.com/tech/article/3361482/worlds-first-ai-model-earn-perfect-score-maths-olympiad-comes-chinas-rednote>
- [665] IT Home (16 Aug. 2026): Huawei Ascend completes “0-day” inference deployment and optimization of dots3-note on Atlas 800 A3 and Atlas 900 A3 SuperPoD via vLLM-Ascend; vLLM-Ascend project documentation carries a Dots3-Note deployment tutorial validated on Atlas 800I A3. <https://docs.vllm.ai/projects/ascend/en/latest/tutorials/models/Dots3-Note.html>
- [666] Wall Street Journal via Investing.com (17 Jun. 2026) and Bloomberg (15 Jun. 2026): Xiaohongshu preparing a Hong Kong IPO at a valuation above \$70B, recent private transactions implying more than \$50B; more than 400 million monthly active users; 2025 net income above \$2B. <https://www.investing.com/news/stock-market-news/chinas-xiaohongshu-prepares-hong-kong-i-po-at-over-70-bln-valuation--wsj-4746470>
- [667] Xiaomi MiMo, “MiMo-V2.5 open-sourced” (27 Apr. 2026): MiMo-V2.5 (310B/15B active MoE, native omni-modal text/image/video/audio input, 1M context) and MiMo-V2.5-Pro (1.02T/42B active, text, 1M context, 27T training tokens), both MIT; MiMo-V2-Flash (309B/15B, MIT) Dec. 2025. <https://huggingface.co/XiaomiMiMo/MiMo-V2.5>
- [668] Vercel, “Open-weight models surge to 29% of volume, price per token flattens,” AI Gateway production index (13 Jul. 2026, June 2026 data): “Open-weight models ran 29% of gateway tokens in June on just under 4% of spend”; $\sim 11\%$ of tokens in April; DeepSeek 22.6% of June tokens, within two points of Google; GLM 5.2 took 76% of its family’s tokens within two weeks of launch; Anthropic 61% of spend on 32% of tokens; roughly one in eight enterprise customers runs an open-weight model. <https://vercel.com/blog/ai-gateway-production-index-july-2026>
- [669] 1 TB NVMe SSD retail, pre-surge: $\sim \$65$ – $\$85$ in Q1 2025 rising to $\$75$ – $\$100$ by Q1 2026; a 1 TB drive $\sim \$45$ in late 2025 had reached $\sim \$90$ by mid-2026; a 2 TB drive $\sim \$120$ – $\$150$ a year earlier (mid-2025) reached $\$300$ – $\$480$ by Jul. 2026. Storage Disk Prices, SSD Price History; Tech Insider, “SSD Prices Double as NAND Shortage Hits Gaming PCs.” <https://storagediskprices.com/ssd-price-history/>, <https://tech-insider.org/ssd-prices-nand-shortage-2026/>
- [670] US Energy Information Administration, Electric Power Monthly (May 2026): average industrial price 8.71 ¢/kWh, commercial 13.54 ¢/kWh. <https://www.eia.gov/electricity/monthly/update/end-use.php>
- [671] API list prices, August 2026, per million tokens input/output: Claude Opus 5 $\$5/\25 (Sonnet 5 $\$2/\10); GPT-5.6 Sol $\$5/\30 , Luna $\$0.20/\1.20 after the 30 Jul. cut; Gemini 3.1 Pro $\$2/\12 to 200K context; Grok 4.6 $\$2/\6 ; Kimi K3 $\$3/\15 ($\$0.30$ cached); Qwen3.8-Max $\$2/\6 ; DeepSeek V4-Pro from 16 Aug. $\$1.32/\3.96 peak and $\$0.66/\1.98 off-peak, V4-Flash $\$0.44/\1.32 peak. Vendor pricing pages. <https://platform.kimi.ai/docs/pricing/chat-k3>

- [672] Hot Chips 2026, Stanford University, 23–25 August 2026; 31 technical presentations across memory, server CPU, GPU/accelerator, network and tutorial tracks. Deck numbers cited in this book follow the conference numbering. Where a figure appears only in a deck it is marked as vendor-presented and unaudited; where a vendor withheld a figure that exists elsewhere, the outside source is named. <https://hotchips.org/>
- [673] Hot Chips 2026 deck 09, Frans Sijstermans (NVIDIA), “RISC-V CPUs in NVIDIA Servers: CUDA’s Third Host ISA”: NVIDIA states the requirements for a RISC-V CPU to host CUDA alongside x86 and Arm—RVA23 compliance plus the RISC-V Server SoC and Server Platform specifications, Boot and Runtime Services (ratified August 2025) and ACPI; CUDA-specific asks reduce to coherent PCIe I/O and full-speed PCIe peer-to-peer with standard root-complex ordering. A second tier for NVLink Fusion requires NVIDIA’s NVLink-C2C soft IP behind C2C hubs speaking the CHI coherence protocol, plus DOCA and NCCL support. SiFive appears on the NVLink Fusion partner slide alongside Arm, Fujitsu and Intel. The requirement document itself is stated to be available to partners only. Reported by Chester Lam, Chips and Cheese (24 Aug. 2026). <https://chipsandcheese.com/p/hot-chips-2026-cuda-targets-risc>
- [674] Hot Chips 2026 deck 07, Krste Asanović (RISC-V International, SiFive, UC Berkeley), “Update on RISC-V Standards and Adoption: Profiles and Platforms”: RVA23 ratified October 2024 with vector and hypervisor mandatory; RVA23.1 near ratification; Server Platform specification ratified May 2026. *No ratified matrix extension exists*: one fast-track dot-product proposal (Zvbdota) and three competing task groups—IME (matrix-in-vector), VME (vector-matrix, adding an accumulator file) and AME (separate matrix state)—with ratification targets spanning February 2026 to April 2027. New “optimization guidance” options concede that a functional-only ISA string does not tell a vendor how to ship a competitive core. <https://riscv-europe.org/summit/2026/media/proceedings/2026-06-10-RISC-V-Summit-Europe-09h00-ASANOVIC%20C4%86-slides.pdf>
- [675] Hot Chips 2026 deck 08, Gordan Markuš (Canonical), “Enterprise Open Source on RISC-V and Convergence on RVA23”: Ubuntu’s riscv64 baseline was RVA20 from 20.04 LTS through 25.04 and is RVA23 from 25.10; 26.04 LTS carries RVA23 as the long-term baseline, with RVA20 silicon retaining a supported path on 24.04 LTS for up to ten years under Ubuntu Pro. The readiness number: roughly 95% of archive packages build on riscv64 but only **55.1%** of tested packages pass, against 92.3% on amd64 and 92.5% on arm64, on roughly one seventh the test coverage—attributed to emulation pending server-grade RVA23 hardware. RISE board farm runs on Alibaba and SpacemiT hardware. <https://canonical.com/blog/canonical-and-ubuntu-risc-v-a-2025-retro-and-looking-forward-to-2026>
- [676] SiFive, “BigSky SF-2U870” (24 Aug. 2026): the first rackable enterprise RISC-V server—32 P870-D cores at 2.0 GHz, 256 GB DDR5-5600, 64 lanes of PCIe Gen5, 2U, fully RVA23 compliant, shipping with Ubuntu 26.04 LTS and RHEL 10, and able to act as a CUDA head node for NVIDIA GPU clusters. Available in limited quantities; no price, no SPEC result and no named customer tape-out disclosed. Futurum’s assessment is that BigSky is “primarily a software porting vehicle rather than revenue-generating hardware,” valuable because it moves the multi-year software phase of every SiFive-IP-based custom SoC onto hardware a customer can rack today. <https://www.hpcwire.com/off-the-wire/sifive-brings-risc-v-to-rackable-servers-with-bigsky-platform/>
- [677] Qualcomm’s acquisition of Ventana Micro Systems (announced Dec. 2025; terms undisclosed; Veyron V1 never reached market and V2 slipped from 2025 to 2026), and Reuters reporting (16 Jun. 2026) that Qualcomm was in talks to acquire Tenstorrent for up to \$10B—unconfirmed as of the currency date. Of three Western RISC-V server contenders, one shipped a server, one was absorbed without shipping its first product, and one is a rumoured acquisition target. https://www.theregister.com/2025/12/10/qualcomm_riscv_arm_ventana/
- [678] Hot Chips 2026 deck 01, Jim Handy (Objective Analysis), “Memory: Feeding AI’s Voracious Hunger for Data”: lowest spot DRAM price up $7.1\times$ and NAND up $6.9\times$ from the mid-2025 trough; combined DRAM and NAND revenue of the five major suppliers from a $\sim\$15\text{B}$ quarterly trough in 2Q23 to $\sim\$180\text{B}$ in 2Q26, almost entirely price-driven rather than bit-driven; hyperscale property and equipment spending from $\sim\$30\text{B}$ per quarter in 1Q21 to $\sim\$170\text{B}$ in 2Q26; DRAM wafer starts flat for over a decade and new capacity taking more than two years to install. The deck’s summary slide is a photograph of banknotes: the 2026 memory wall is priced, not clocked. Conflict disclosed on slide 7: the data underpin a paid report.

- [679] Hot Chips 2026 decks 02 (Micron), 03 (Samsung) and 04 (SK hynix): the same JEDEC HBM4 generation is quoted by SK hynix at 8 Gb/s and 2.0 TB/s per stack, by Micron at 11 Gb/s, and by Samsung at 3.0 TB/s—a $1.5\times$ spread in per-stack bandwidth within one nominal generation, so a bill of materials specifying “HBM4” does not specify a part. Micron: an eight-stack HBM4 package carries 12,344 mm² of memory silicon, more than $8\times$ the die area of a typical GPU. Samsung deck 28: memory rose from 52% to 63% of AI-chip component spend between 1Q24 and 4Q25.
- [680] Hot Chips 2026 deck 03, Sangwook Han (Samsung), “HBM Base Die: How HBM Will Evolve Using Advanced Logic Processes”: from HBM4 the base die is fabricated on a 4 nm foundry logic process, which turns the memory vendor into a chip designer. The three-phase roadmap is reclaim xPU area by absorbing the memory controller into the base die; fill the free base-die area with repair, telemetry, test and processing elements; then delete the interposer entirely by bonding the cell-array stack directly onto the xPU (“zHBM”). Custom HBM is explicitly *not* JEDEC-standardised: a customer-specific logic die sharing standard cell-array stacks. Strategically this inverts the value chain—the memory vendor comes to own an interface that belonged to the accelerator designer.
- [681] Hot Chips 2026 decks 23 (AMD, Steve Scott, “System Architecture of the MI400 Series”), 36 (Microsoft, Maia 200), 37 (SambaNova SN50) and 33 (Meta MTIA 400): AMD builds a 72-GPU load/store domain out of UALink-over-Ethernet on merchant 512-port 200G switch ASICs, deleting the scale-up NIC from the tray; Microsoft runs one Ethernet protocol from chip to a 6,144-accelerator two-tier Clos, with slide 16 titled “Why Unified?” and listing a standard switch ecosystem as the reason; SambaNova uses ten integrated 800G Ethernet ports per socket and two merchant switches for a 64-socket domain; Meta uses RoCE on dedicated I/O chiplets. Google’s optically switched torus is proprietary but is not sold; NVIDIA’s NVLink 6 (130 TFLOPS of in-network compute) is the only proprietary scale-up fabric a third party can buy. Of eight parts with an off-package scale-up fabric, six build it from Ethernet-class silicon.
- [682] Hot Chips 2026 deck 30, Hemal Shah (Broadcom), “Thor Ultra”: every mechanism the Ultra Ethernet Consortium argued for shipped in production silicon this year—packet spraying, out-of-order placement with in-order delivery, selective retransmission, receiver-driven credit congestion control, packet trimming, programmable transport—but as *RoCEv2 extensions*. Broadcom brands its transport eRoCE and describes it as “MRC++,” preserving the verbs ecosystem so xCCL, NCCL, RCCL and MPI run unmodified: “the ecosystem cost is zero.” Neither Broadcom’s nor NVIDIA’s deck presents a UET transport or a libfabric stack, and NVIDIA’s networking deck does not mention the UEC at all. The consortium standardised the mechanism set and lost the wire format.
- [683] Hot Chips 2026 deck 39, Richard Ho, Ravi Narayanaswami and Chris Leary (OpenAI), “You Can Just Build Things ... Chips”: Jalapeño, OpenAI’s inference accelerator built with Broadcom—13.4 PFLOP/s mxfp4, 216 GiB HBM4 at 15.4 TB/s, 700 W, a 128-chip scale-up domain at 600 GB/s and a 2,048-chip pod over a half-flattened two-level Broadcom Tomahawk 6 Clos. The organising metric is output tokens per second per all-in utility megawatt at matched interactivity; per-chip throughput and chip count are declared explicit non-goals. Architectural argument against a specialised fleet: “dark silicon is cheaper than idle accelerators.” More than half the core was written in the XLS hardware language with an internal model in the loop. Vendor-supplied comparisons against GB200/GB300 and MI355X are unaudited and exclude speculative decoding. <https://openai.com/index/openai-broadcom-jalapeno-inference-chip/>
- [684] OpenAI and Broadcom, “Strategic collaboration to deploy 10 gigawatts of OpenAI-designed accelerators” (13 Oct. 2025), signed as a term sheet with no financial terms disclosed; deployment from 2H 2026 through 2029. Jalapeño unveiled 24 June 2026 as the first product. Greg Brockman: “By building our own chip, we can embed what we’ve learned from creating frontier models directly into the hardware.” Foundry, process node, unit cost and unit volume are not disclosed in any primary source. <https://investors.broadcom.com/news-releases/news-release-details/openai-and-broadcom-announce-strategic-collaboration-deploy-10>
- [685] Hot Chips 2026 deck 38, Norman P. Jouppi and Sridhar Lakshmanamurthy (Google), “Eighth Generation TPU Family: Two Chips for the Agentic Era”: for the first time Google shipped a training part (8t: 9,600 chips, optically switched 3D torus, 6 HBM3E stacks) and an inference part (8i: 1,152 chips, seven-hop Boardfly hierarchy, 8 HBM3E stacks, 384 MB SRAM, collectives moved into the ICI I/O die) *concurrently* rather than alternately, on the stated ground that pre-training, distillation, sampling, MoE and agentic serving now need different interconnect topologies. One

- XLA/JAX/PyTorch stack targets both. Claimed $> 2\times$ perf/TCO and $2\times$ perf/W generation over generation, without a stated workload. No per-chip FLOPS, power, node or die size published. <https://www.servethehome.com/google-tpu-8i-for-inference-and-tpu-8t-for-training-announced/>
- [686] Sundar Pichai, Alphabet Q1 2026 earnings call (29 Apr. 2026): “As TPU demand grows from AI labs, capital markets firms, and high-performance computing applications, we’ll begin to deliver TPUs to a select group of customers in their own data centers in a hardware configuration to expand our addressable market opportunity”—a small share of TPU hardware revenue in 2026, the majority in 2027. Alphabet raised 2026 capital expenditure guidance to \$180–190B in the same call. Google and Broadcom expanded their Anthropic partnership on 6 April 2026 to multiple gigawatts of next-generation TPU capacity from 2027. <https://www.datacenterdynamics.com/en/news/google-to-sell-tpus-to-a-select-group-of-customers-for-their-data-centers/>
- [687] Microsoft, “Maia 200: the AI accelerator built for inference” (26 Jan. 2026): TSMC 3 nm, 140B+ transistors, 216 GB HBM3E at 7 TB/s, 272 MB on-die SRAM, >10 PFLOPS FP4, 750 W, clusters to 6,144 accelerators over one Ethernet protocol; deployed in US Central. The headline claim is economic, not architectural: “30% better performance per dollar than the latest generation hardware in our fleet today.” NVIDIA is not mentioned. <https://blogs.microsoft.com/blog/2026/01/26/maia-200-the-ai-accelerator-built-for-inference/>
- [688] Meta, MTIA roadmap (announced Mar. 2026, Broadcom co-designed; Hot Chips 2026 deck 33): four generations on roughly a six-month cadence—MTIA 300 (800 W, in production), MTIA 400 (1,200 W, 288 GB HBM3E at 9.2 TB/s, 6 PF FP8/MX8, 12 PF MX4, in lab test), MTIA 450 (early 2027) and MTIA 500 (later 2027). Meta has deployed “hundreds of thousands” of earlier MTIA parts for inference. Meta diverges from the OCP microscaling standard deliberately: 16 elements per shared exponent rather than 32, plus an MX8S 16×16 FP8 block format created for training. <https://www.tomshardware.com/tech-industry/semiconductors/meta-reveals-four-new-mtia-chips-built-for-ai-inference>
- [689] Amazon Web Services, Trainium3 general availability (2 Dec. 2025): first 3 nm Trainium, 2.517 PFLOPS FP8 and 144 GB HBM3E at 4.9 TB/s per chip, UltraServers to 144 chips, scaling to a million-chip cluster; AWS states more than one million Trainium processors deployed. Matt Garman: “Trainium already represents a multi-billion-dollar business.” <https://www.hpcwire.com/2025/12/02/aws-brings-the-trainium3-chip-to-market-with-new-ec2-ultraservers/>
- [690] AWS and NVIDIA (26 Aug. 2026): two million additional NVIDIA GPUs (Blackwell Ultra, Rubin, Rubin Ultra) across AWS in 2027–2028, plus 100,000 GPUs for US government AI factories—and, materially, *Trainium4 will adopt NVIDIA NVLink Fusion and NVIDIA custom HBM*, placing AWS’s own accelerator inside NVIDIA’s rack-scale interconnect and memory stack. Garman: “Customers want the freedom to choose the best tools.” This is the strongest single piece of evidence that NVIDIA is absorbing the custom-silicon wave rather than being displaced by it. <https://press.aboutamazon.com/aws/2026/8/aws-and-nvidia-to-deliver-2-million-additional-gpus-and-next-generation-infrastructure-for-agentic-and-physical-ai>
- [691] Broadcom, second quarter fiscal 2026 results (3 Jun. 2026): total revenue \$22.187B, up 48%; AI semiconductor revenue \$10.8B, up 143% year on year; third-quarter guidance of more than \$16.0B in AI semiconductors, up about 200%. Hock Tan attributes growth to “increasing demand for custom AI accelerators and AI networking.” Annualising the guided run rate ($> \$64$ B a year) against NVIDIA’s data-centre run rate ($\sim \$356$ B a year on the July-quarter actual) places merchant custom-ASIC revenue at roughly 11–15% of the combined total, the range set by whether actuals or guidance are annualised on each side [M]. Note that NVIDIA’s \$108.0B third-quarter figure is *total-company* guidance; NVIDIA publishes no segment guidance. <https://investors.broadcom.com/news-releases/news-release-details/broadcom-inc-announces-second-quarter-fiscal-year-2026-financial>
- [692] NVIDIA, “NVIDIA partners with Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR to establish AI compute infrastructure financing platforms to mobilize over \$500 billion of third-party capital” (10 Aug. 2026). The agreements are memoranda of understanding, “subject to execution of the final agreements.” Jensen Huang: “We began by building chips; today, we are helping create a new class of productive, investable infrastructure: AI factories” and “NVIDIA compute is . . . fungible and transferable across customers and operators.” David Solomon: “We’re excited for the new opportunity to create a market for credit backed by NVIDIA compute.” The widely reported 25% residual-value share appears in Huang’s own remarks and press reporting, *not*

- in the release or in any filing. <https://nvidianews.nvidia.com/news/nvidia-partners-with-apollo-blackrock-blackstone-brookfield-goldman-sachs-and-kr-to-establish-ai-compute-infrastructure-financing-platforms-to-mobilize-over-500-billion-of-third-party-capital>
- [693] US Securities and Exchange Commission, Office of Structured Finance interpretive position (29 Jul. 2026): securities in data-centre securitizations are not “asset-backed securities” under Exchange Act §3(a)(79), because data centres are “tangible, physical assets that endure beyond the tenor of the securities” rather than self-liquidating financial assets—placing them outside Regulation AB and outside the 5% credit-risk-retention requirement of Regulation RR. The ruling preceded NVIDIA’s financing-platform announcement by twelve days. Alston & Bird (3 Aug. 2026); Seward & Kissel (10 Aug. 2026). <https://www.alston.com/en/insights/publications/2026/08/data-center-securitizations-sec-regulations>
- [694] NVIDIA Corporation, Form 8-K filed 17 August 2026 (Items 1.01, 2.03, 7.01): multiple residual value guaranties with SB Energy Corp. on data-centre leases at the Portsmouth Site, Pike County, Ohio, whose tenant is an OpenAI Group PBC affiliate. “NVIDIA’s aggregate payment obligation is cumulatively capped at \$105 billion for its initial commitment under the Agreements.” Approximately 4.25 GW of IT load initially with an option on a further 3.8 GW; twenty-year maximum term; *terminates early if OpenAI achieves a satisfactory credit rating*; OpenAI must reimburse and indemnify NVIDIA for all amounts paid. First phases effective in fiscal 2029. NVIDIA also invests \$1.5B in SB Energy and becomes exclusive AI compute infrastructure provider for the site. Earlier press reports of a \$250B commitment are superseded by the filed figure. <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000069/nvda-20260817.htm>
- [695] NVIDIA, second quarter fiscal 2027 results (quarter ended 26 Jul. 2026; announced 26 Aug. 2026): revenue \$96.2B, up 106% year on year; data-centre revenue \$89.0B, up 117%; GAAP and non-GAAP gross margin 75.0%; GAAP net income \$59.688B; third-quarter guidance \$108.0B $\pm$ 2% with gross margin 74.0%, excluding China data-centre compute revenue. GAAP diluted EPS (\$2.46) *exceeds* non-GAAP (\$2.22) because GAAP other income of \$7.773B is almost entirely gains on equity securities—that is, roughly an eighth of GAAP net income for the quarter is mark-ups on stakes in companies that buy NVIDIA’s chips. Shares closed up 8.7% on 27 August, the largest single-day gain since April 2025. <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000073/q2fy27pr.htm>
- [696] NVIDIA, CFO Commentary filed as Exhibit 99.2 to the Q2 FY2027 results (26 Aug. 2026), and Form 10-Q for the period ended 26 July 2026: maximum gross guarantee exposure \$108.5B, of which \$105B is the SB Energy residual value guarantees and \$3.5B land, power and shell guarantees for AI cloud partners’ lease obligations in the event of their default; supply and capacity purchase commitments \$279B (from \$119B the prior quarter, memory-driven); AI cloud revenue-share commitments \$36B; future equity investment commitments \$25B; data-centre leases held for third-party reassignment \$20B; non-marketable equity securities \$47.898B and publicly held equity securities \$47.740B. <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000073/q2fy27cfocommentary.htm>
- [697] NVIDIA’s revenue-sharing and credit-support model for AI clouds, introduced July 2026: typically six-year take-or-pay arrangements in which the operator retains rental revenue up to a backstop floor and NVIDIA takes a share of the excess; counterparties reported to include Sharon AI (~40,000 GB300) and Firmus. Earlier and separate: CoreWeave’s 8-K of 15 September 2025 obliging NVIDIA to purchase residual unsold capacity through 13 April 2032, capped at \$6.3B; a \$1.5B arrangement with Lambda; a 9.3% passive stake in Nebius (Schedule 13G, 20 Jul. 2026). NVIDIA’s own framing: the model gives it “a recurring, usage-linked earnings stream.” SemiAnalysis (6 Jul. 2026). <https://newsletter.semianalysis.com/p/nvidia-gpu-debt-backstop-unleashes>
- [698] Reuters and *The Wall Street Journal* (27 Aug. 2026), one day after earnings: NVIDIA paused parts of the July revenue-sharing and credit-support programme, reportedly after employees flagged potential antitrust scrutiny and after friction over the extent to which NVIDIA could dictate which customers its partners rented to. NVIDIA’s on-record statement: “The new business model we introduced in July that opens up compute access to the fast-growing AI ecosystem is still in place and continues to evolve due to high demand.” <https://whbl.com/2026/08/27/nvidia-pause-s-revenue-sharing-deals-with-ai-cloud-companies-wsj-reports/>

- [699] NVIDIA Q2 FY2027 earnings call, 26 August 2026, prepared remarks: “We recognize the scale of this support, and we know some will call this circular financing. We see it differently. . . . Independent capital still underwrites every deal on its own merits. *We’re not making loans.* In this model, we get paid twice—once on the hardware sale, and again through the share of rental revenue.” Huang, separately: “We’ve invested nearly \$50 billion in the Frontier AI Labs. . . the only regret that I have is that I didn’t invest more and sooner”; and, on the Ohio structure, “Is this circular financing? No, and OpenAI is paying the lease.” <https://seekingalpha.com/article/4940563-nvidia-corporation-nvda-q2-2027-earnings-call-transcript>
- [700] Huang, same call: “The world will need both closed models and open models. And both closed models and open models are skyrocketing in use. . . . *Nearly all open models run on NVIDIA* . . . our position in open models is very good, because the CUDA ecosystem is literally everywhere.” Compare Nader Shaikh, Director of Developer Technology, NVIDIA, at the Six Five Summit (27 Aug. 2026): “We build Nemotron because we know that this is going to be a multi-model world, and we see where there’s the bottleneck”—arguing that durable value accrues to proprietary data rather than to model weights, which is why NVIDIA releases weights *and* training data. <https://www.investing.com/news/transcripts/nvidia-at-the-six-five-summit-open-models-reshape-ai-strategy-93CH-4880247>
- [701] NVIDIA Nemotron 3 family (15 Dec. 2025): Nano 30B-A3B, Super ~100B-A10B and Ultra ~500B-A50B, released with *weights, training datasets and reinforcement-learning libraries*. Huang: “Open innovation is the foundation of AI progress. With Nemotron, we’re transforming advanced AI into an open platform.” Nemotron 3.5 Lightning (11 Aug. 2026), 30B MoE with 3B active, shipped under the OpenMDW-1.1 licence with training data and recipes, on Hugging Face, ModelScope, OpenRouter, SageMaker JumpStart and Google Cloud. The Nemotron Pretraining Collection exceeds ten trillion tokens; NVIDIA also publishes open robotics, protein-structure and driving datasets. Nemotron 3 Ultra was the sixth-largest model by weekly tokens on OpenRouter at the currency date. <https://nvidianews.nvidia.com/news/nvidia-debuts-nemotron-3-family-of-open-models>
- [702] NVIDIA open-sourced the CUDA Tile IR and cuTile compiler stack under Apache 2.0 with LLVM exceptions (25 Dec. 2025)—a Tile MLIR dialect, Python bindings, bytecode representation and conformance suite—describing it as “the largest and most comprehensive update to the CUDA platform since it was invented two decades ago.” TensorRT-LLM is Apache 2.0. At Hot Chips 2026 NVIDIA stated it was extending CUDA to cover the Groq-derived LP30 LPU as a supported target; outlets characterise the depth of that integration differently, from “an LPU extension to CUDA” to “a fully-supported target.” <https://github.com/NVIDIA/cuda-tile>
- [703] Hot Chips 2026 deck 34, Igor Arsovski and Santosh Raghavan (NVIDIA), “Think Fast: LPU Accelerator for Heterogeneous Compute”: the Groq-derived LP30—no DRAM, 0.5 GB of on-die SRAM per chip aggregating to 128 GB per 256-chip rack at ~40 PB/s and 315 PFLOPS FP8, a cycle-exact VLIW ISA on a single global clock with no memory barriers and no kernel boundaries, every chip its own router at 350 ns SRAM to SRAM. NVIDIA acquired Groq’s assets and key staff for a reported \$20B, completed late December 2025. The commercially significant slide is the disaggregation set: prefill on Rubin and decode on LP30, attention on Rubin and feed-forward on LP30, Rubin verifying while LP30 drafts—claimed ~3× to ~5× interactivity over Rubin alone. NVIDIA is publicly designing a system in which its own GPU is not the right chip for part of the inference pipeline. Groq 3 LPX entered full production 24 August 2026. <https://www.servethehome.com/nvidias-groq-3-lpu-accelerators-for-heterogeneous-ai-compute-at-hot-chips-2026/>
- [704] Morgan Stanley initiation of credit coverage on NVIDIA (24 Aug. 2026), rating Neutral: total credit exposure of roughly \$200B by the end of 2028, of which about \$170B off balance sheet—approximately \$38B of minimum rating adjustments, \$68B of residual-value support tied to the \$500B framework and \$64B of revenue-sharing and credit support for non-hyperscalers—against projected net cash of about \$520B in fiscal 2029. The note coins “balance-sheet-as-a-service.” Component breakdown reported secondhand; the original note was not obtained.
- [705] Dave Vellante, “Nvidia’s Jensen Huang just made the AI buildout too big to fail,” SiliconANGLE (15 Aug. 2026, updated 21 Aug.): “Independent capital can fund more factories. But independent capital is not independent demand”—the requirement for credit enhancement is itself evidence that markets do not yet accept compute as enduring collateral. Three simultaneous warning signals nominated: GPU rental prices falling, utilisation weakening, financing terms tightening. Concludes

- the buildout is not too big to fail but “too interconnected to fail quietly.” For the opposing structural reading see NVIDIA’s own rebuttal of vendor-financing comparisons (Nov. 2025): “unlike Lucent, NVIDIA does not rely on vendor financing arrangements to grow revenue.” <https://siliconangle.com/2026/08/15/nvidias-jensen-huang-just-make-ai-buildout-big-fail/>
- [706] AMD, second quarter 2026 results (4 Aug. 2026): revenue \$11.5B, up 50%; data-centre revenue \$6.7B, up 107% year on year and 58% of the total; GAAP gross margin 54%; third-quarter guidance approximately \$13B. Form 10-Q for the period ended 27 June 2026 discloses \$4.5B of leases not yet commenced and **\$4.1B in data-centre lease guarantees**, plus data-centre leases with \$9.5B of aggregate future payments beginning in 2027–2028. <https://ir.amd.com/news-events/press-releases/detail/1295/amd-reports-second-quarter-2026-financial-results>
- [707] AMD’s warrant instruments: a warrant for up to 160 million shares at \$0.01 issued to OpenAI (6 Oct. 2025), vesting in tranches tied to 6 GW of purchases *and* escalating AMD share-price thresholds; and an equivalent 160-million-share warrant issued to Meta under an amended master purchase agreement (Form 8-K filed 24 Feb. 2026), vesting from 1 GW to 6 GW with price thresholds escalating to \$600 per share. Separately, up to \$5B of equity and 2 GW of MI450 for Anthropic (22 Jul. 2026). As of the August 2026 filing, no warrant shares had vested or been issued. The customer is being paid, in the vendor’s own equity, to buy the vendor’s product. <https://www.sec.gov/Archives/edgar/data/2488/000000248826000045/amd-20260223.htm>
- [708] AMD’s openness posture: ROCm 7 positioned explicitly against CUDA; Instella and Instella-MoE released as *fully* open 3B models—weights, data and training code. UALink 1.0 released 2025 with 2.0 due in 2026 adding in-network computing, chaired by AMD’s Kurtis Bowman; the promoter group is AMD, Broadcom, Cisco, Google, HPE, Intel, Meta and Microsoft, and NVIDIA is not a member. <https://rocm.blogs.amd.com/artificial-intelligence/introducing-instella-3B/README.html>
- [709] SpaceX “Terafab” (confirmed by SpaceX and reported by Reuters, 6 Aug. 2026): \$16.8B initial investment in a vertically integrated logic and memory manufacturing, packaging and test site in Grimes County, Texas, with a minimum of 3,000 jobs; output targeted at Tesla Optimus, Cybercab and SpaceX space-based data centres. Musk announced the project in March 2026 with a stated ambition of up to 50× the annual AI-chip output of major suppliers; a May 2026 regulatory filing showed a proposed \$55B initial and up to \$119B total across phases. Intel entered a pact with Tesla, SpaceX and xAI (Apr. 2026) for advanced logic and packaging, which one analyst characterises as “likely more an Intel fab expansion with Tesla, SpaceX, and xAI as anchor customers.” Process node and timelines are not disclosed in the August confirmation, and the figures have moved three times. SpaceX absorbed xAI in an all-stock merger at a combined \$1.25T valuation (2 Feb. 2026). <https://techcrunch.com/2026/08/06/tesla-and-spacex-will-invest-16-8b-t-o-start-building-terafab-chip-factory-in-texas/>
- [710] Tesla’s dual-foundry arrangement for AI5 and AI6, confirmed by Musk (4 Nov. 2025): variants of the same design at TSMC and Samsung. Samsung’s Taylor, Texas 2 nm fab reached AI5 tape-out (reported 13 Jul. 2026), with volume production targeted for 2H 2027. Dojo was shut down in August 2025—Musk called it “an evolutionary dead end”—and restarted as Dojo3 in January 2026 with the stated goal of space-based AI compute. <https://www.tomshardware.com/tech-industry/se-miconductors/musk-confirms-tesla-ai5-and-ai6-will-be-made-at-samsung-and-tsmc>
- [711] Zhipu AI / Z.ai, GLM-5.3 (announced 13–14 Aug. 2026; API from 18 Aug.; open weights released 28 Aug. 2026): ~744B total parameters with ~40B active, 78 layers, 256 routed experts with 8 active per token, 1M context. The licence is neither Apache 2.0 nor MIT but a custom “glm-5.3” licence whose reported terms require a security review for model-as-a-service operators above \$10B of trailing revenue. The model card documents Ascend NPU deployment paths through vLLM-Ascend, xLLM and SGLang. Company-reported benchmarks lead on vulnerability discovery (CyberGym 84.5 against GPT-5.6 Sol’s 83.6) and trail roughly two to one on exploit development (ExploitGym 105/130 against 216/293); Z.ai’s claim of 2,436 vulnerabilities found across 269 projects is unaudited. <https://huggingface.co/zai-org/GLM-5.3>
- [712] Zhipu AI / Z.ai, GLM-5.3-Flash (released 26 Aug. 2026; revealed on 27 August to have been the stealth model “Ox Alpha”): 320B total parameters with 18B active, hybrid sparse and linear attention, the first natively multimodal model of the GLM-5 series, trained on a 30T-token multimodal corpus, 1M context—and released under **MIT**. Claimed ~3× compute and ~4.4× KV-cache reduction against full GLM-5.3. List price \$0.15/\$0.50 per million tokens input/output,

- roughly $7.5\times$ below GLM-5.3. Hugging Face records 346,516 downloads in the first month with 62 quantized derivatives; the model is listed on the Dell Enterprise Hub.
<https://huggingface.co/zai-org/GLM-5.3-Flash>
- [713] Artificial Analysis intelligence index, launch week (v4.1.1, late August 2026; superseded by the v4.3 re-basing of 7 September, on which GLM-5.3 scores 45 and Flash 42): GLM-5.3 at 60, second of 111 evaluated models; **GLM-5.3-Flash at 57, fourth of 111**, at \$0.09 cost per task against GLM-5.3's \$1.40. Frontier closed models: Claude Opus 5 (max) 63; Claude Fable 5 62; Claude Opus 5 (high), GPT-5.6 Sol (max) and Grok 4.6 (high) 61. The measured gap between the best open-weight Chinese model and the best closed model is three index points; Flash at 57 ties GPT-5.6 Sol (high) and exceeds Gemini 3.7 Flash. Caveat: the Flash score circulated in Z.ai's own materials before appearing on the public board. <https://artificialanalysis.ai/models/glm-5-3-flash>
- [714] Unsloth quantization measurements for GLM-5.3-Flash: UD-IQ1_S 93.1 GB retaining 70.9% top-1 accuracy; UD-Q2_K_XL 108.7 GB at 78.3%; **UD-IQ3_XXS 120.4 GB at 81.6%**; UD-Q4_K_XL 199.7 GB at 92.2%; BF16 641.6 GB. Minimum combined RAM and VRAM: 100 GB at one bit, 128–150 GB at three bits, 162–210 GB at four bits. Unsloth states three-bit “works on 128GB devices like a Mac or NVIDIA DGX Spark.” Measured decode on four DGX Spark units at native FP8: 28–34 tok/s on short prompts, ~ 43 tok/s at 50K prompts. By contrast full GLM-5.3 at FP8 requires eight H200- or H20-class accelerators. A frontier-adjacent model has arrived on a single workstation-class machine—not on a laptop. <https://unsloth.ai/docs/models/glm-5.3-flash>
- [715] Z.ai's domestic-silicon claim for GLM-5.3-Flash and what it does and does not say. Z.ai's own announcement states the model was “previously previewed as Ox Alpha, *running* entirely on Chinese AI chips”; founder Jie Tang adds “Powered by pure Chinese chips.” Reporting puts the preview cluster at roughly 100,000 domestic accelerators serving 62 trillion tokens over the preview week. Three qualifications are load-bearing. First, the claim is about *serving*, not training: Z.ai has made no training-hardware statement for GLM-5.3 or Flash, and impicator.ai's summary is exact—“served is not trained.” Second, *no chipmaker is named*; Counterpoint's Ivan Lam speculates Ascend plus other domestic vendors, and CNBC reported it could not verify the claim independently. Third, this is *not* a first: Zhipu claimed Ascend-only training for GLM-Image in January 2026 [P, unaudited]; the June 2026 claim for the GLM-5 family is press-derived with no located primary [A/C] (see the entry on that claim). What is new is production-scale inference on a domestic cluster carrying a globally top-ranked workload. Reports of “100 trillion tokens per day” are irreconcilable with router data and should not be used. <https://www.scmp.com/tech/big-tech/article/3365433/zhipu-ai-shares-jump-viral-ox-alpha-model-revealed-glm-53-flash-chinese-chips>
- [716] Prior Zhipu domestic-training claims, for calibration: GLM-Image trained end to end on Huawei Atlas 800T A2 (Kunpeng 920 plus Ascend 910), reported January 2026, with *The Register* noting that no server counts or timings were disclosed; and the GLM-5 family reported in June 2026 as trained on roughly 100,000 Ascend 910B with MindSpore and no NVIDIA at any stage. Neither claim has been independently audited, and for the GLM-5 claim no primary Zhipu statement was located. Against them stands the reporting that DeepSeek's R1 training did not complete reliably on Ascend and reverted to NVIDIA—which is why the training-versus-serving distinction carries weight rather than being a quibble.
https://www.theregister.com/2026/01/15/zhipu_glm_image_huawei_hardware/
- [717] OpenAI, “Advancing the price-performance frontier with GPT-5.6” (30 Jul. 2026): GPT-5.6 Luna—the family's cheapest tier, corresponding to the earlier “nano” class—cut 80%, from \$1.00/\$6.00 to **\$0.20/\$1.20** per million tokens input/output; Terra cut 20% to \$2/\$12; *Sol held unchanged* at \$5/\$30. OpenAI attributes the cut to efficiency at every layer and states that GPT-5.6 Sol autonomously optimised production kernels, “reducing serving costs by 20% and improving token-generation efficiency by 15%.” On 21 August OpenAI cut Sol to \$4/\$20—explicitly “for the next 3 months,” a promotion rather than a repricing. No OpenAI executive has attributed any 2026 cut to Chinese competition on the record.
<https://openai.com/index/advancing-the-price-performance-frontier-with-gpt-5-6/>
- [718] The consequence of the 30 July cut, and the fact most easily missed: at \$0.20/\$1.20 GPT-5.6 Luna is *cheaper than most Chinese open-weight models on list price*—below GLM-5.3 (\$1.40/\$4.40), Kimi K3 (\$3/\$15), Qwen3.8-Max (\$2/\$6) and, on input, below DeepSeek V4-Pro. On a combined million-in, million-out basis: Luna \$1.40, GLM-5.3 \$5.80, Kimi K3 \$18, Claude Opus 5 \$30. Measured elasticity, 27 July to 14 August: Luna usage rose to $13.8\times$ its pre-cut average and Terra

- to $5.6\times$, against $1.1\times$ for the undiscounted Sol, with roughly 75% of the gain taken from competitors rather than cannibalised. The incumbent did not merely match the commodity price; at the bottom of its ladder it went underneath it. <https://openrouter.ai/rankings>
- [719] Caixin Global (14 Aug. 2026): DeepSeek launched V4-Pro and *raised* API prices by as much as 1,100% effective 13 August 2026, introducing peak and off-peak tiering—“a break from the prolonged price war among China’s AI developers.” No rationale was stated. Note the sequencing: the deepest Western cut (30 July) preceded this increase, so neither move can be read as a direct answer to the other, and the simple story in which cheap Chinese models drive Western cuts does not survive the price leader reversing direction mid-cascade. <https://www.caixinglobal.com/2026-08-14/deepseek-launches-v4-pro-and-raises-api-prices-by-as-much-as-1100-102473919.html>
- [720] OpenRouter public rankings, week to 29 August 2026, by tokens routed: Ox Alpha (GLM-5.3-Flash, then unattributed) 20.7T, up 216%; DeepSeek V4 Flash 12.3T; Xiaomi MiMo-V2.5 10.0T; **GPT-5.6 Luna 6.75T, up 31%**; Tencent Hy3 6.62T; **NVIDIA Nemotron 3 Ultra (free) 5.62T, up 27%**; GLM-5.3 Flash 4.62T; **Gemini 3.7 Flash 4.05T, up 156%**. The two fastest-growing Western entries are the two that were discounted, and the sixth-largest model on the board is a chip vendor’s open-weight release. Router volumes are a small and promotion-distorted slice of total inference and should be read as a directional signal. <https://openrouter.ai/rankings>
- [721] Reported unit economics, all via *The Information* and not confirmed by the companies: OpenAI first-quarter 2026 revenue \$5.7B against \$3.7B of cash burn, with adjusted gross margin around 33% against a 46% target, inference cost rising from ~\$8.4B in 2025 to a projected ~\$14.1B in 2026, and a revenue run rate of \$40B by August 2026; Anthropic at a \$65B run rate at the end of July 2026 with *API gross margin reported above 80%*. Sources conflict on the direction of OpenAI’s margin. The relevant caution for this book’s argument: no credible published estimate of loss per token exists, and no analyst asserts that these prices are below cost—“inference loses money” is a model-mix and pricing claim, not a universal one. <https://www.axios.com/2026/08/17/anthropic-revenue-run-rate-ipo-openai>
- [722] Hot Chips 2026 decks 38 (Google) and 39 (OpenAI)—the only two of forty decks claiming AI inside the silicon design flow, both from companies that design accelerators for their own frontier models. Google: an “AI Designer” taking specification to RTL on O(100) TPUs for about a week and an “AI Optimizer” doing RTL and physical-design optimisation on O(100) CPUs for about a week, yielding a 6% power and 5.8% area reduction on the TPU 8t matrix unit (6% more FLOPs in the same budget) and 5.3% area on 8i. OpenAI: measured block-level wins against an “optimized human baseline” and the qualitative claim that “major changes landed up until the day of RTL freeze.” Neither reports full-chip PPA, bug escapes or schedule saving against a baseline, and *nobody claimed AI-driven verification*, which is where most of a chip programme’s cost sits. Separately, OpenAI reports an internal agent driving a DeepSeek MLA kernel from 0.31% to 88.94% of roofline in about forty hours—bring-up software, historically the thing that decides whether a novel accelerator is usable, produced by an agent in days.
- [723] The benchmark surface of Hot Chips 2026: nine decks report inference economics on Chinese open-weight models. d-Matrix reports GLM 5.2 and Kimi K3 serving rates; OXMIQ builds its cost model on Kimi-K2 1T; NVIDIA’s Rubin deck measures agentic throughput on DeepSeek-v4-PRO and its Spectrum-X deck uses DeepSeek V3 training step time as the headline multi-tenancy result; Intel plots every published frontier model from Llama 2 70B to Kimi K2 1T to establish its capacity-versus-bandwidth thesis; SambaNova quotes MiniMax M2.7; OpenAI reports DeepSeek R1 on GB300. US silicon vendors now validate and publish their inference economics against Chinese weights, because those are the frontier-scale weights that can be obtained and published against. No Chinese chip is named anywhere in the conference, and no Chinese vendor presented.
- [724] NVIDIA and Poolside, reported by the *Wall Street Journal* from deal documents and announced in Poolside’s shareholder letter (20–24 Aug. 2026): approximately \$7B in total—**\$6B for a non-exclusive licence to Poolside’s “Model Factory”** training platform and \$1B of equity at a \$12B pre-money valuation—together with offers to **109 Poolside engineers and researchers to join NVIDIA and work on Nemotron**, NVIDIA’s open-weight family. Poolside remains independent, its founders do not join, and the licence is non-exclusive. The stated purpose is to build a leading US open-weight model against DeepSeek and Moonshot. Earlier and separate: a reported \$500M–\$1B NVIDIA investment inside a \$2B round (Oct. 2025) alongside “Project

- Horizon,” a 2 GW West Texas site with CoreWeave. Sources differ on whether the announcement date is 20 or 24 August. <https://www.forbes.com/sites/jonmarkman/2026/08/24/nvidia-pays-pools-ide-6b-to-license-its-model-factory-and-109-workers/>
- [725] NVIDIA and Hugging Face—*reported, not concluded* [A]: *The Information* reported on 26 August 2026 that NVIDIA was closing on an acquisition of Hugging Face for about \$12.9B, and as of the currency date neither company had confirmed it, with *Business Insider* reporting that talks had not produced a signed agreement. *As of the currency date neither company had confirmed it*, and *Business Insider* reported that talks had not produced a signed agreement and “could still atomize”; TechCrunch notes that NVIDIA’s silence is itself unusual for a company that moves quickly to correct reports it considers inaccurate. Context: NVIDIA participated in Hugging Face’s 2023 Series B at a \$4.5B valuation and had a standing model-hosting and DGX Cloud partnership; Hugging Face reportedly *declined* a \$500M NVIDIA investment at a \$7B valuation in 2025, on the ground that it did not want a dominant investor able to sway its decisions. Hugging Face revenue is around \$150M a year. Unlike NVIDIA’s licence-plus-acquihiere structures, an equity acquisition would require Hart–Scott–Rodino review. Treat as reported, not concluded [A]. <https://techcrunch.com/2026/08/26/nvidia-closes-in-on-hugging-face-acquisition/>
- [726] The pattern of licence-plus-acquihiere transactions that leave the target nominally independent and, on the reported structure, outside Hart–Scott–Rodino merger review: **Groq**, ~\$20B (25 Dec. 2025)—a non-exclusive inference-technology licence with roughly 90% of staff including chief executive Jonathan Ross and president Sunny Madra moving to NVIDIA, Groq continuing under a new chief executive; **Enfabrica**, ~\$900M (2026); **Poolside**, ~\$7B (Aug. 2026) [724]. Roughly \$27B in nine months. Senators Warren and Blumenthal wrote to NVIDIA on 23 March 2026: “by licensing its technology and hiring its most important employees, NVIDIA has effectively acquired Groq in all but name,” and “we are concerned that this takeover could stifle competition, further entrenching NVIDIA’s dominance in the AI chip industry.” An earlier letter of 4 February 2026 from Warren, Wyden and Blumenthal asked the FTC and DOJ to investigate reverse-acquihiere deals at Meta, Google and NVIDIA. No enforcement action had been taken by the currency date. <https://www.warren.senate.gov/newsroom/press-releases/warren-blumenthal-question-whether-nvidias-20-billion-groq-deal-is-attempt-to-avoid-antitrust-laws/>
- [727] NVIDIA’s Q2 FY2027 CFO Commentary (26 Aug. 2026) discloses a change in reporting structure that most coverage missed: “**Gaming**” has ceased to be a reported market platform, with GeForce revenue folded into a new **Edge Computing** line (\$7,198M, +27% year on year), and **Data Center is now split into Hyperscale (\$48,710M, +102%) and AI Clouds, Industrial & Enterprise (\$40,313M, +138%)**—management stating that the latter “will represent roughly half of our data center business” going forward. Networking was described on the call as a record quarter growing 18% sequentially, but *no networking dollar figure was published*. Shipments of Hopper products to China were under 1% of data-centre revenue and *no China data-centre compute revenue is assumed in the outlook*. Approximately \$26.0B was returned to shareholders in the quarter with \$99.0B remaining authorised. On the call management guided to roughly **70% revenue growth in fiscal 2028**, against a sell-side consensus near 44%—a forward statement that does not appear in the press release. https://s201.q4cdn.com/141608511/files/doc_financials/2027/Q227/Q2FY27-CFO-Commentary.pdf
- [728] Jensen Huang, Q2 FY2027 press release (26 Aug. 2026): “AI has reached its inflection point. It’s doing useful work. Its tokens are productive and profitable. Now, compute is revenue. And demand is accelerating. This time last year, one lab alone was driving the buildout; today, we have a golden age of new AI labs and startups, multiple frontier labs scaling in parallel, *a thriving open-model ecosystem* and physical AI coming online.” On the call, on OpenAI: “OpenAI’s existing and planned commitments represent approximately **12 gigawatts** of NVIDIA compute” through 2030. On custom accelerators: “many of these XPUs are inference specific chips for one cloud or one service. NVIDIA is an entire AI factory platform that spans the entire AI lifecycle that you can use in any cloud,” and, of the frontier laboratories designing their own silicon, “I have 100% confidence that our technology will continue to be extraordinary for them. . . They’re going to use it.” <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027>
- [729] OpenAI, “Jalapeño: first results” (25 Aug. 2026): vendor-run comparisons claiming **1.5–1.9× more AI work per watt at peak throughput and 1.7–3.6× lower end-to-end latency** against GB200/GB300 and MI355X, measured on GPT-OSS 120B, DeepSeek R1 670B and Kimi K2.5. Two qualifications are load-bearing. First, deployment: OpenAI states it will “begin

- deploying Jalapeño within OpenAI’s compute infrastructure by the end of the year”—*the part was not deployed at the currency date*. Second, and decisively for how the challenge should be read, the same post says: **“We will continue to widely deploy accelerators from NVIDIA and other partners for both training and inference workloads.”** SemiAnalysis notes the comparison sets HBM4 parts against older HBM3E systems and calls it “somewhat incomplete and unfair.” Neither OpenAI’s nor Broadcom’s announcement names NVIDIA; the anti-NVIDIA framing is analyst- and press-supplied. <https://openai.com/index/jalapeno-first-results/>
- [730] Analyst assessments of Jalapeño: Alexander Harrowell (Omdia), “this is the biggest competitive threat to NVIDIA,” noting that roughly half of AI infrastructure spending comes from hyperscalers running or able to run custom-silicon programmes; Adrien Sanchez (Yole), that a hyperscaler-designed chip can now match or beat Blackwell-class GPUs on inference efficiency, threatening NVIDIA’s inference margins specifically. Against them, Daniel Newman (Futurum) calls the “NVIDIA killer” framing “exhausting” and credits Broadcom rather than OpenAI for the engineering. Jensen Huang’s own response, on CNBC, 26–27 August 2026: “Lots of projects get started. Lots of projects get canceled. . . I don’t have to take anything personally. . . We have the supply chain and the technology scale to be their largest supplier.” <https://www.ibtimes.com/custom-ai-chips-are-coming-nvidia-openais-jalapeno-could-biggest-warning-yet-3806846>
- [731] OpenAI executive departures, and the direction of the talent flow. **Chris Malone**, OpenAI’s Head of Data Centers, departed in the week of 17 August 2026 (reported 25 August) after sixteen months, in what TechCrunch describes as a continuing stream of high-profile exits; OpenAI said it had recently reorganised its infrastructure organisation. Malone had never worked at NVIDIA, and the “head of data centers” title was OpenAI’s alone: he spent close to five years at Meta as a *distinguished engineer* and nearly twelve years at Google working on data-centre infrastructure and IT hardware research and development, per Data Center Dynamics’ account of his career. OpenAI’s infrastructure now reports to **Sachin Katti**, recruited in November 2025 from **Intel**, where he was Chief Technology and AI Officer, “to work on building out the compute infrastructure for AGI.” NVIDIA’s data-centre leadership is intact—Ian Buck remains VP of Hyperscale and HPC—and the senior technical flow in 2026 ran *toward* NVIDIA, not away from it: Groq’s leadership and roughly 90% of its staff in December 2025, and 109 Poolside researchers in August 2026 [726, 724]. <https://techcrunch.com/2026/08/25/openai-loses-a-top-data-center-exec-a-s-stream-of-high-profile-departures-continues/>
- [732] Tencent Hunyuan, *Hy4 preview* (28 Aug. 2026): open-sourced, **770B total / 49B active** mixture-of-experts with a context window above one million tokens, distributed on Hugging Face, Tencent Cloud TokenHub and OpenRouter at \$0.834/\$2.501 per million tokens. Tencent’s own evaluation across 163 experts and 203 engineering tasks scores it 2.99 of 4 against GLM-5.3 at 2.92 and Kimi K3 at 2.94—vendor-reported and unaudited. It is the largest open-weight release of the last week of August 2026, and three of the four open-weight releases in that window were Chinese. <https://technode.com/2026/08/28/tencent-open-sources-hy4-preview-with-770b-parameters-and-a-1m-token-context/>
- [733] DeepSeek’s first external funding round, reported by *The Information* and widely relayed (16–17 June 2026): approximately \$7.4B at a valuation above \$50B, with the state-backed National AI Industry Investment Fund (the “Big Fund”) as the sole holder of direct equity and voting rights, Tencent and CATL in a non-voting limited partnership under a five-year lock-up, and Liang Wenfeng reported to have invested roughly RMB 20B personally. A late-August report of a \$74B valuation ahead of a 2027 onshore listing conflicts with this and is not used [A]. <https://www.forbes.com/sites/anishasircar/2026/06/17/deepseek-just-raised-74-billion-heres-the-catch/>
- [734] *Opinions of the State Council on Deepening the Implementation of the “AI+” Initiative* (Guofa [2025] No. 11), signed 21 August 2025 and released 26 August 2025. The document sets out six key actions, and **“AI+ science and technology” is the first of them**—ahead of industry, consumption, welfare, governance and global cooperation. Its three science sub-clauses direct the acceleration of “AI-driven new research paradigms” and of “from 0 to 1” discovery; the construction of *scientific large models* (kexue da moxing) and AI-upgraded research platforms; and—the analytically important one—the integrated development of “AI-driven technology R&D, engineering realisation and product landing,” which chains discovery explicitly to engineering and to product. The document’s numeric targets are *diffusion* targets (over 70% penetration of new-generation intelligent terminals and agents by 2027, over 90% by 2030); **it sets no numeric target for science**. https://www.gov.cn/zhengce/content/202508/content_7037861.htm

- [735] Ministry of Science and Technology and the National Natural Science Foundation of China, launch of the “AI for Science” special deployment (the AI-for-science special deployment), 27 March 2023. It names four basic disciplines—mathematics, physics, chemistry, astronomy—and four applied priorities—drug development, gene research, agricultural breeding, and new materials; commits to domain-specific AI4S platforms, to the national next-generation AI public computing platform, to heterogeneous integration of high-performance and intelligent computing centres, and to scientific data sharing. It is a ministerial deployment of work rather than a numbered regulation, and Western sources that describe it as a “plan” overstate its form. The traceable competitive funding instrument that followed is the National Key R&D Programme’s “Disruptive Technology Innovation” special programme, whose 2025 AI-for-science call (13 May 2025) names five directions: frontier science, life and health, new materials, engineering design, and other AI-driven research-system technologies. https://www.gov.cn/xinwen/2023-03/27/content_5748495.htm
- [736] The Party’s *Recommendations* for the 15th Five-Year Plan (adopted at the Fourth Plenum, released 28 October 2025) contain, in section 14, the operative instruction “**use artificial intelligence to lead a transformation of the scientific research paradigm**”—use artificial intelligence to lead a transformation of the scientific research paradigm—and an official interpretation built around that phrase carries commentary from CAS academician E Weinan. A precision point that cuts against the enthusiastic reading: *the phrase does not appear in the Outline (gangyao) adopted by the National People’s Congress in March 2026*. In the Outline, AI appears under compute, data and diffusion (Chapter 12 on compute, algorithm and data supply, which directs the study and construction of “ultra-large-scale intelligent computing clusters” and government purchase and leasing of compute; Chapter 13 on digital empowerment), while basic research appears in Chapter 8 *without* AI. Anyone writing that the 15th Five-Year Plan makes AI-for-science a central pillar of the state plan is reading the Party recommendations rather than the plan. https://www.gov.cn/zhengce/202511/content_7049042.htm
- [737] M. Bianchini, V. Di Girolamo, J. Ravet and D. Arranz, *Artificial Intelligence in Science—Promises or Perils for Creativity?*, European Commission, 7 April 2025 (ISBN 978-92-68-25300-7): a bibliometric study of roughly three million papers across eighty fields, 2000–2022, finding that **China overtook both the United States and the European Union in AI-assisted scientific output from 2016**, and produced more than 25,000 AI-assisted papers in 2022 against roughly 15,000 for the EU and 12,000 for the US—with advantages reported on novelty and citation impact as well as volume. The finding is the European Commission’s own, which is why it is used here rather than a Chinese source. https://research-and-innovation.ec.europa.eu/knowledge-publications-tools-and-data/publications/all-publications/artificial-intelligence-science-promises-or-perils-creativity_en
- [738] R. Arcesati, MERICS (4 March 2026): roughly half of the project categories in the National Natural Science Foundation of China’s current call foreground AI for science, with individual project budgets of EUR 240,000–600,000 and topics including generative AI for fusion reactor design and biological component design for extreme-environment biomanufacturing [A]. <https://merics.org/en/comment/science-funding-call-provides-insight-chinas-ai-priorities>
- [739] Shanghai AI Laboratory, *Intern-S1* (released 26 July 2025; arXiv:2508.15763): 241B total parameters—a 235B Qwen3 mixture-of-experts language backbone with a 6B InternViT vision encoder—pre-trained on 5T tokens of which **2.5T are scientific-domain**, released under **Apache 2.0** on Hugging Face, ModelScope and GitHub. Self-reported ChemBench 83.4, MatBench 75.0, MSEarthMCQ 65.7, MicroVQA 63.9; comparative claims against Western frontier models are vendor-reported and no independent replication was located. The successor platform, Shusheng Duanyan (Intern Duanyan), announced 17 July 2026, is a five-layer scientific discovery stack with more than two hundred cross-disciplinary agents across six domains and a dry-lab/wet-lab closed loop—and *no open-source release was stated for it*. <https://huggingface.co/internlm/Intern-S1>
- [740] ByteDance AI4Science, *Protenix* (Apache 2.0; v1.0.0 released 5 February 2026 at 368M parameters, v2 on 8 April 2026 at 464M): described as the first fully open-source model to outperform AlphaFold3 across diverse benchmark sets at matched training-data cutoff, model scale and inference budget, evaluated with the accompanying PXMeter toolkit over more than six thousand complexes against AlphaFold3, Boltz-1 and Chai-1. Code, weights, training and inference pipelines and a web server are all released. The significance for this book is not the benchmark but the licence: the open-weight ratchet of Chapter 10 has reached structural biology, and it reached it from a Chinese laboratory. <https://github.com/bytedance/Protenix>

- [741] The DeepModeling community and DP Technology (Shenshi Keji, founded 2018; chief scientist Zhang Linfeng): DeePMD-kit, the deep-potential molecular dynamics package (*J. Chem. Phys.* 159, 054801, 2023); **DPA-2**, “a large atomic model as a multi-task learner” (*npj Computational Materials* 10, 293, 19 December 2024), covering 73 elements across 18 pre-training datasets spanning alloys, semiconductors, battery materials, drug molecules, ferroelectrics, solid-state electrolytes and catalysis, with density-functional labels from VASP, Gaussian and the domestic ABACUS code; the ABACUS electronic-structure package; and the **OpenLAM** initiative (launched 1 January 2024, “Conquer the Periodic Table!”). The community reports more than ten thousand members across more than twenty countries. DP Technology raised a Series C above RMB 800M (about US\$114M) in December 2025 in which three of six named investors are Beijing municipal state funds, and reports more than three million users across over a thousand universities and research organisations plus corporate clients including PetroChina, CATL and BYD [A]. Code, weights and datasets are open; on the evidence located this is among the most genuinely open scientific-AI ecosystems anywhere, Western or Chinese. <https://www.nature.com/articles/s41524-024-01493-2>
- [742] Chinese Academy of Sciences, *ScienceOne*, unveiled at the World AI Conference, 26–28 July 2025, built by twelve CAS institutes including Automation, the Computer Network Information Centre, the National Science Library, the Academy of Mathematics and Systems Science and the Institute of High Energy Physics. It is built on *Chinese open-source foundation models* and integrates AlphaFold and MatterGen as tools; it powers a digital-cell platform, particle-collider simulation and telescope coordination. Performance claims are self-reported; parameters, training compute and open-source status were not disclosed. https://english.cas.cn/newsroom/cas_media/202507/t20250728_1048587.shtml
- [743] The institutional layer of Chinese AI for science: the **AI for Science Institute, Beijing** (Beijing Kexue Zhineng Yanjiuyuan, founded September 2021 under CAS academician E Weinan), which runs the DeepModeling community and has shipped the DPA potential family, the ABACUS density-functional code, the DeepFlame combustion platform and the AIS-Square model-sharing platform; the **Shanghai AI Laboratory**, the most productive shipper of scientific models; the **Shanghai Academy of AI for Science** (2023, with Fudan), which produces the FuXi weather and climate family; **BAAI**’s 2026 “Wujie” series (Brain μ for neuroscience, OpenComplex 2.5 for biomolecules, Physis for physical world modelling); **Pengcheng Laboratory**; and **Zhejiang Lab**, whose Three-Body Computing Constellation launched its first twelve compute-carrying satellites in May 2025 against a stated target of roughly a thousand by around 2030. https://www.ncsti.gov.cn/kjdt/ztbd/szn/snhc/202501/t20250103_191555.html
- [744] Pengcheng Cloudbrain (Pengcheng Yunnao), the first fully domestic large-scale intelligent computing system, built by Pengcheng Laboratory with Huawei, Peking University and Tsinghua, awarded the **National Science and Technology Progress Award, First Prize** in the 2025 cycle (announced 8 July 2026). It is reported to support training for more than half of China’s domestic foundation models and to have guided the construction of twenty-five intelligent computing centres. The award is the clearest available signal of how the Chinese state ranks *domestic scientific computing infrastructure* against the models that run on it: the prize went to the machine. https://www.stdaily.com/web/gdxw/2026-07/08/content_544307.html
- [745] The FuXi (Fuxi) end-to-end weather and climate system (Fudan University and the Shanghai Academy of AI for Science, with the National Climate Center on the Fengshun sub-seasonal model), selected among China’s 2025 Top Ten Meteorological Science and Technology Advances (announced 21 March 2026): China’s first purely data-driven end-to-end forecasting system, ingesting satellite remote sensing directly, operationally deployed across Chinese meteorological departments and integrated into an early-warning platform serving forty-three countries. Developer-reported claims—medium-range ensembles beating ECMWF on 98.1% of variables by CRPS, sixty-day forecasts in three minutes—are not independently replicated, and the open-weight status of the current system is not stated. By contrast, Huawei’s *Pangu-Weather* has been checked by third parties in the peer-reviewed literature (*Remote Sensing* 17(2):191, 2025; *Meteorological Applications*, 2025; *Weather*, 2025, on both Pangu-Weather and FengWu), which makes weather the one Chinese AI-for-science domain with genuine external validation. <https://news.fudan.edu.cn/2026/0323/c1247a148495/page.htm>
- [746] Evidence that the Chinese AI-for-science programme is aimed at industry rather than at publication. Of the ten joint results Shanghai AI Laboratory announced at the 2026 World AI

- Conference, three are aerospace design with a state aerospace research institute and a commercial rocket firm, electroplating-materials formulation with Xiamen University, and an analogue chip-design agent with Shanghai Jiao Tong University claiming five- to tenfold sizing-optimisation efficiency and a halving of layout-planning cycles. Three of the ten carry foreign co-authors (Columbia, Cambridge, University College London), so the programme is not hermetically sealed. The “AI+” initiative’s science pillar makes the industrial intent textual: its second sub-clause chains AI-driven research to “engineering realisation and product landing” [734].
<https://www.shlab.org.cn/news/5444277>
- [747] Reuters exclusive (Liam Mo and Fanny Potkin, 26 Aug. 2026), three people familiar with the talks: Moonshot AI is negotiating *revenue-sharing* agreements with Microsoft, Amazon and Google under which the three US clouds would host Kimi K3, Moonshot seeking **up to 30% of revenue from K3-related services**, terms said to align with those offered to large customers of the open-weight model; discussions “at an early stage” with “no certainty they will result in agreements”; revenue split, data access and token auditing unresolved; Moonshot has already signed similar agreements with smaller cloud platforms. Microsoft, Google and AWS declined comment; Moonshot did not respond. It would be the first such pact between a Chinese laboratory and a major US cloud. The word in every account is *revenue*, not profit, and no source attributes to the talks any motive beyond monetization. <https://finance.yahoo.com/technology/ai/articles/exclusive-chinas-moonshot-talks-microsoft-075340033.html>
- [748] Moonshot AI, Kimi K3 licence (Hugging Face, [moonshotai/Kimi-K3/LICENSE](https://huggingface.co/moonshotai/Kimi-K3/blob/main/LICENSE)): “If the Licensee or any of its affiliates operates a Model as a Service business, and the aggregate revenue of the Licensee and its affiliates exceeds 20 million US dollars... over any consecutive 12 months, the Licensee must enter into a separate agreement with Moonshot AI before using the Software or its derivative works for any commercial purpose,” plus a display requirement above 100M monthly active users or \$20M monthly revenue. The clause is new relative to K2. It is the legal lever behind the cloud negotiations: a hyperscaler cannot serve K3 as it serves DeepSeek or Qwen without contracting with the laboratory, which is why K3, unlike K2.5, is not on Bedrock. Chinasoft International disclosed a “token revenue-sharing and joint innovation agreement” with Moonshot on 20 July 2026, percentages undisclosed.
<https://huggingface.co/moonshotai/Kimi-K3/blob/main/LICENSE>
- [749] Reuters (7 Aug. 2026, two people familiar): Alibaba will require large commercial users of the open-weight Qwen3.8-Max to share revenue, the rate under negotiation and the design explicitly modelled on Moonshot’s; the Qwen3.8-Max licence’s US\$50M model-as-a-service threshold is the instrument. <https://finance.yahoo.com/technology/ai/articles/alibaba-plans-revenue-sharing-next-131417995.html>
- [750] What the US hyperscalers already serve, from their own documentation: DeepSeek-R1 on Amazon Bedrock Marketplace from 30 January 2025 and fully managed from 10 March 2025; six open-weight models added fully managed on 10 February 2026—DeepSeek V3.2, MiniMax M2.1, GLM 4.7, GLM 4.7 Flash, Kimi K2.5, Qwen3 Coder Next; Google Vertex model-as-a-service for Kimi K2 Thinking (13 Nov. 2025), MiniMax M2, DeepSeek-V3.2, GLM 4.7 and GLM 5 (10 Feb. 2026); DeepSeek in Microsoft Foundry’s “sold by Azure” category. AWS’s third-party model terms state that DeepSeek and Qwen serverless models “are sold by AWS” and that customer data is not shared with model providers. No cloud discloses any payment to a Chinese laboratory for serving permissively licensed weights, and Bedrock lists DeepSeek V3.2 at roughly two to five times the laboratory’s own price. Kimi K3 is reachable in Foundry only through Fireworks AI as a partner model, at \$3.30/\$16.50. <https://aws.amazon.com/legal/bedrock/third-party-models/>
- [751] G. Webster and Y. Xu, “Kimi’s tightrope: openness and revenue,” *Transpacifica* (27 Jul. 2026)—“with revenue comes regulability”: downloading open weights left it ambiguous whether any transaction with a Chinese laboratory occurred, but “if a US company needs a contract with Moonshot to provide the inference tokens... the picture looks different”; under a supply-chain rule “it’s harder to argue there’s no transacting going on,” and if Moonshot were sanctioned “the contract to provide inference is clearly an instance of doing business.” The only serious analysis located runs against the reading that hosting on US clouds makes a US restriction harder; no source was found that supports it. Separately: Axios (20 Jul. 2026) reported a draft executive order that would require US companies hosting Chinese models to guarantee security and accept liability, not issued at the currency date.
<https://herecomes.transpacifica.net/p/kimis-tightrope-openness-and-revenue>

- [752] FY2026 National Defense Authorization Act §1532 (enacted 18 Dec. 2025): the Department of Defense must exclude “covered AI”—defined by reference to DeepSeek and High-Flyer, entities on the Consolidated Screening List or the 1260H list, and companies domiciled in covered nations with unmitigated foreign control—and contractors may not use it in DoD work, with case-by-case waivers. The only enacted federal restriction on Chinese models at the currency date; it does not distinguish self-hosted from cloud-served. The *No Adversarial AI Act* (S.2177/H.R.4142, introduced 25 Jun. 2025) remained in committee.
<https://www.congress.gov/bill/119th-congress/senate-bill/2177/text>
- [753] Reuters (7 Jul. 2026, three sources): China’s Ministry of Commerce met Alibaba, ByteDance and Z.ai over the preceding month to discuss options including barring public release of the most advanced models, domestic-only access, curbs on overseas access to unreleased systems, criminalizing model leakage and screening foreign capital; a panel of legal scholars proposed a tiered filing, security-review and domestic-only scheme; “officials have decided nothing yet,” any curbs might apply only to future models, and no timeline was given. Bloomberg (26 May 2026) reported passport-surrender and travel restrictions extended from DeepSeek to Moonshot, StepFun, ByteDance and others. Note that Zhipu’s two-week hold on the GLM-5.3 weights in August was, on every account located, the company’s own safety review and not a government one.
<https://time.com/article/2026/07/07/china-ai-models-alibaba-bytedance/>
- [754] NSA, FBI and CISA, joint cybersecurity advisory AA26-251a (9 Sep. 2026), naming DeepSeek, Moonshot, Alibaba, MiniMax, StepFun and Z.AI as conducting distillation campaigns against US frontier models, with traffic routed through cloud providers and third-party aggregators. The advisory concerns distillation, not hosting.
<https://www.cisa.gov/news-events/cybersecurity-advisories/aa26-251a>
- [755] Bloomberg (4 Sep. 2026, anonymous sources), relayed by TechNode and Mobile World Live: DeepSeek plans to order at least **160,000 Huawei Ascend 950DT** accelerators for a gigawatt-scale inference site at Ulanqab, Inner Mongolia, with delivery paced by Huawei’s and the domestic HBM supply chain’s capacity and expected to take up to a year; training is reported to remain on NVIDIA hardware. The implied cost of roughly \$2.6B is a derived estimate. The order matters to this book less for what DeepSeek gets than for what it shows about supply: a single domestic customer absorbs a year of Ascend output.
<https://technode.com/2026/09/07/deepseek-huawei-ascend-950dt-chips-inner-mongolia/>
- [756] *The Information* (28 Aug. 2026), relayed by Tom’s Hardware: the Commerce Department is drafting a rule to block Chinese firms from renting NVIDIA compute in third countries, reportedly triggered by Moonshot training Kimi K3 on servers in Thailand; the draft may reach industry in September and its legal authority is contested. Not published at the currency date [A].
<https://www.tomshardware.com/tech-industry/policy/new-us-export-controls-reportedly-target-chinese-access-to-remote-ai-servers-trump-admins-cut-down-ai-diffusion-rule-could-be-share-d-with-industry-as-soon-as-september>
- [757] Reuters and HK01 (3 Sep. 2026): Moonshot AI has made a confidential A1 filing for a Hong Kong listing, reportedly to raise about \$3B at a pre-money valuation near \$50B; the company declined comment [A]. Reported annualized revenue rose from \$100M in March to \$300M in June 2026 (*The Information*, unverified), with no geographic split disclosed. <https://technode.com/2026/09/03/moonshot-ai-reportedly-submits-confidential-hong-kong-ipo-filing/>
- [758] Artificial Analysis, Intelligence Index v4.2 (4 Sep. 2026) and **v4.3 (7 Sep. 2026)**: two re-basings in one week (Terminal-Bench v4.0 replacing v2.1, AutomationBench-AA replacing τ^3 -Banking, GPQA Diamond dropped, private-set weight 45%). On v4.3: Claude Fable 5.1 (max) 53, GPT-6 Astra (max) 53, Claude Opus 5 (max) 51, Muse Spark 1.3 48, GPT-5.6 Sol 47, GLM-5.3 (max) 45, **GLM-5.3-Flash 42**, Qwen3.8-2.4T-A95B 40, GPT-5.6 Luna 38, DeepSeek V4 Pro 36, DeepSeek V4 Flash 35. The launch-week score of 57 for GLM-5.3-Flash (and 60 for GLM-5.3, 63 for Opus 5) was on v4.1.1. Any index number in this book carries its version; cross-version comparisons are meaningless. <https://artificialanalysis.ai/articles/artificial-analysis-intelligence-index-v4-3>
- [759] Artificial Analysis independent sub-scores for GLM-5.3-Flash on v4.3: Terminal-Bench v4.0 **33%** against 12% for GPT-5.6 Luna and 12% for DeepSeek V4 Flash (GLM-5.3: 42%); AutomationBench-AA 60% (Luna 50%); but GDP.pdf 15% (Luna 24%), CritPt 15% (Luna 21%), output 61–71 tokens/s against Luna’s 120, and 34.5s to the first answer token; cost per task \$0.25 against Luna \$0.18 and GLM-5.3 \$2.01. Vendor-reported figures on the model card

- (Terminal-Bench 2.1 84.3, DeepSWE v1.1 63.4) are not independently replicated, and no SWE-bench Verified, SWE-bench Pro, LiveCodeBench or Aider result exists from any source. The model is, on independent evidence, the cheapest point on the price–agentic-coding frontier and not a general replacement for the flagship tier.
<https://artificialanalysis.ai/models/comparisons/glm-5-3-flash-vs-gpt-5-6-luna>
- [760] OpenRouter rankings, data through 8 September 2026, by tokens: Tencent Hy4 preview 18.7T; GPT-5.6 Luna 14.5T; DeepSeek V4 Flash 12.4T; **GLM-5.3-Flash 12.2T**; DeepSeek V4 Flash 0423 5.0T; MiniMax M3 (free) 4.3T; Hy3 3.7T; Nemotron 3 Ultra (free) 3.7T; MiMo-V2.5 3.4T; GLM-5.3 3.4T. Of the ten, the active-parameter counts where disclosed are 49B, 13B, 18B and below—the volume tier is a sparse-mixture tier. Tom’s Hardware (4 Sep. 2026), citing an OpenRouter engineer: platform token volume up more than $25\times$ in a year and doubled in the last month; Claude Fable 5 in “the low billions” of monthly output tokens against roughly 12T each for Luna and GLM-5.3-Flash. Router volumes remain a small, promotion-distorted slice of global inference. <https://openrouter.ai/rankings>
- [761] Vercel AI Gateway: open-weight models at 36% of July 2026 tokens on 8.6% of spend, a single-day reading of about 62% of tokens on 22 August 2026, self-hosting break-even “past 50 million tokens per day,” 81% of tokens on models absent six months earlier; the top four frontier laboratories still take 95% of spend, Anthropic 61–65% of spend at roughly $4.4\times$ the average token price. The barbell—volume in the open sparse tier, spend in the closed frontier—is the single most important fact about the model market at the currency date. <https://vercel.com/i/open-weight-models>
- [762] P. Belcak et al. (NVIDIA Research), “Small Language Models are the Future of Agentic AI,” arXiv:2506.02153 (2 Jun. 2025): the argument that most agentic sub-tasks are repetitive, narrow and non-conversational, and are therefore better served by small models with a large model held in reserve—together with NVIDIA’s own subsequent agentic flagship, Nemotron 3 Super at 120B, which is not small. Against it, the “Hierarchy of Agentic Capabilities” study (arXiv:2601.09032, Jan. 2026): on 150 realistic workplace tasks “weaker models struggle with fundamental tool use and planning,” and even the best models fail about 40%. Morgan Stanley (3 Aug. 2026) on the rent consequence: open weights “may really kill... the idea that the foundation model layer can sustain high rents”—not compute demand. <https://arxiv.org/abs/2506.02153>
- [763] On the widely repeated claim that the GLM-5 family was trained on roughly 100,000 Huawei Ascend 910B with no NVIDIA hardware at any stage (June 2026): no primary Zhipu document stating it was located. It circulates through TechTimes (21 Jun.), Tom’s Hardware (30 Jun.) and Decrypt (18 Jun.) without a cited source; the GLM-5 repository carries no hardware statement; and Huawei’s own Ascend announcement of 12 February 2026 describes “0-day” *inference* support and post-release training guides, not a training claim. The one Zhipu domestic-training claim with a primary is GLM-Image (January 2026), with no accelerator count. This book therefore grades the GLM-5 training claim as press-derived [A/C] and does not stack later claims on it.
<https://www.hiascend.com/activities/dynamic-news/648>
- [764] World Artificial Intelligence Cooperation Organization, founding agreement signed in Shanghai on 16 July 2026 (Wang Yi for China; the UN Secretary-General attended): an intergovernmental organization headquartered in Shanghai with a Council of members and a Secretariat; civilian domain only; funded by member contributions scaled to Universal Postal Union classes; enters into force on three ratifications; one-year withdrawal notice. Twenty-nine signatories at founding, reported to include Russia, Brazil, South Africa, Indonesia, Kazakhstan, Pakistan, Malaysia, Serbia, Kenya and most of ASEAN; no official roster was retrievable, and India, the Gulf states other than Oman, Turkey and Egypt are absent. Xi Jinping’s WAIC keynote of 17 July pledged 5,000 training places for developing-country personnel over five years, regional “application cooperation centres” and a meteorological warning system for thirty countries—and *no* compute, financing, datacentre or open-weights commitment. An academic mapping (arXiv:2606.23860) scores WAICO lowest on formalization of fifteen comparable bodies: “a name, a host city, and a statement of purpose.” The earlier *Global AI Governance Action Plan* (26 Jul. 2025) and *AI+ International Cooperation Initiative* (23 Sep. 2025) likewise contain no infrastructure or financing clause.
https://english.www.gov.cn/news/202607/17/content_WS6a59a226c6d00ca5f9a0c432.html
- [765] Bloomberg (26 Aug. 2026, anonymous sources), relayed by The Next Web, Al Jazeera (4 Sep.) and Al-Monitor: Huawei has bid to build Egyptian government AI datacentres—about 1,408 Ascend 950-series accelerators for a training cloud and roughly 600 for two inference clusters, a

- twelve-month build, reportedly with iFlytek—which would be the first confirmed Ascend export anywhere; the US State Department contacted NVIDIA, AMD and Microsoft to assemble a counter-offer and warned that Ascend use may breach US rules, Egypt having needed a US licence for advanced AI chips since 2023. No price, no financing terms, nothing signed; Xi Jinping’s state visit to Cairo (31 Aug.–2 Sep.) produced no announced AI agreement. <https://www.aljazeera.com/news/2026/9/4/as-the-china-us-ai-race-reaches-cairo-egypt-faces-a-strategic-decision>
- [766] Bloomberg via *The Edge* (7 Sep. 2026) and SCMP (8 Sep.): Telekom Malaysia is evaluating Huawei Ascend 910C as the backbone of a RM2B (about US\$494M) sovereign-AI project to keep government, military and intelligence data in-country; quantity unclear; US officials warned in 2025 that 910C use could violate US export rules; Malaysia accepted export-control commitments in its 2025 reciprocal trade agreement with the United States; unsigned. In May 2025 a deputy minister announced and within hours retracted a 3,000-chip Ascend deployment. <https://www.scmp.com/week-asia/politics/article/3366792/will-malaysias-potential-huawei-ai-deal-risk-clash-over-us-trade-pact>
- [767] Brazil, 21 August 2026: an R\$2.3B (about US\$444M) AI supercomputing programme of which R\$1.3B (about US\$251M) goes to Huawei and iFlytek for a Rio de Janeiro machine on Ascend and a Portuguese-language model, from July 2027, with a further R\$1B tender for a Rio Grande do Norte machine expected to go to NVIDIA; financed by Brazil’s own FNDCT; 3,000 professionals to be trained over five years. “The strategy is not to depend on a single company, technology or country.” The only contracted sovereign-AI deployment on Chinese silicon outside China located at the currency date—and it is hedged, half-NVIDIA and Brazil-financed. <https://www.aljazeera.com/economy/2026/8/21/brazil-launches-ai-supercomputer-push-while-balancing-us-and-chinese-tech>
- [768] Kazakhstan–China package of 16 July 2026 (Tokayev in Shanghai, “\$15B, seventy-plus agreements”): a strategic partnership between the Ministry of AI and Digital Development and Huawei; a Samruk-Kazyna agreement to acquire Huawei equipment; and a Kazakhtelecom–Hengtong agreement on “basic principles” for a 1 GW “Data Center Valley” at Ekibastuz, first facility targeted for June 2027; project-level values undisclosed; a mix of binding agreements and memoranda. Pakistan’s Punjab government agreed six datacentres and a thousand “AI master trainers” with Huawei on 28 August 2026, with no money, chips or contract described. <https://www.intellinews.com/kazakhstan-china-sign-15bn-worth-of-deals-as-tokayev-visits-shanghai-455262/>
- [769] Executive Order 14320, “Promoting the Export of the American AI Technology Stack” (23 Jul. 2025), and the Commerce Department’s American AI Exports Program: full-stack packages (hardware, data, models, security, applications) with federal financing referrals—direct loans, guarantees, equity, political-risk insurance—and an Export–Import Bank “ExportAI” framework under consideration from May 2026; 78 consortium proposals received by 30 June 2026 and *no packages designated* at the currency date. Neither the Chinese nor the American export offer had been delivered anywhere by 9 September 2026; the contest in Cairo is between two offers, neither of which yet exists in the field. <https://www.trade.gov/press-release/commerce-statement-application-conclusion-next-steps-american-ai-exports-program>
- [770] On the supply side of any Ascend export: Beijing’s reported requirement that state-funded compute run at least 80% domestic silicon under a roughly RMB 2T national compute programme (Bloomberg, 9 Jun. 2026, relayed) [C]; a SemiAnalysis-derived estimate that 2026 domestic HBM supports only some 250,000–300,000 Ascend 910C-equivalents [A/C]; DeepSeek’s 160,000-unit order alone reported to take up to a year [755]; and BIS guidance of 13 May 2025 that using Ascend “anywhere in the world” risks violating US controls, extended on 1 June 2026 to state that licence requirements follow PRC-headquartered parents abroad. Huawei was reported in July 2025 to be pitching low thousands of 910B to the UAE, Saudi Arabia and Thailand with no finalized deals; the Gulf went to US suppliers. No PRC policy authorizing or restricting Ascend export was located. <https://www.aljazeera.com/economy/2026/6/1/us-says-ban-on-ai-chip-shipments-applies-to-chinese-firms-outside-china>
- [771] Hugging Face, “The state of open models, summer 2026”: 151,448 Qwen derivative models, about 2.6 times Meta’s; roughly 39.6 million monthly GGUF downloads for Qwen against 7.5 million for Llama. The diffusion mechanism that actually operates in the non-aligned world is this one—weights that cost the recipient nothing and lock the recipient into nothing—and it is a

- different instrument from the hardware-plus-service-plus-debt bundle that built Huawei's telecom position. <https://huggingface.co/blog/state-of-open-models-summer-2026>
- [772] Huawei Investment & Holding, *Annual Report 2025*: revenue RMB 880.9B (about US\$126B), net profit RMB 68.0B, ICT infrastructure RMB 375.0B (+2.6%), R&D RMB 192.3B (21.8% of revenue); "more than 170 countries and regions" of operation. Earlier: 2006 US\$8.5B; 2009 US\$21.8B; 2020 peak RMB 891B before the Honor divestiture and sanctions; 2021 RMB 636B; 2024 RMB 862.1B. <https://www.huawei.com/en/annual-report/2025>
- [773] Dell'Oro Group telecom-equipment share series as reported: Huawei 27.7% (2018), 28.1% (1H19), 29% (2Q19, Nokia 15.7%, Ericsson 13.1%), **31% (1H20)**; and **31% again in 1H25** (Nokia 13%, Ericsson 12%, ZTE 10%)—*21% outside China and 40% outside North America*. Dell'Oro (17 Feb. 2026): the RAN market stabilized in 2025 after losing about a fifth of its value in 2022–24, with Huawei and Nokia gaining and vendor concentration at a ten-year high. Omdia (Aug. 2025): Huawei the leading RAN vendor in three of five regions (Asia-Pacific, Middle East and Africa, Latin America), third in Europe, absent from North America; 2024 sales only 4% below 2020. <https://www.delloro.com/news/ran-market-stabilized-in-2025/>
- [774] N. Ahrens, *China's Competitiveness: Myth, Reality, and Lessons for the United States and Japan—Case Study: Huawei*, CSIS (Feb. 2013): China Development Bank credit facilities for Huawei's customers of US\$10B (2004) raised to **US\$30B (December 2009)**, with Export–Import Bank support of about US\$600M, and roughly US\$10B drawn by 2013; pricing about 30% below Western competitors in developing markets and about 12% below Russian rivals; the rural-first entry (more than two hundred staff across Heilongjiang counties against three or four for Ericsson) and the sequence Russia, Thailand, Brazil, South Africa, Ethiopia (2003), Nigeria (US\$200M CDMA, 2005) before Europe; parity with Ericsson by 2011. https://csis-website-prod.s3.amazonaws.com/s3fs-public/legacy_files/files/publication/130215_competitiveness_Huawei_casestudy_Web.pdf
- [775] C.-W. Yap, "State support helped fuel Huawei's global rise," *Wall Street Journal* (25 Dec. 2019): up to US\$75B of state support—about US\$46B in loans, credit lines and other financing from state lenders, about US\$25B in tax savings 2008–18, plus grants and discounted land—and, per analysts and customers, prices undercutting rivals by some 30%. Huawei disputed the characterization. Counter-data: a 2017 European study found Huawei's prices only slightly lower than rivals' in Europe, and a 2019 study found them about 25% *higher* in China. The 2011 EU inquiry put the US\$30B CDB line at roughly ten times Huawei's operating cash flow. <https://www.lightreading.com/business-management/no-one-wants-to-talk-about-huawei-s-state-subsidies>
- [776] LexisNexis IPlytics, *5G SEP Report* (Jan. 2025): Huawei first in declared 5G patent families and first in 3GPP technical contributions over 2015–24 (Ericsson second, Nokia third, Qualcomm fourth, ZTE fifth); Chinese-headquartered firms above 40% of declared 5G families. The standards position is the endgame move of the telecom campaign, executed inside a Western-founded body rather than around it. <https://www.lexisnexisip.com/wp-content/uploads/2025/01/2025-LexisNexis-5G-SEP-Report.pdf>
- [777] The pushback and what it removed: Australia (Aug. 2018); US Entity List (16 May 2019); United Kingdom (14 Jul. 2020: no new equipment after 2020, removal by end-2027, up to £2B and up to three years' delay); Sweden (Oct. 2020); EU 5G toolbox (Jan. 2020) and the Commission's proposed mandatory phase-out of high-risk suppliers (20 Jan. 2026). US "rip and replace": about US\$5B appropriated, 53 of 126 projects (42%) complete by June 2026 against a missed May 2026 deadline. Strand Consult: Chinese vendors' share of European 5G sites about 32% in 2Q22 and 4Q24—flat—with Germany at 59%; the United Kingdom fell from 41% to about 20%; of 23 states that legislated, seven imposed restrictions. The bans bit in the Five Eyes, the Nordics and Japan, and the May 2020 US foundry rule—a semiconductor chokepoint, not a network-security finding—is what triggered the British reversal. <https://www.lightreading.com/5g/huawei-has-hardly-been-weakened-in-european-5g-data-shows>
- [778] E. Dou, *House of Huawei: The Secret History of China's Most Powerful Company* (Portfolio, Jan. 2025), and K. Chan's review, "Huawei the Hydra," *Phenomenal World* (1 May 2025), for the industrial-policy reading; and, on Nortel, S. Cooper, *Global News* (Aug. 2020): the 2004 breach traced to Shanghai (1,488 documents including optical and 3G/4G research), Huawei's roughly 40% underbid on the 2008 Telus/BCE contract, and the caution that a 2014 University of Ottawa

study found no evidence that the intrusion caused Nortel's 2009 bankruptcy.
<https://phenomenalworld.org/reviews/huawei-the-hydra/>